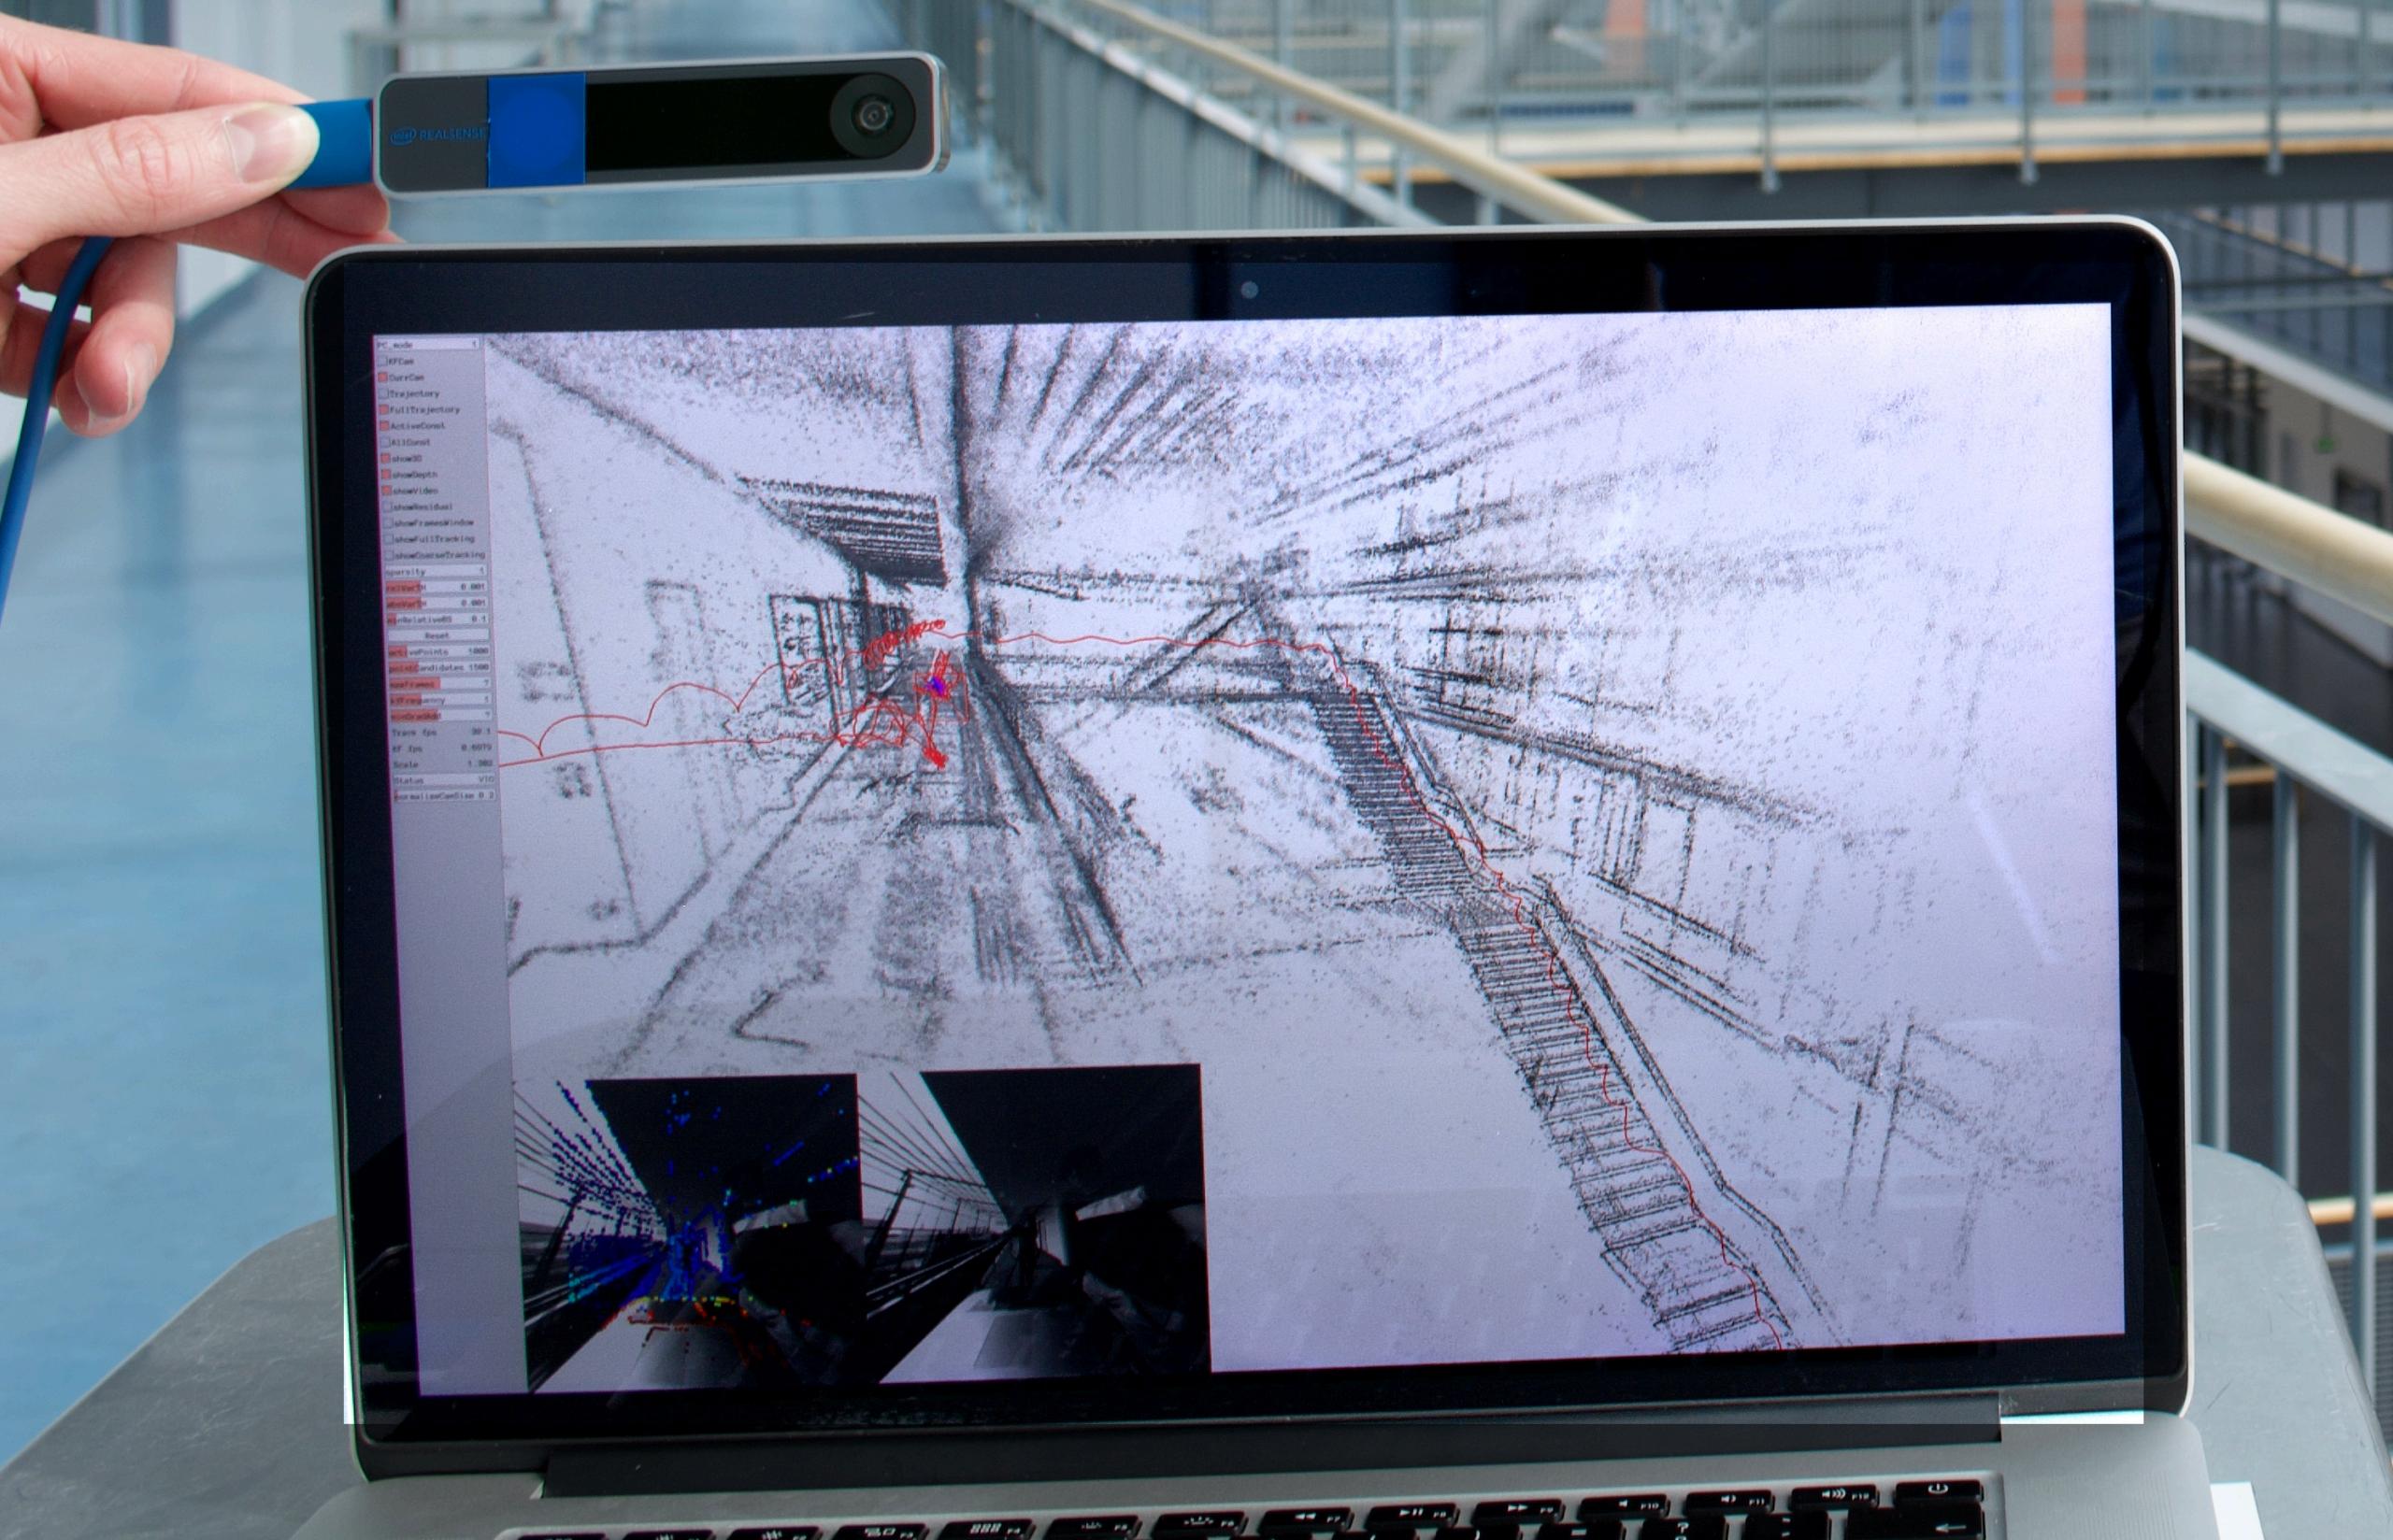


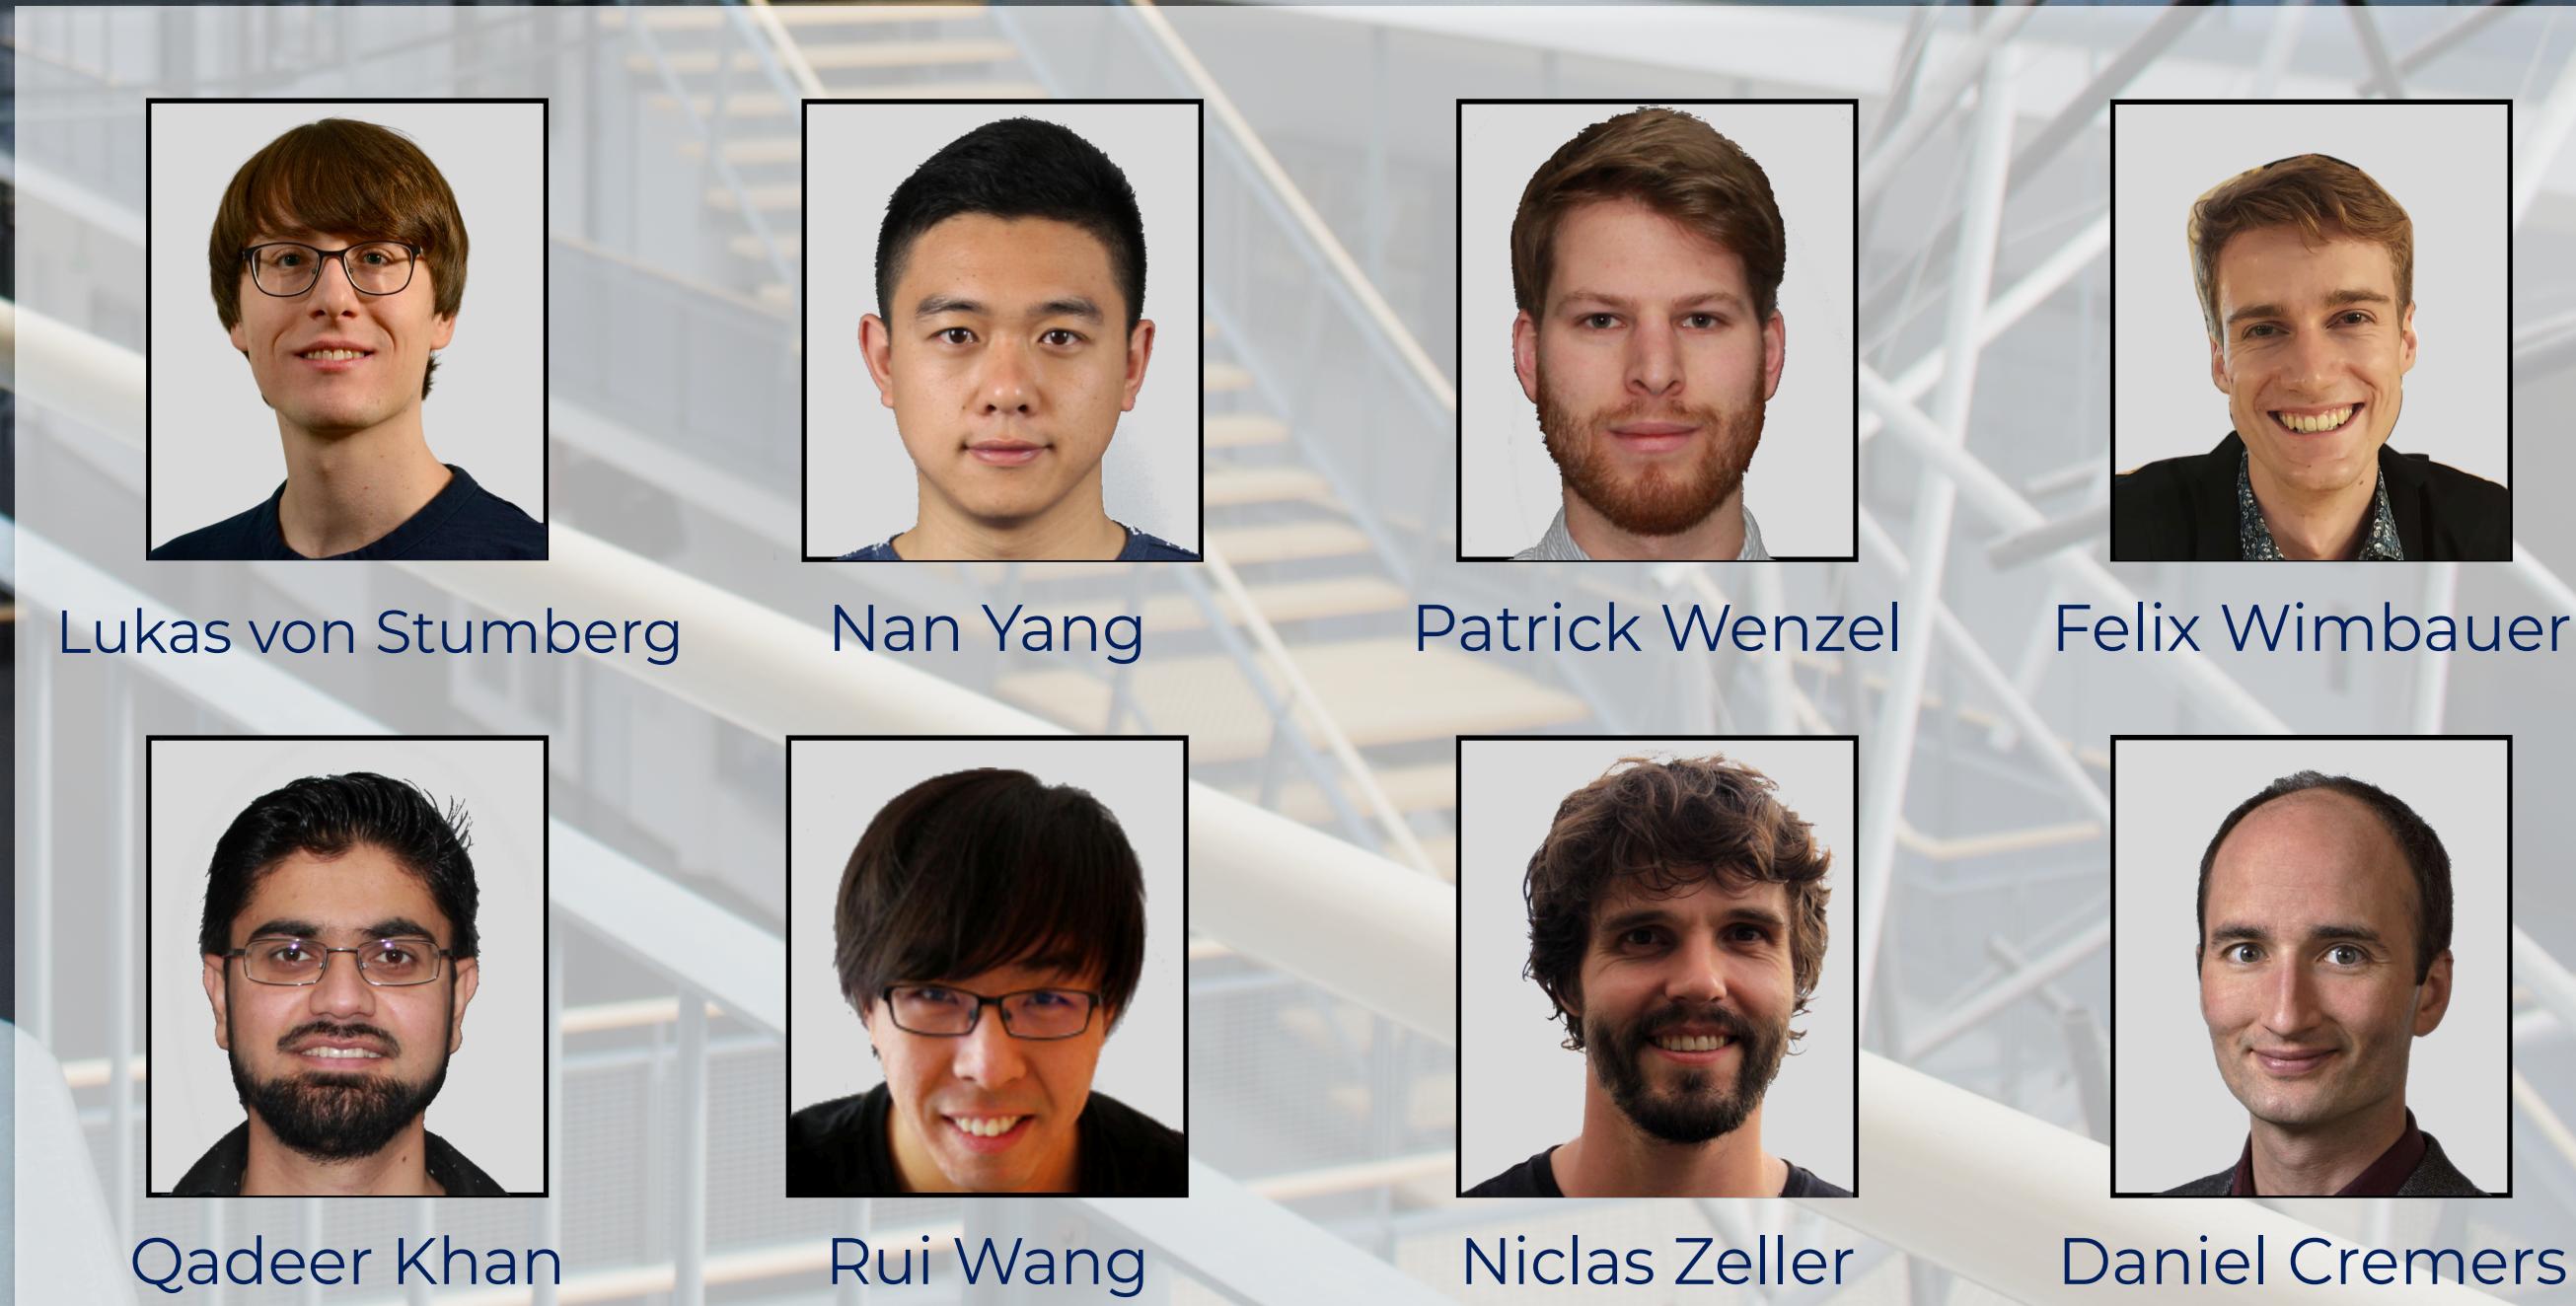
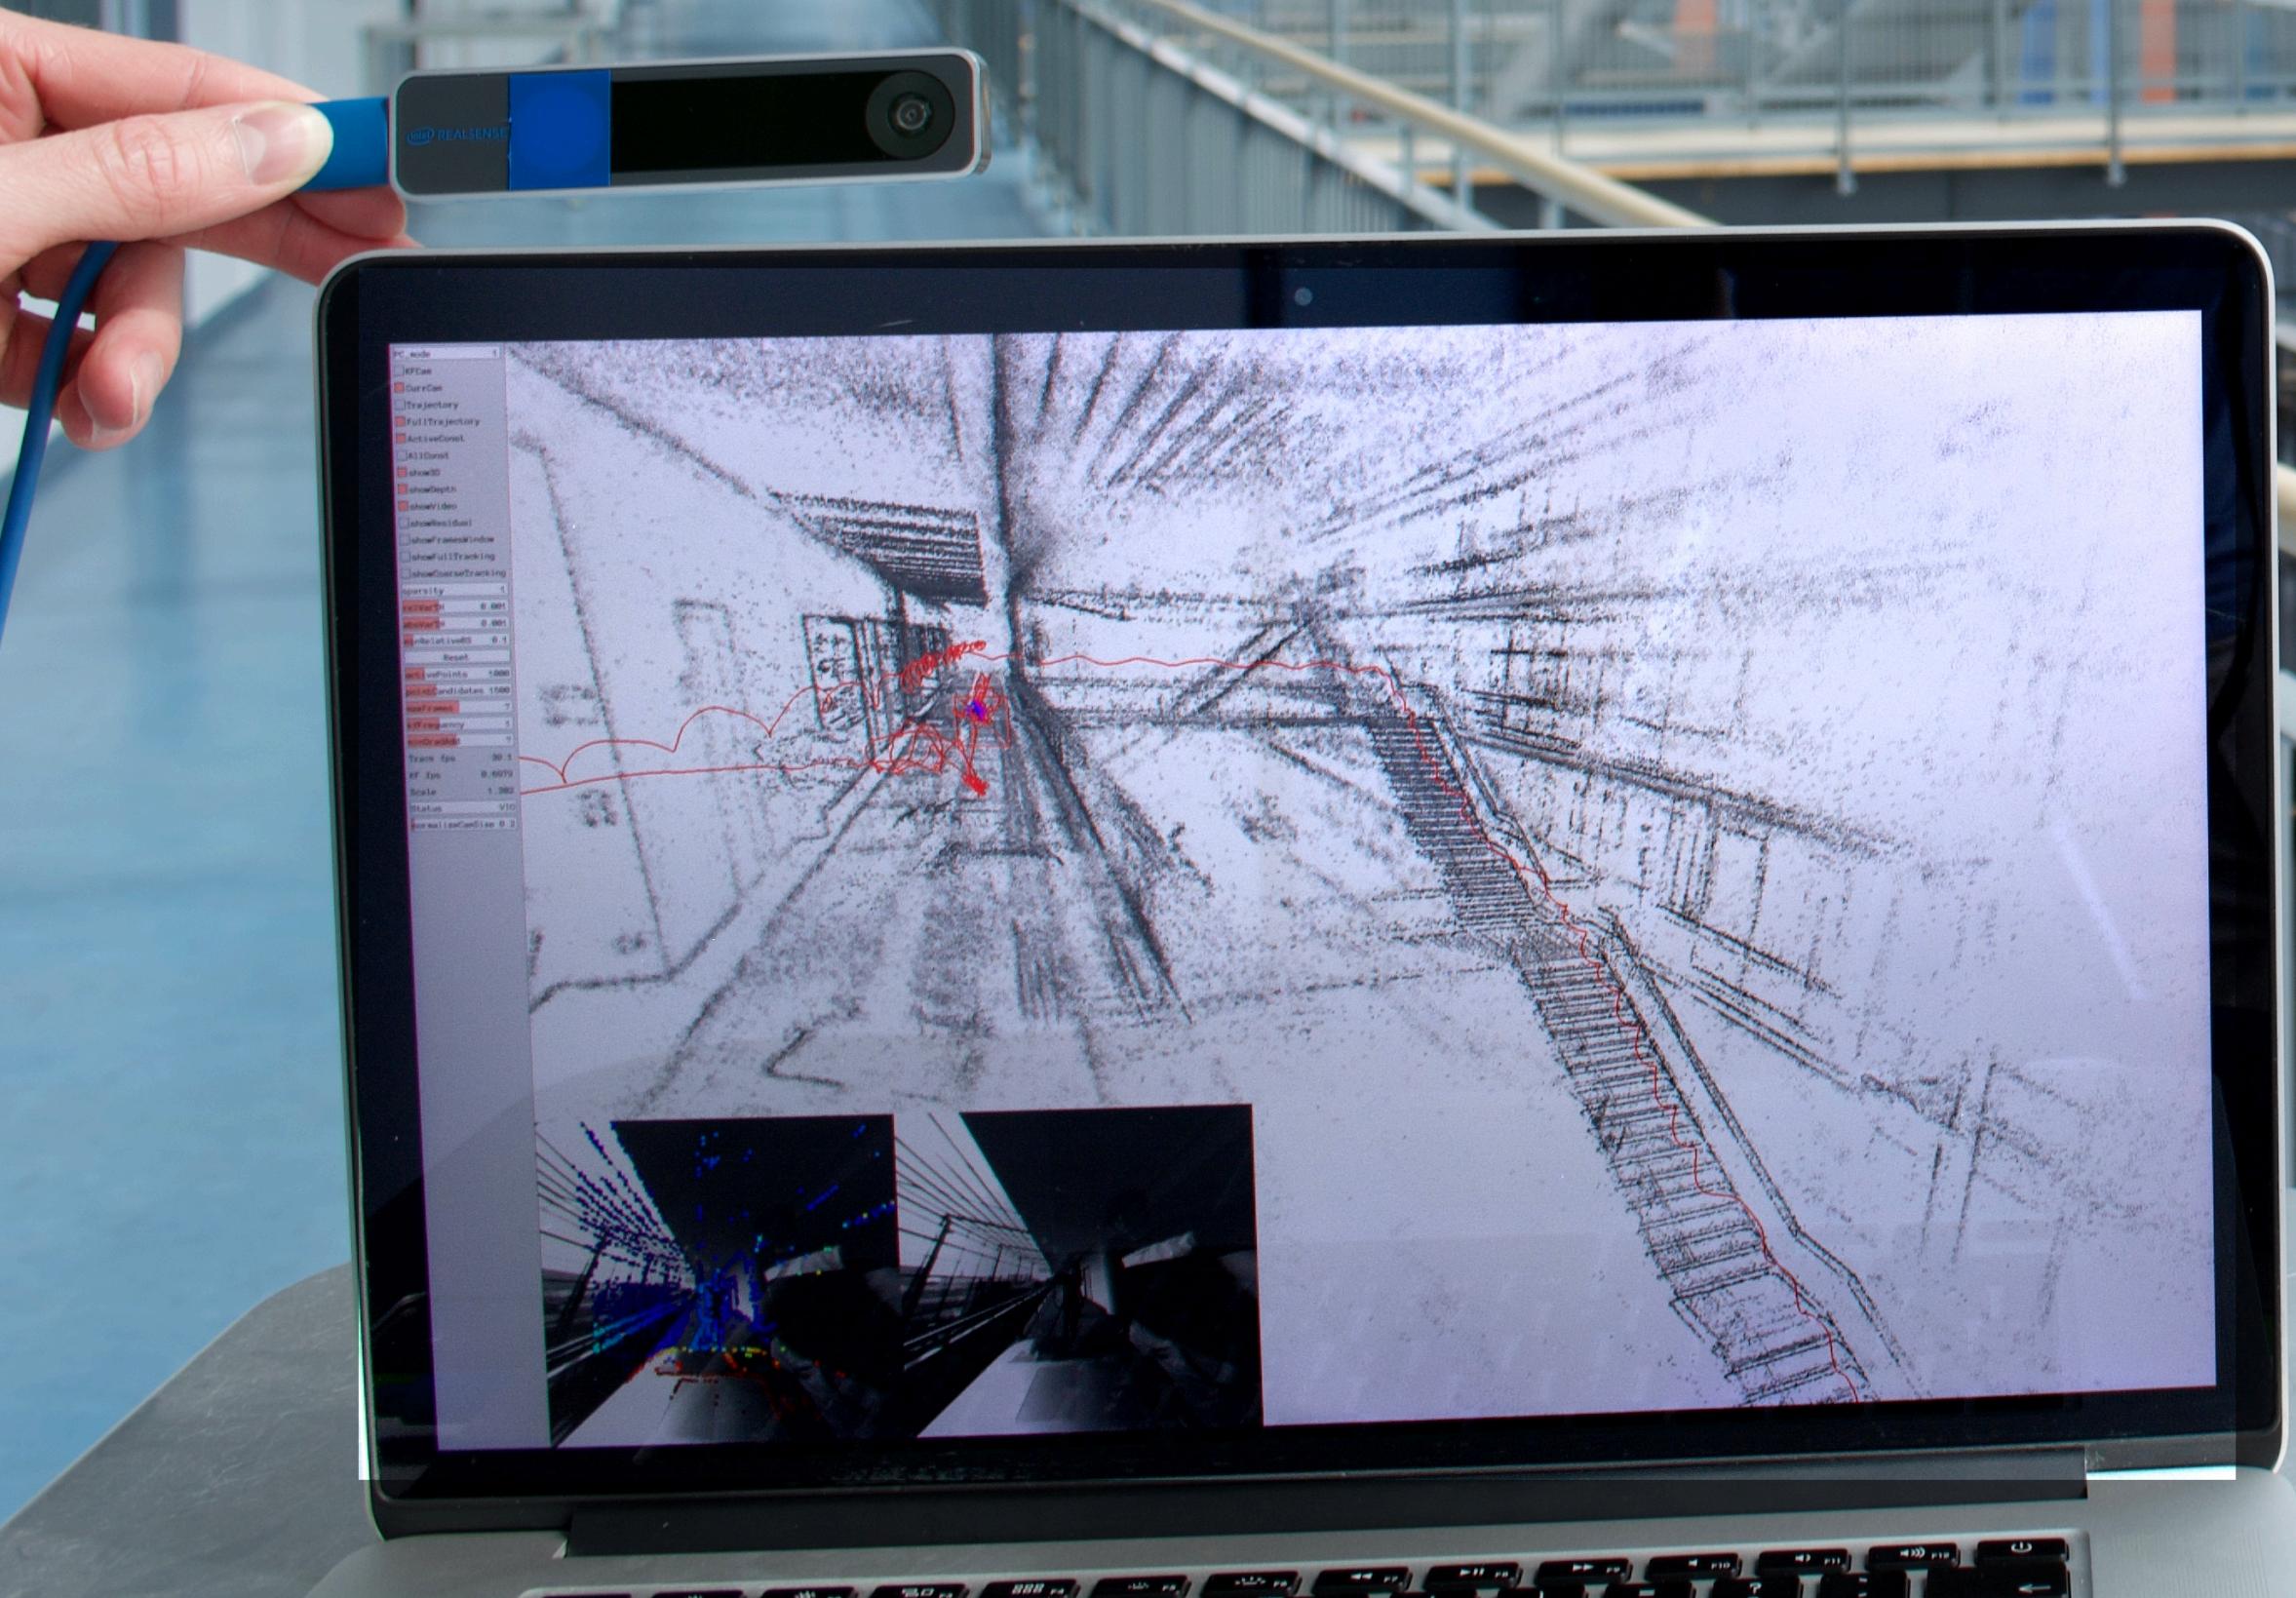
Visual SLAM From Optimization to Learning

Lukas von Stumberg

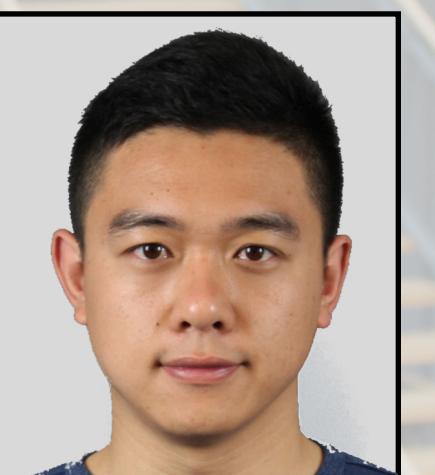


Visual SLAM From Optimization to Learning

Lukas von Stumberg



Lukas von Stumberg



Nan Yang



Patrick Wenzel



Felix Wimbauer



Qadeer Khan



Rui Wang



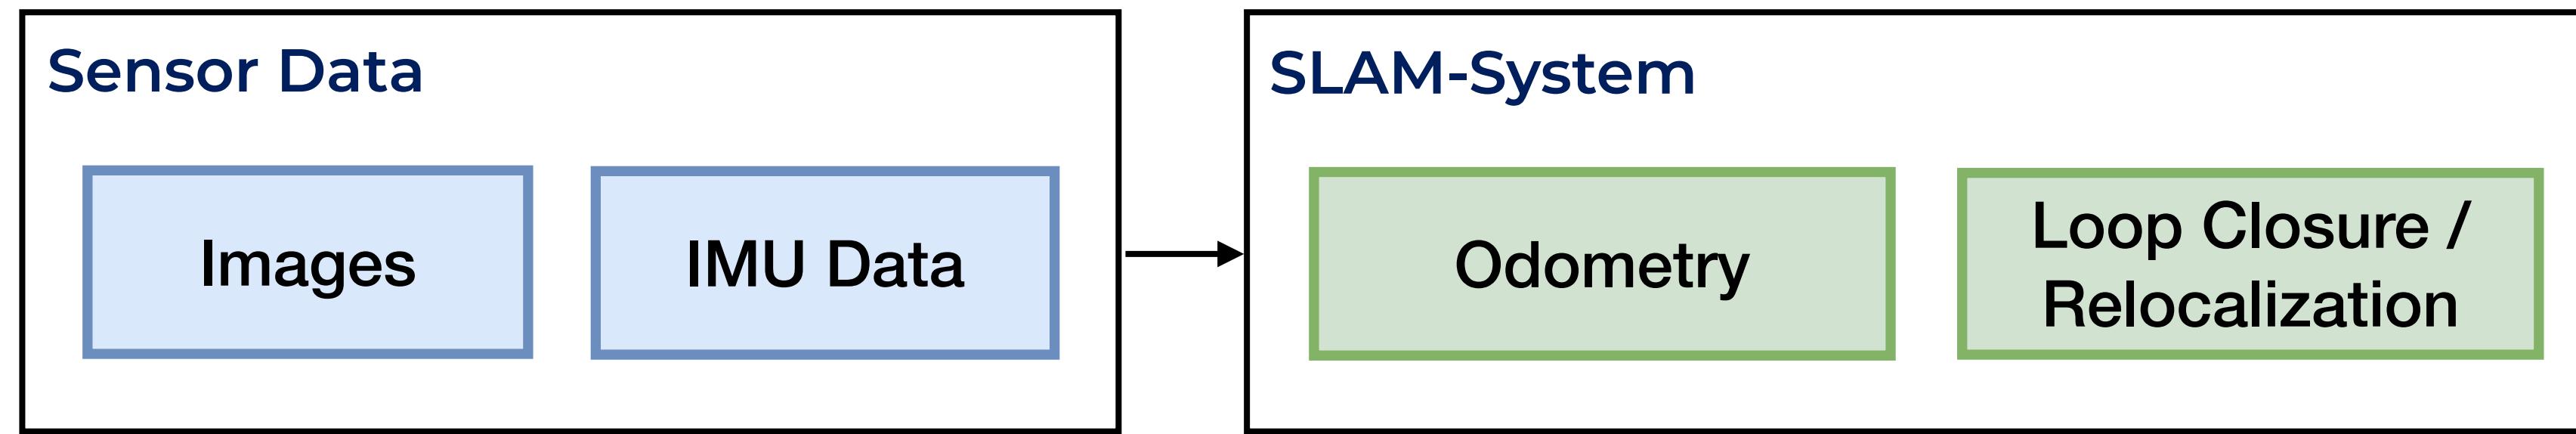
Niclas Zeller



Daniel Cremers

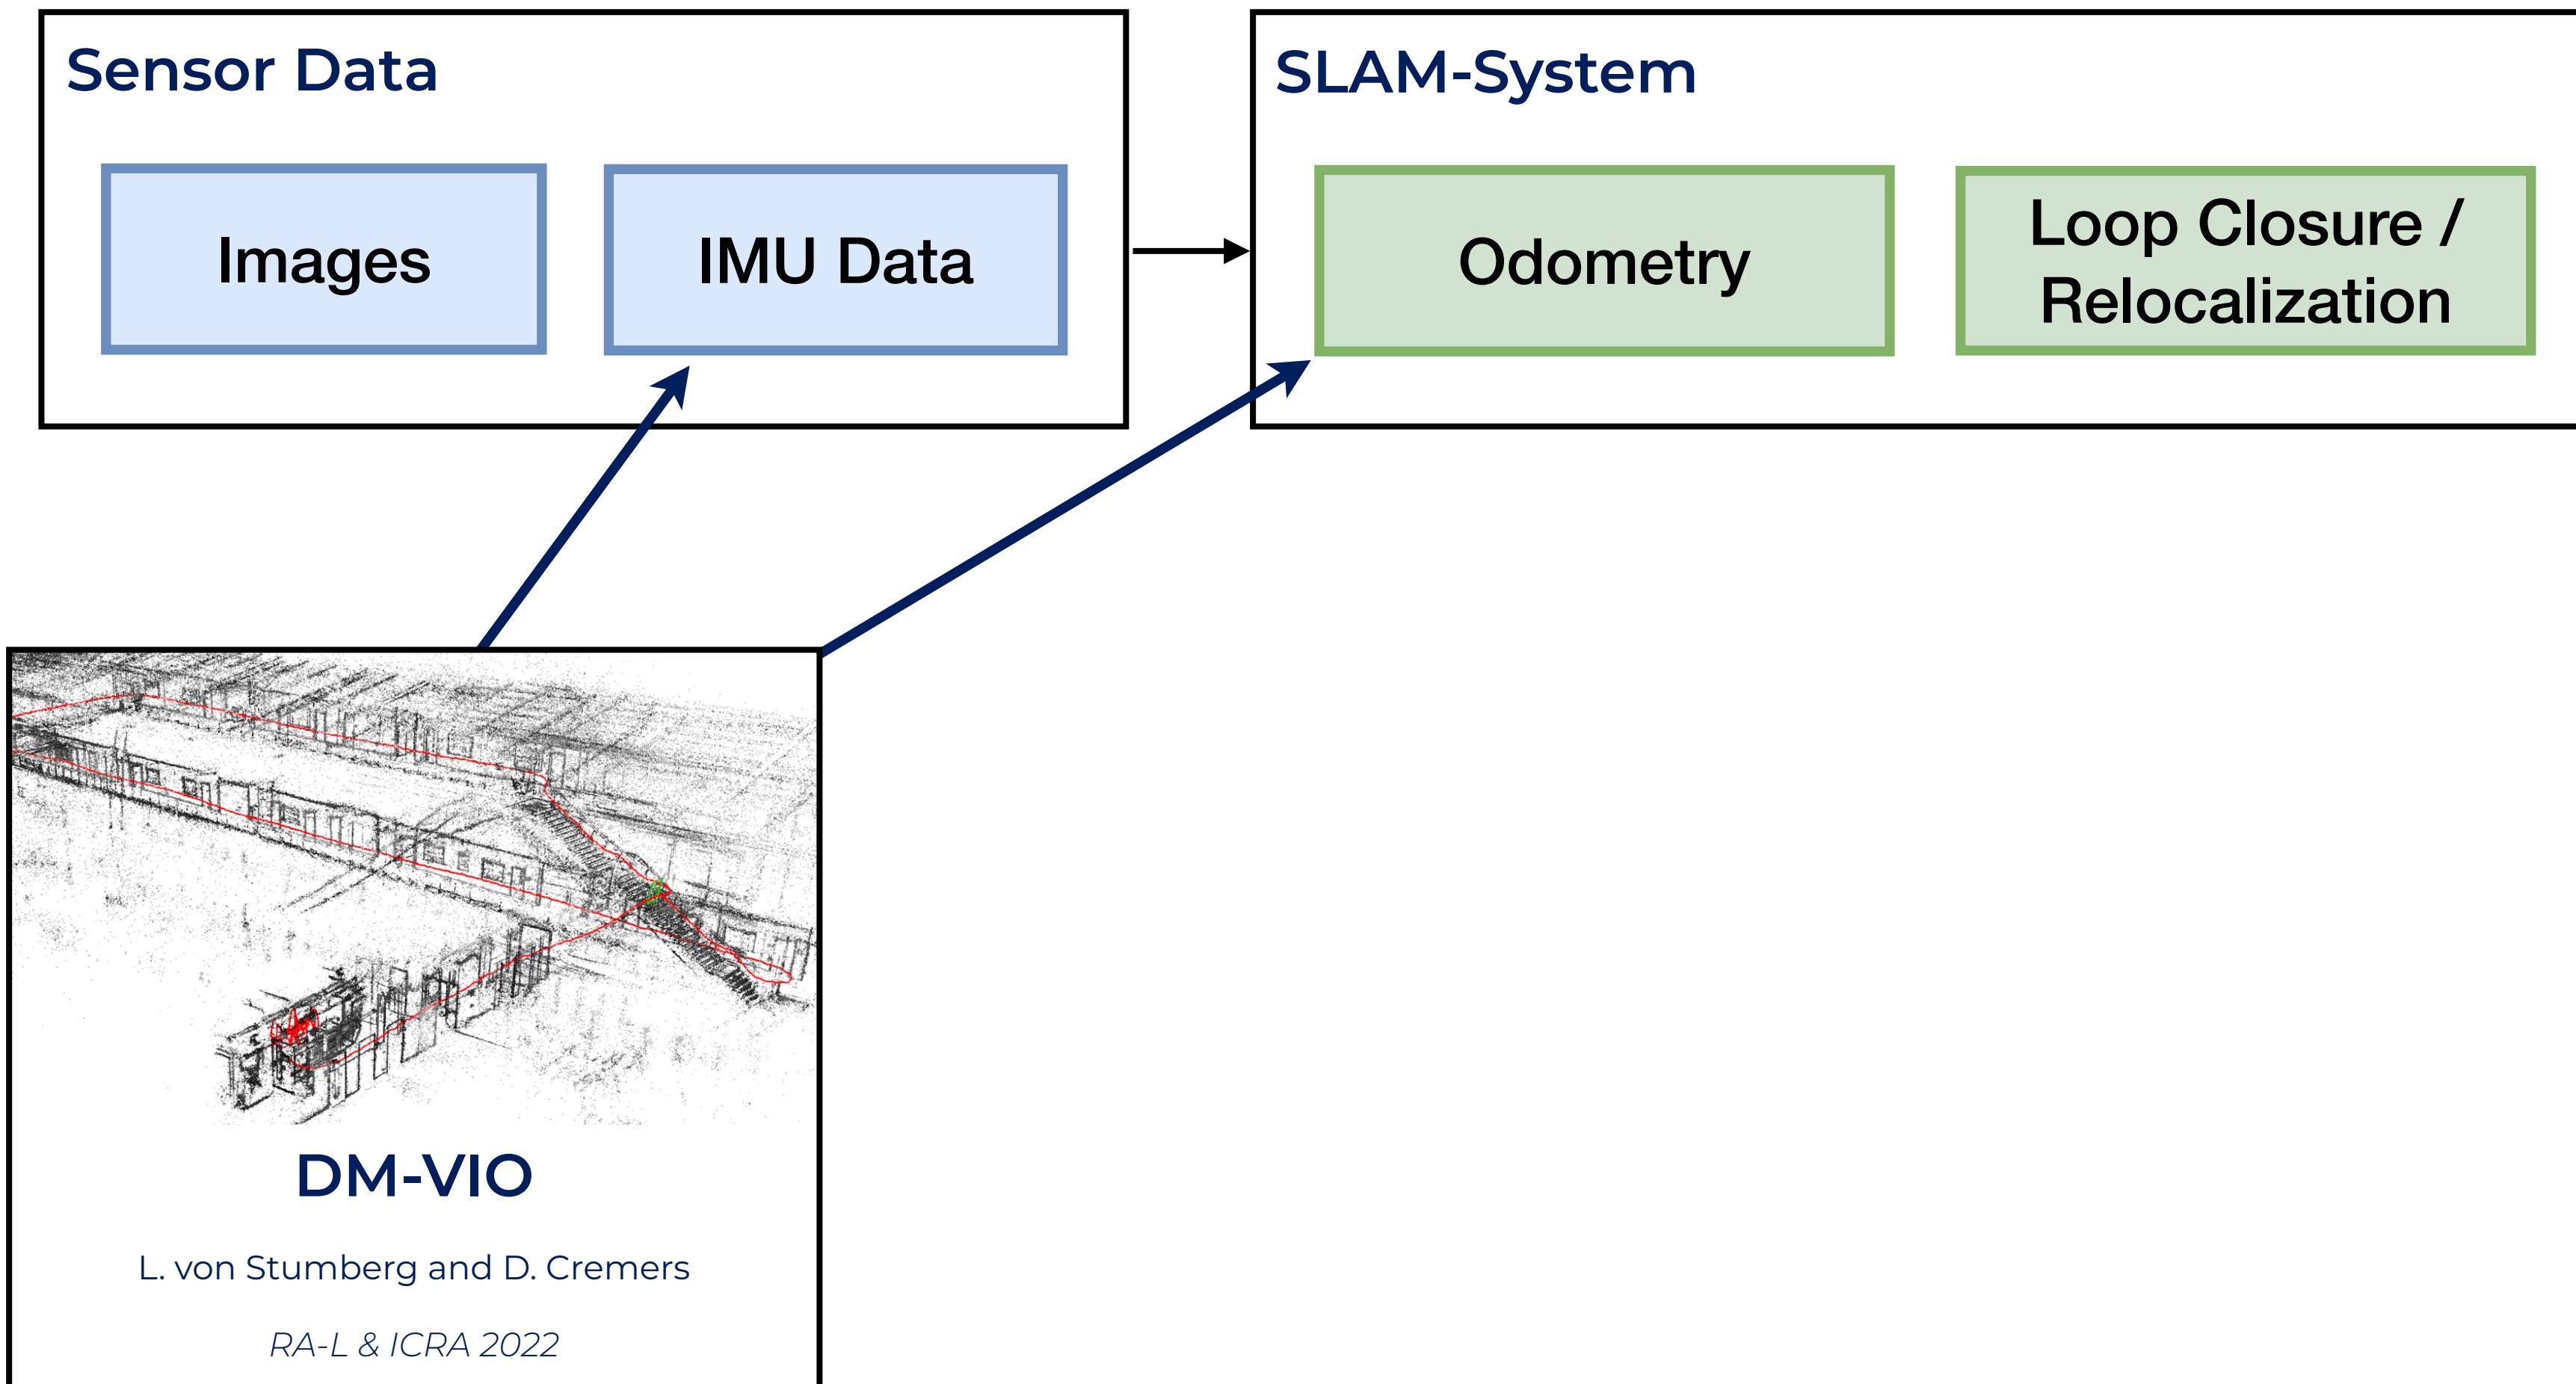
Visual SLAM

From Optimization to Learning



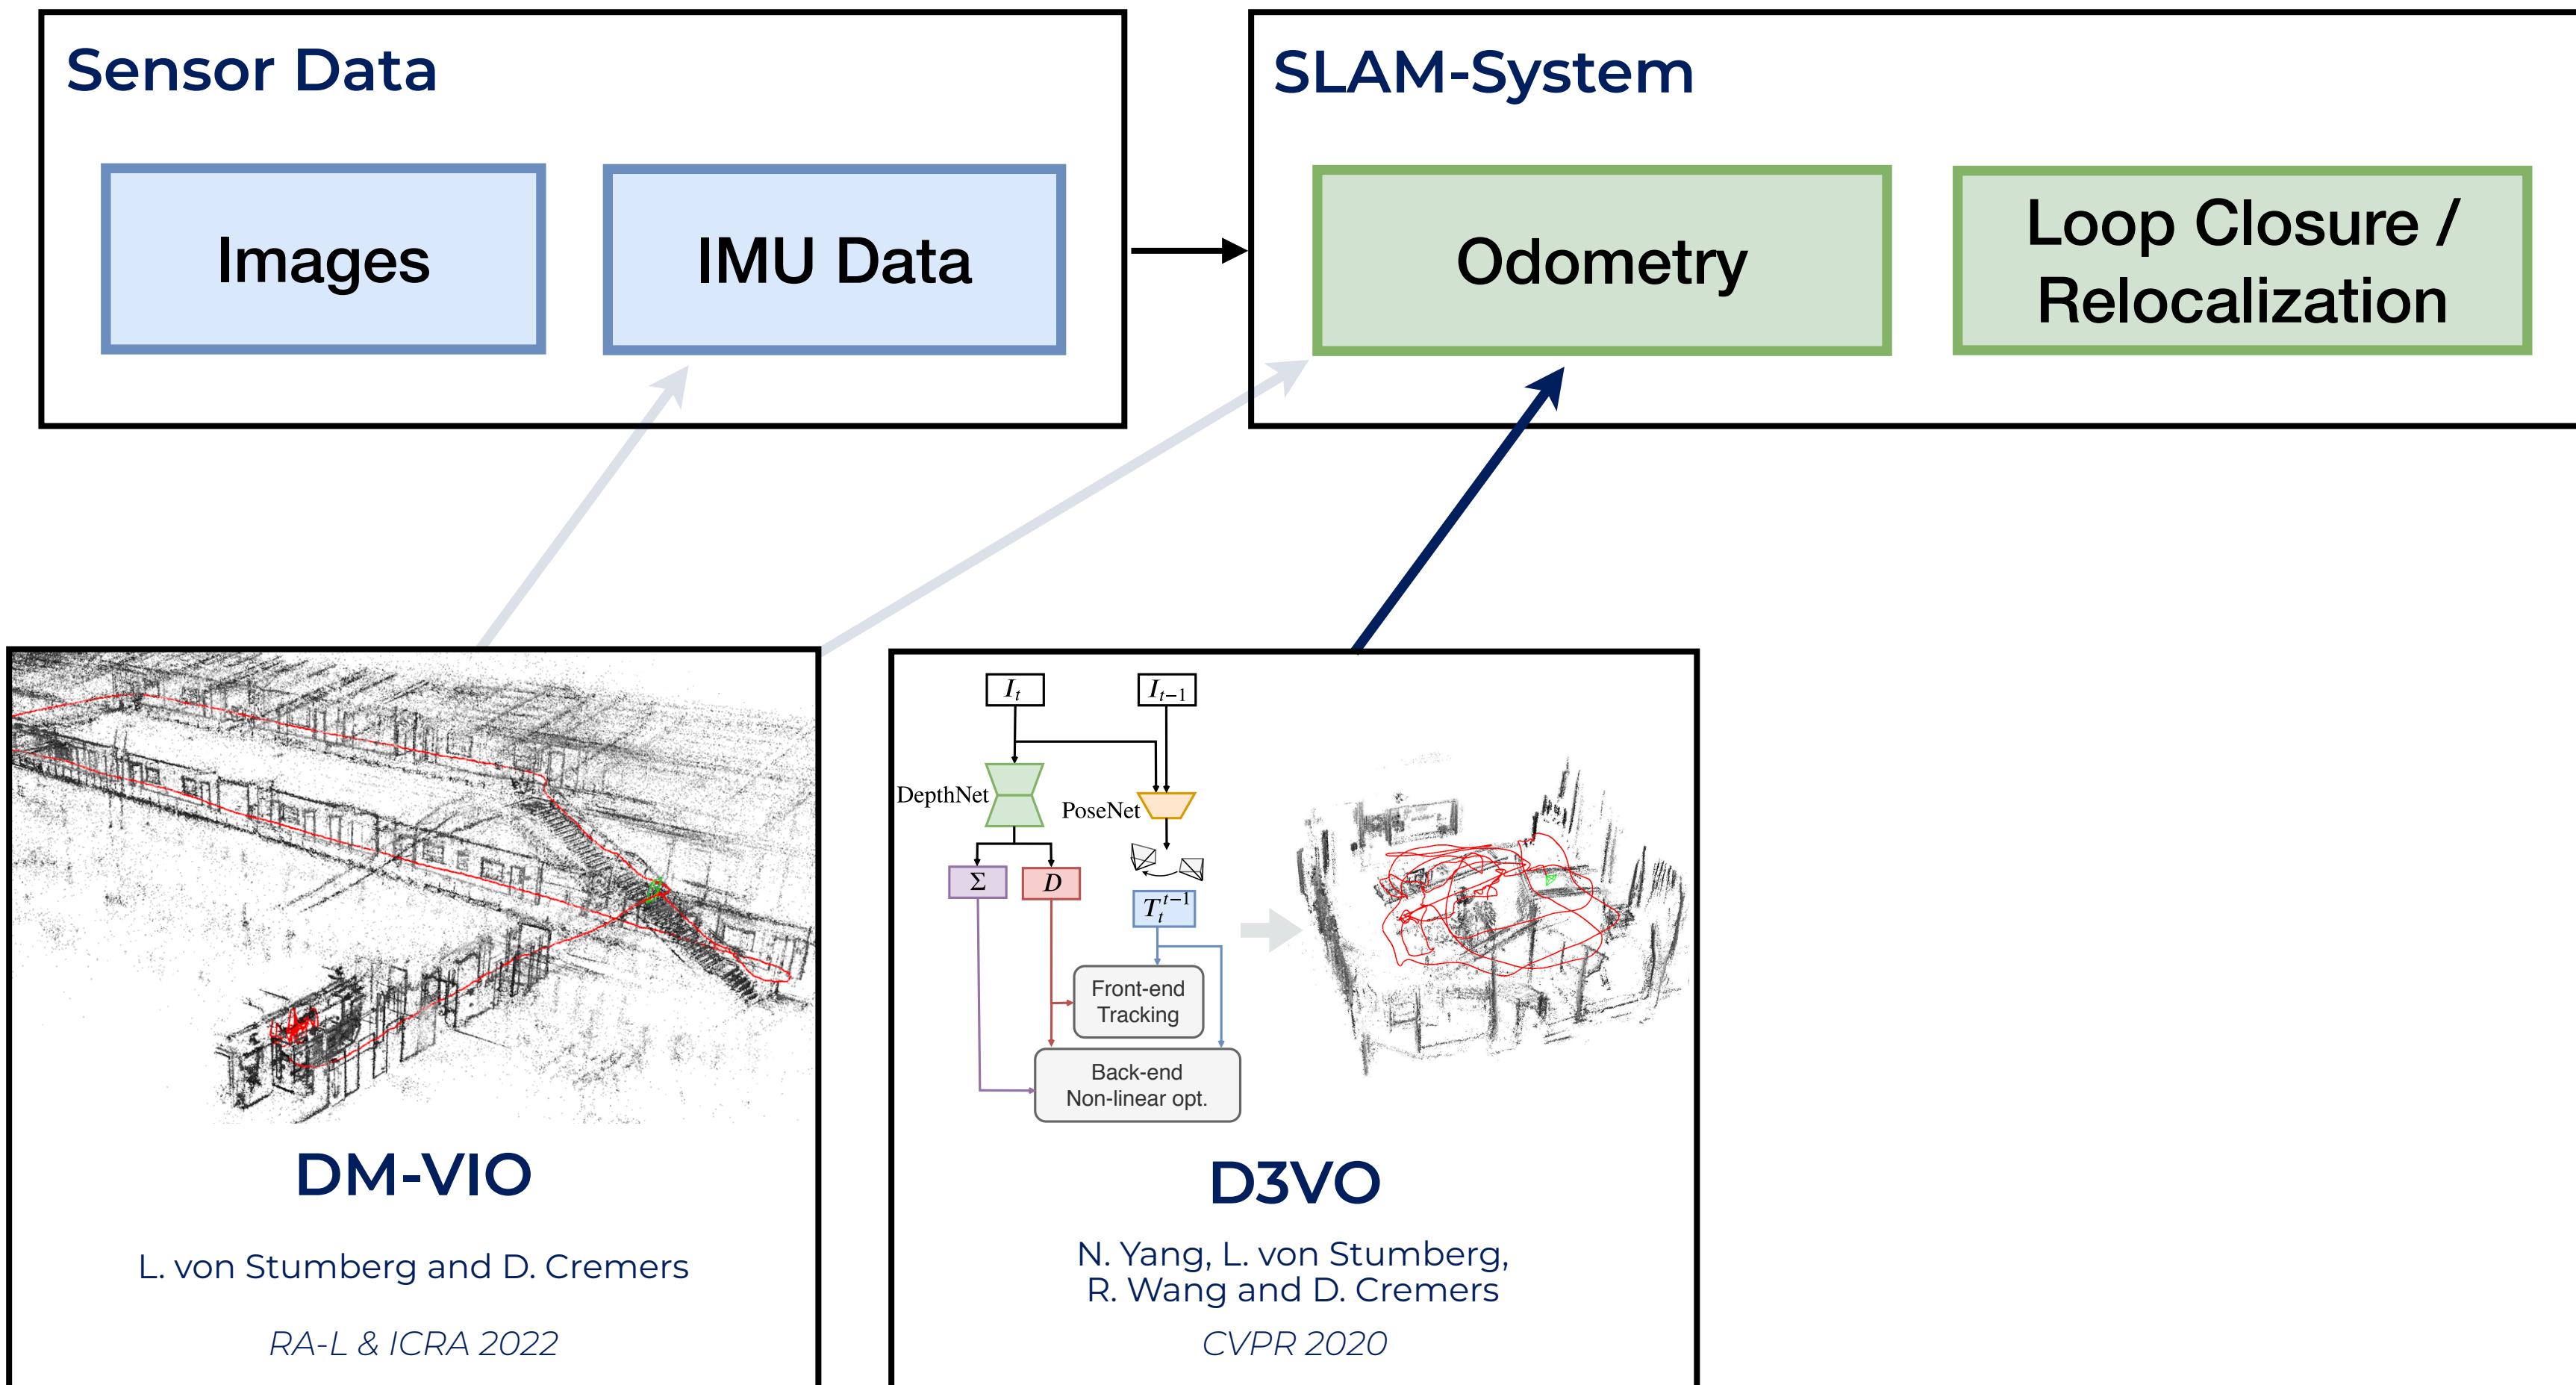
Visual SLAM

From Optimization to Learning



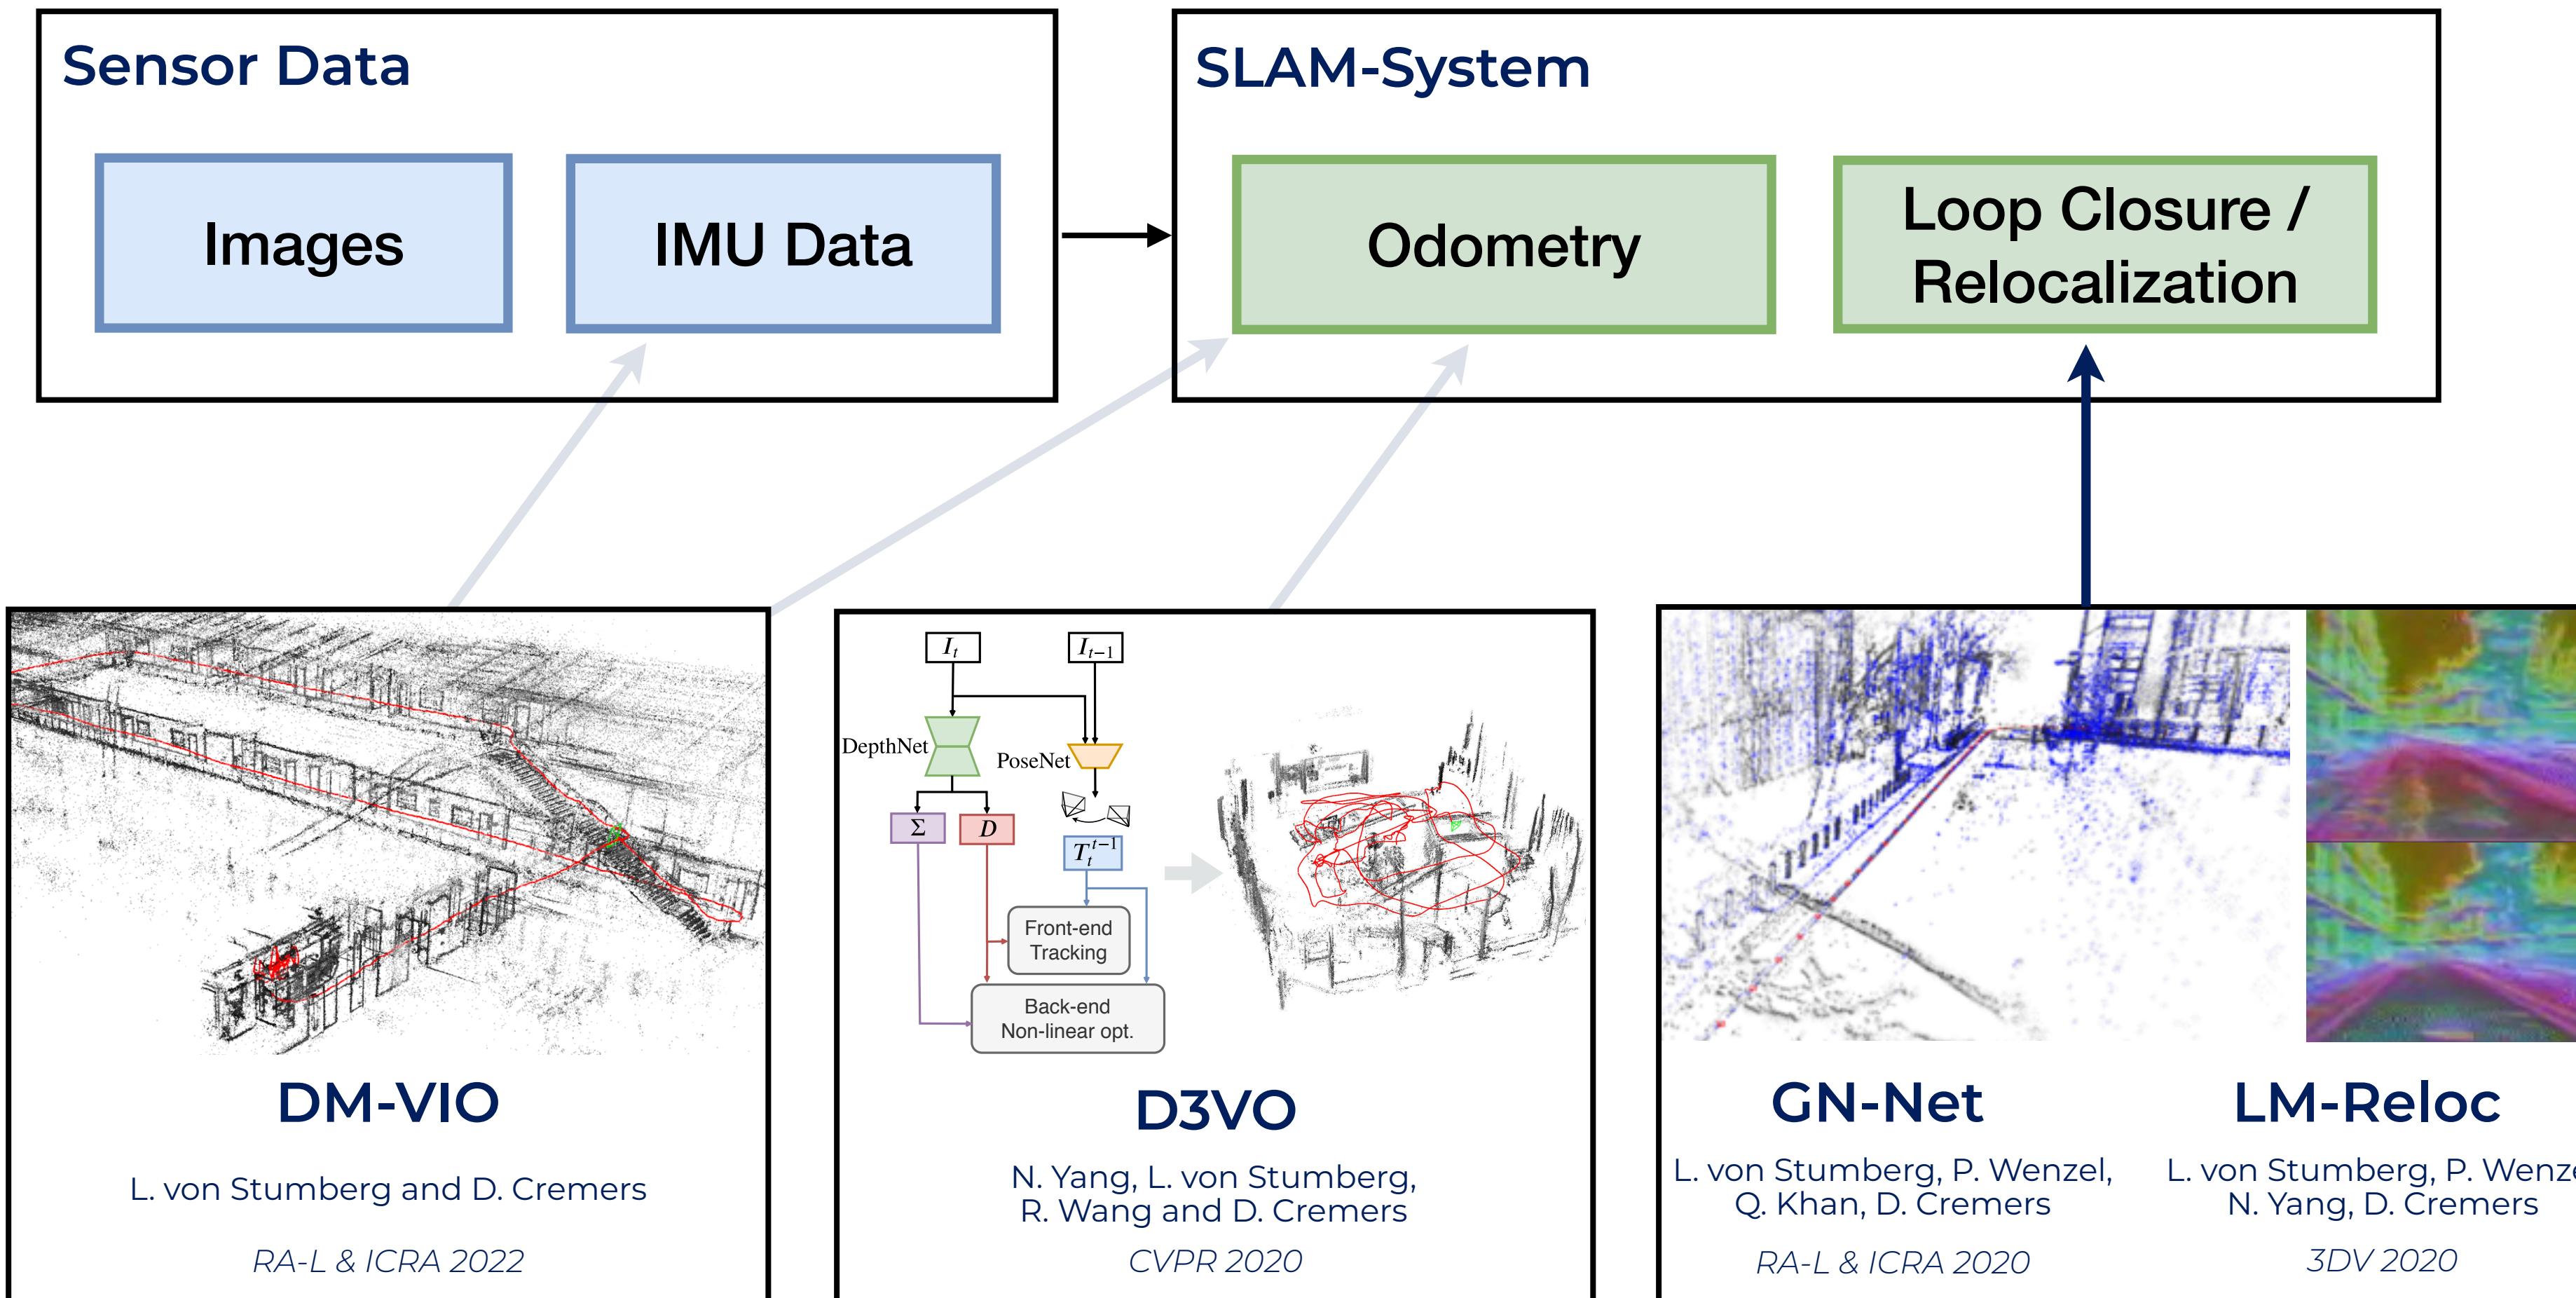
Visual SLAM

From Optimization to Learning



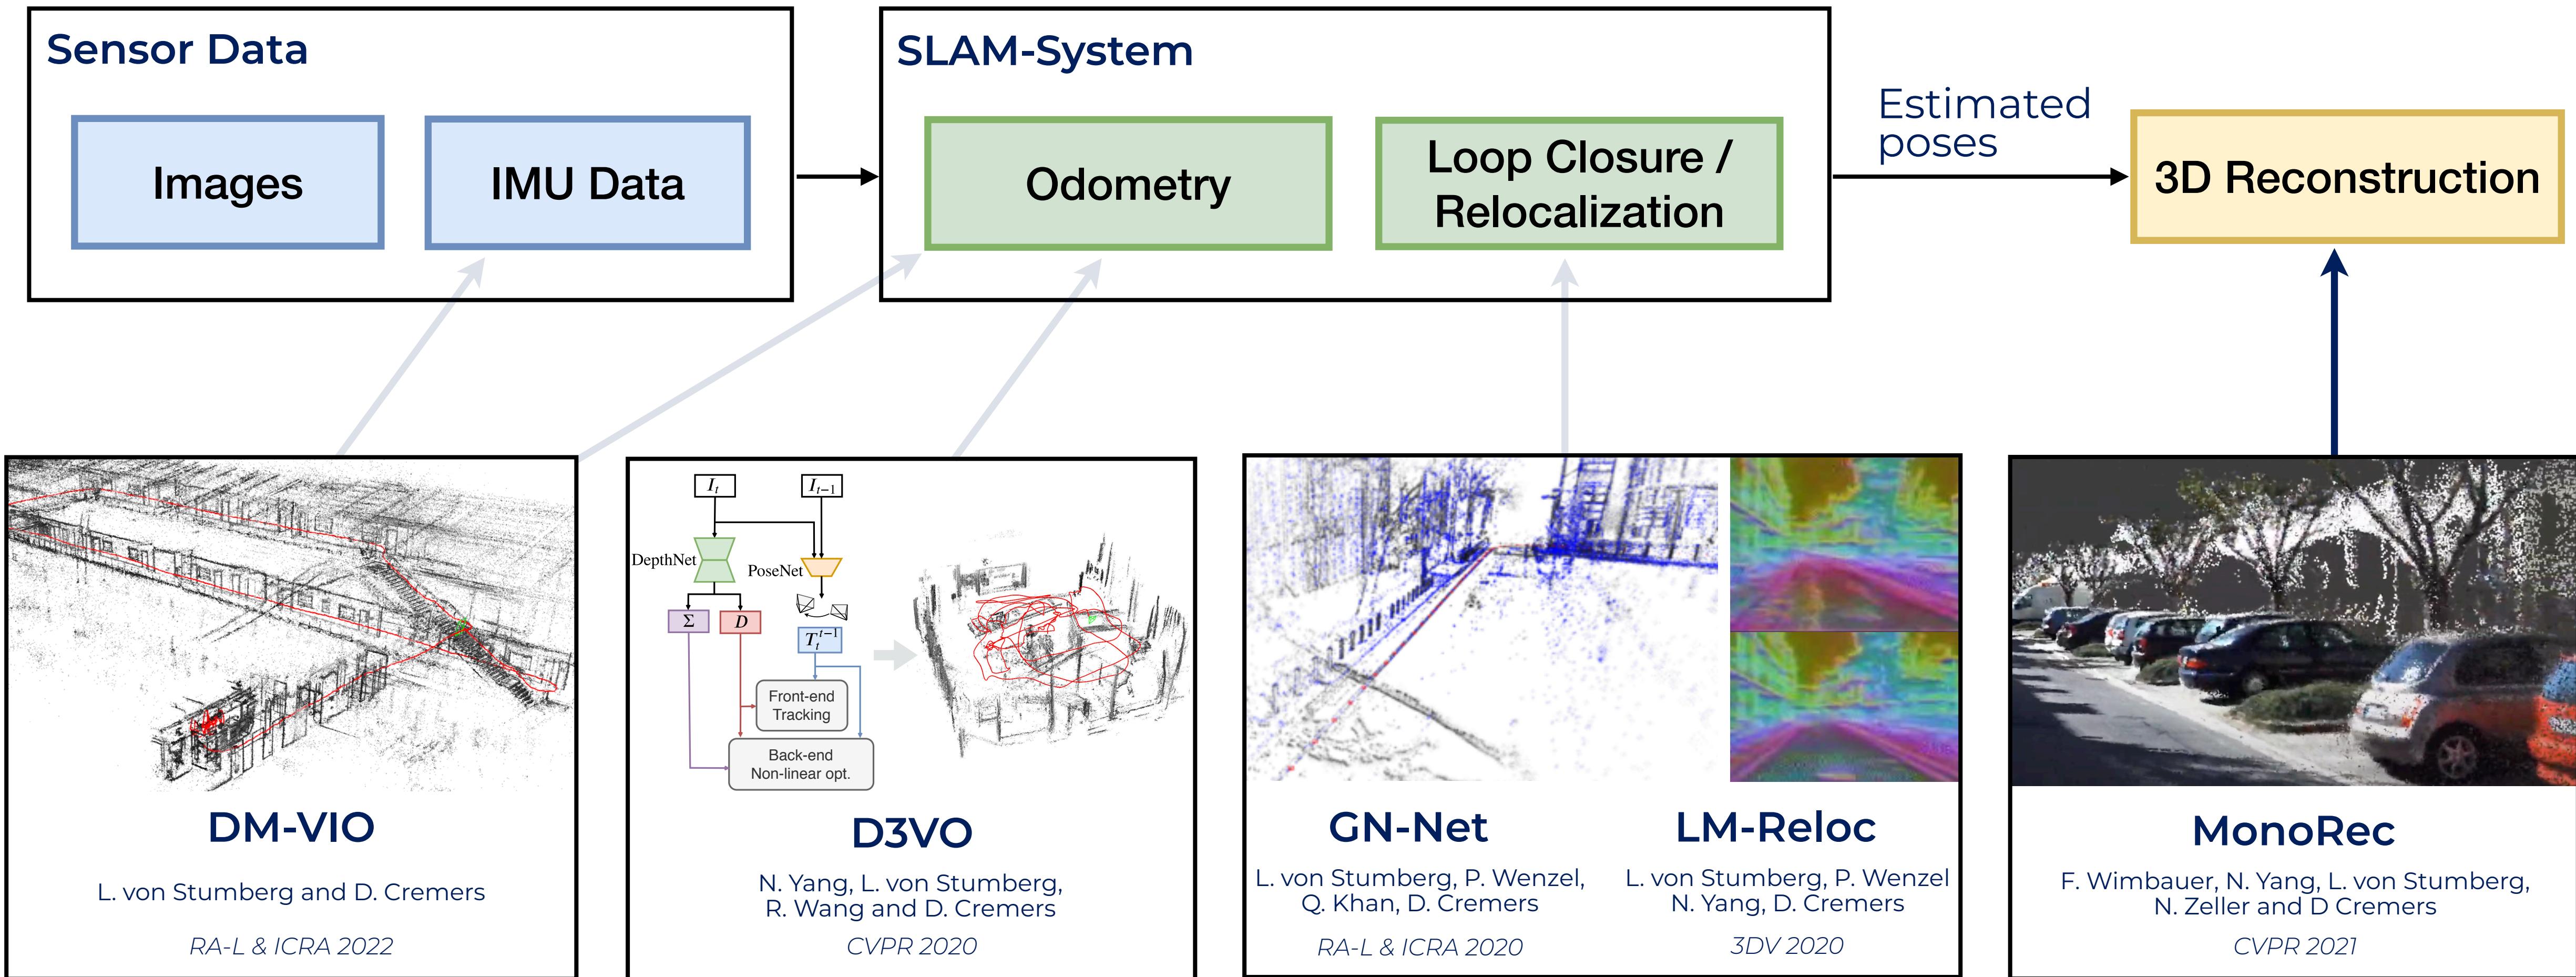
Visual SLAM

From Optimization to Learning



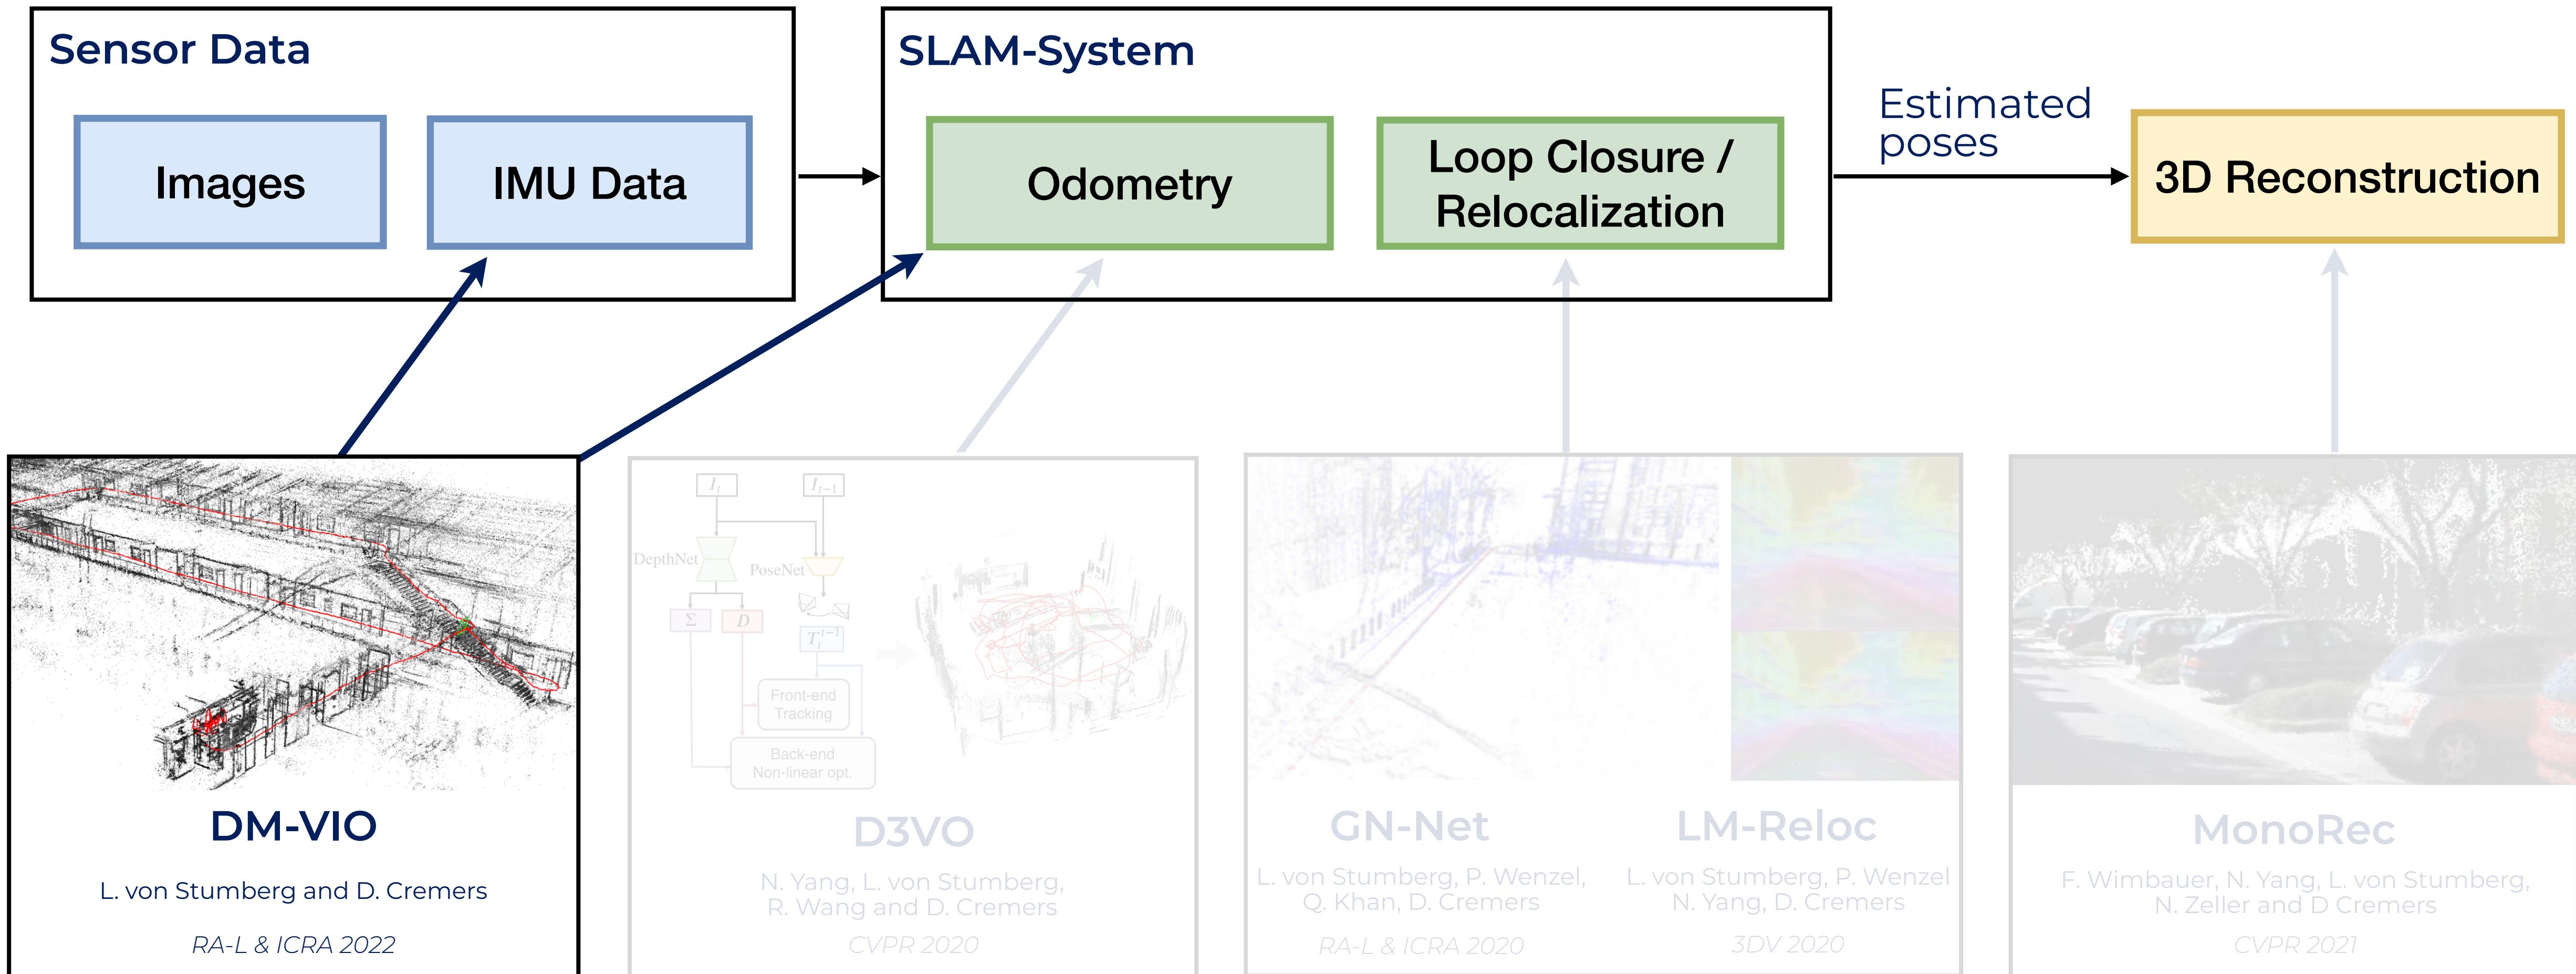
Visual SLAM

From Optimization to Learning



Visual SLAM

From Optimization to Learning



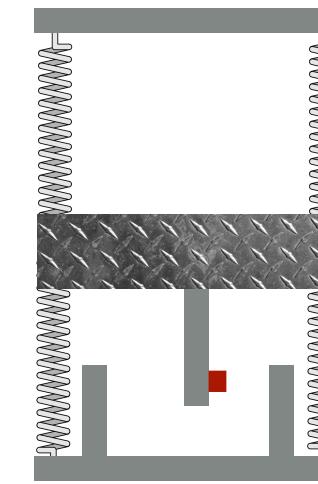
Visual-Inertial Odometry

Track motion of the system using

Camera



IMU (Inertial Measurement Unit)



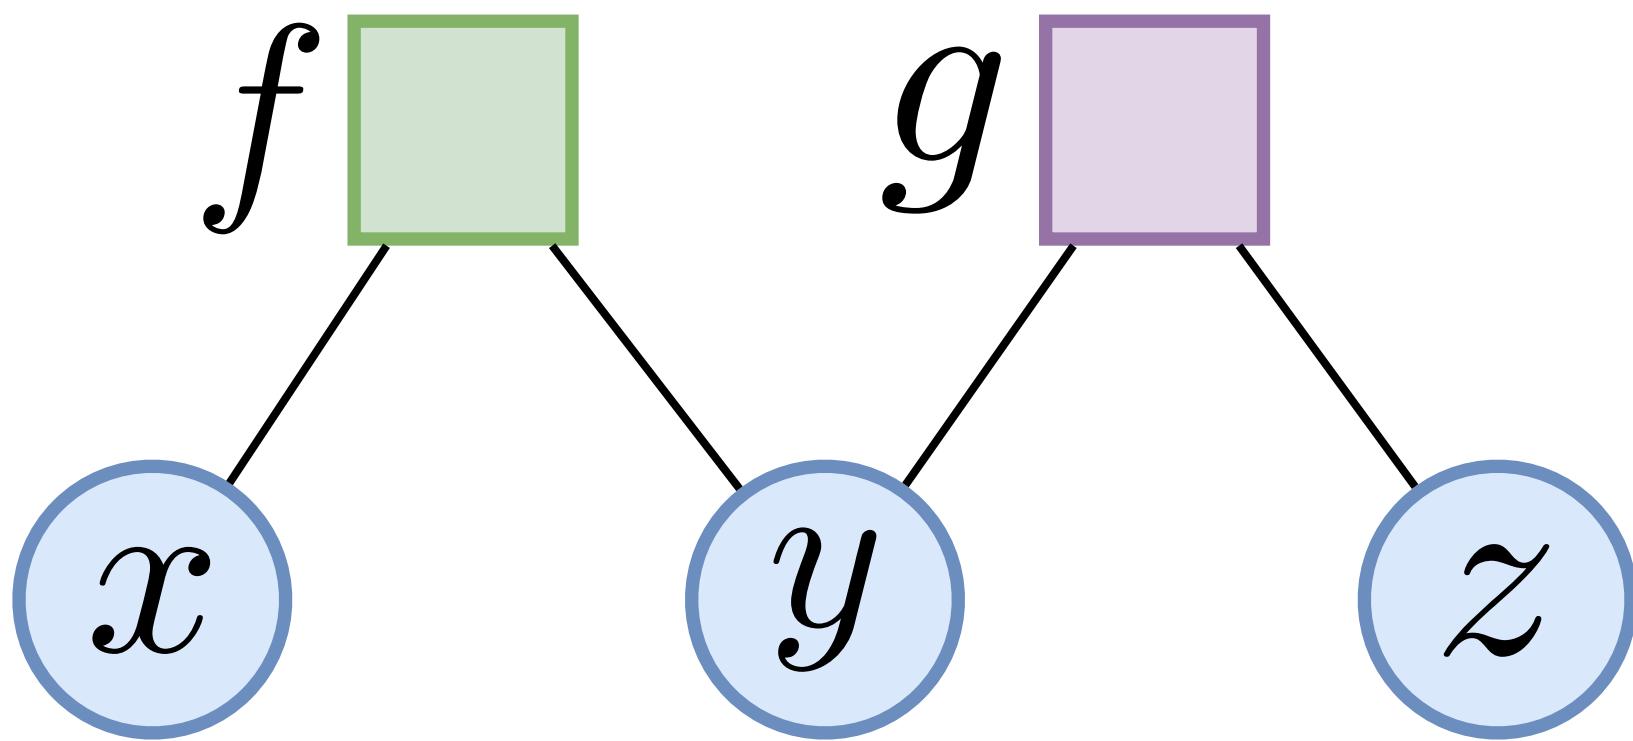
accelerometer + gyroscope

$$E = E_{\text{visual}} + E_{\text{inertial}}$$

Related work:

A. Mourikis and S. Roumeliotis. A multi-state constraint Kalman filter for vision-aided inertial navigation. ICRA 2007

Factor graphs



- **Energy function:**

$$E = f(x, y) + g(y, z)$$

Bundle Adjustment

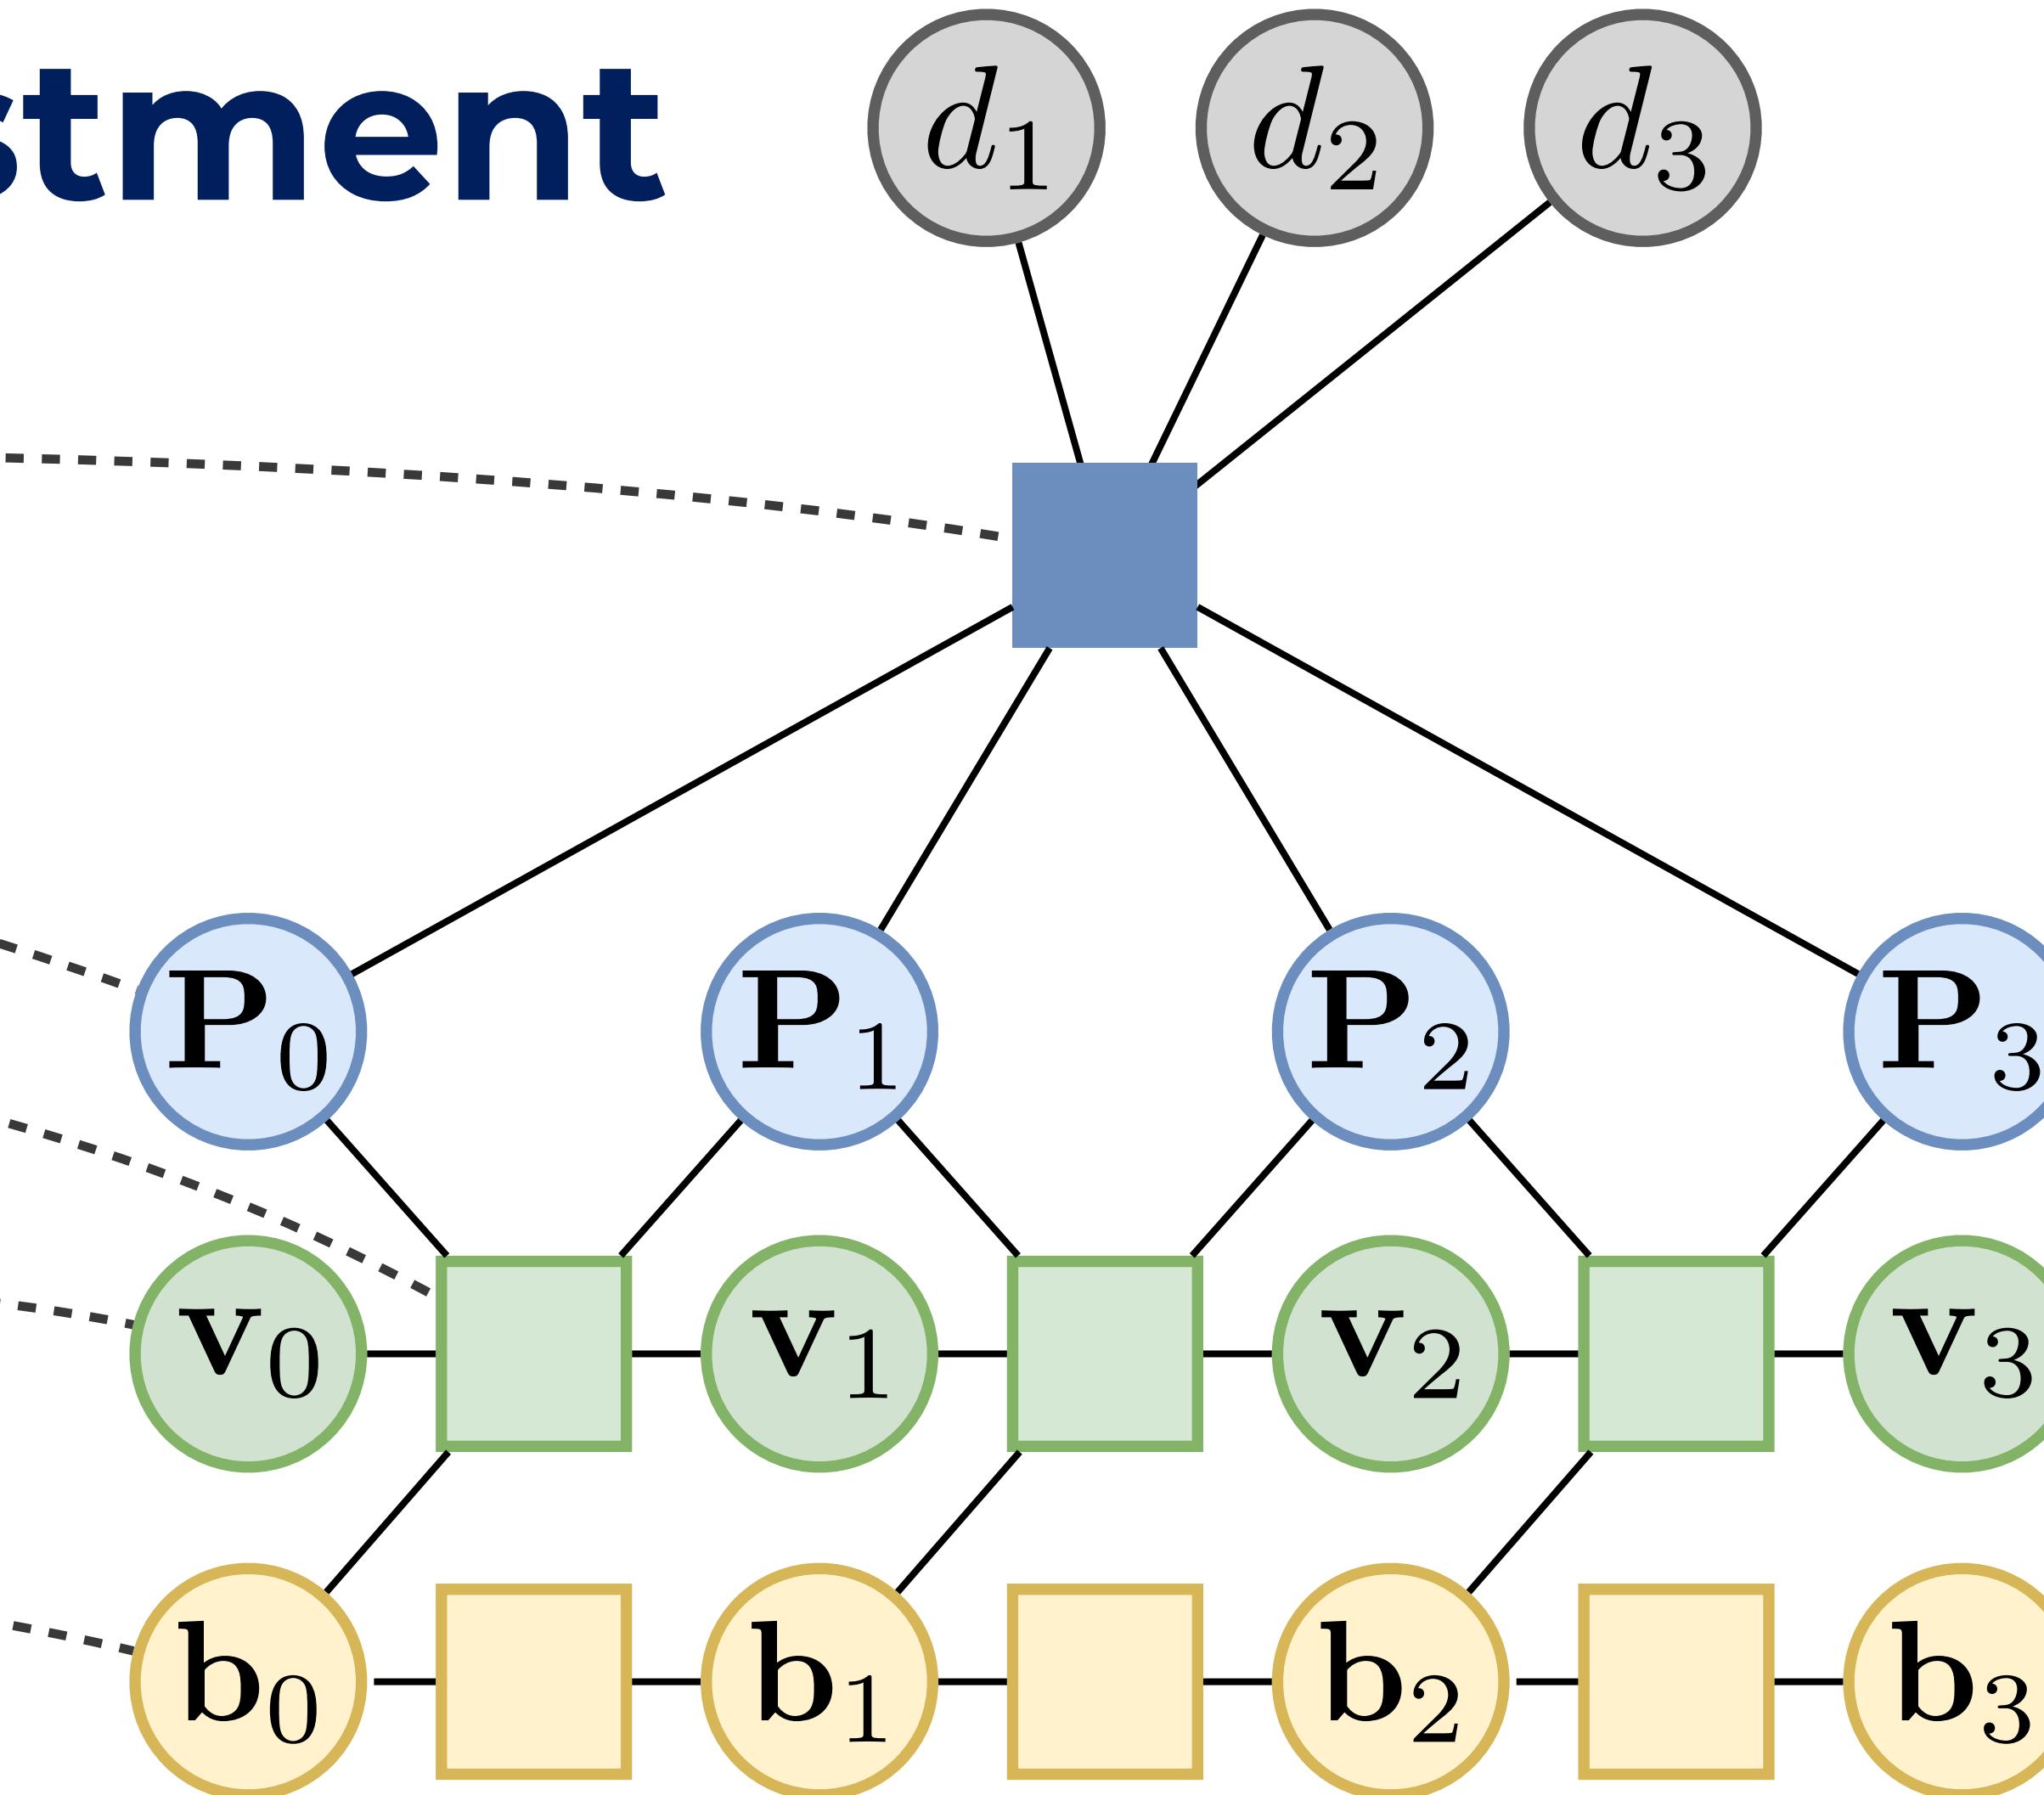
photometric error

poses

inertial error

velocities

biases



Related work:

C. Forster, L. Carlone, F. Dellaert, and D. Scaramuzza.

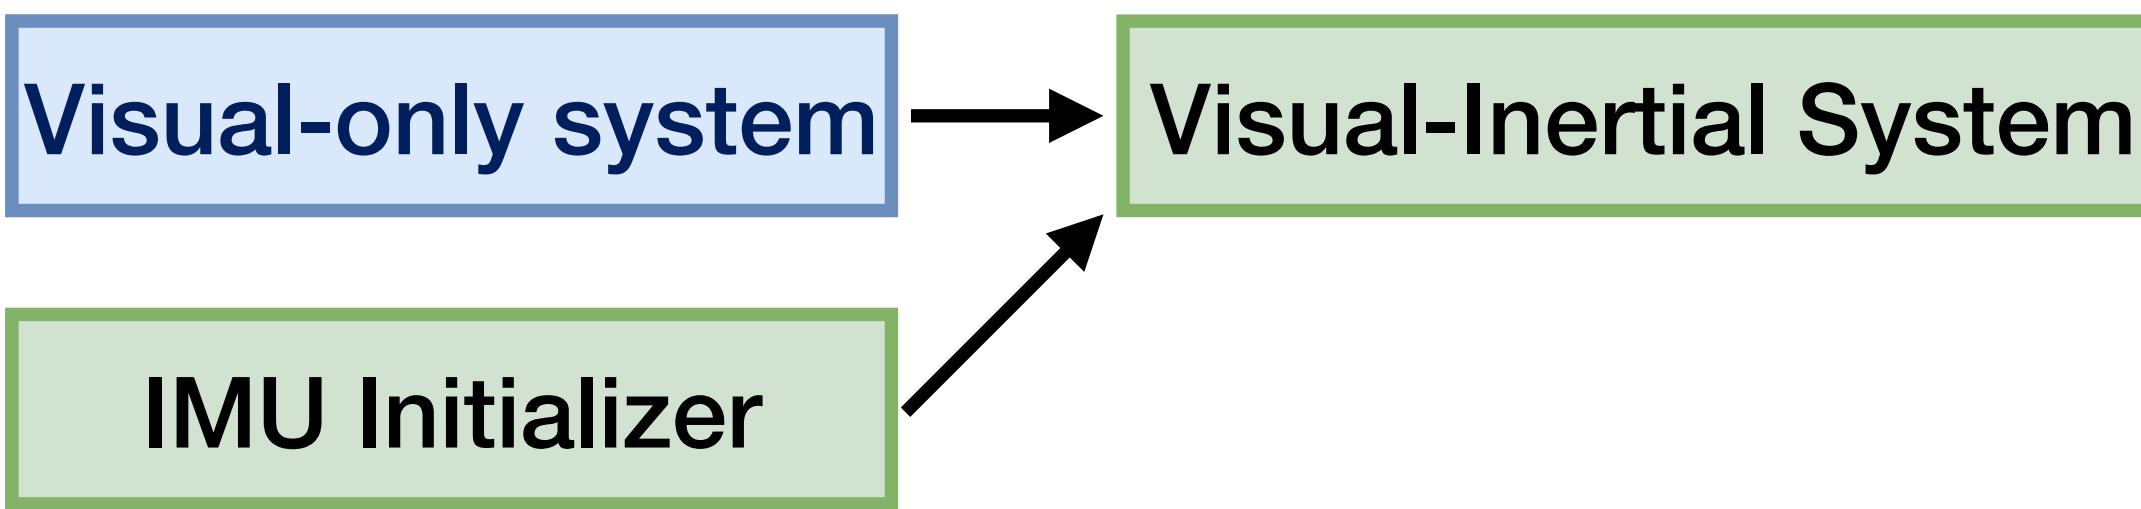
Reliable integration on manifold for efficient visual-inertial

multi-sensor odometry, RSS 2015, Sparse Odometry, TPAMI 2017

$$E(\mathbf{s}) = E_{\text{photo}}$$

IMU Initialization

Most systems (e.g. ORB-SLAM 3)



- Trade-off when to initialize. If too early the scale might be inaccurate

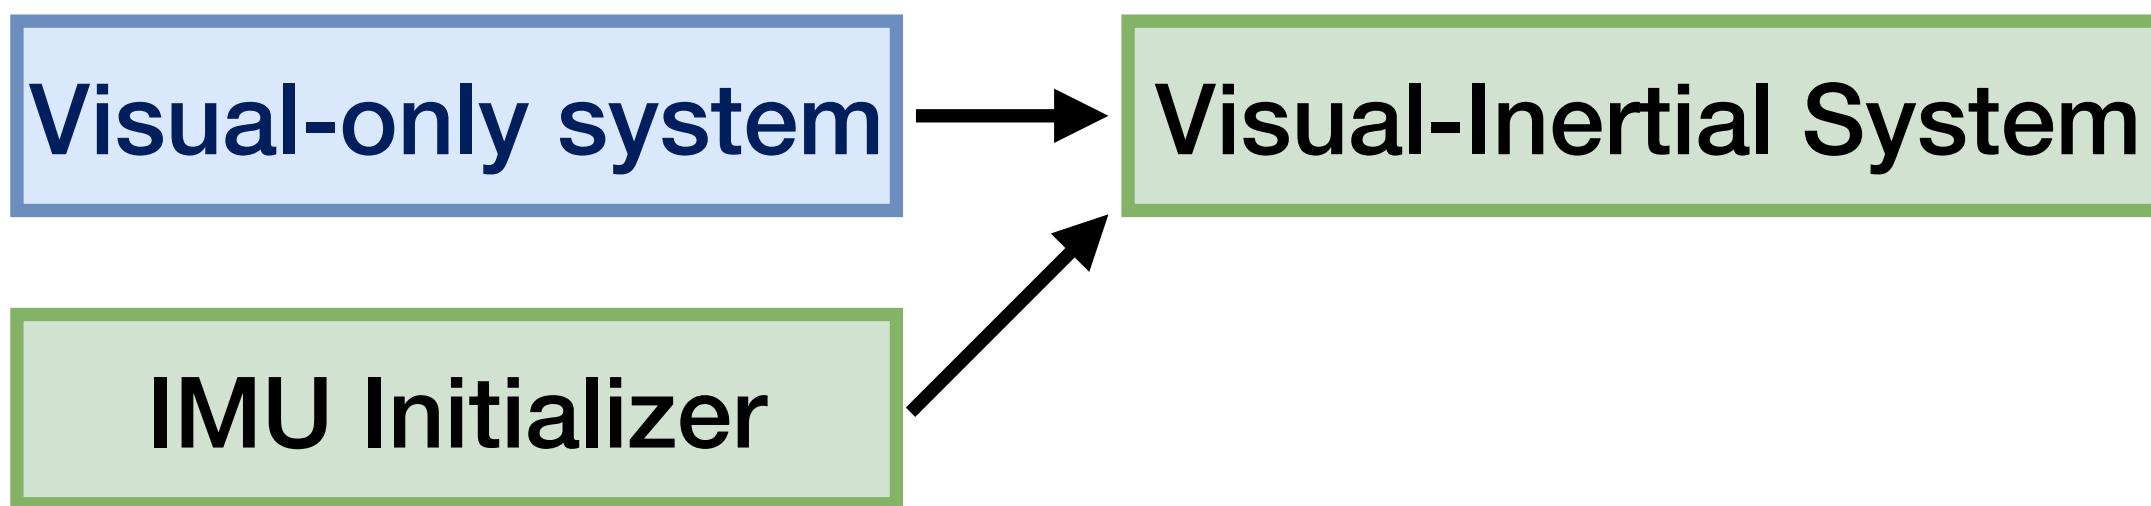
VI-DSO (previous work)

Visual-Inertial System
with scale optimization +
special marginalization

- Pro: Use IMU data immediately
- Might fail if initial scale is very off
- Scale convergence takes longer

IMU Initialization

Most systems (e.g. ORB-SLAM 3)



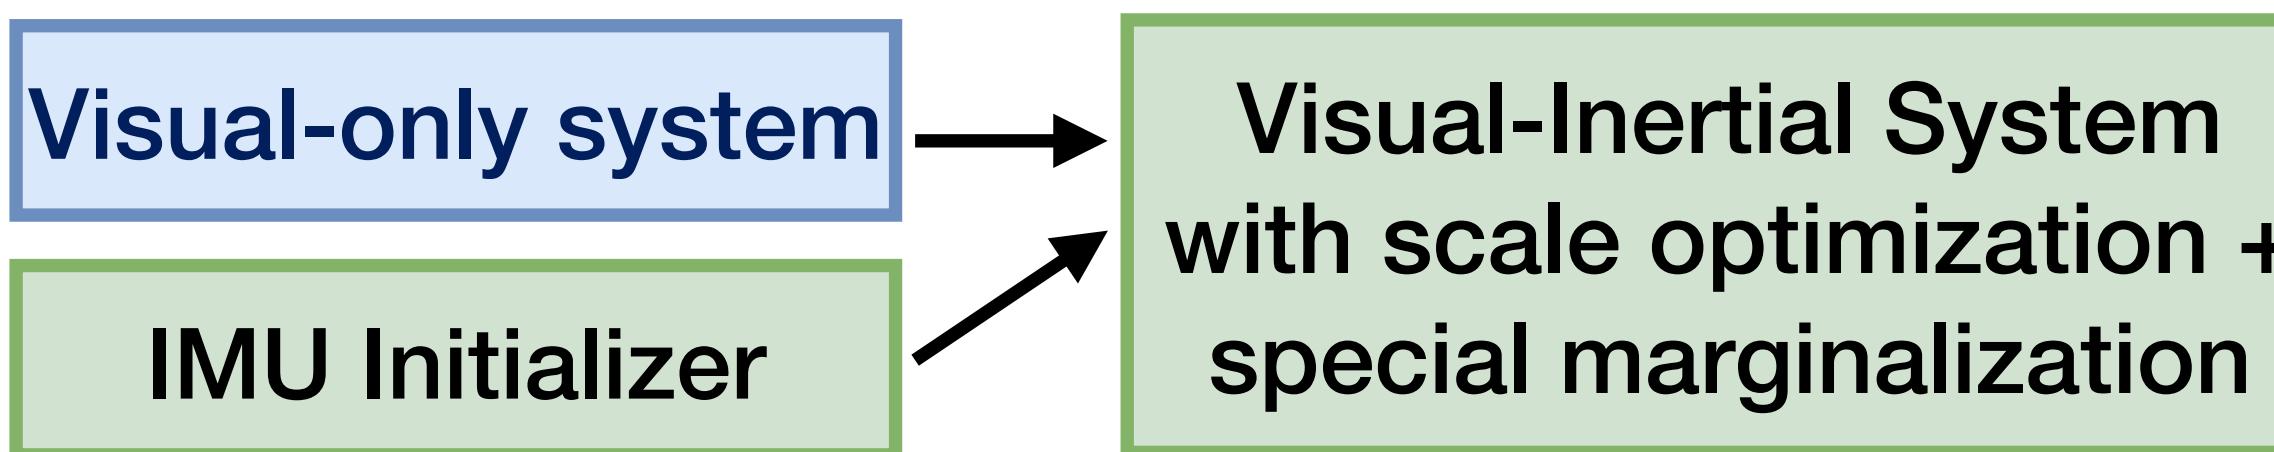
- Trade-off when to initialize. If too early the scale might be inaccurate

VI-DSO (previous work)

Visual-Inertial System
with scale optimization +
special marginalization

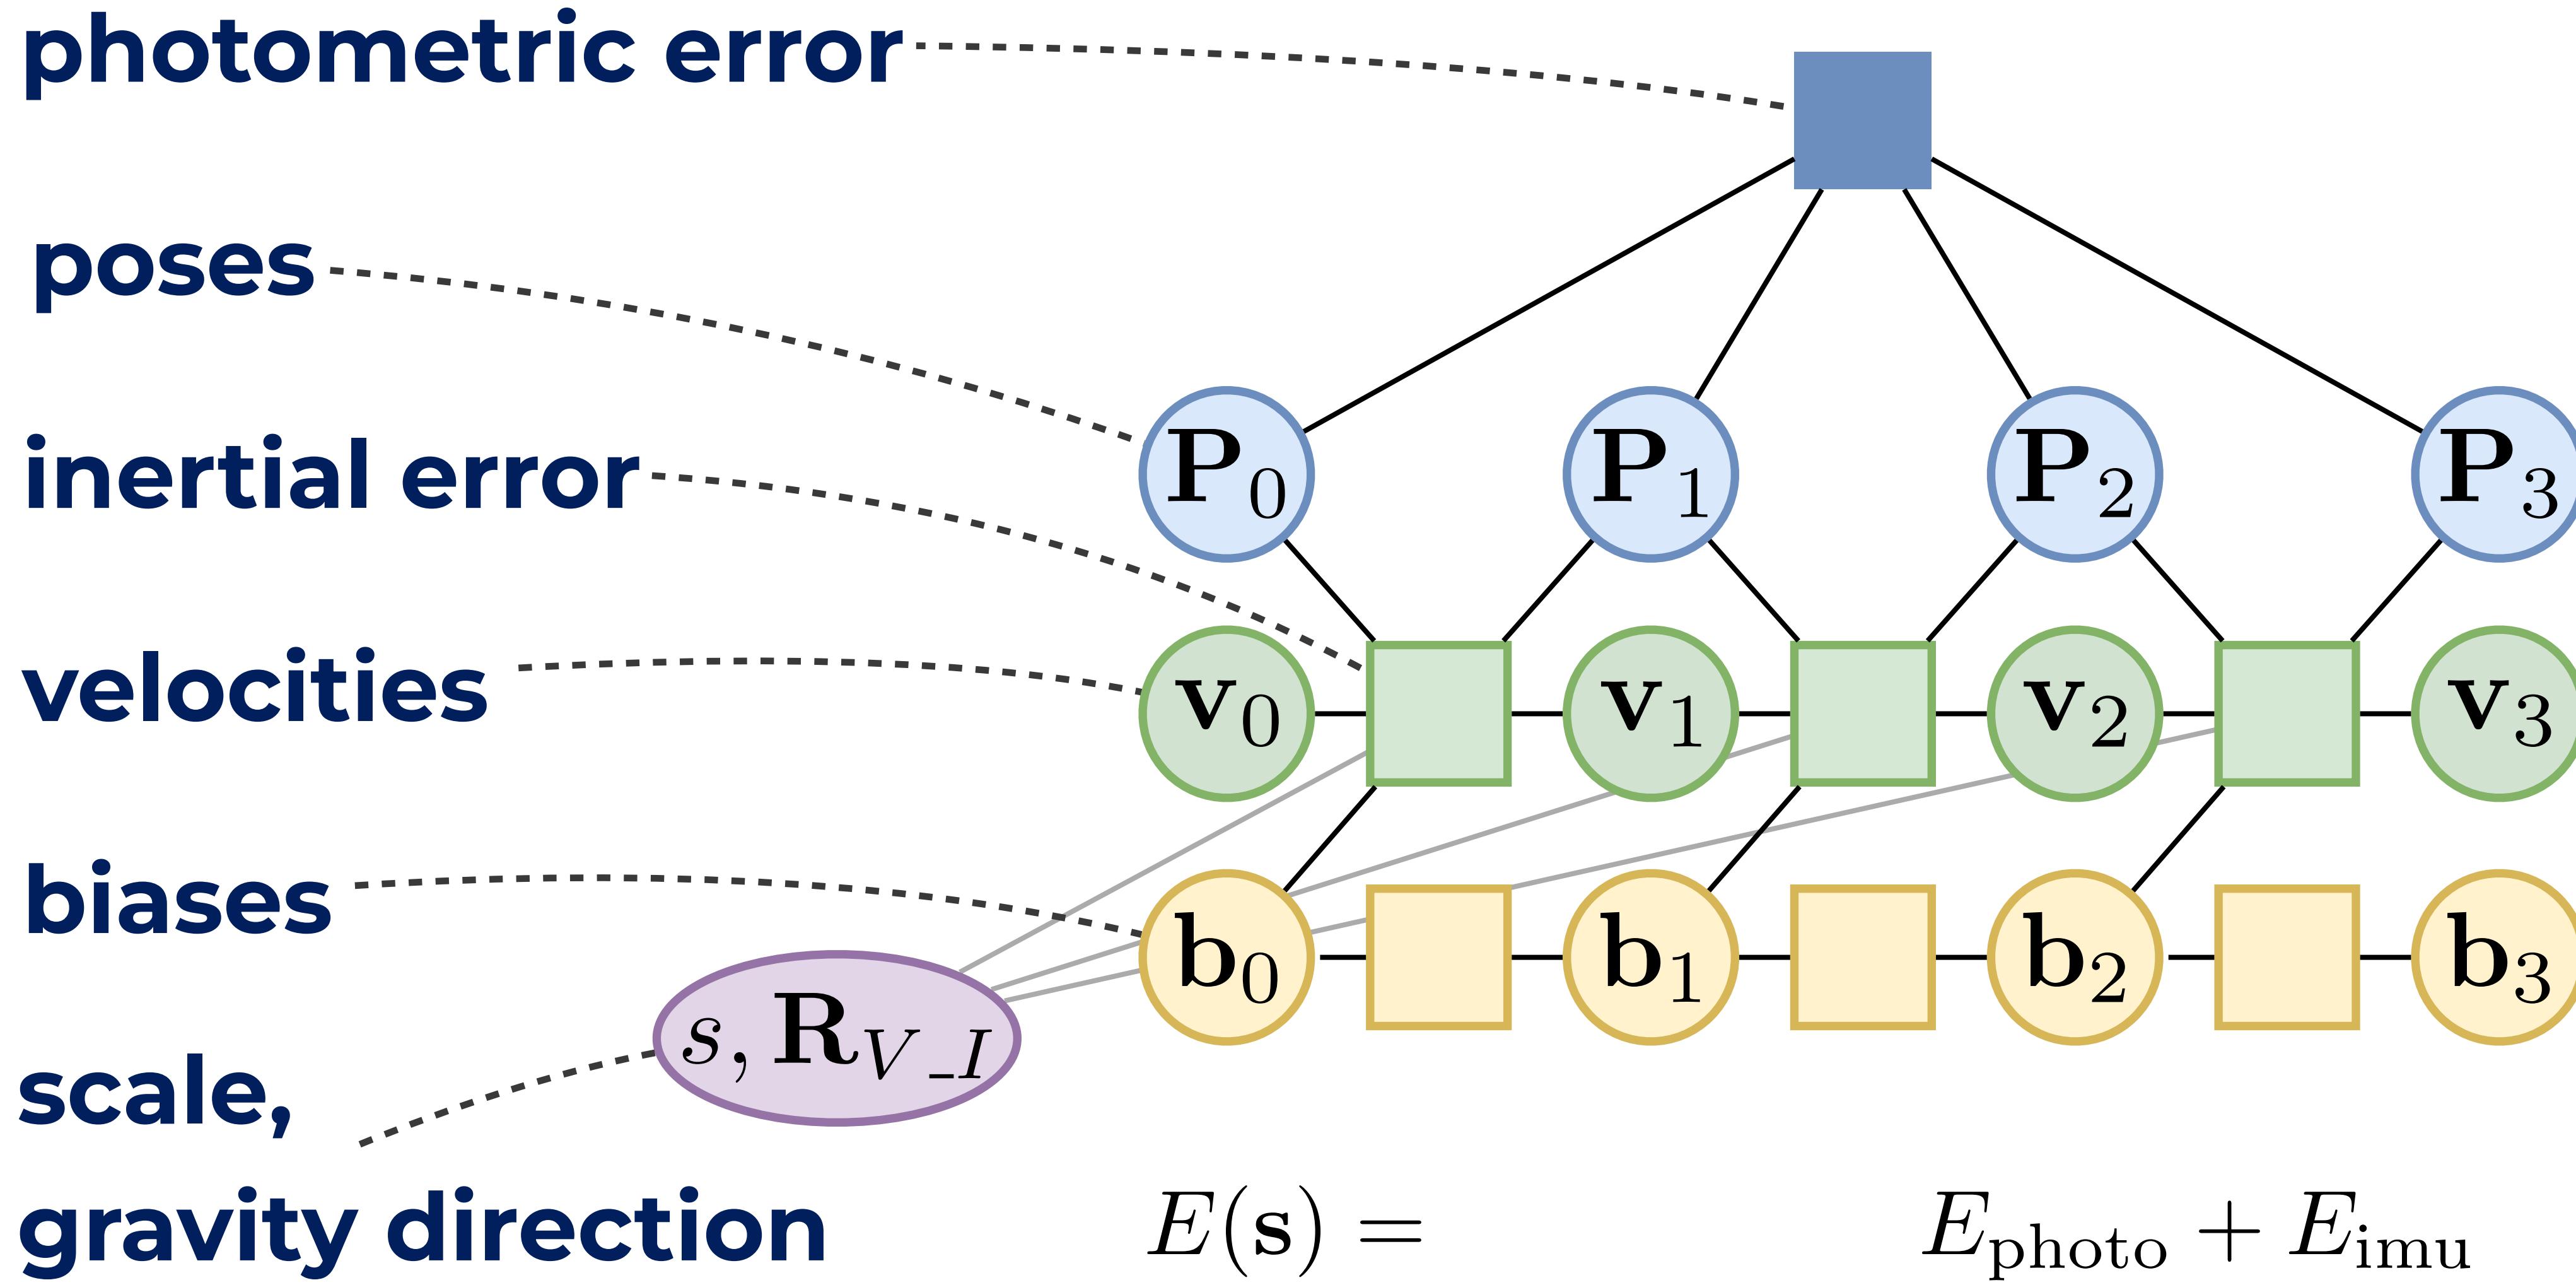
- Pro: Use IMU data immediately
- Might fail if initial scale is very off
- Scale convergence takes longer

DM-VIO



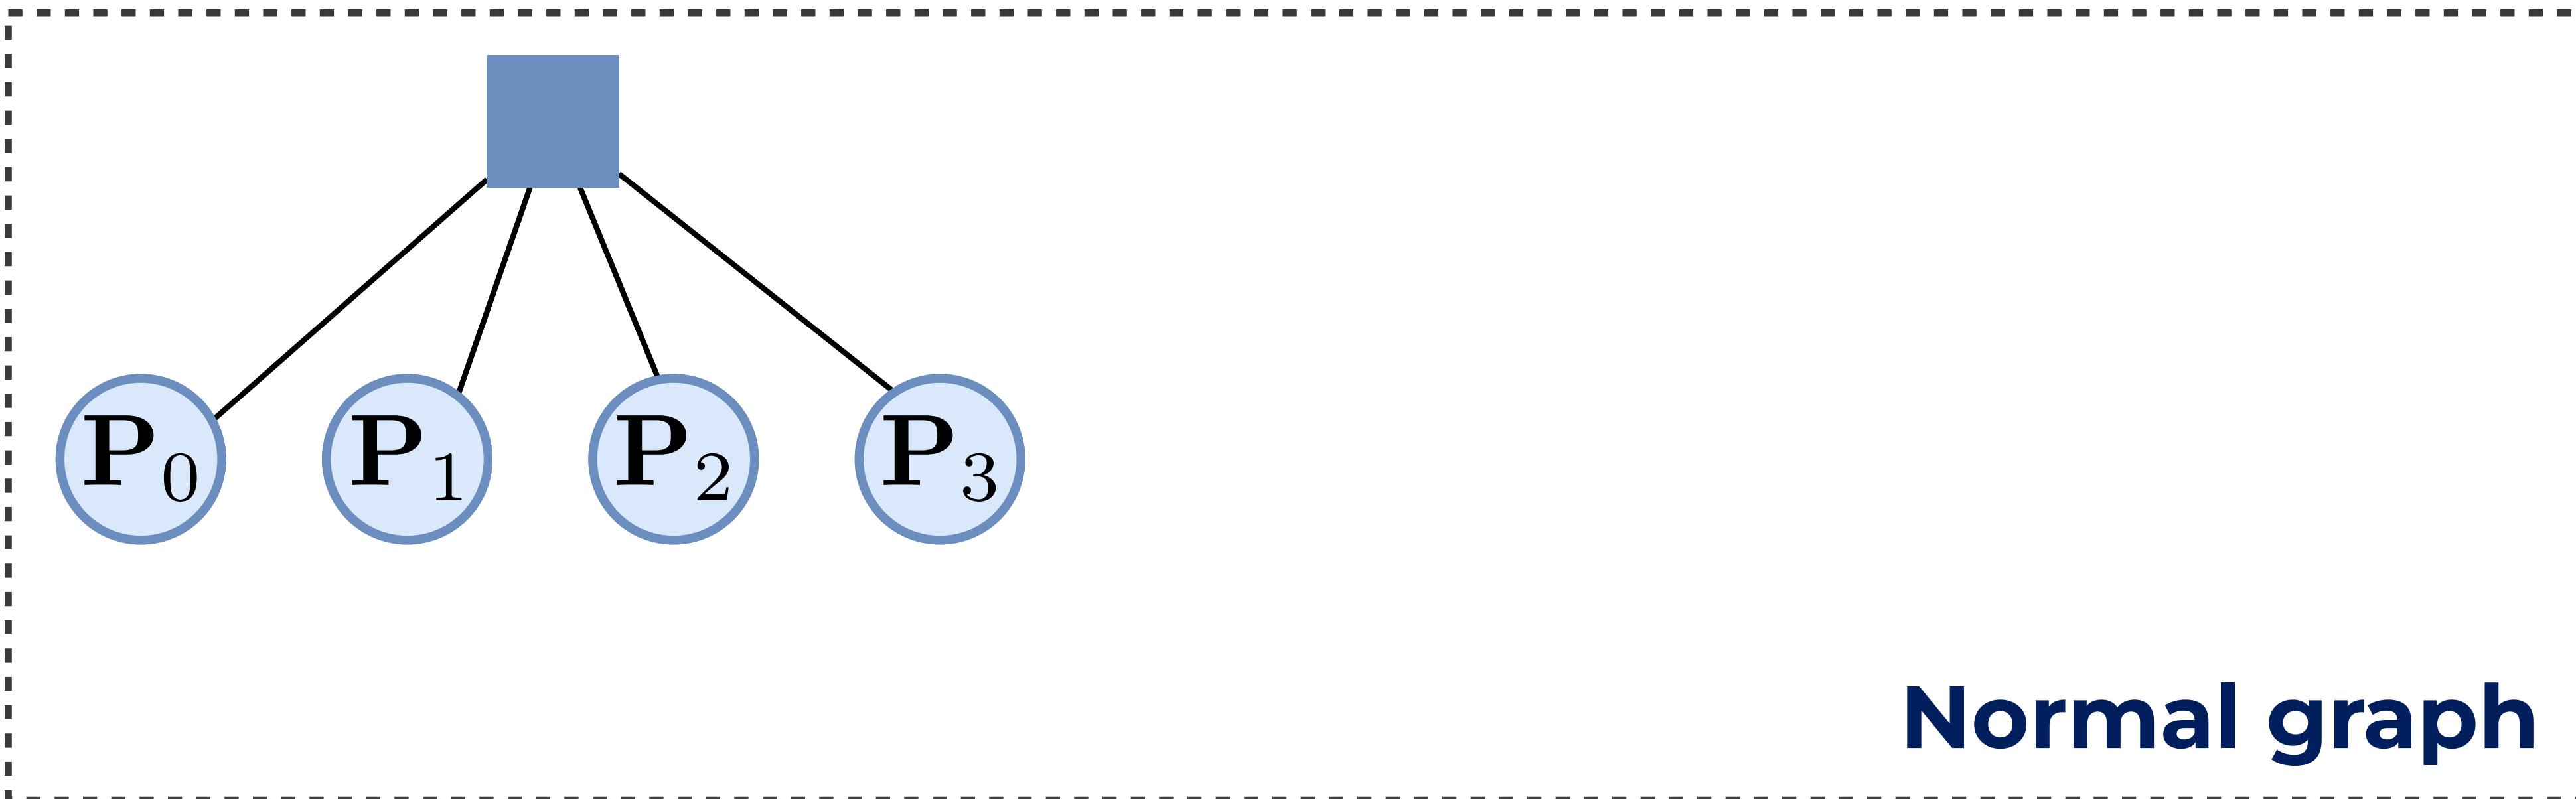
- We can initialize early
- Works in general environments
- Converges quickly as soon as scale is observable

Bundle Adjustment

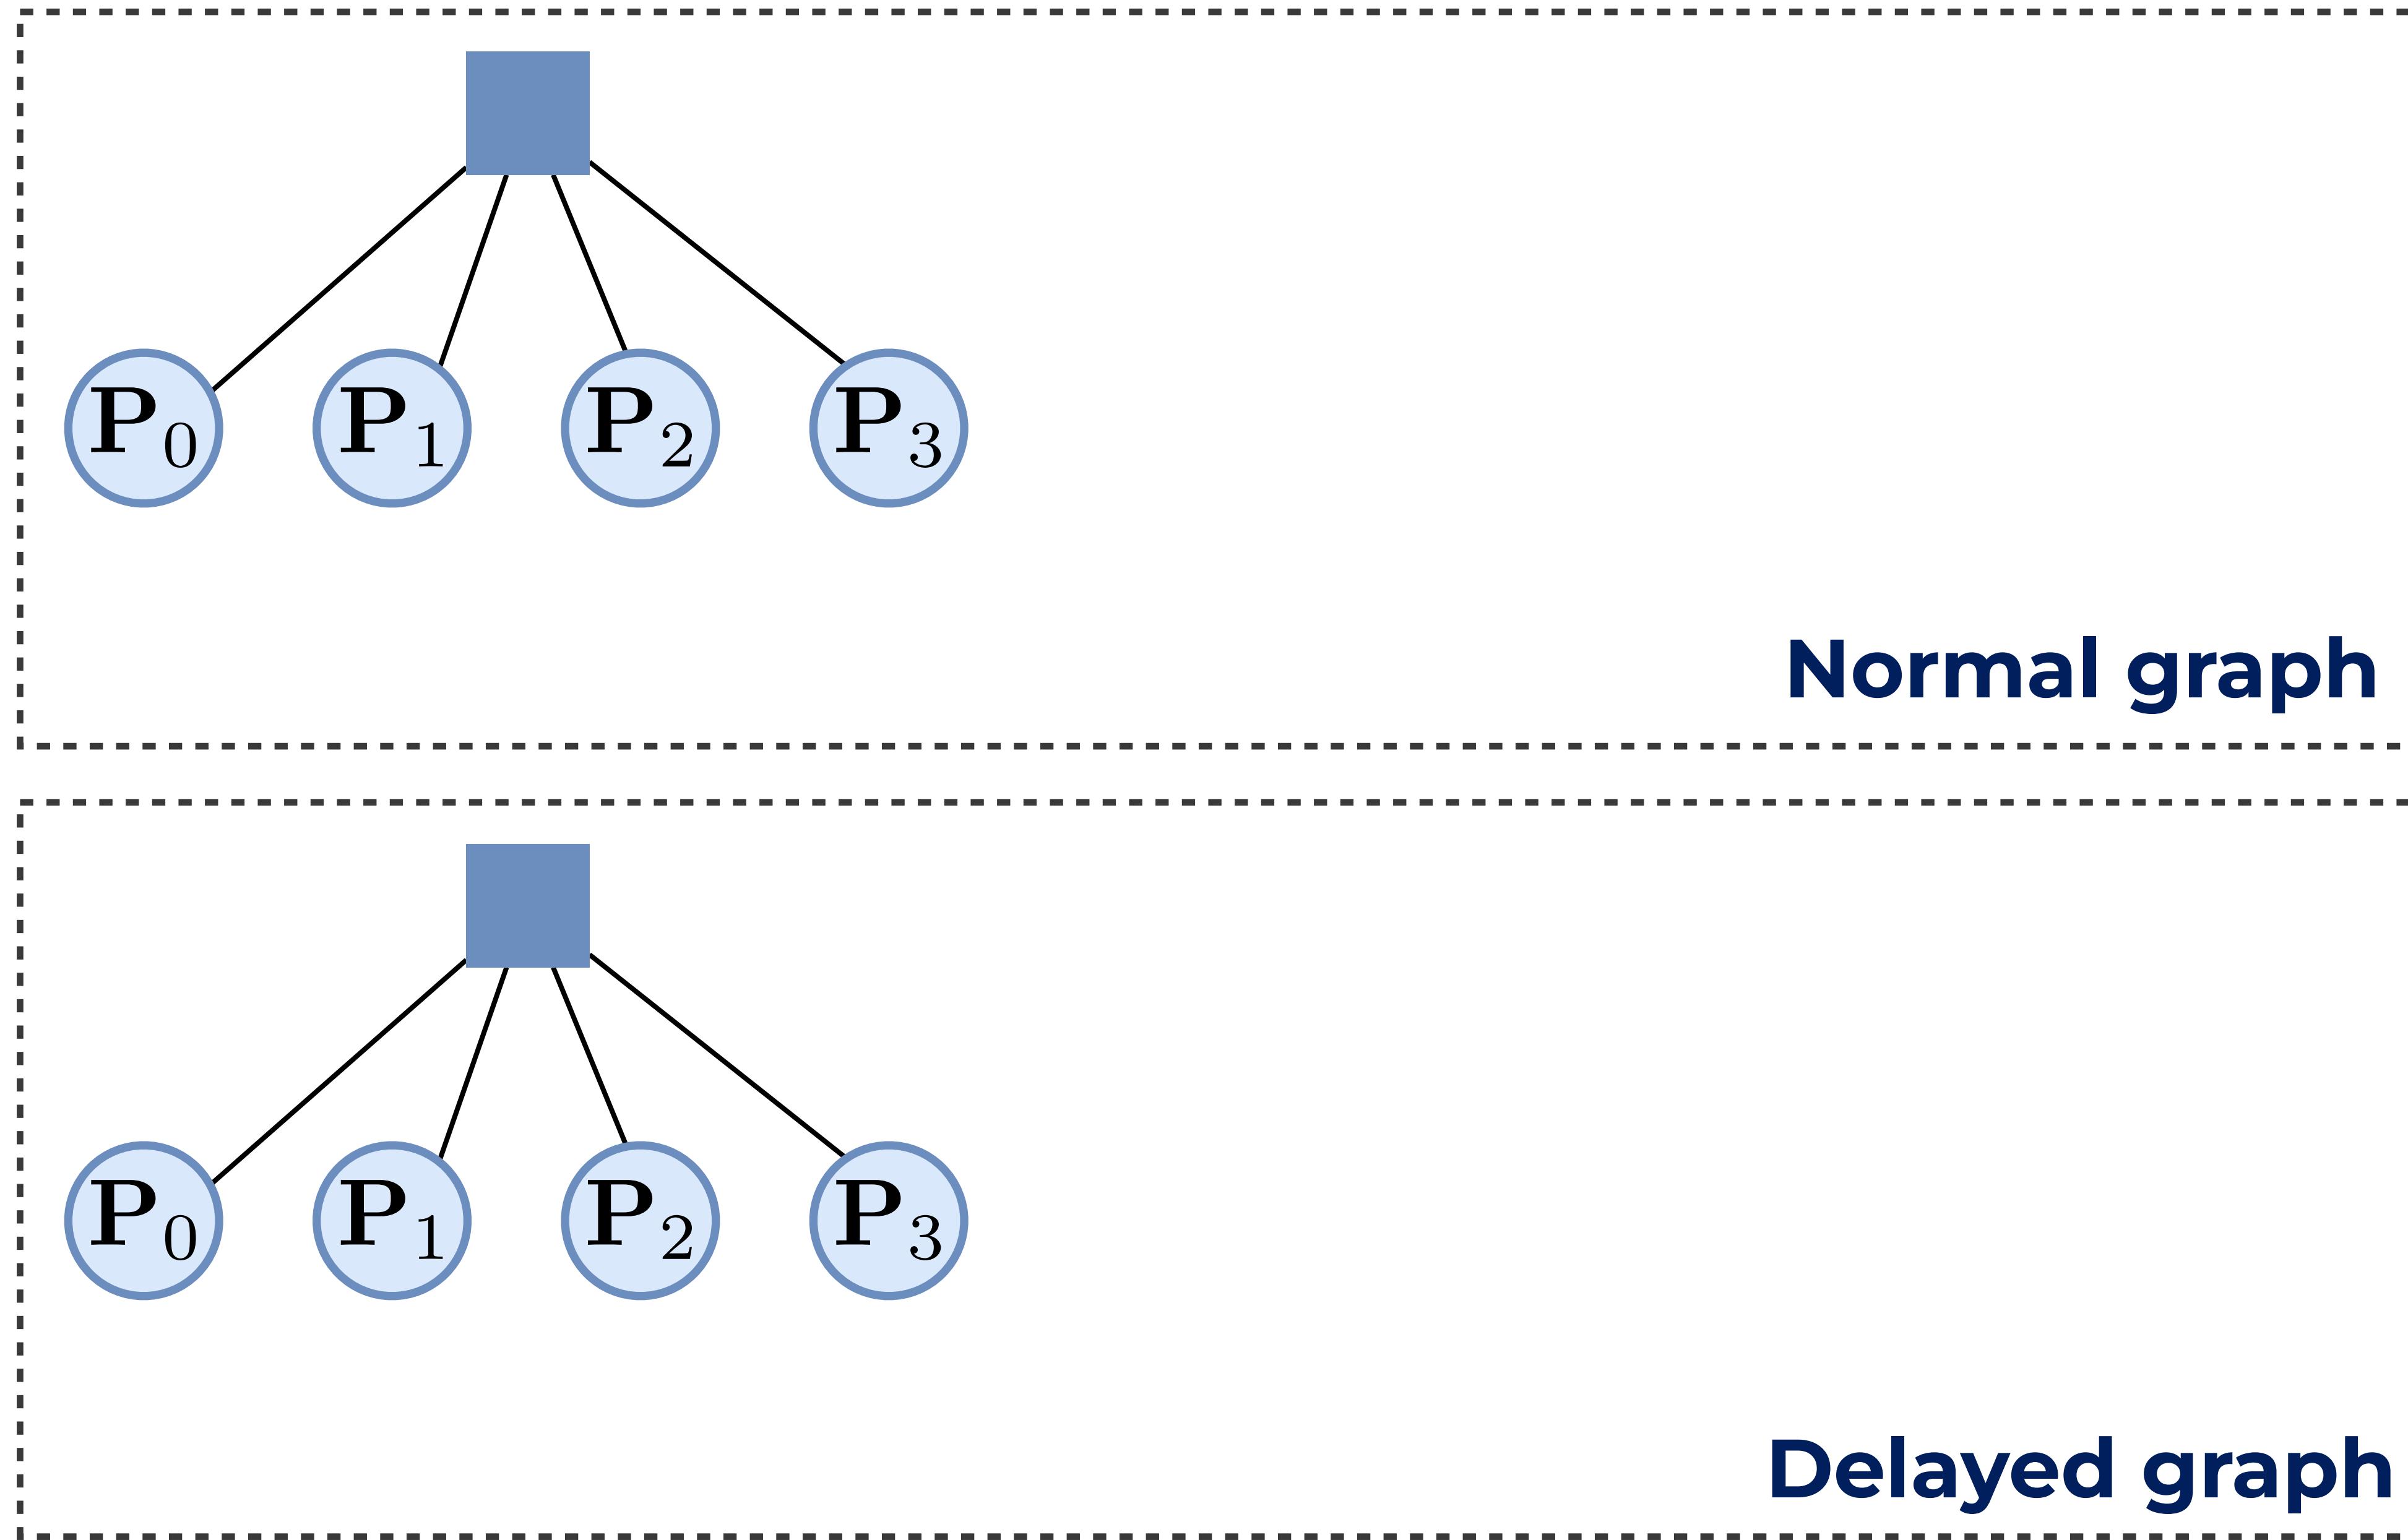


Marginalization order: $P_1 \rightarrow P_2 \rightarrow P_3 \rightarrow P_4 \rightarrow P_0$

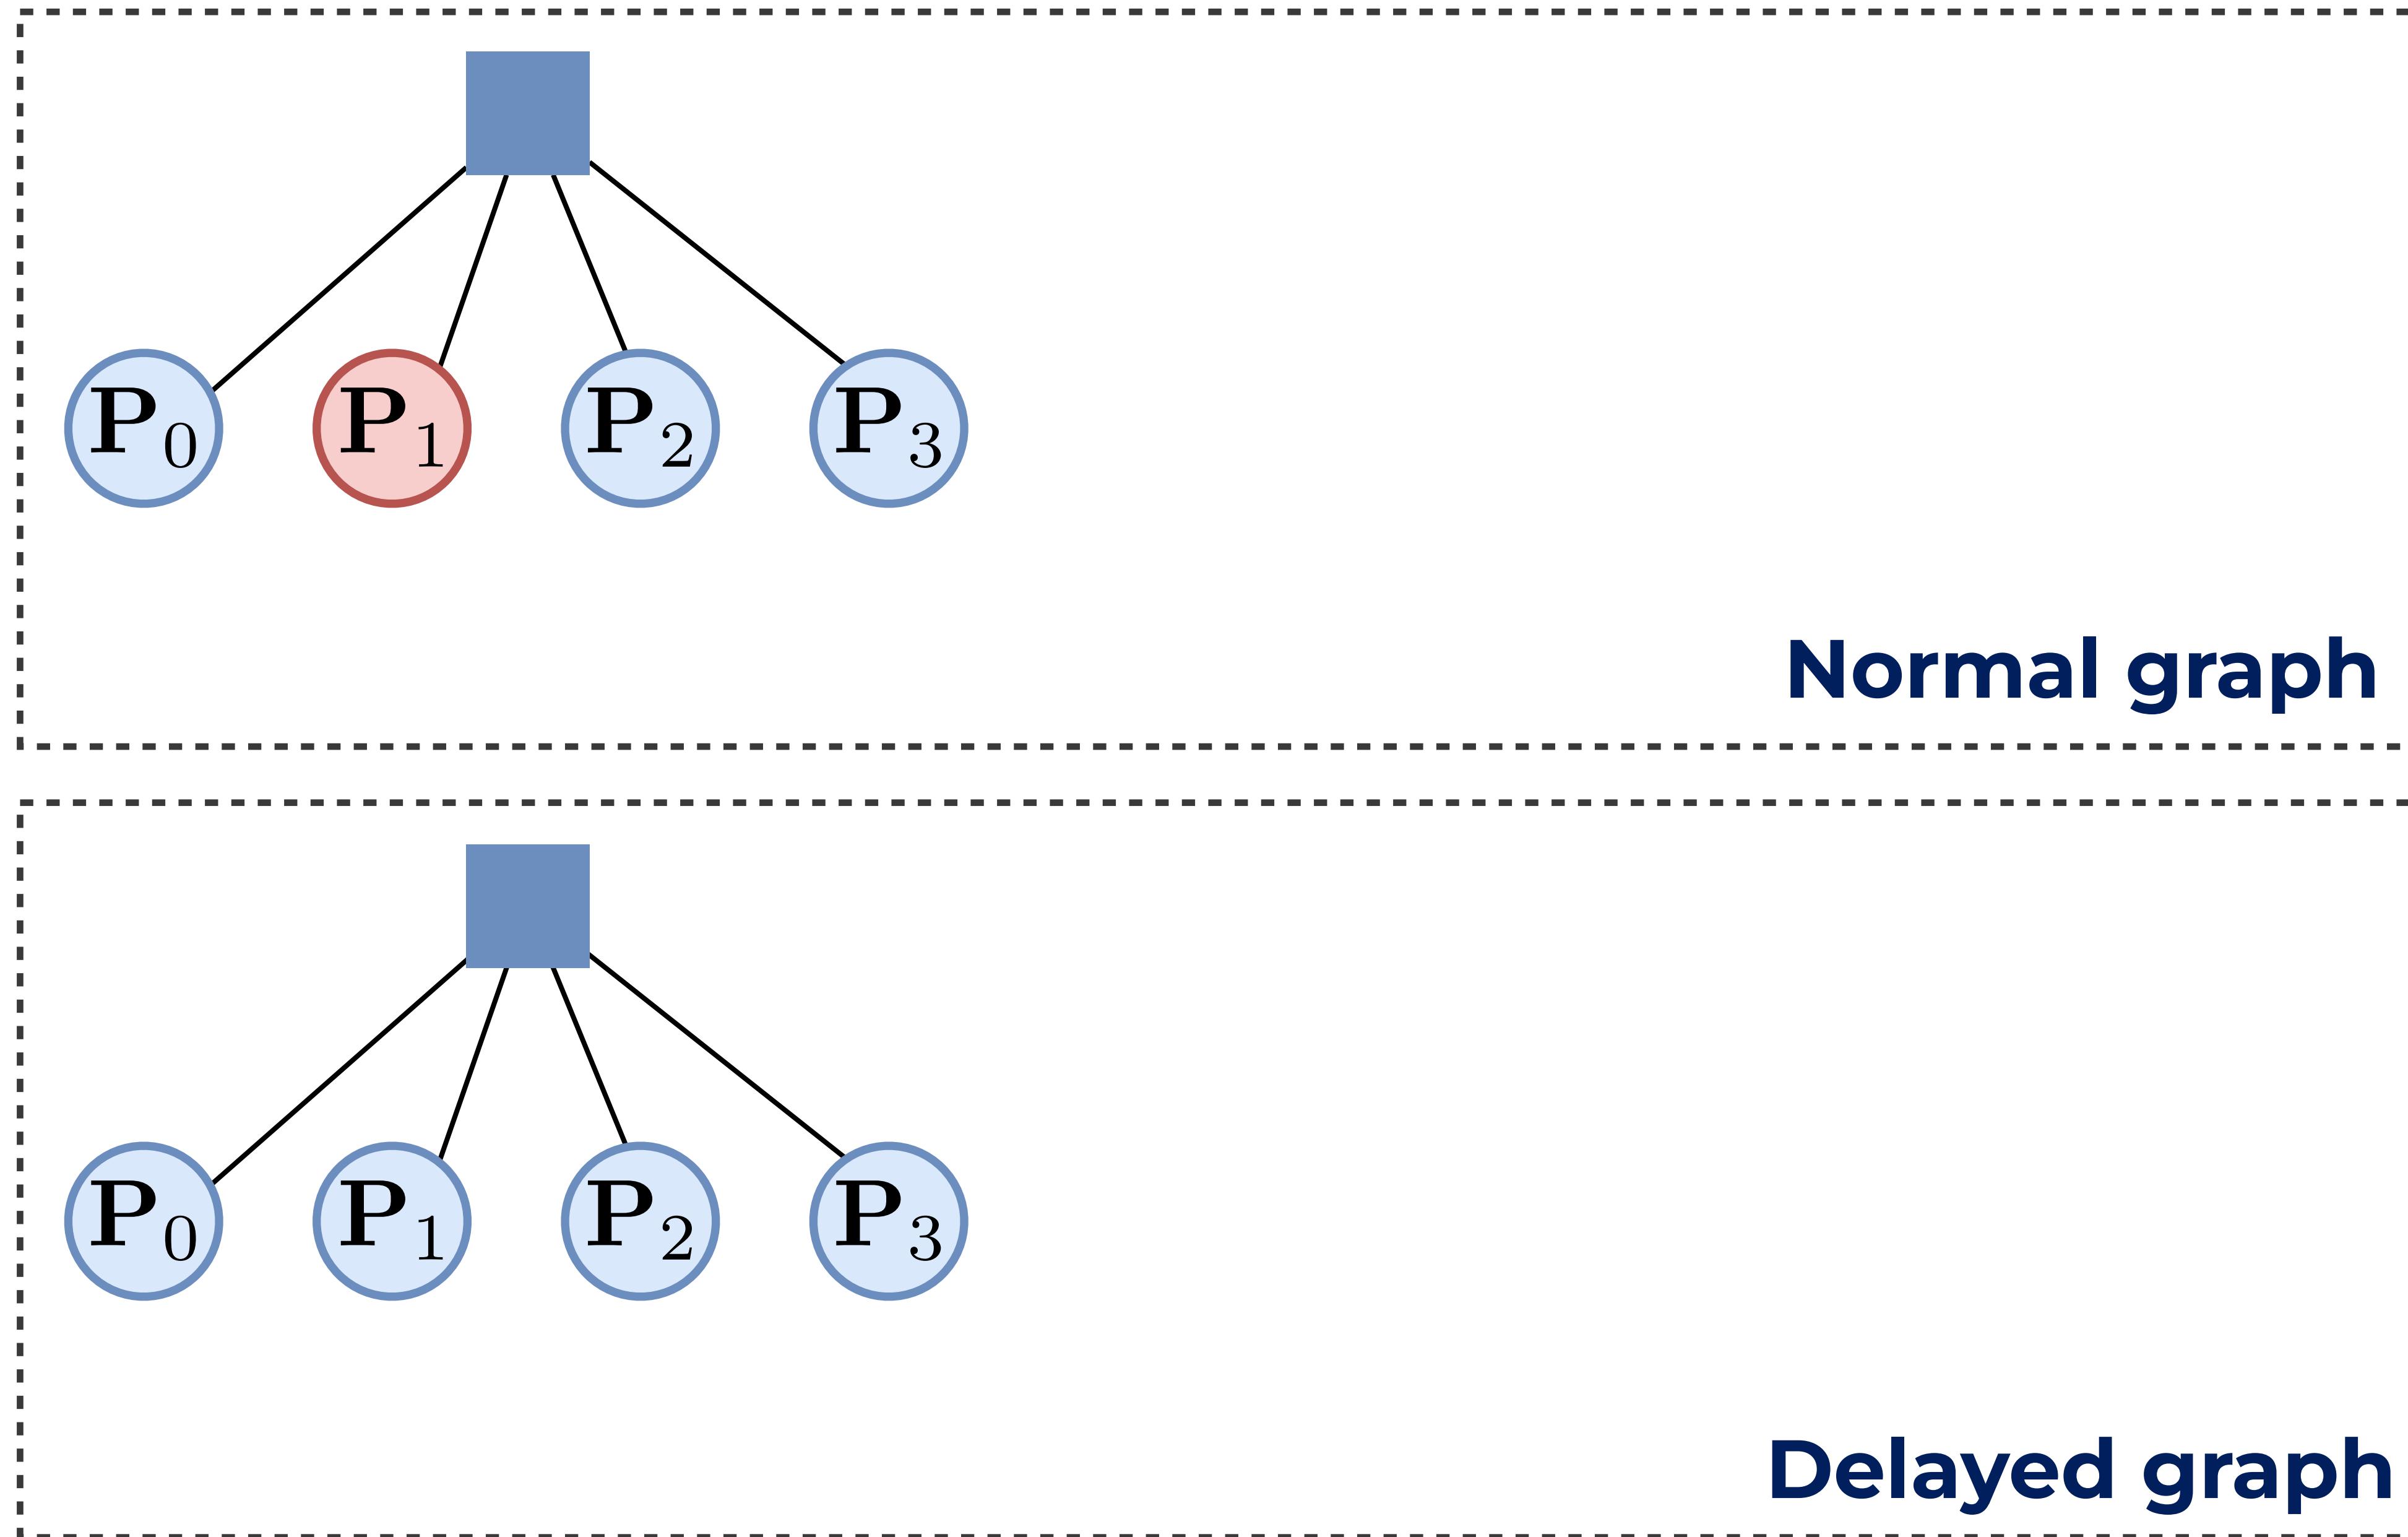
Marginalization order: $P_1 \rightarrow P_2 \rightarrow P_3 \rightarrow P_4 \rightarrow P_0$



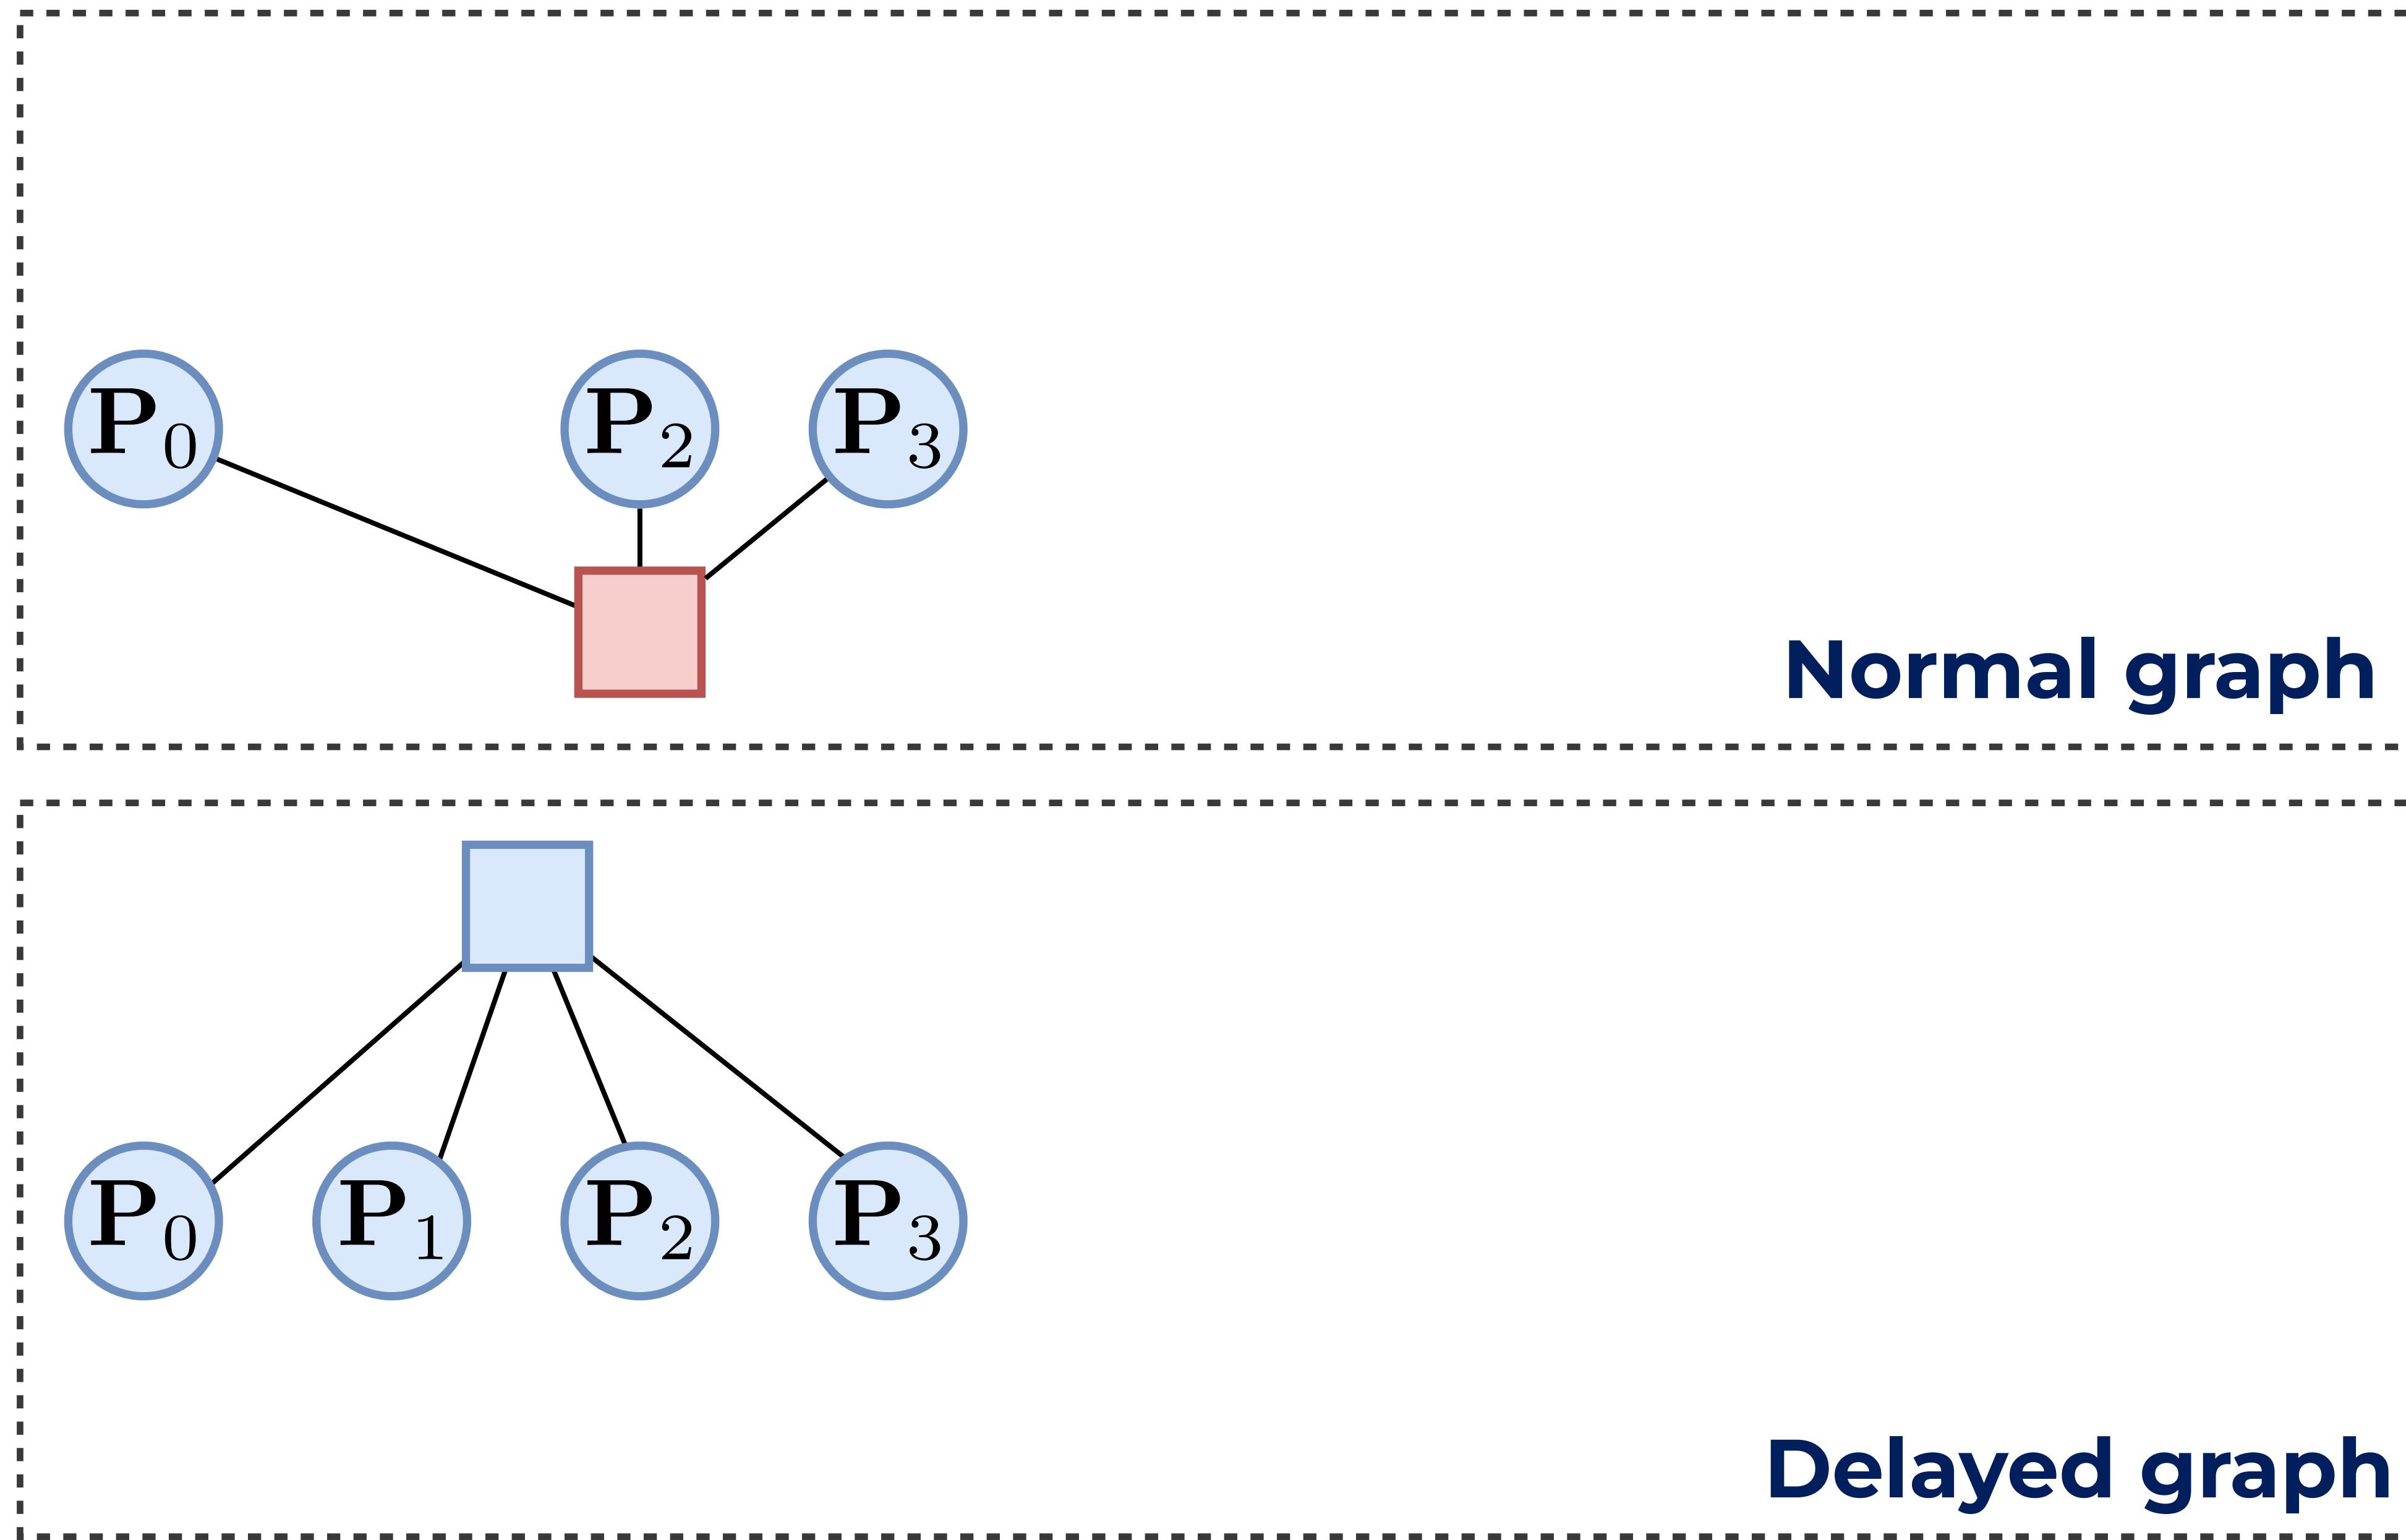
Marginalization order: $P_1 \rightarrow P_2 \rightarrow P_3 \rightarrow P_4 \rightarrow P_0$



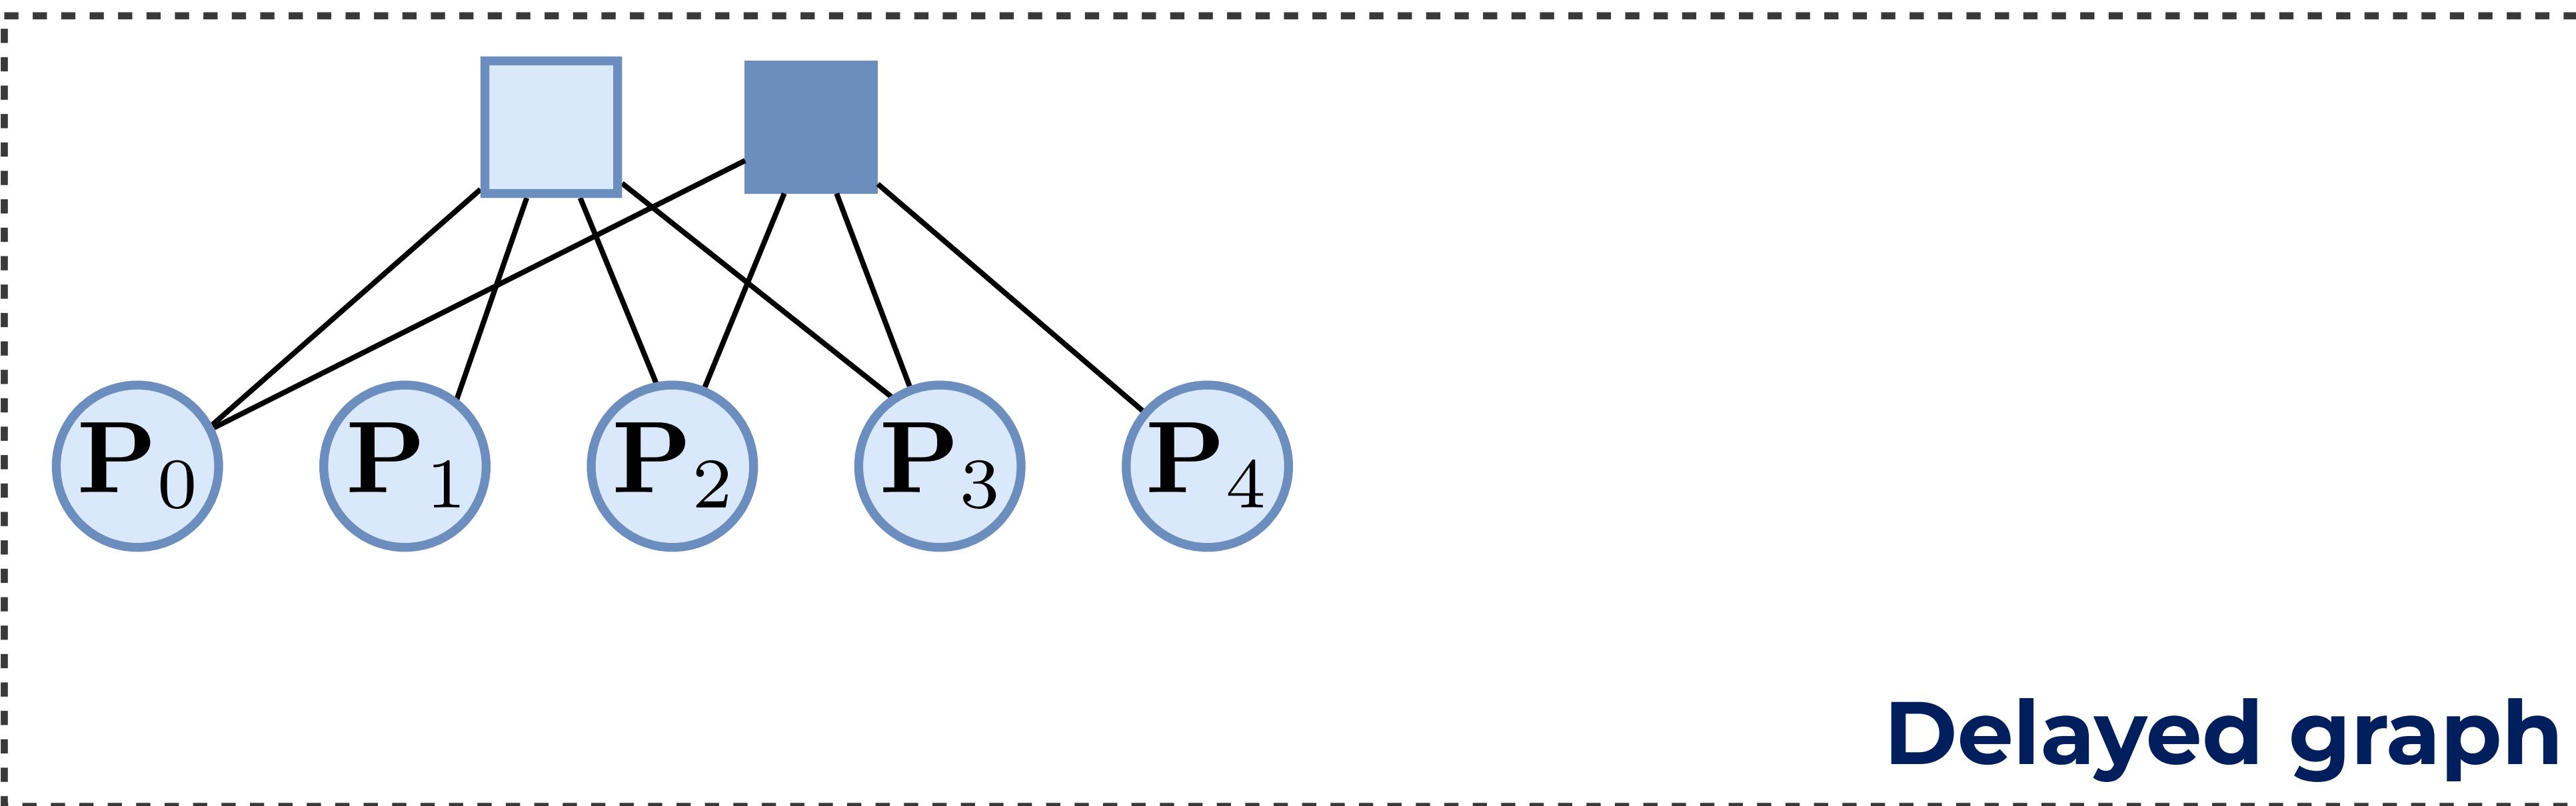
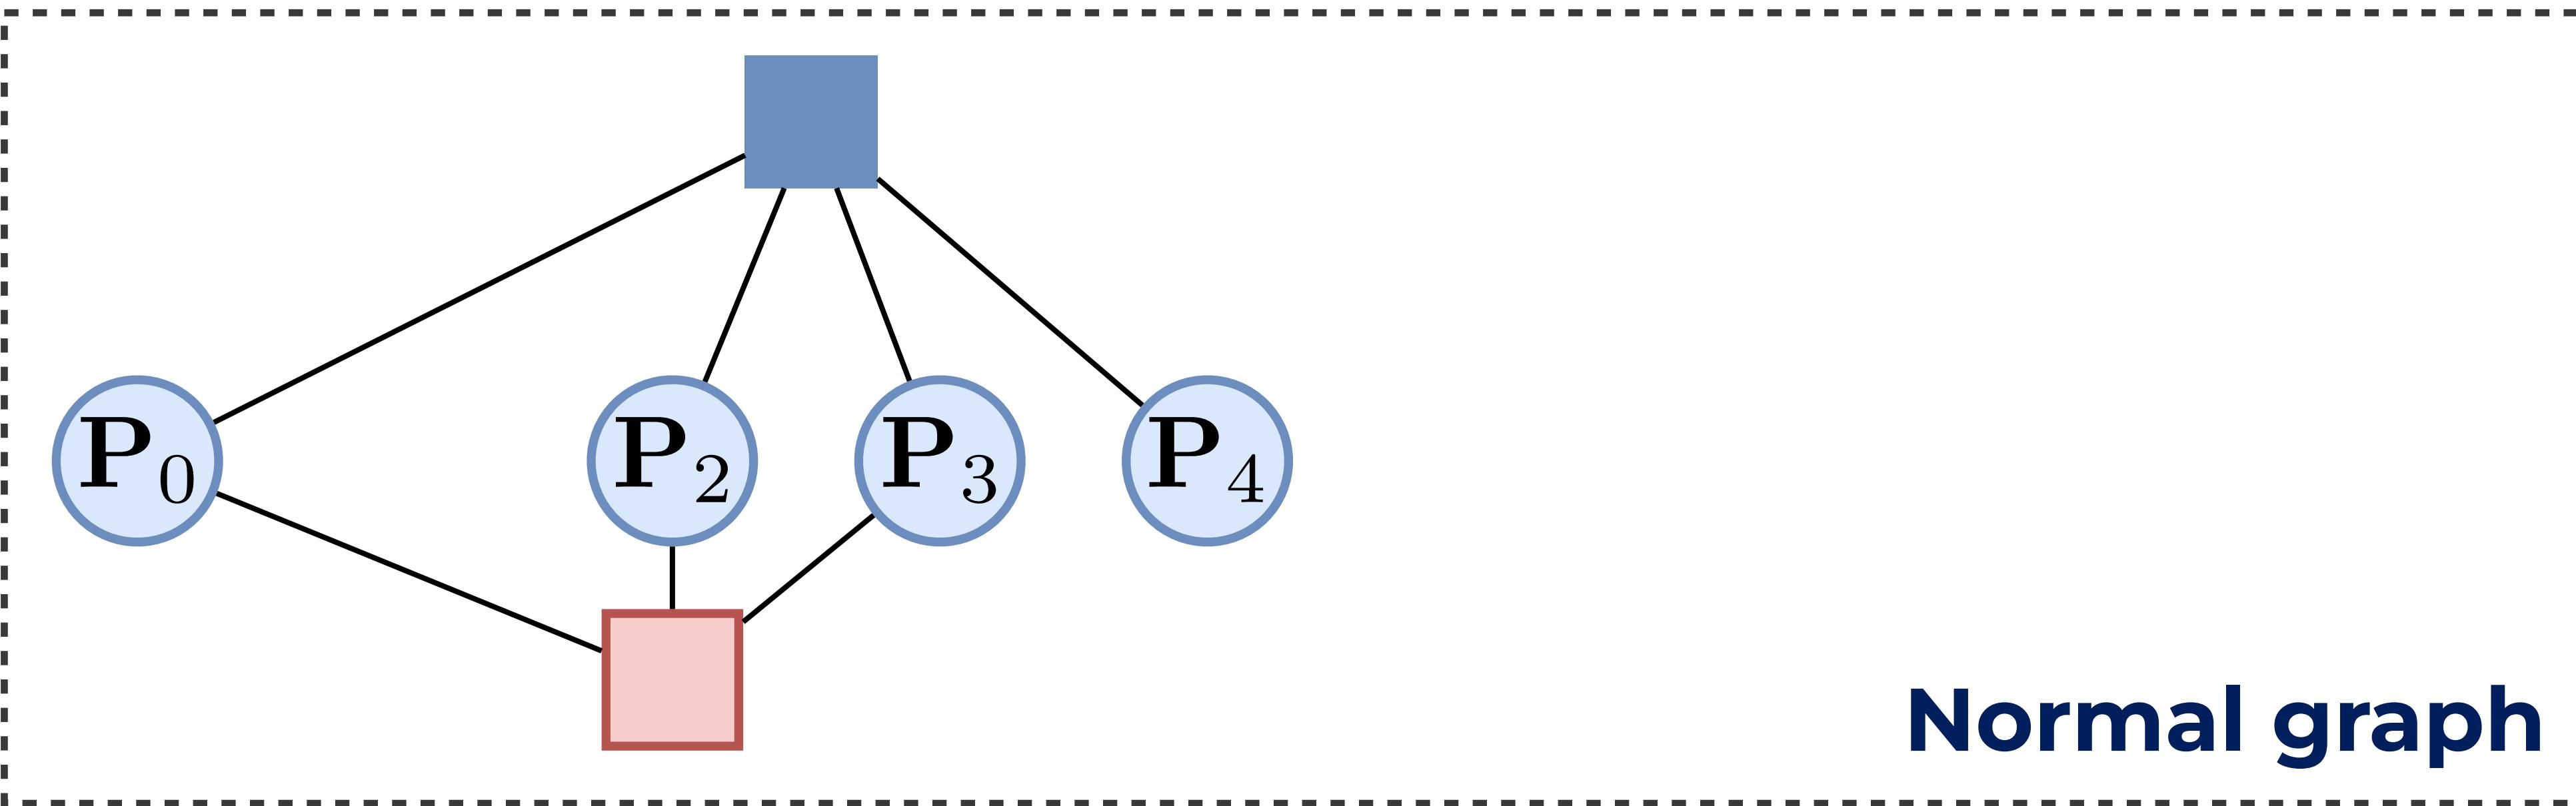
Marginalization order: $P_1 \rightarrow P_2 \rightarrow P_3 \rightarrow P_4 \rightarrow P_0$



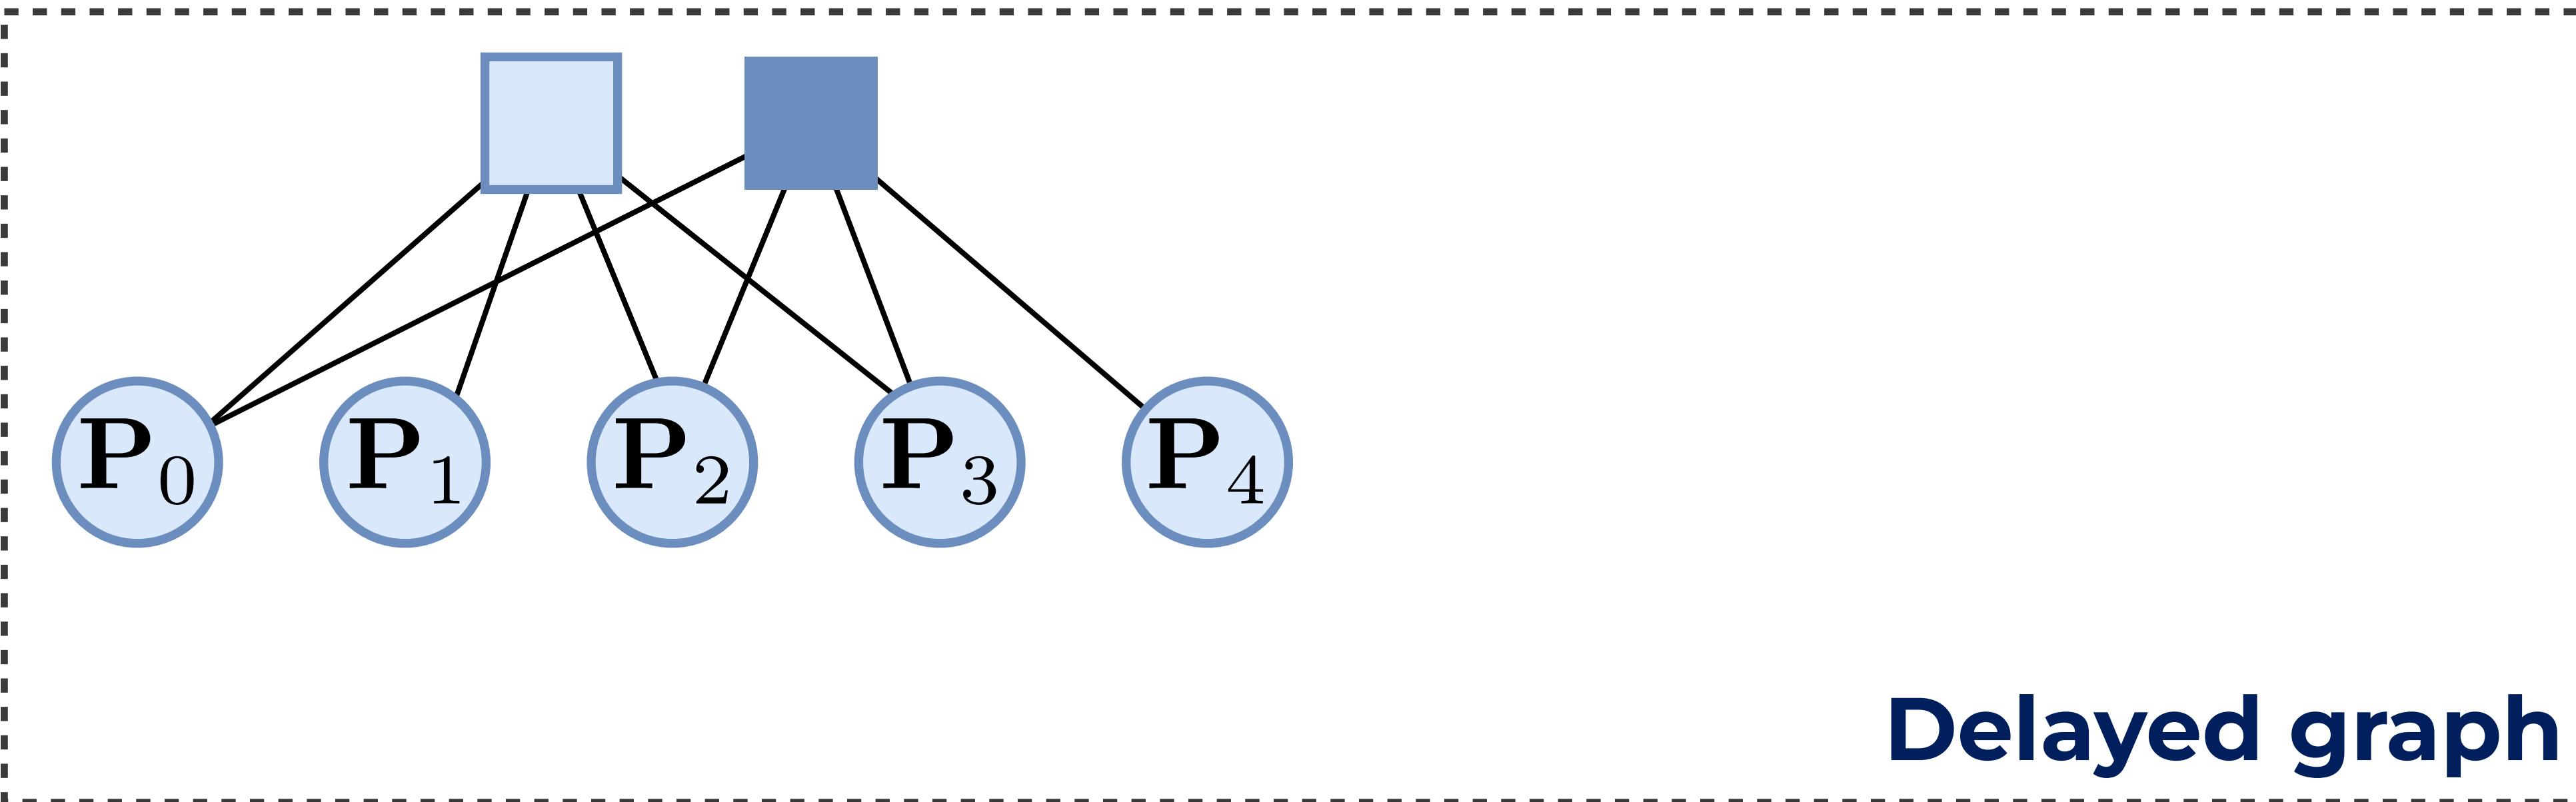
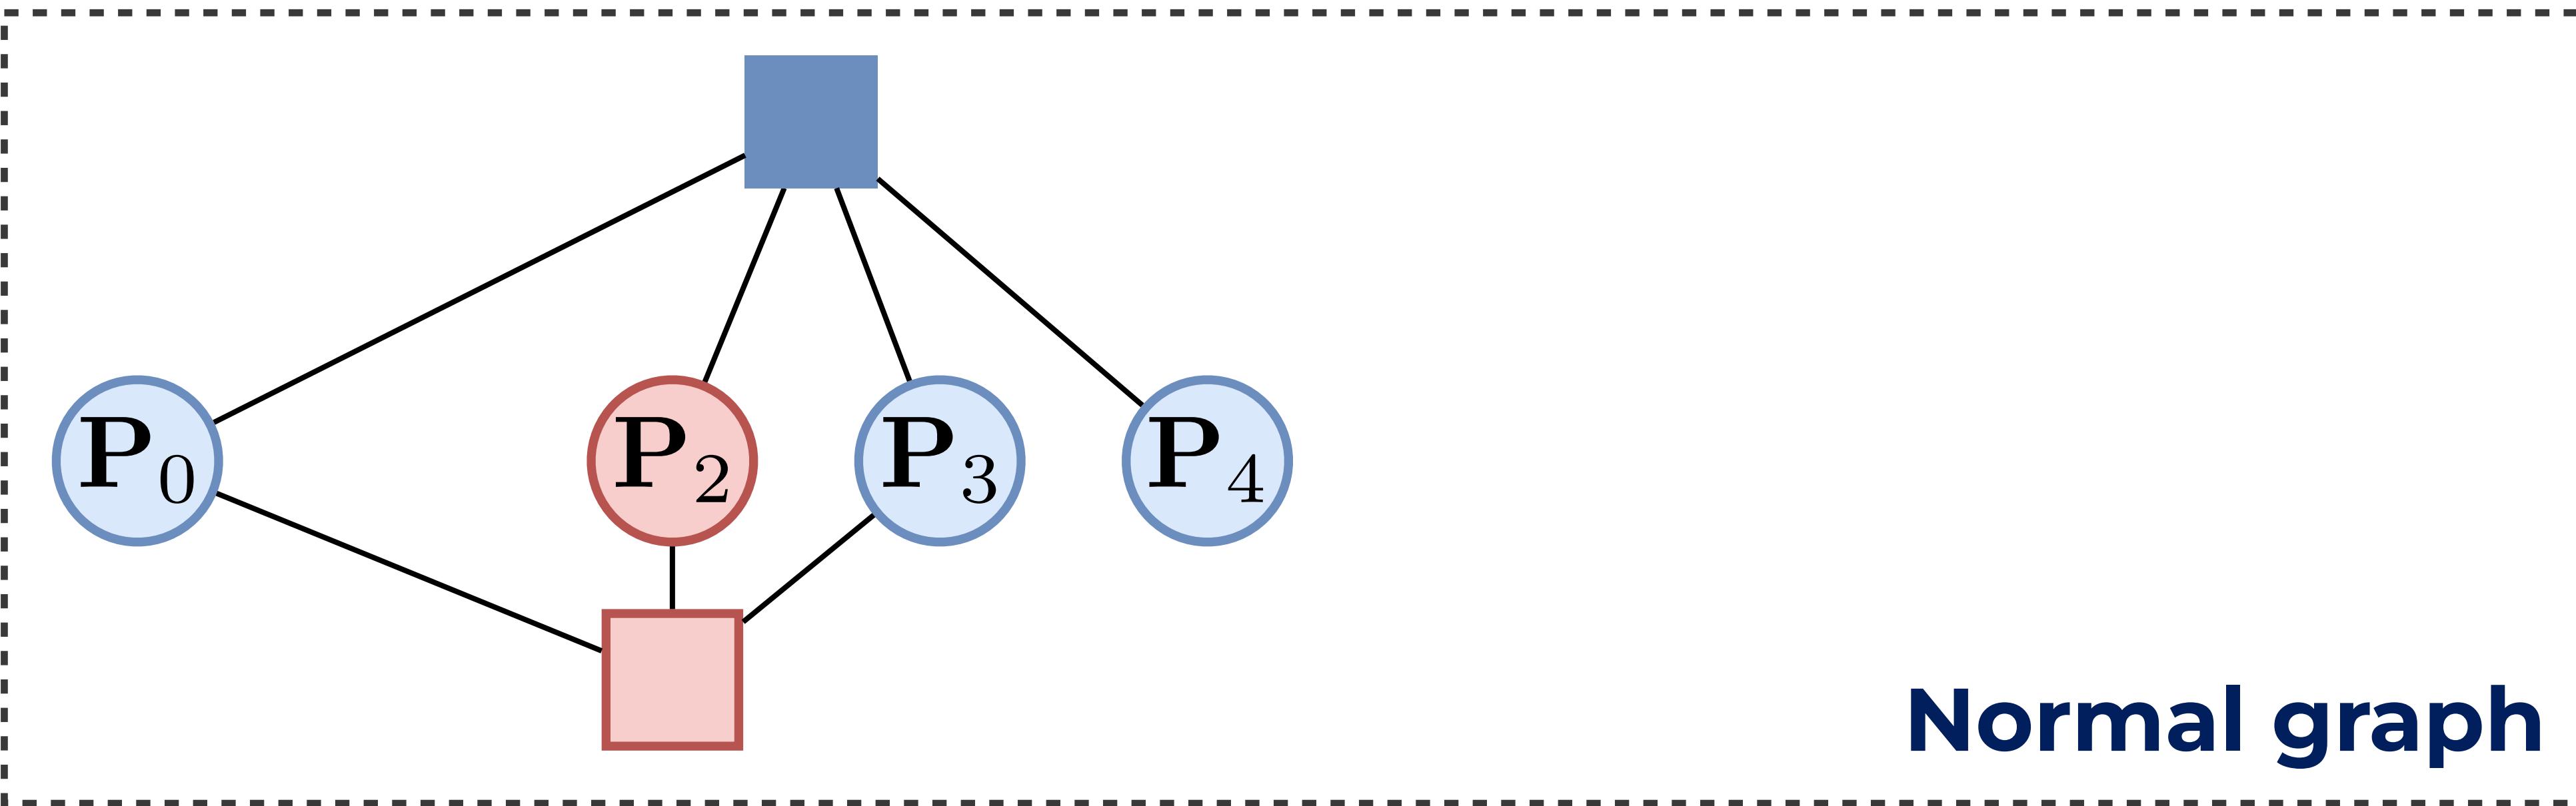
Marginalization order: $P_1 \rightarrow P_2 \rightarrow P_3 \rightarrow P_4 \rightarrow P_0$



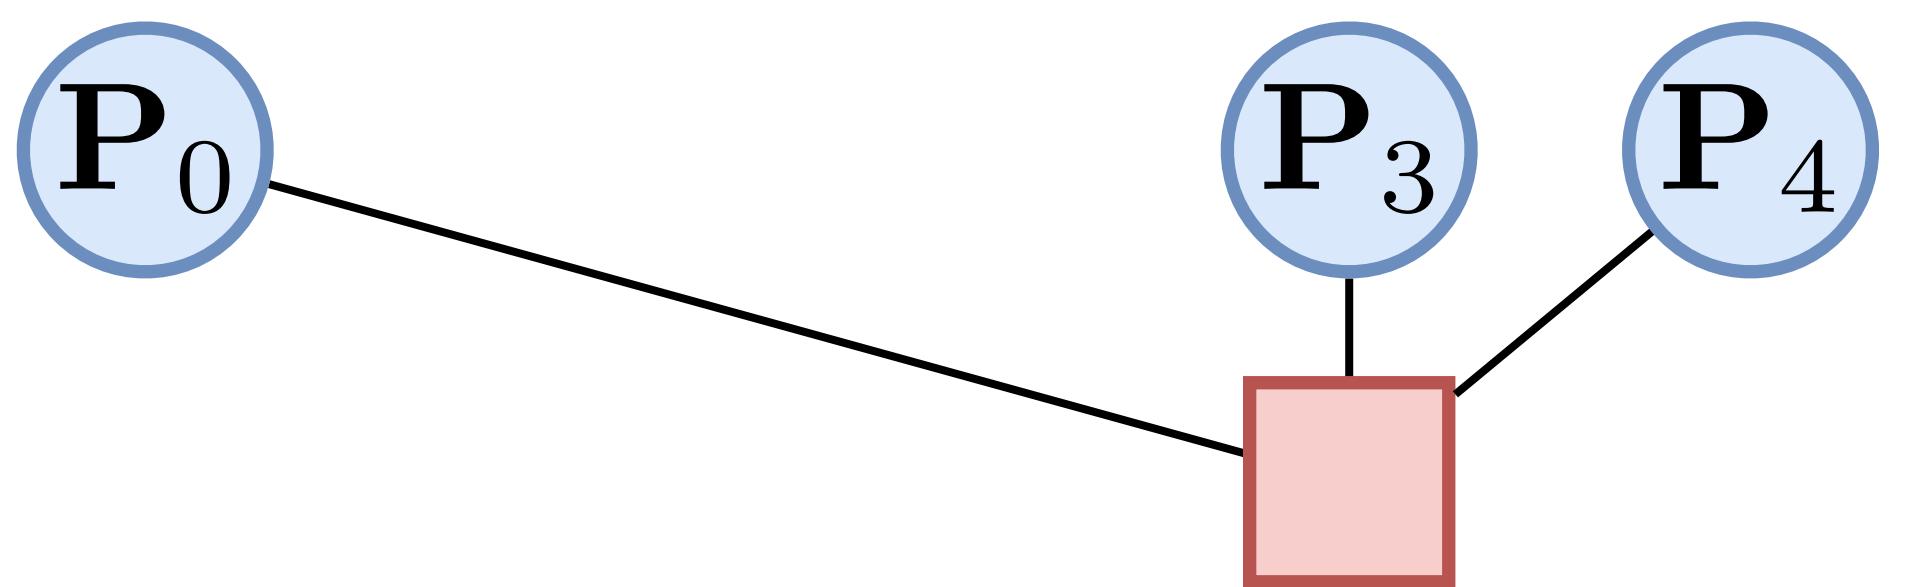
Marginalization order: $P_1 \rightarrow P_2 \rightarrow P_3 \rightarrow P_4 \rightarrow P_0$



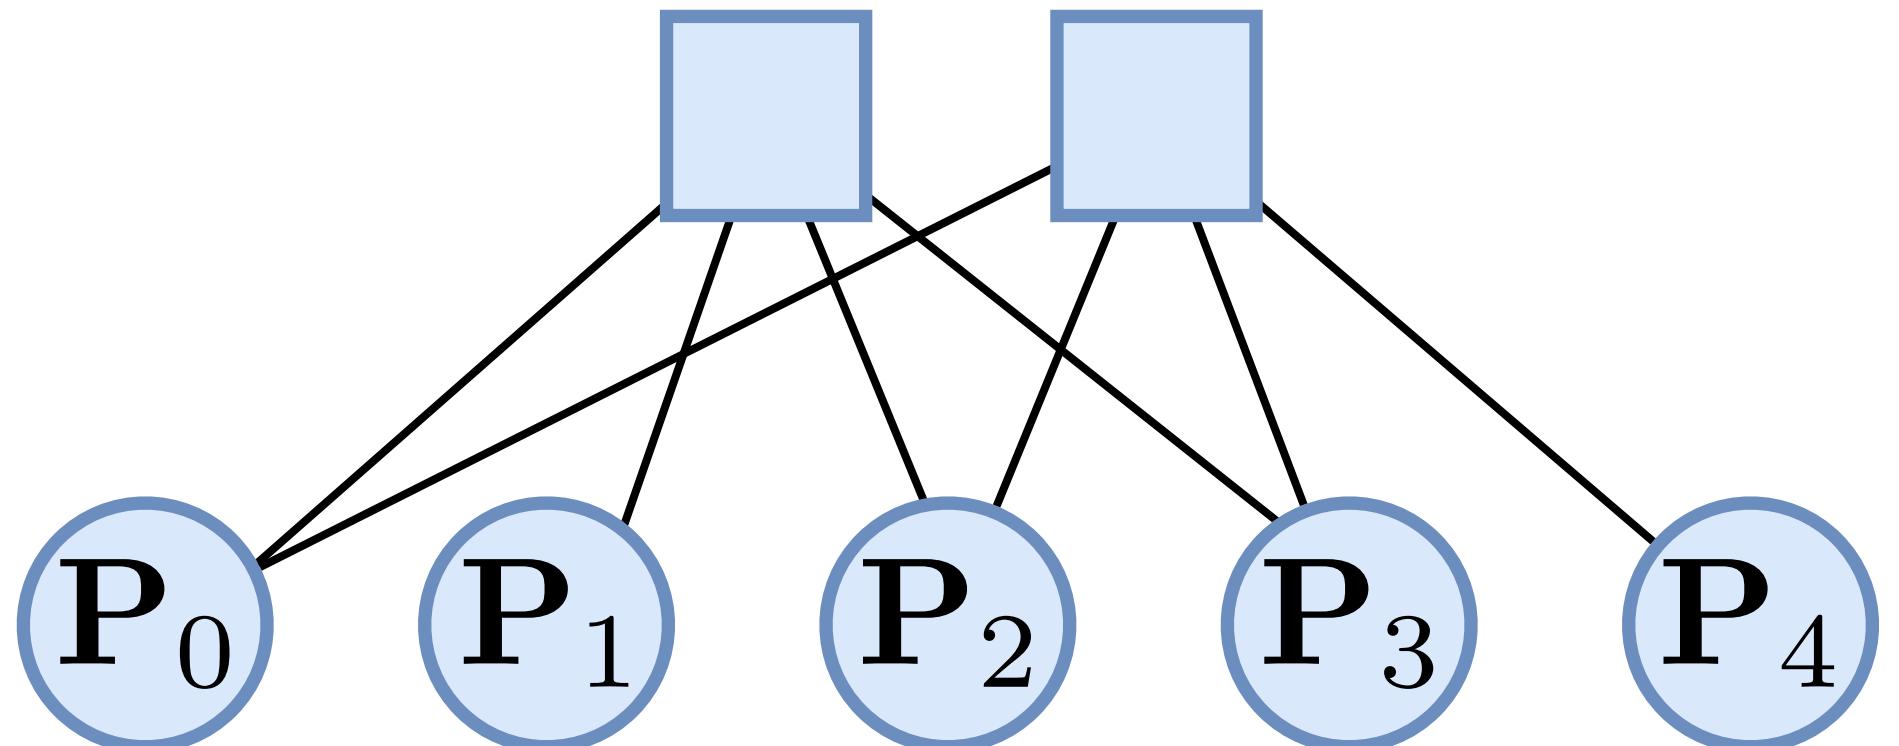
Marginalization order: $P_1 \rightarrow P_2 \rightarrow P_3 \rightarrow P_4 \rightarrow P_0$



Marginalization order: $P_1 \rightarrow P_2 \rightarrow P_3 \rightarrow P_4 \rightarrow P_0$



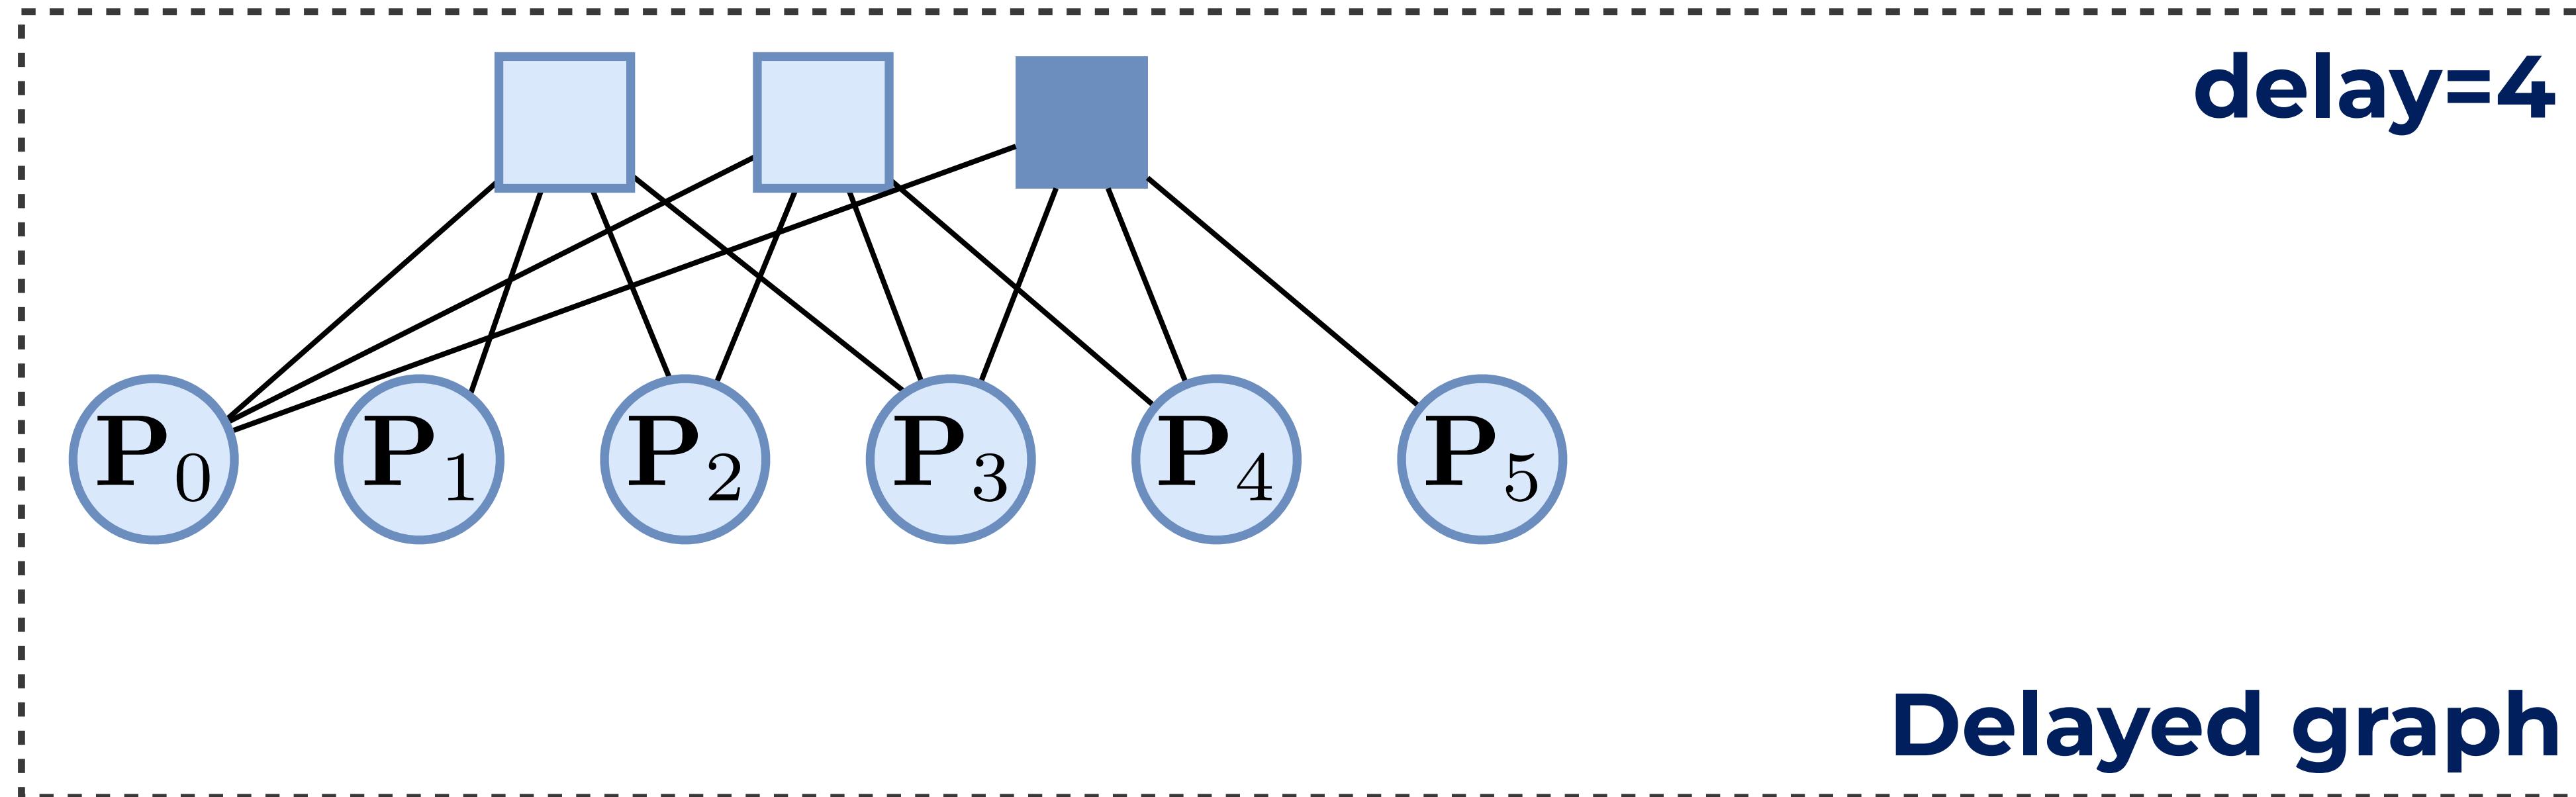
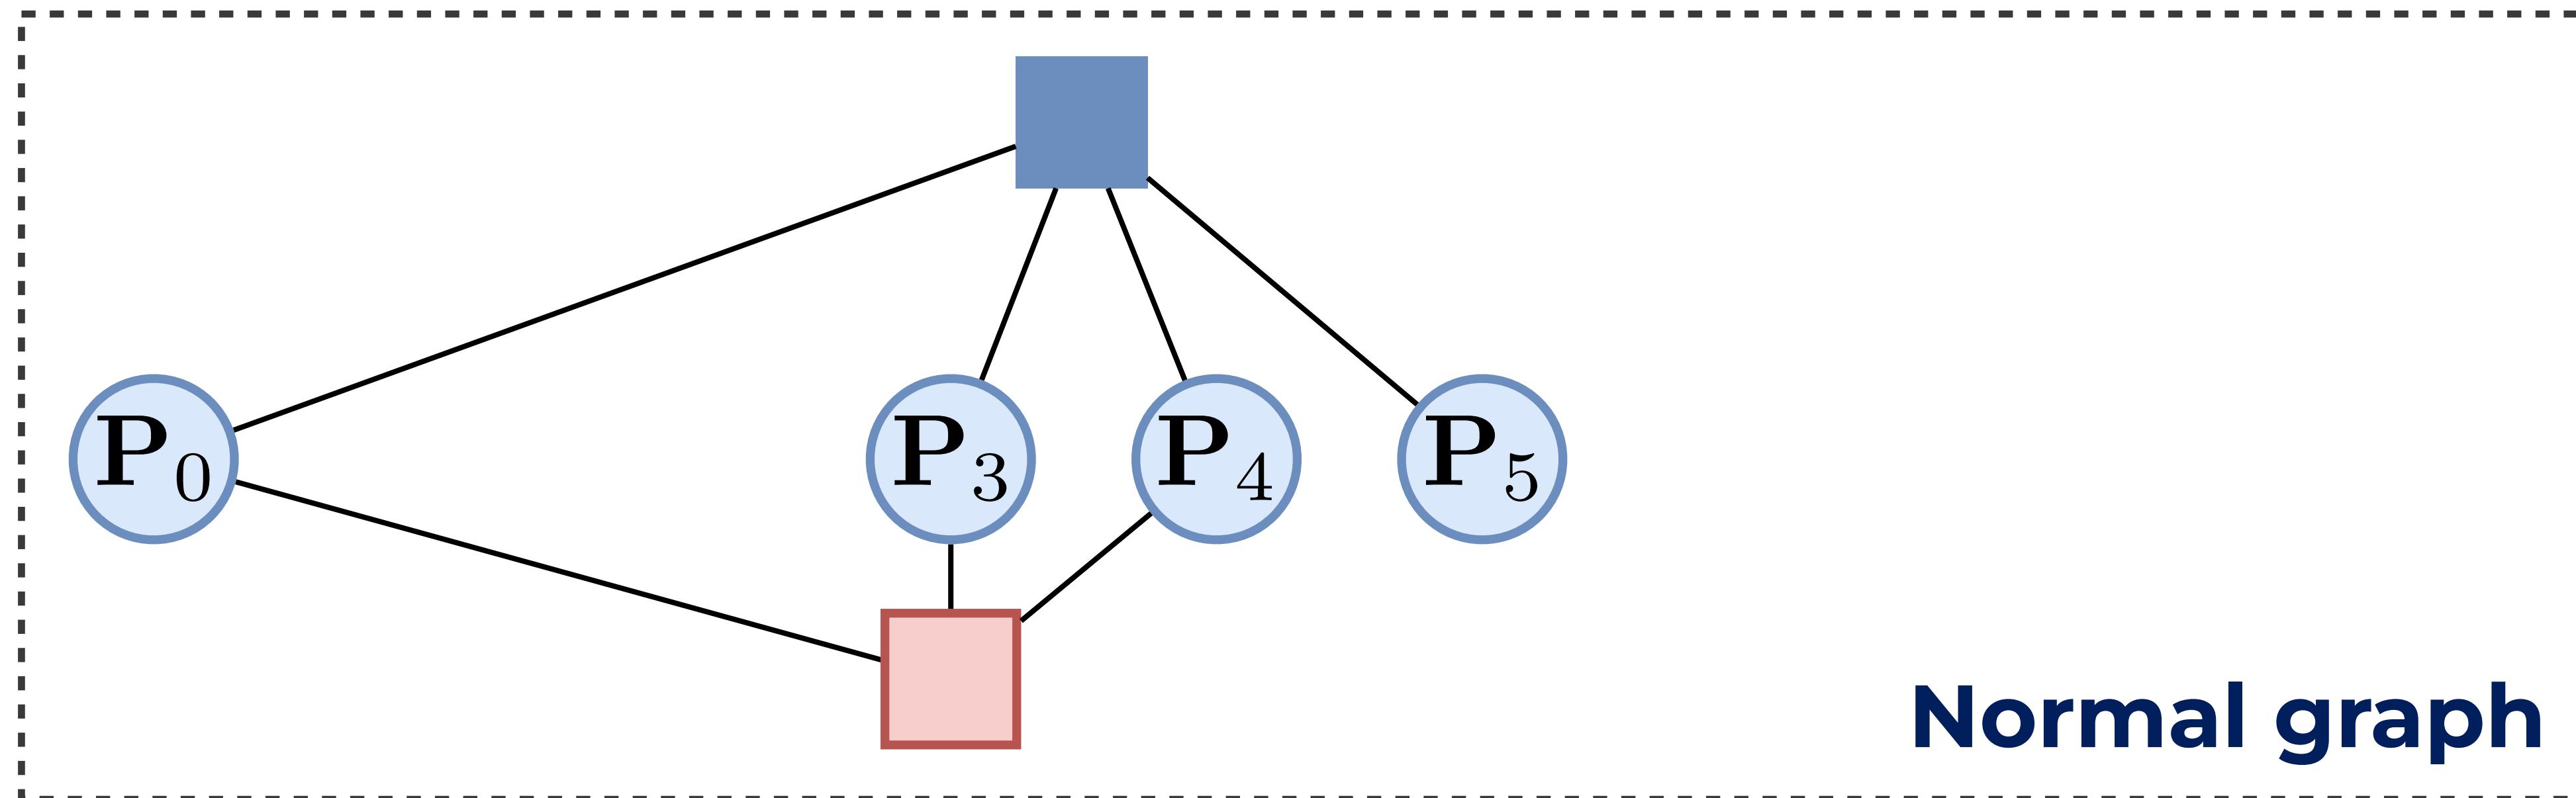
Normal graph



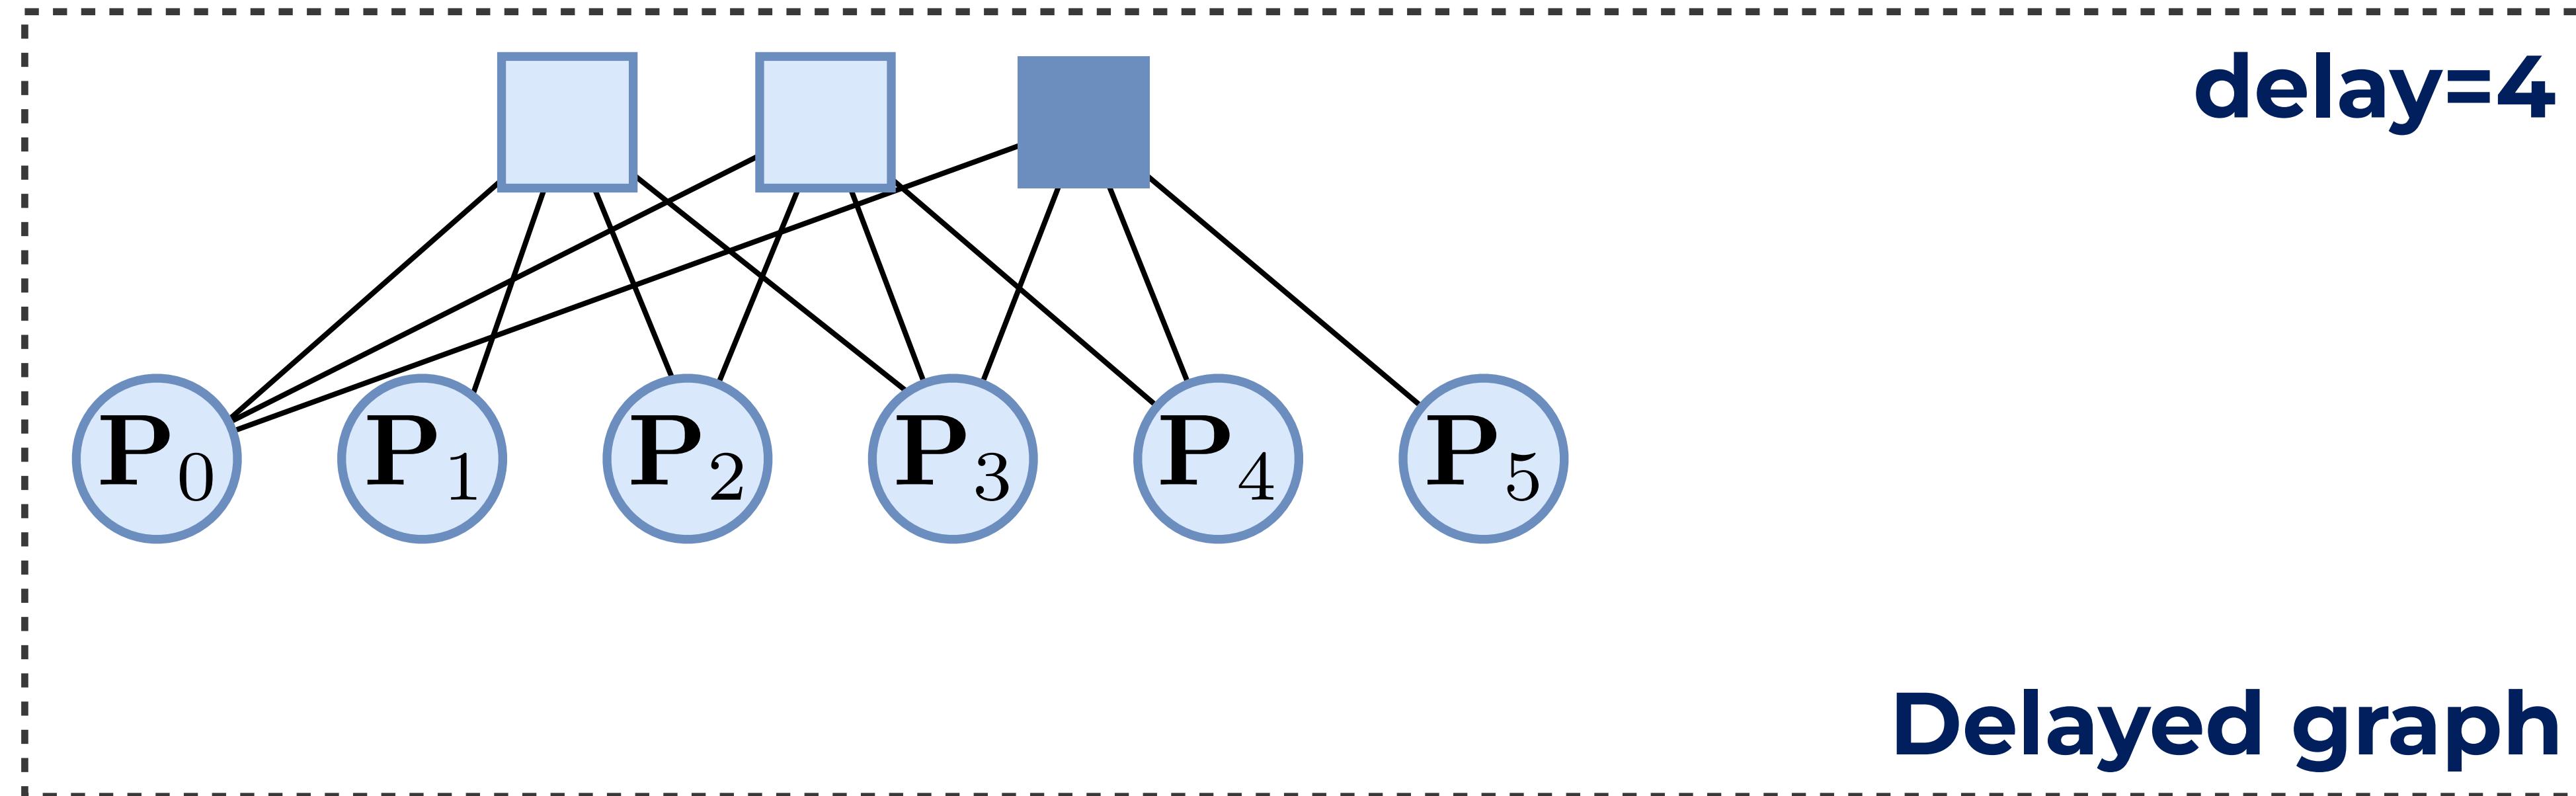
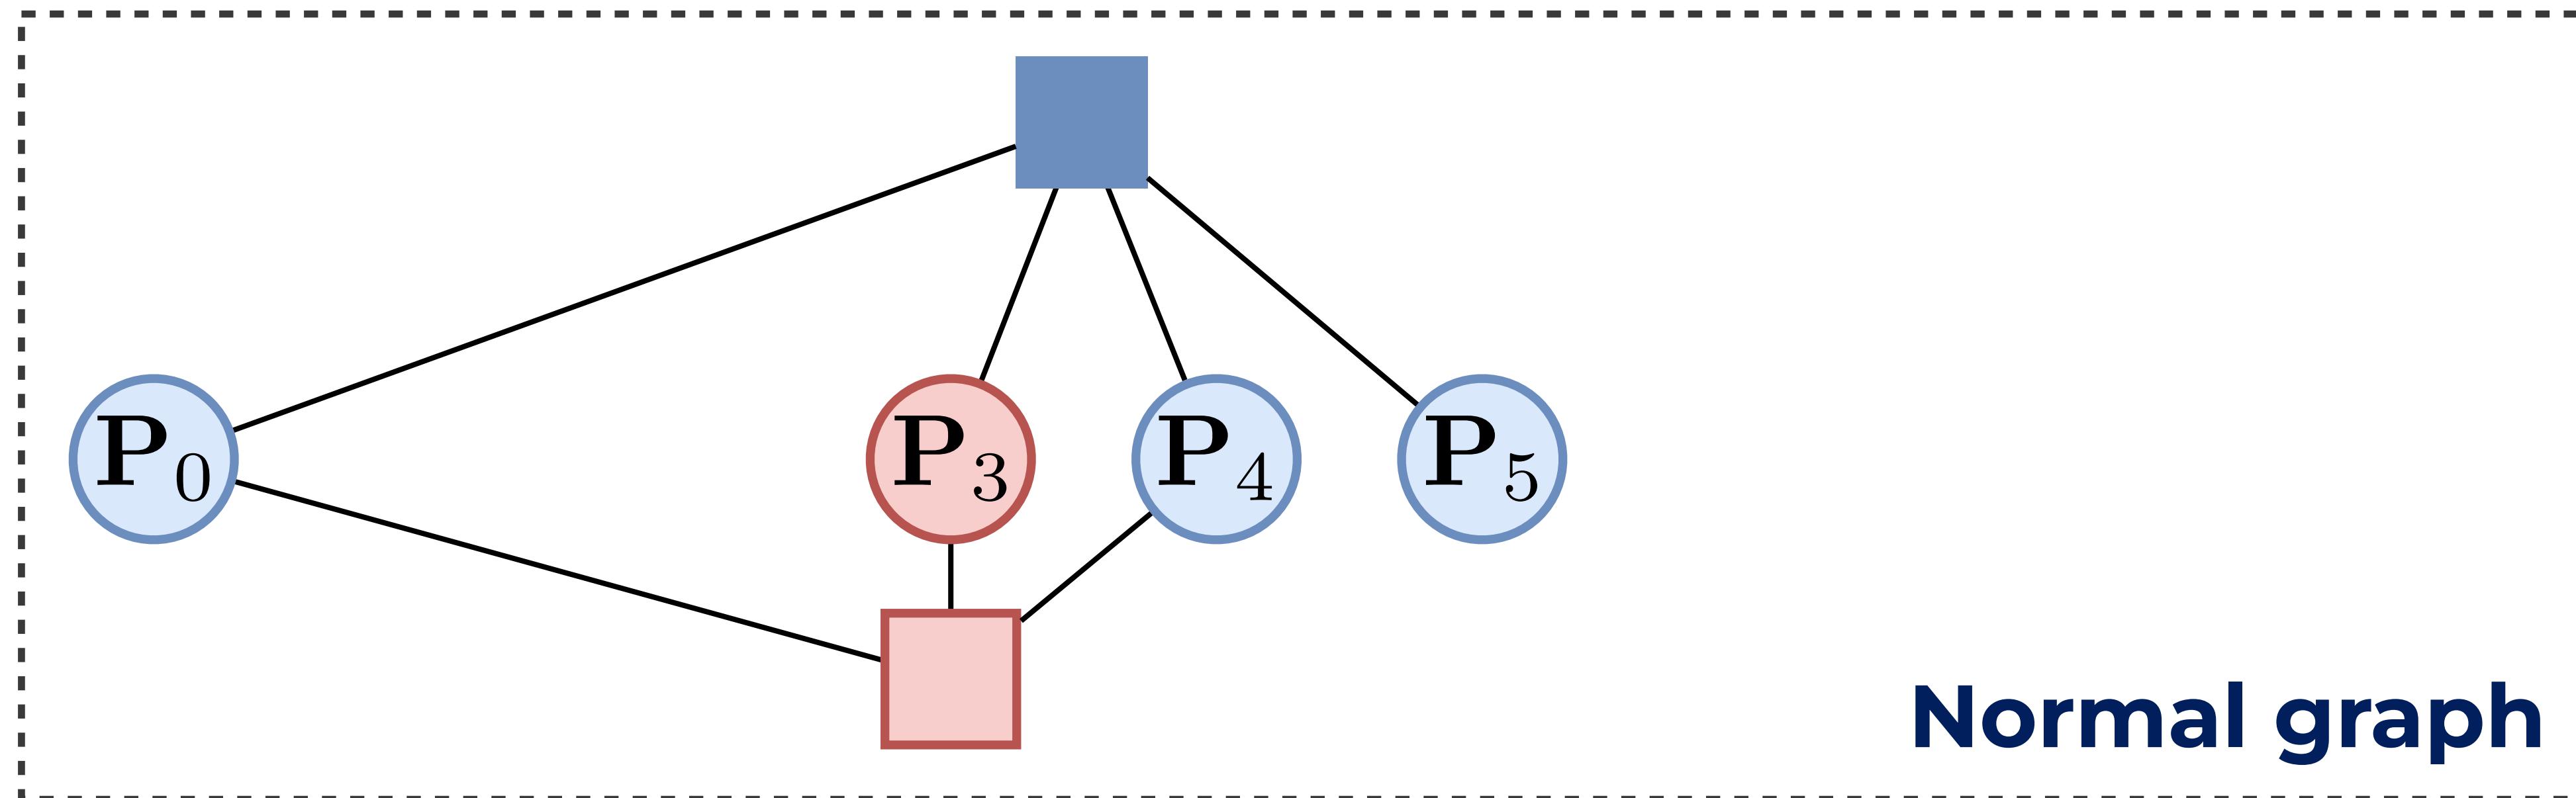
delay=4

Delayed graph

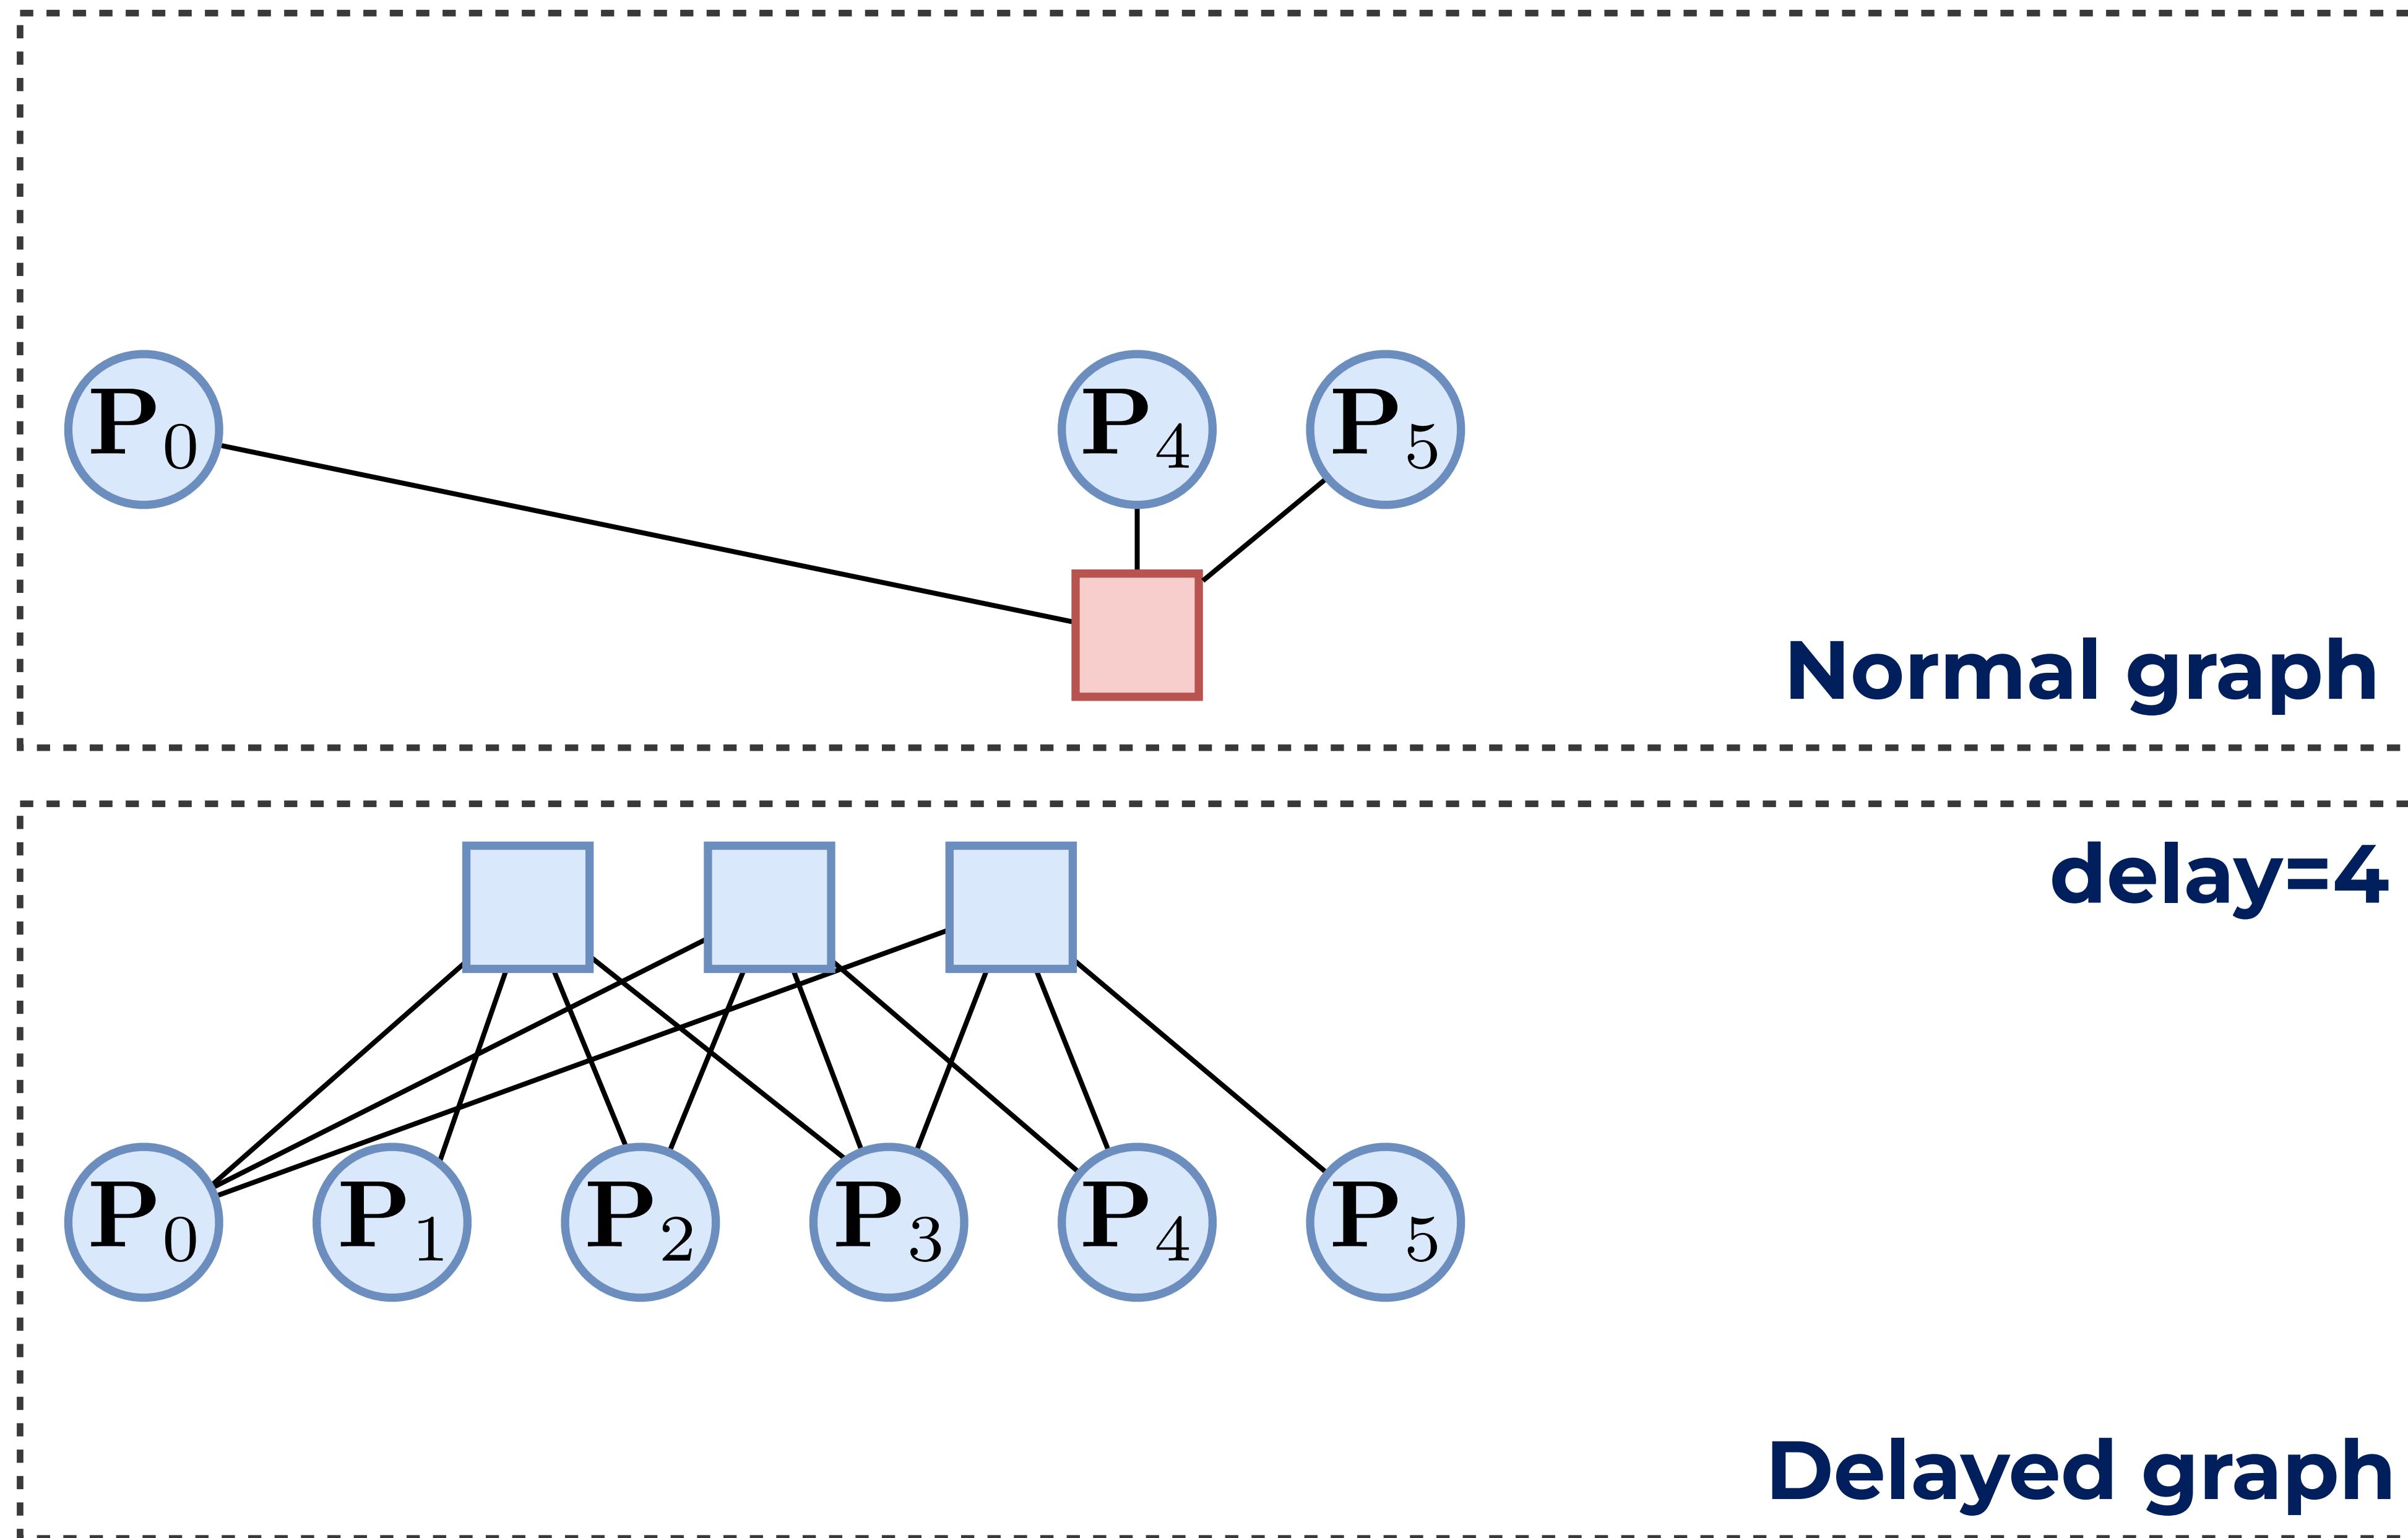
Marginalization order: $P_1 \rightarrow P_2 \rightarrow P_3 \rightarrow P_4 \rightarrow P_0$



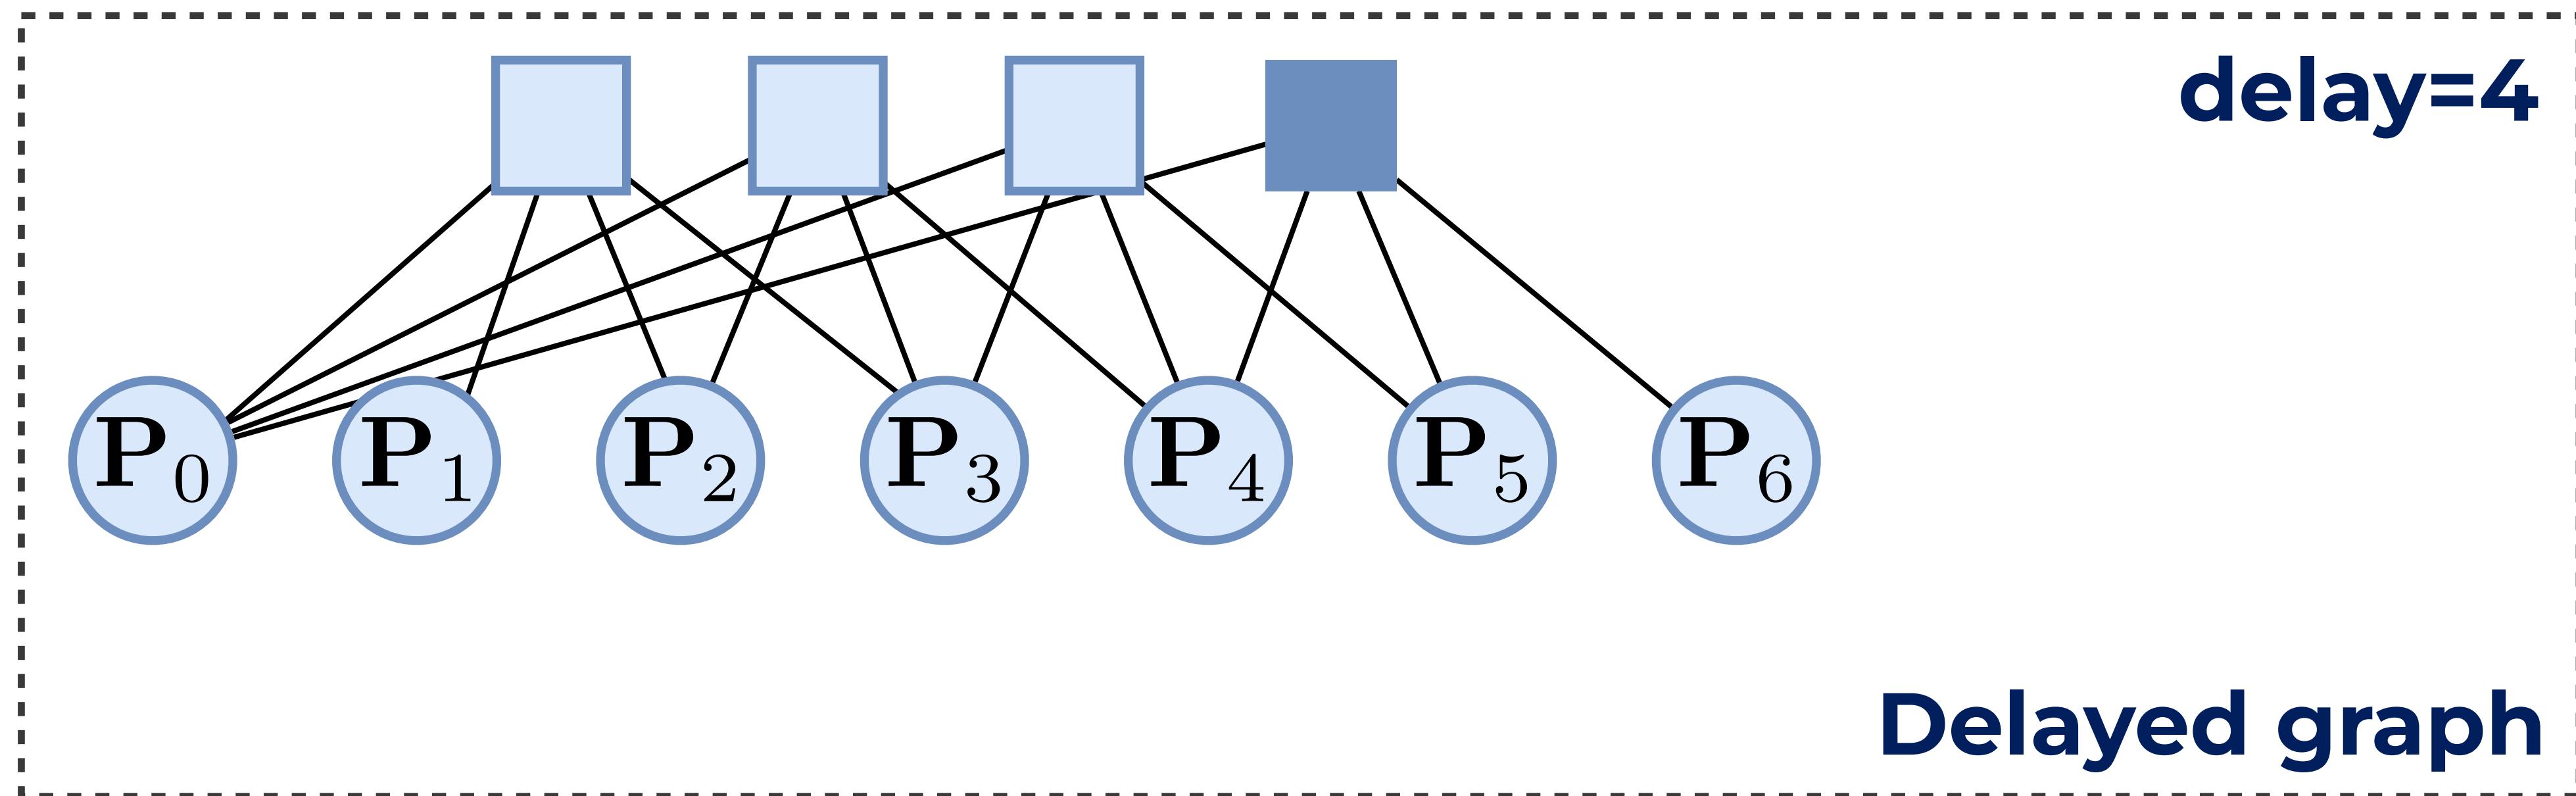
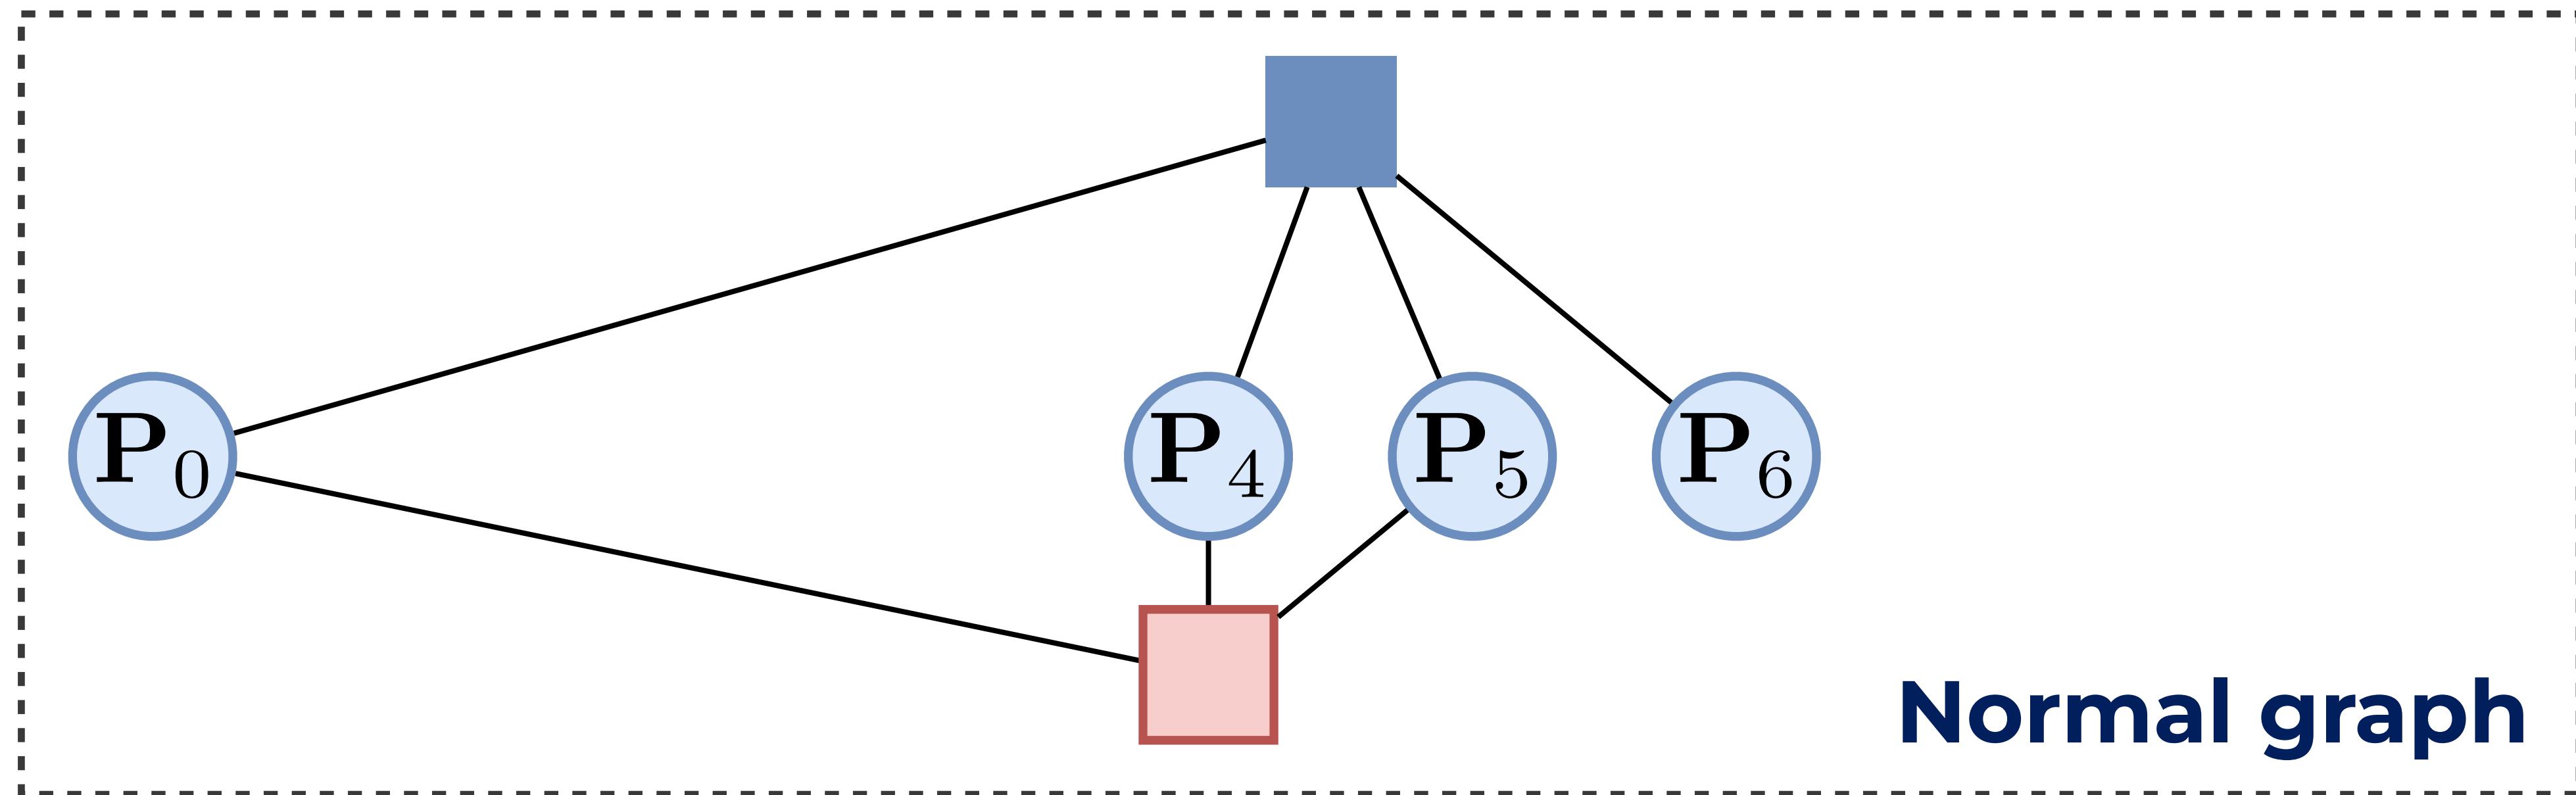
Marginalization order: $P_1 \rightarrow P_2 \rightarrow P_3 \rightarrow P_4 \rightarrow P_0$



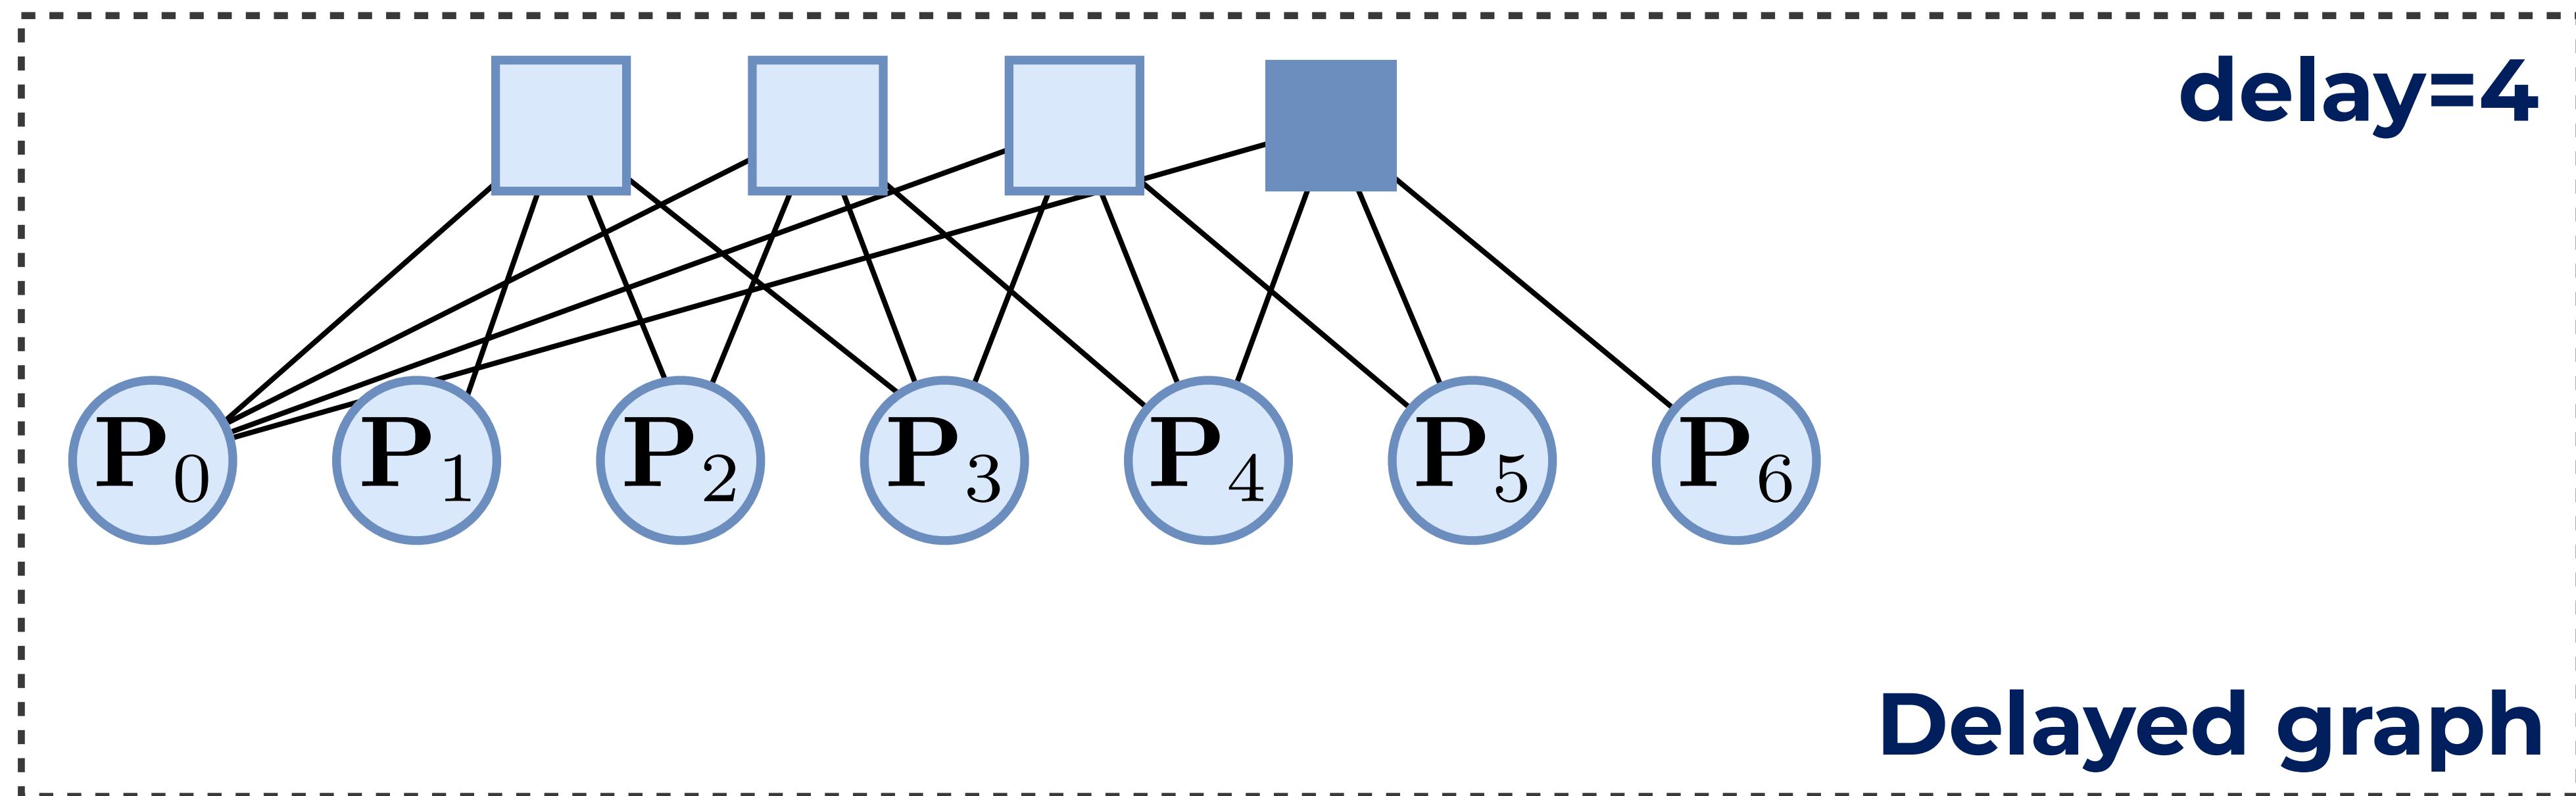
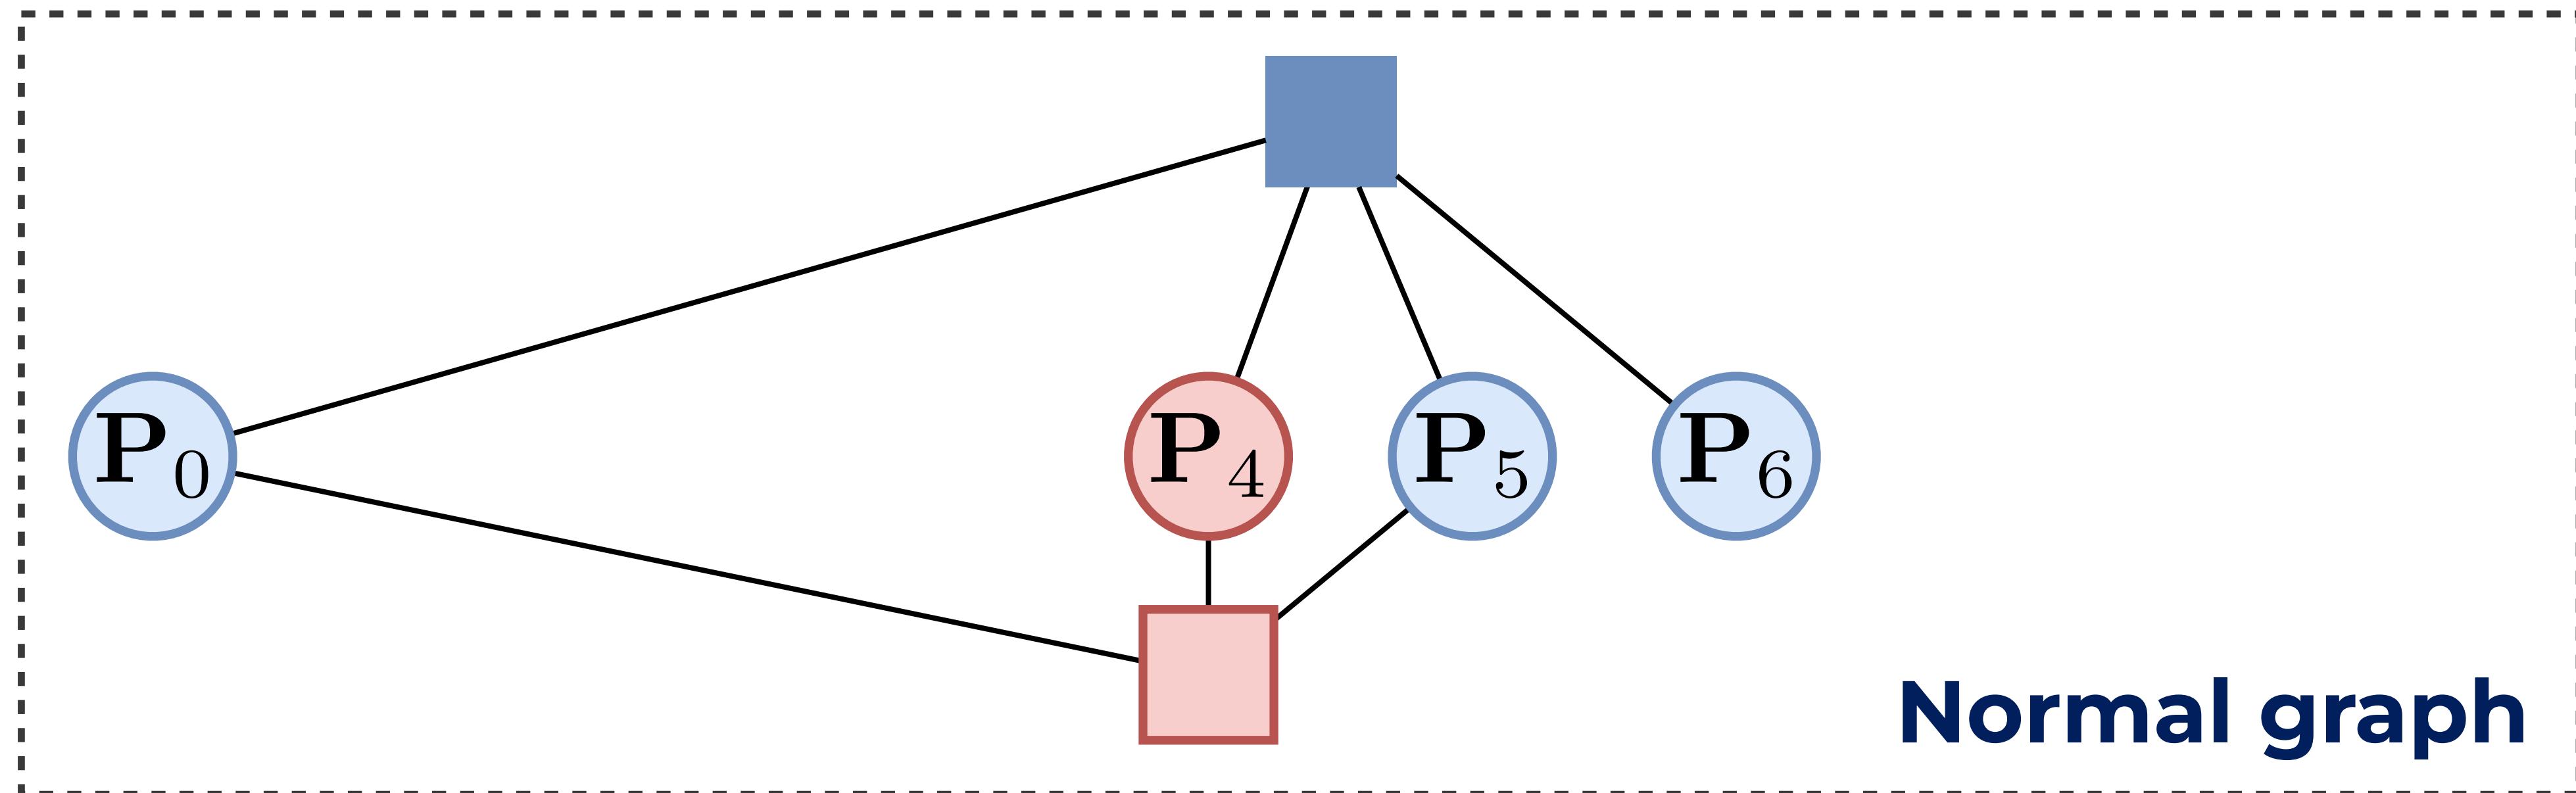
Marginalization order: $P_1 \rightarrow P_2 \rightarrow P_3 \rightarrow P_4 \rightarrow P_0$



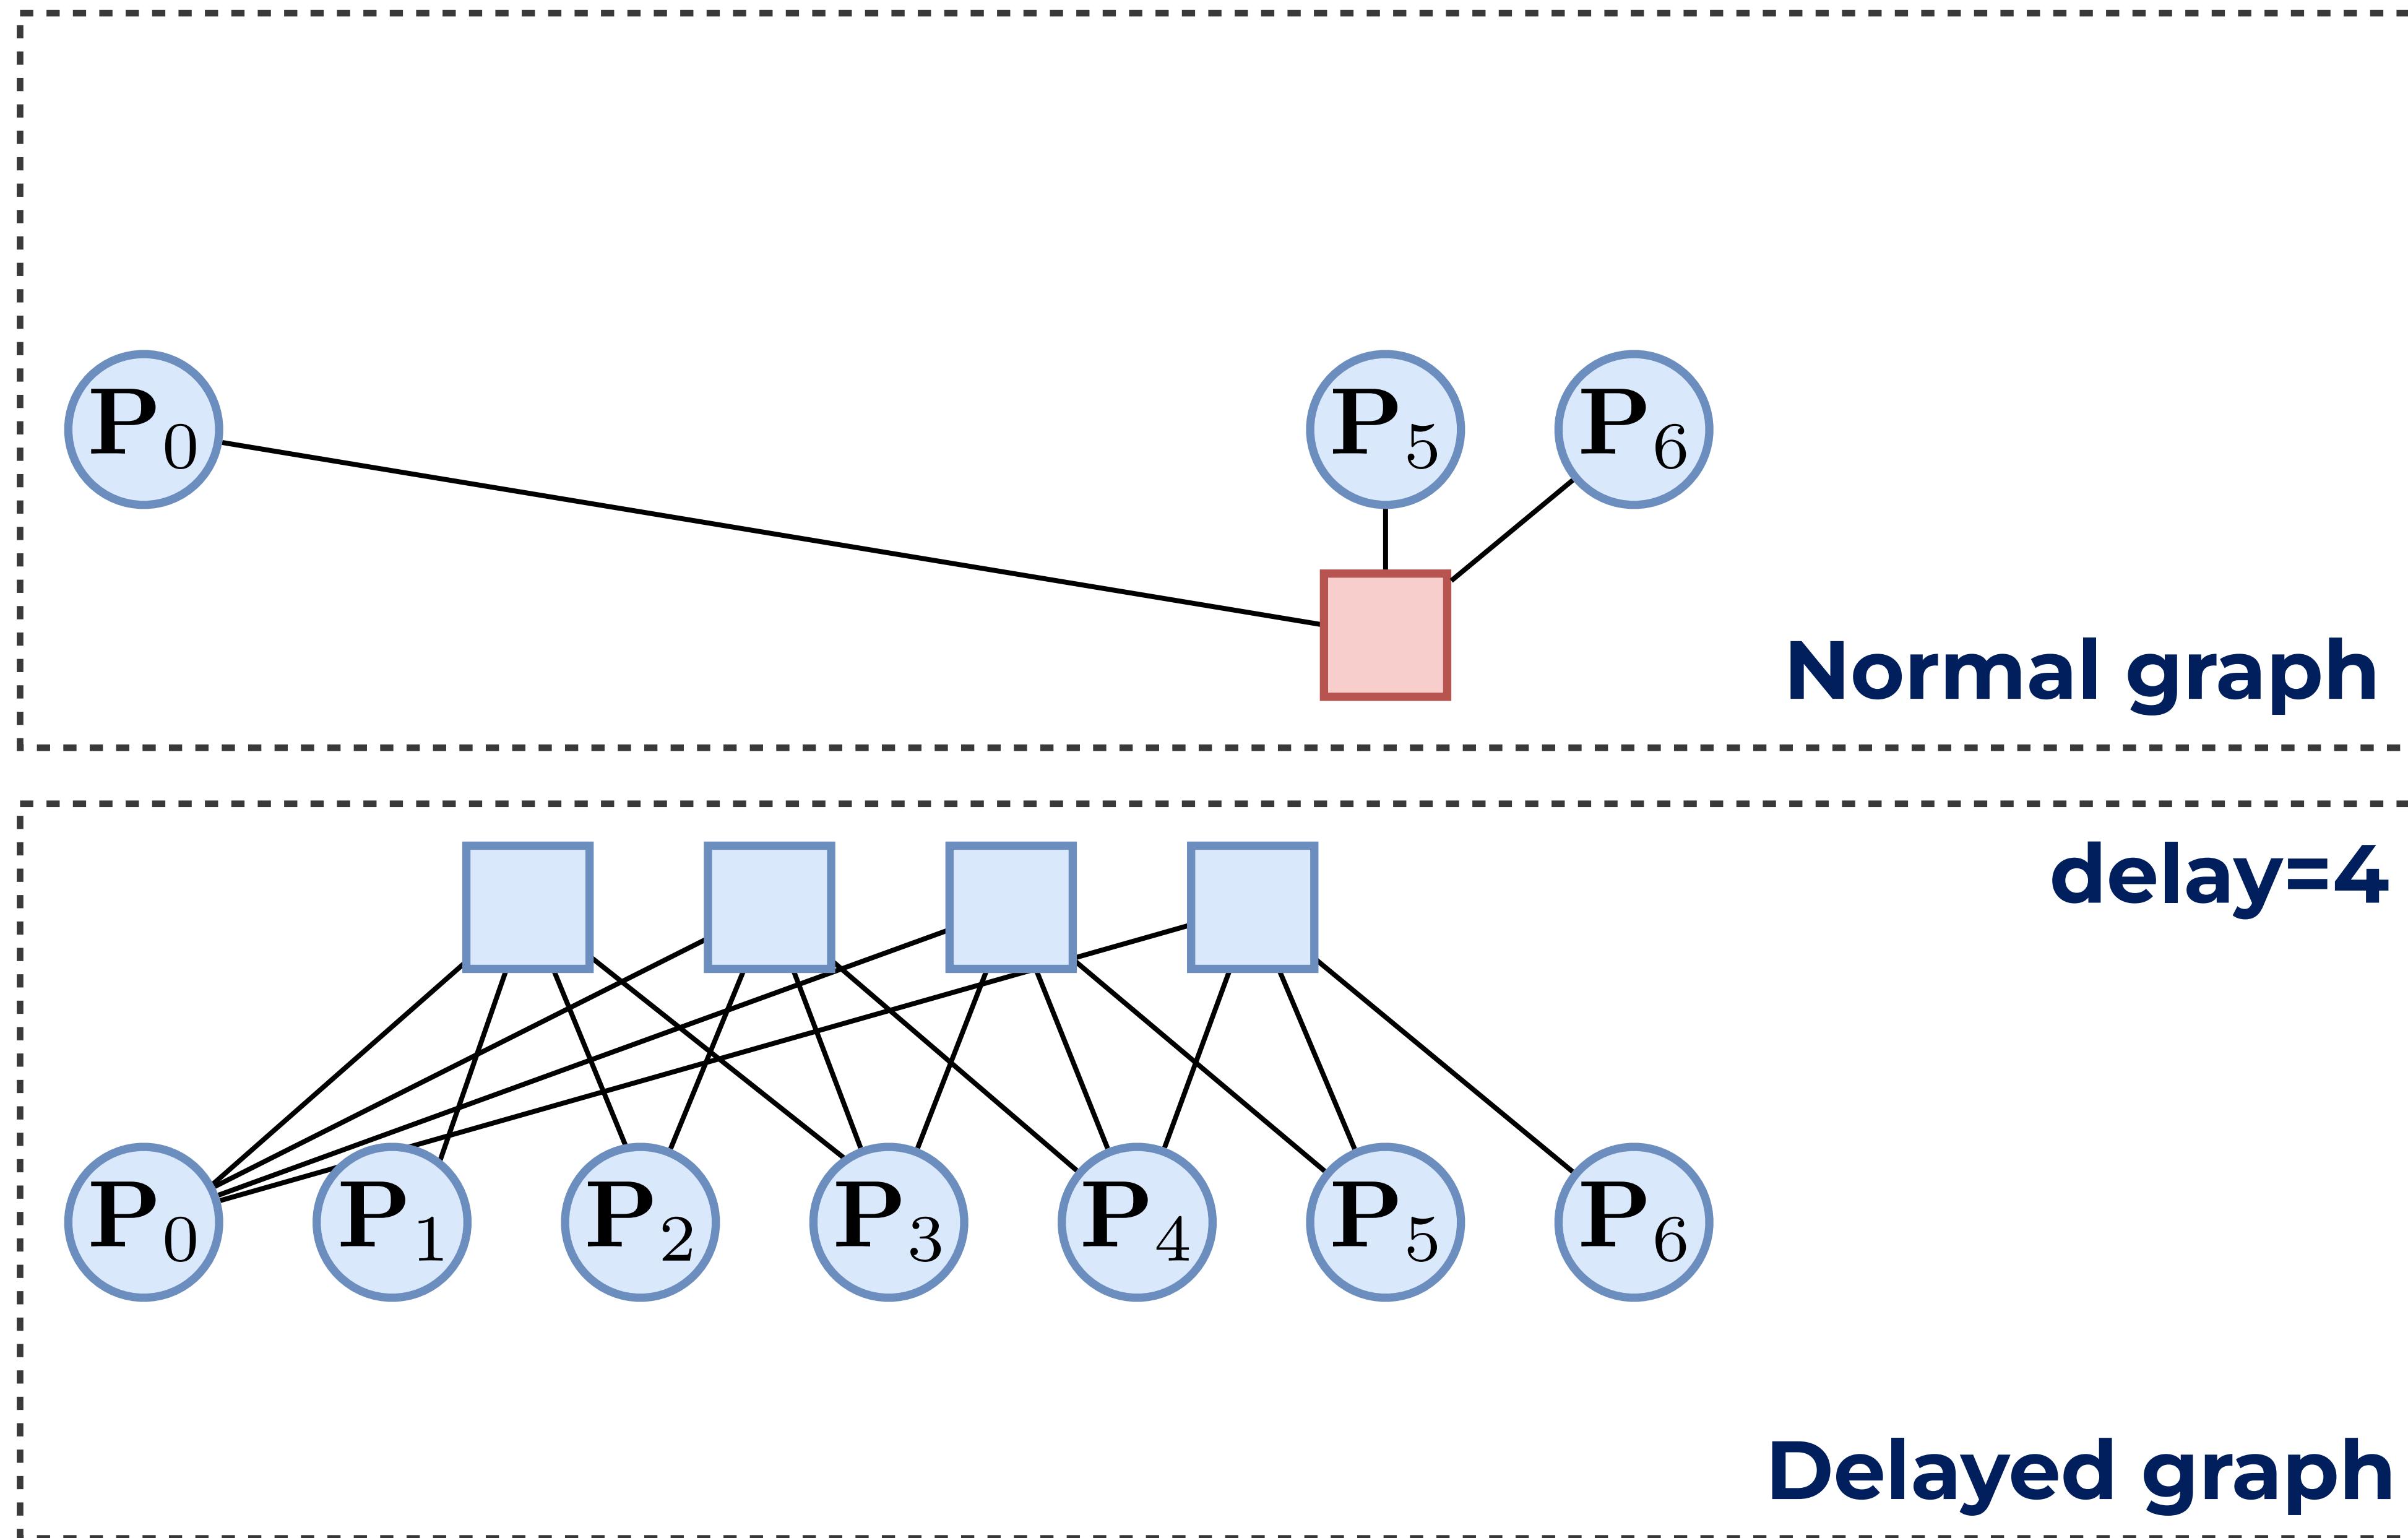
Marginalization order: $P_1 \rightarrow P_2 \rightarrow P_3 \rightarrow P_4 \rightarrow P_0$



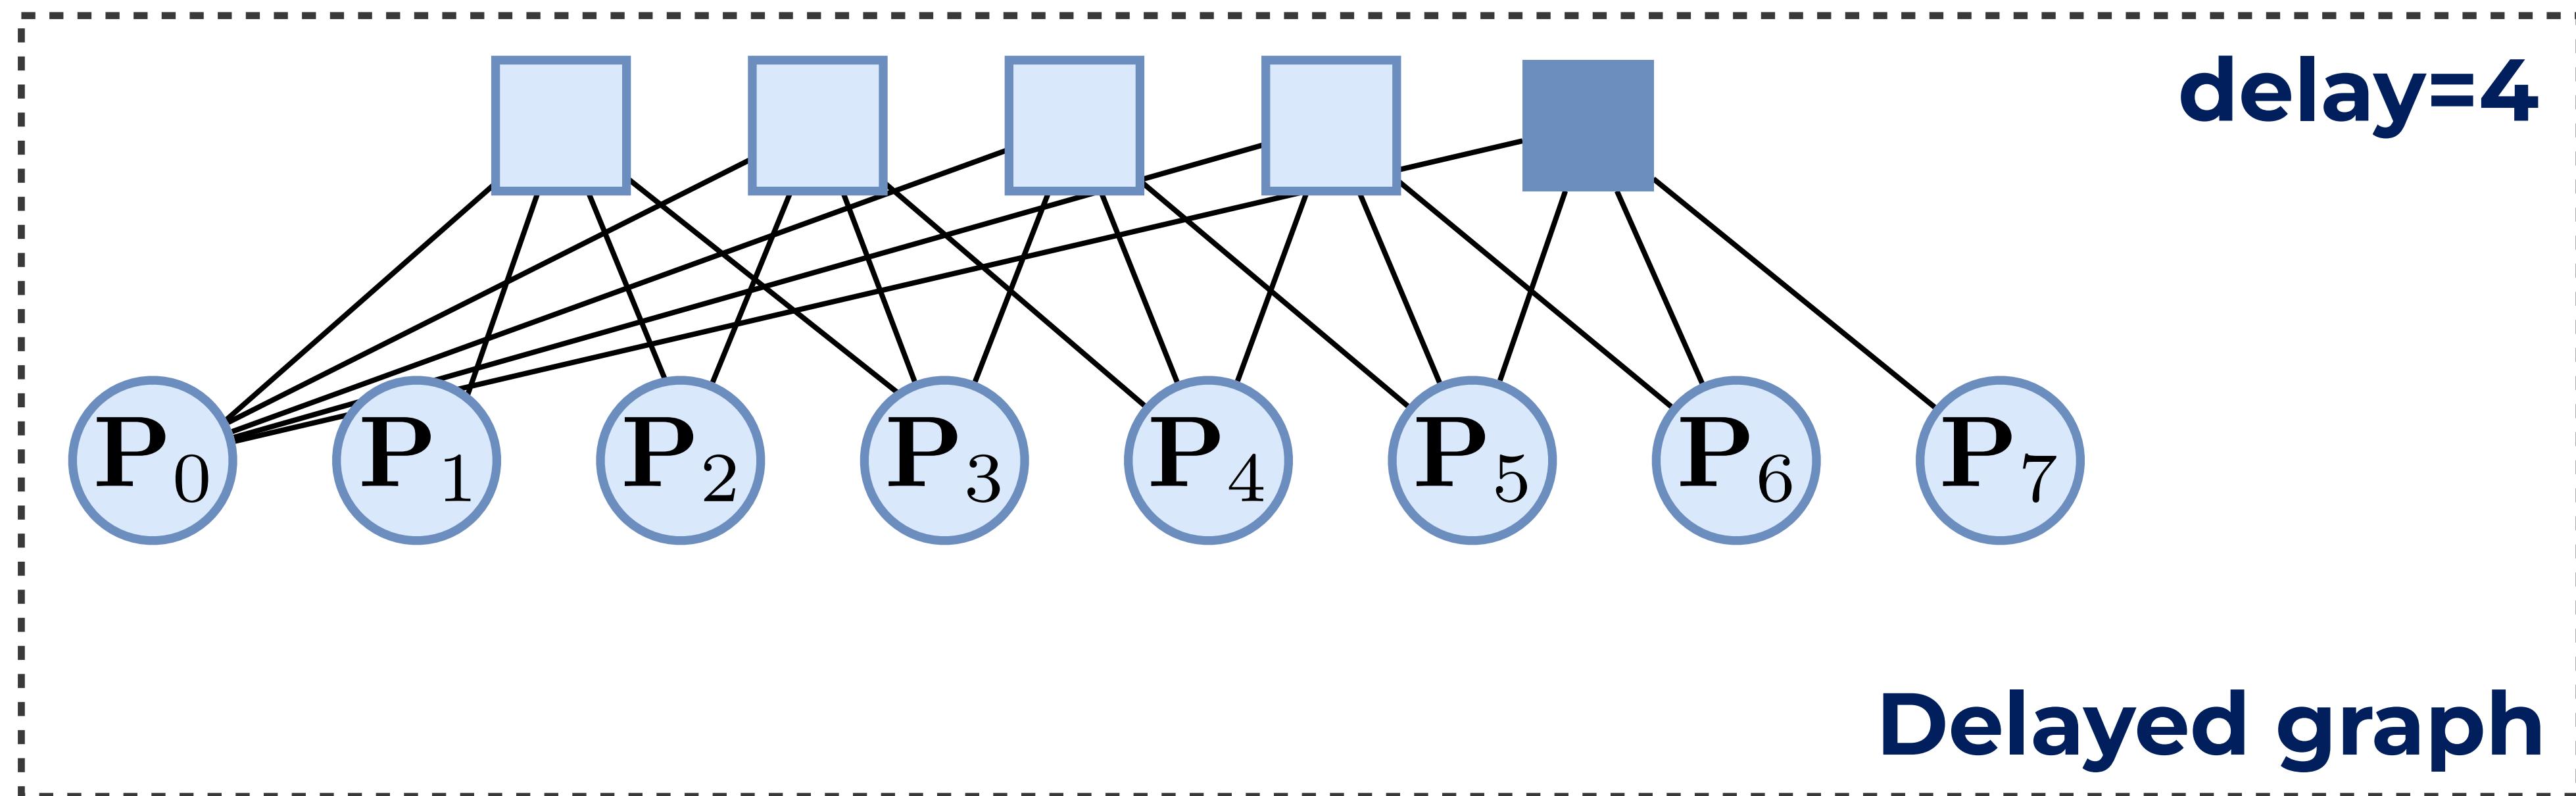
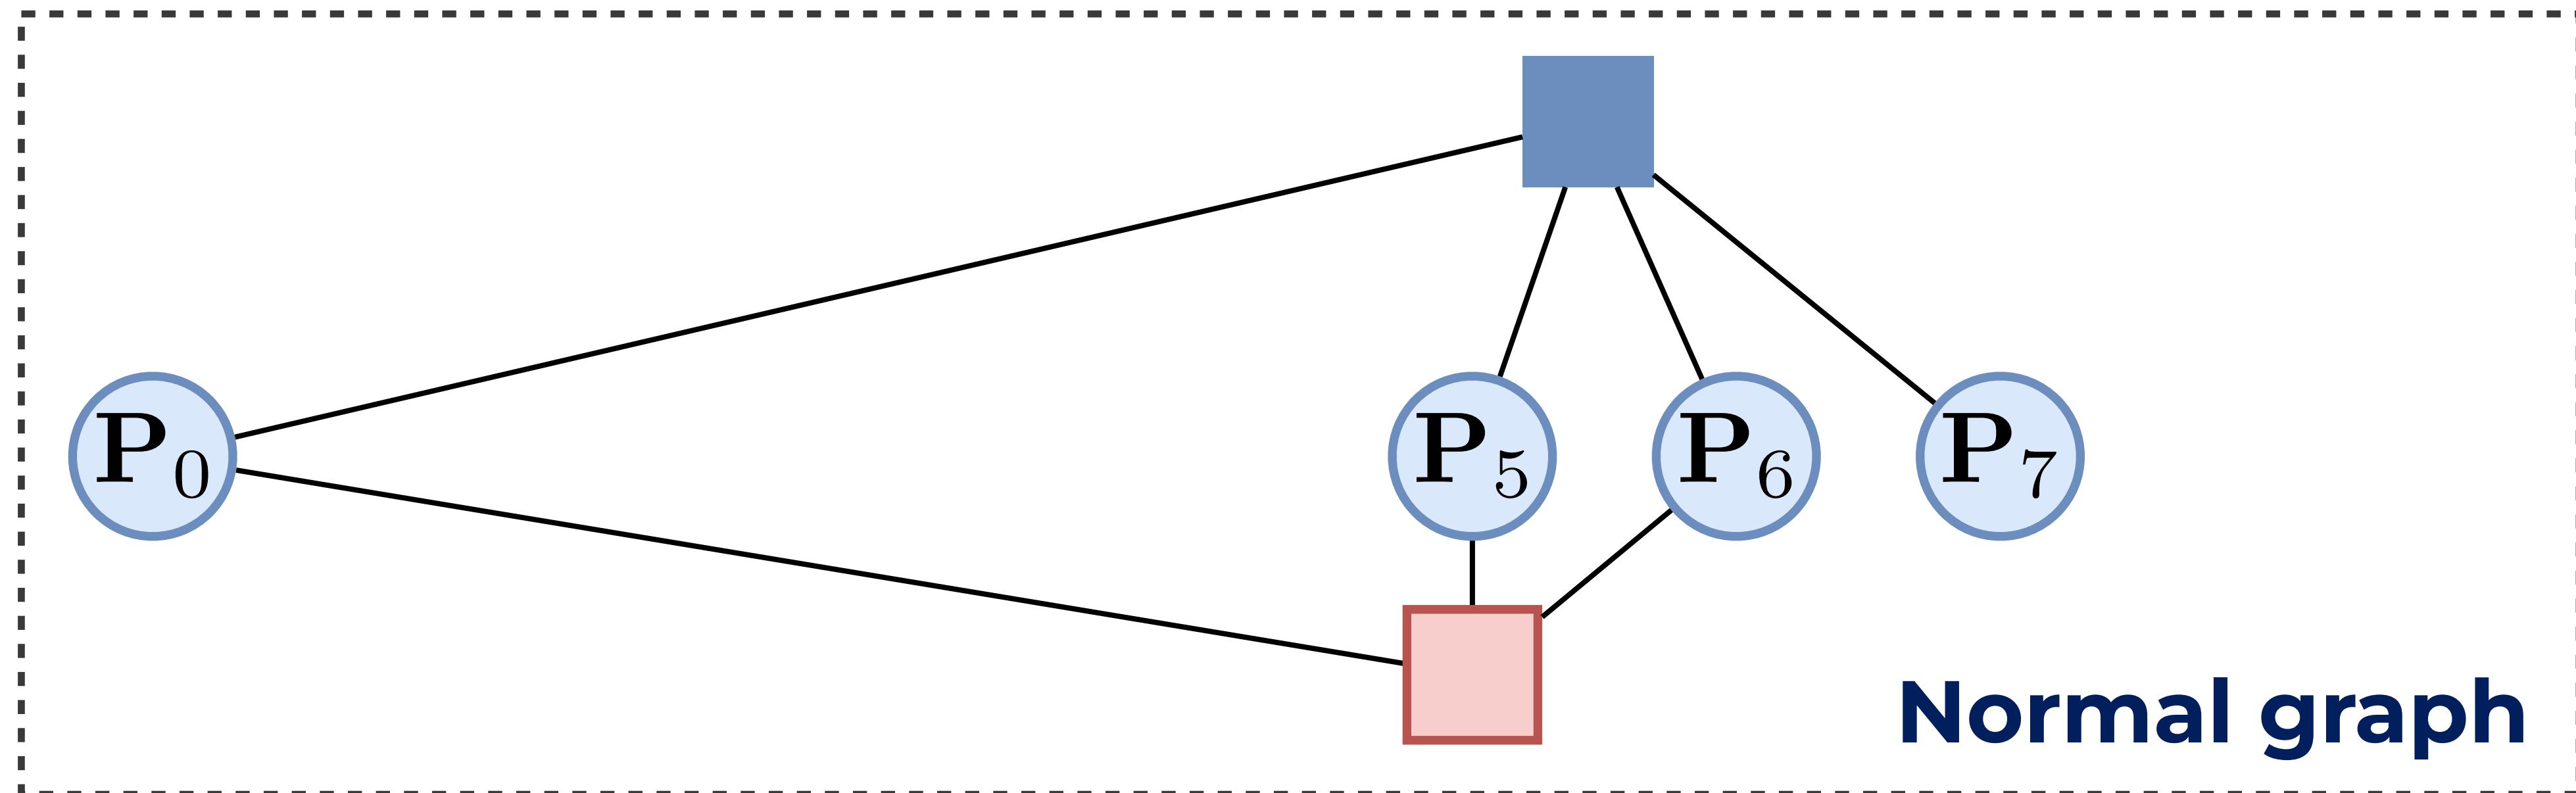
Marginalization order: $P_1 \rightarrow P_2 \rightarrow P_3 \rightarrow P_4 \rightarrow P_0$



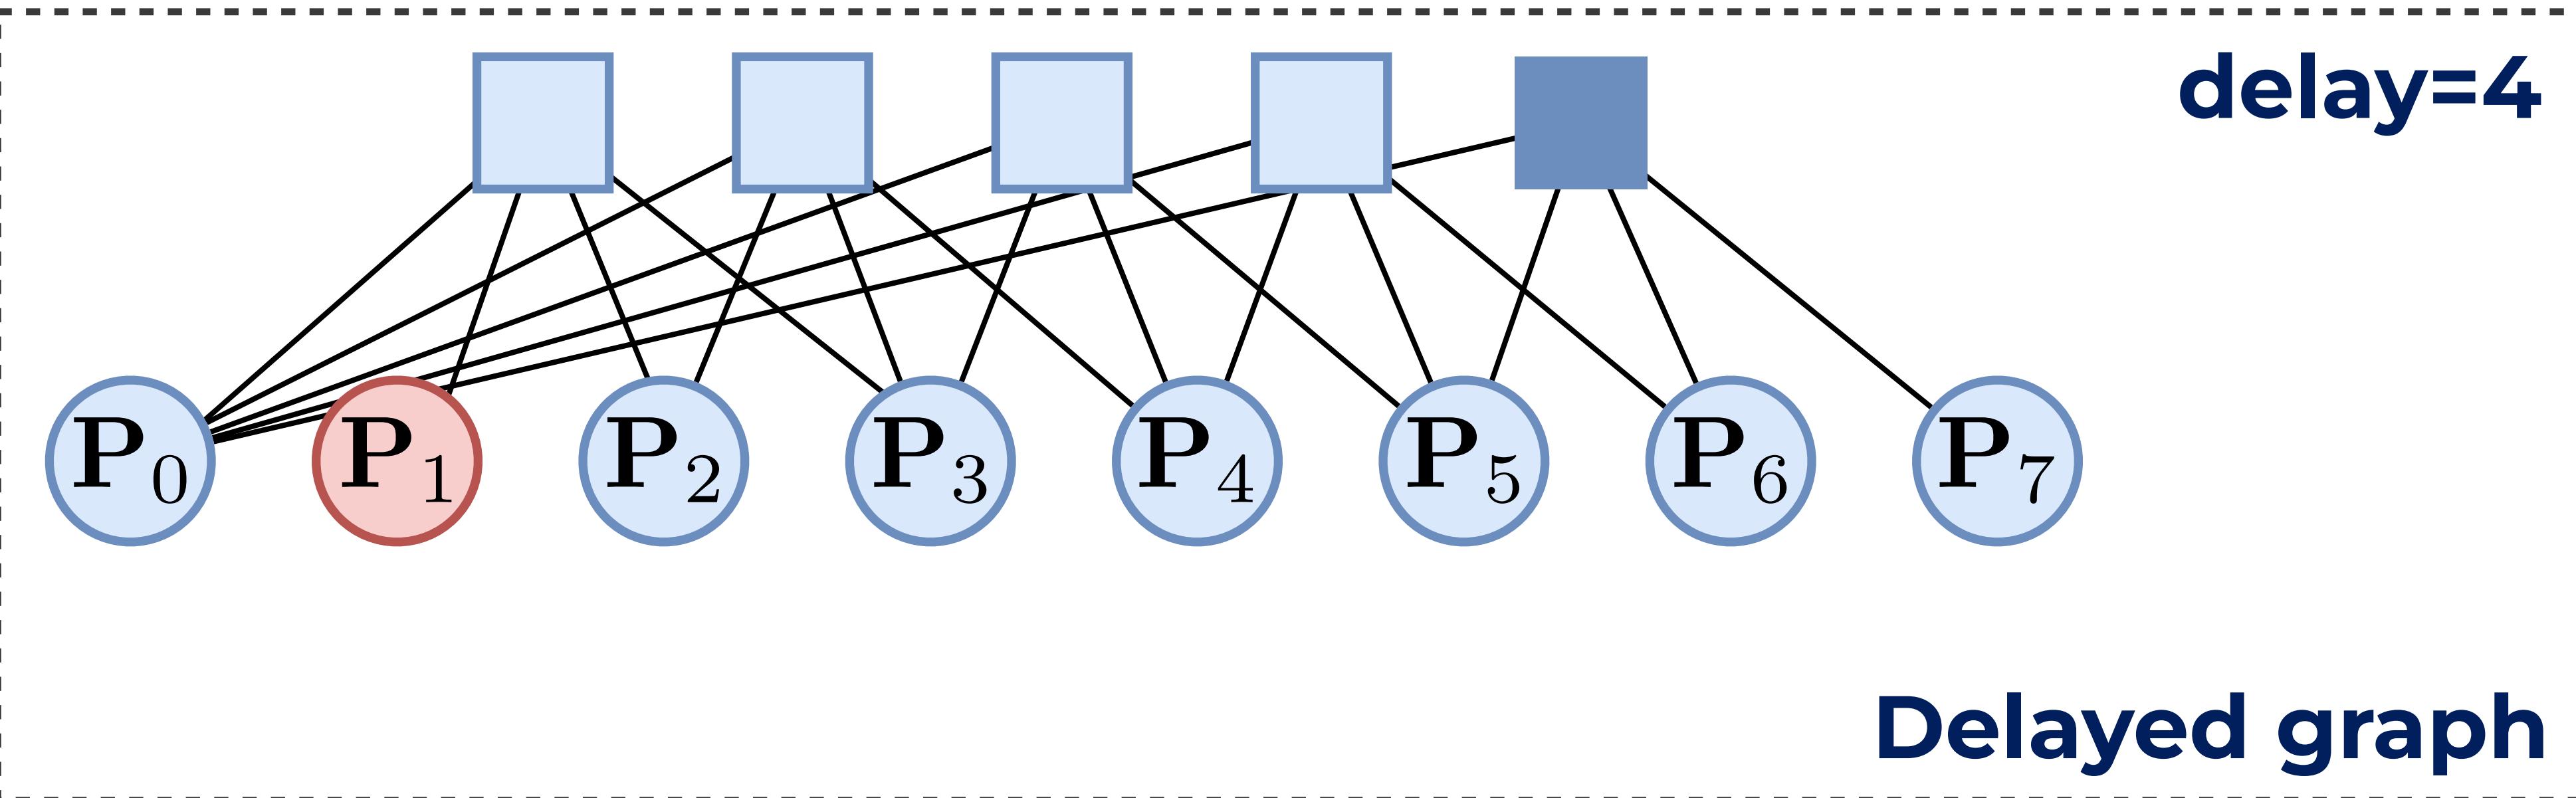
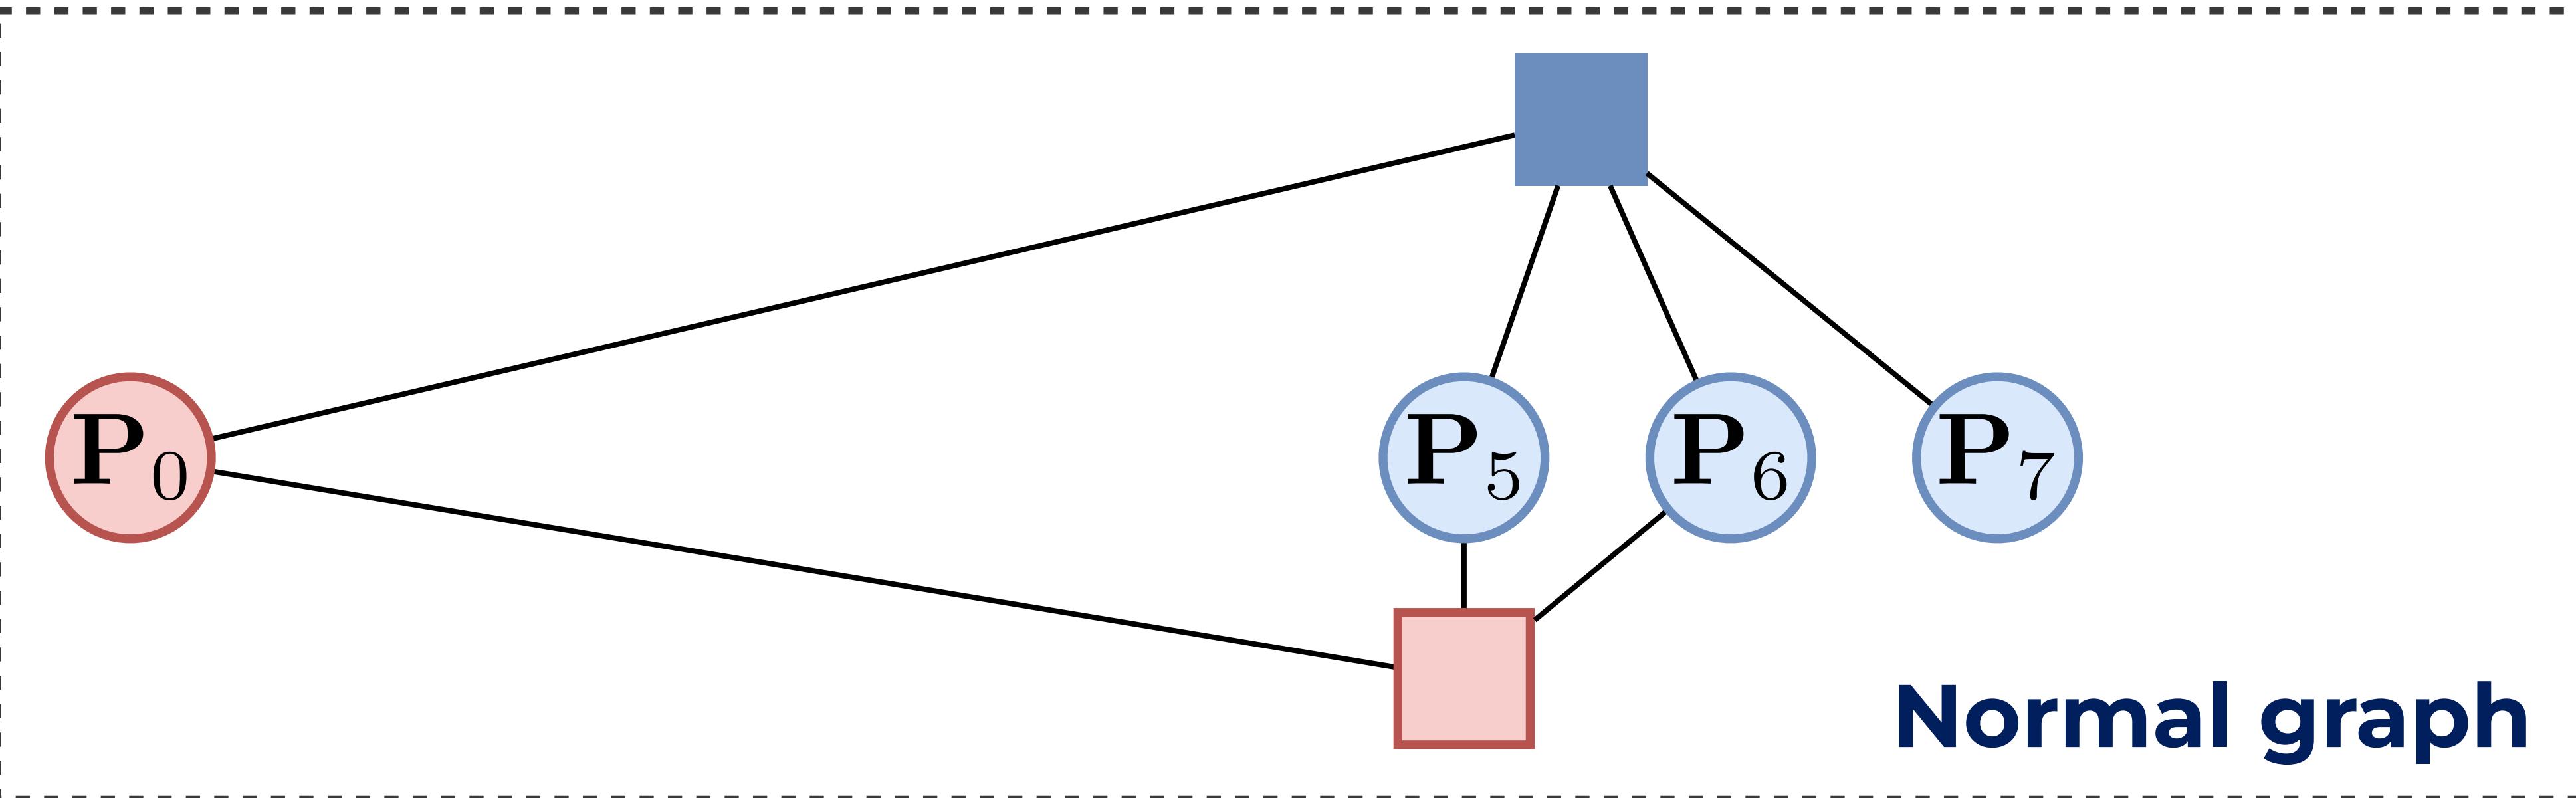
Marginalization order: $P_1 \rightarrow P_2 \rightarrow P_3 \rightarrow P_4 \rightarrow P_0$



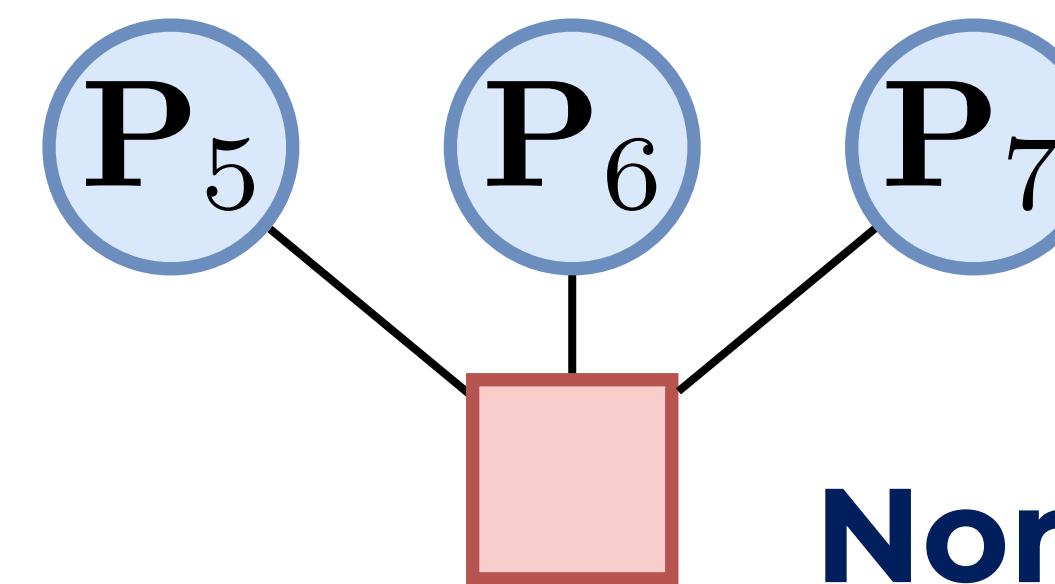
Marginalization order: $P_1 \rightarrow P_2 \rightarrow P_3 \rightarrow P_4 \rightarrow P_0$



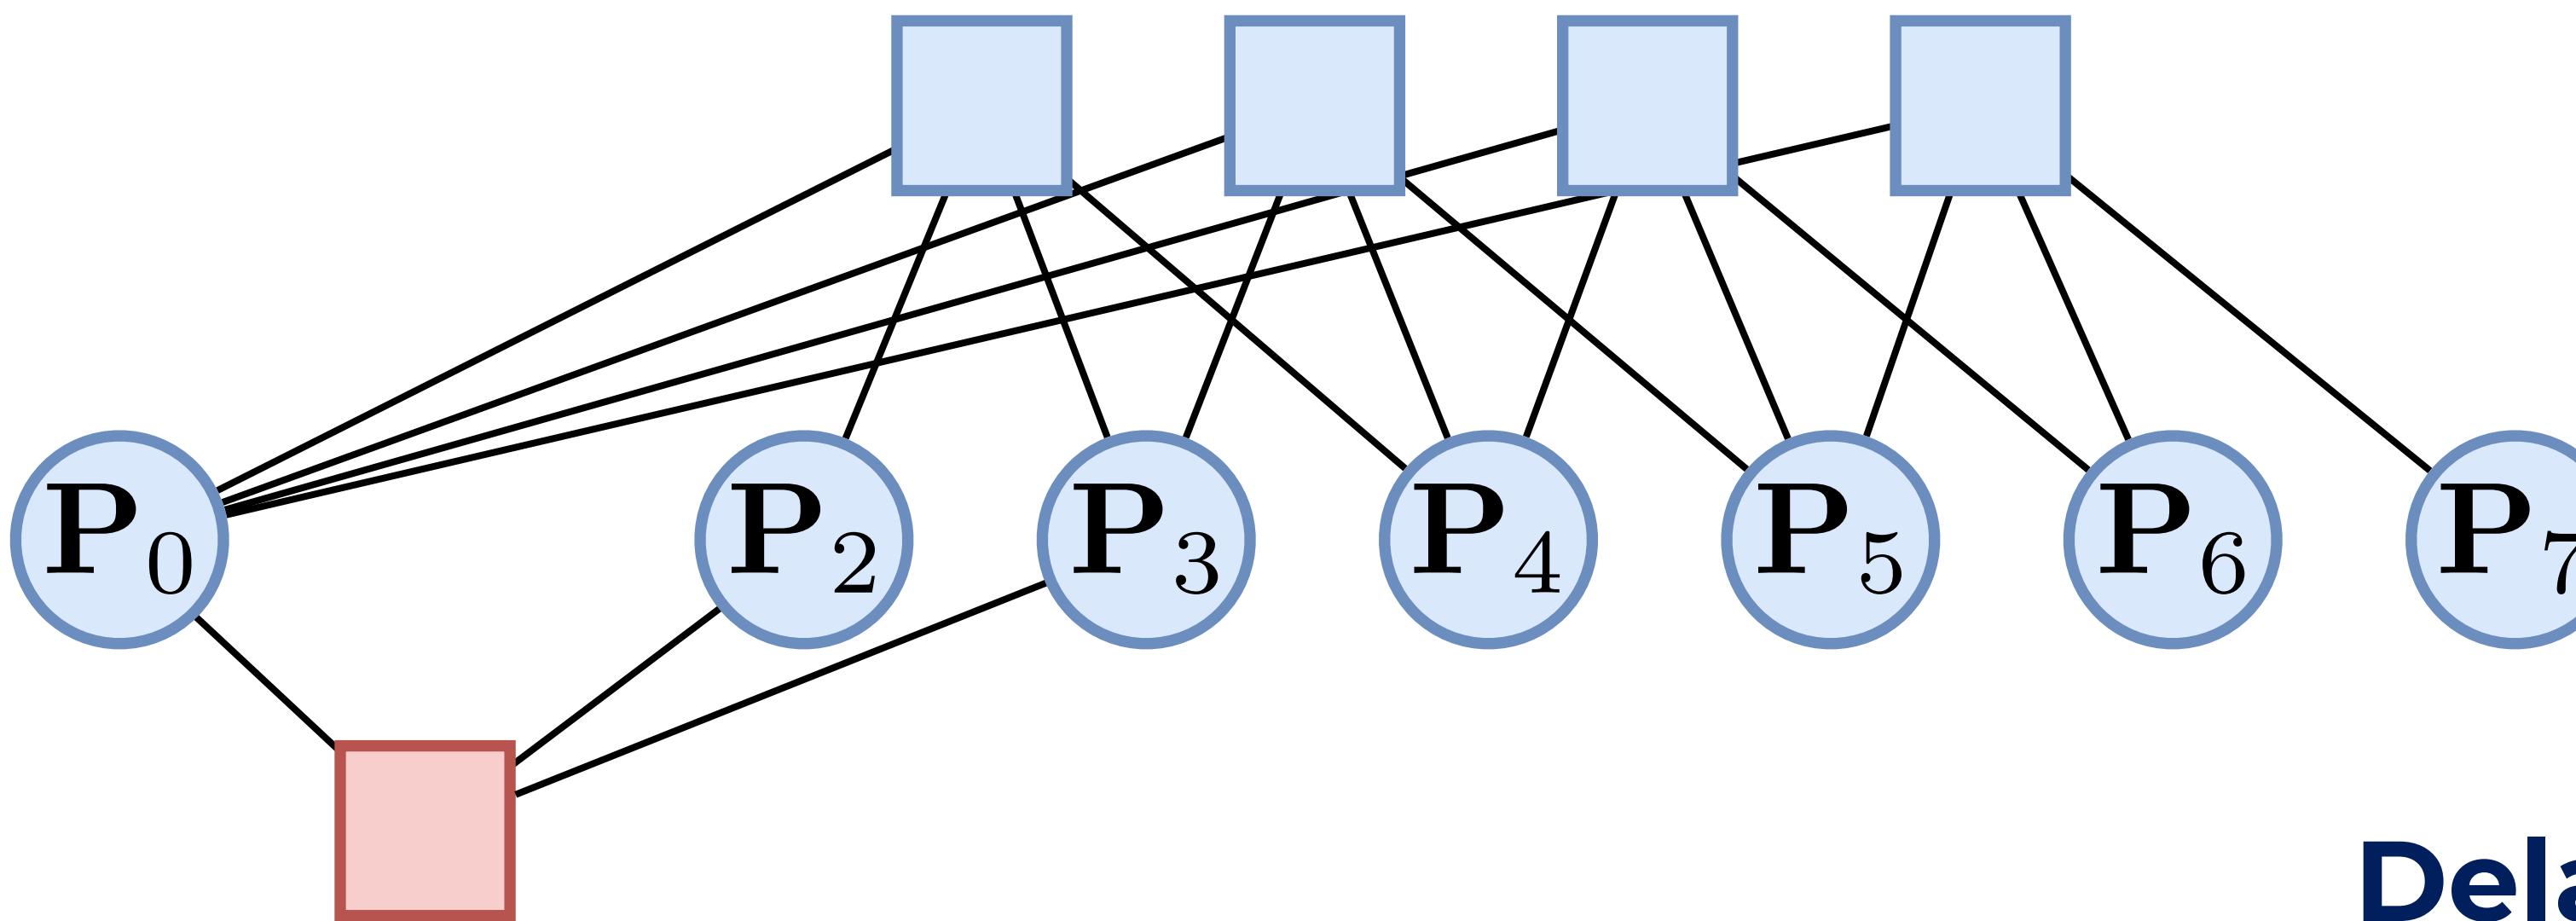
Marginalization order: $P_1 \rightarrow P_2 \rightarrow P_3 \rightarrow P_4 \rightarrow P_0$



Marginalization order: $P_1 \rightarrow P_2 \rightarrow P_3 \rightarrow P_4 \rightarrow P_0$

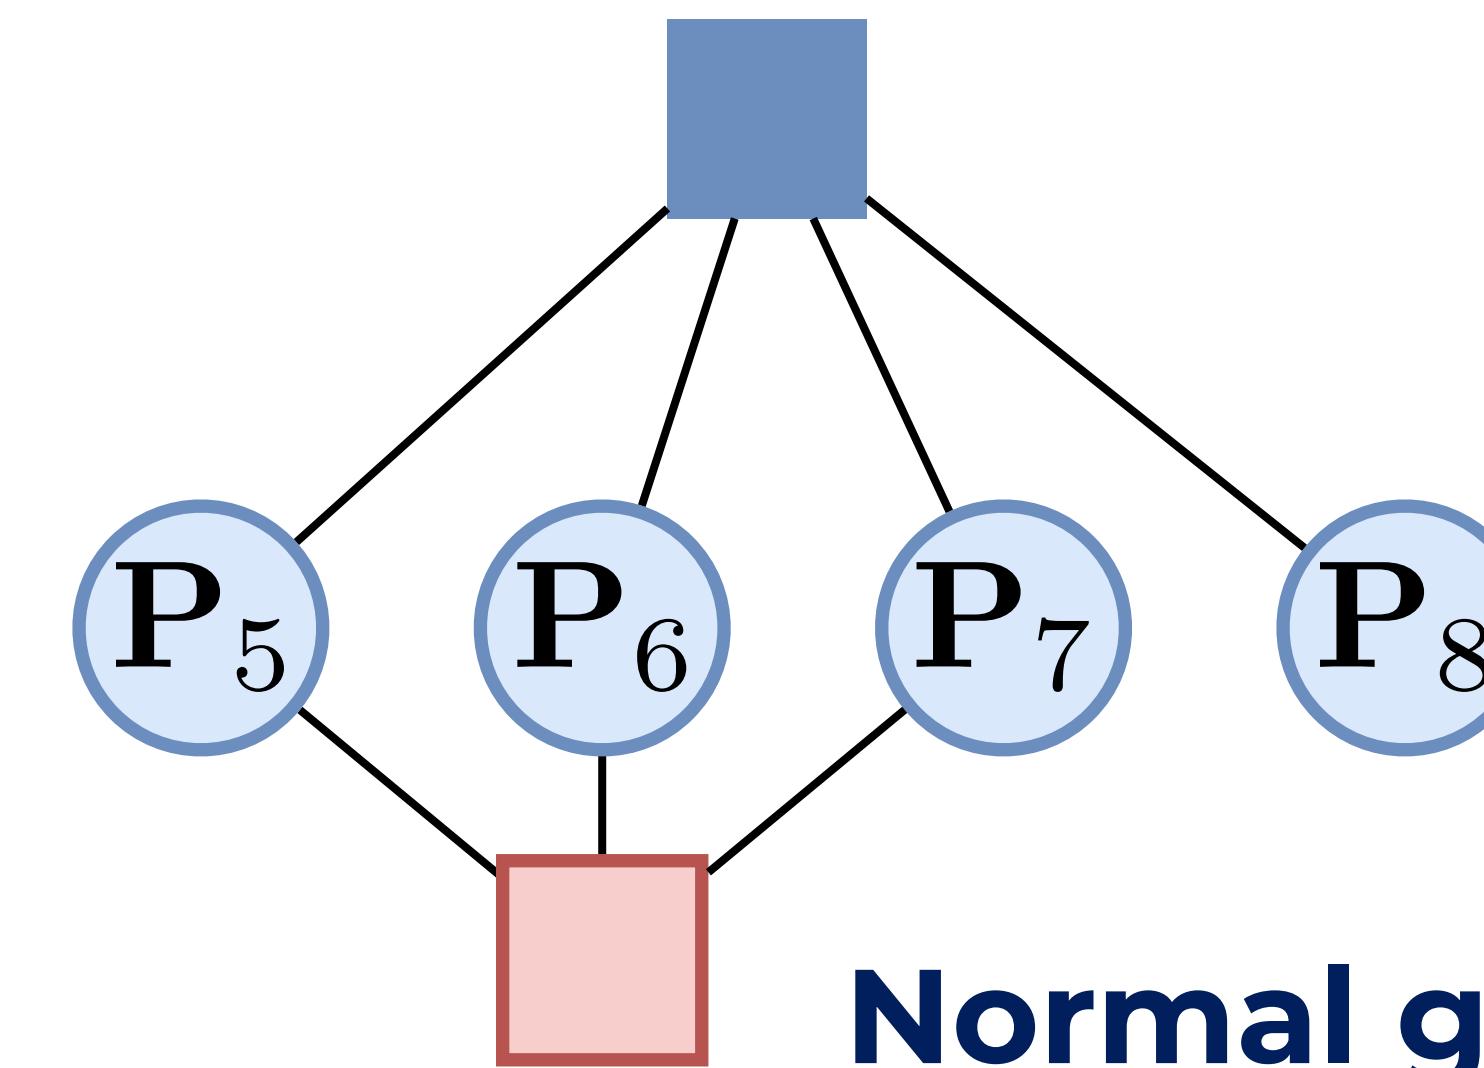


Normal graph

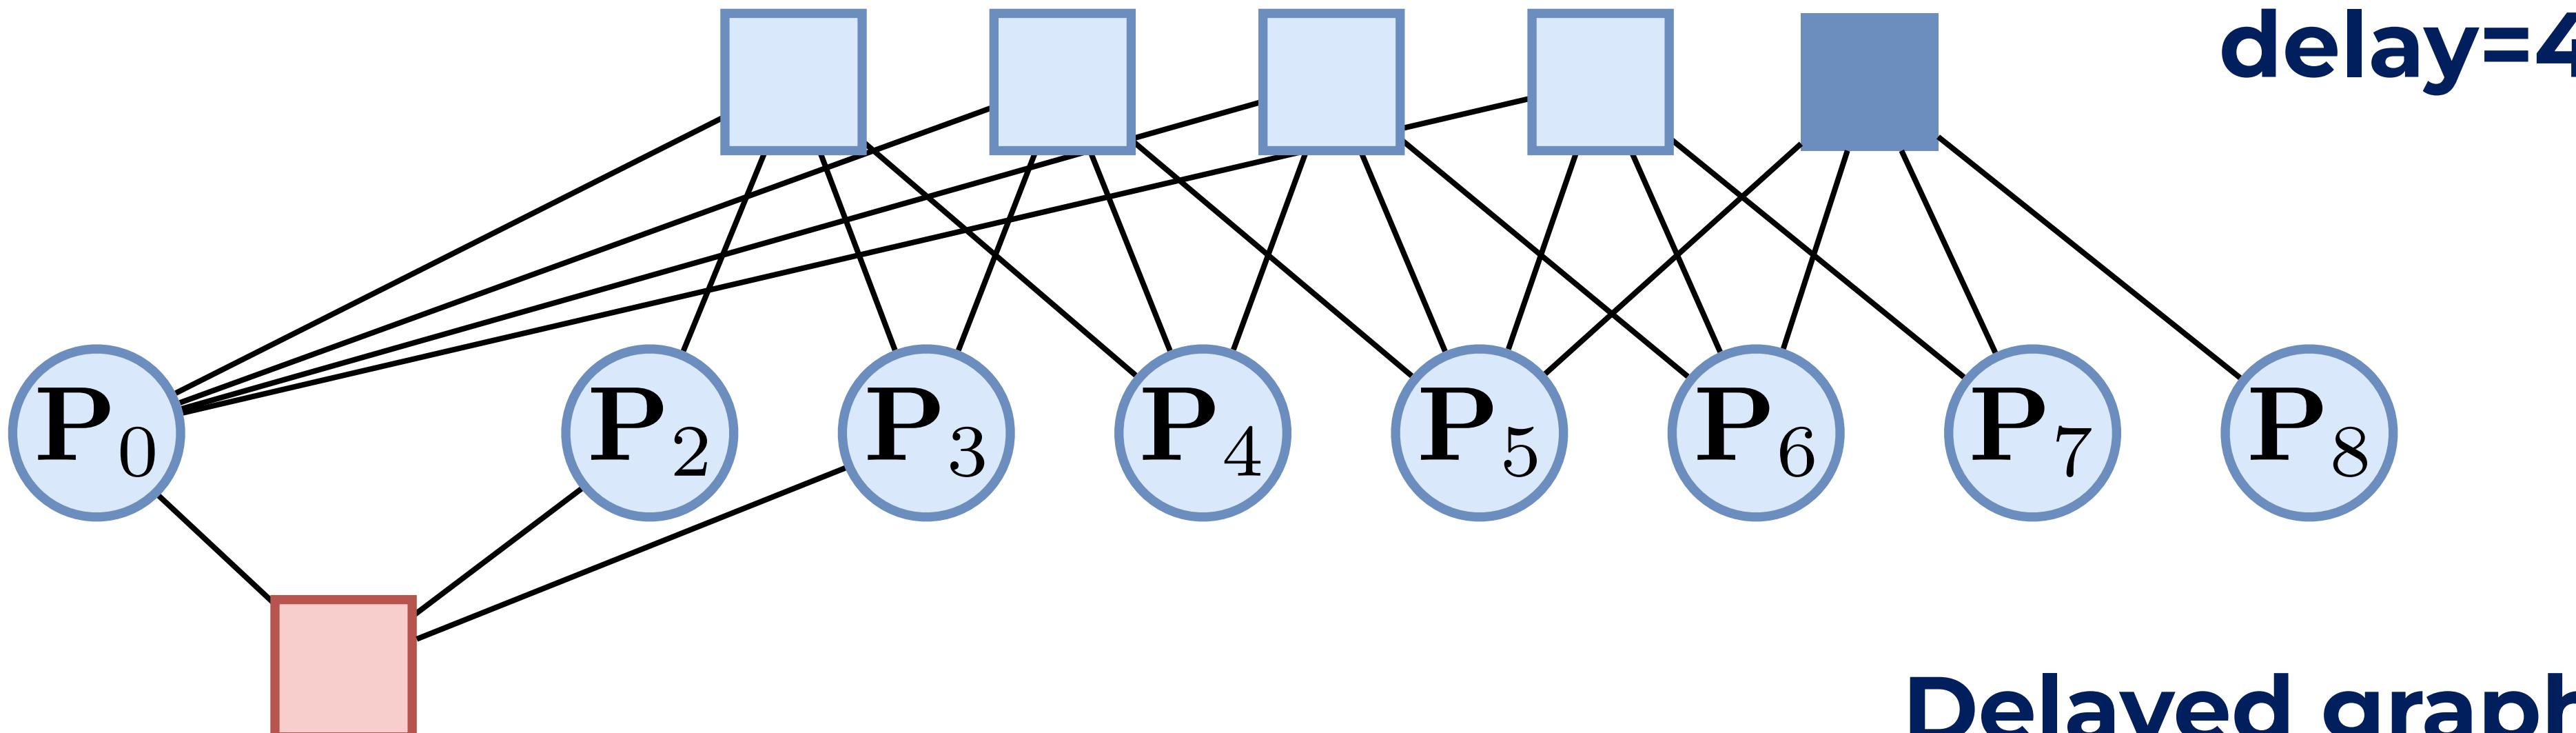


Delayed graph

Marginalization order: $P_1 \rightarrow P_2 \rightarrow P_3 \rightarrow P_4 \rightarrow P_0$

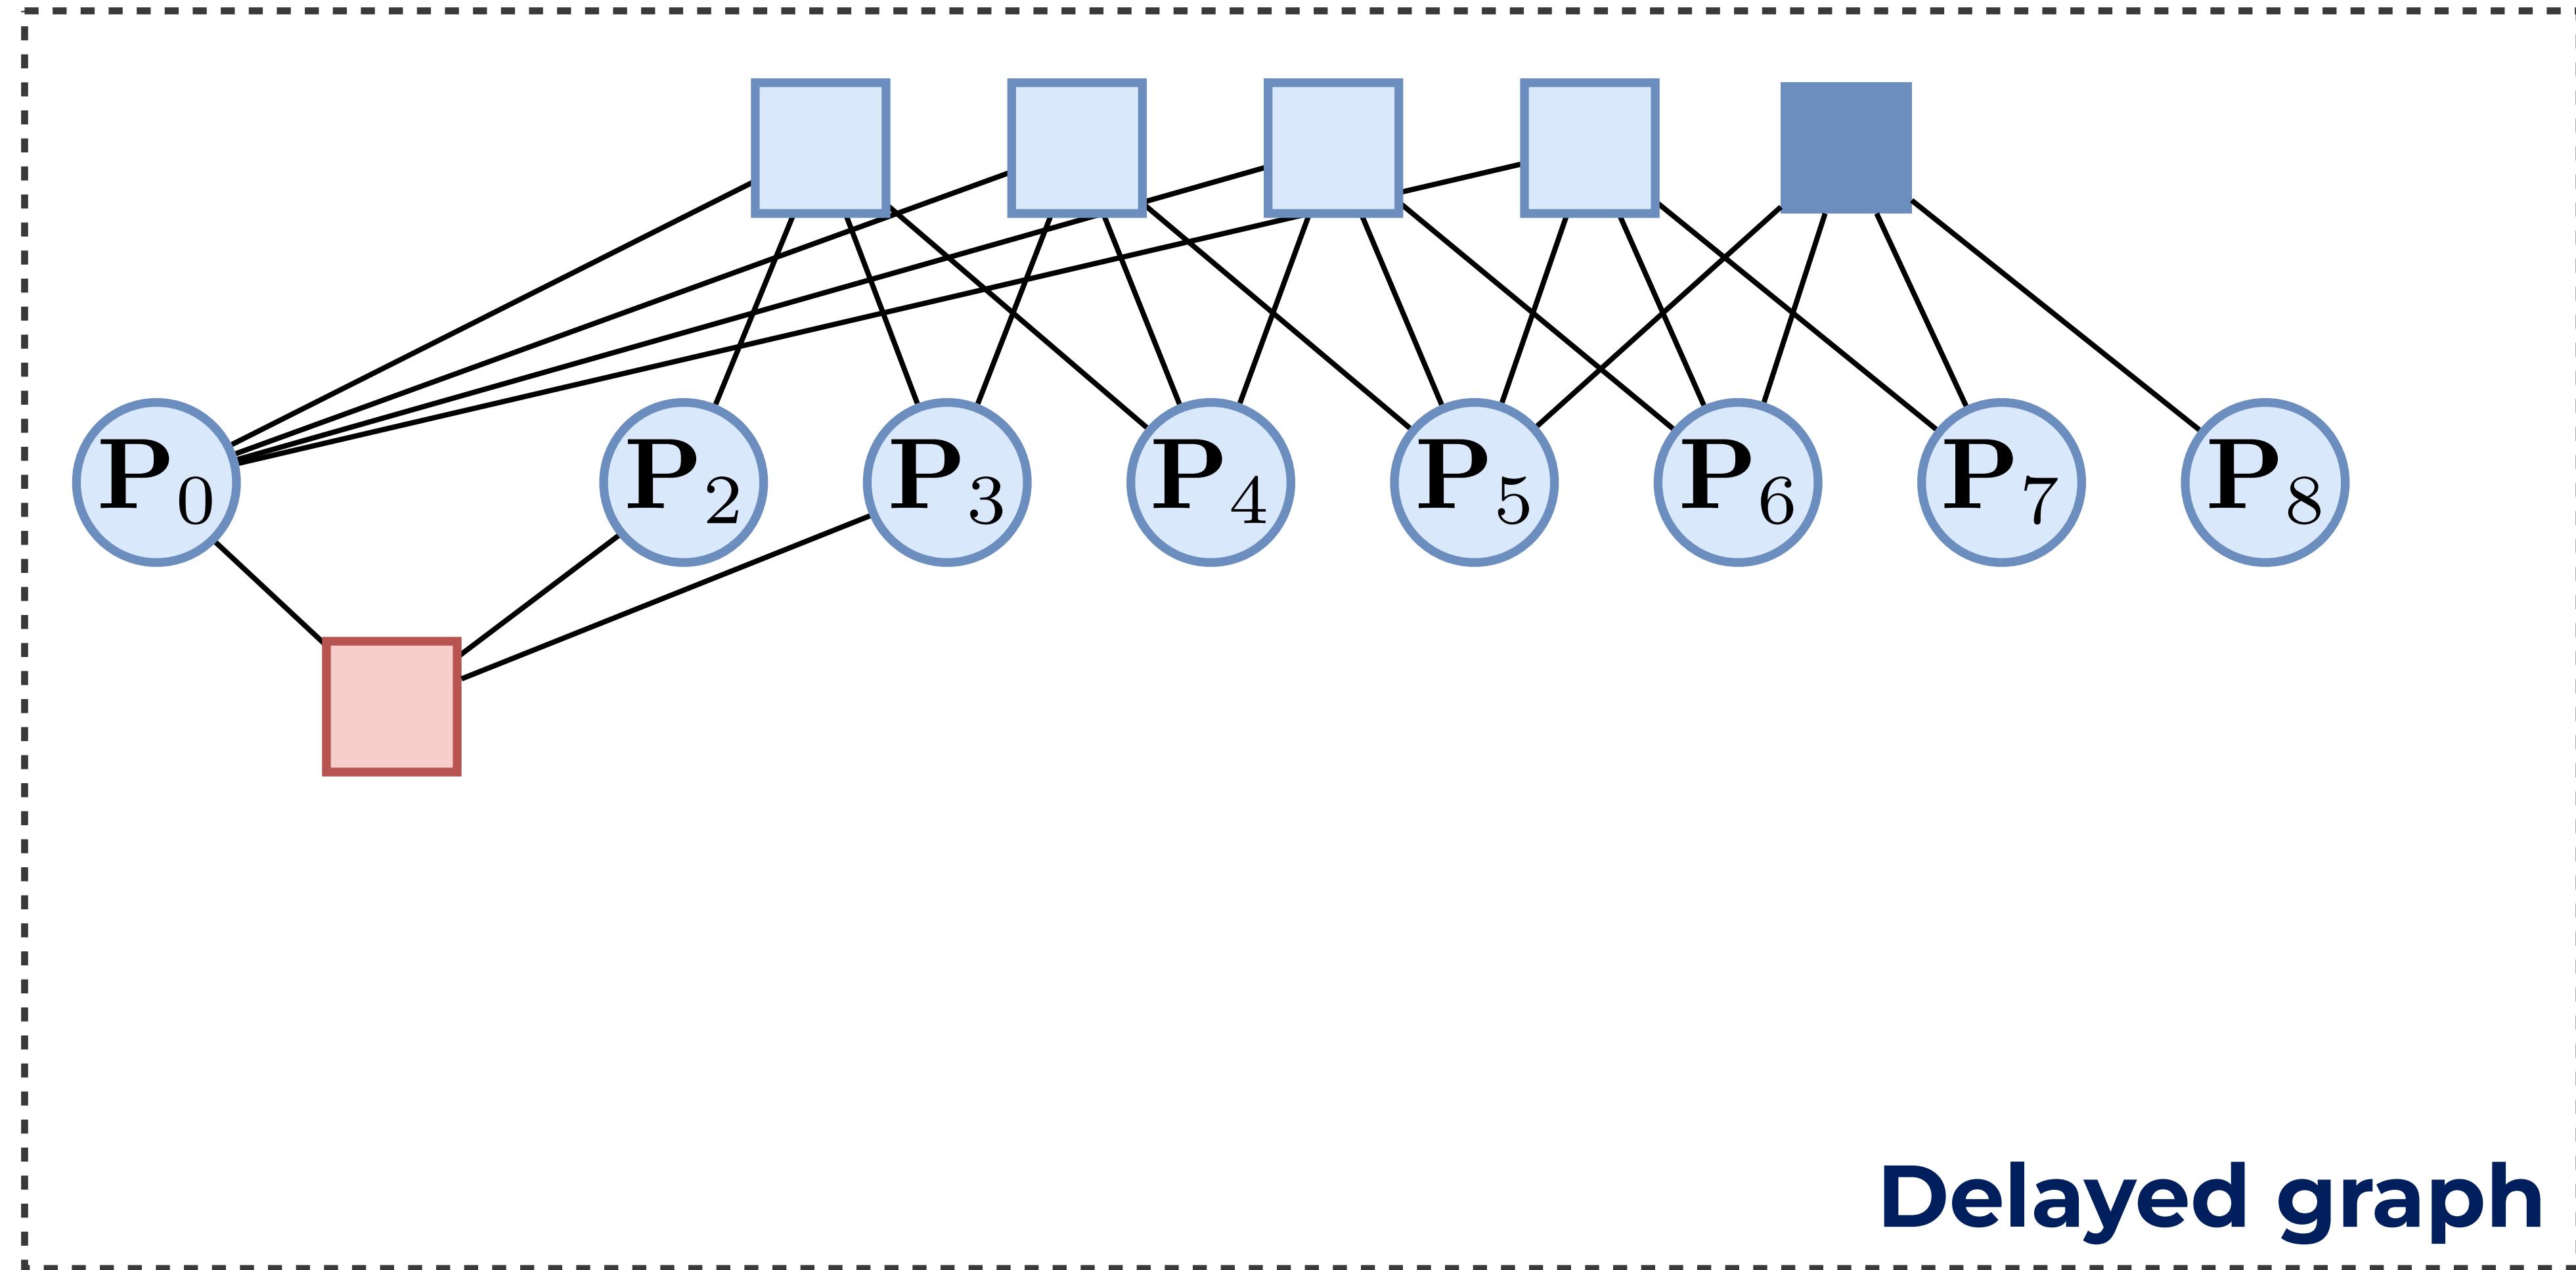


Normal graph

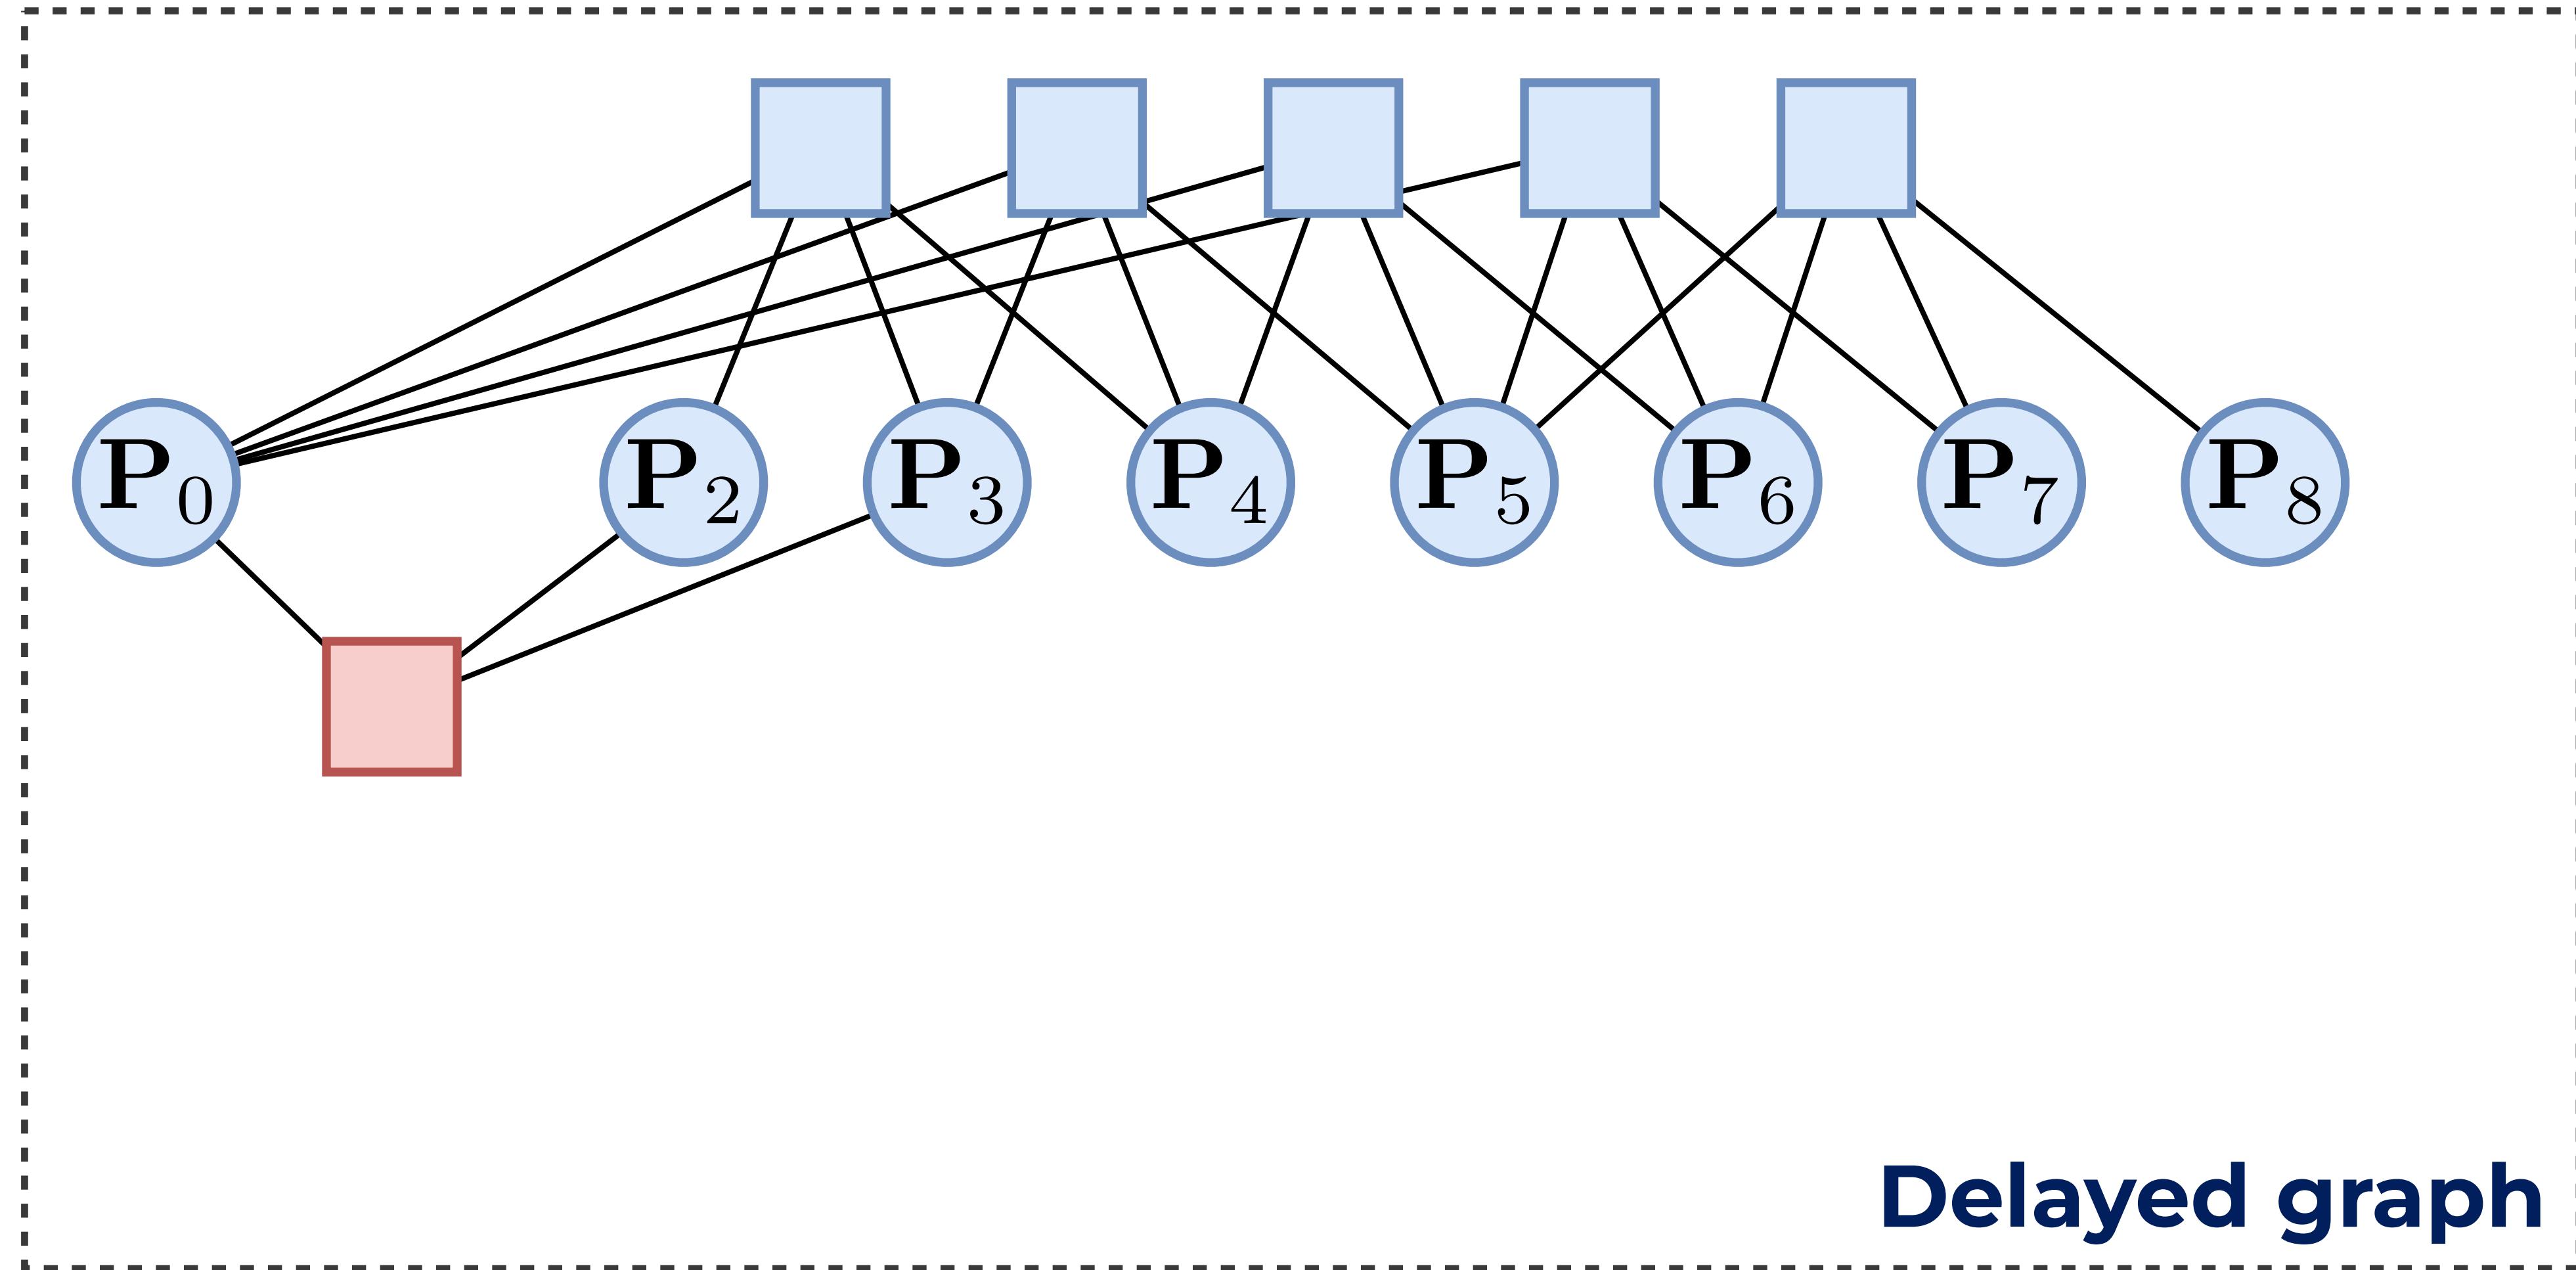


Delayed graph

Pose Graph Bundle Adjustment (PGBA)

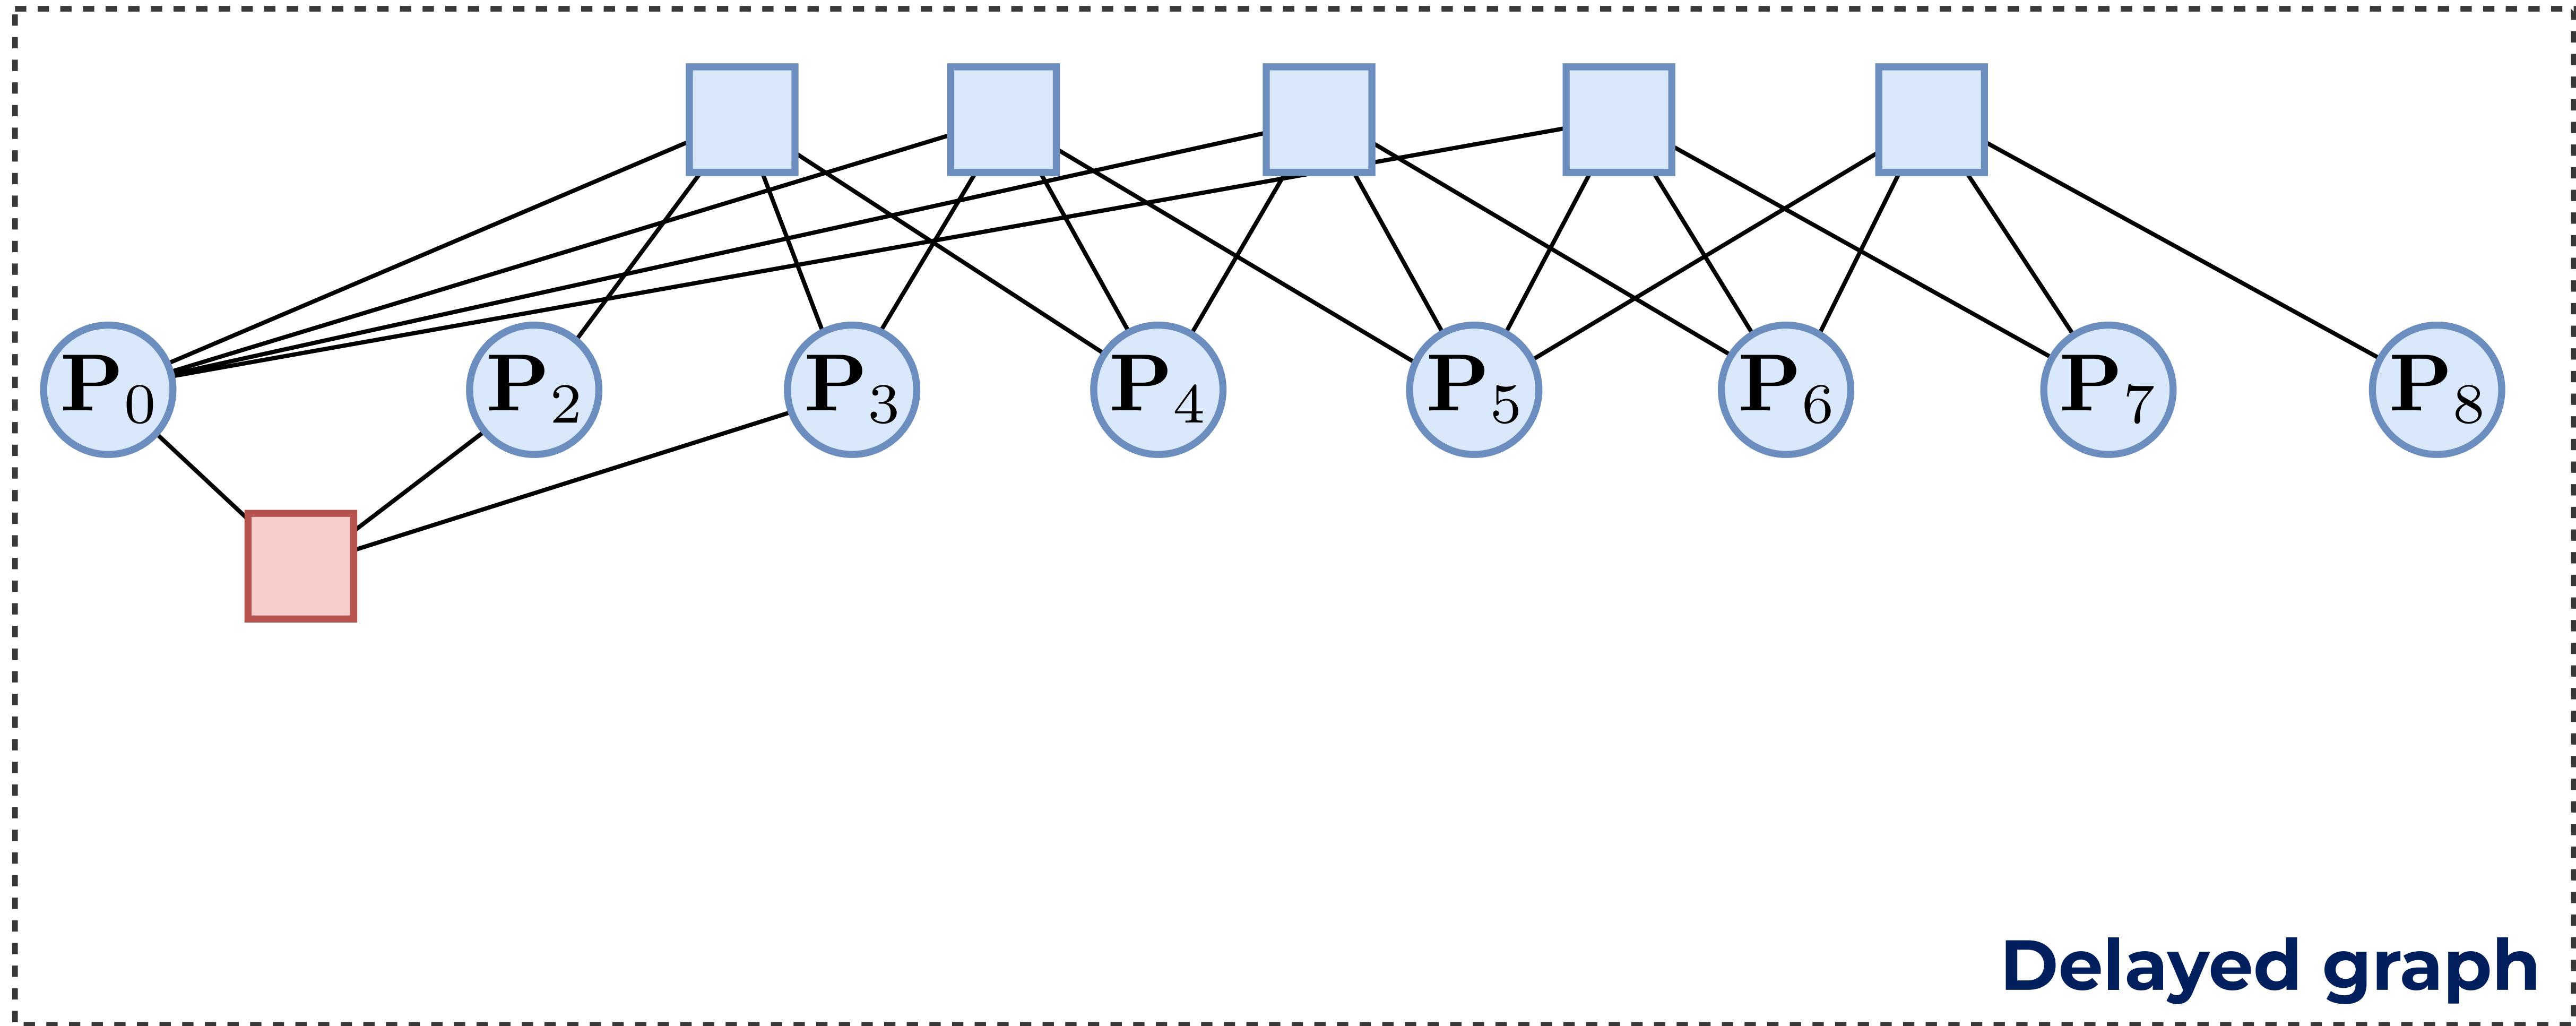


Pose Graph Bundle Adjustment (PGBA)



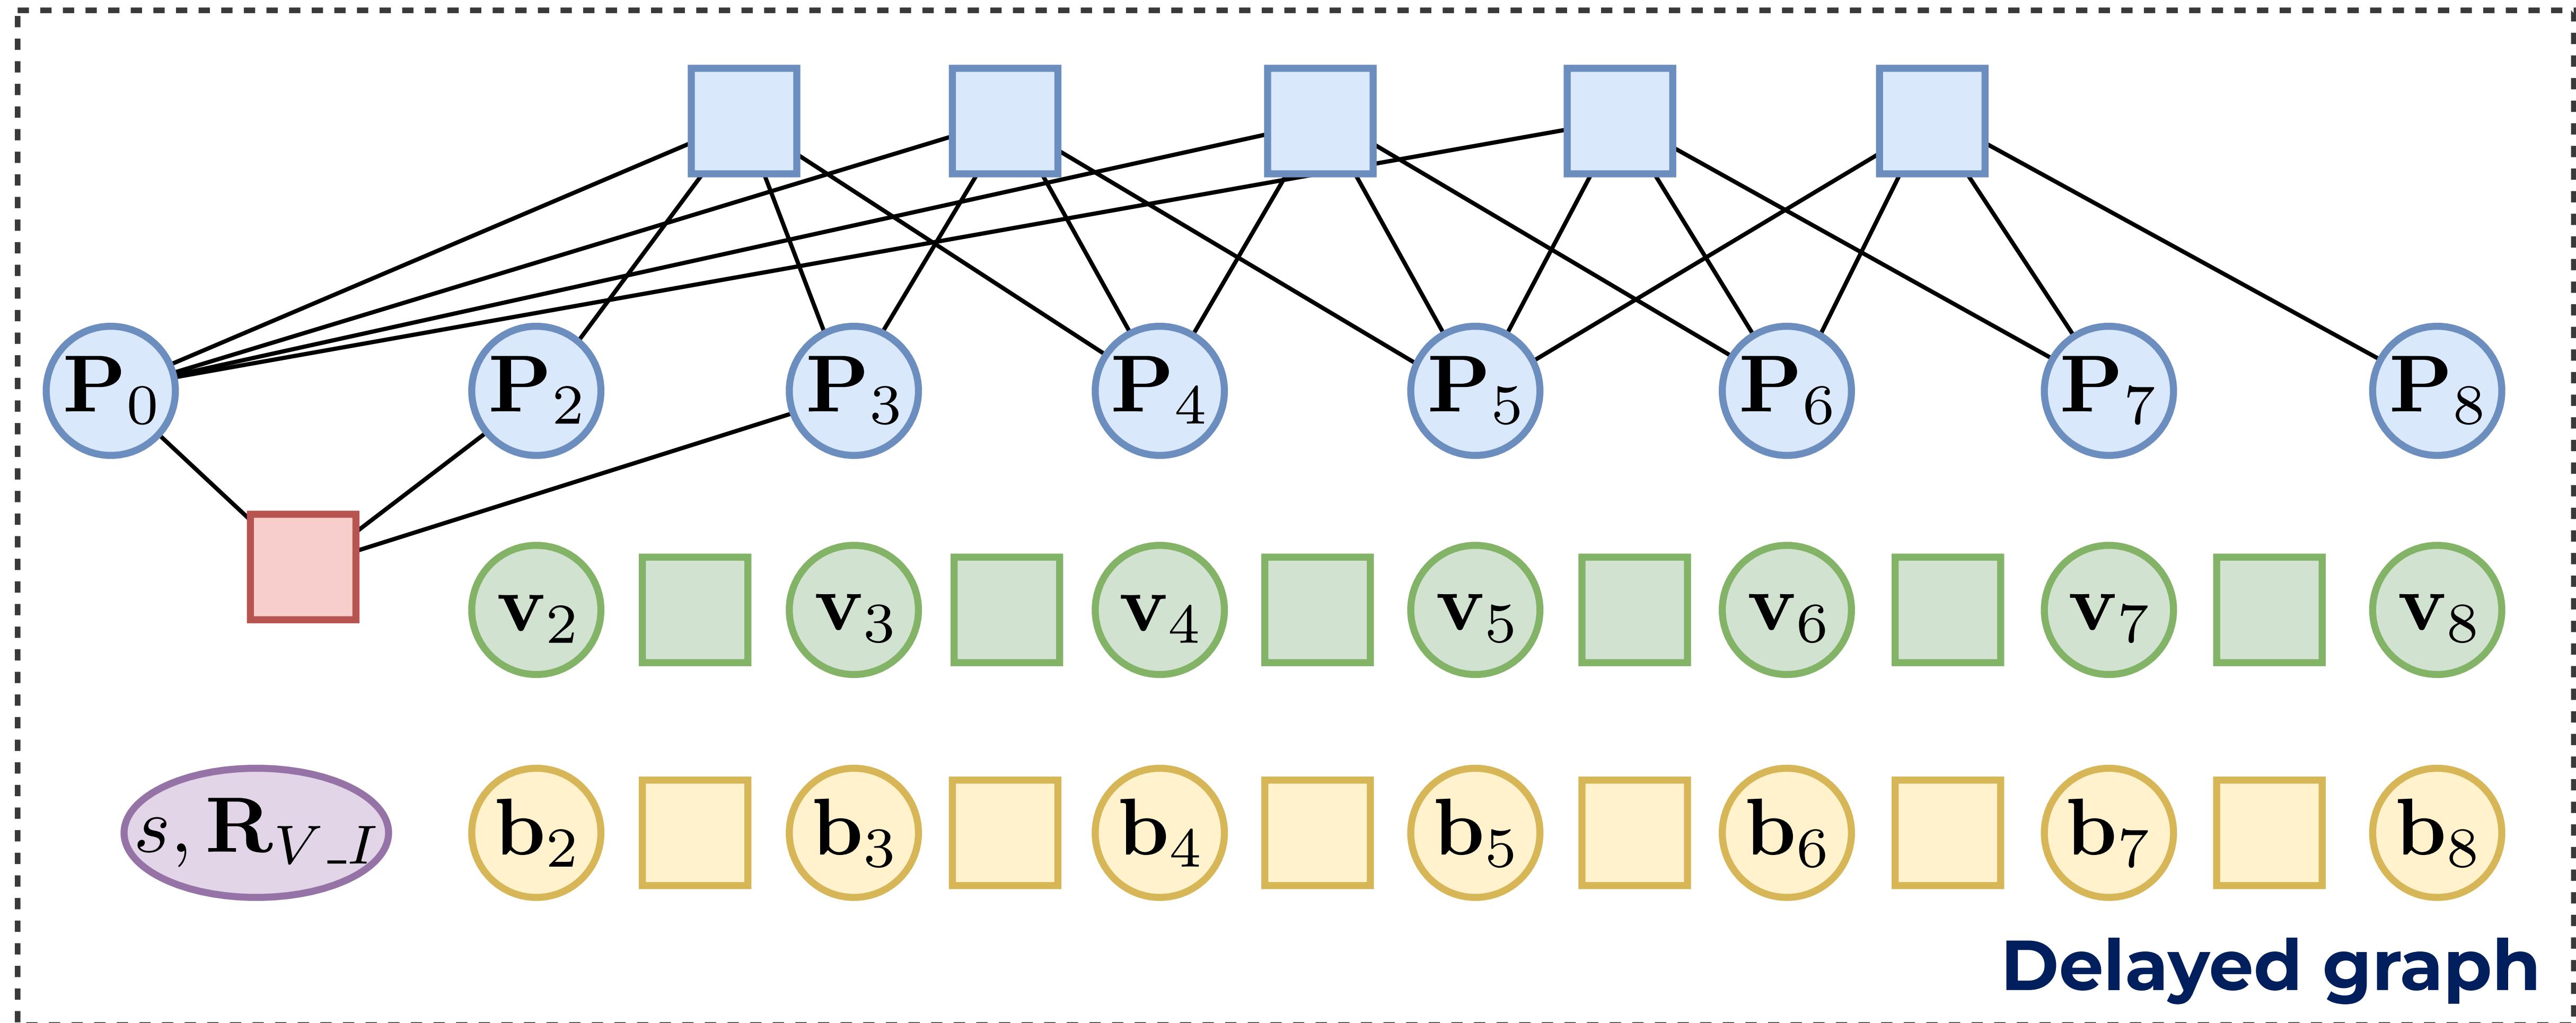
Delayed graph

Pose Graph Bundle Adjustment (PGBA)

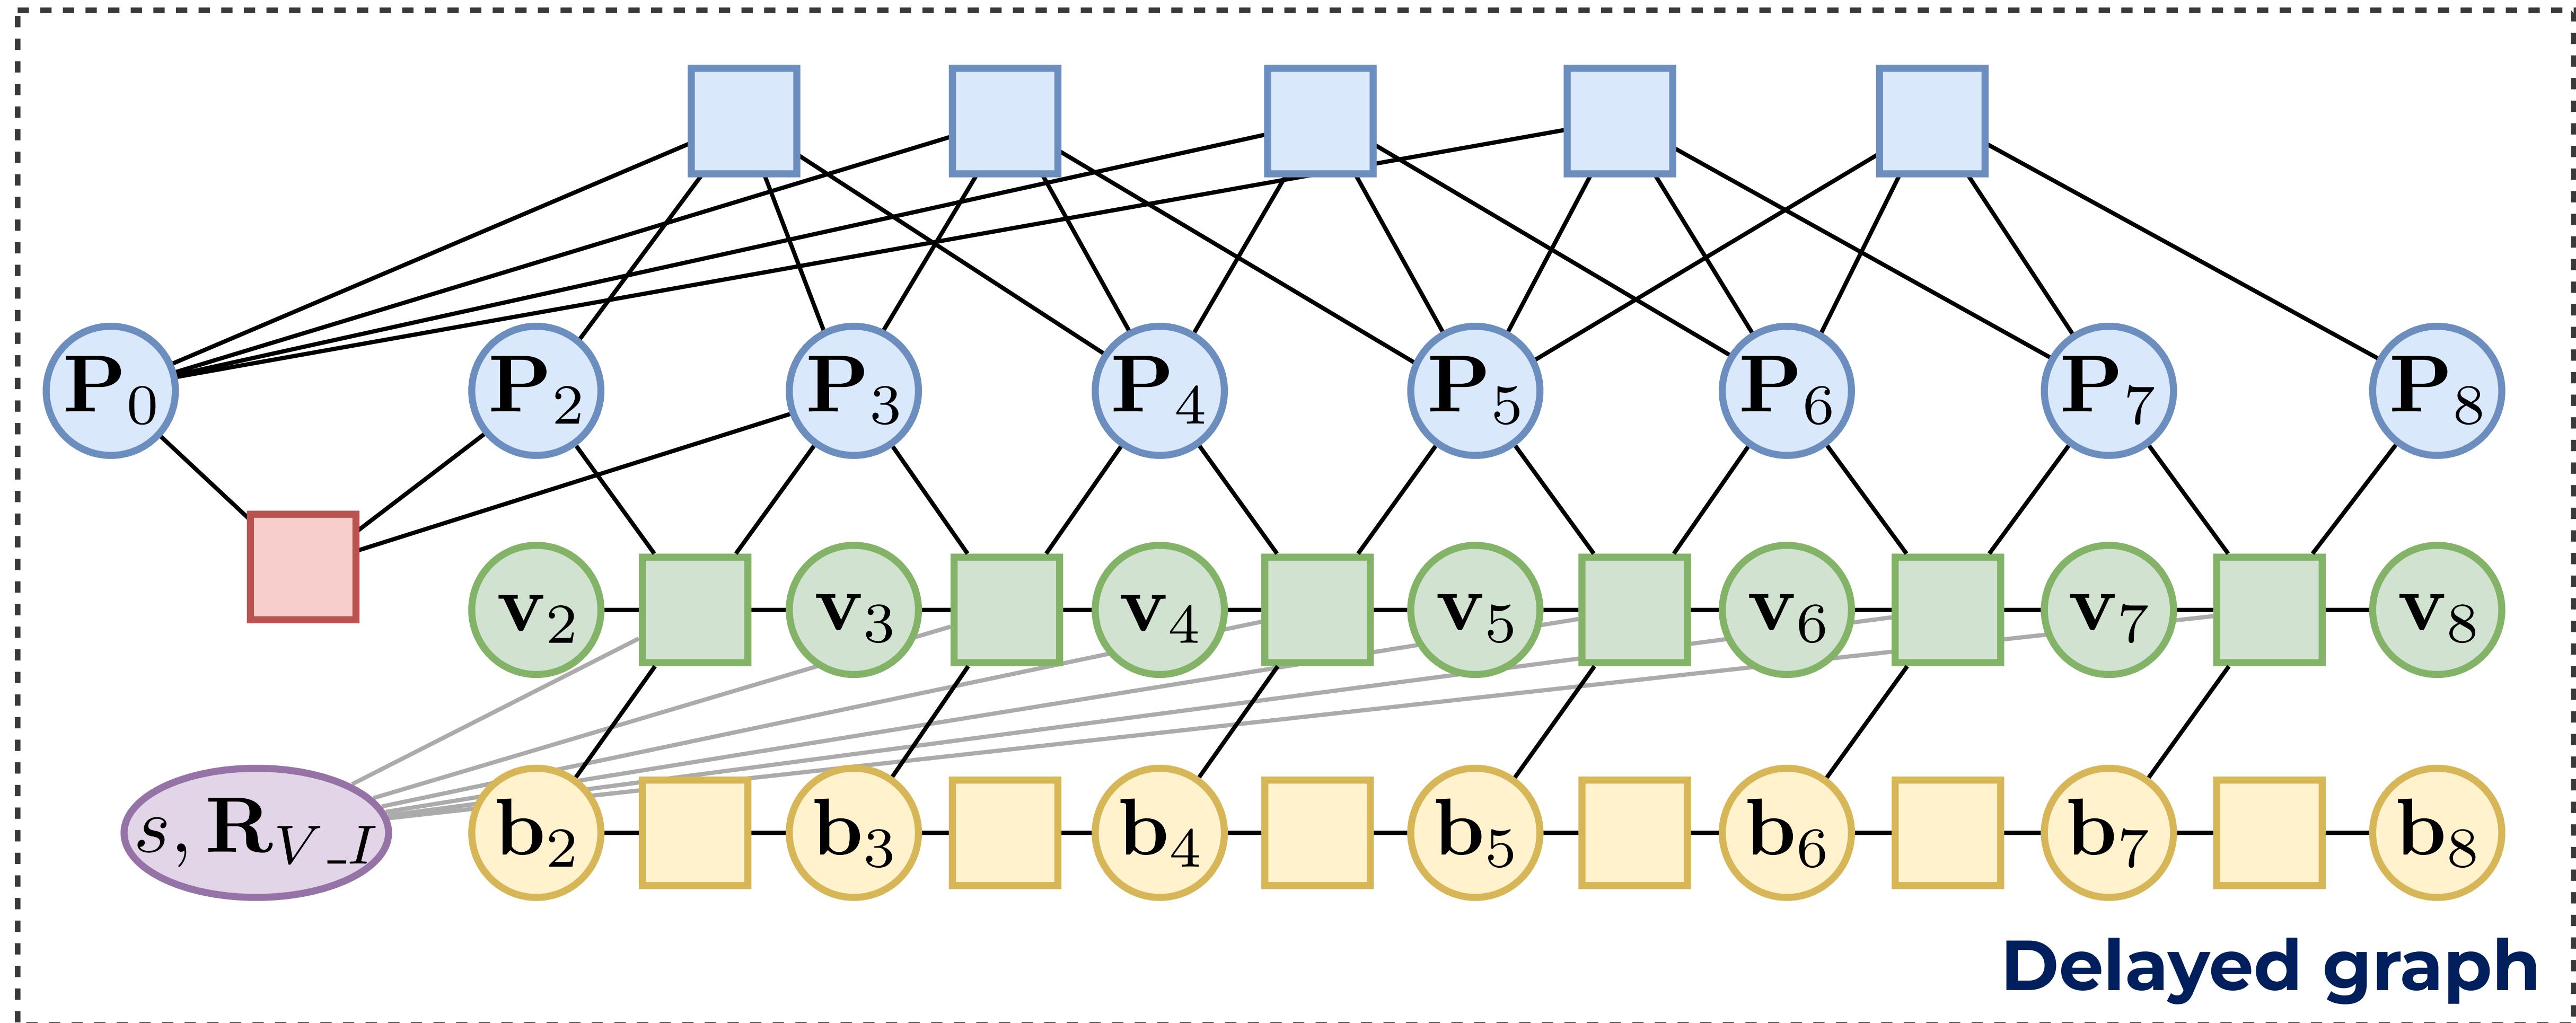


Delayed graph

Pose Graph Bundle Adjustment (PGBA)

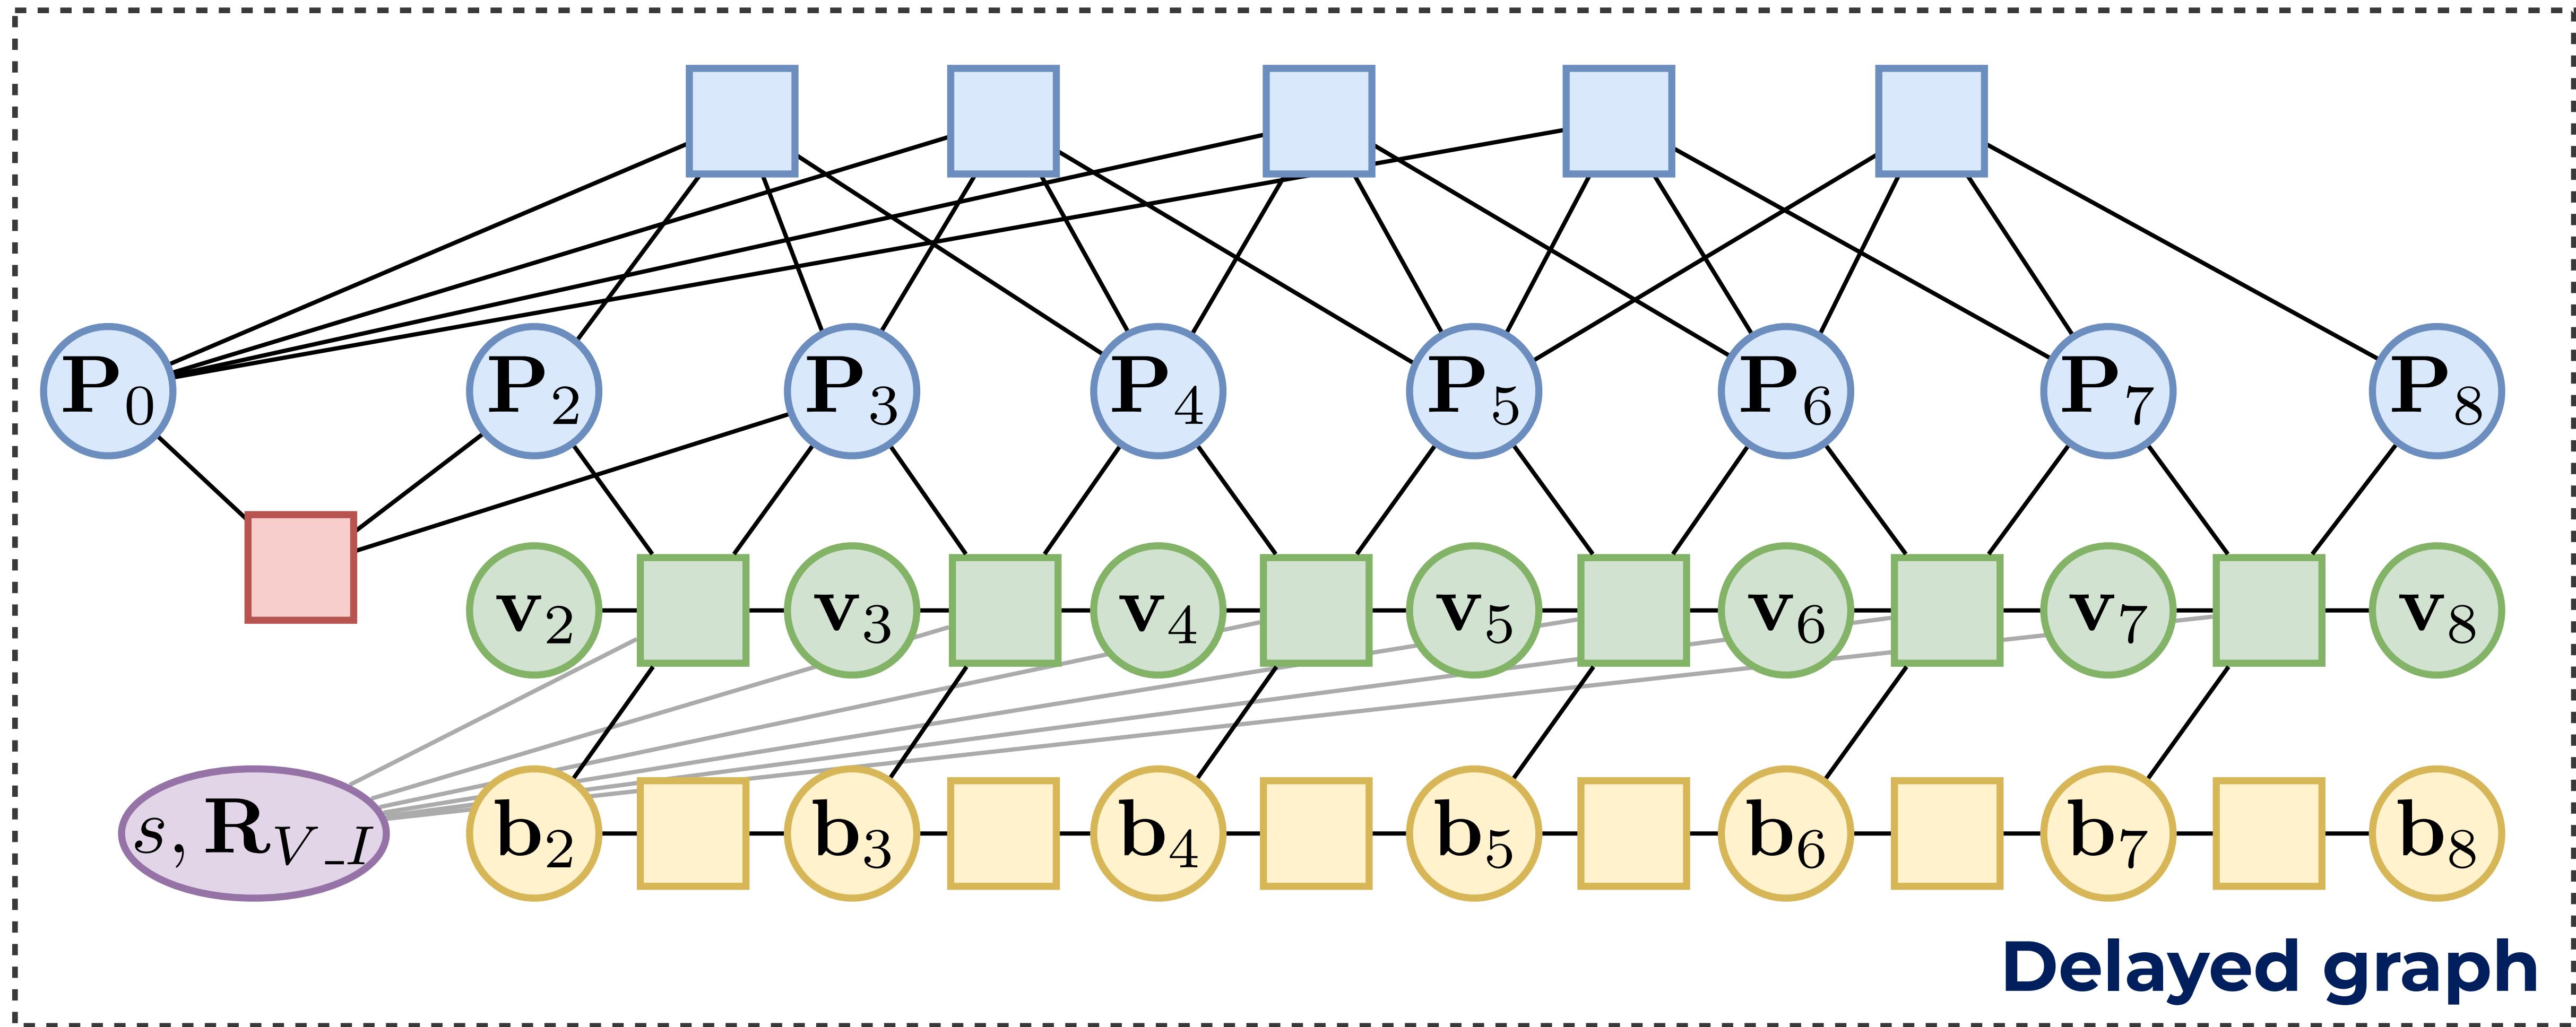


Pose Graph Bundle Adjustment (PGBA)



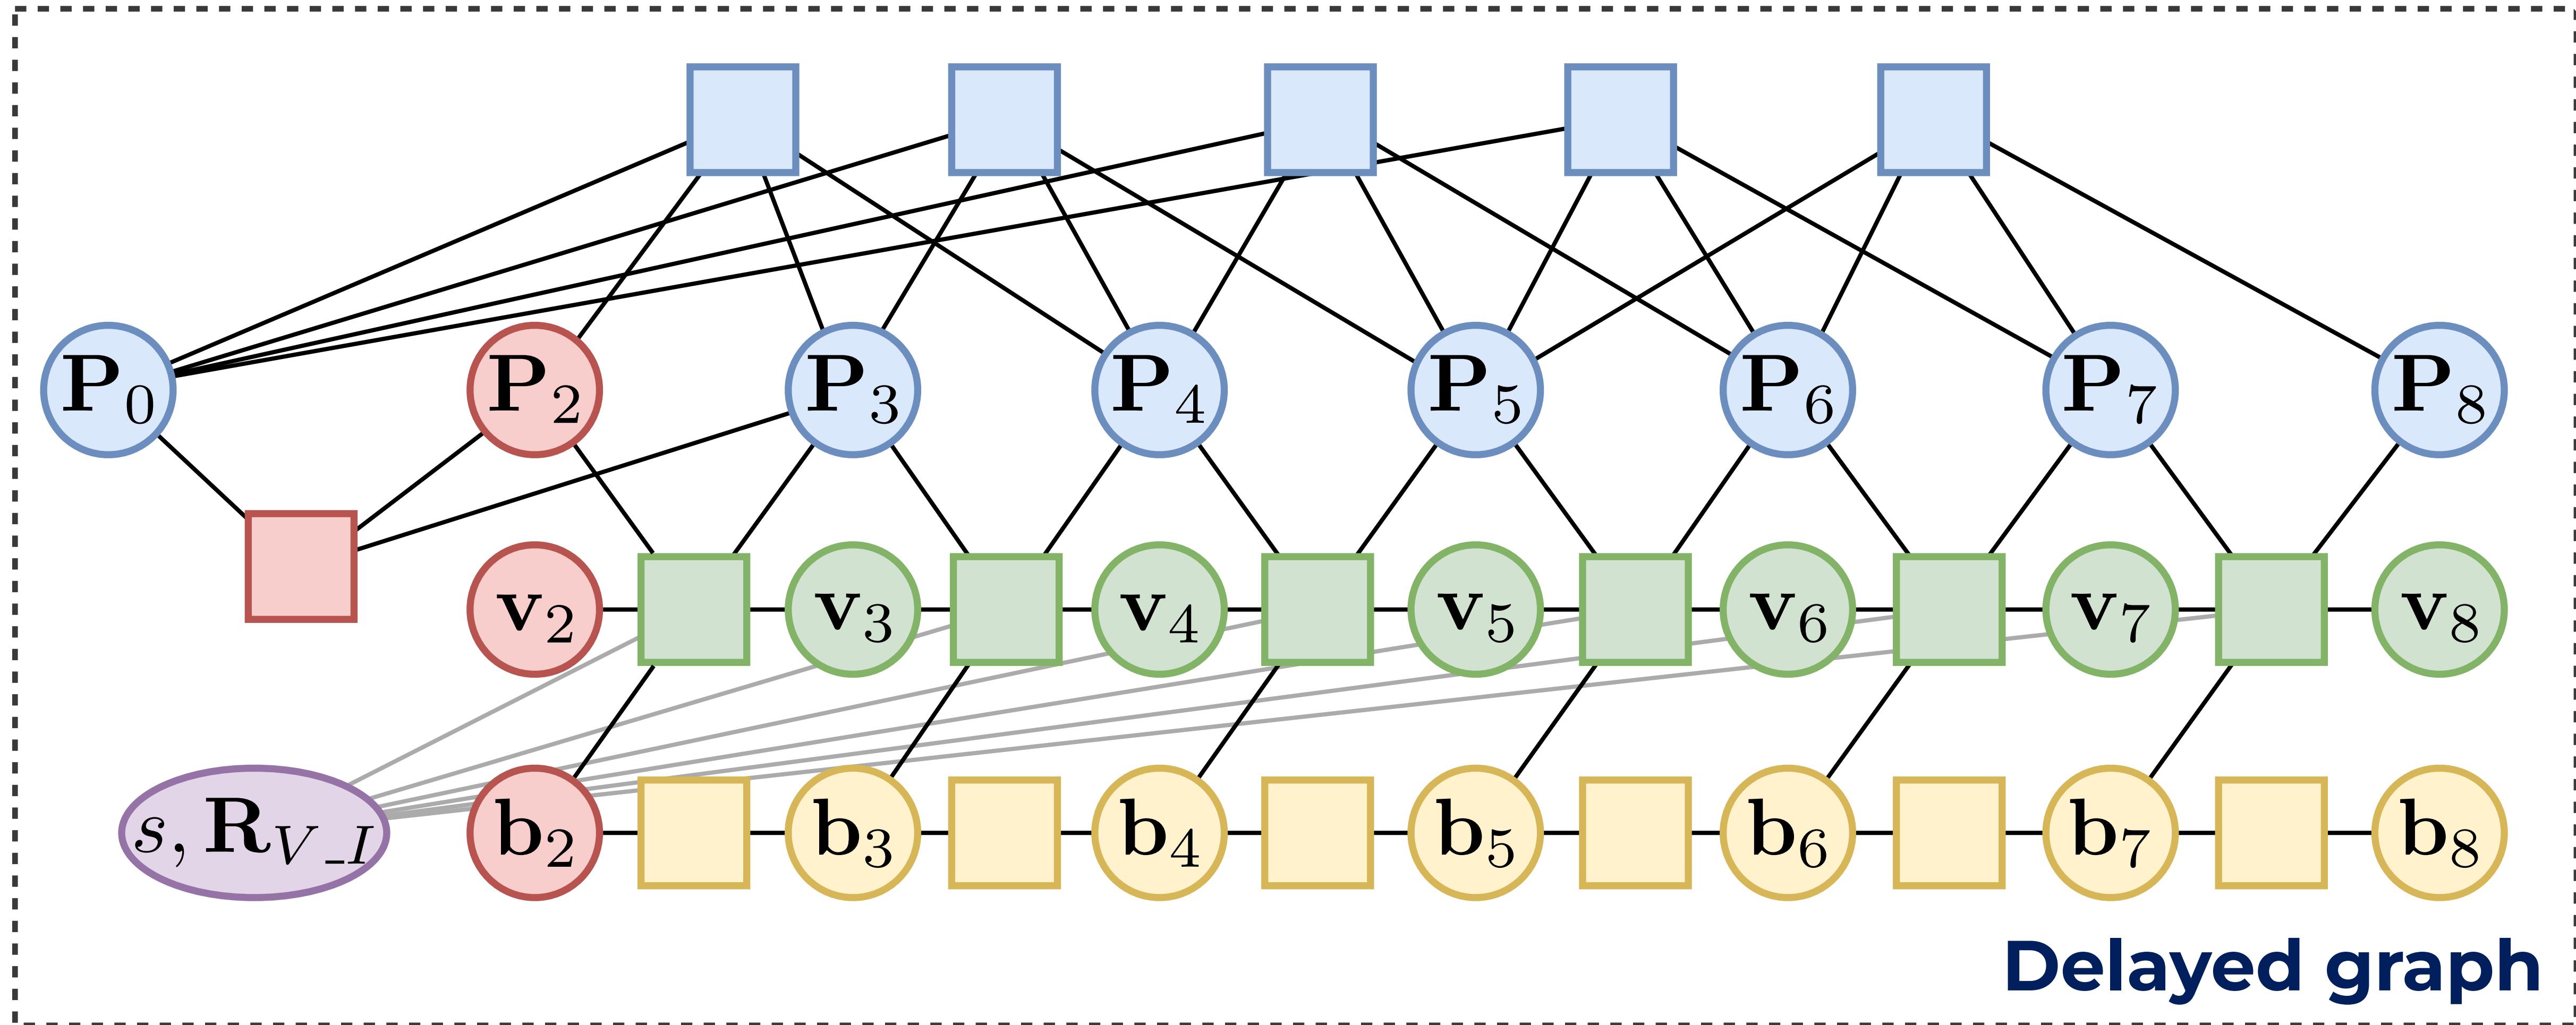
Marginalization order: $P_1 \rightarrow P_2 \rightarrow P_3 \rightarrow P_4 \rightarrow P_0$

Readvancing



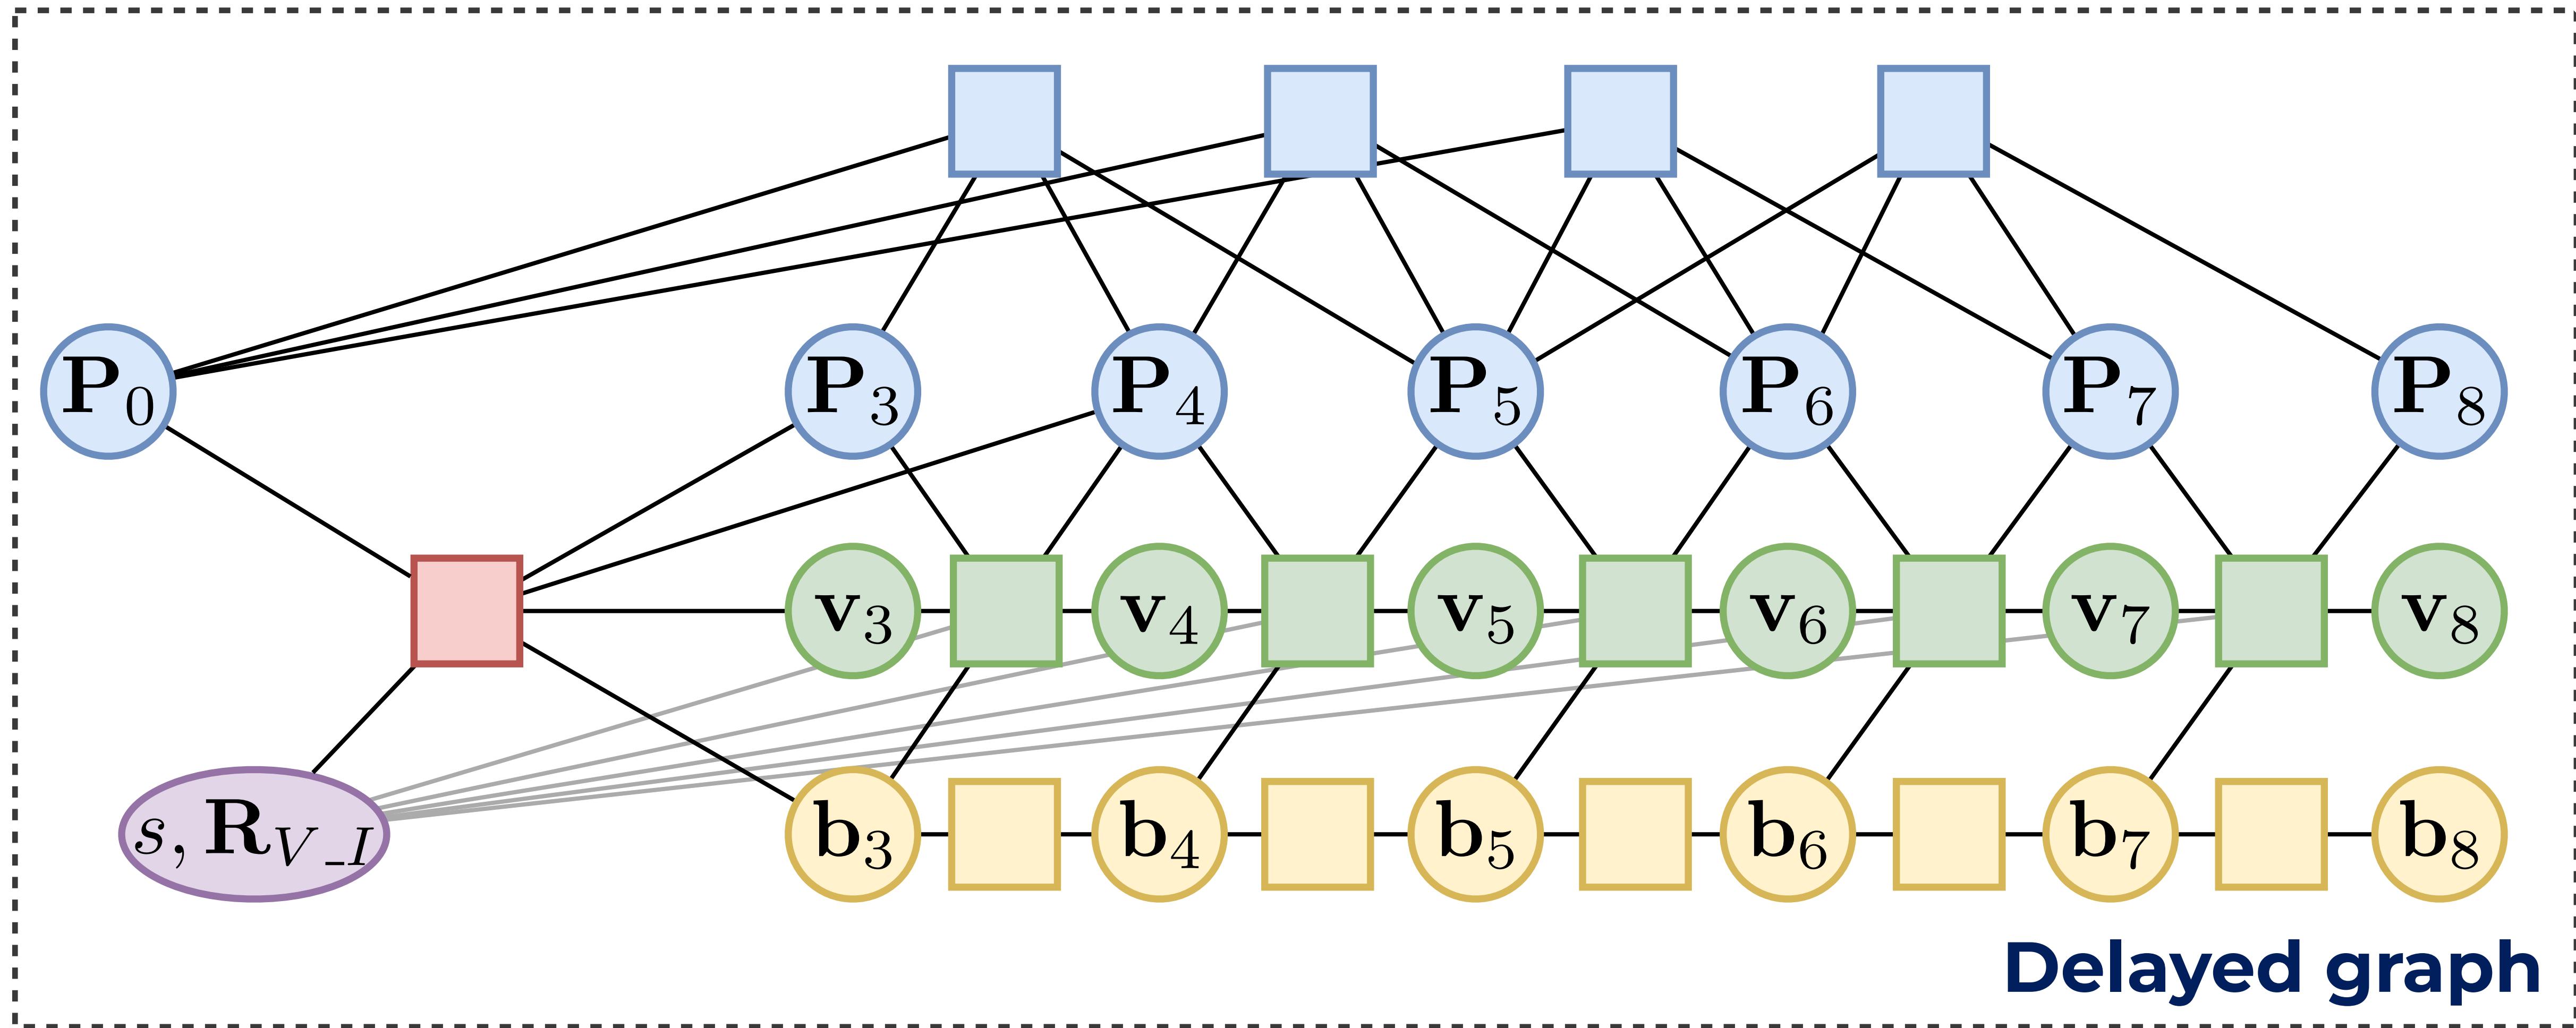
Marginalization order: $P_1 \rightarrow P_2 \rightarrow P_3 \rightarrow P_4 \rightarrow P_0$

Readvancing



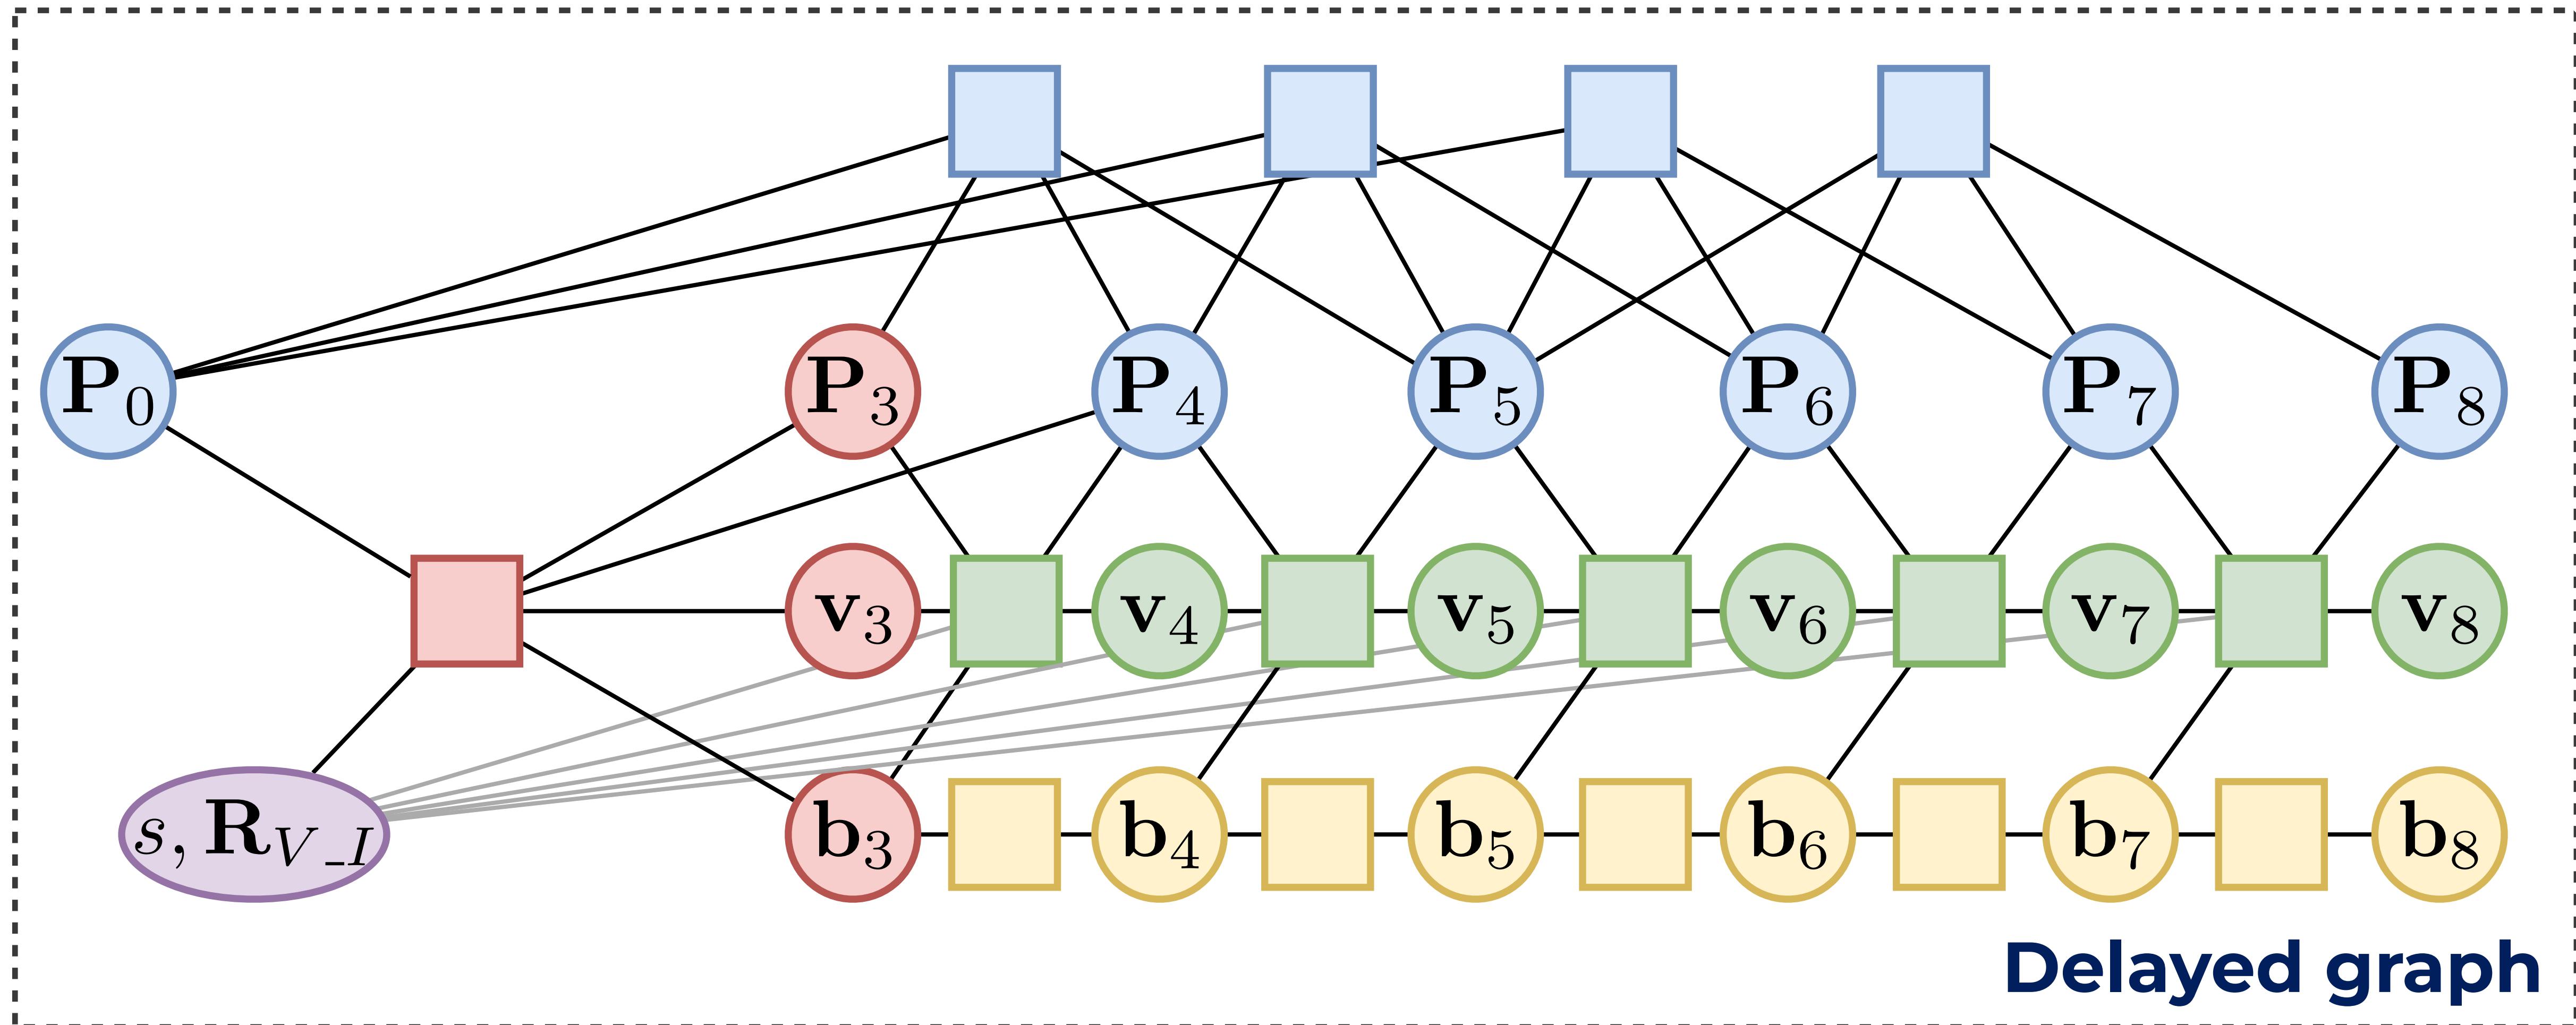
Marginalization order: $P_1 \rightarrow P_2 \rightarrow P_3 \rightarrow P_4 \rightarrow P_0$

Readvancing



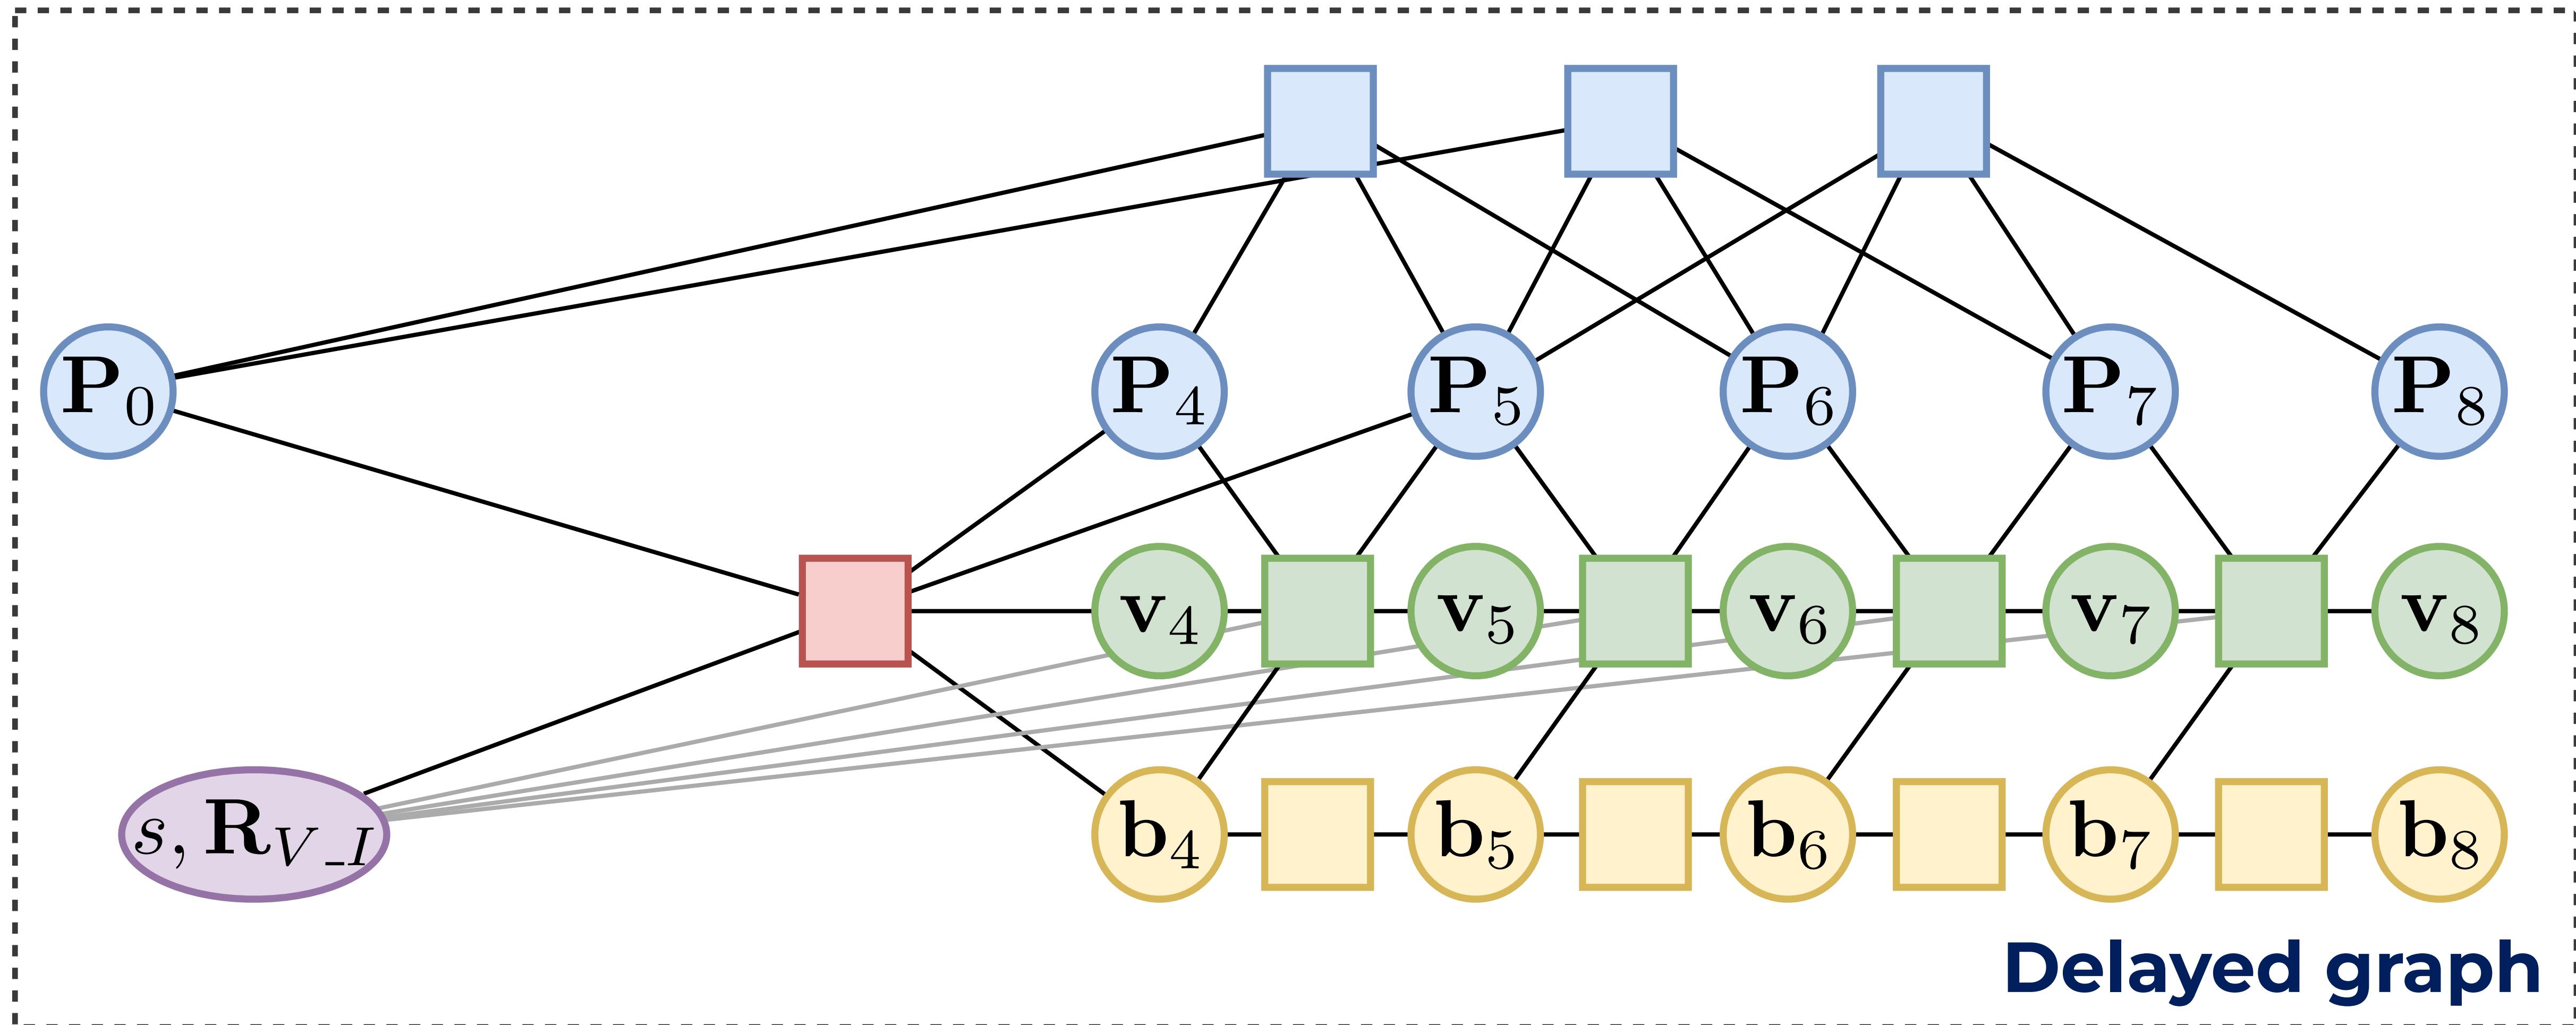
Marginalization order: $P_1 \rightarrow P_2 \rightarrow P_3 \rightarrow P_4 \rightarrow P_0$

Readvancing



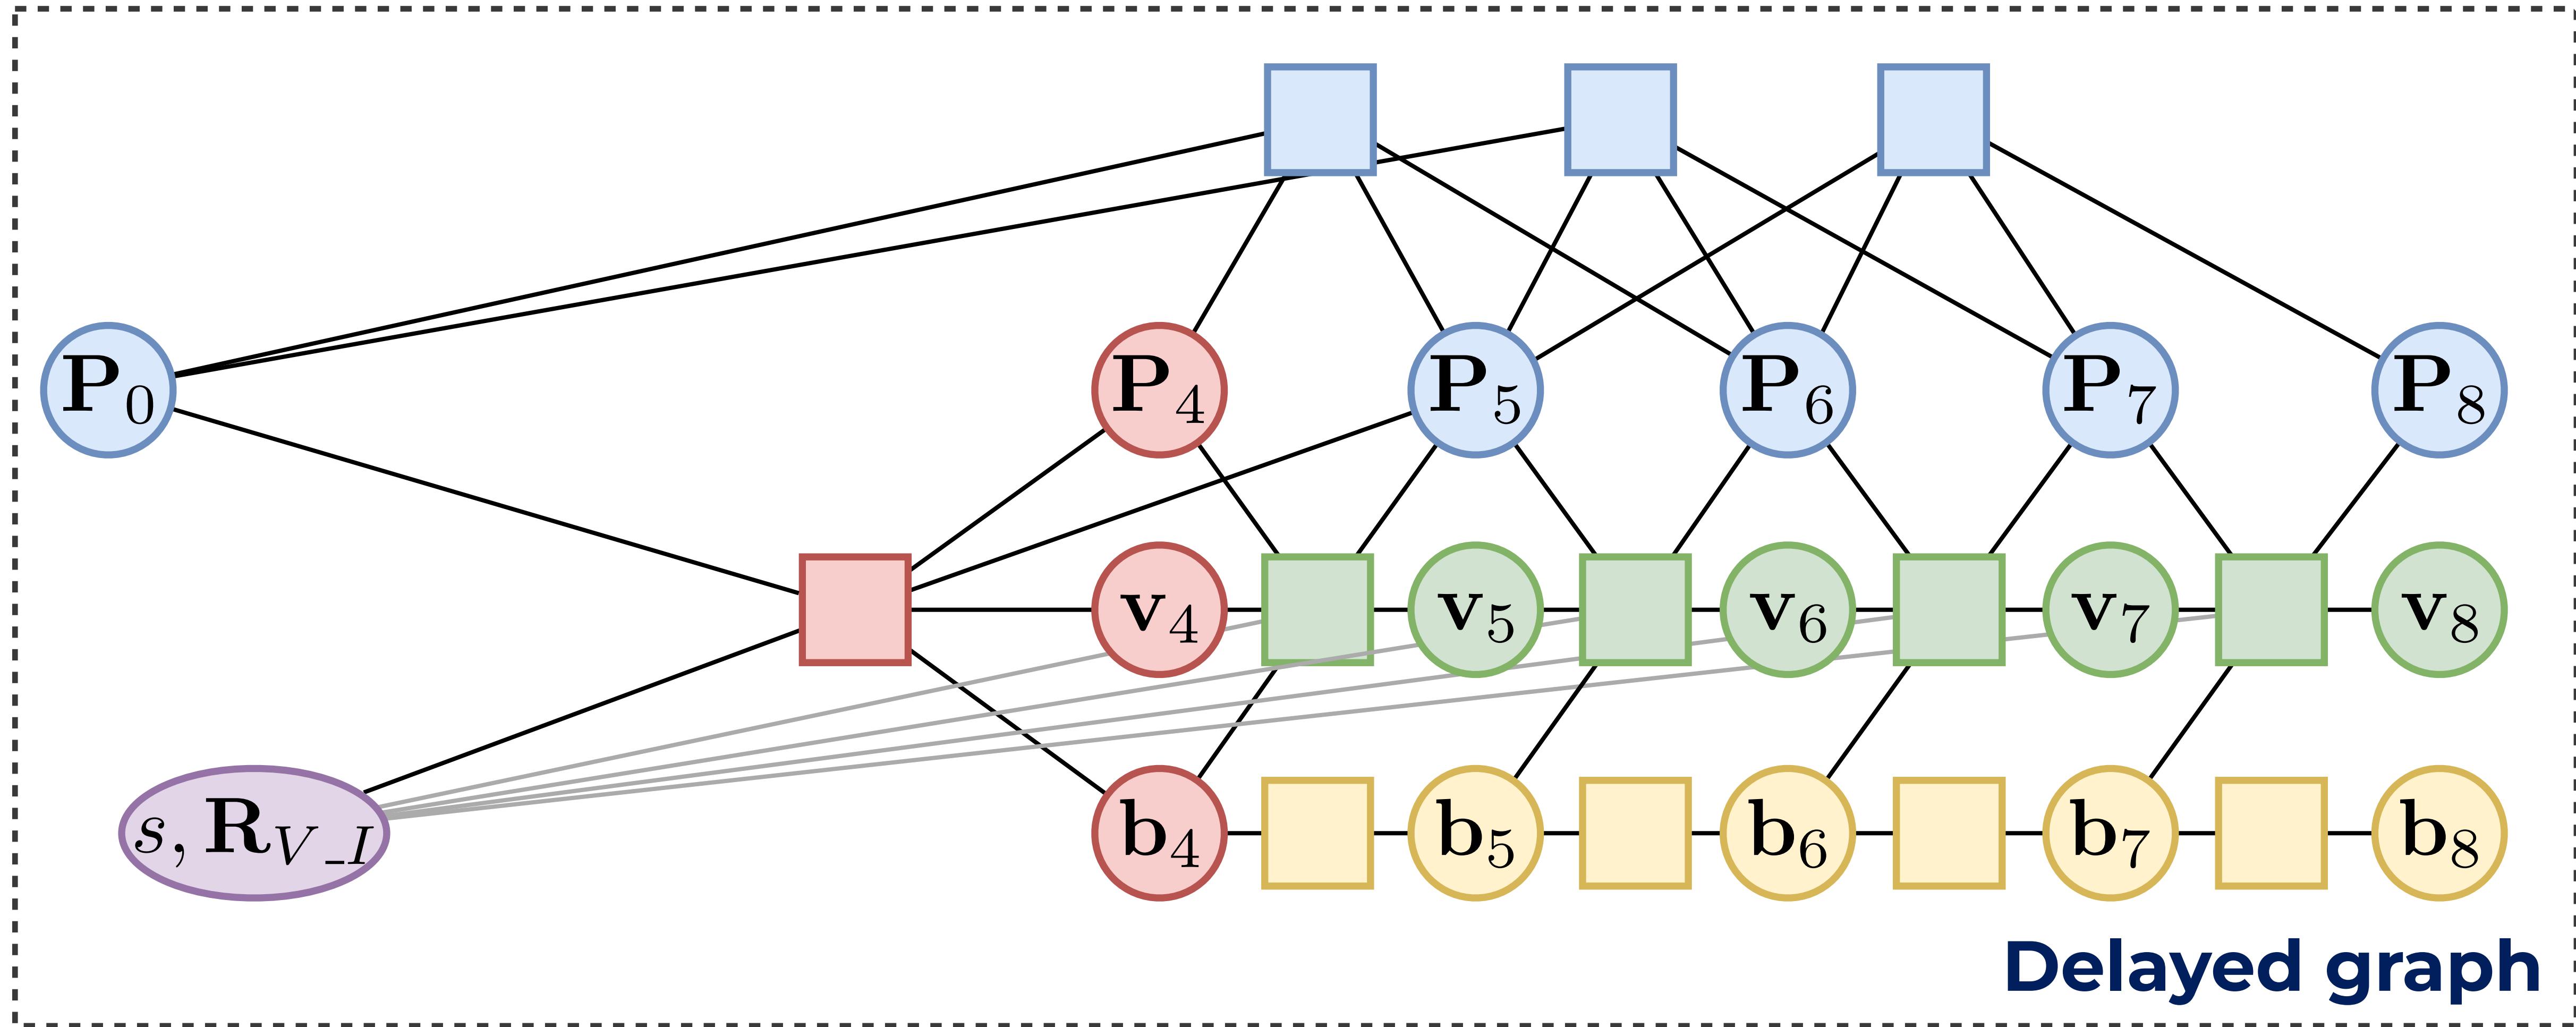
Marginalization order: $P_1 \rightarrow P_2 \rightarrow P_3 \rightarrow P_4 \rightarrow P_0$

Readvancing



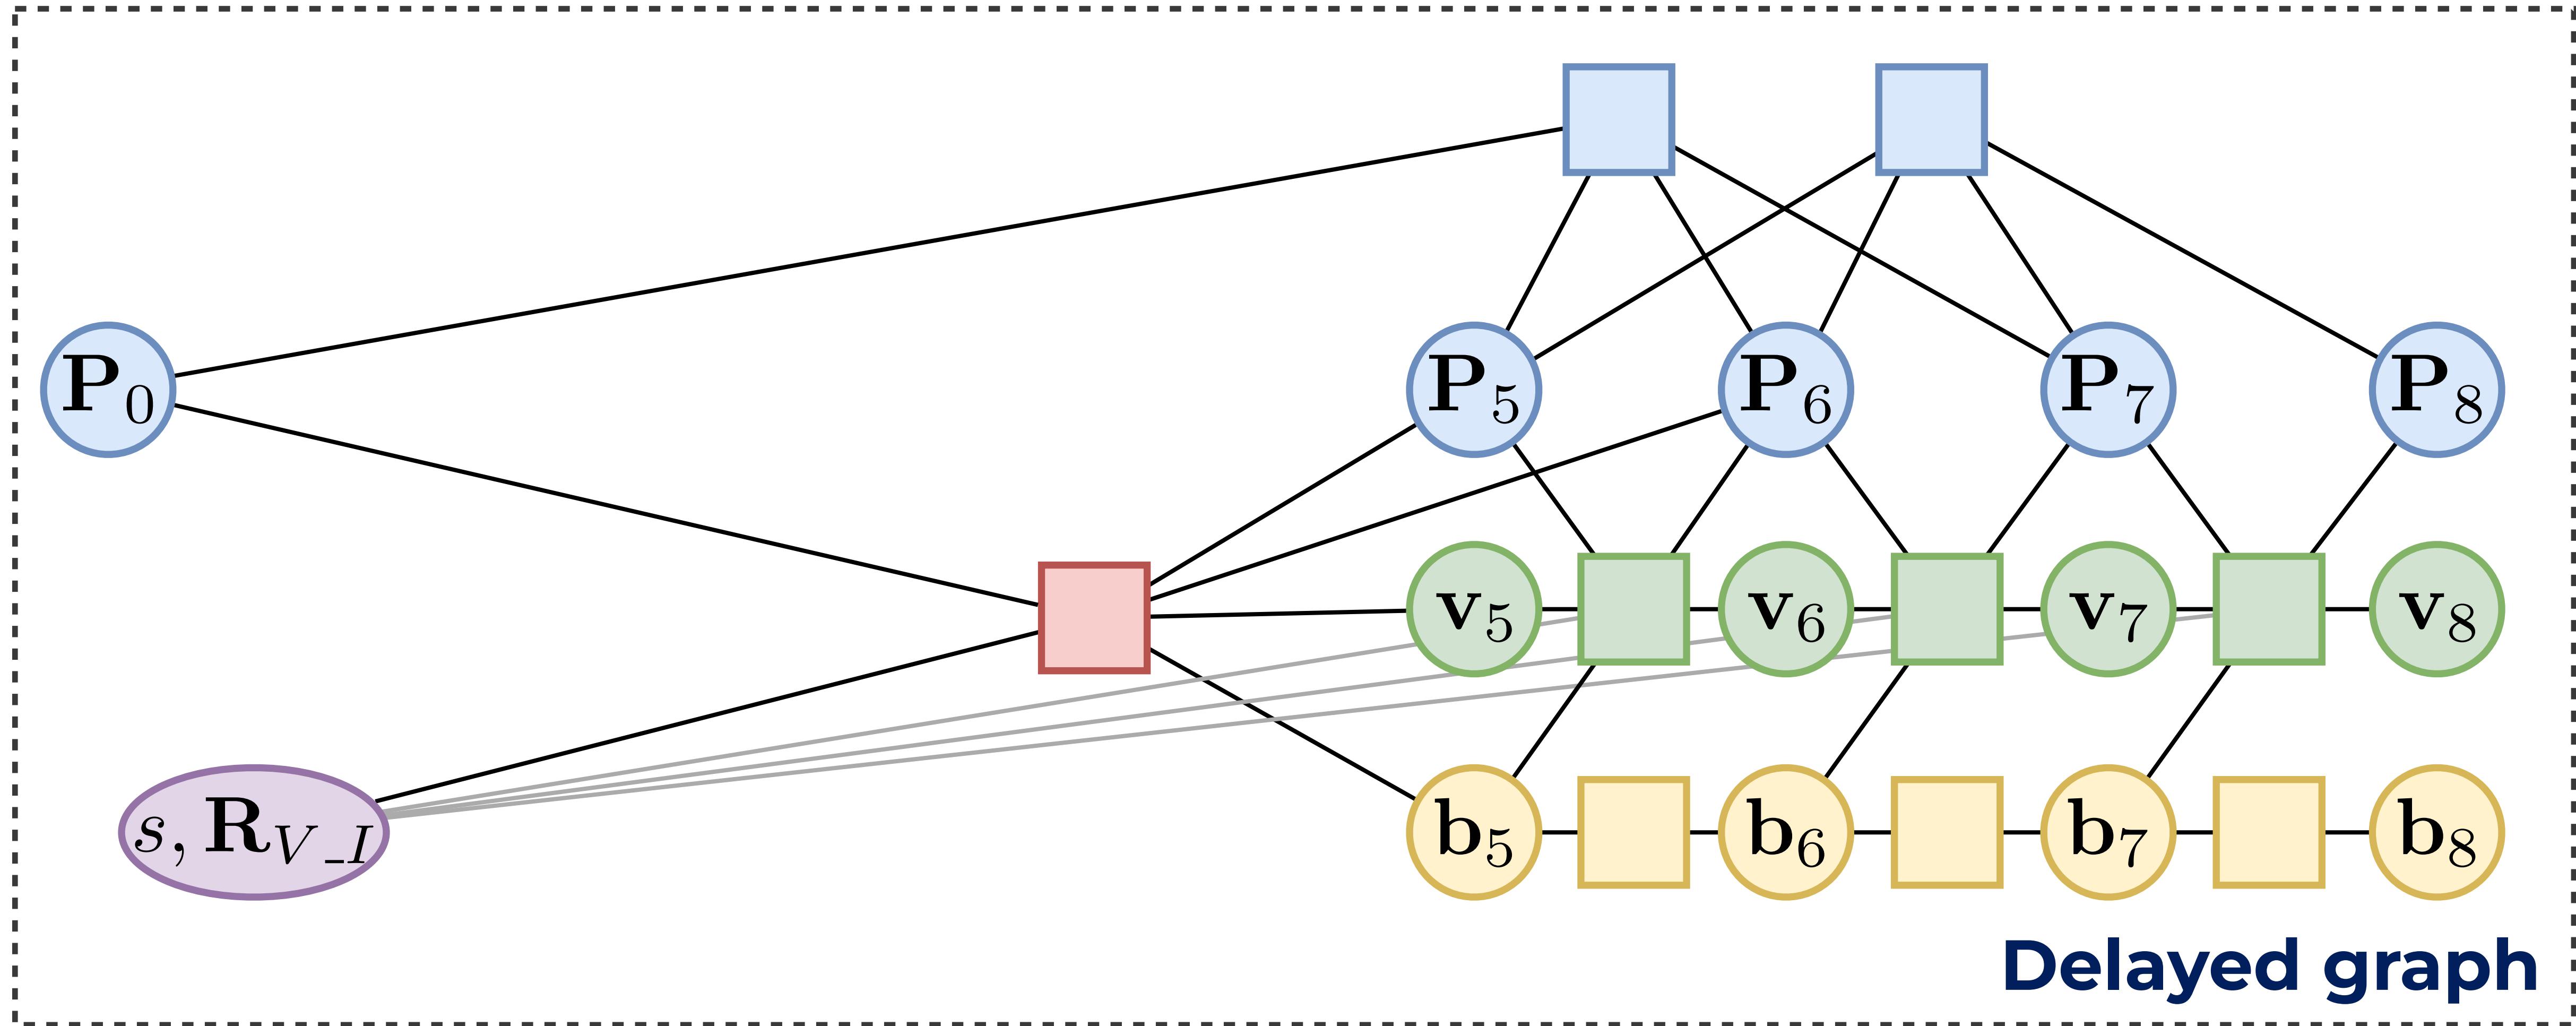
Marginalization order: $P_1 \rightarrow P_2 \rightarrow P_3 \rightarrow P_4 \rightarrow P_0$

Readvancing



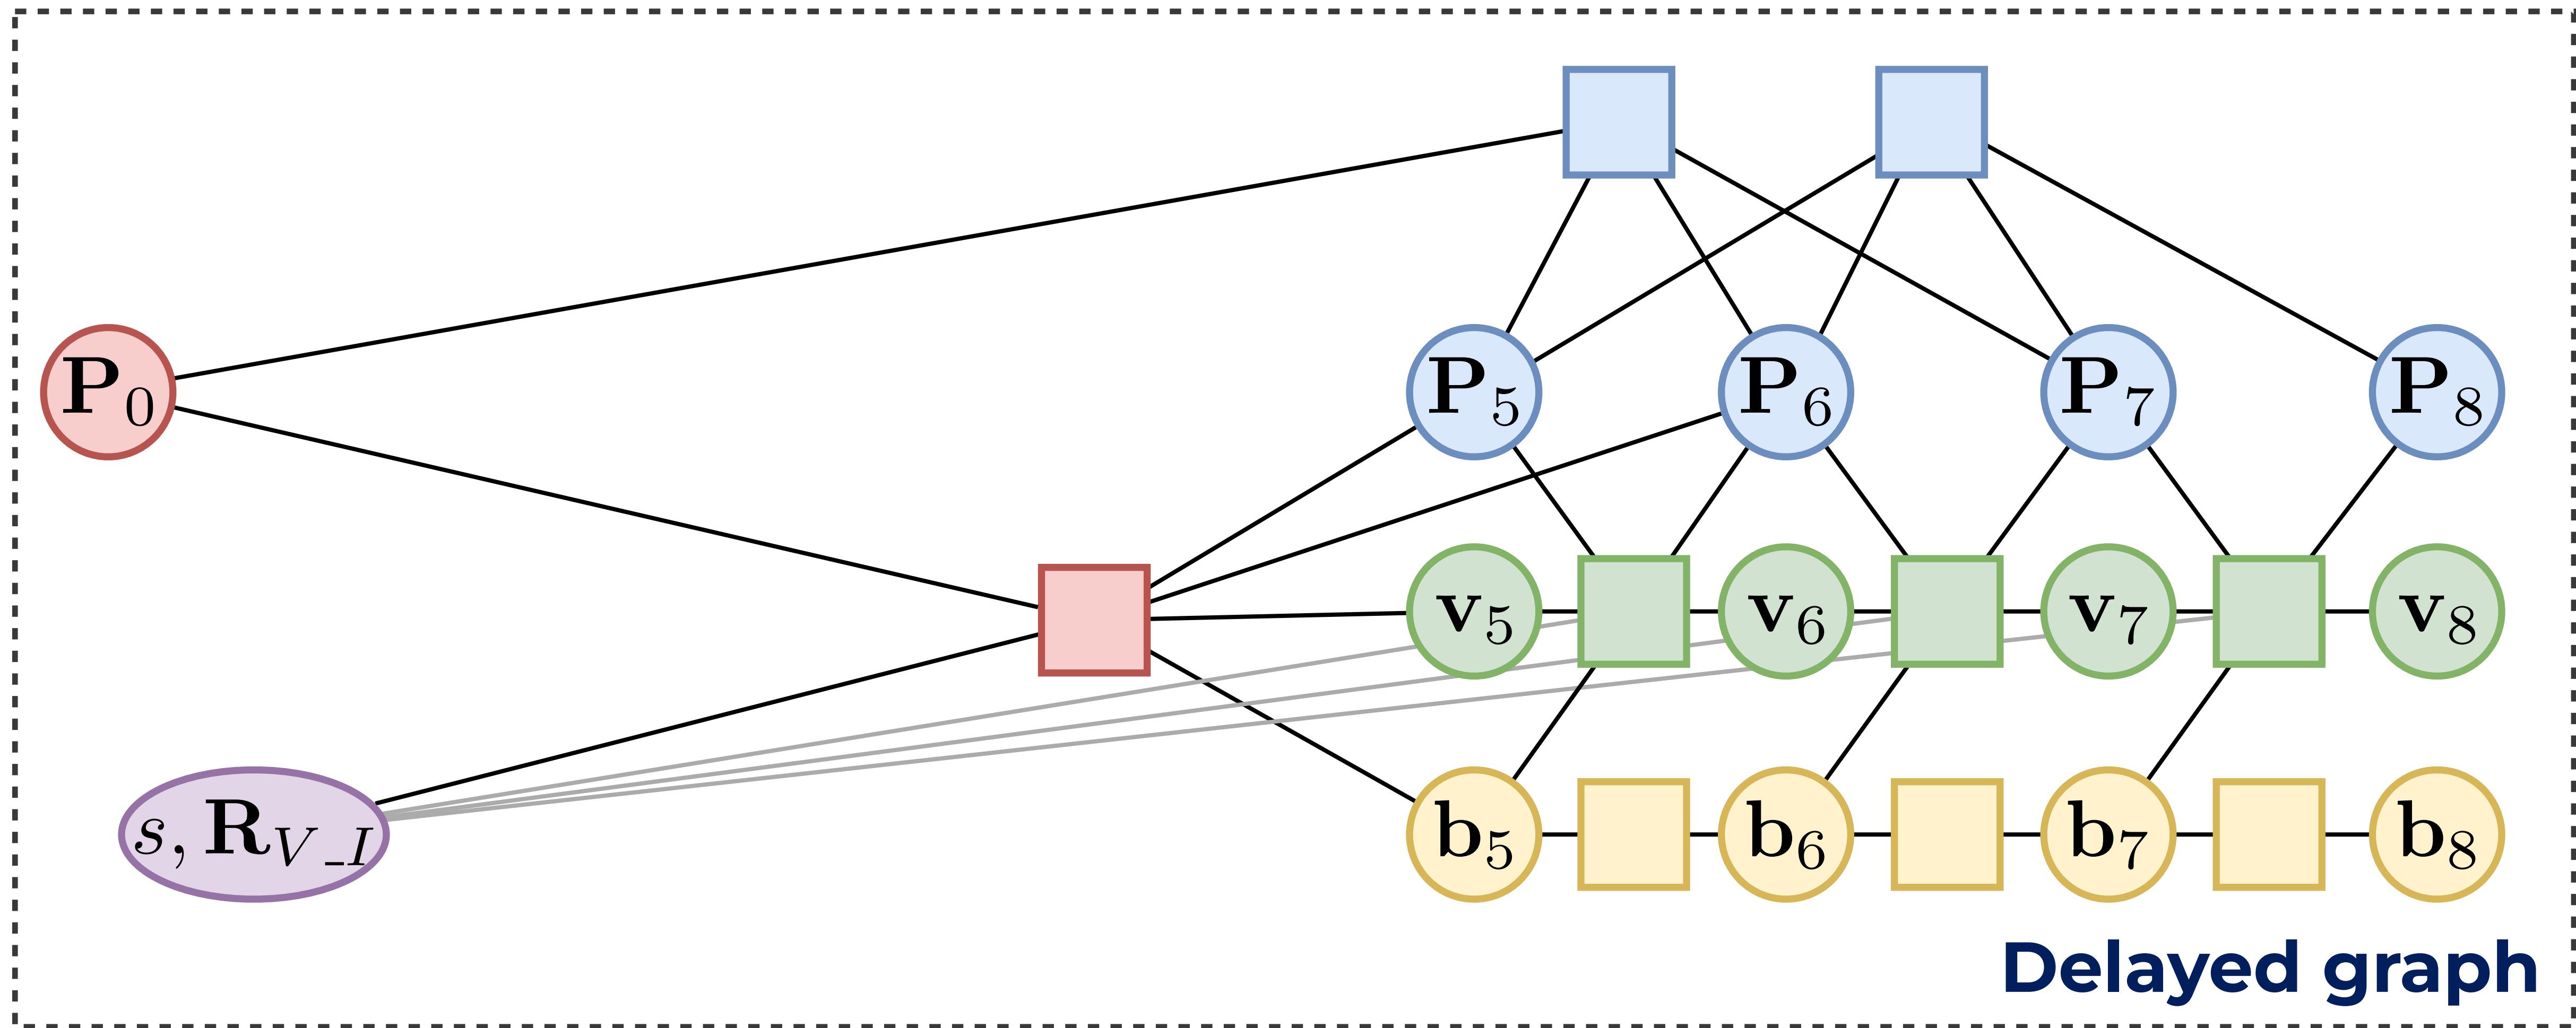
Marginalization order: $P_1 \rightarrow P_2 \rightarrow P_3 \rightarrow P_4 \rightarrow P_0$

Readvancing



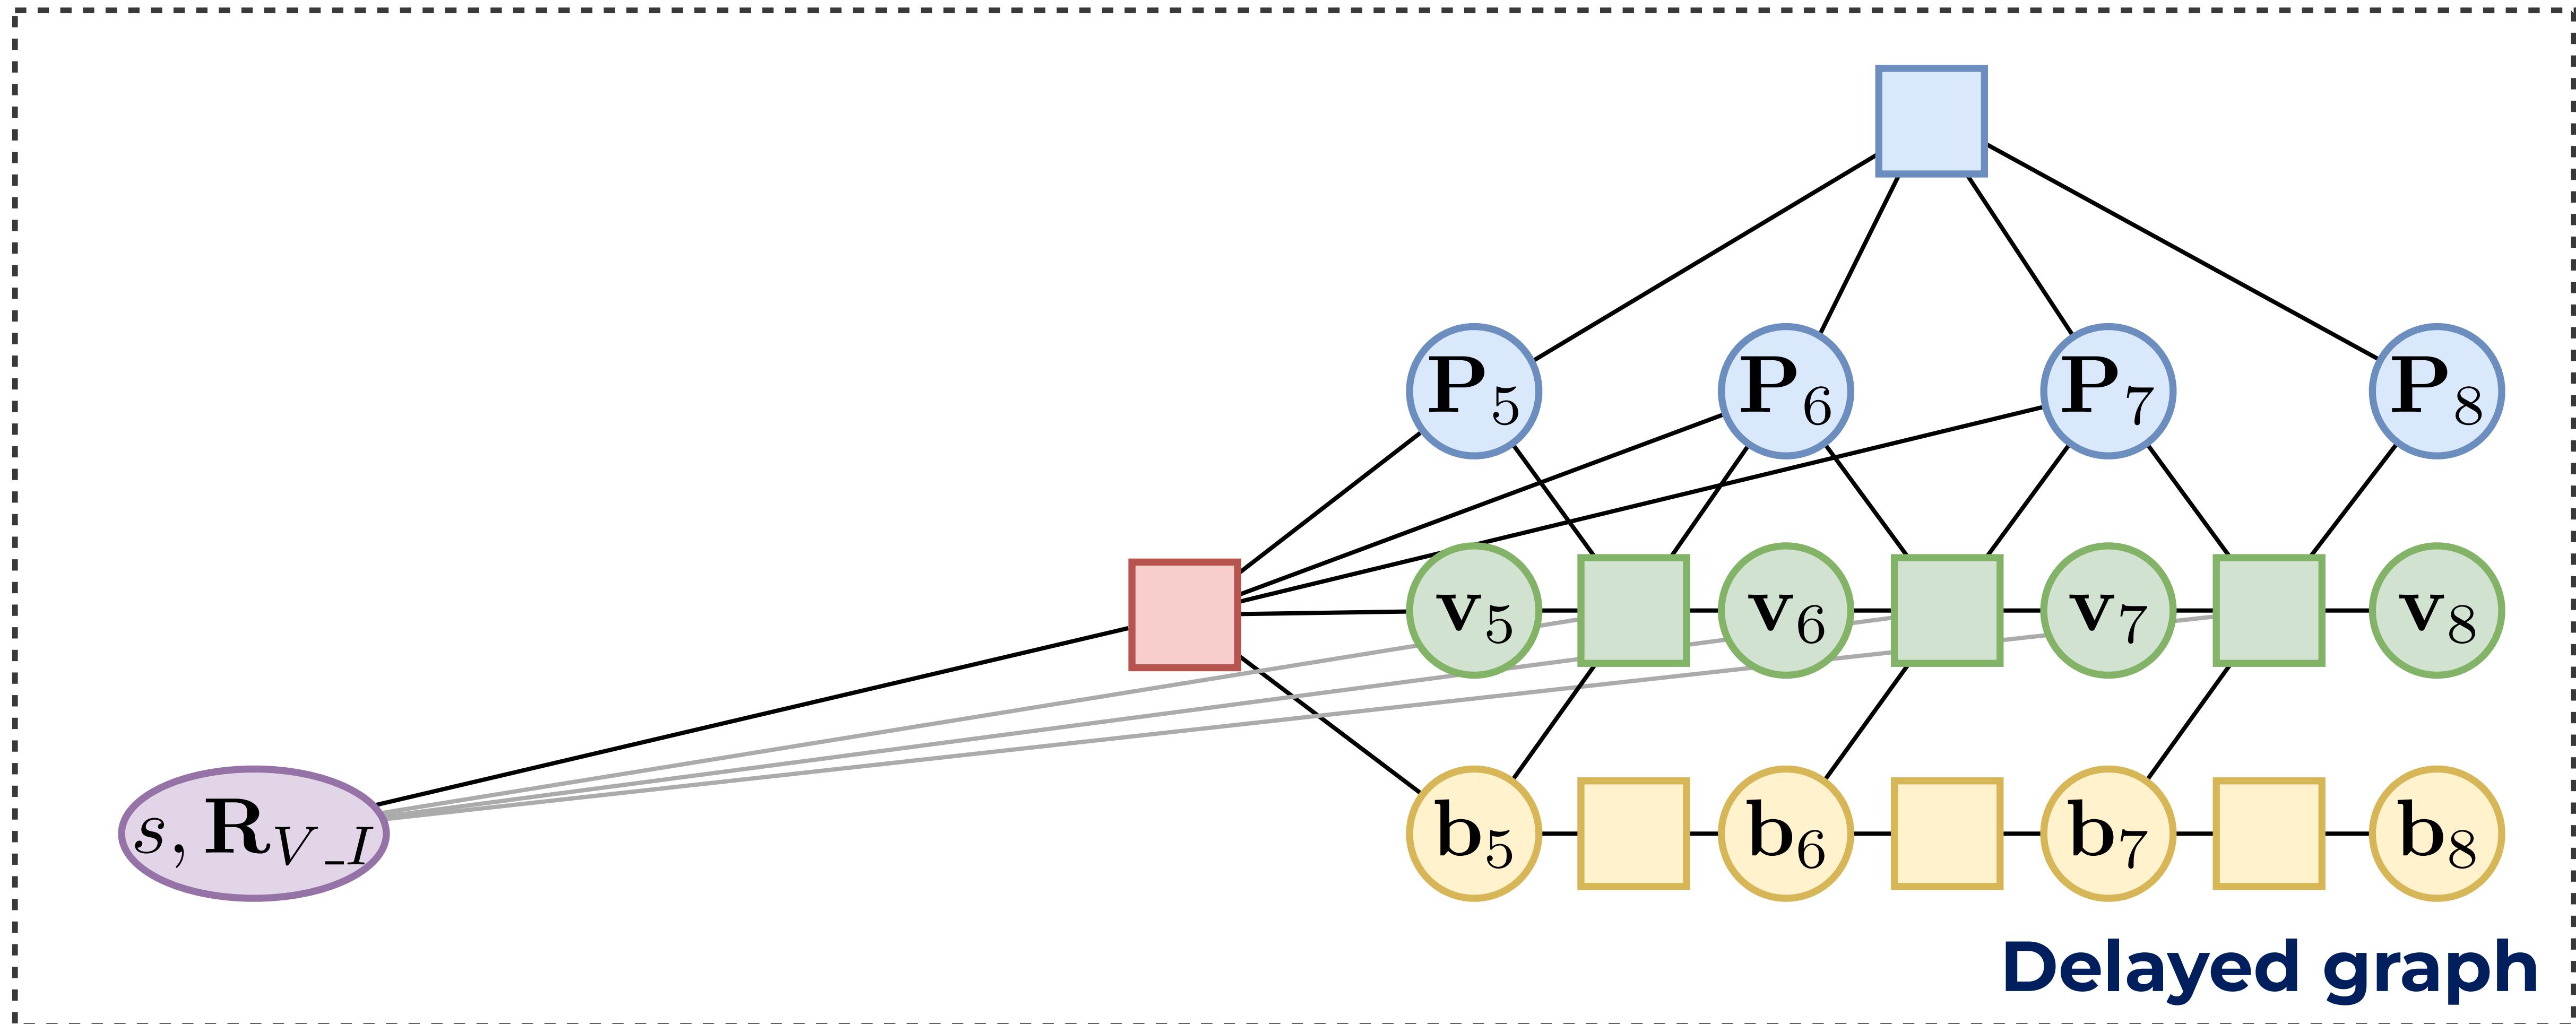
Marginalization order: $P_1 \rightarrow P_2 \rightarrow P_3 \rightarrow P_4 \rightarrow P_0$

Readvancing



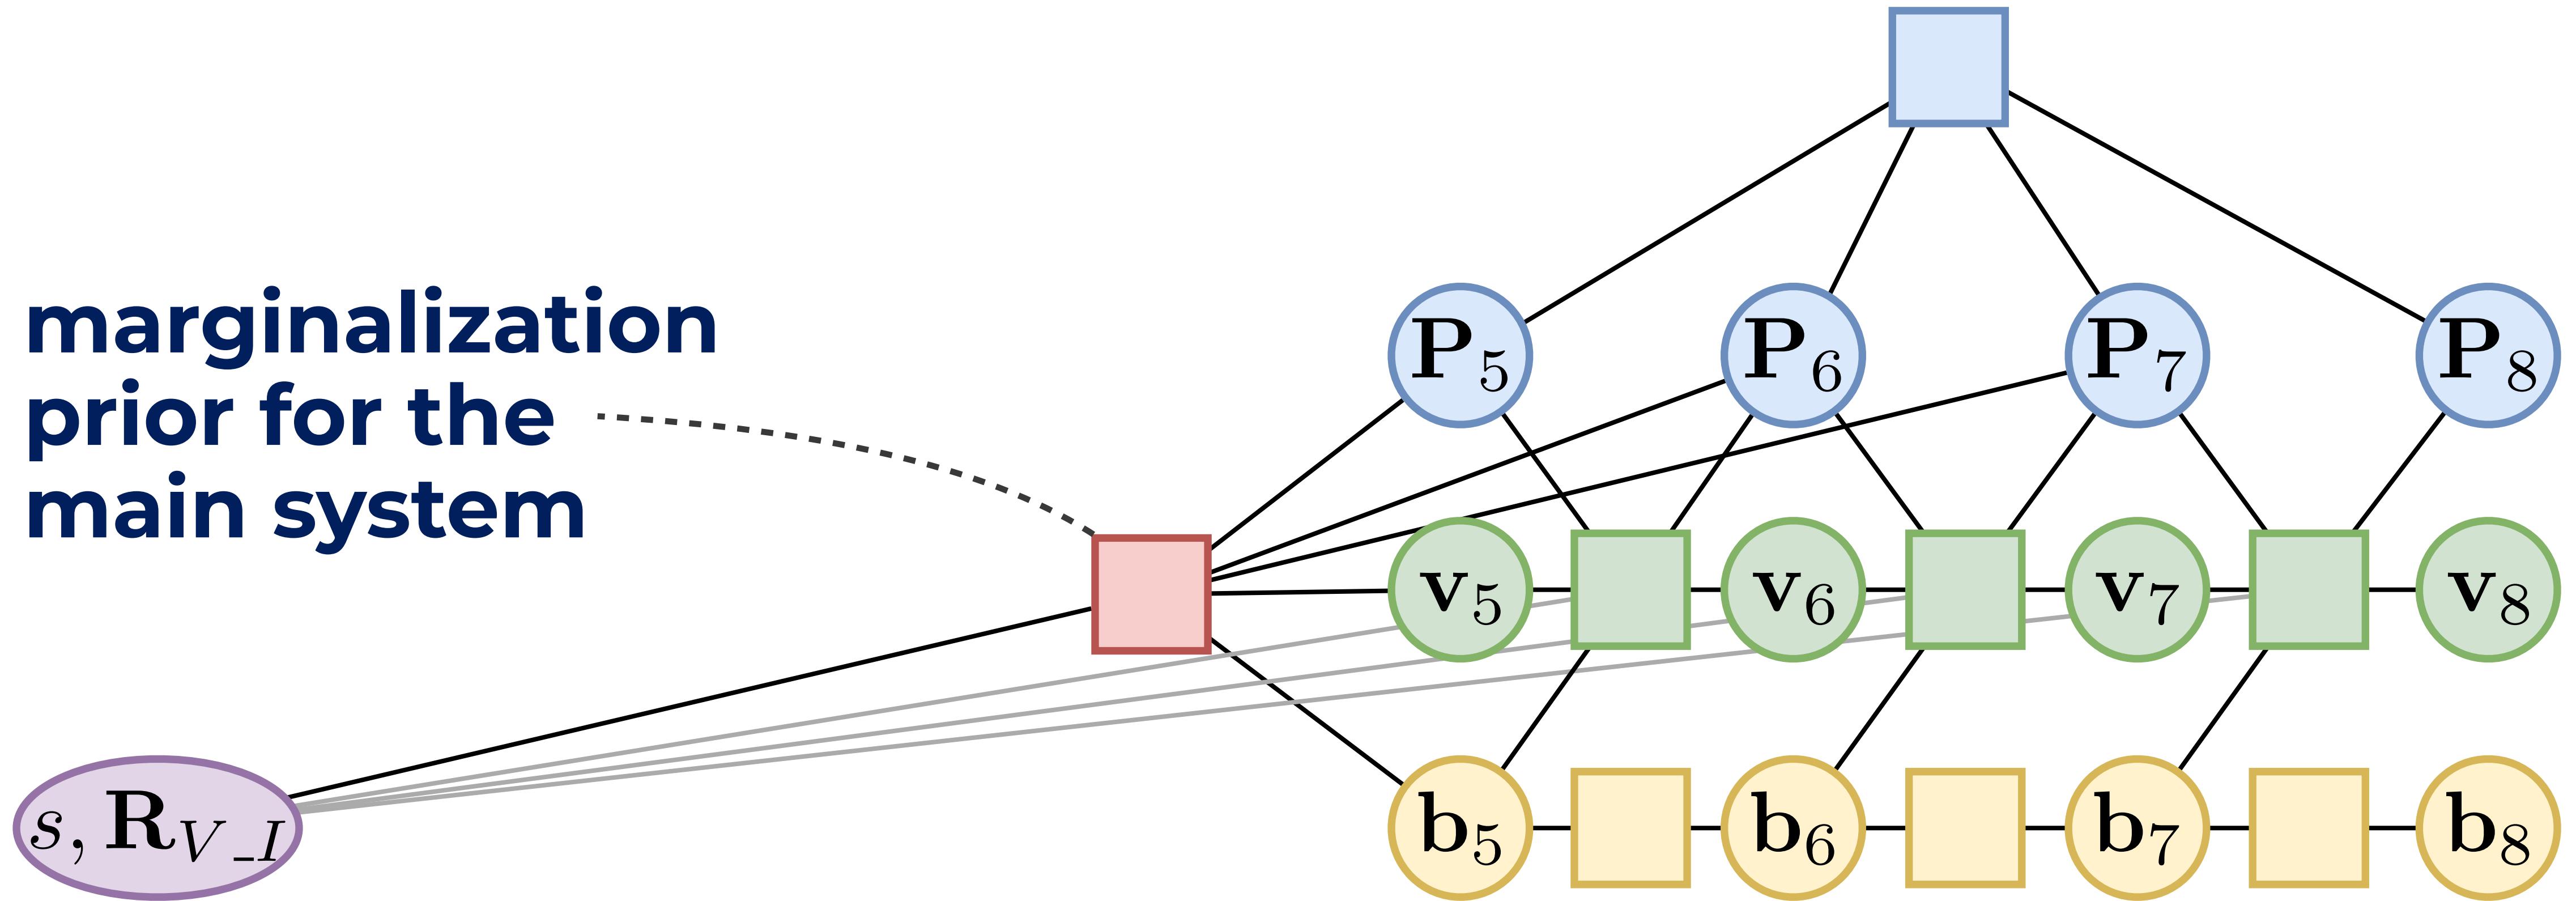
Marginalization order: $P_1 \rightarrow P_2 \rightarrow P_3 \rightarrow P_4 \rightarrow P_0$

Readvancing



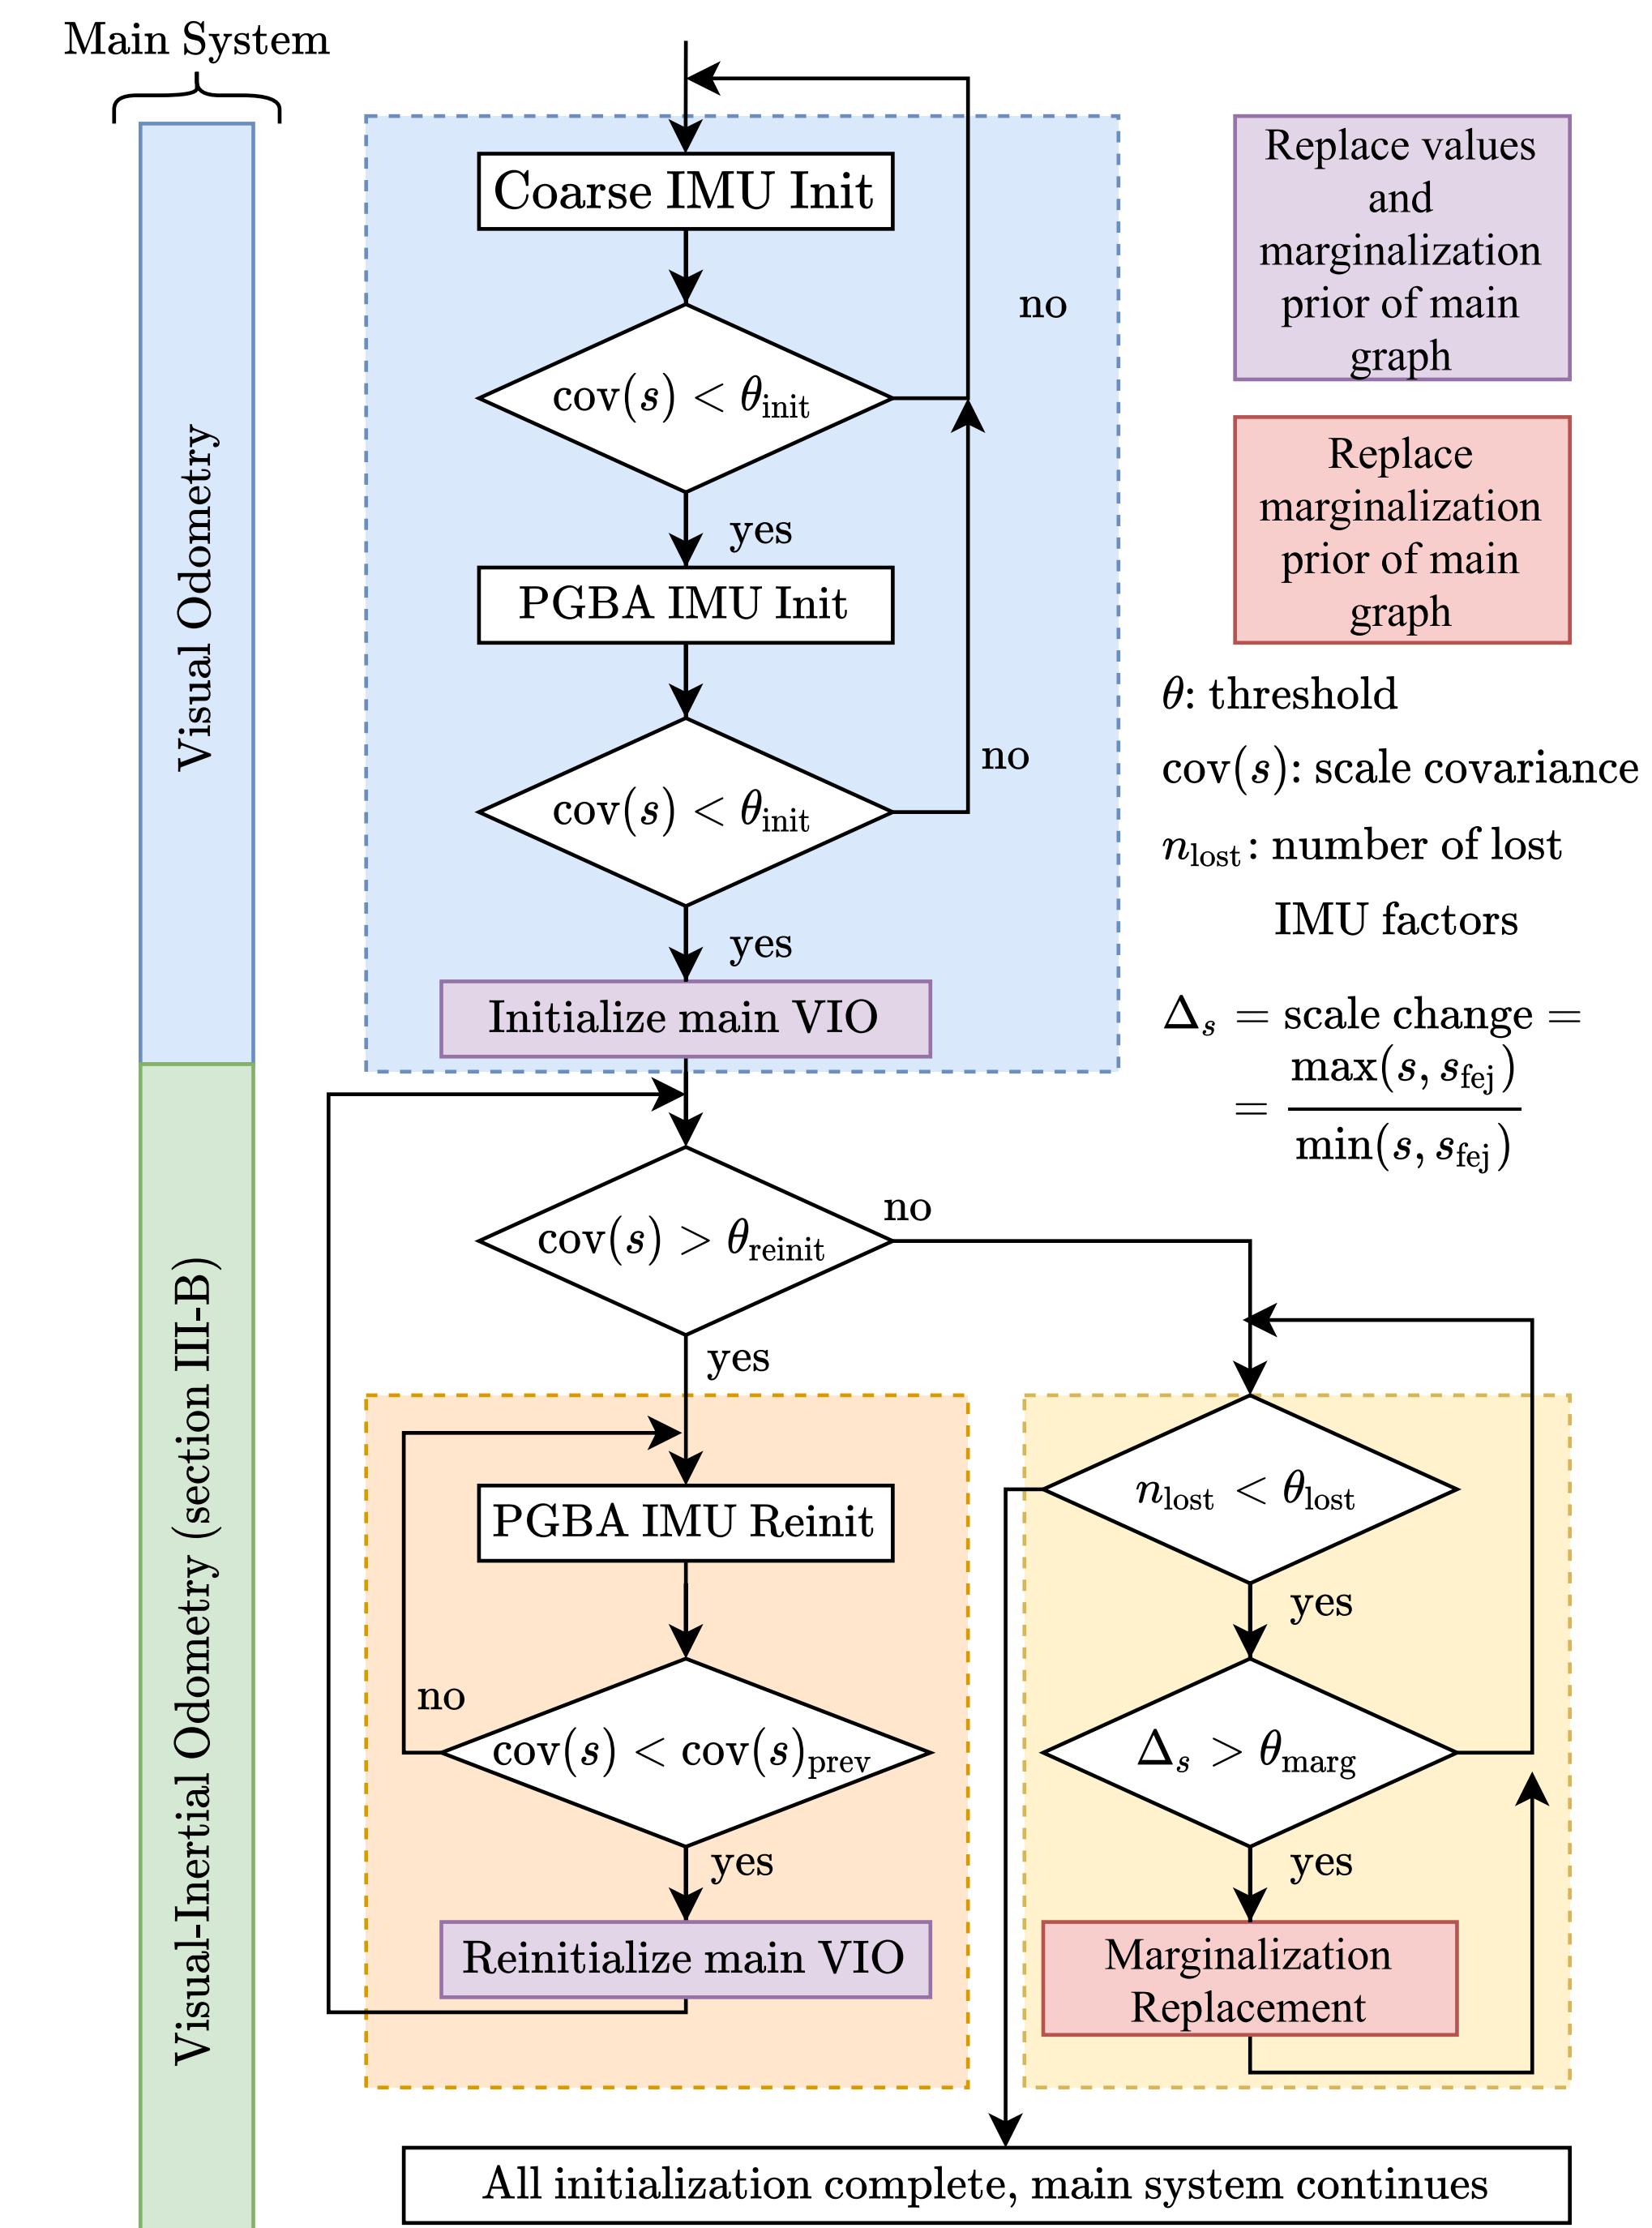
Readvancing

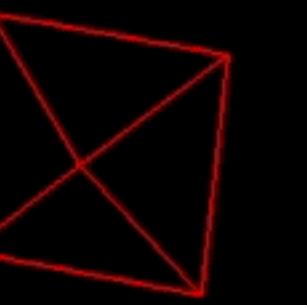
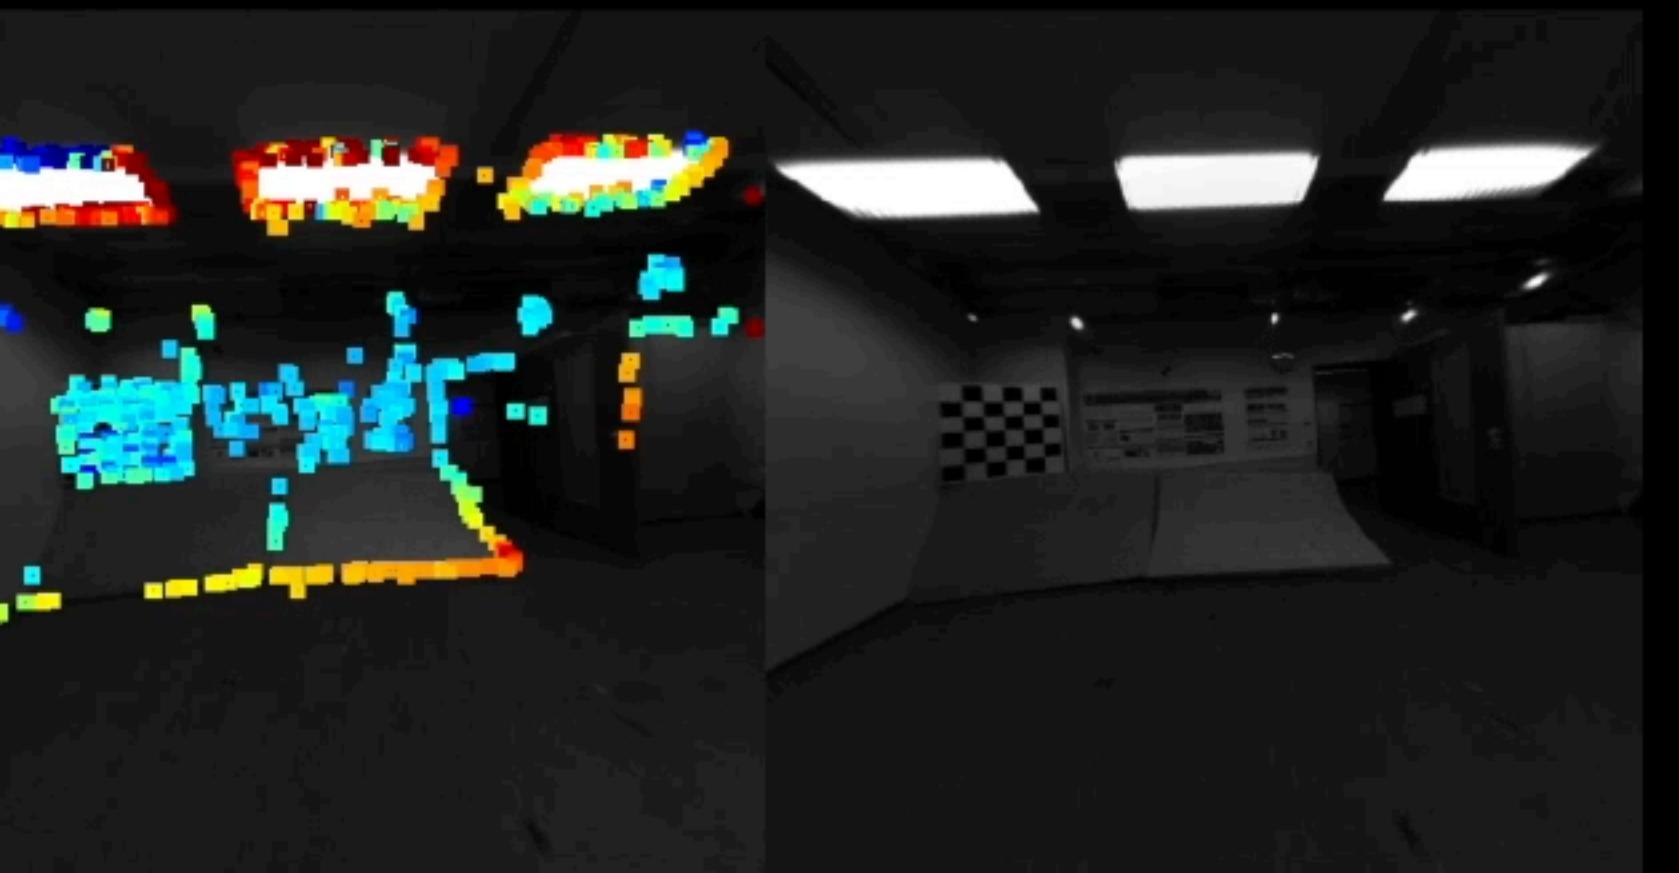
**marginalization
prior for the
main system**

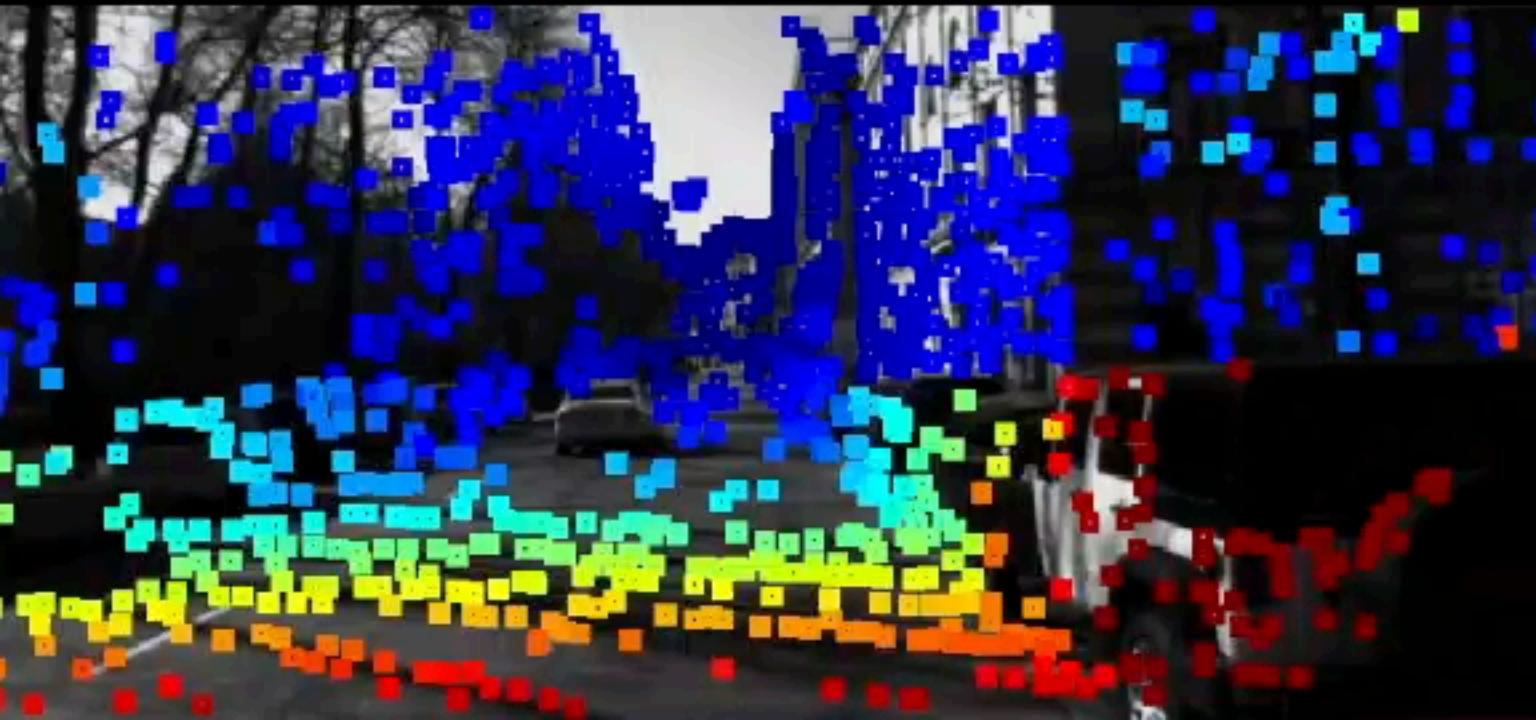


Main graph after IMU initialization

Multi-stage IMU initializer



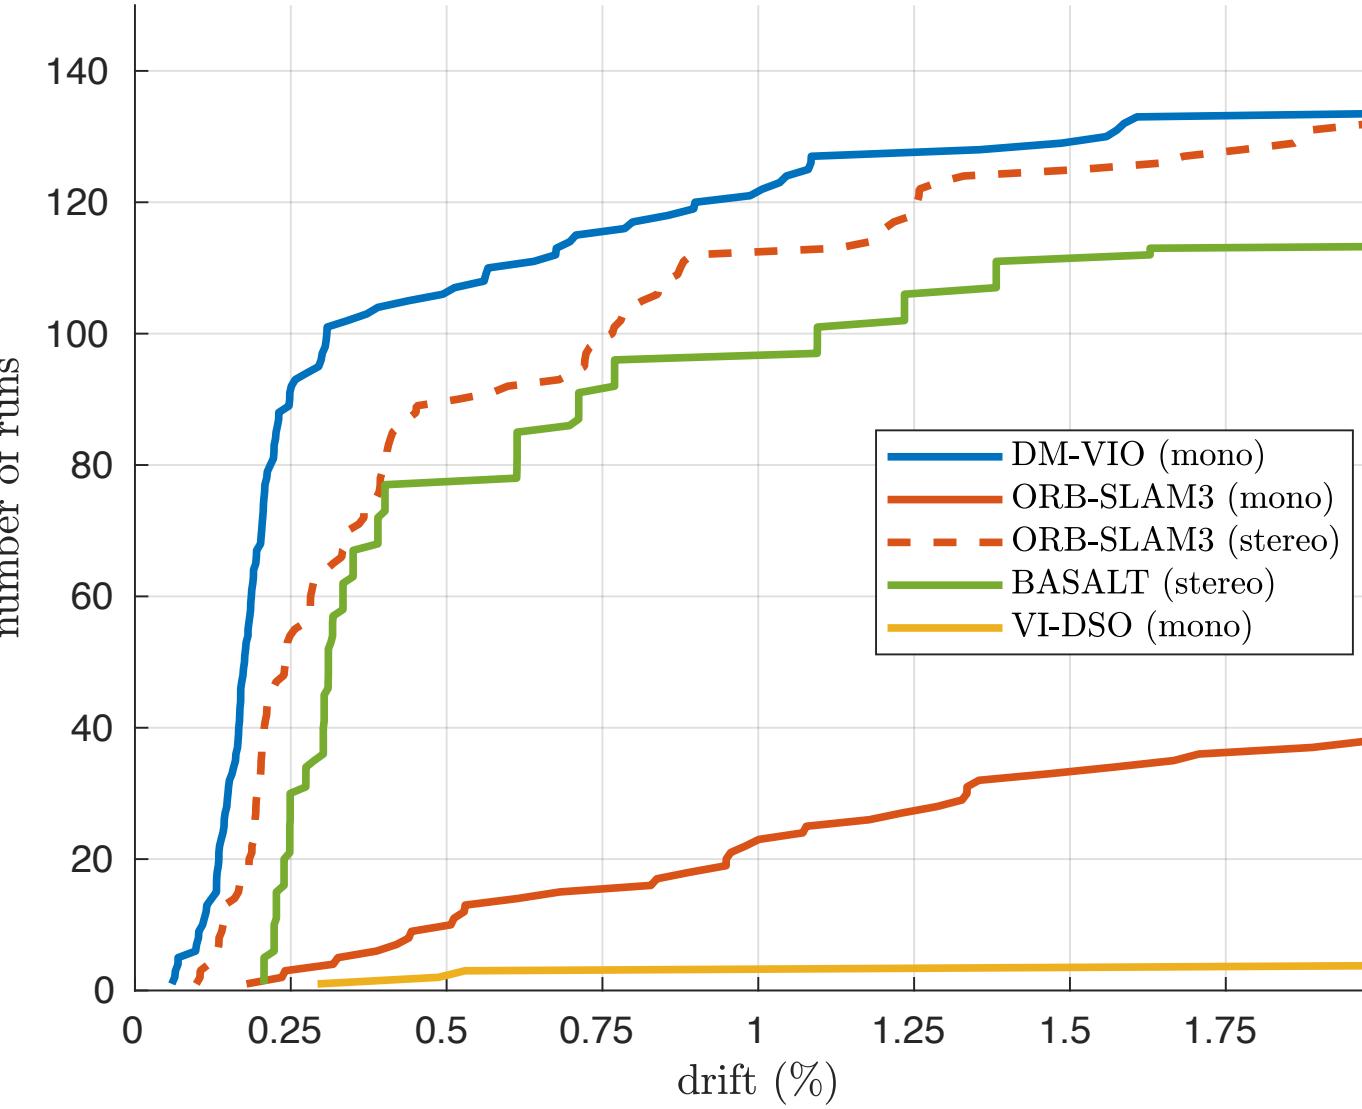




4Seasons
oldtown_2021-02-25_12-34-08

Results

4Seasons



EuRoC

Sequence		MH1	MH2	MH3	MH4	MH5	V11	V12	V13	V21	V22	V23	Avg
MCSKF ² [1] (M)	RMSE	0.42	0.45	0.23	0.37	0.48	0.34	0.20	0.67	0.10	0.16	1.13	0.414
OKVIS ¹ [19] (M)	RMSE	0.33	0.37	0.25	0.27	0.39	0.094	0.14	0.21	0.090	0.17	0.23	0.231
ROVIO ² [18] (M)	RMSE	0.21	0.25	0.25	0.49	0.52	0.10	0.10	0.14	0.12	0.14	0.14	0.224
VINS-Mono [3] (M)	RMSE	0.15	0.15	0.22	0.32	0.30	0.079	0.11	0.18	0.080	0.16	0.27	0.184
Kimera [21] (S)	RMSE	0.11	0.10	0.16	0.24	0.35	0.05	0.08	0.07	0.08	0.10	0.21	0.141
Online VIO [23] (M)	RMSE	0.14	0.13	0.20	0.22	0.20	0.05	0.07	0.16	0.04	0.11	0.17	0.135
VI-DSO [6] (M)	RMSE	0.062	0.044	0.117	0.132	0.121	0.059	0.067	0.096	0.040	0.062	0.174	0.089
	Scale Error (%)	1.1	0.5	0.4	0.2	0.8	1.1	1.1	0.8	1.2	0.3	0.4	0.7
BASALT [20] (S)	RMSE	0.07	0.06	0.07	0.13	0.11	0.04	0.05	0.10	0.04	0.05	-	0.072
DM-VIO (M)	RMSE	0.065	0.044	0.097	0.102	0.096	0.048	0.045	0.069	0.029	0.050	0.114	0.069
	Scale Error (%)	1.3	0.9	0.4	0.2	0.4	0.4	1.0	0.3	0.02	0.6	0.8	0.6

TUM-VI

Sequence	ROVIO stereo	VINS mono	OKVIS stereo	BASALT stereo	DM-VIO mono	length [m]
corridor1	0.47	0.63	0.33	0.34	0.19	305
corridor2	0.75	0.95	0.47	0.42	0.47	322
corridor3	0.85	1.56	0.57	0.35	0.24	300
corridor4	0.13	0.25	0.26	0.21	0.13	114
corridor5	2.09	0.77	0.39	0.37	0.16	270
magistrale1	4.52	2.19	3.49	1.20	2.35	918
magistrale2	13.43	3.11	2.73	1.11	2.24	561
magistrale3	14.80	0.40	1.22	0.74	1.69	566
magistrale4	39.73	5.12	0.77	1.58	1.02	688
magistrale5	3.47	0.85	1.62	0.60	0.73	458
magistrale6	X	2.29	3.91	3.23	1.19	771
outdoors1	101.95	74.96	X	255.04	123.24	2656
outdoors2	21.67	133.46	73.86	64.61	12.76	1601
outdoors3	26.10	36.99	32.38	38.26	8.92	1531
outdoors4	X	16.46	19.51	17.53	15.25	928
outdoors5	54.32	130.63	13.12	7.89	7.16	1168
outdoors6	149.14	133.60	96.51	65.50	34.86	2045
outdoors7	49.01	21.90	13.61	4.07	5.00	1748
outdoors8	36.03	83.36	16.31	13.53	2.11	986
room1	0.16	0.07	0.06	0.09	0.03	146
room2	0.33	0.07	0.11	0.07	0.13	142
room3	0.15	0.11	0.07	0.13	0.09	135
room4	0.09	0.04	0.03	0.05	0.04	68
room5	0.12	0.20	0.07	0.13	0.06	131
room6	0.05	0.08	0.04	0.02	0.02	67
slides1	13.73	0.68	0.86	0.32	0.31	289
slides2	0.81	0.84	2.15	0.32	0.87	299
slides3	4.68	0.69	2.58	0.89	0.60	383
avg drift%	16.83*	1.700	0.815*	0.939	0.472	normalized

master ▾

2 branches 0 tags

Go to file

Code ▾

Lukas von Stumberg Parallelized the most time-intensive part of CoarselInitializer::c... 3b5319a on Jun 6 9 commits

cmake	Initial commit with the first version of the code.	last year
configs	Added live-version of DM-VIO for the Realsense T265 camera. It will ...	last year
doc	Added live-version of DM-VIO for the Realsense T265 camera. It will ...	last year
src	Parallelized the most time-intensive part of CoarselInitializer::calcR...	last month
test	Added live-version of DM-VIO for the Realsense T265 camera. It will ...	last year
thirdparty	Initial commit with the first version of the code.	last year
.gitmodules	Initial commit with the first version of the code.	last year
CMakeLists.txt	Added follow-cam mode to the GUI to smoothly follow the camera tr...	2 months ago
LICENSE	Initial commit with the first version of the code.	last year
README.md	Added follow-cam mode to the GUI to smoothly follow the camera tr...	2 months ago

README.md

DM-VIO: Delayed Marginalization Visual-Inertial Odometry

[Paper](#) | [Video](#) | [Project Page](#)

When using this project in academic work, please consider citing:

About

Source code for the paper DM-VIO:
Delayed Marginalization Visual-Inertial
Odometry

Readme

GPL-3.0 license

Activity

758 stars

53 watching

142 forks

Report repository

Code online:
[https://github.com/
lukasvst/dm-vio](https://github.com/lukasvst/dm-vio)



Releases

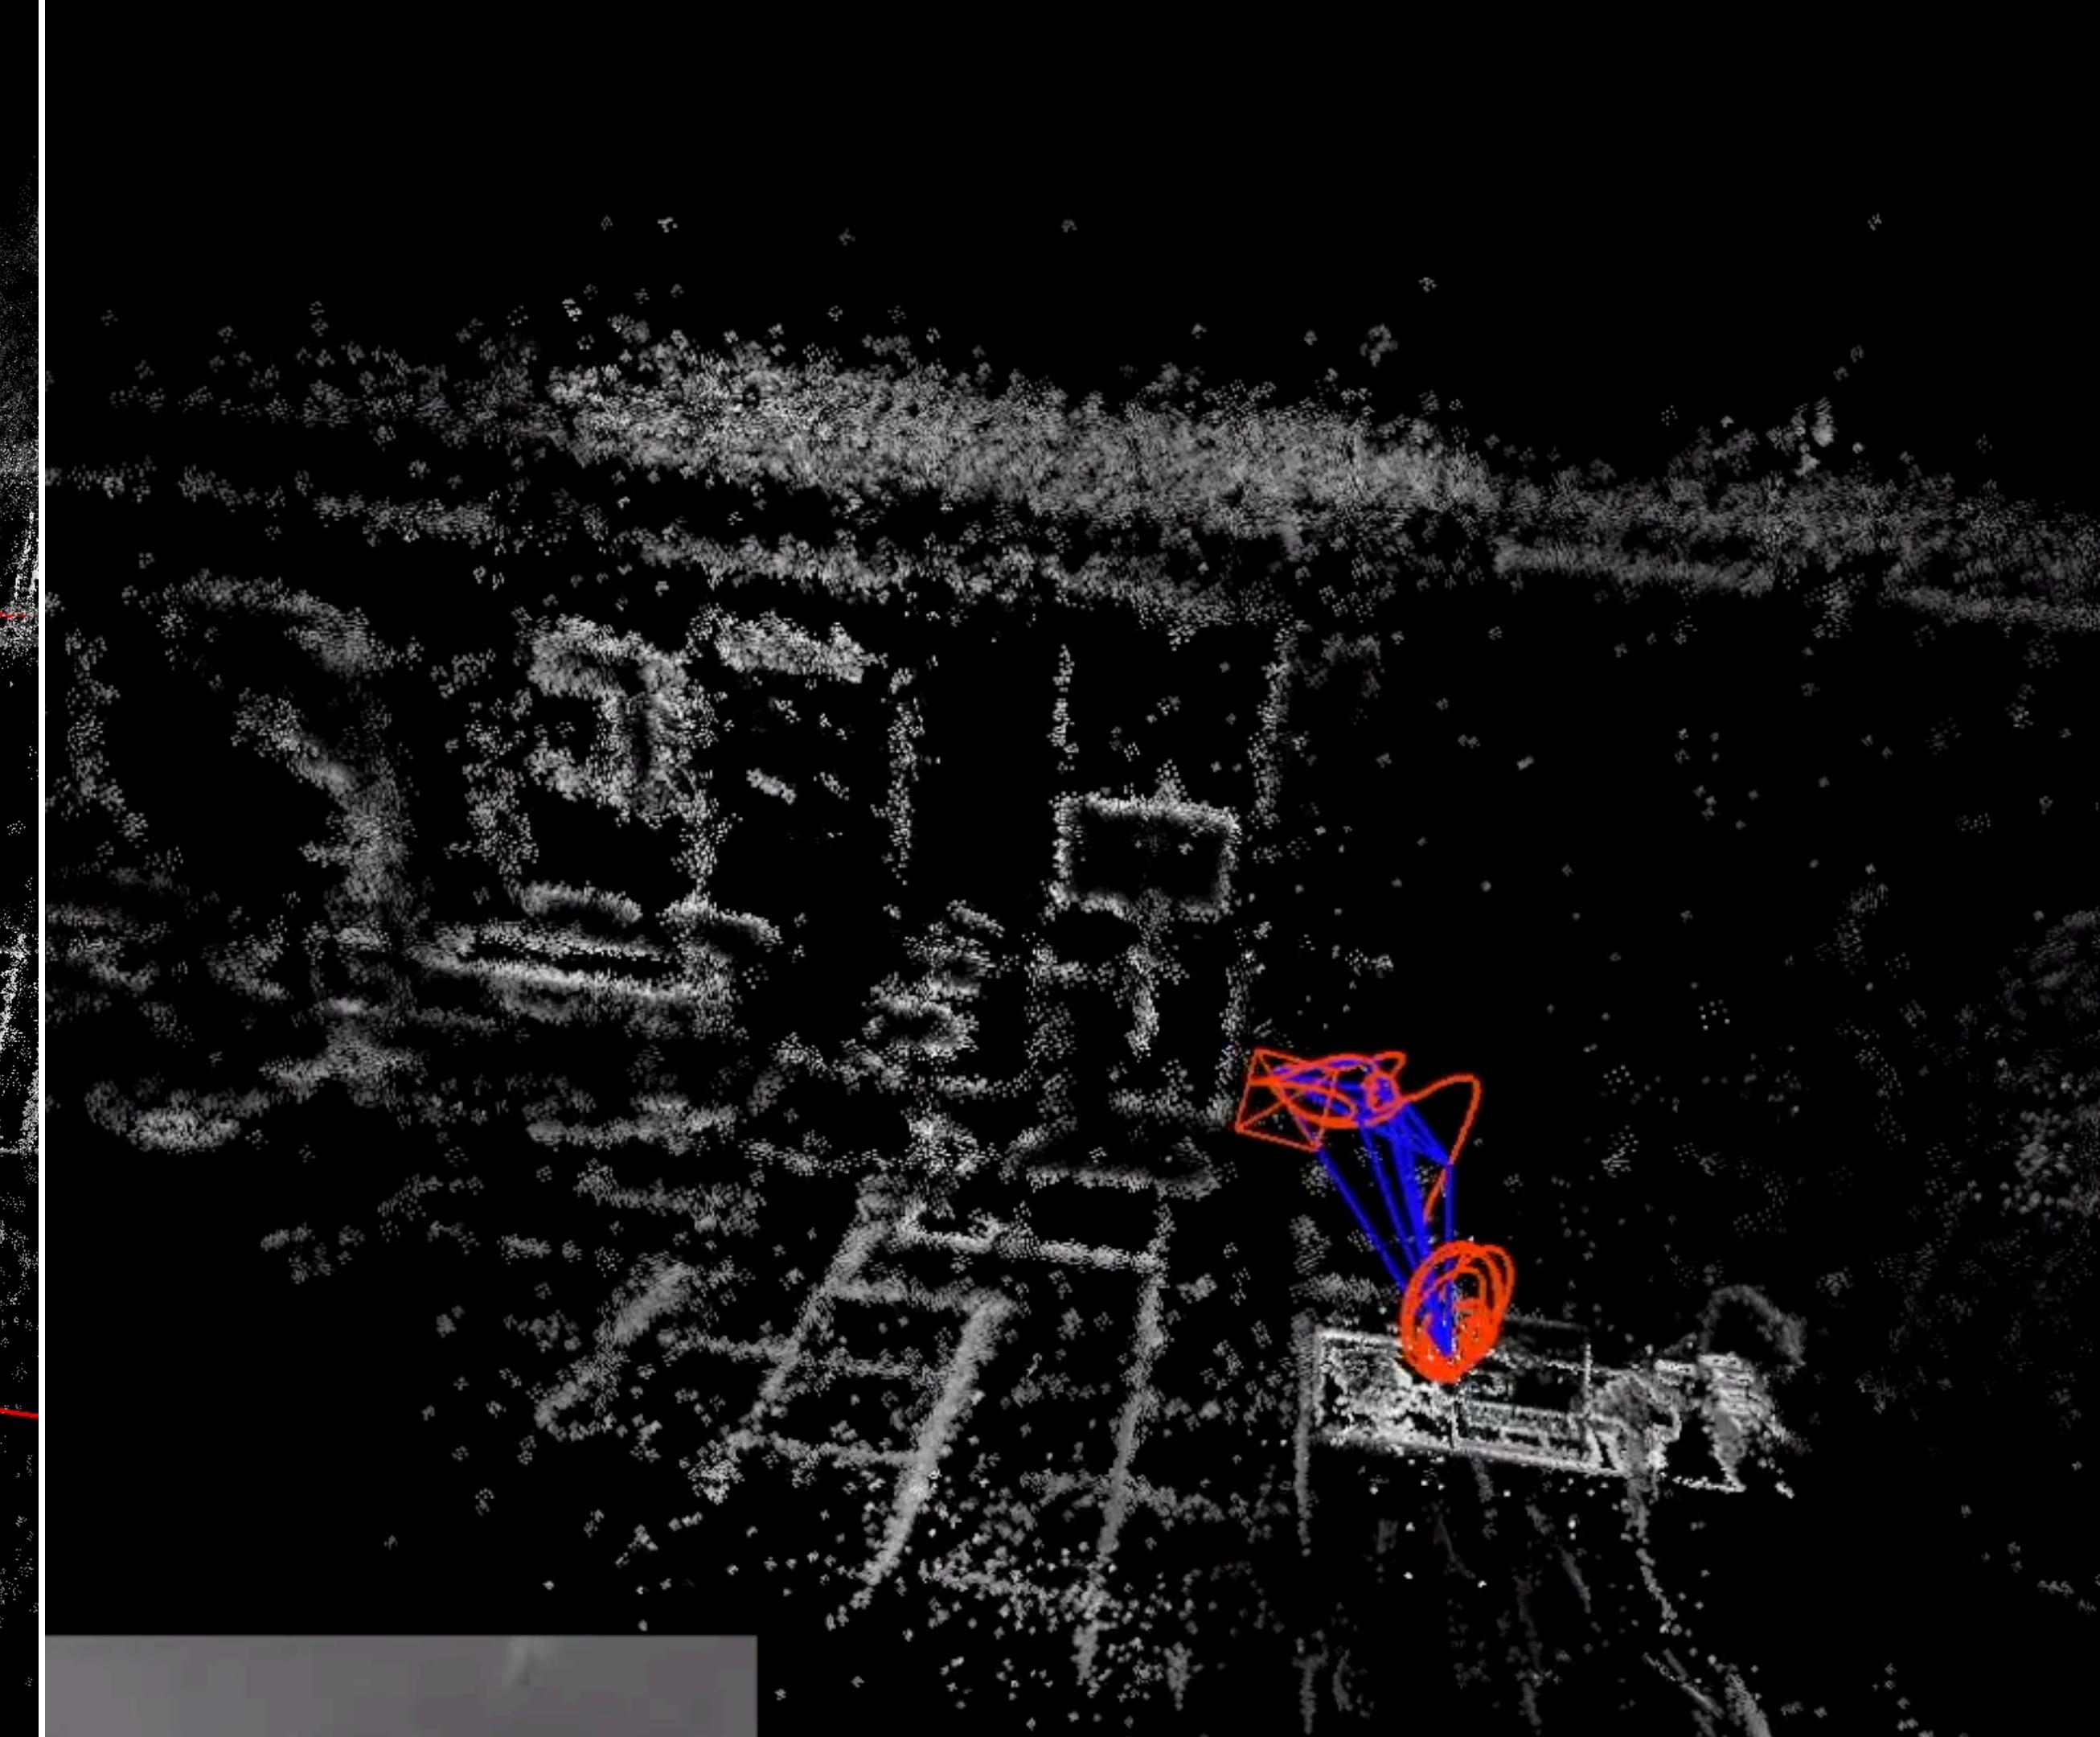
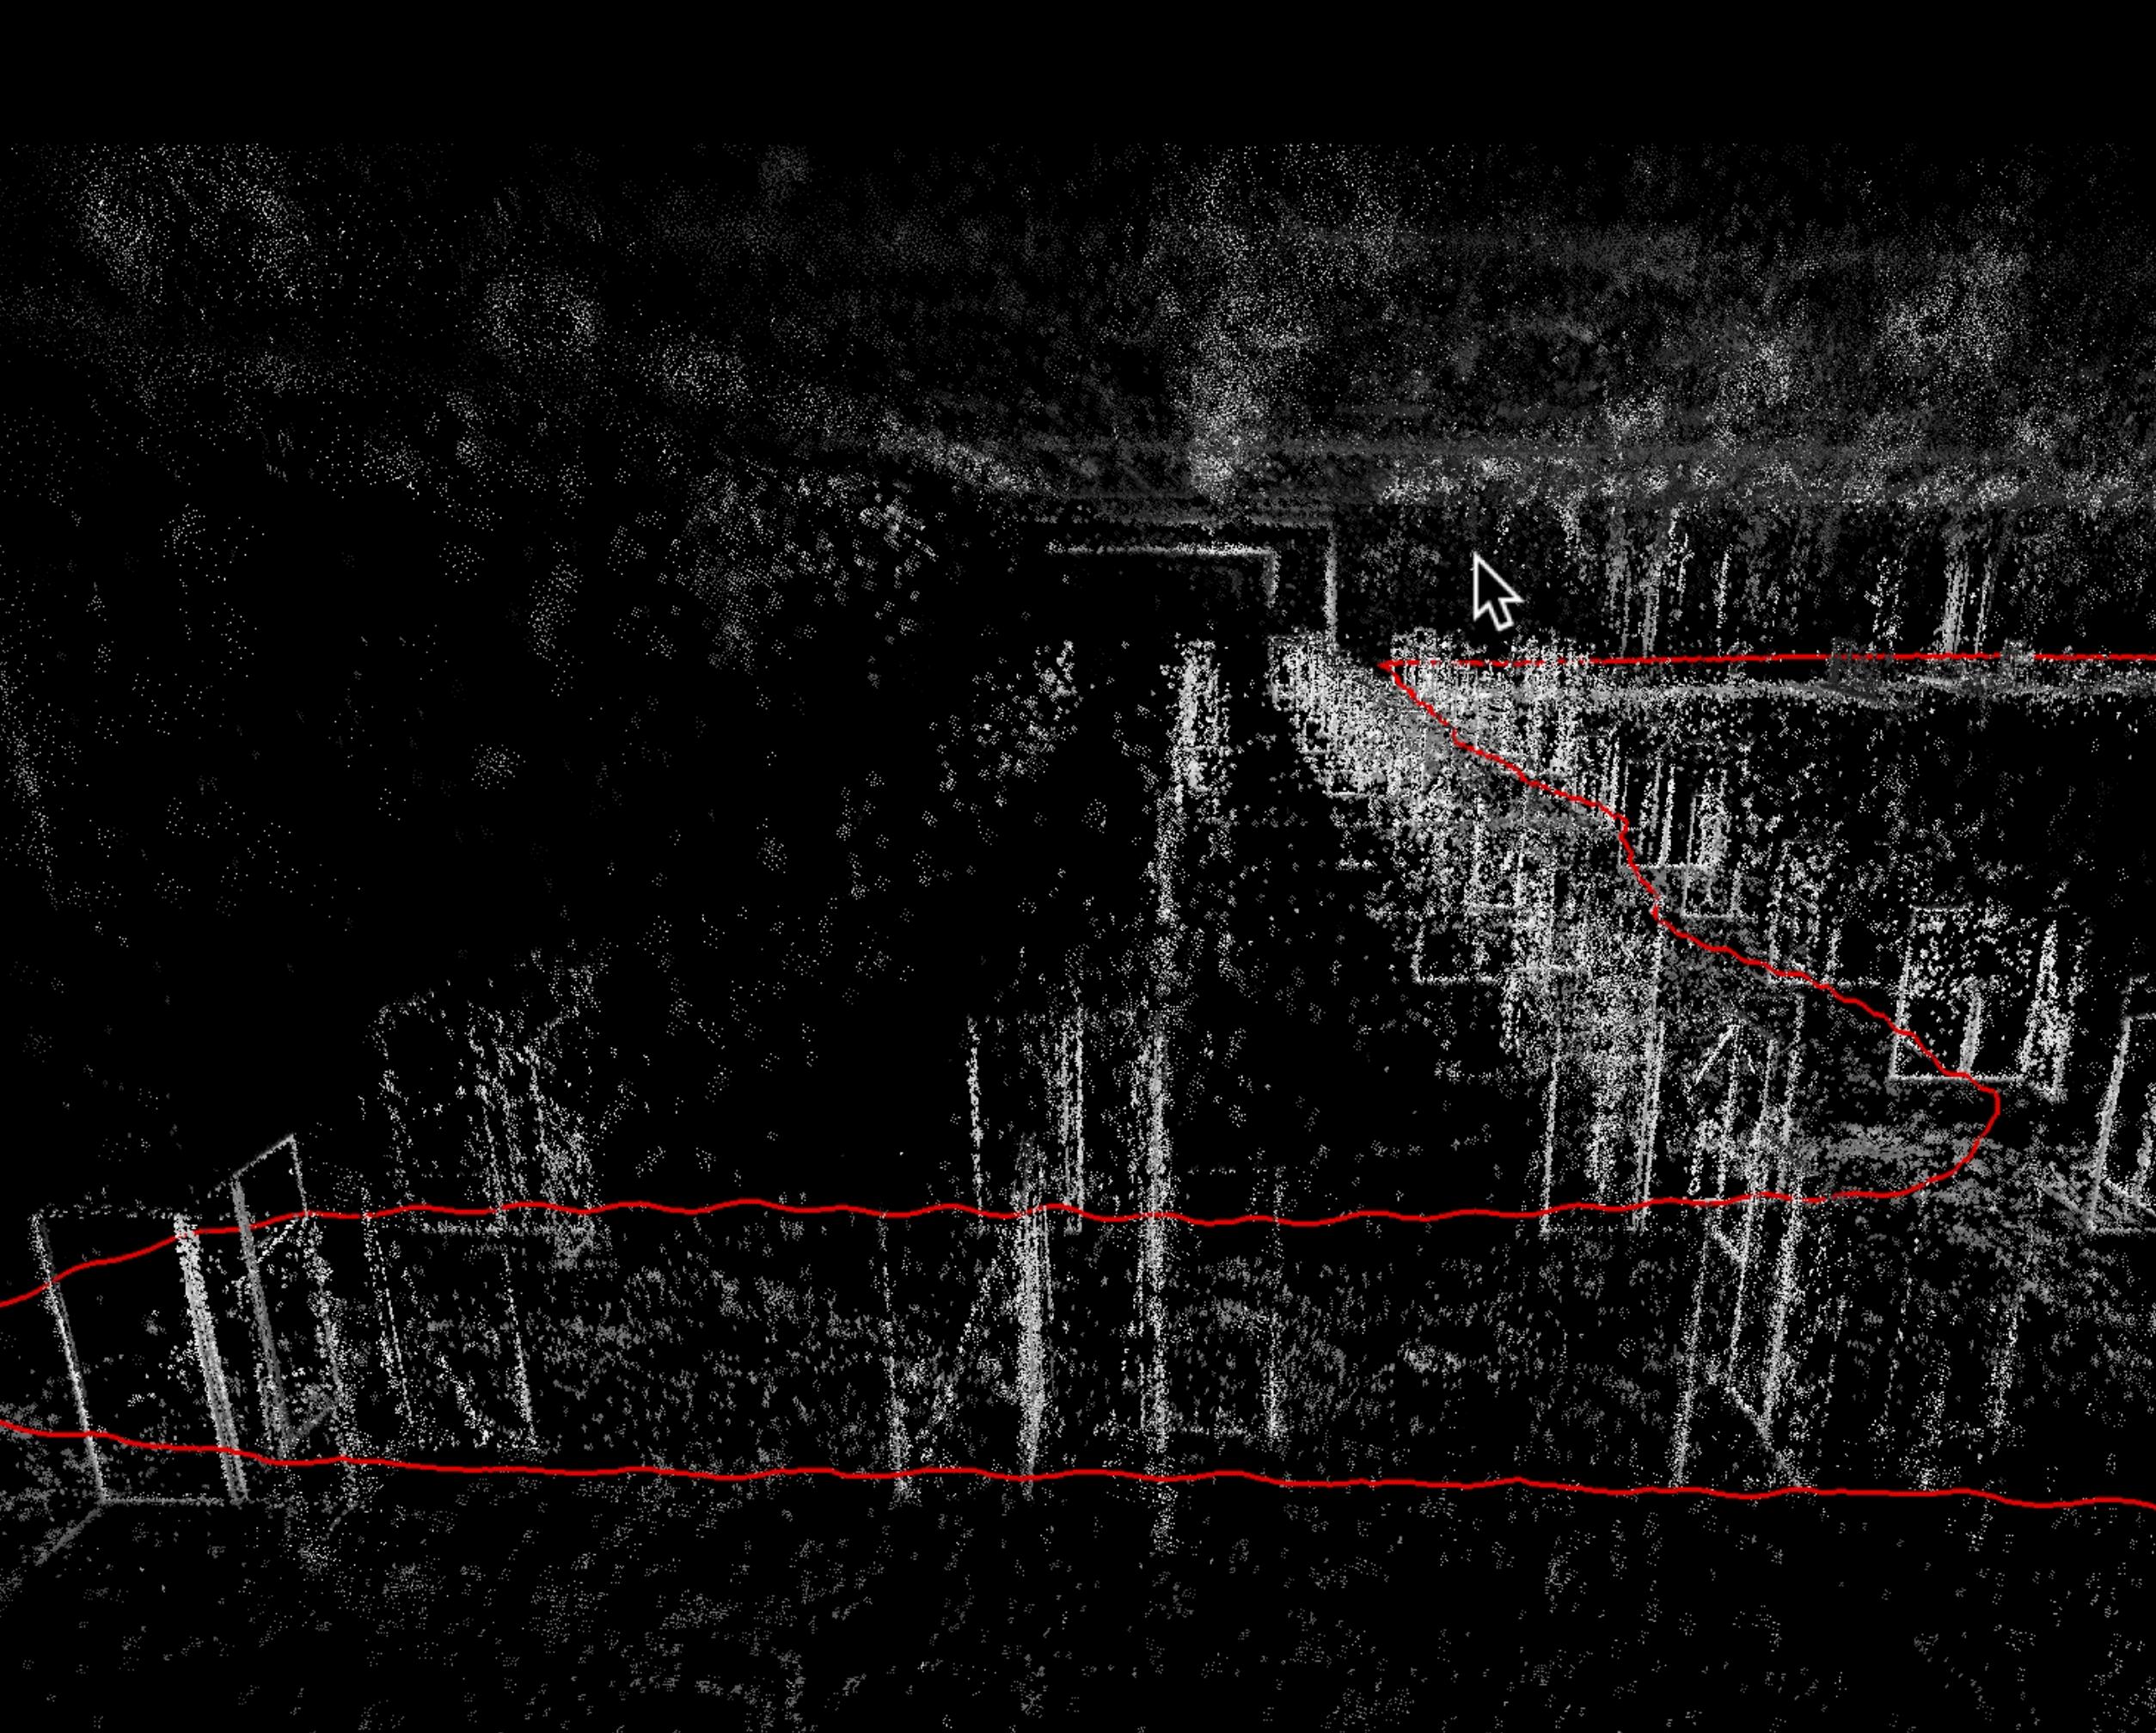
No releases published

Packages

No packages published

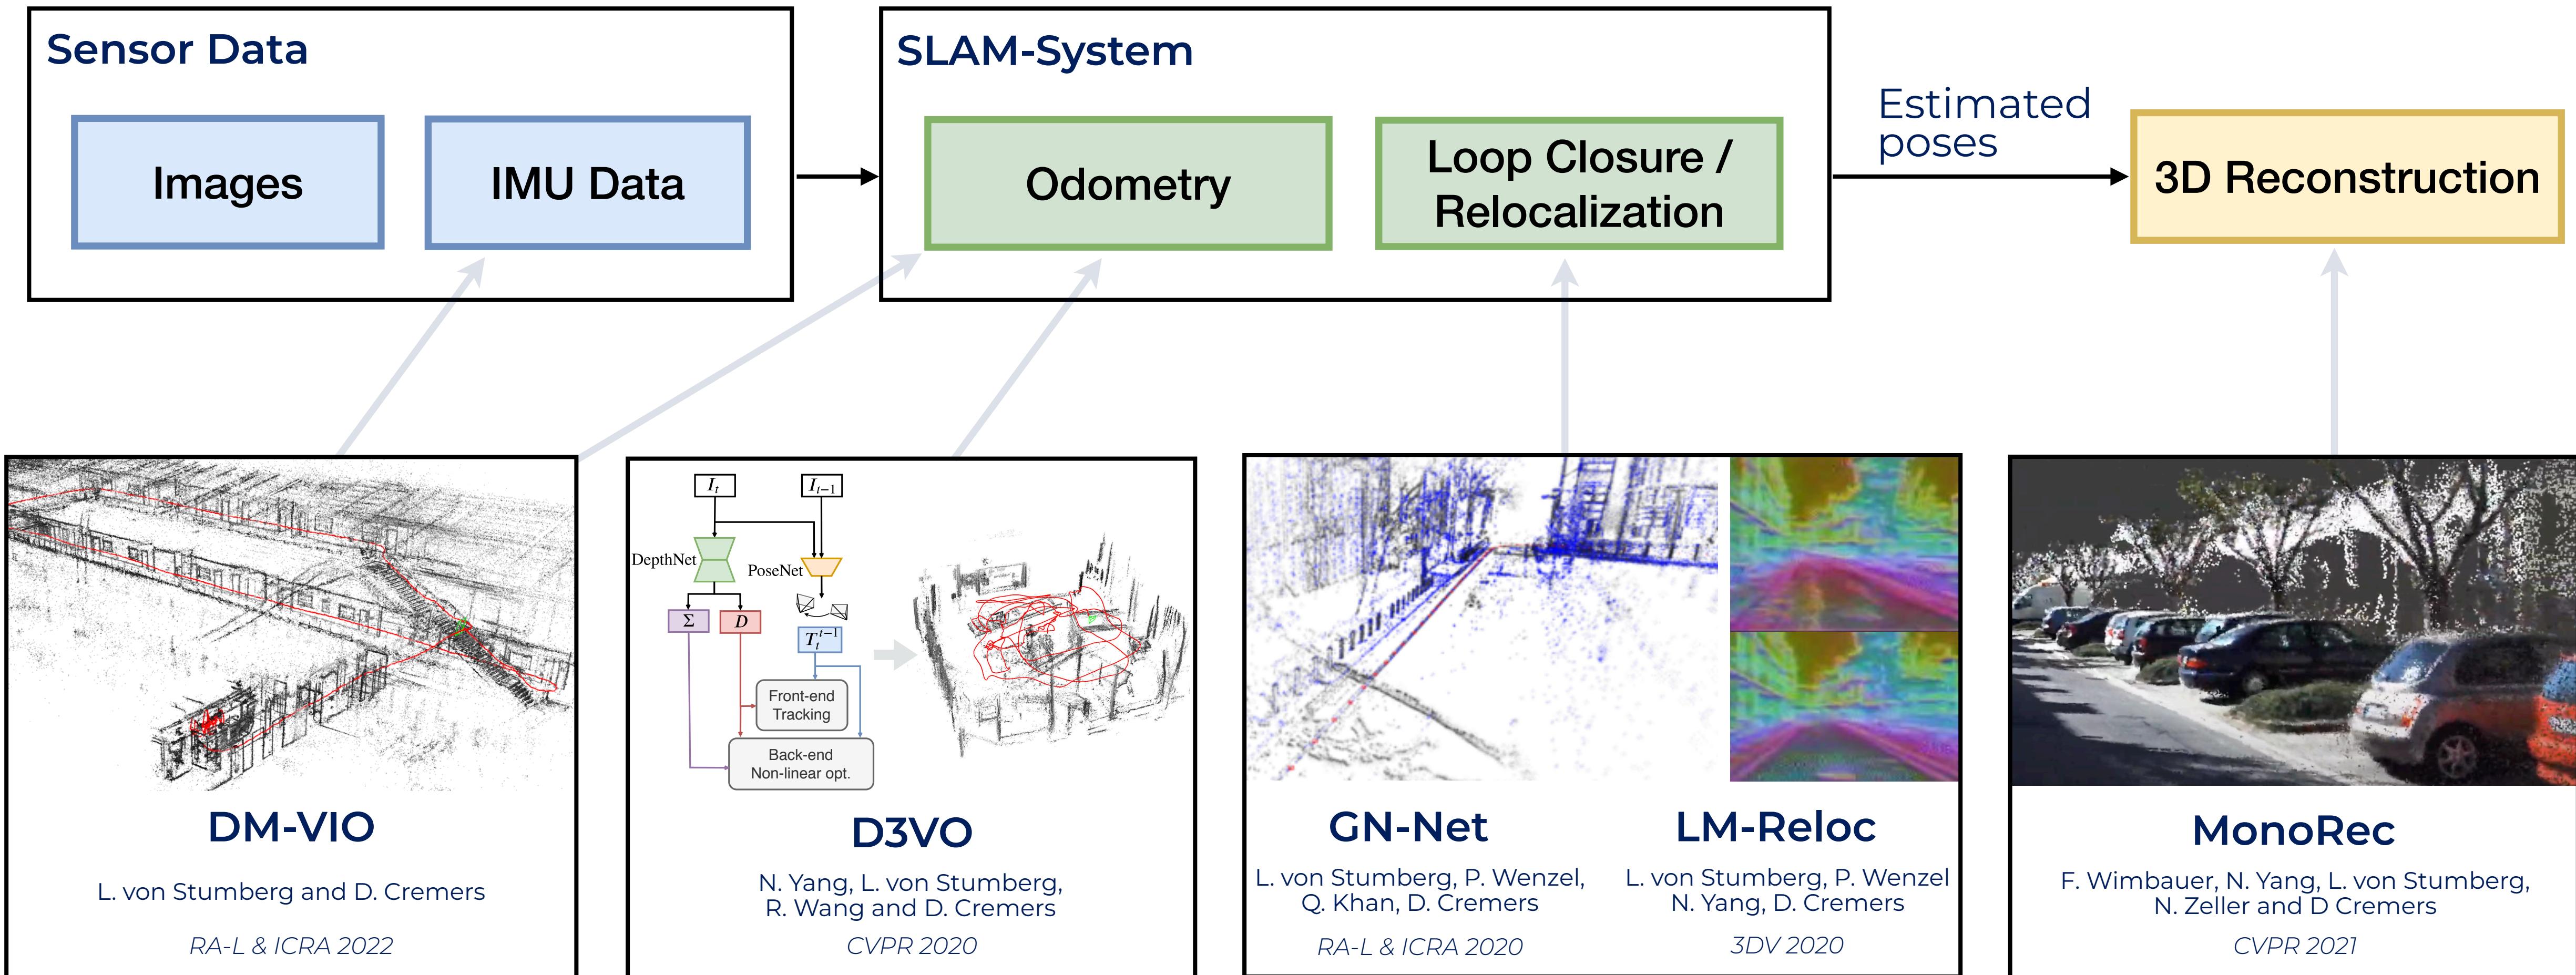
Languages

C++ 98.7% CMake 1.3%



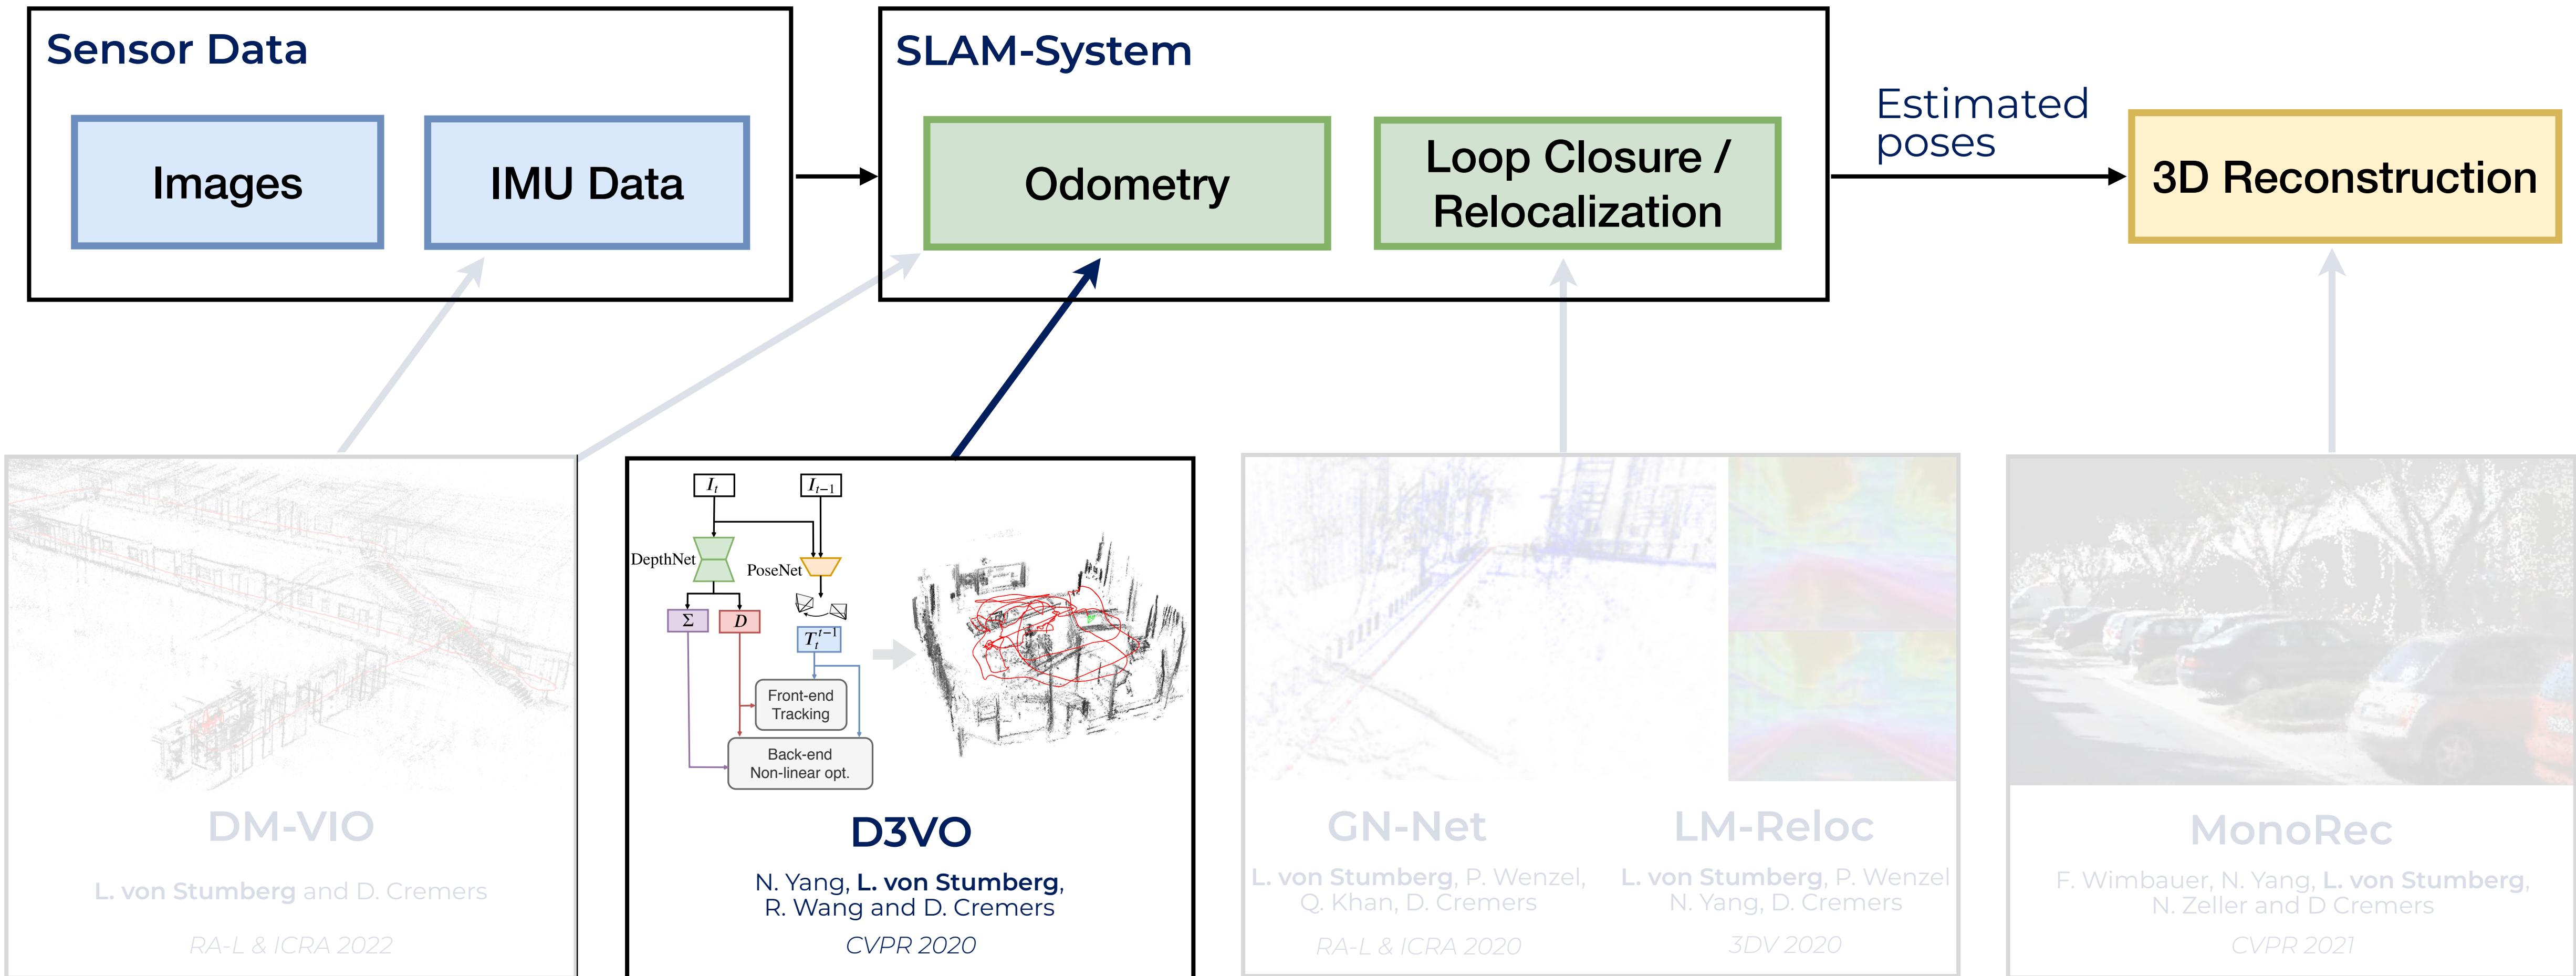
Visual SLAM

From Optimization to Learning



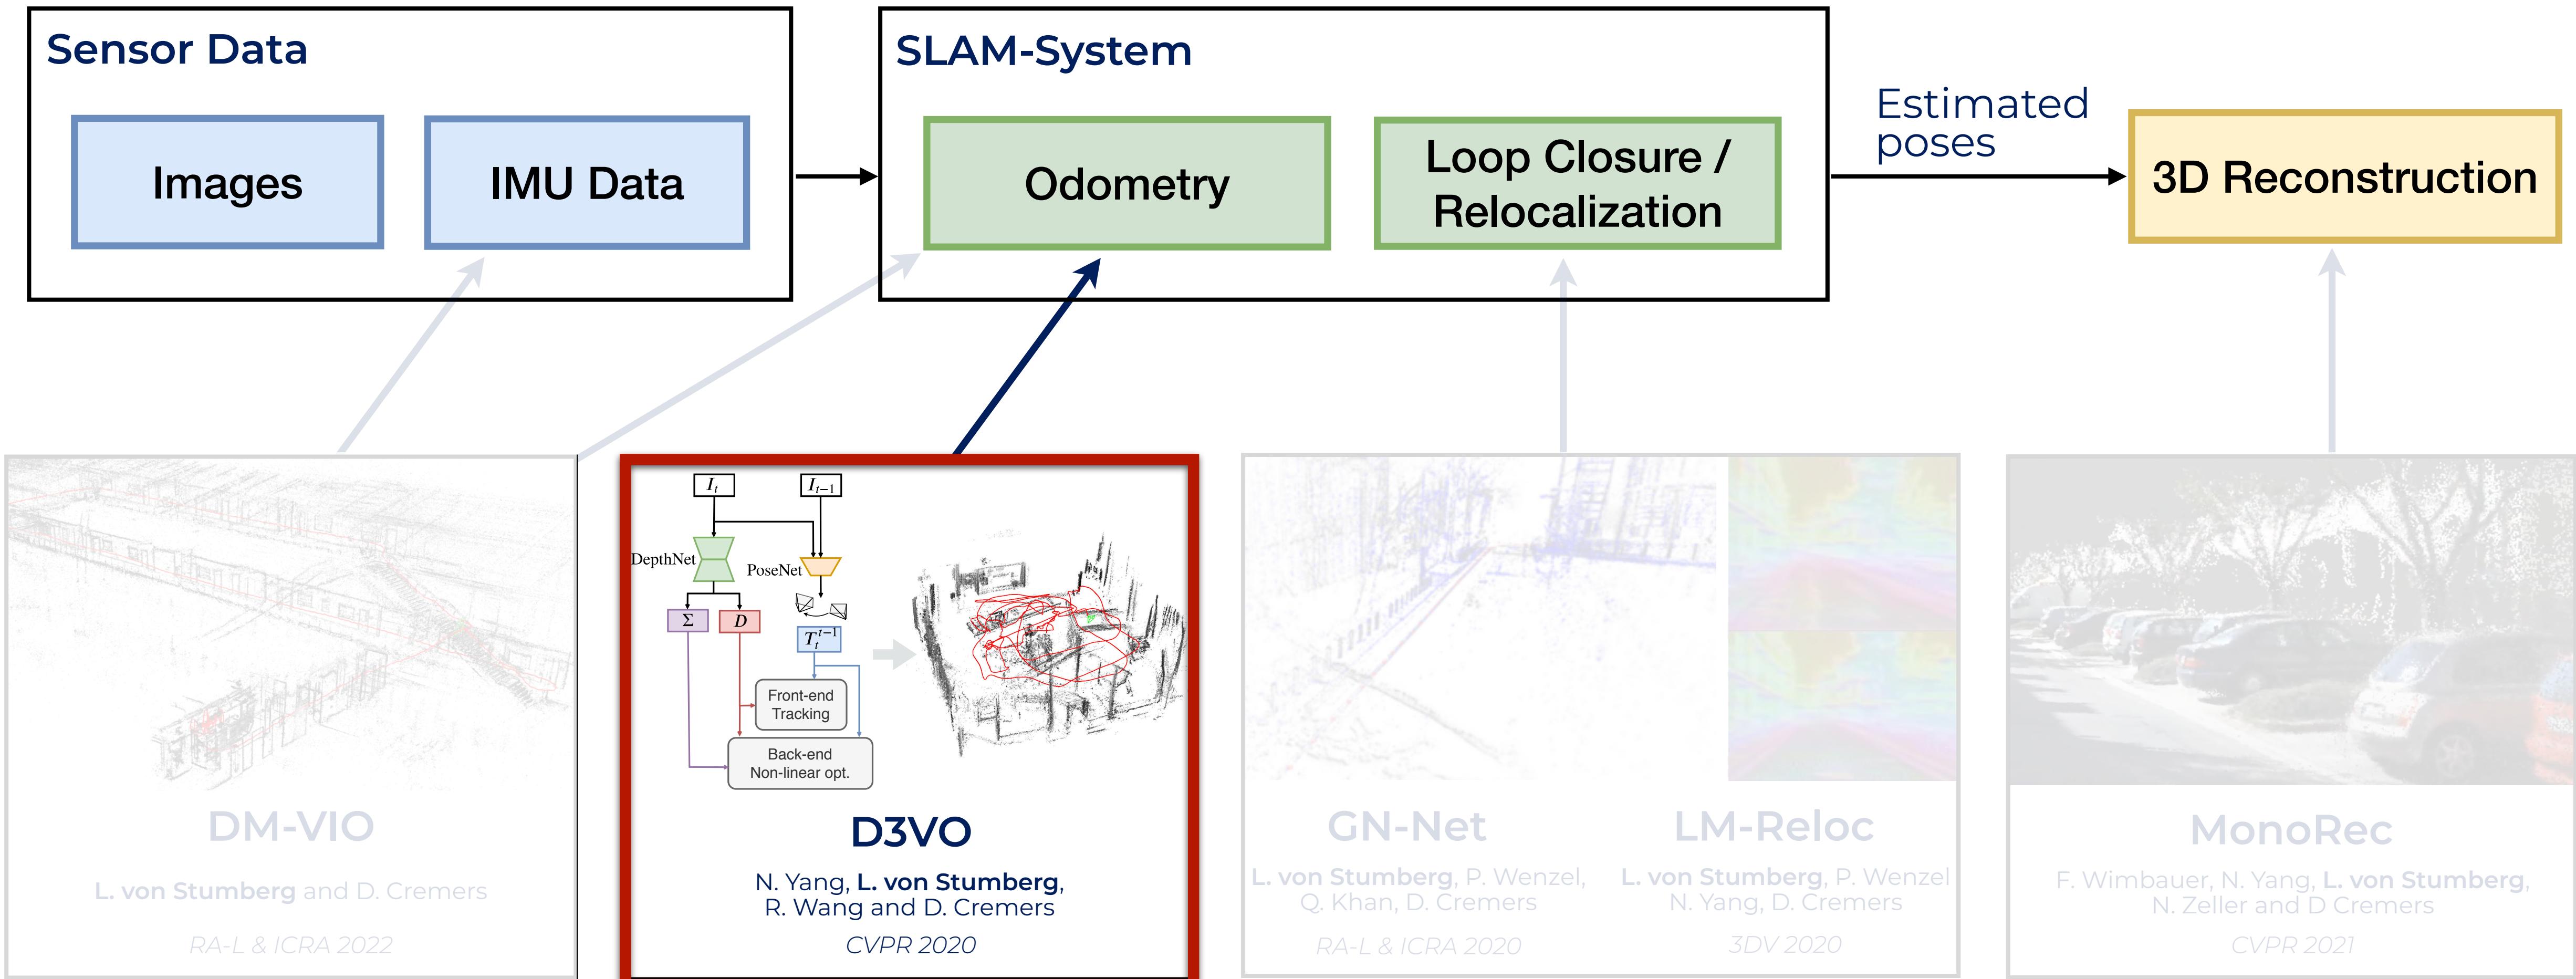
Visual SLAM

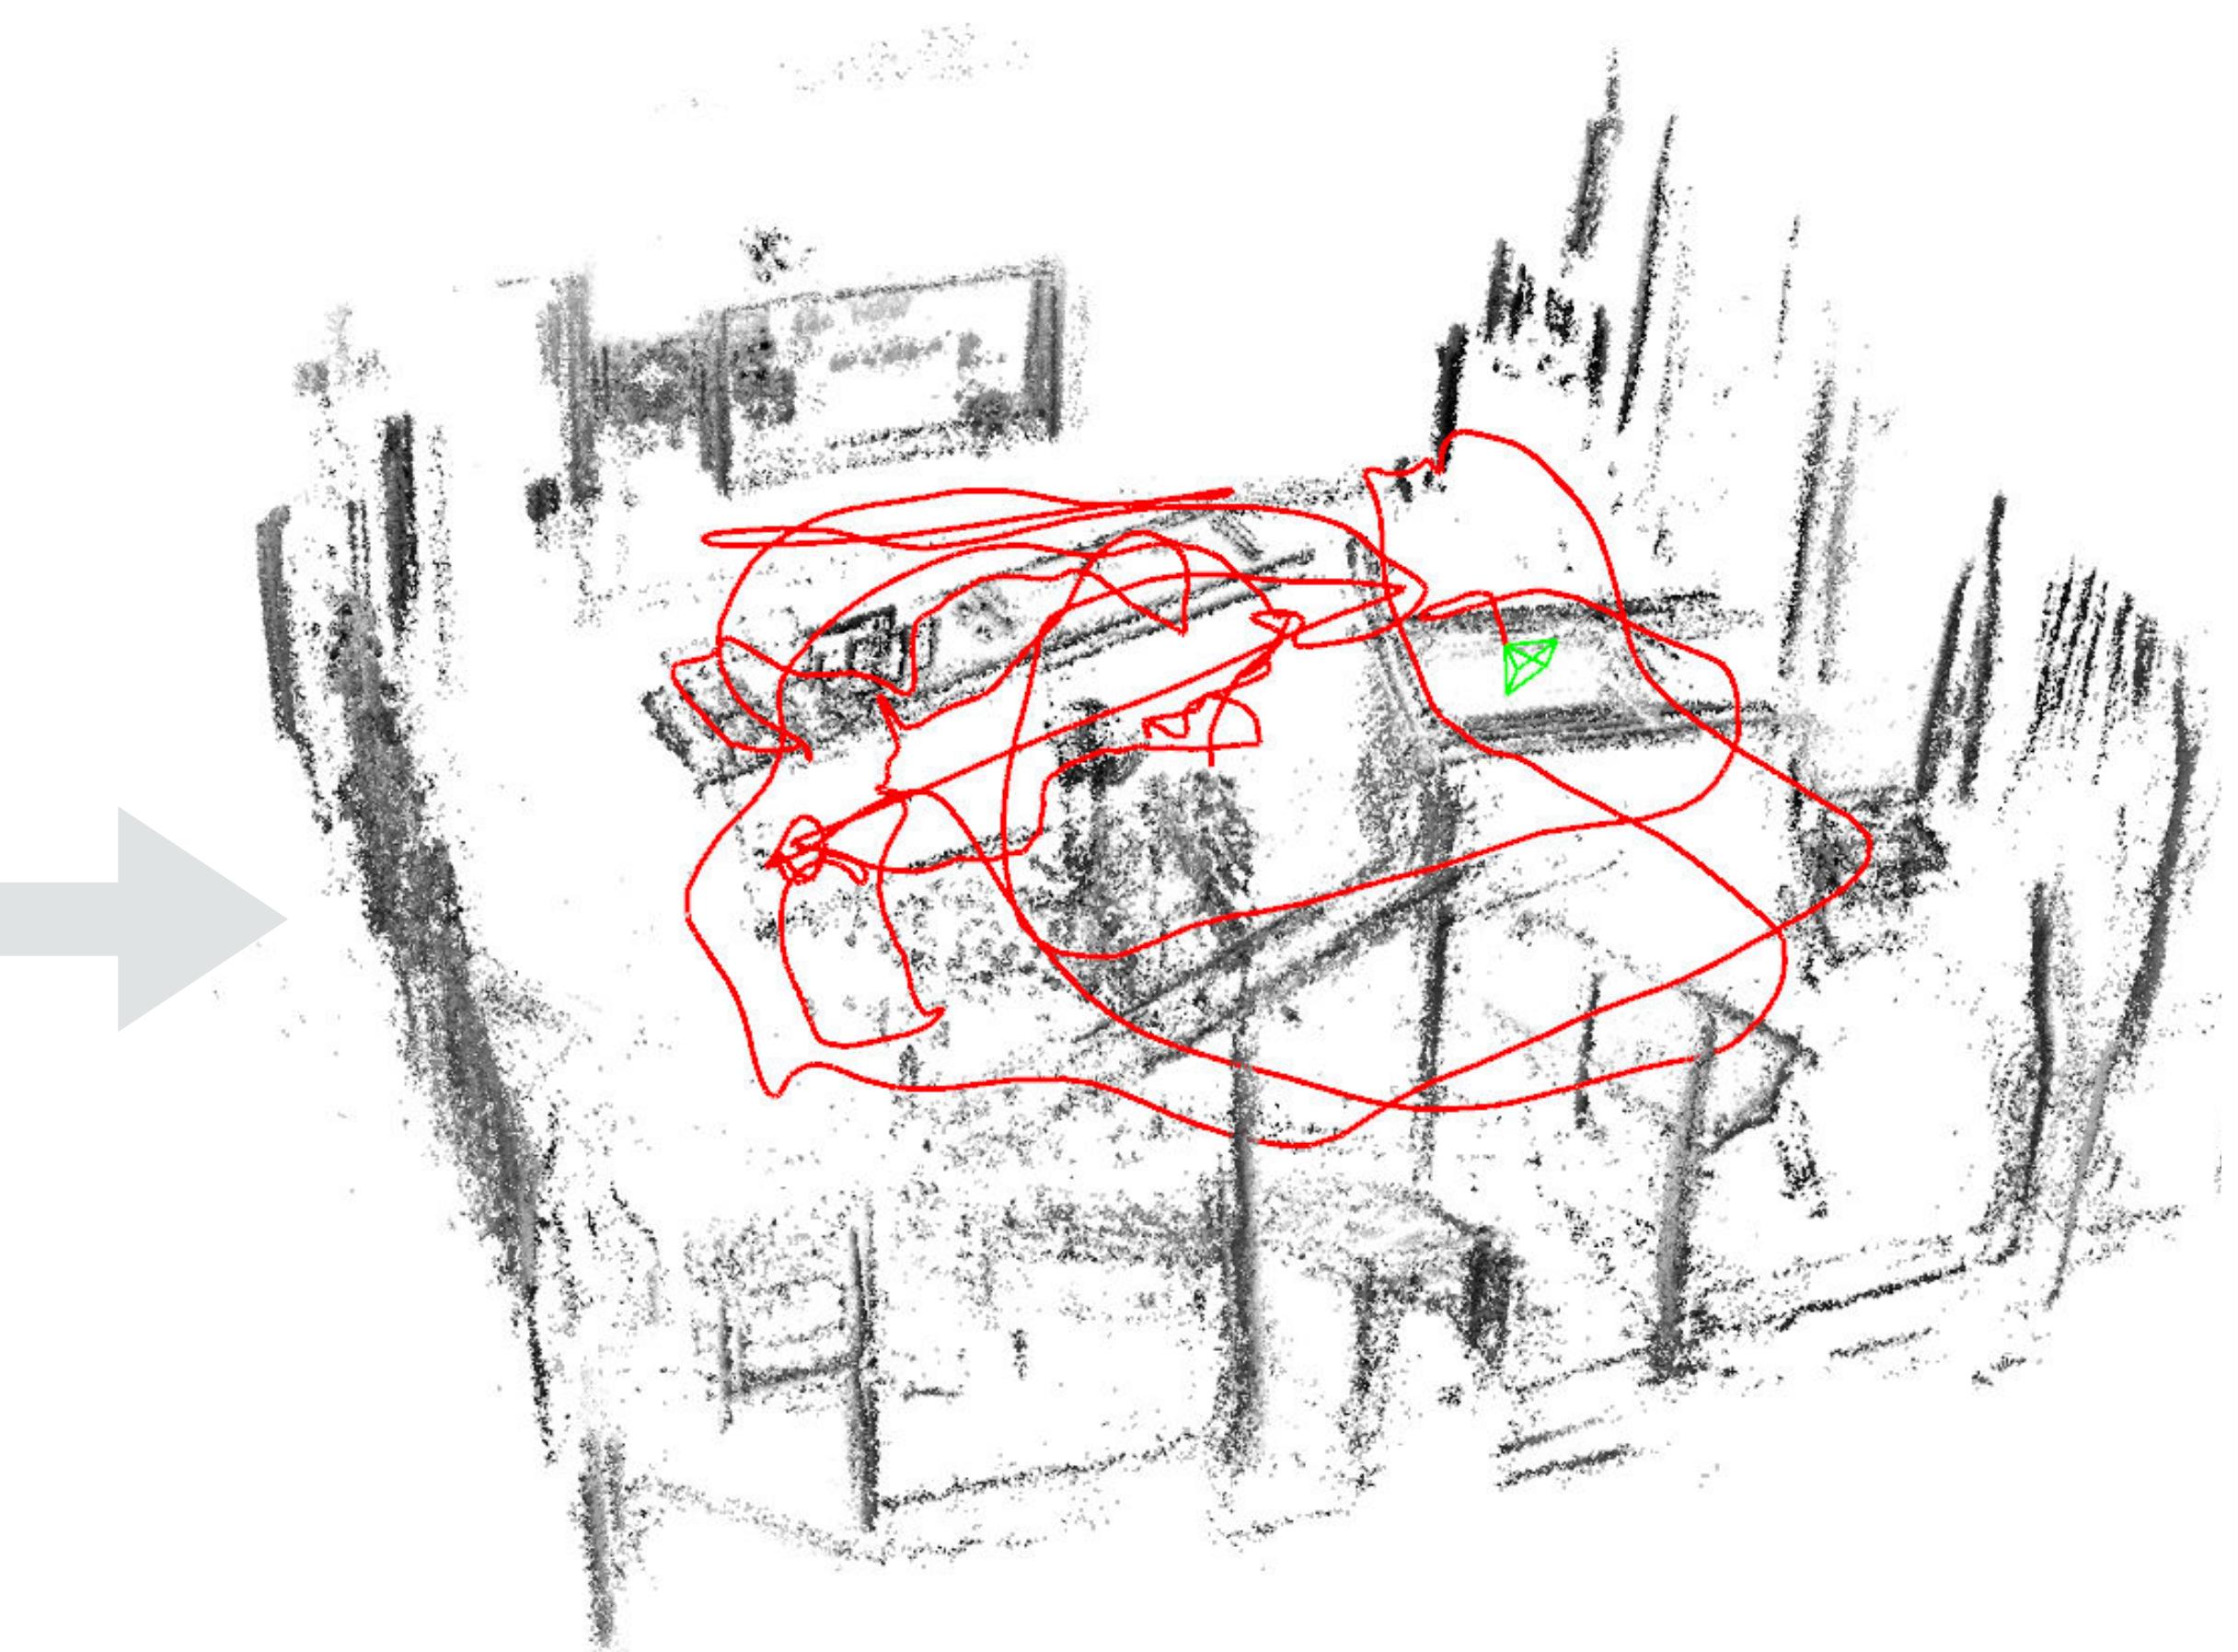
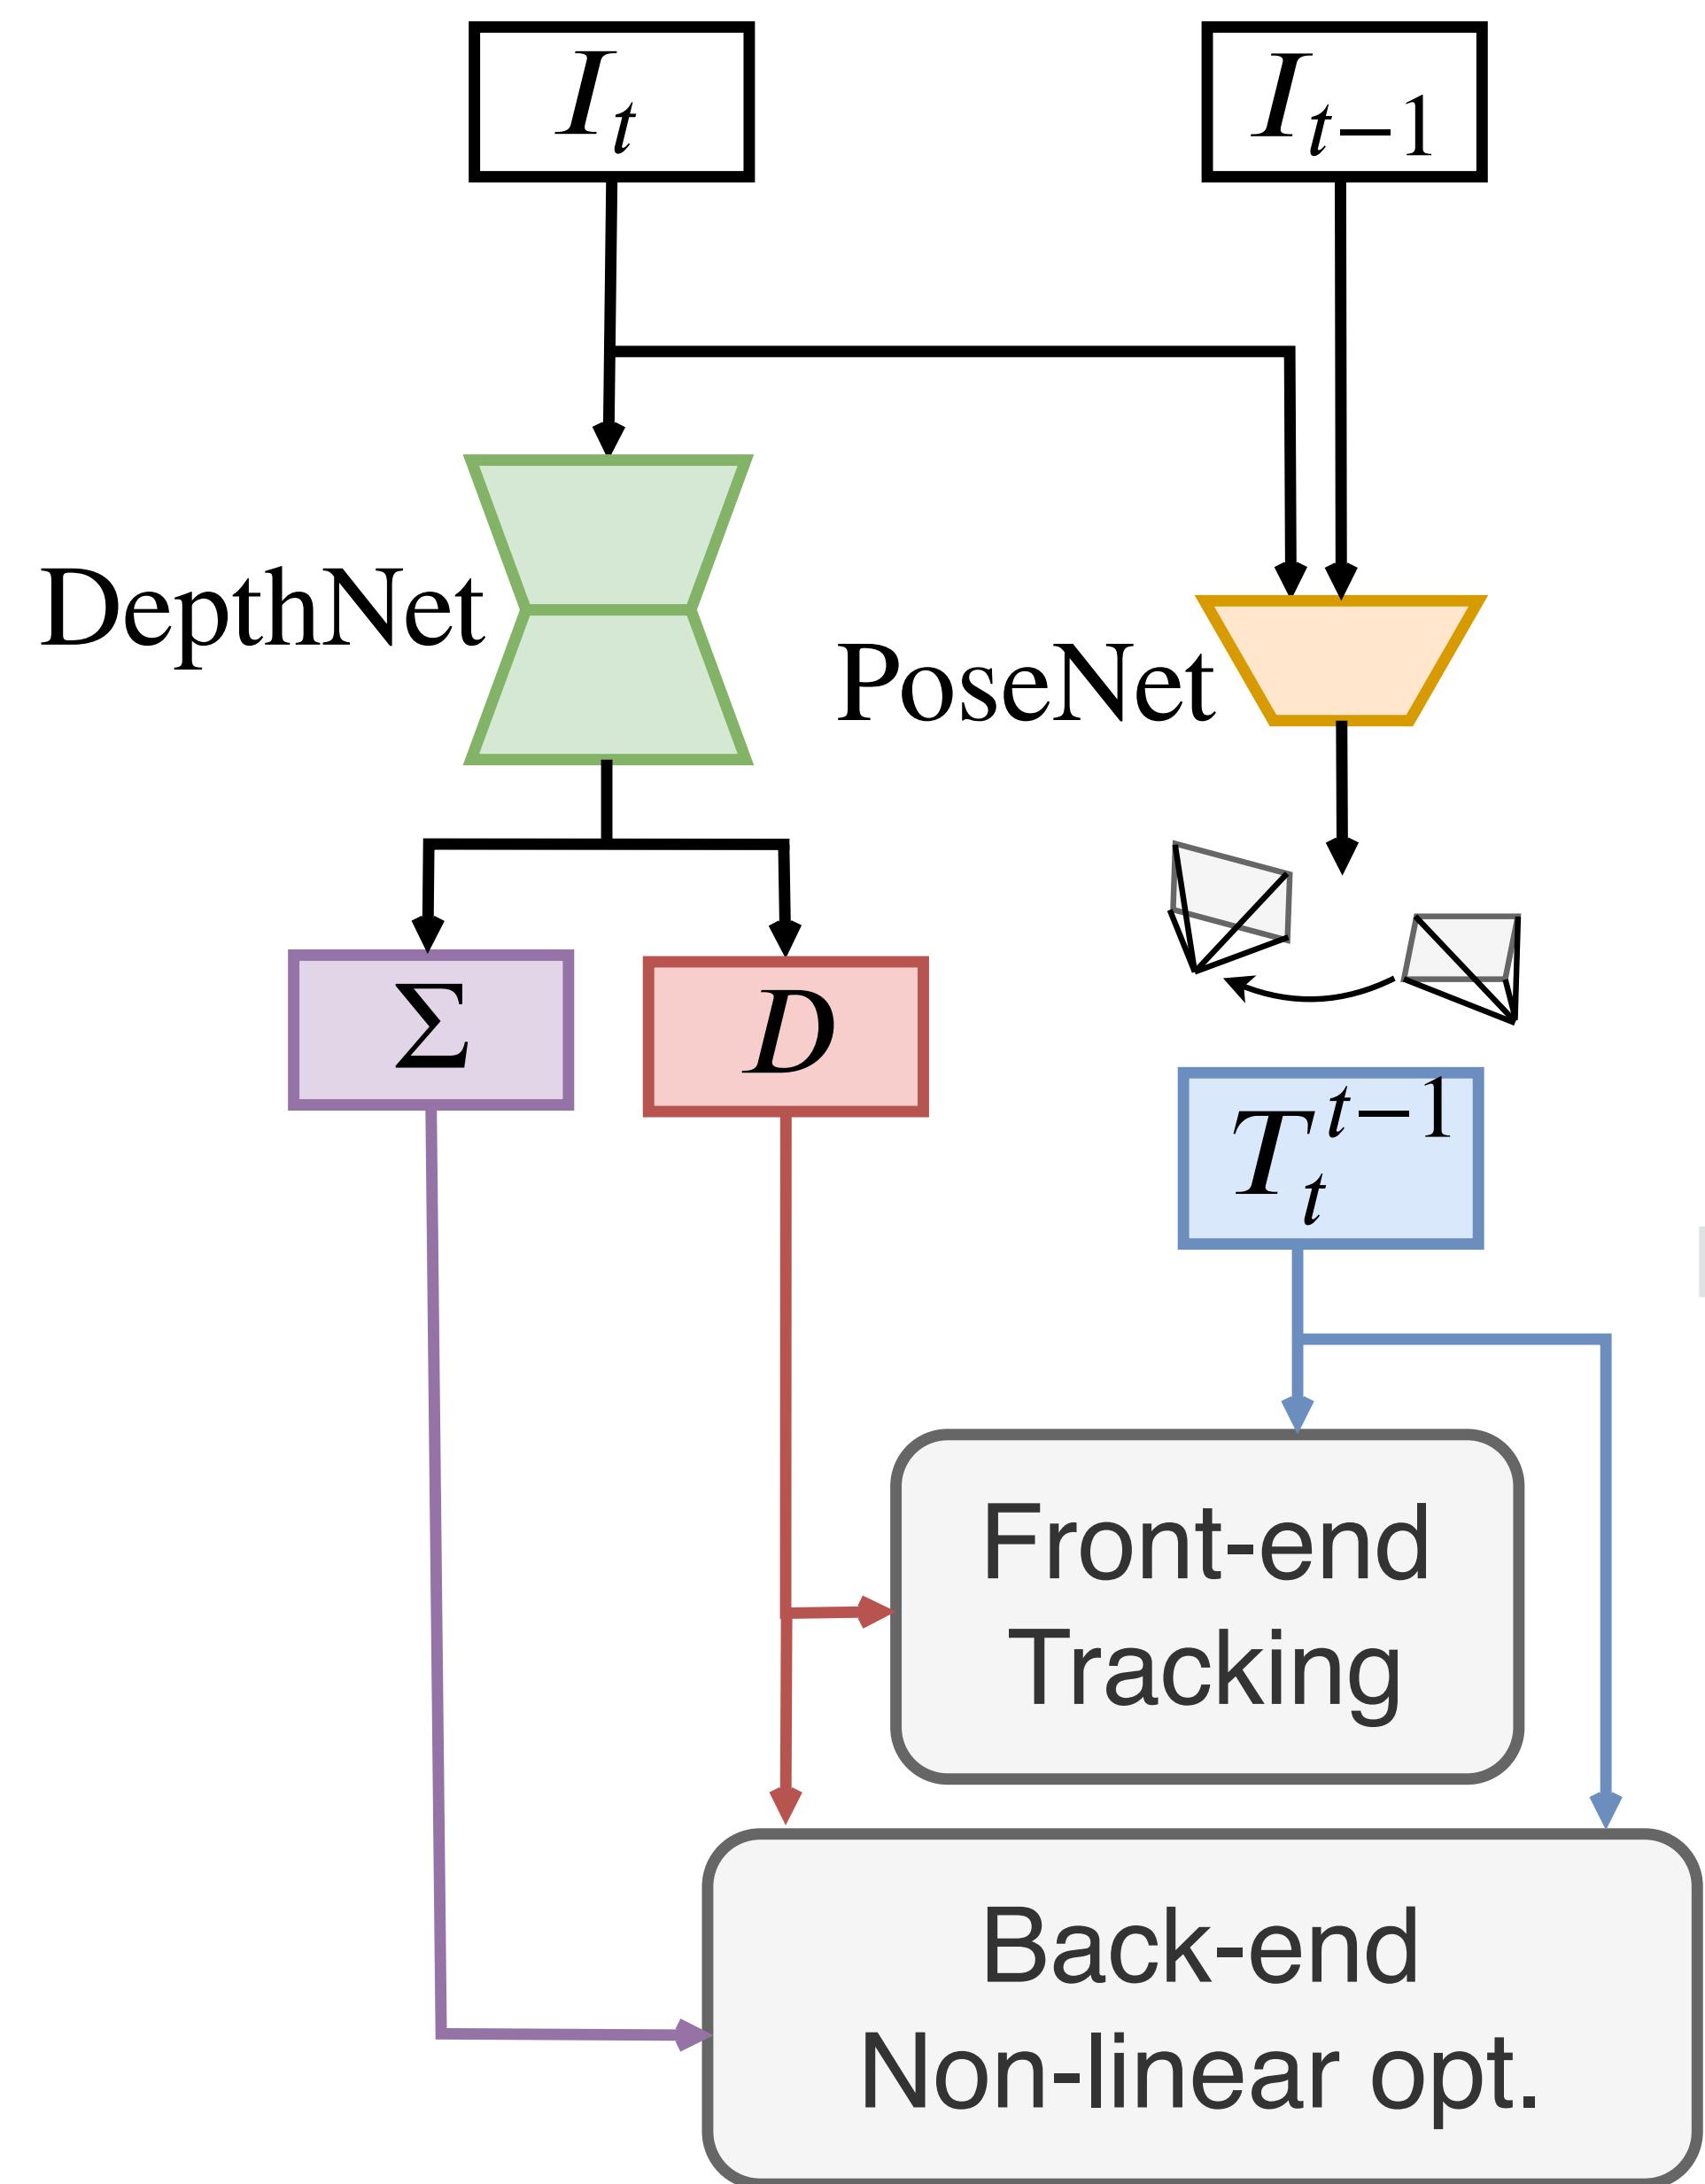
From Optimization to Learning

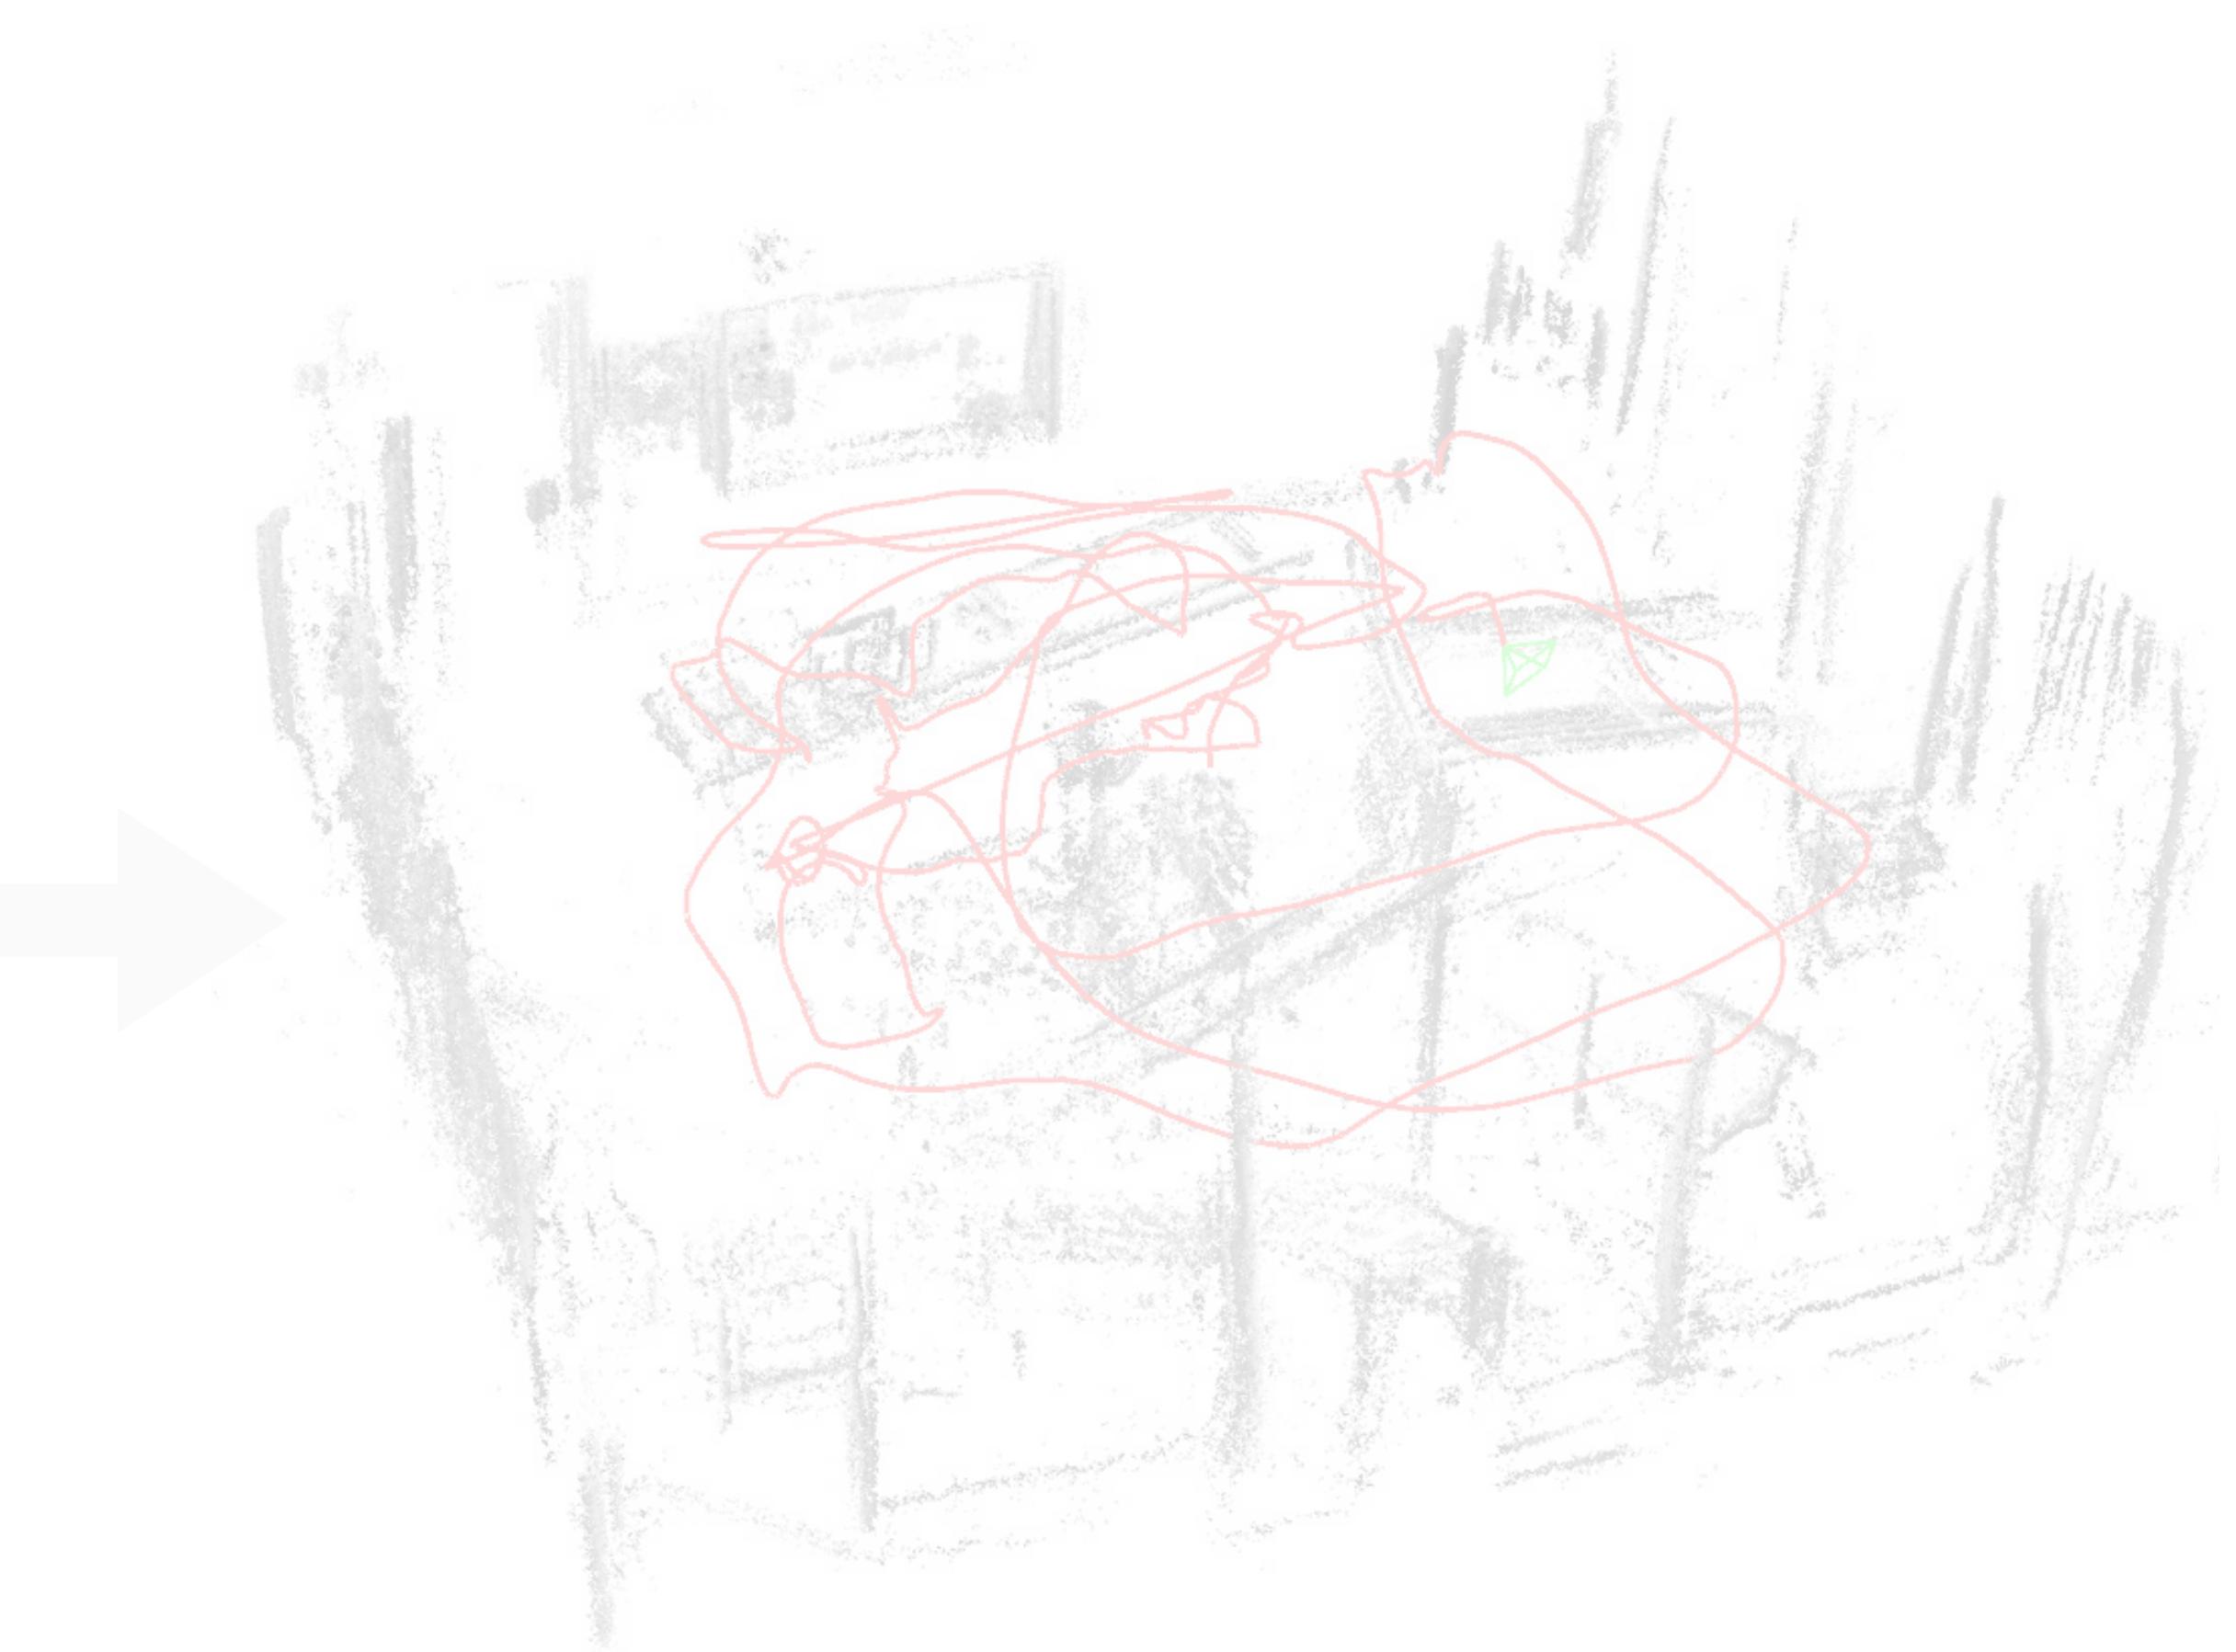
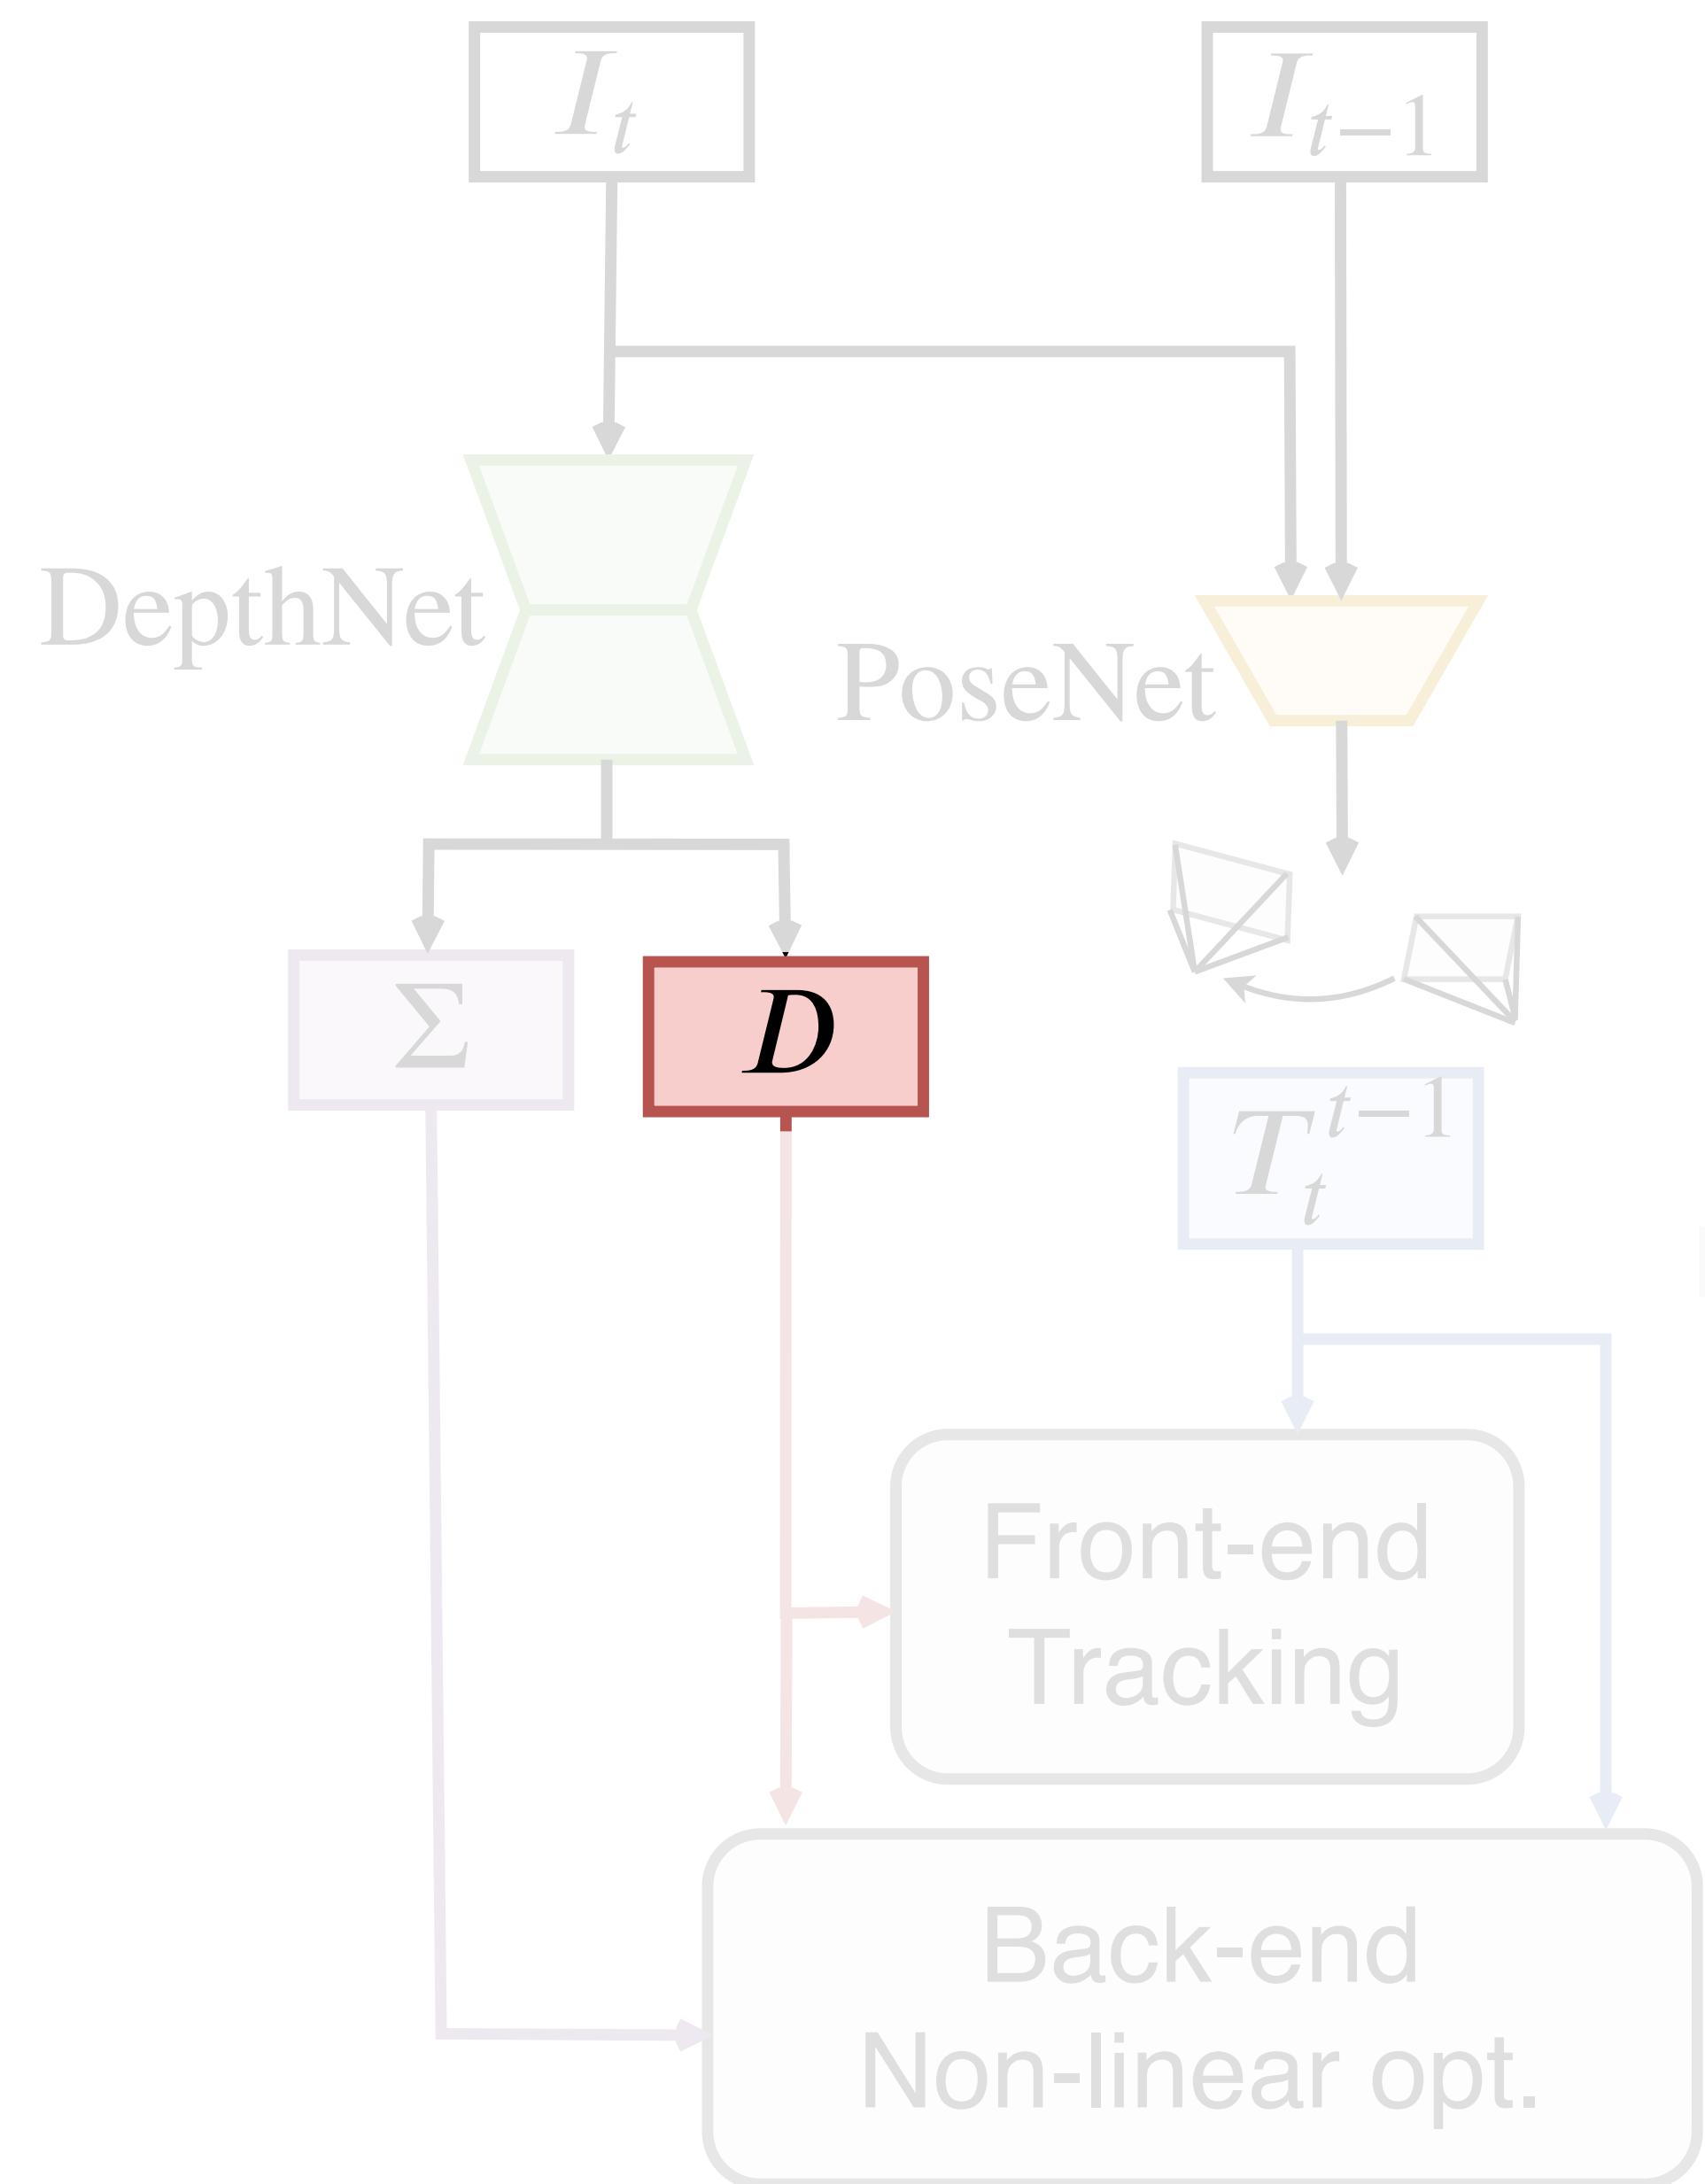


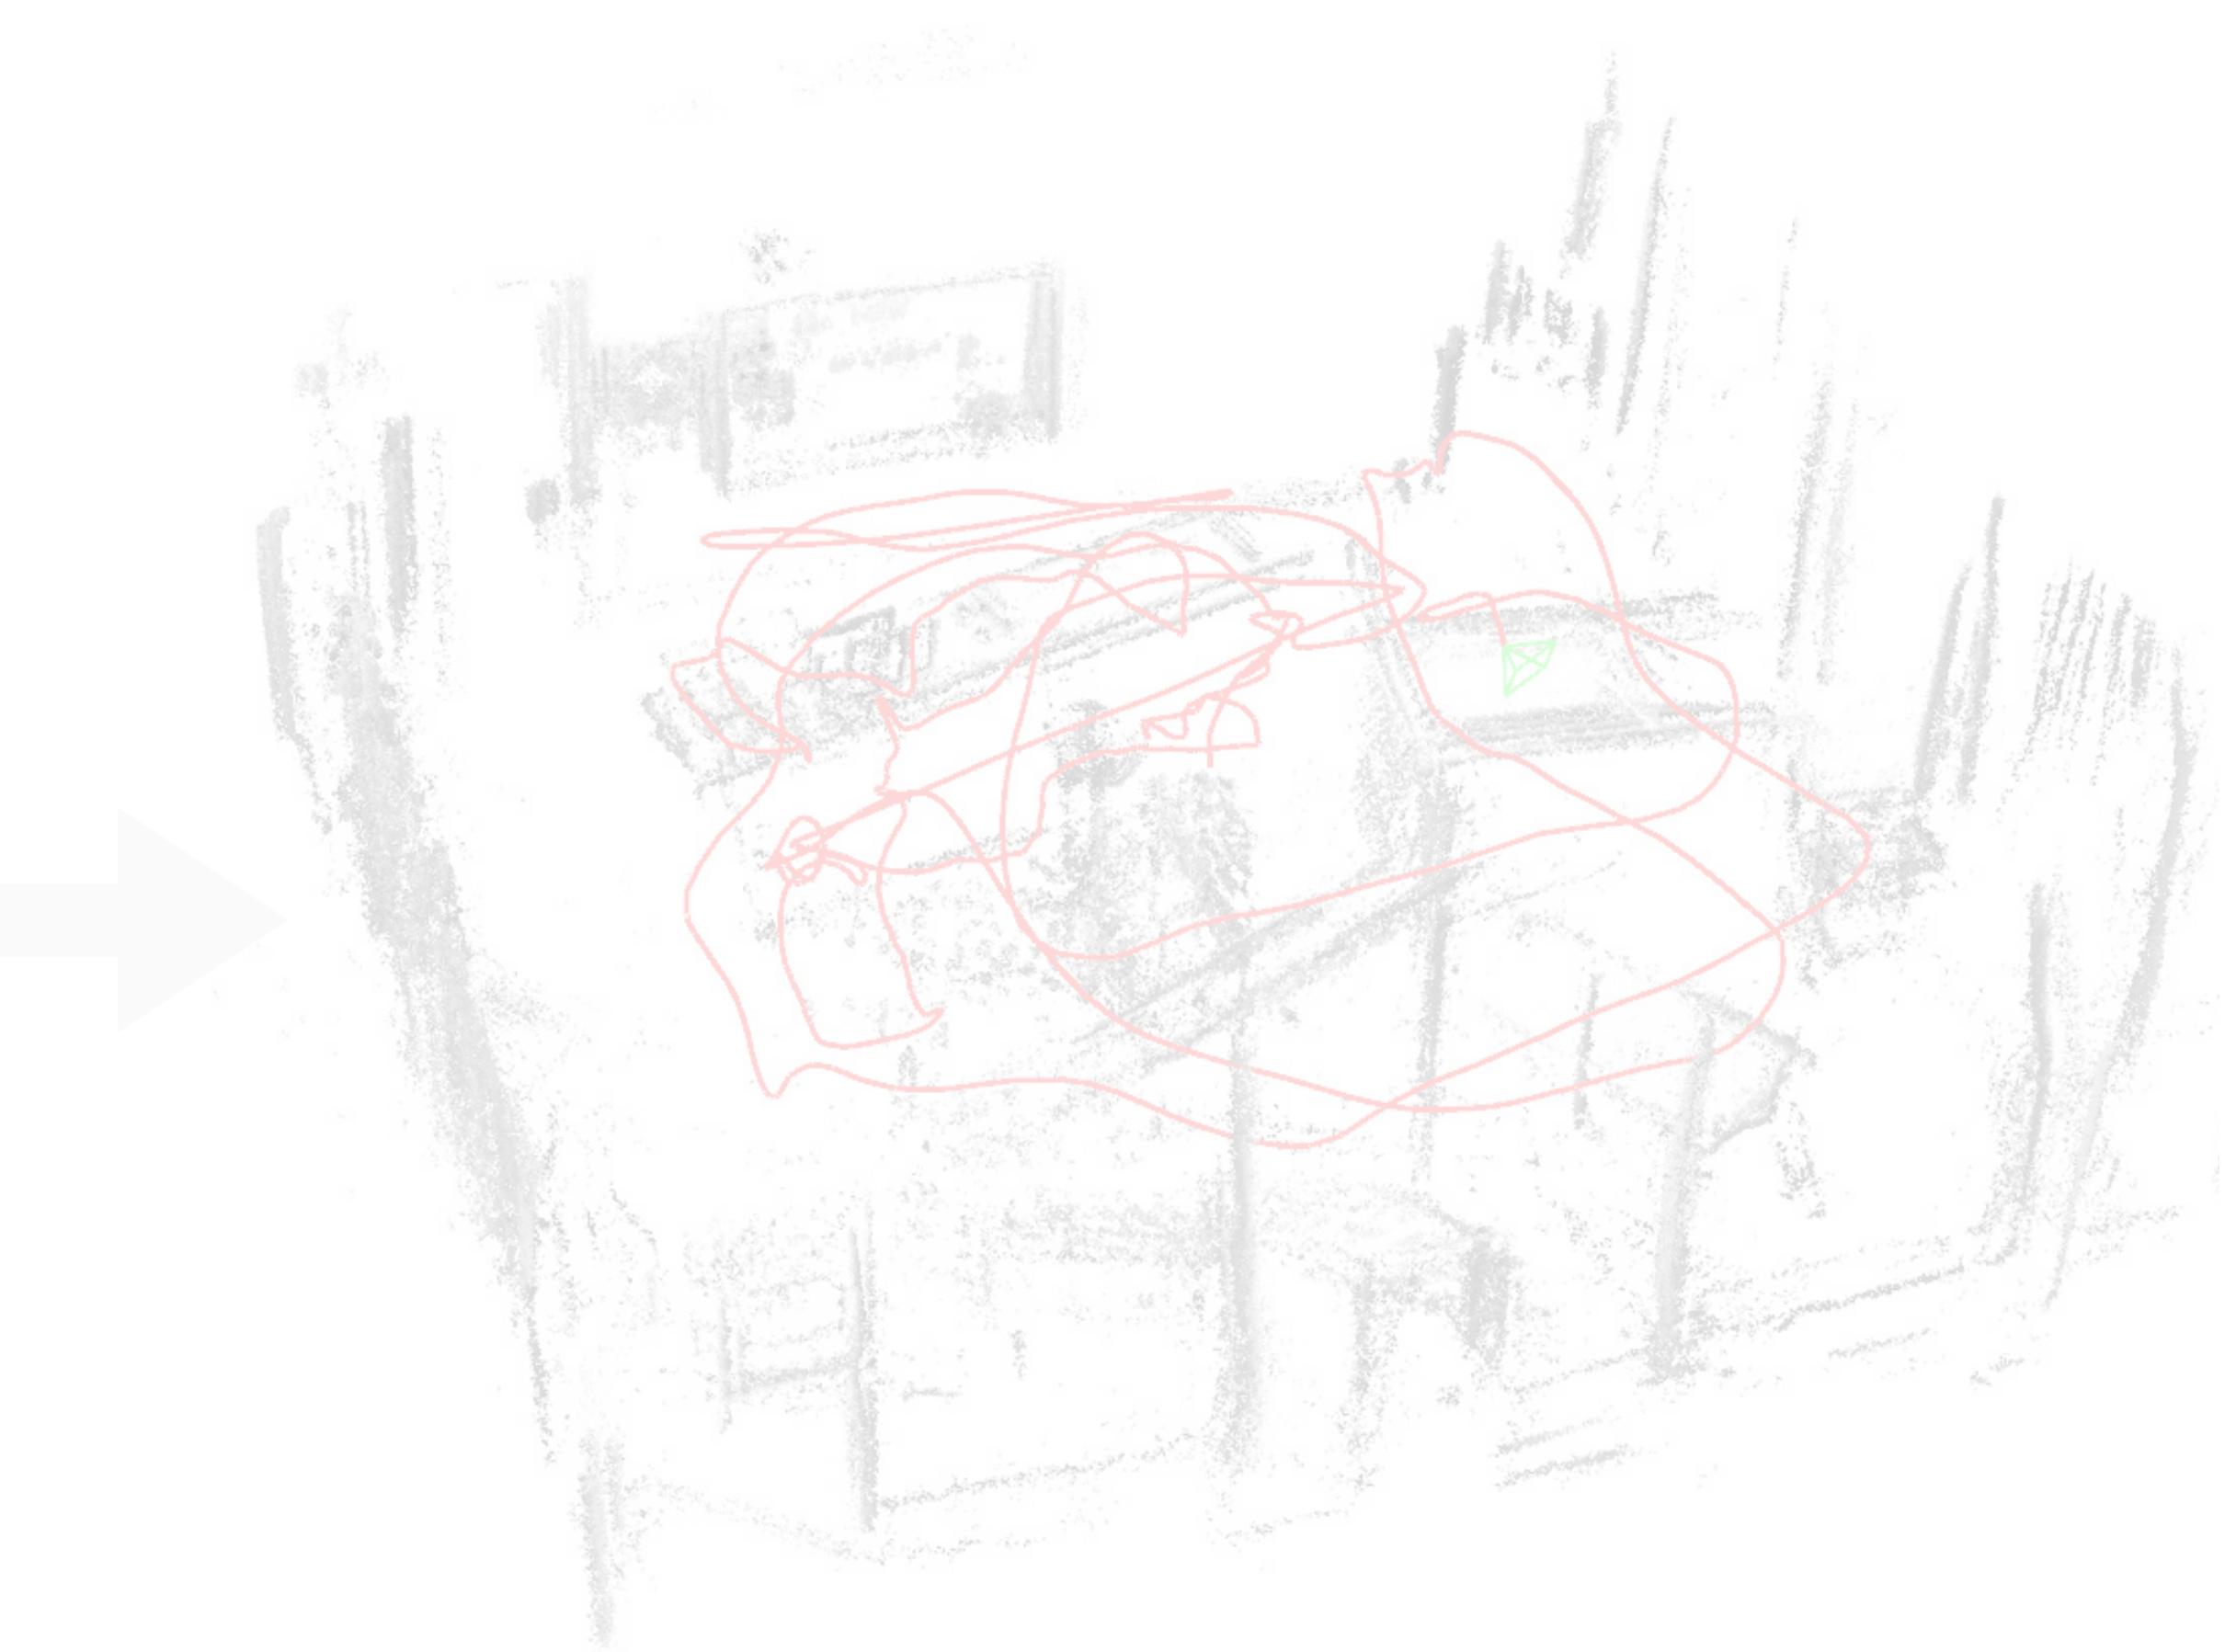
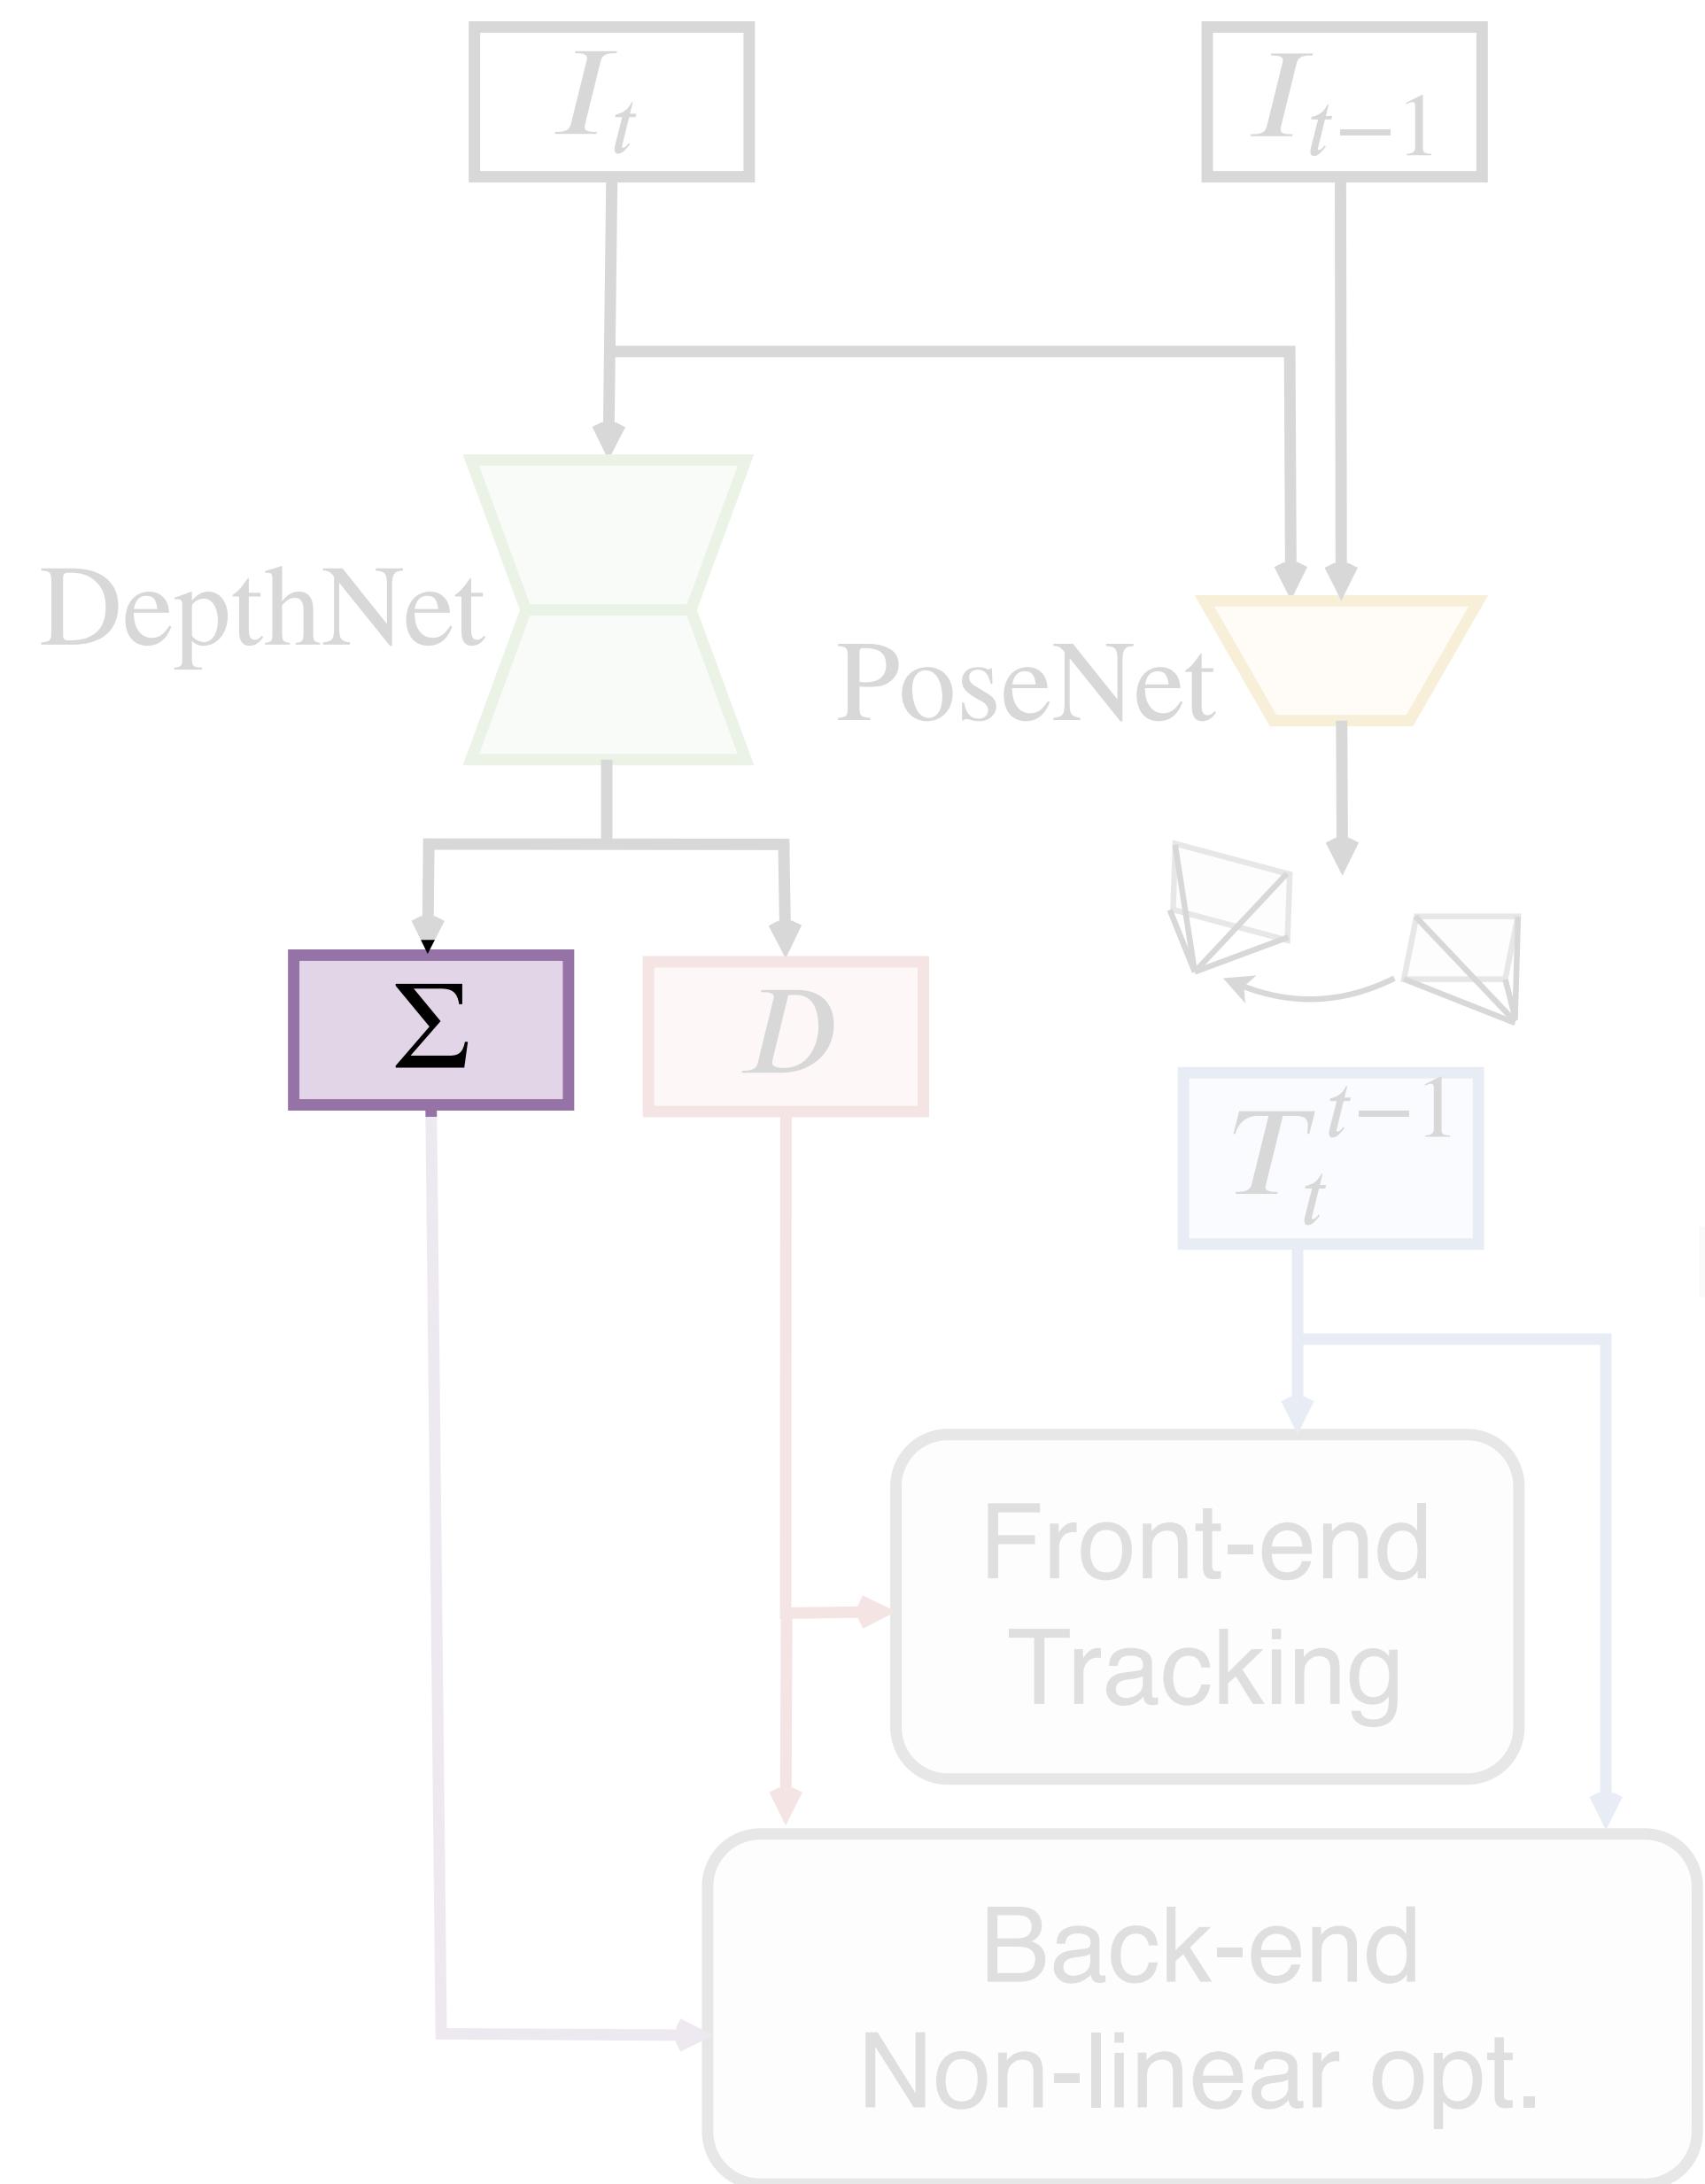
Visual SLAM

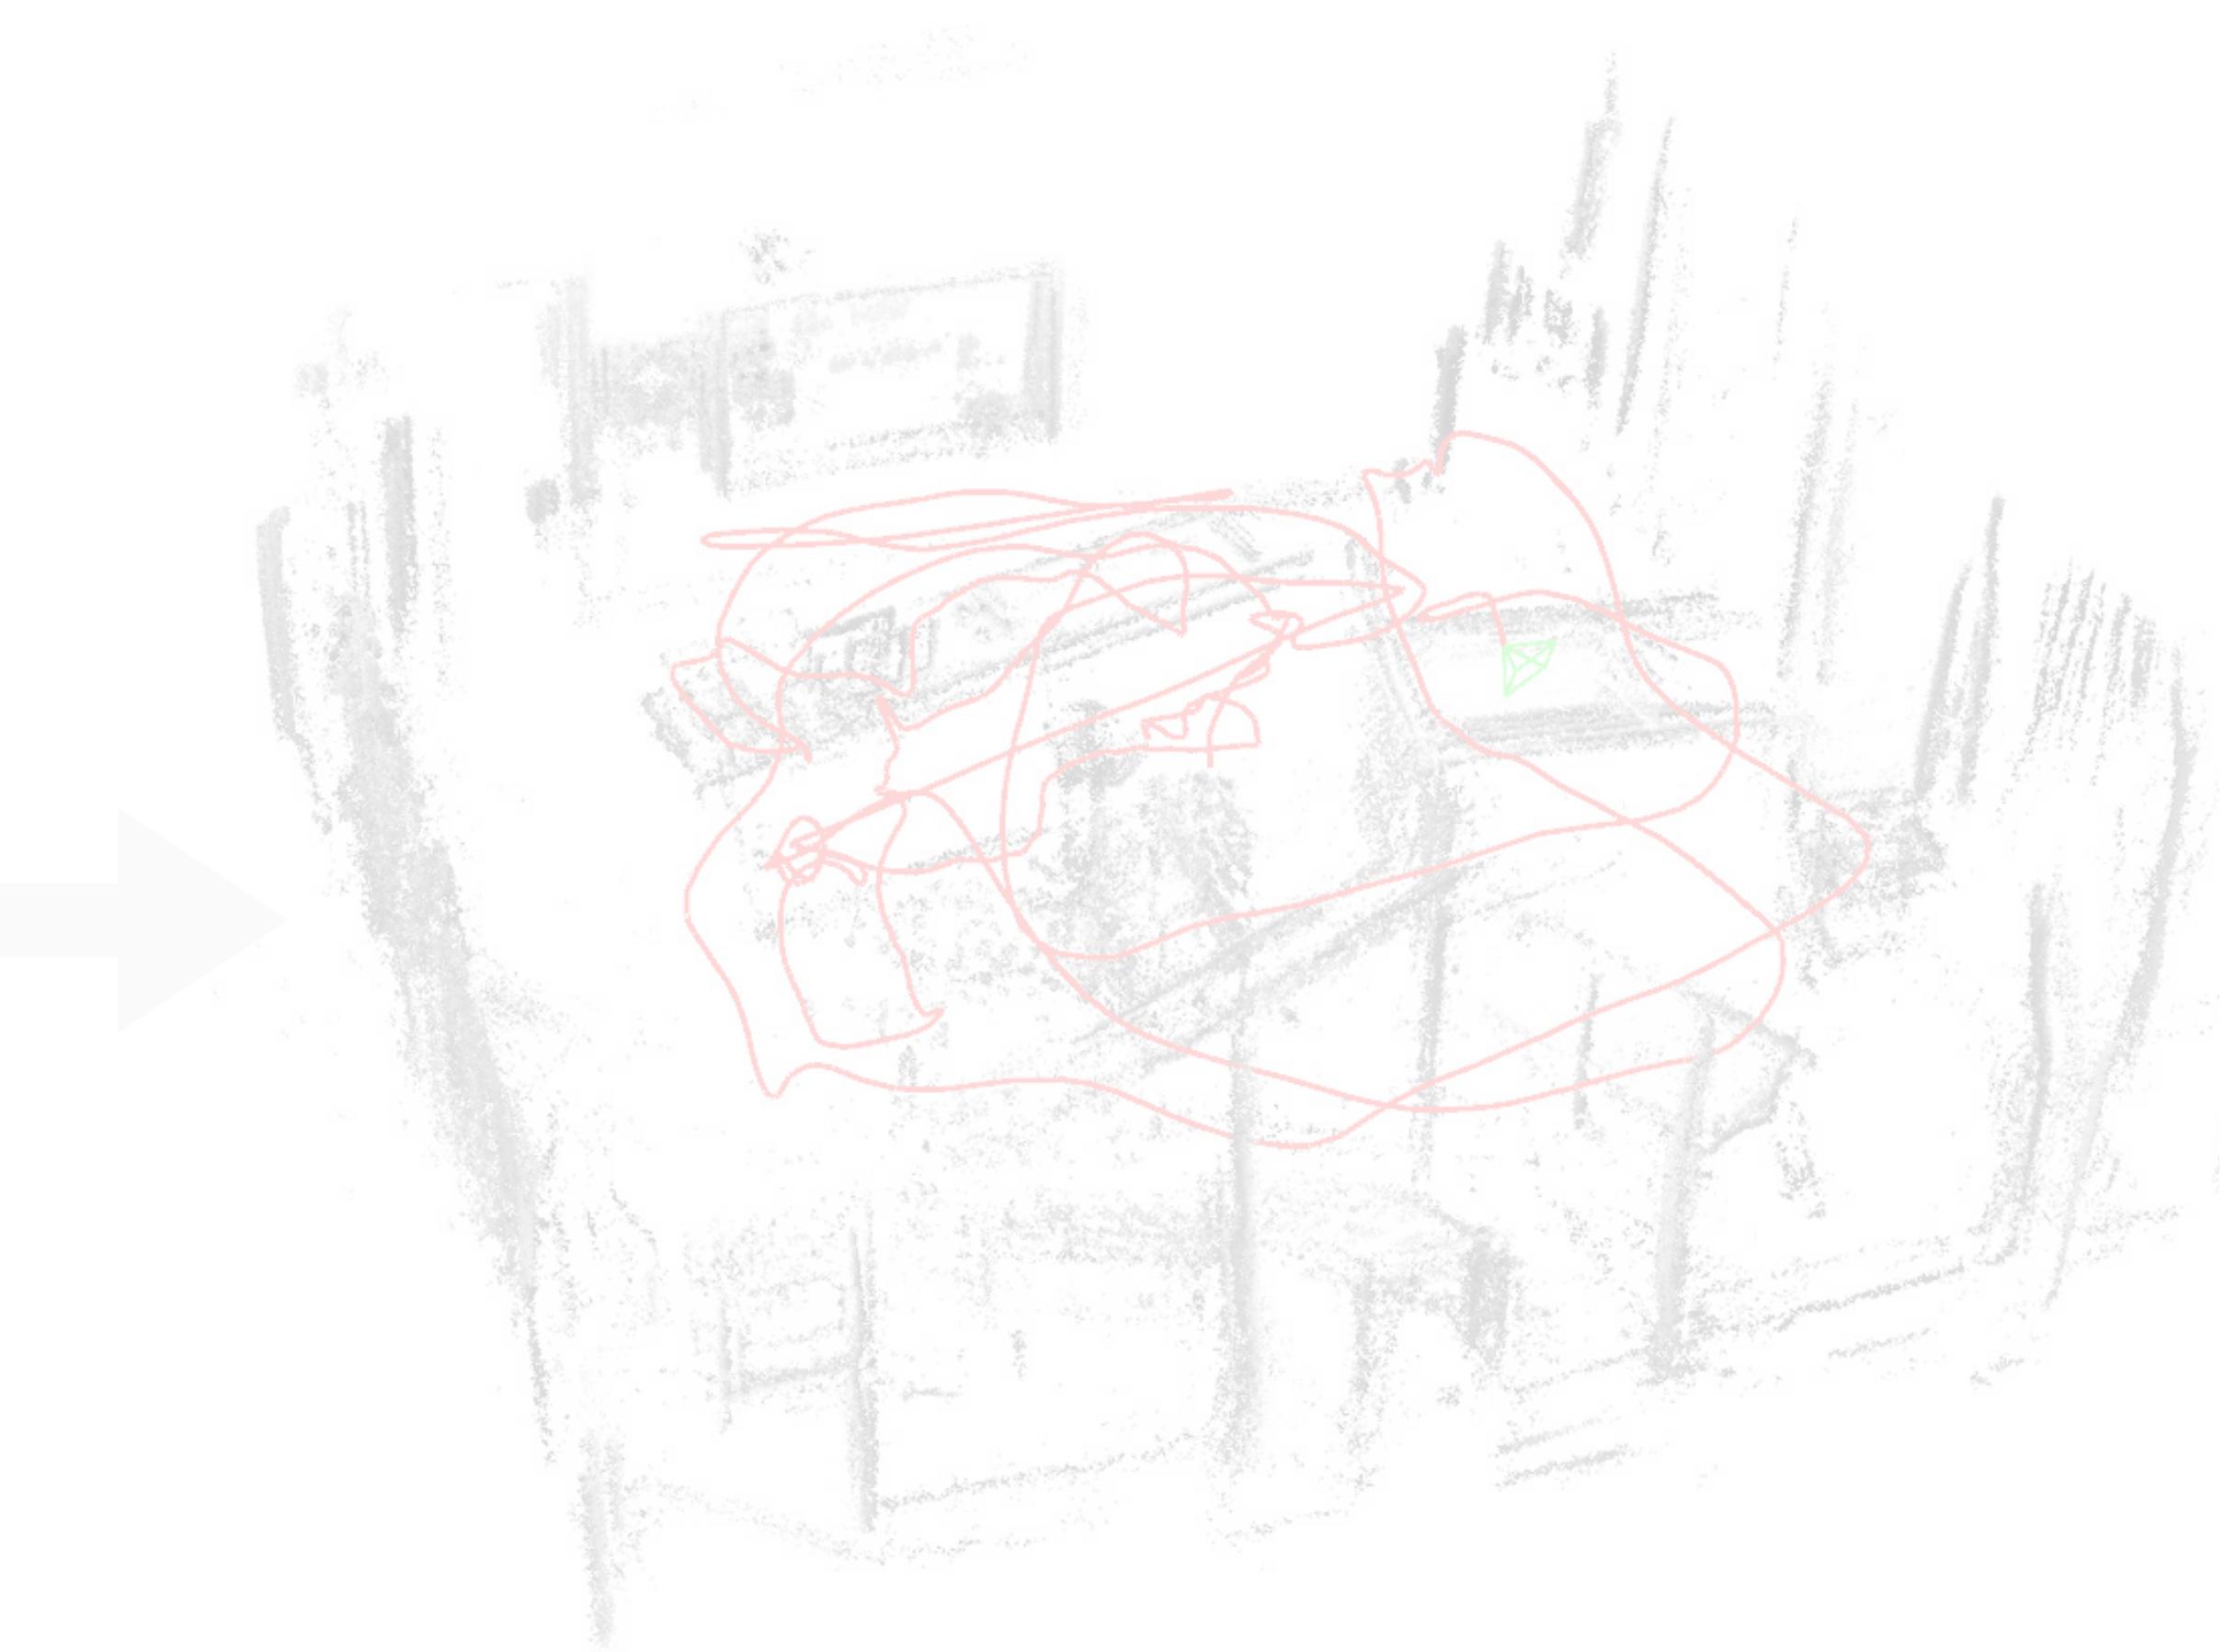
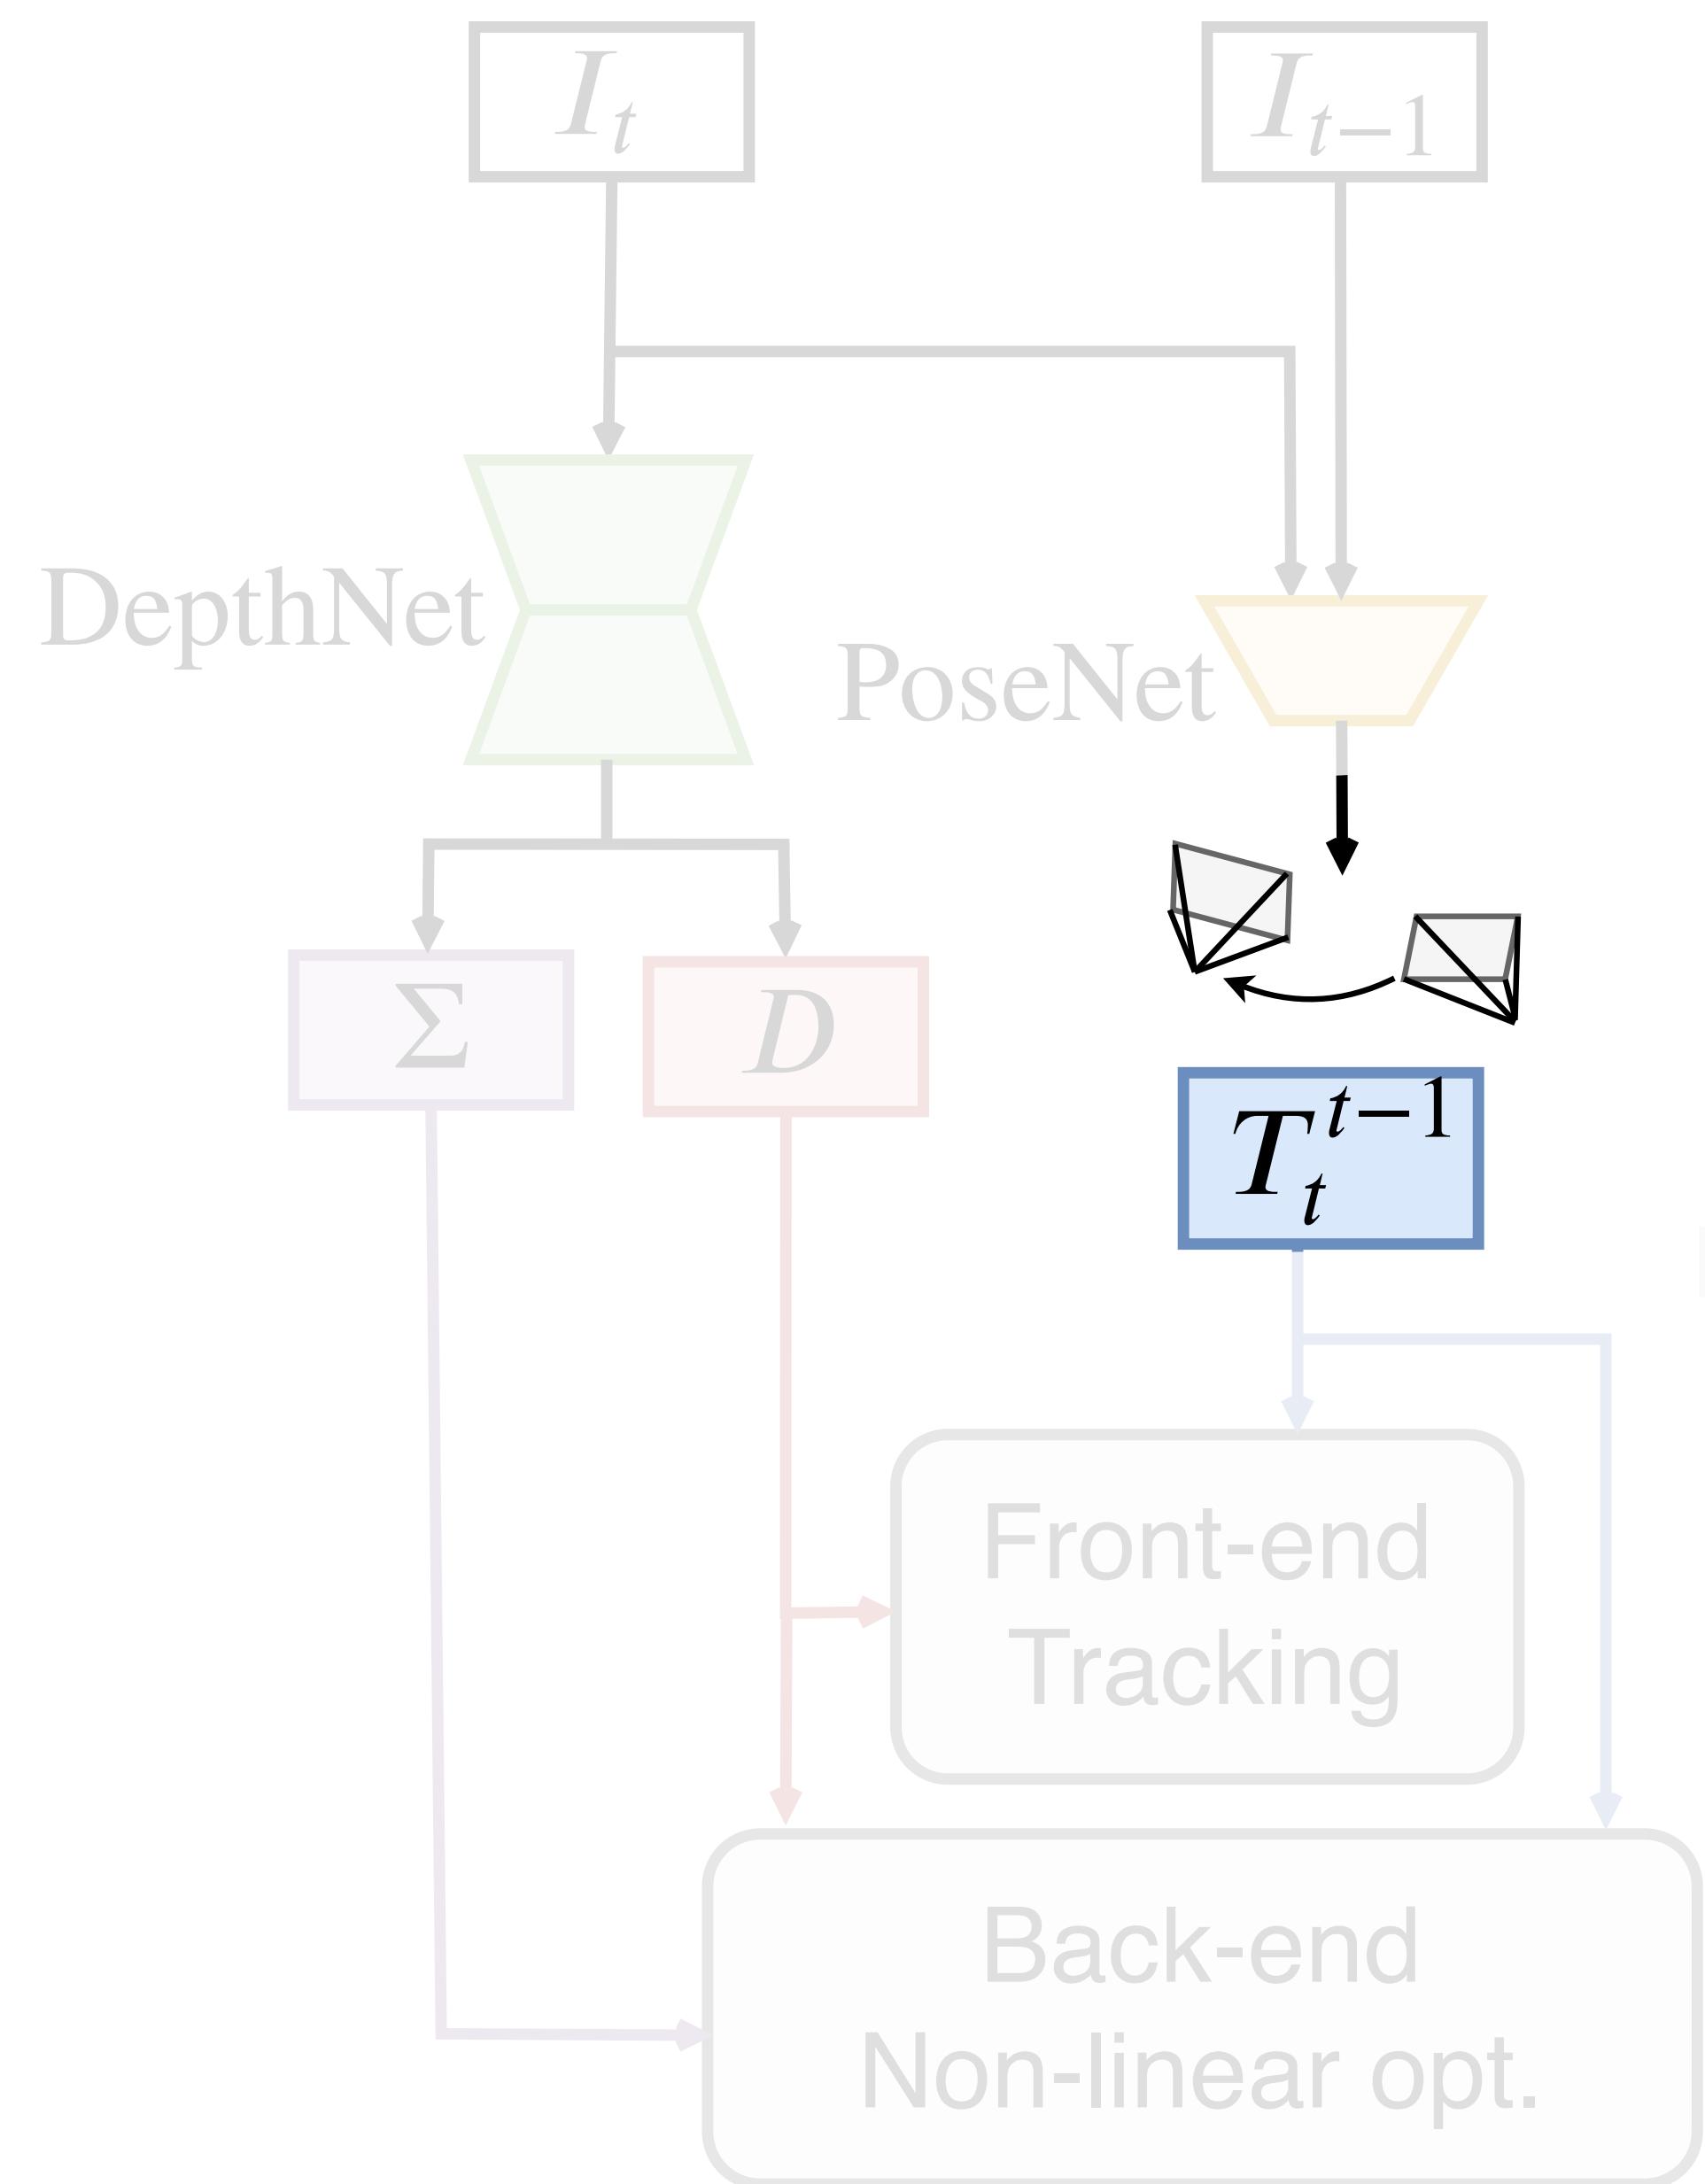
From Optimization to Learning

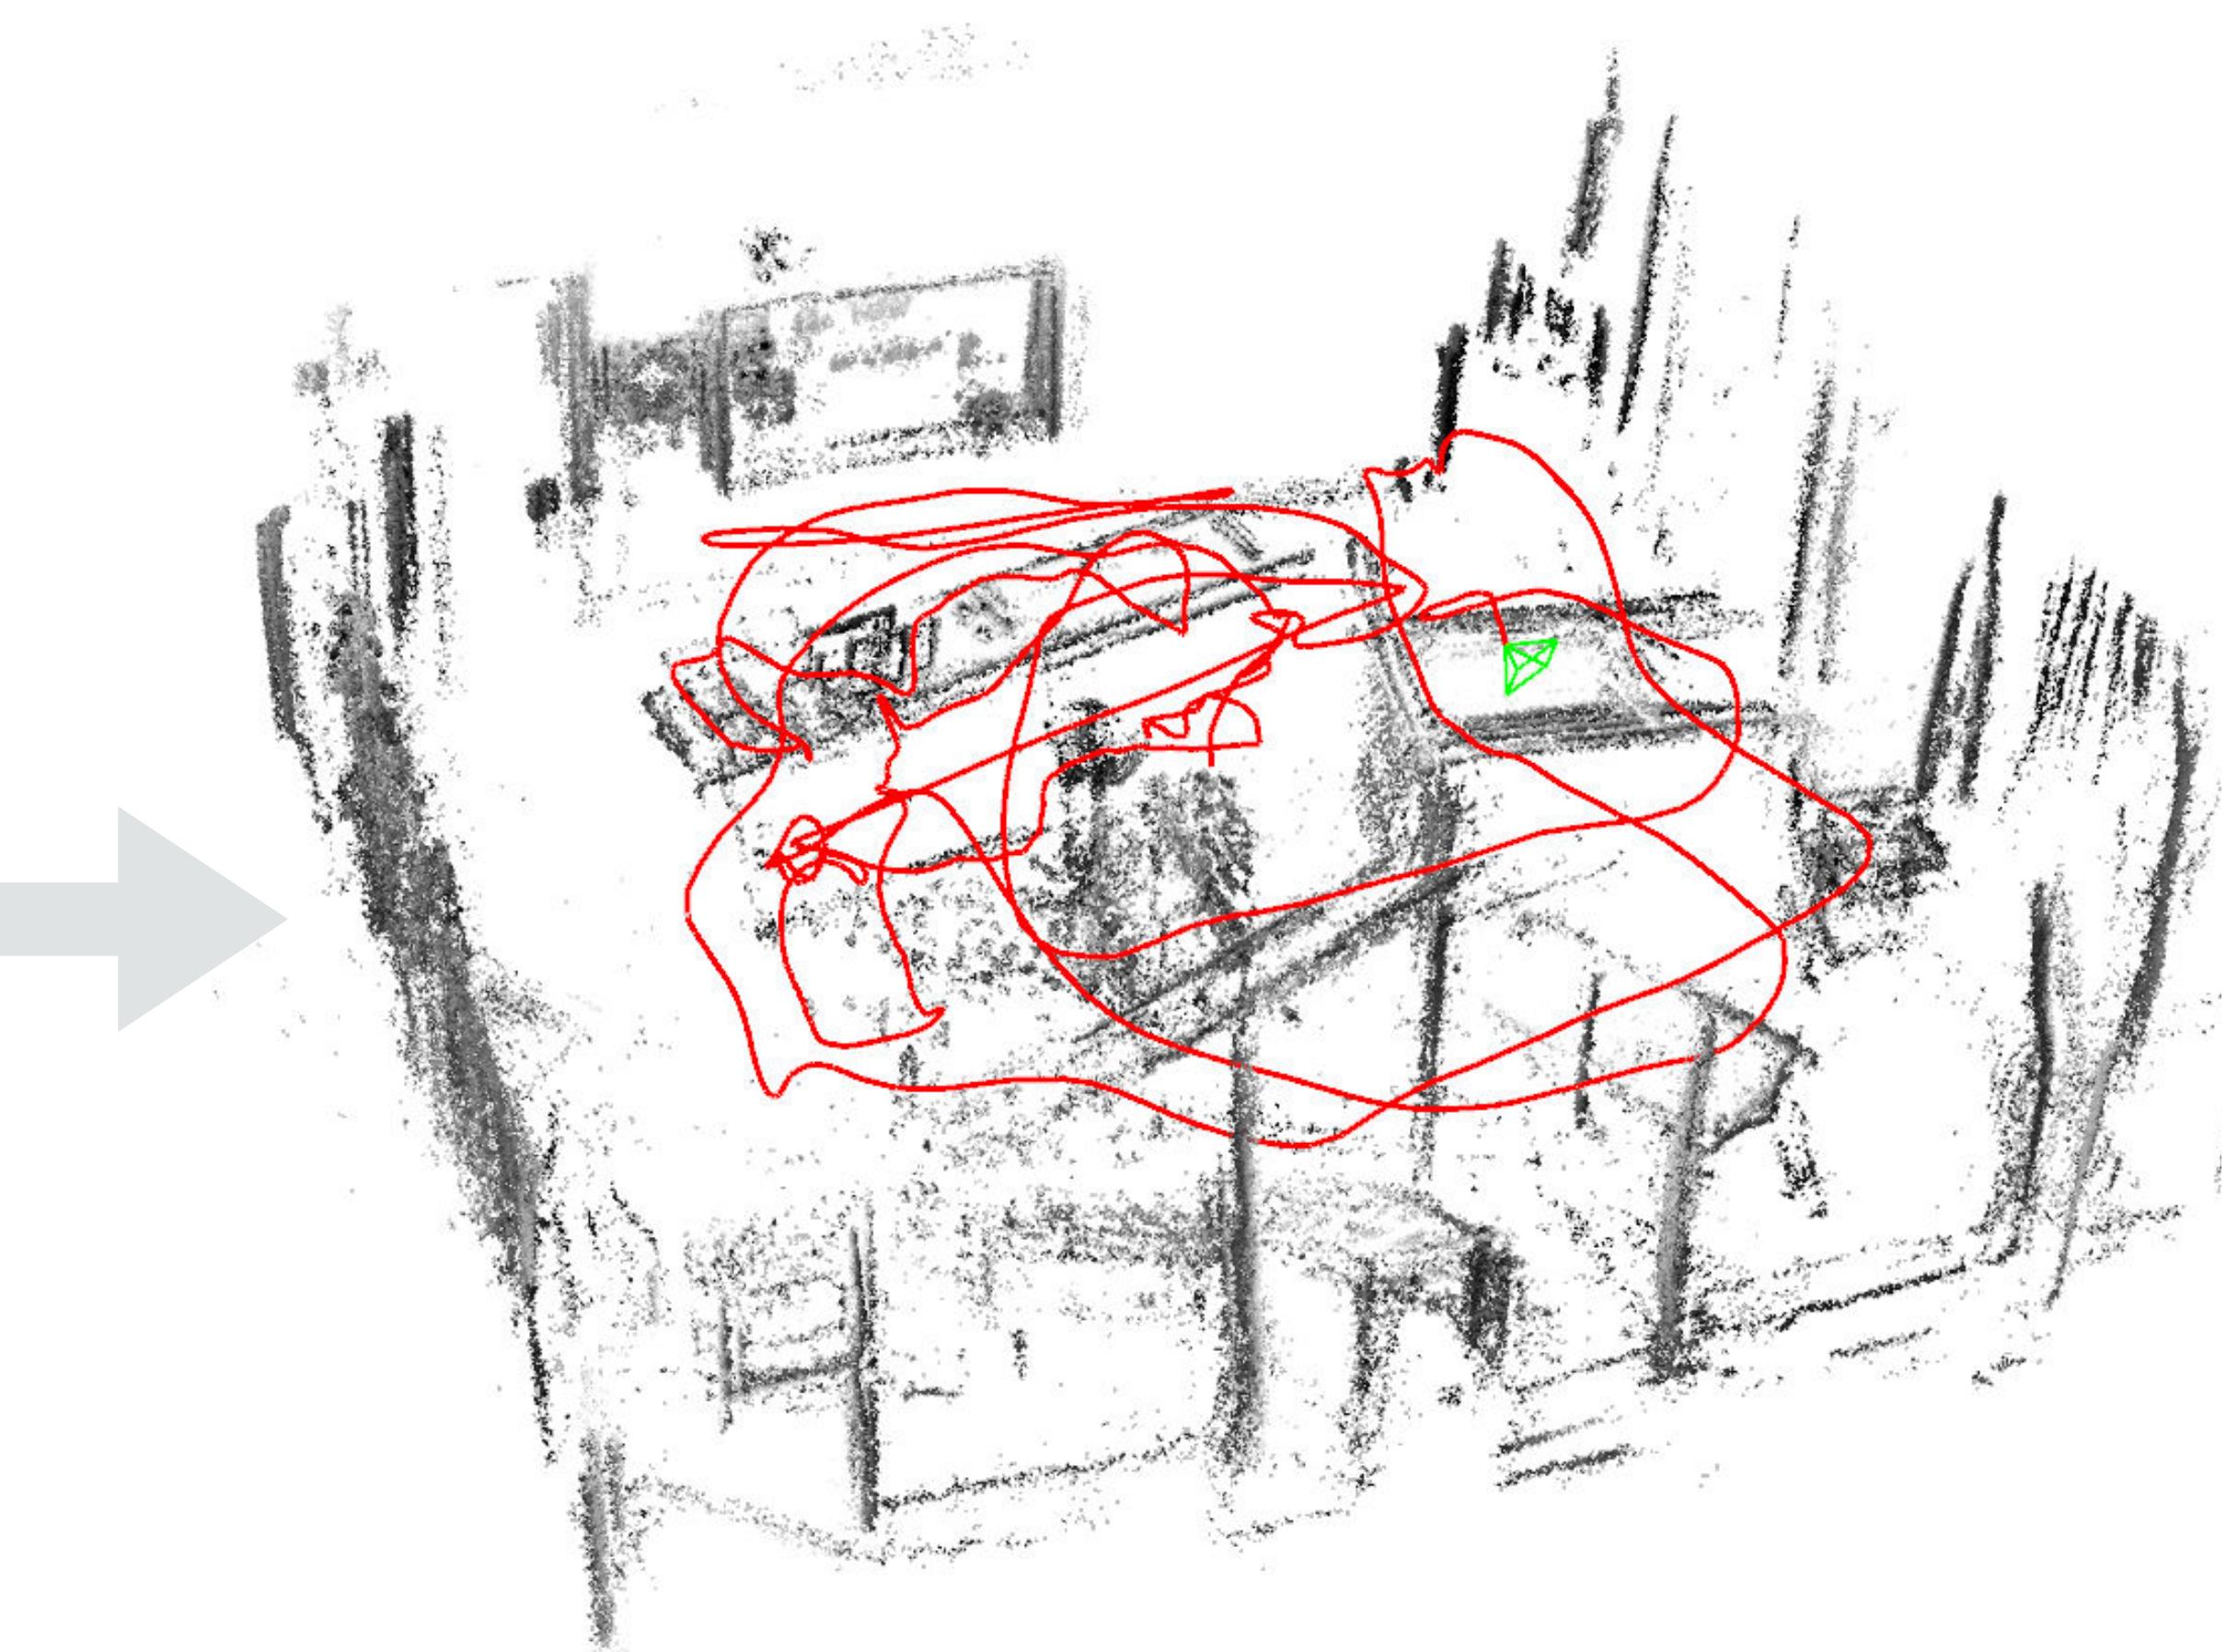
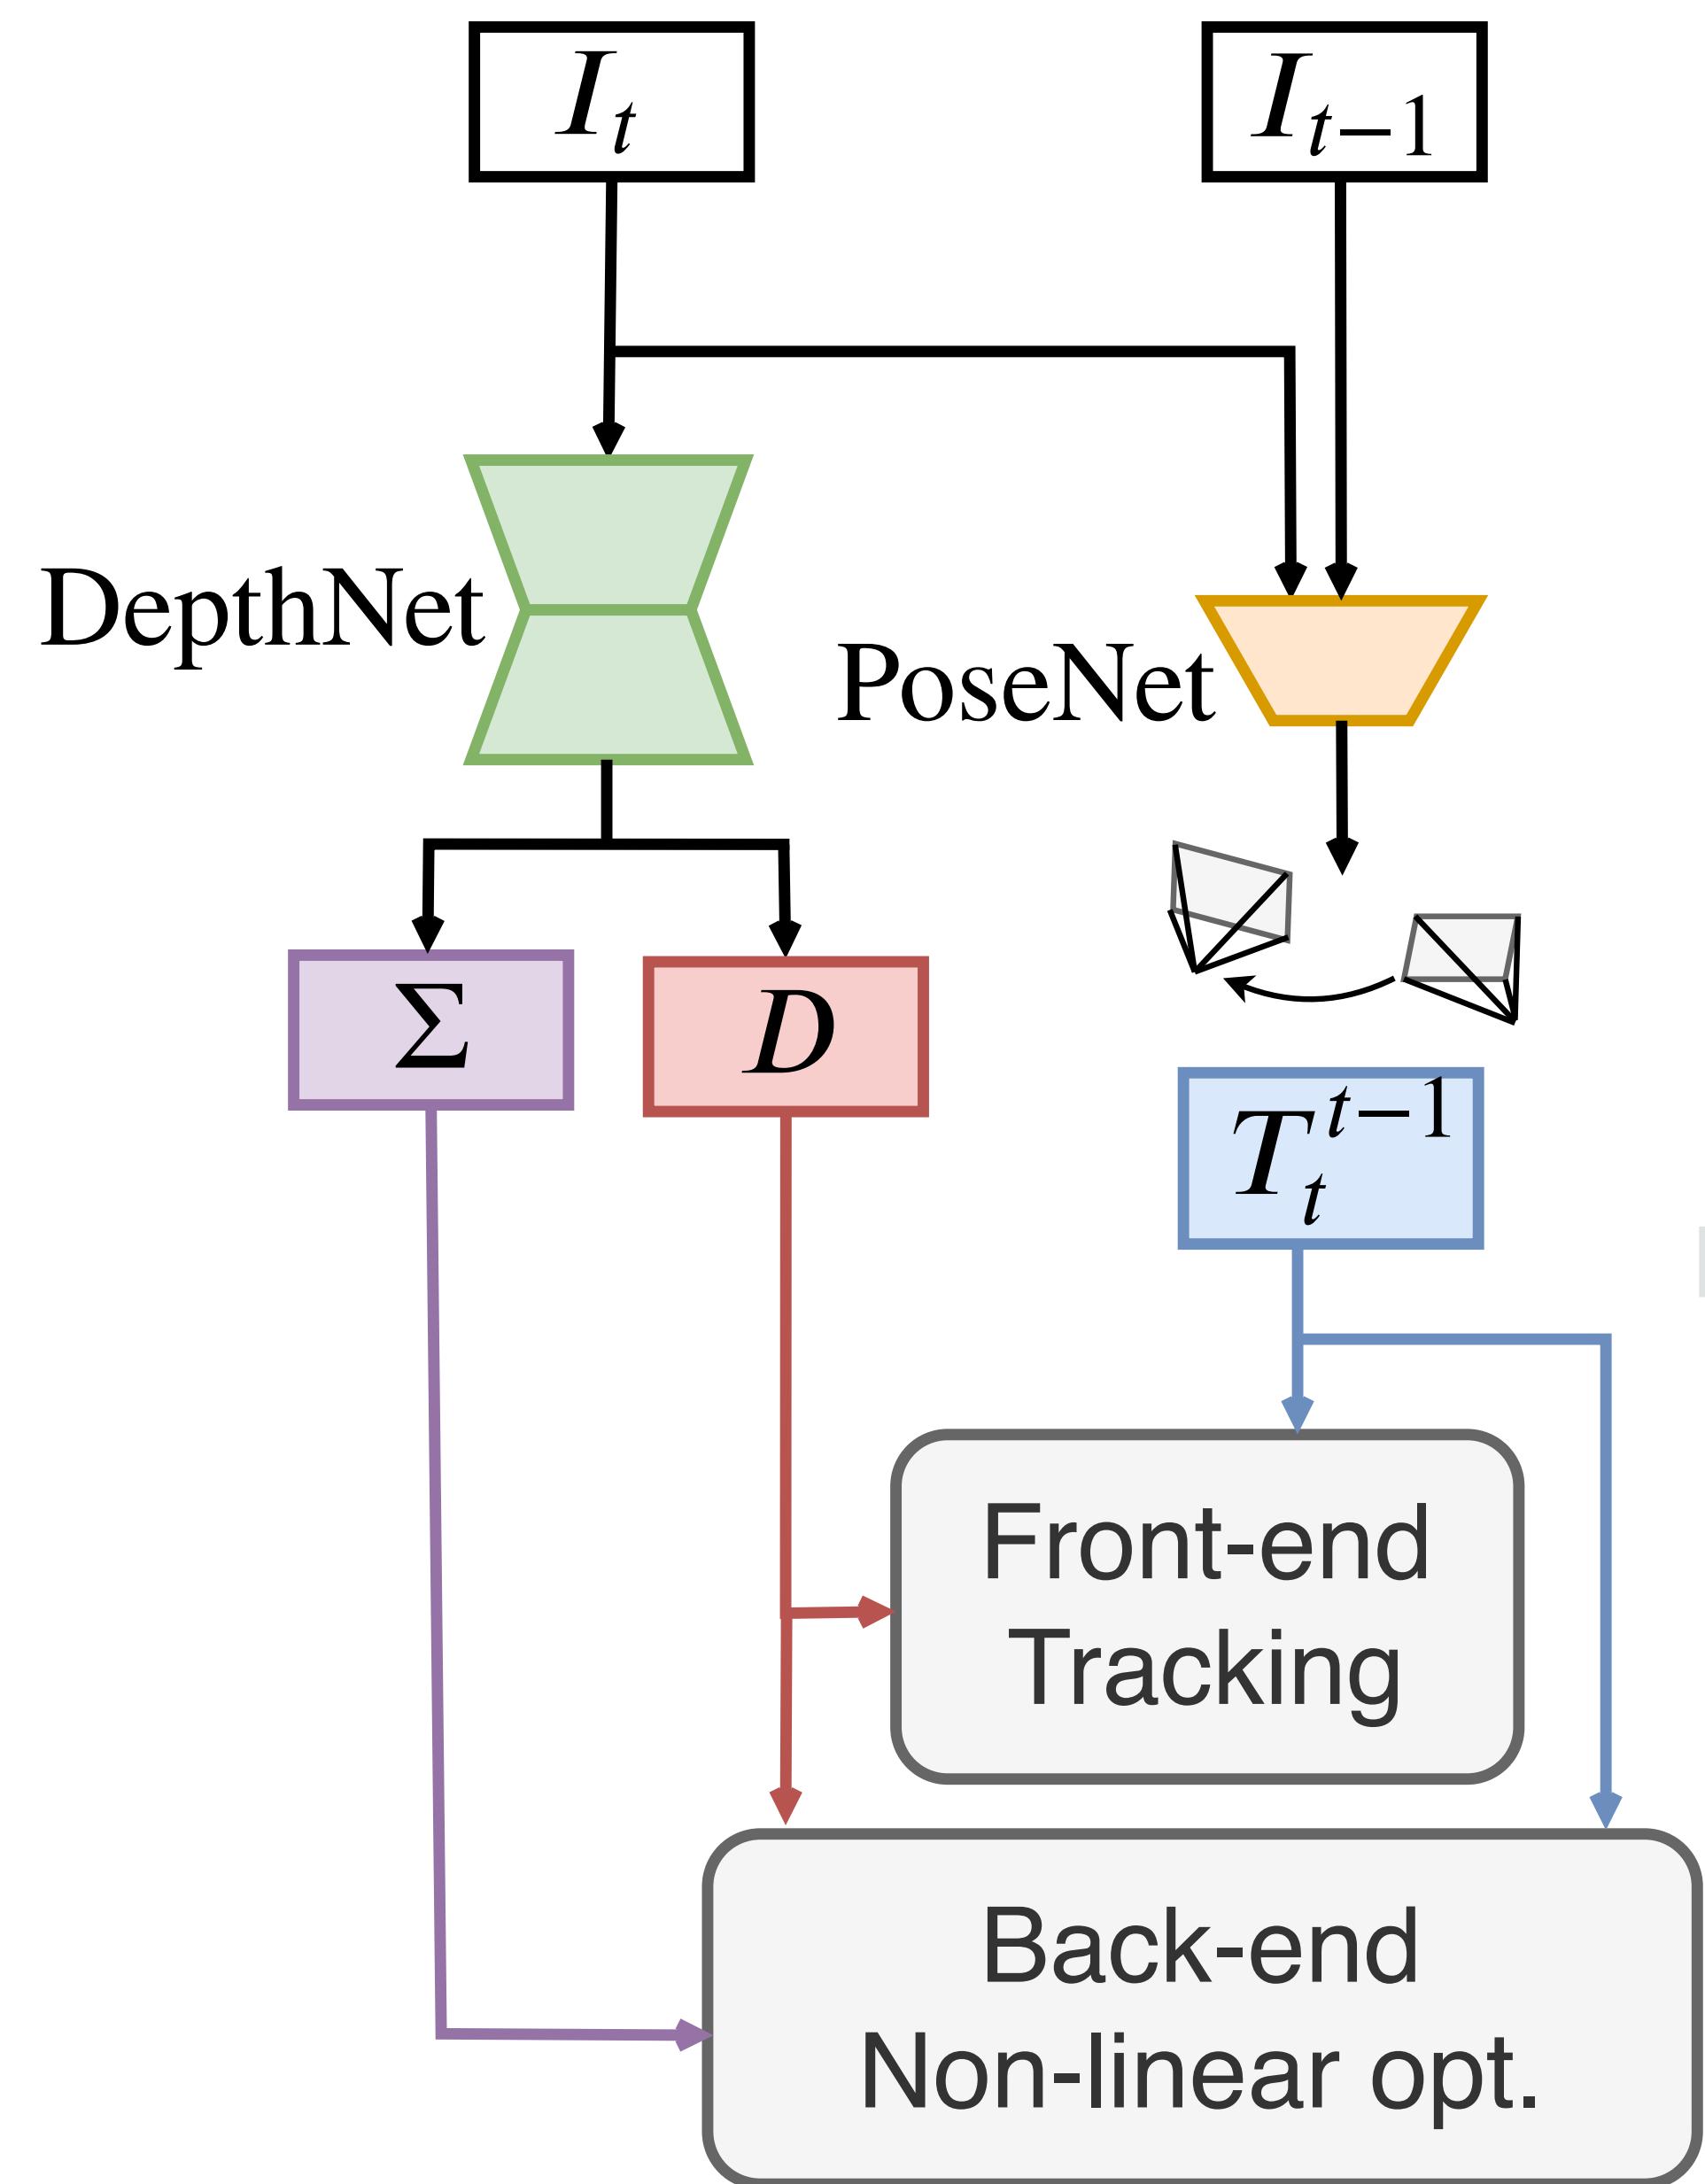


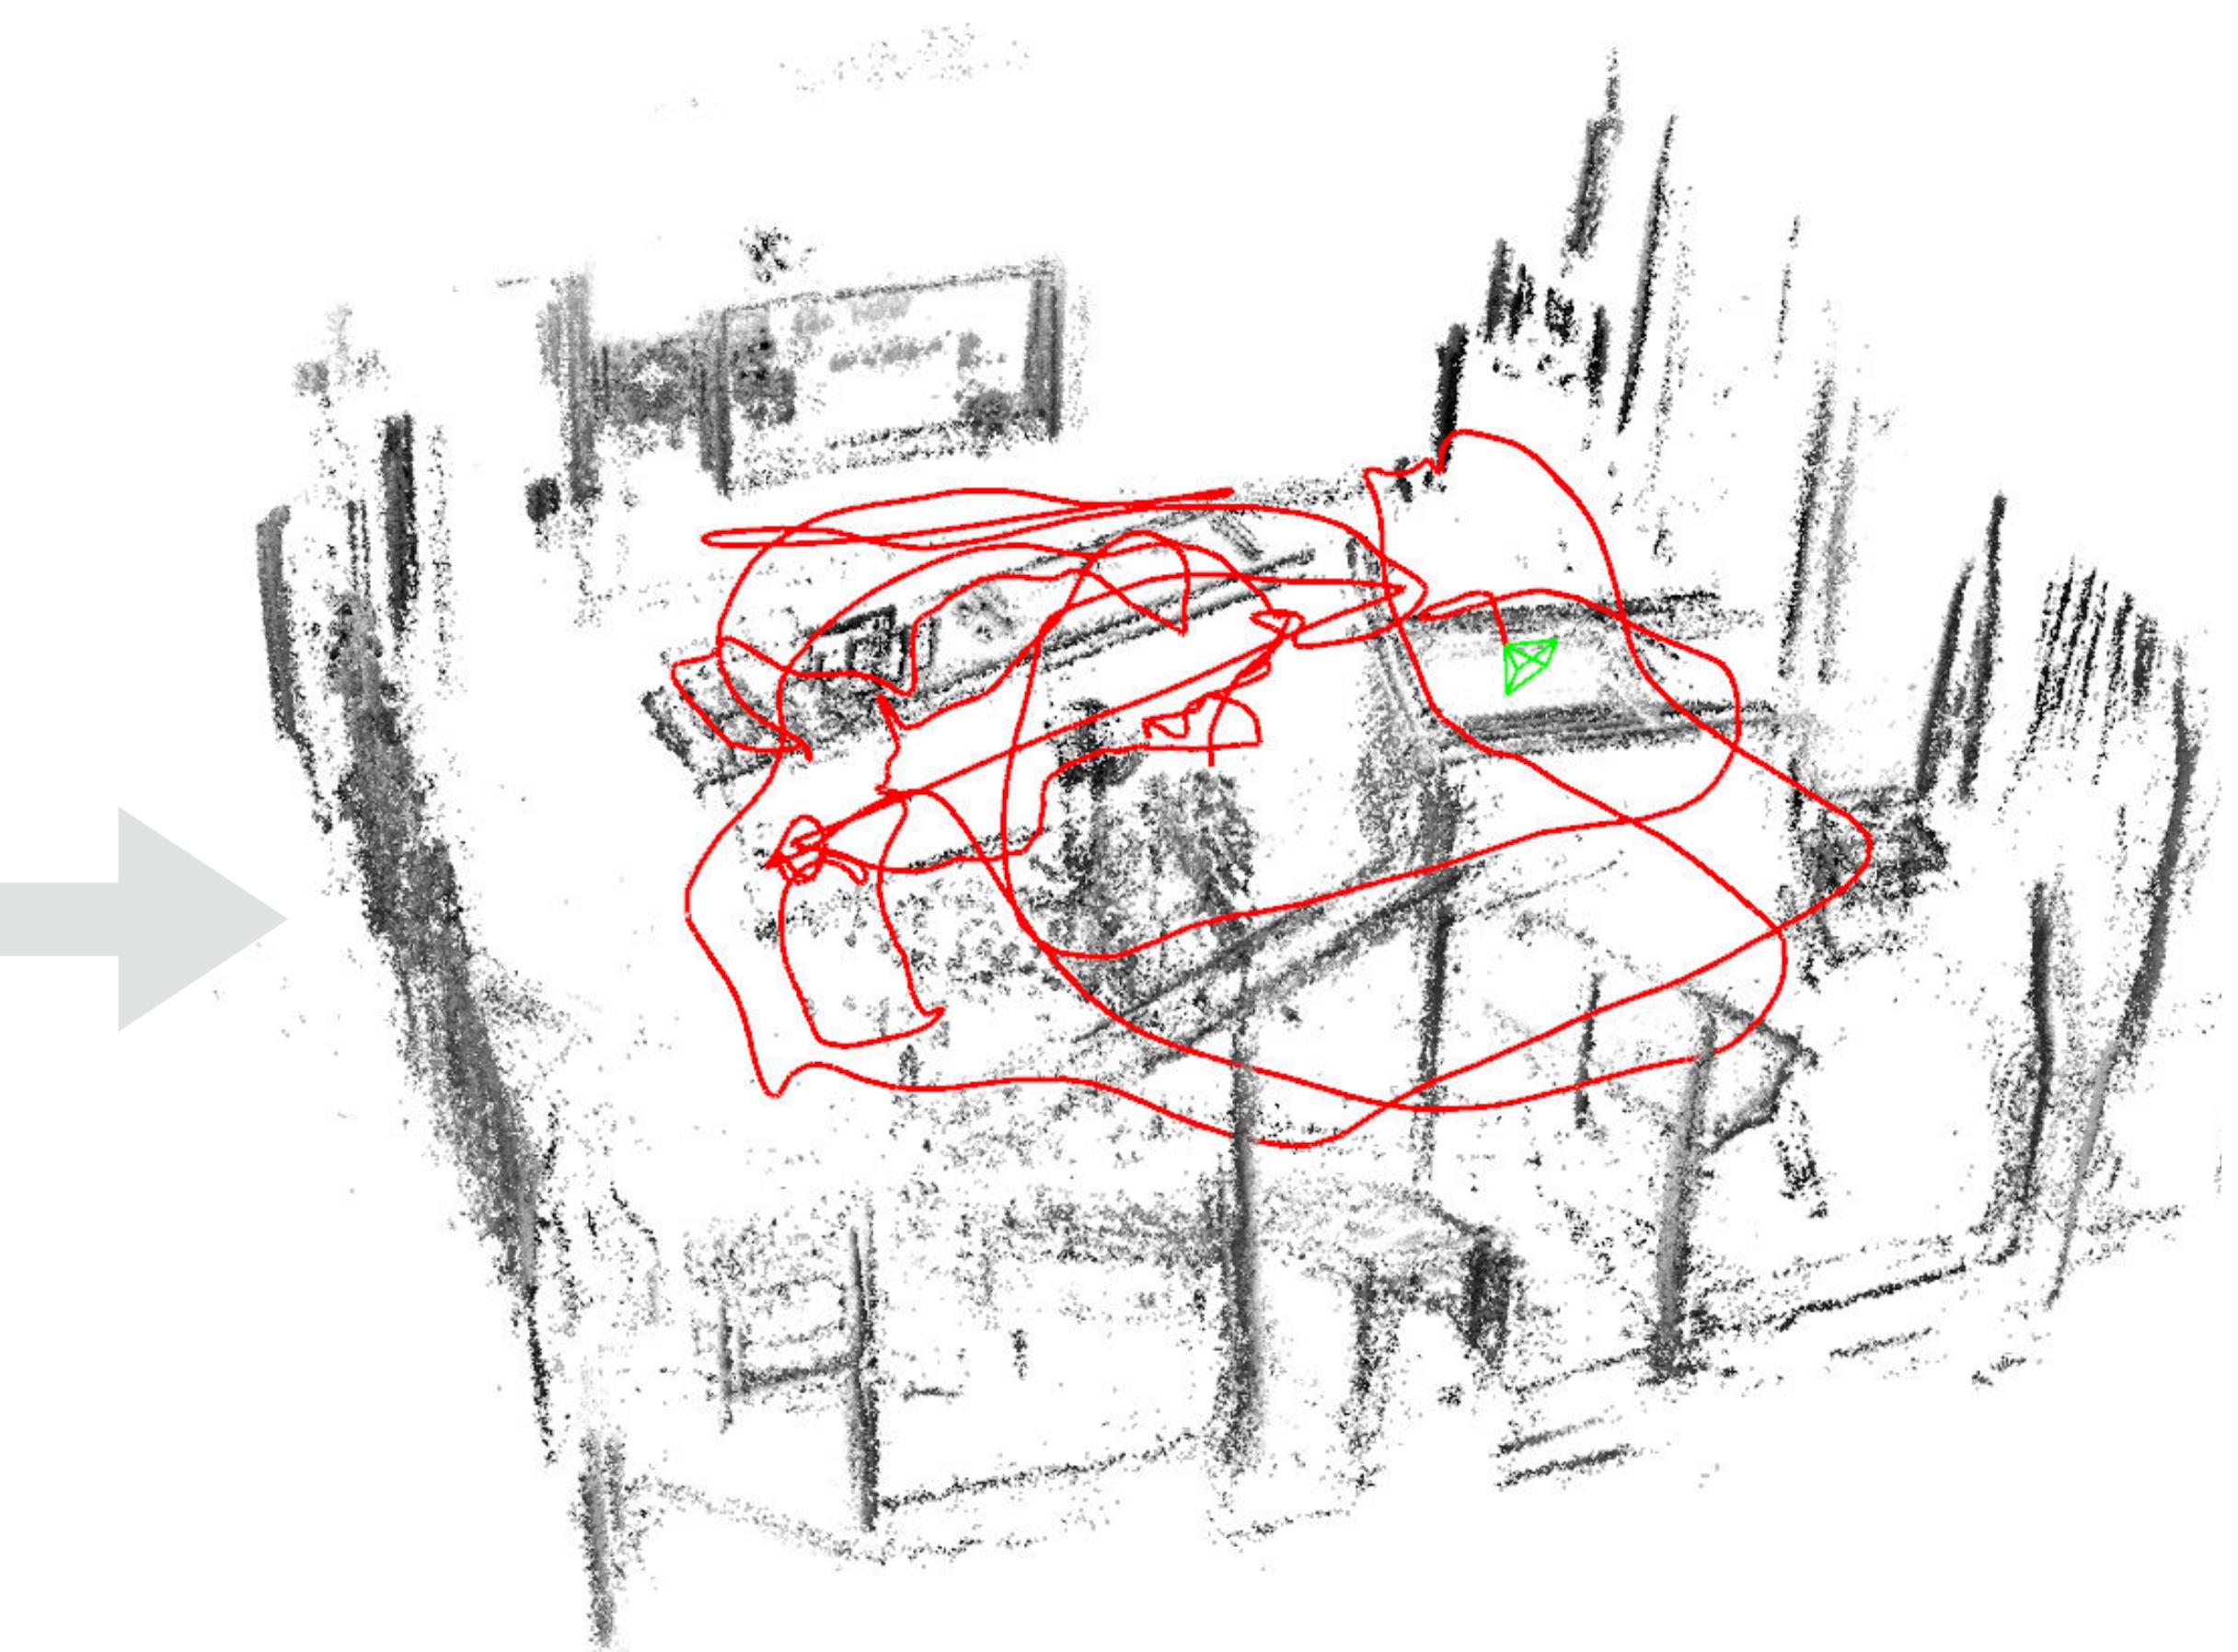
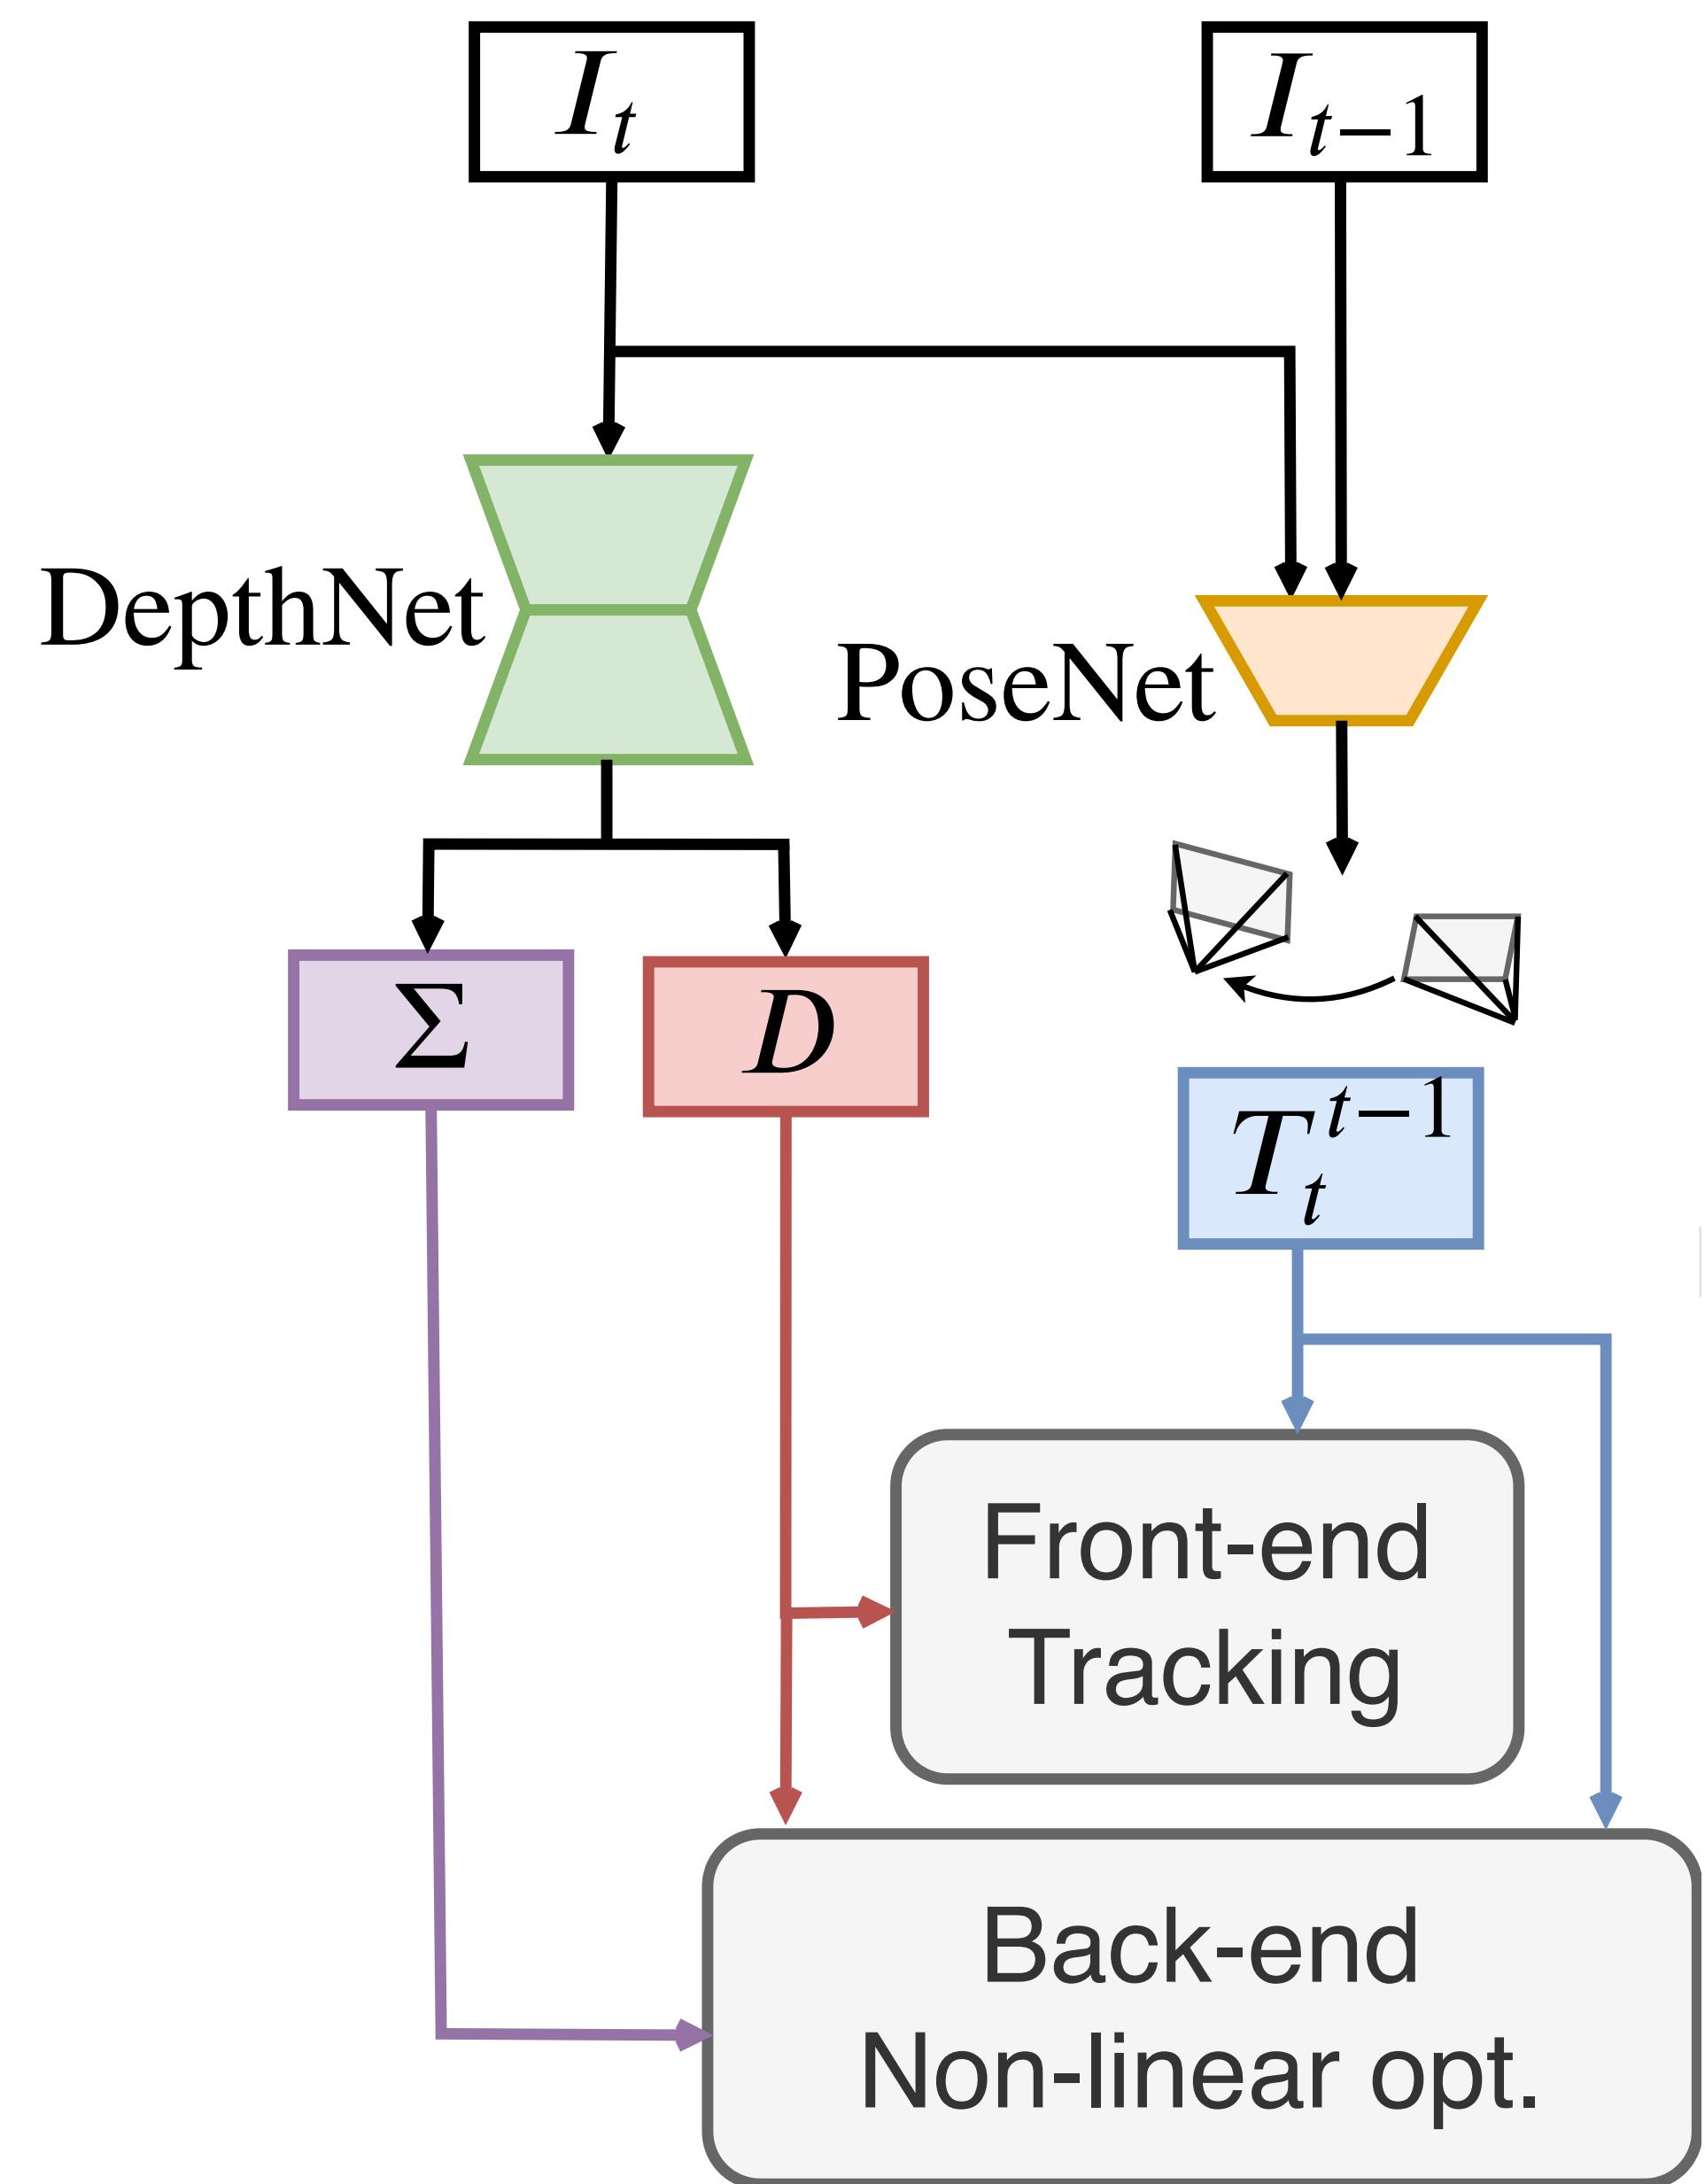


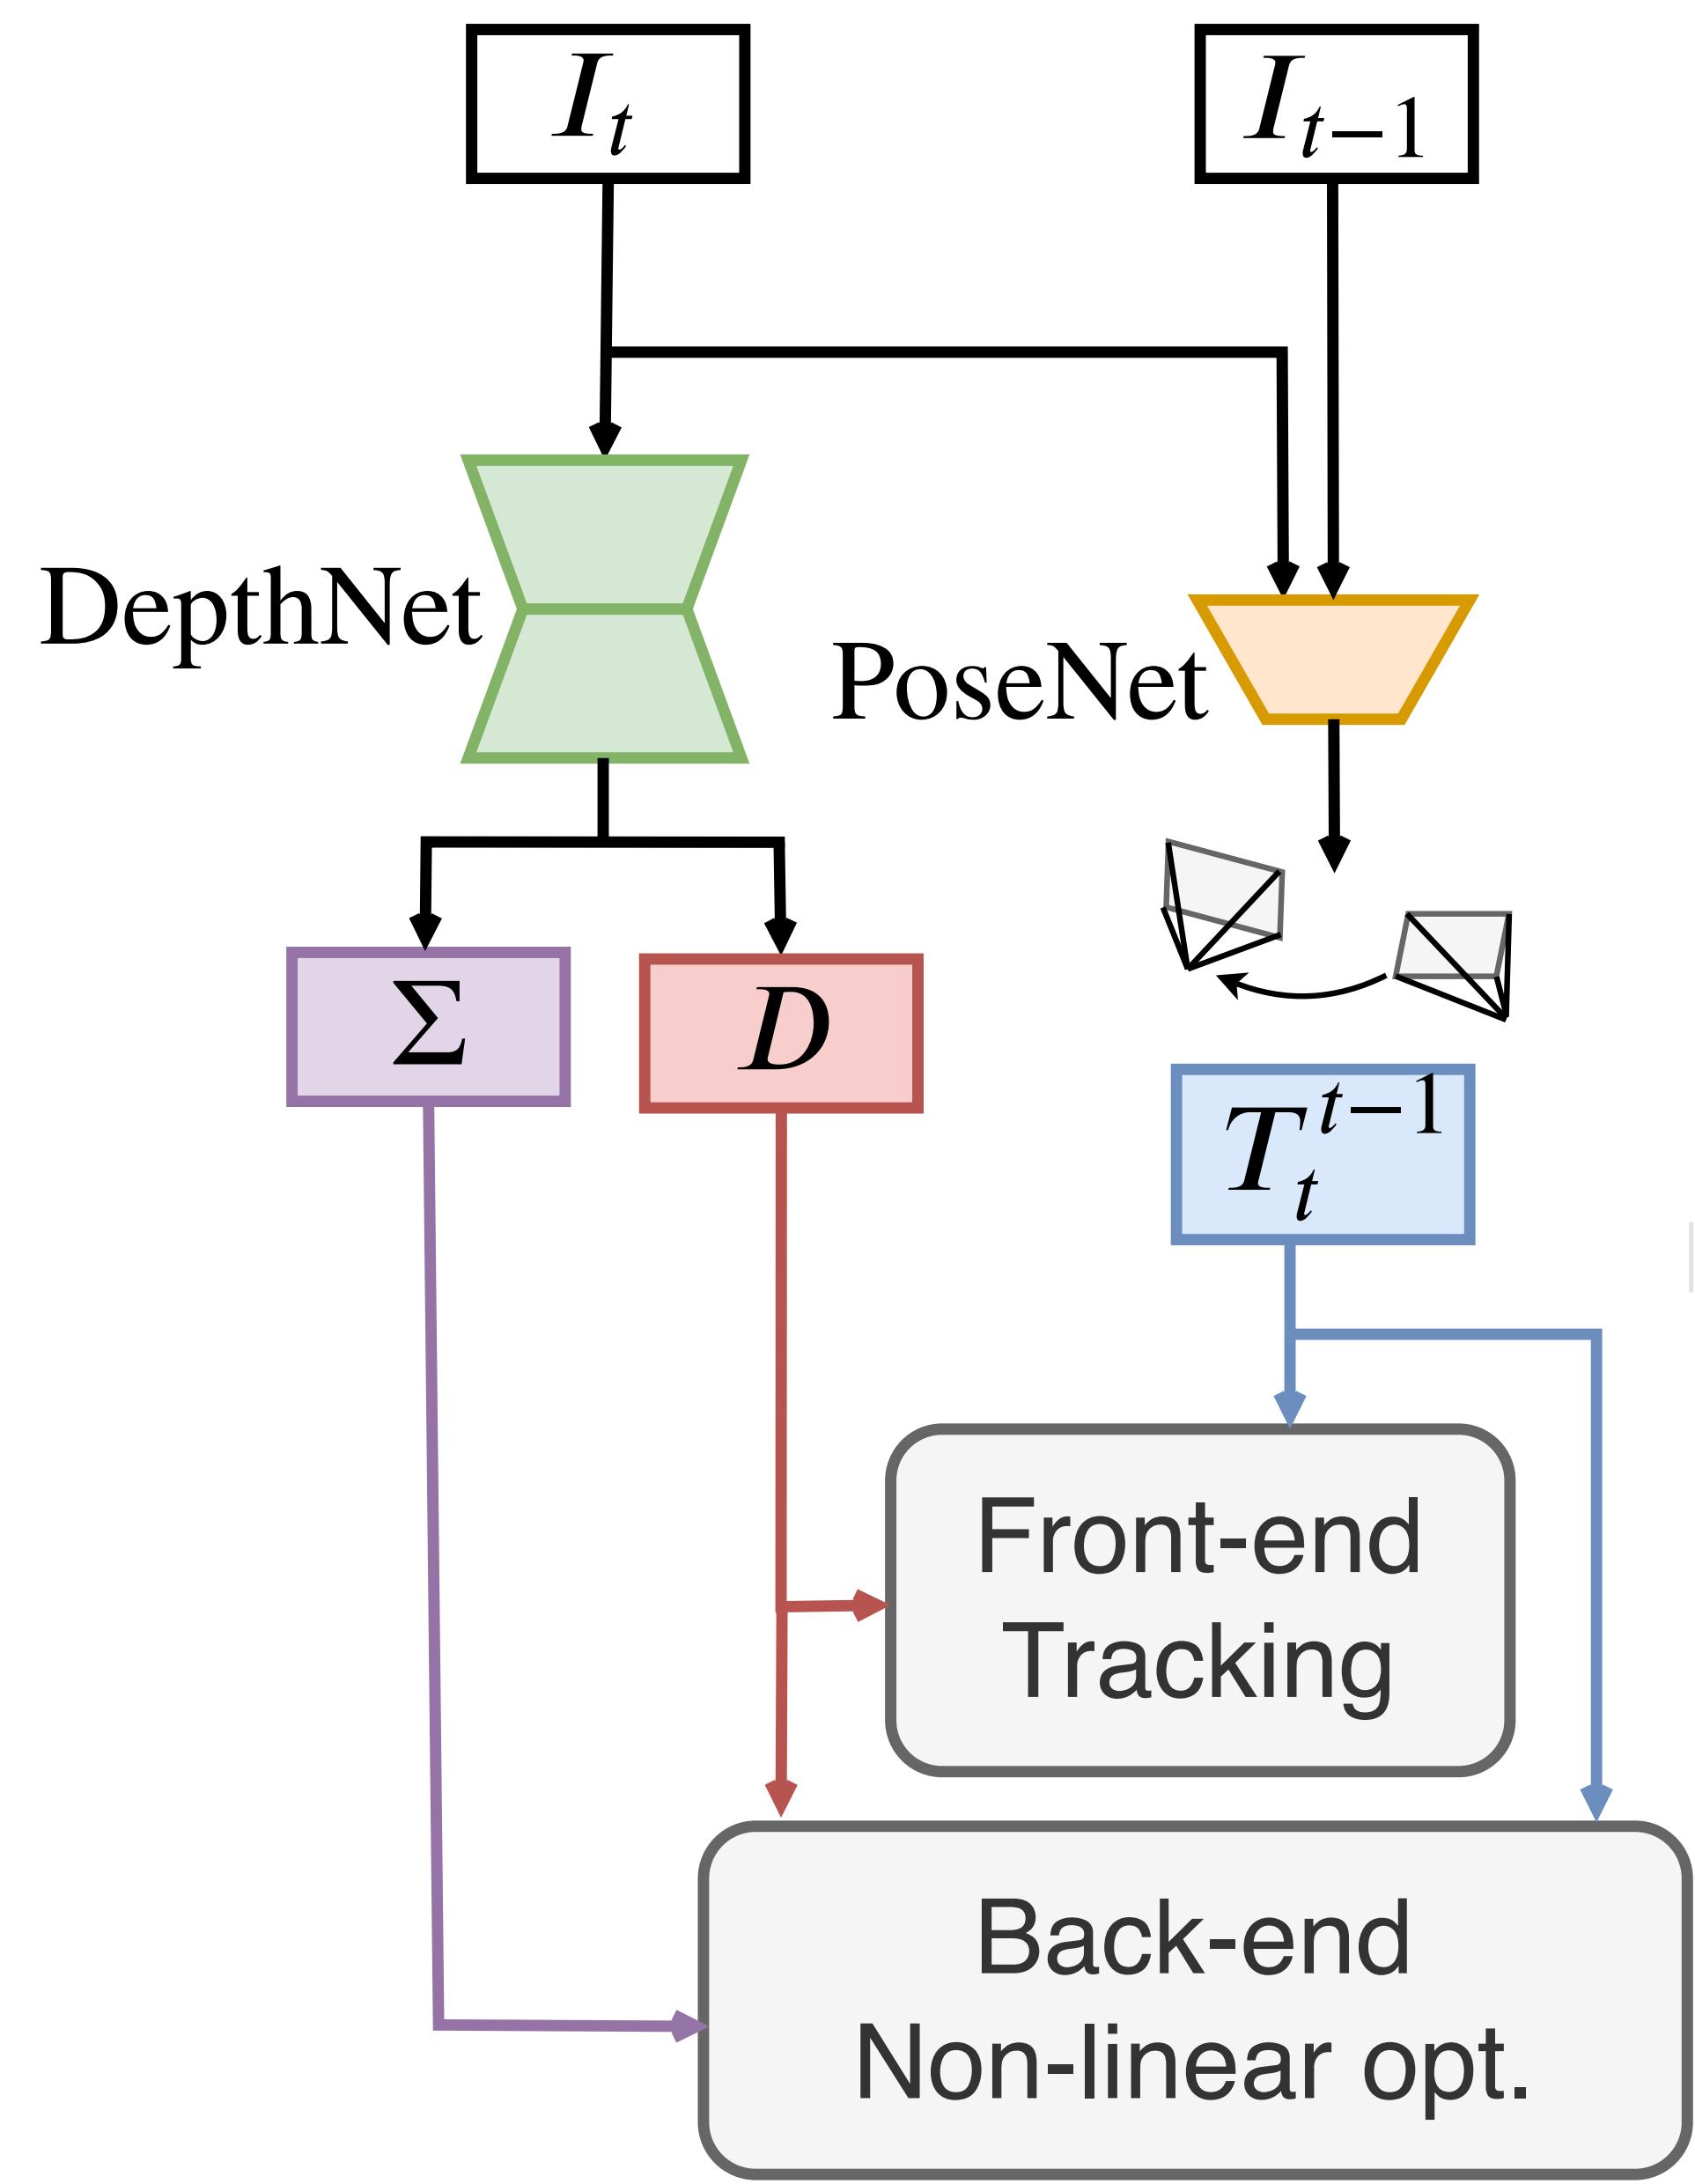




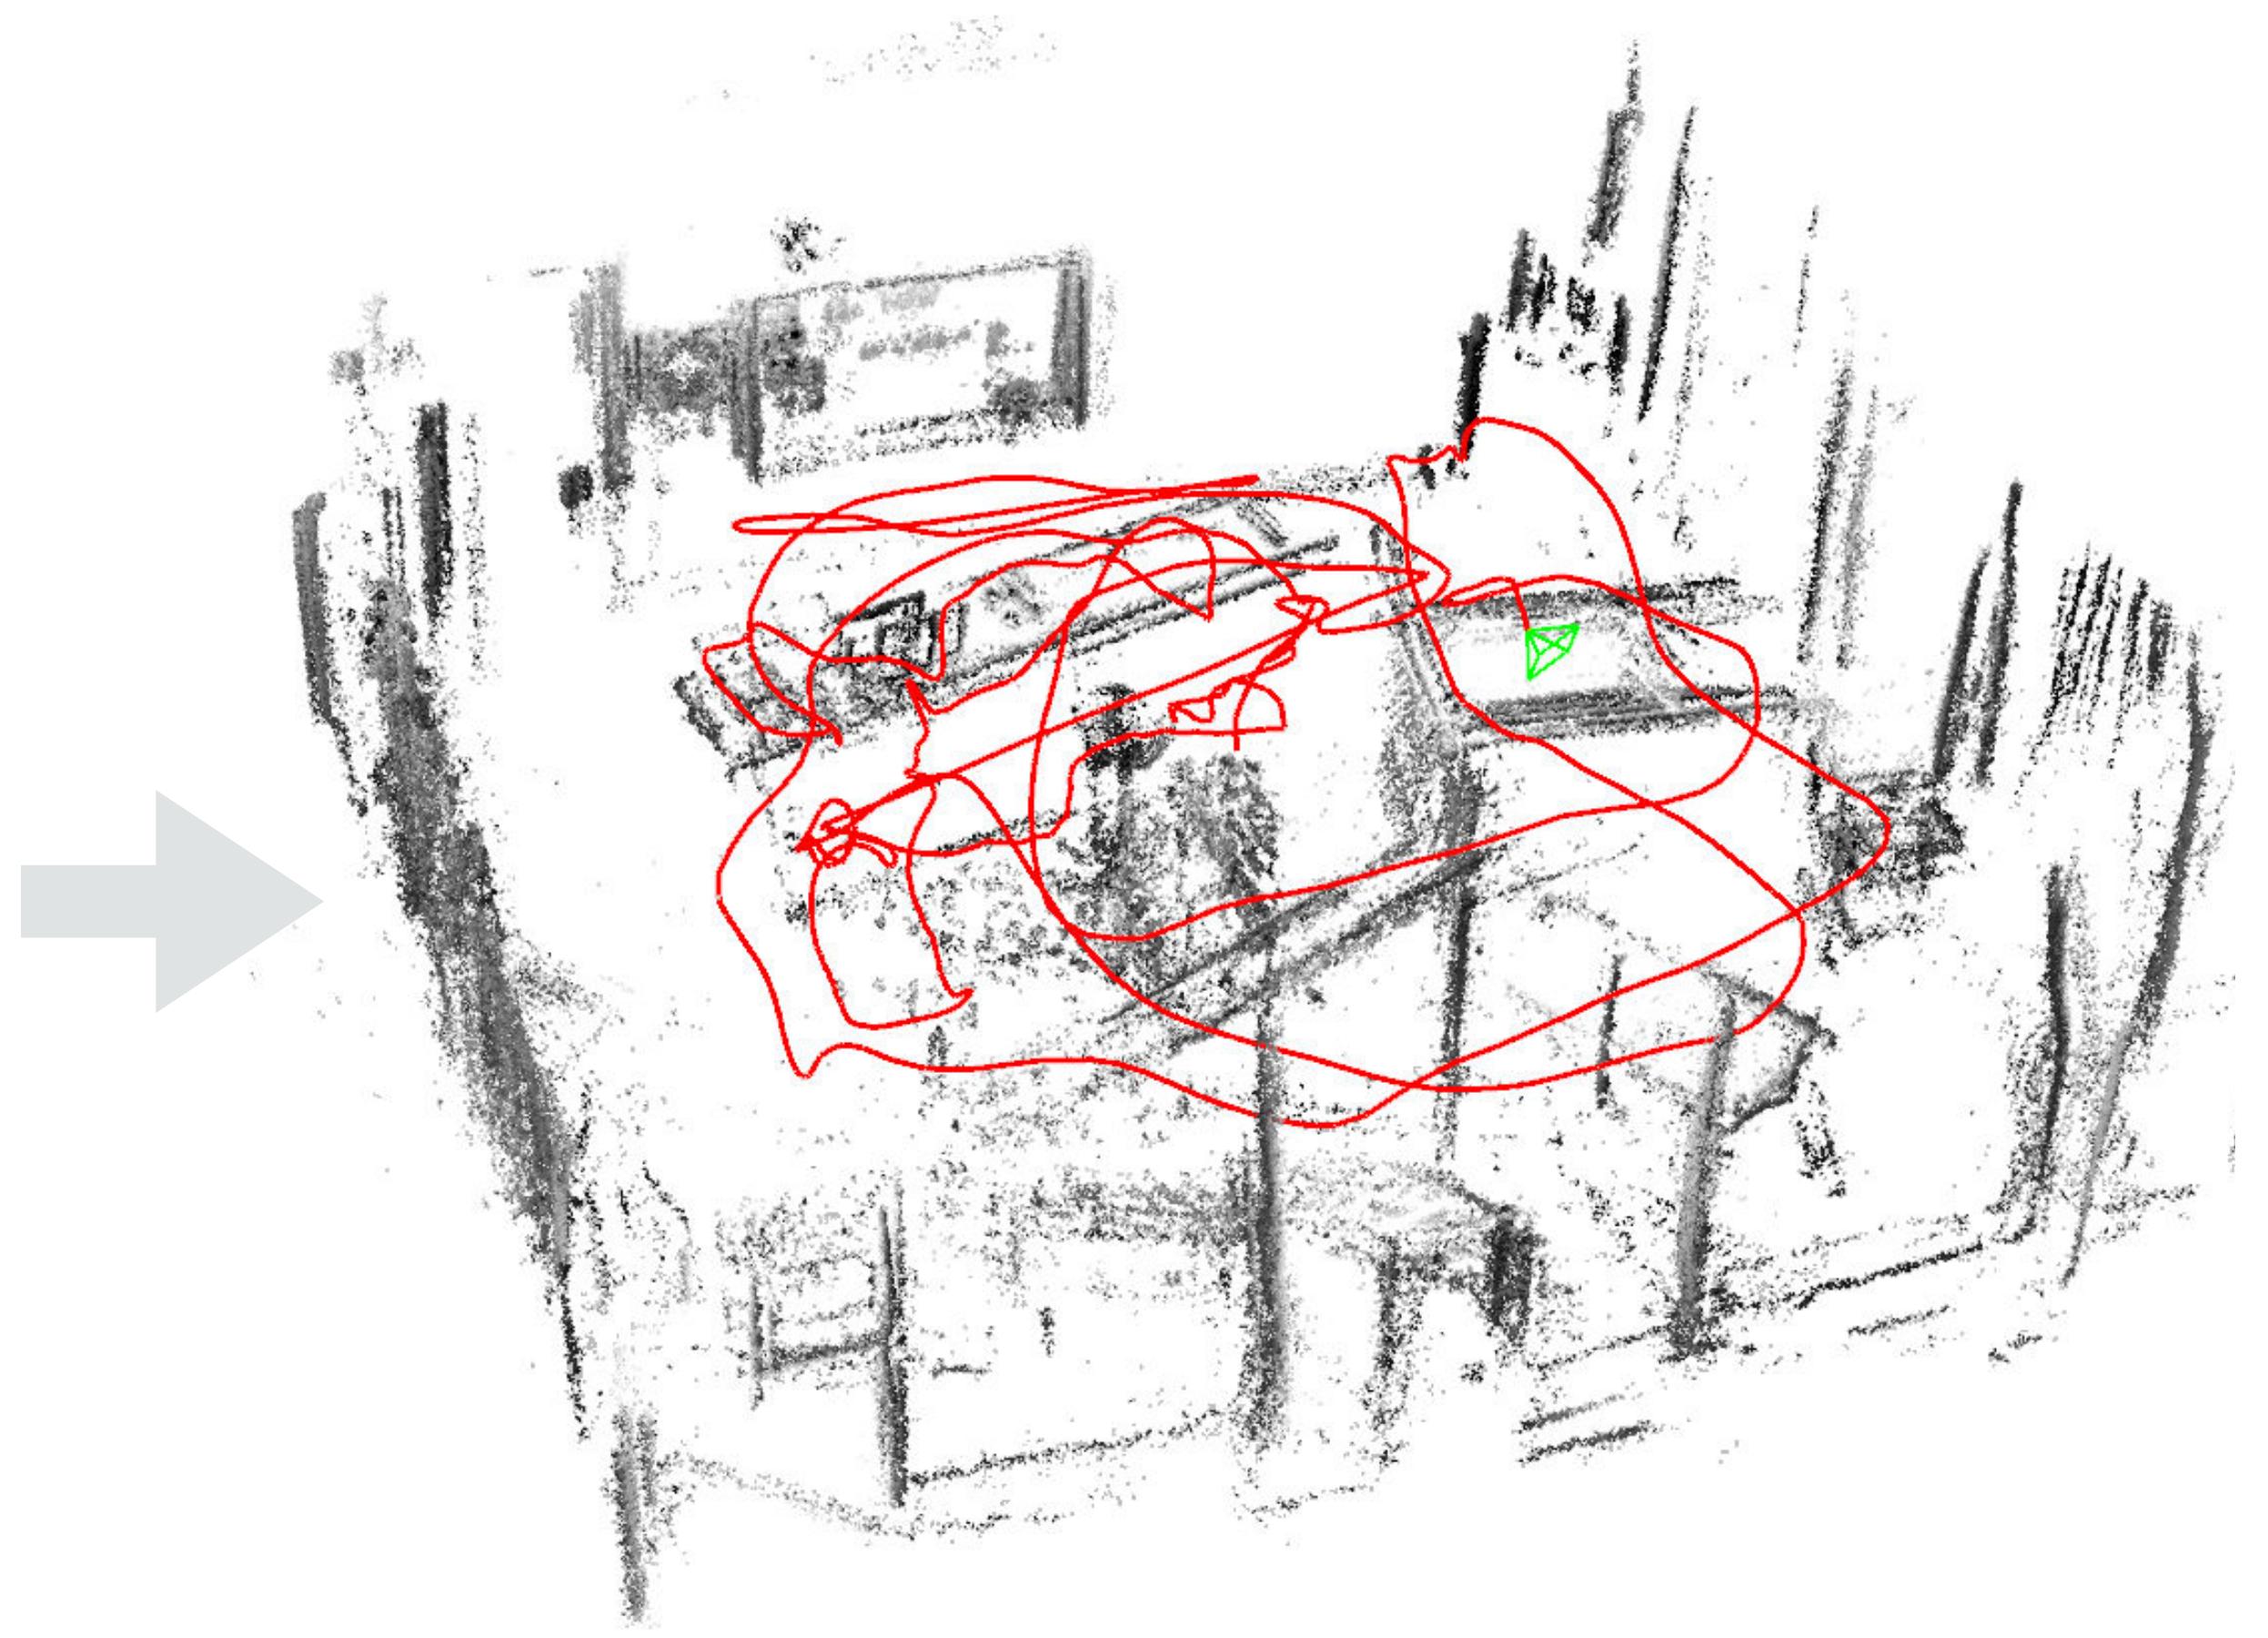






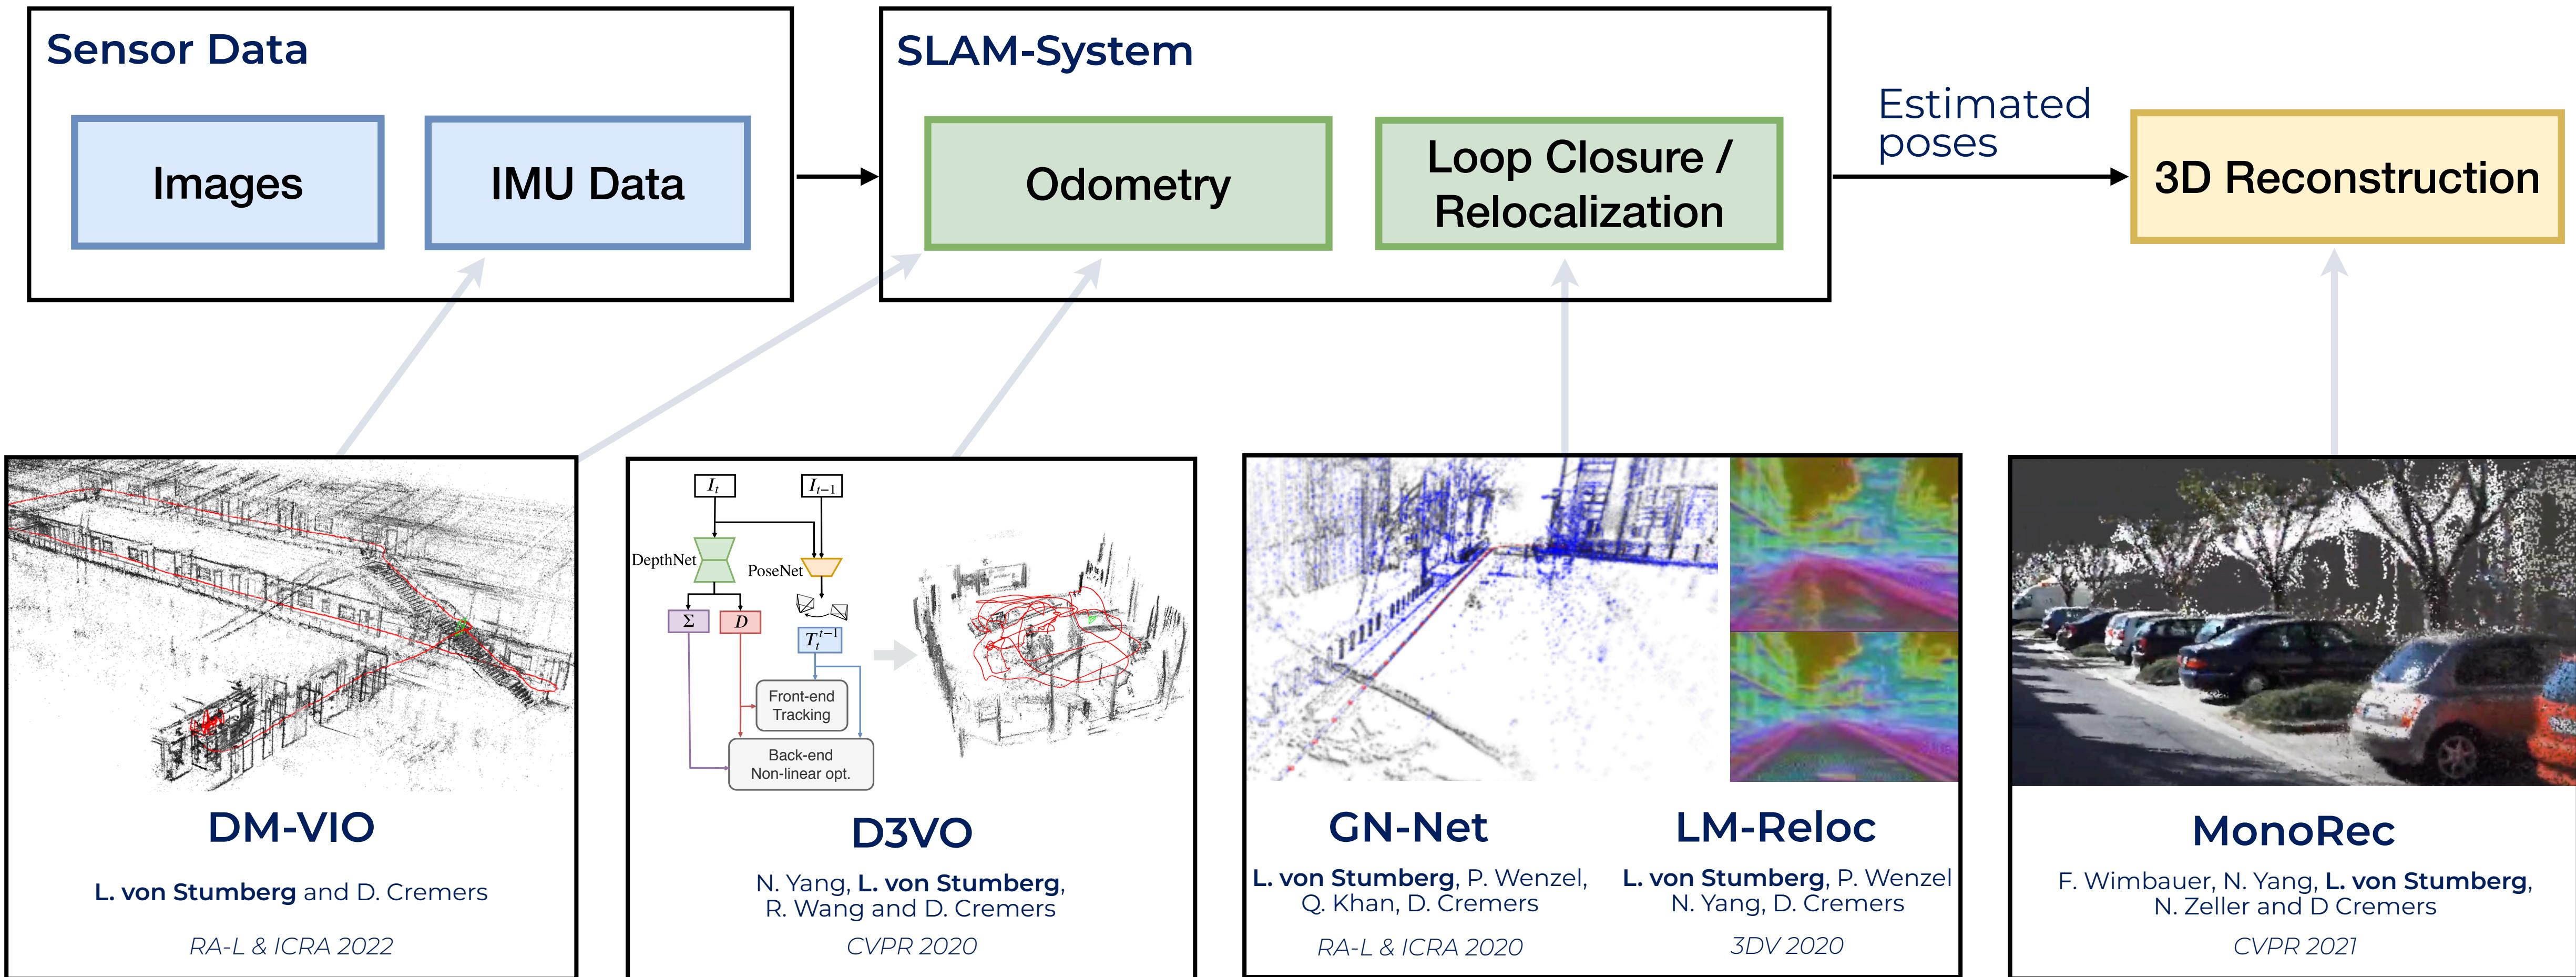


	M03	M05	V103	V202	V203	mean
M	DSO [15]	0.18	0.11	1.42	0.12	0.48
	ORB [24]	0.08	0.16	1.48	1.72	0.72
S+I	VINS [175]	0.23	0.19	0.11	0.10	-
	OKVIS [40]	0.23	0.36	0.13	0.17	-
	Basalt [34]	0.06	0.12	0.10	0.05	-
D3VO		0.08	0.09	0.11	0.05	-
						0.08



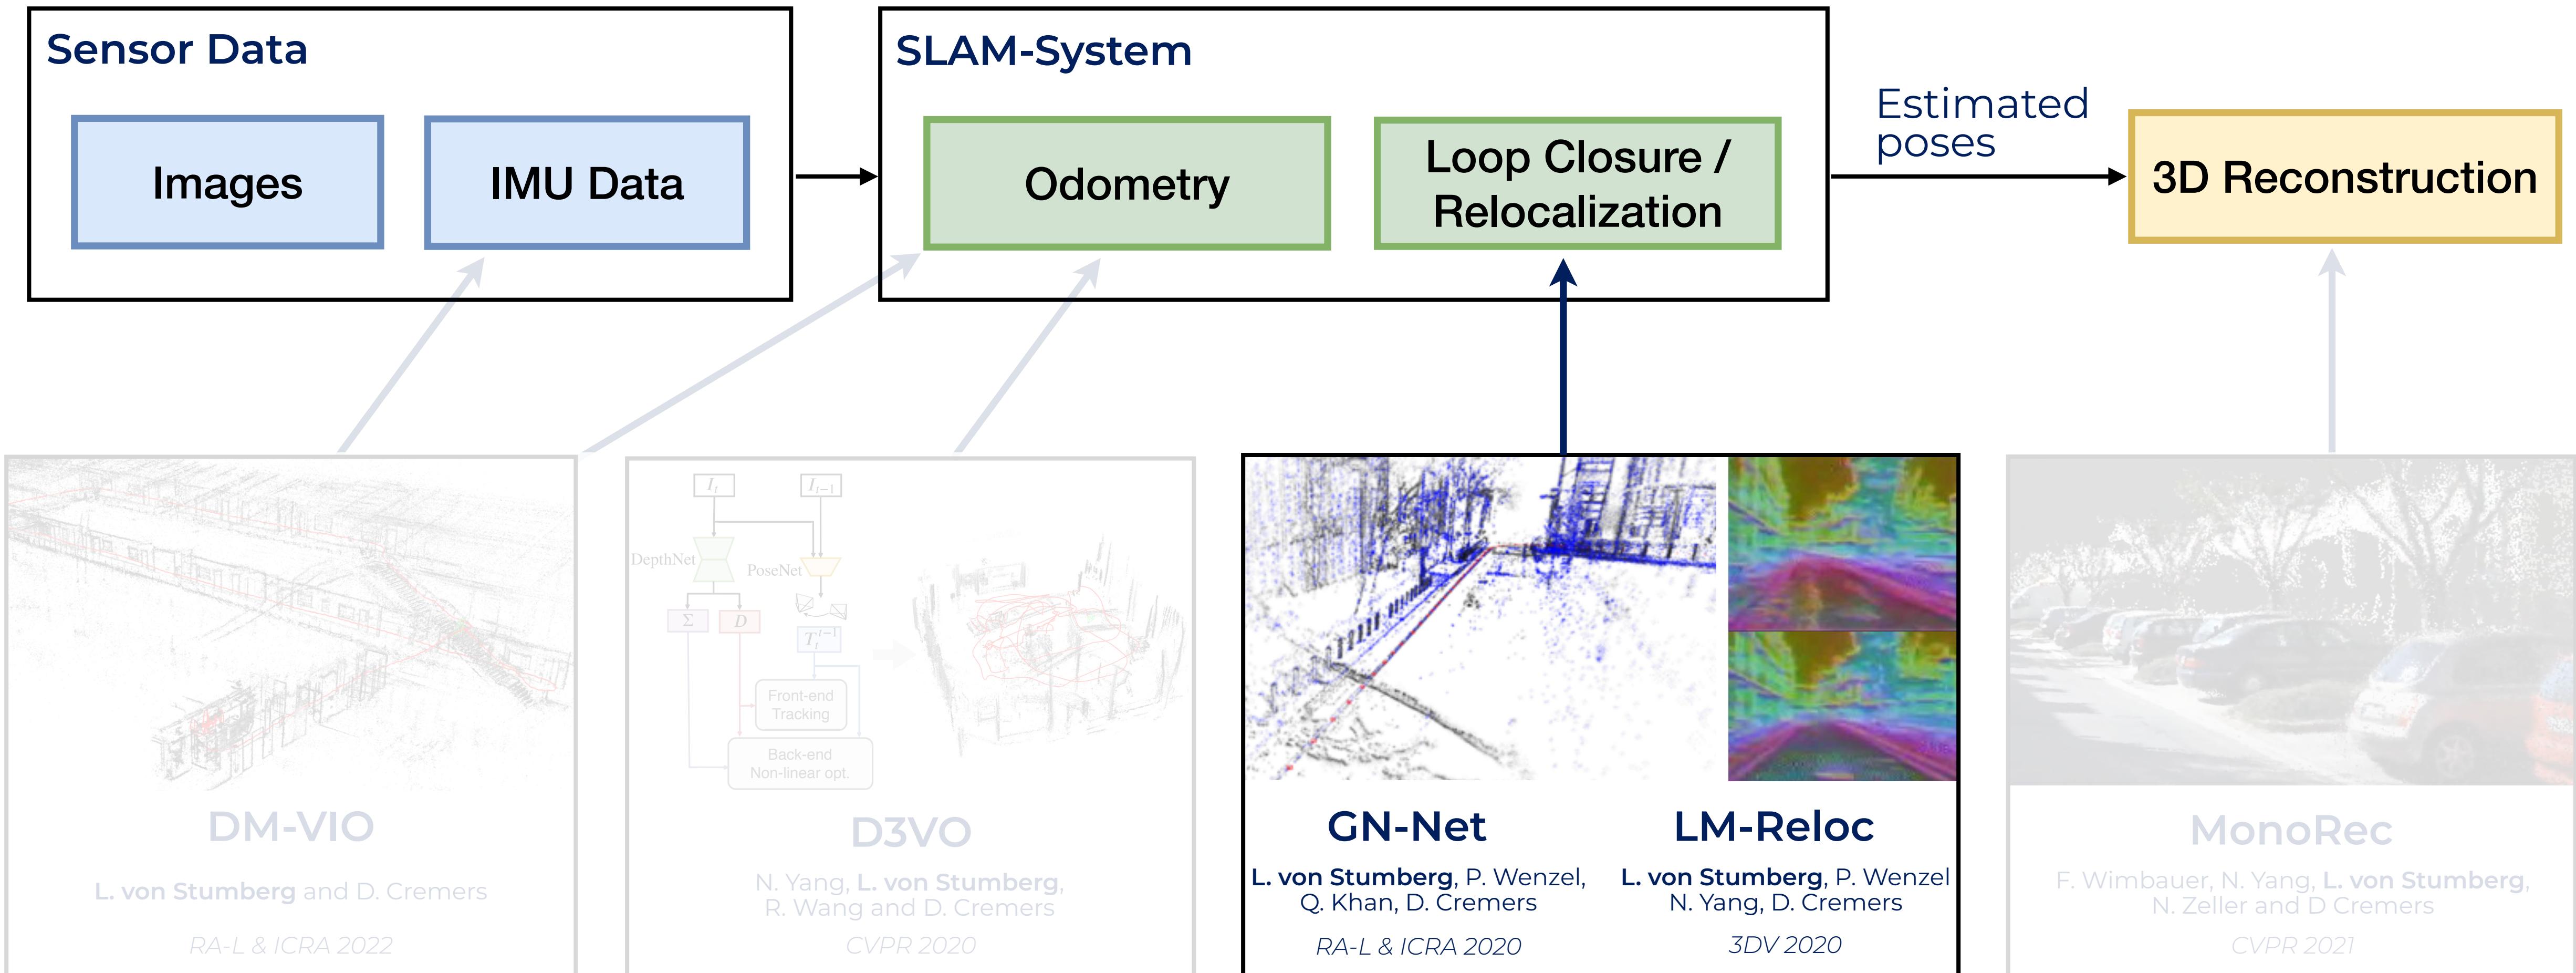
Visual SLAM

From Optimization to Learning



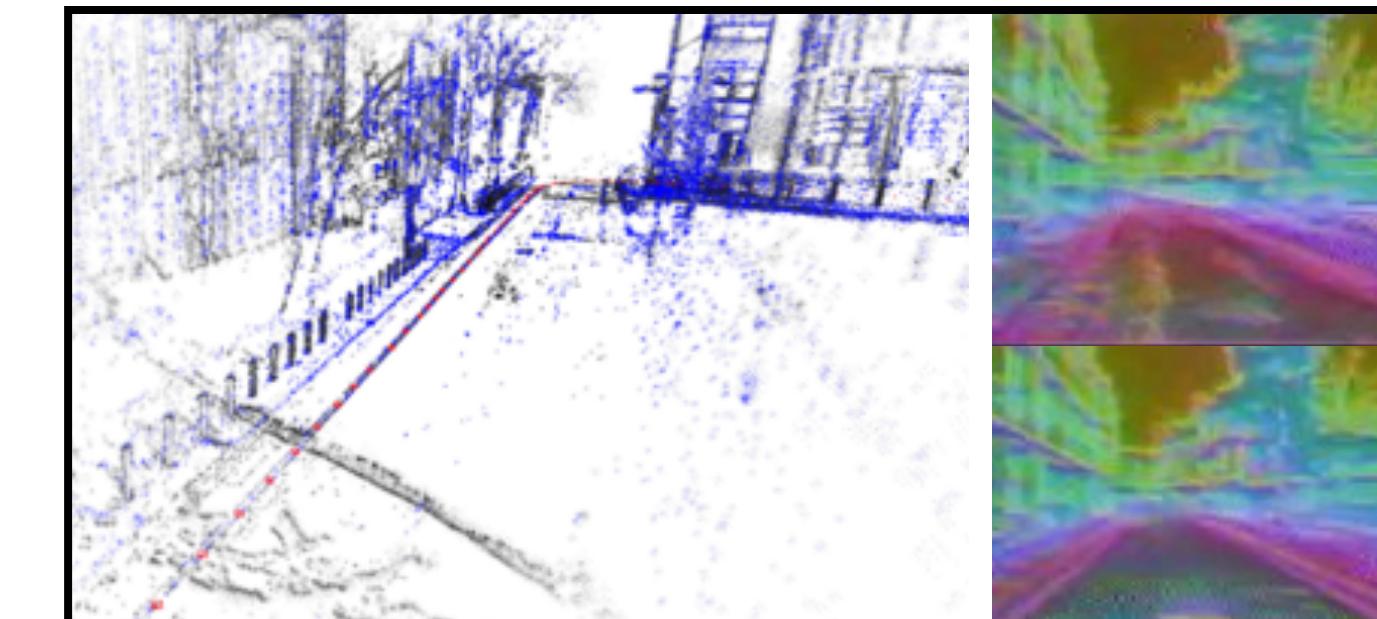
Visual SLAM

From Optimization to Learning



Visual SLAM

From Optimization to Learning



GN-Net

L. von Stumberg, P. Wenzel,
Q. Khan, D. Cremers

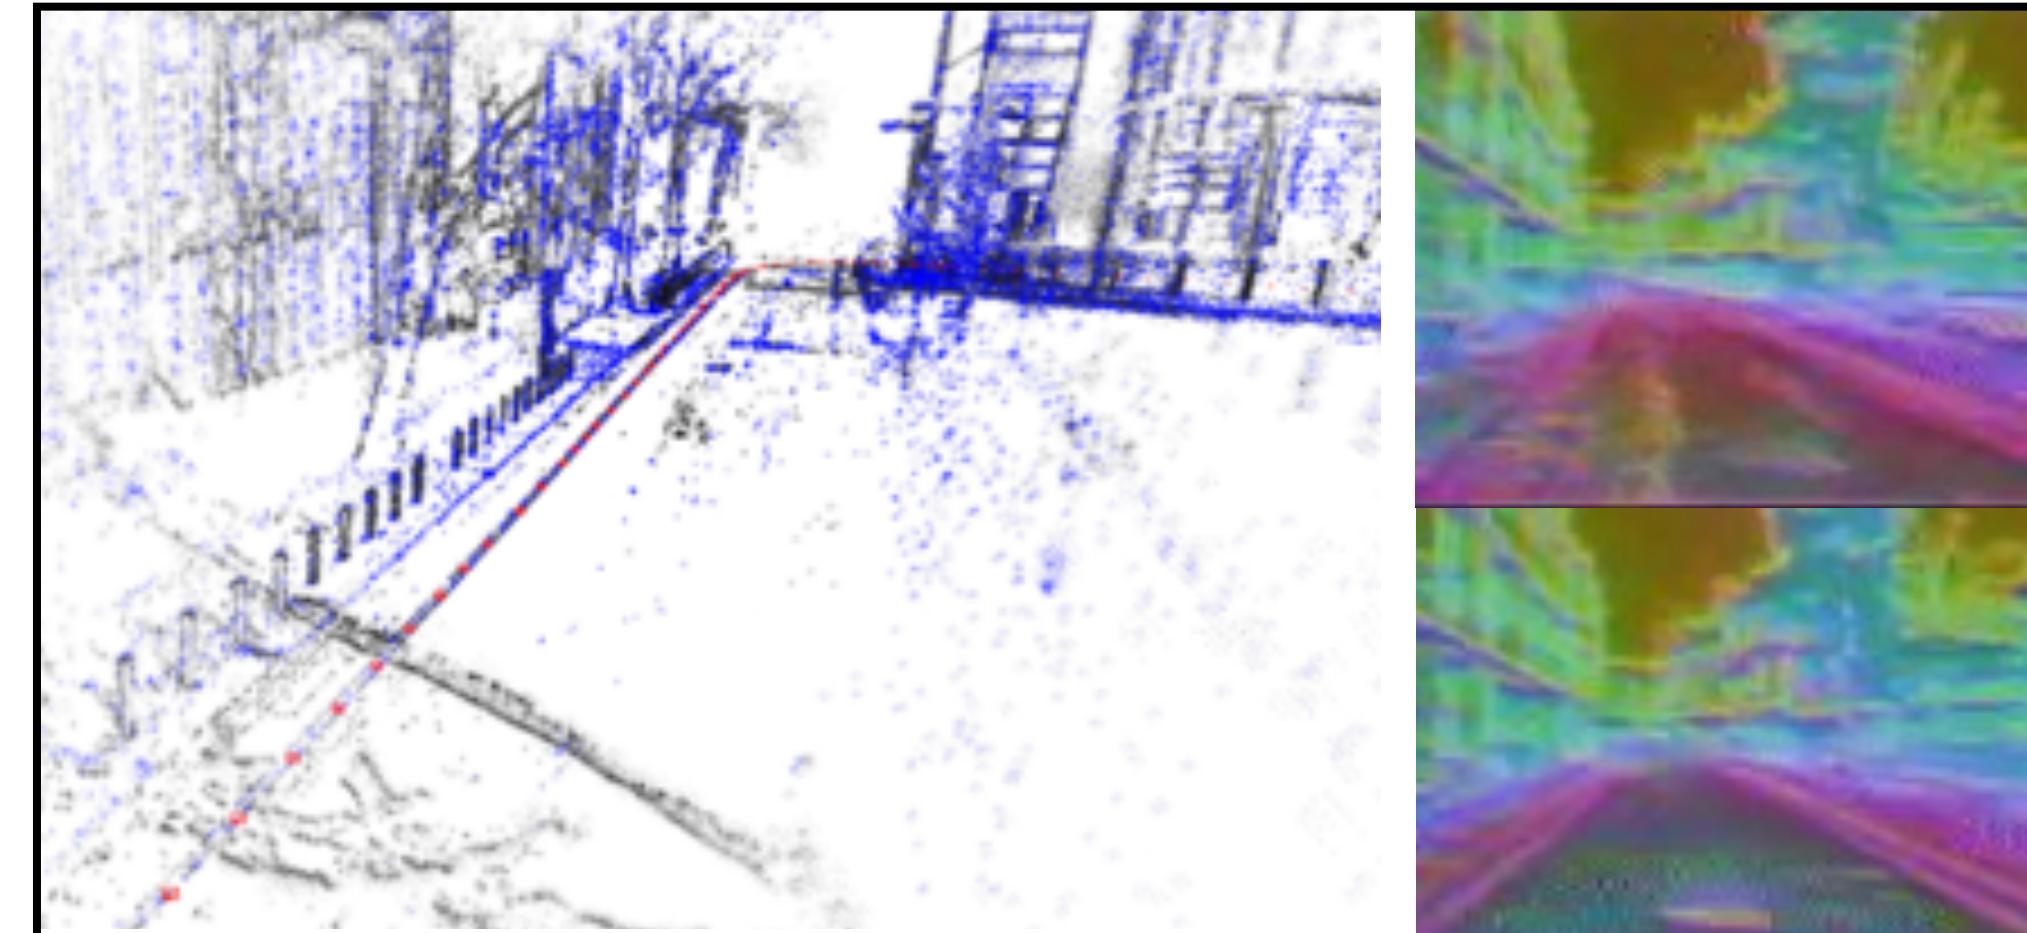
RA-L & ICRA 2020

LM-Reloc

L. von Stumberg, P. Wenzel
N. Yang, D. Cremers

3DV 2020

Visual SLAM From Optimization to Learning



GN-Net

L. von Stumberg*, P. Wenzel*,
Q. Khan, D. Cremers

RA-L & ICRA 2020

LM-Reloc

L. von Stumberg*, P. Wenzel*
N. Yang, D. Cremers

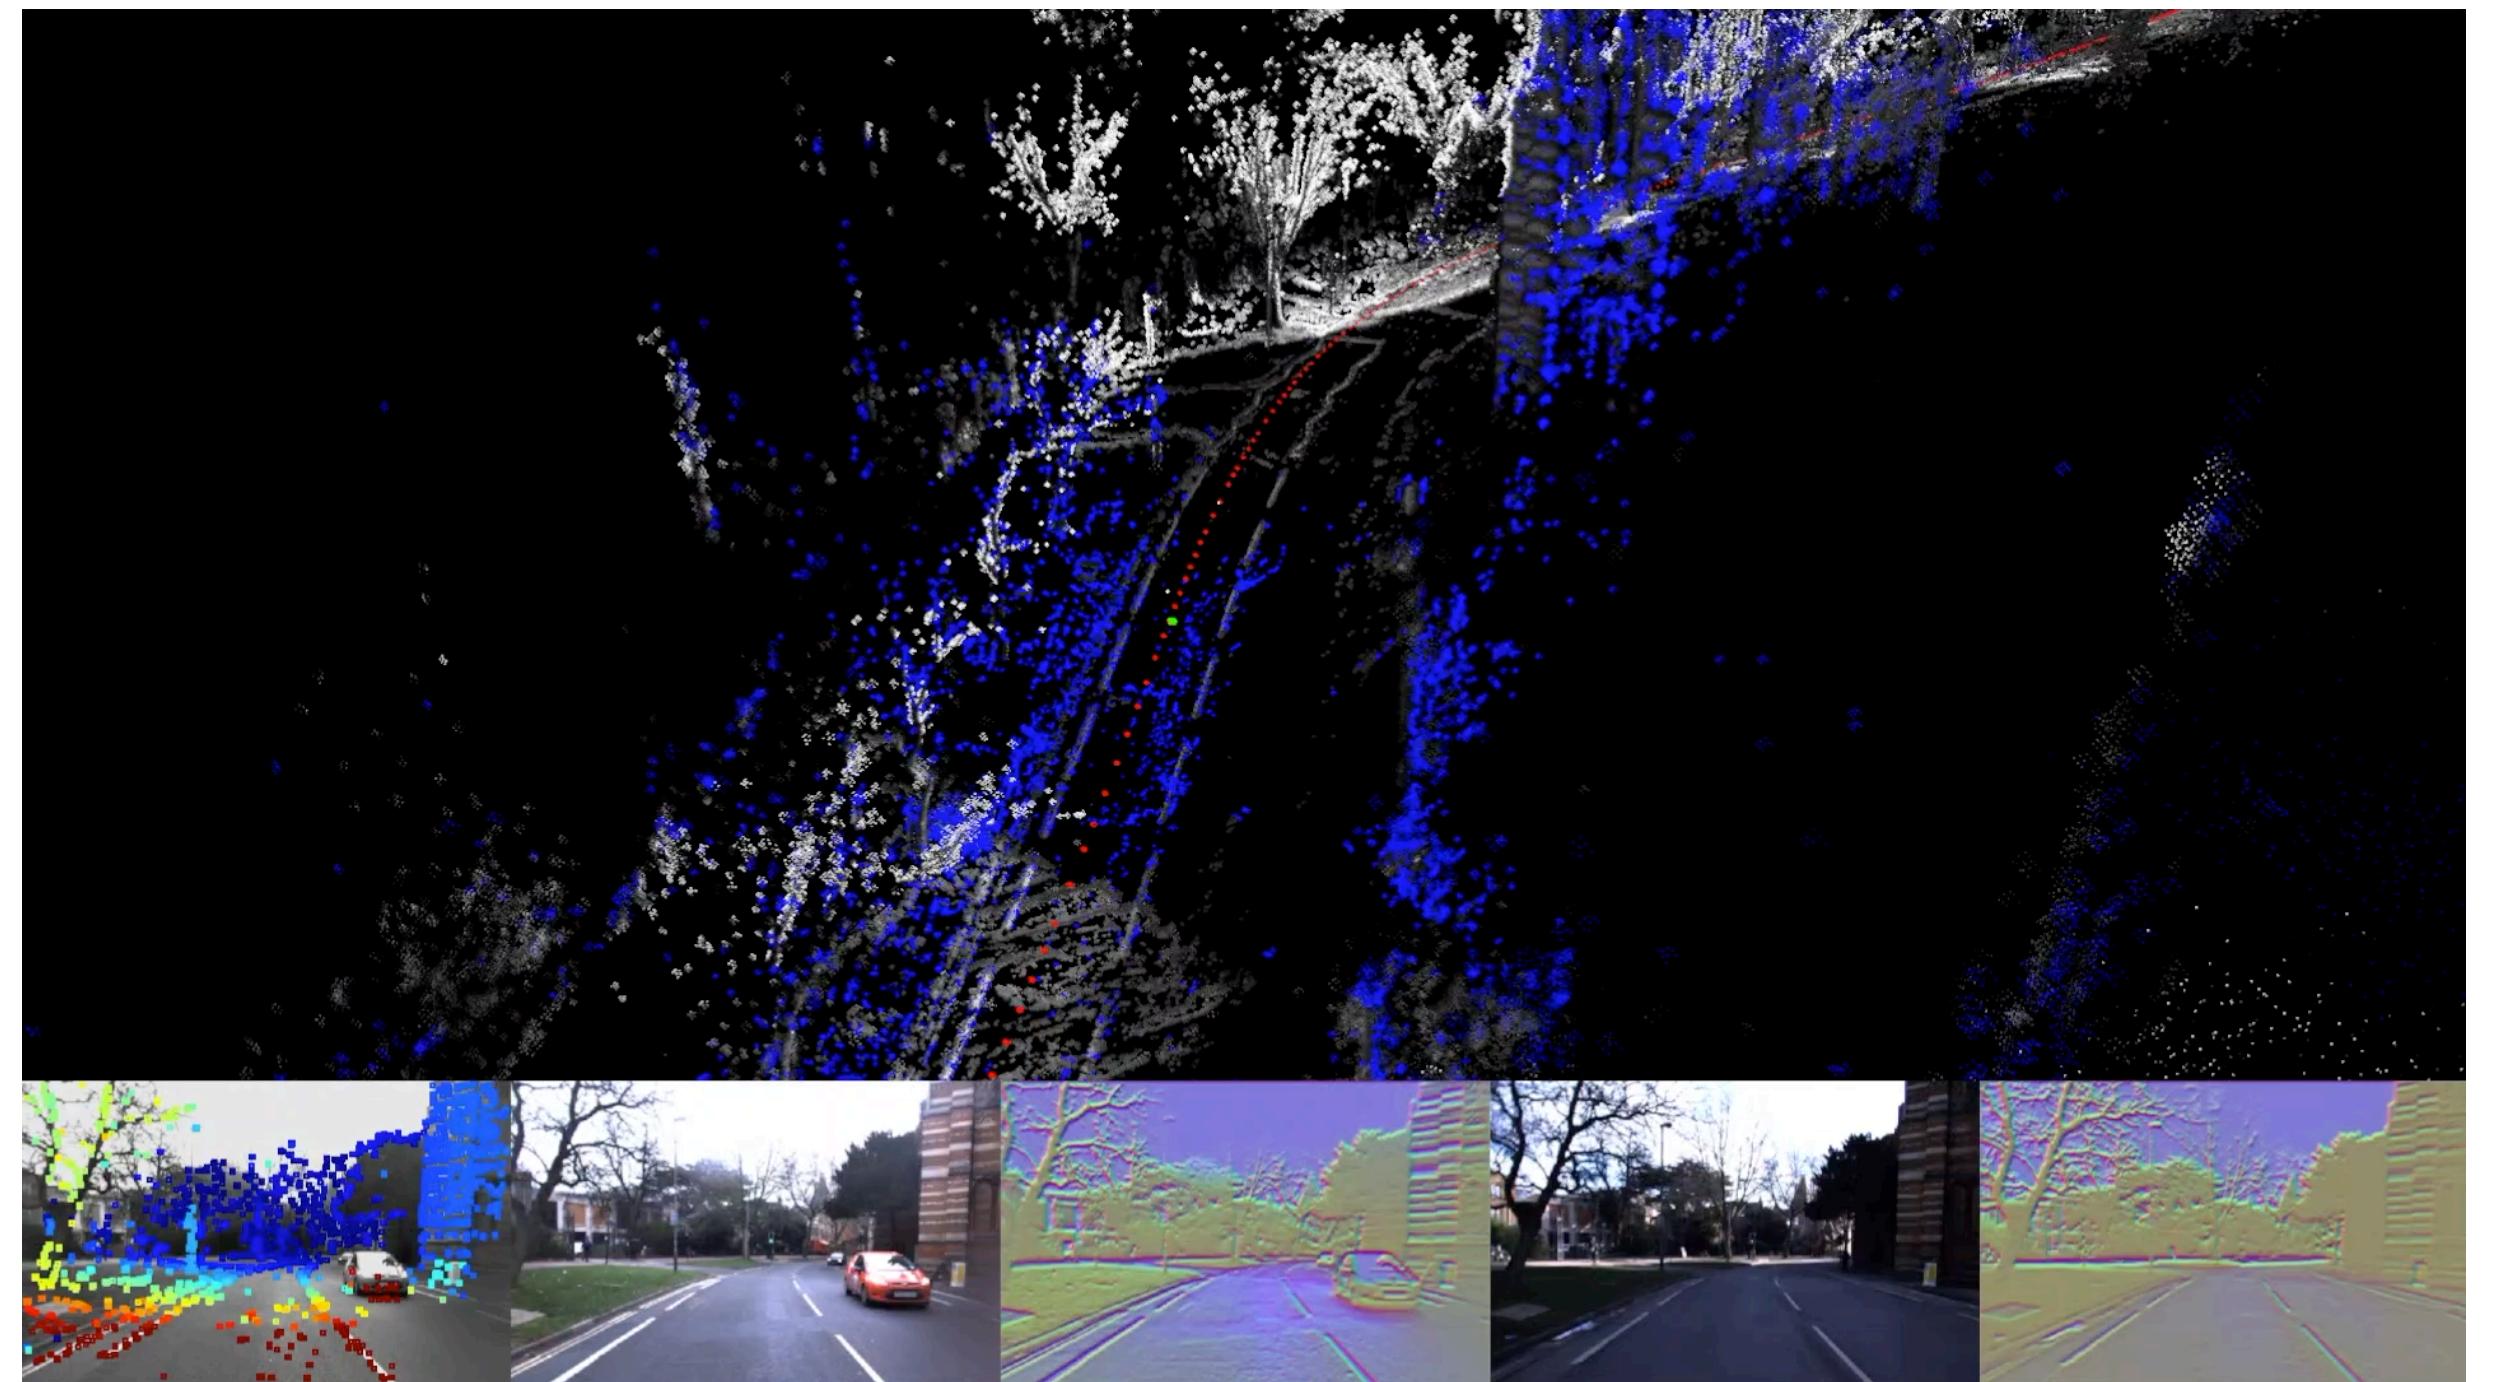
3DV 2020

*equal contribution

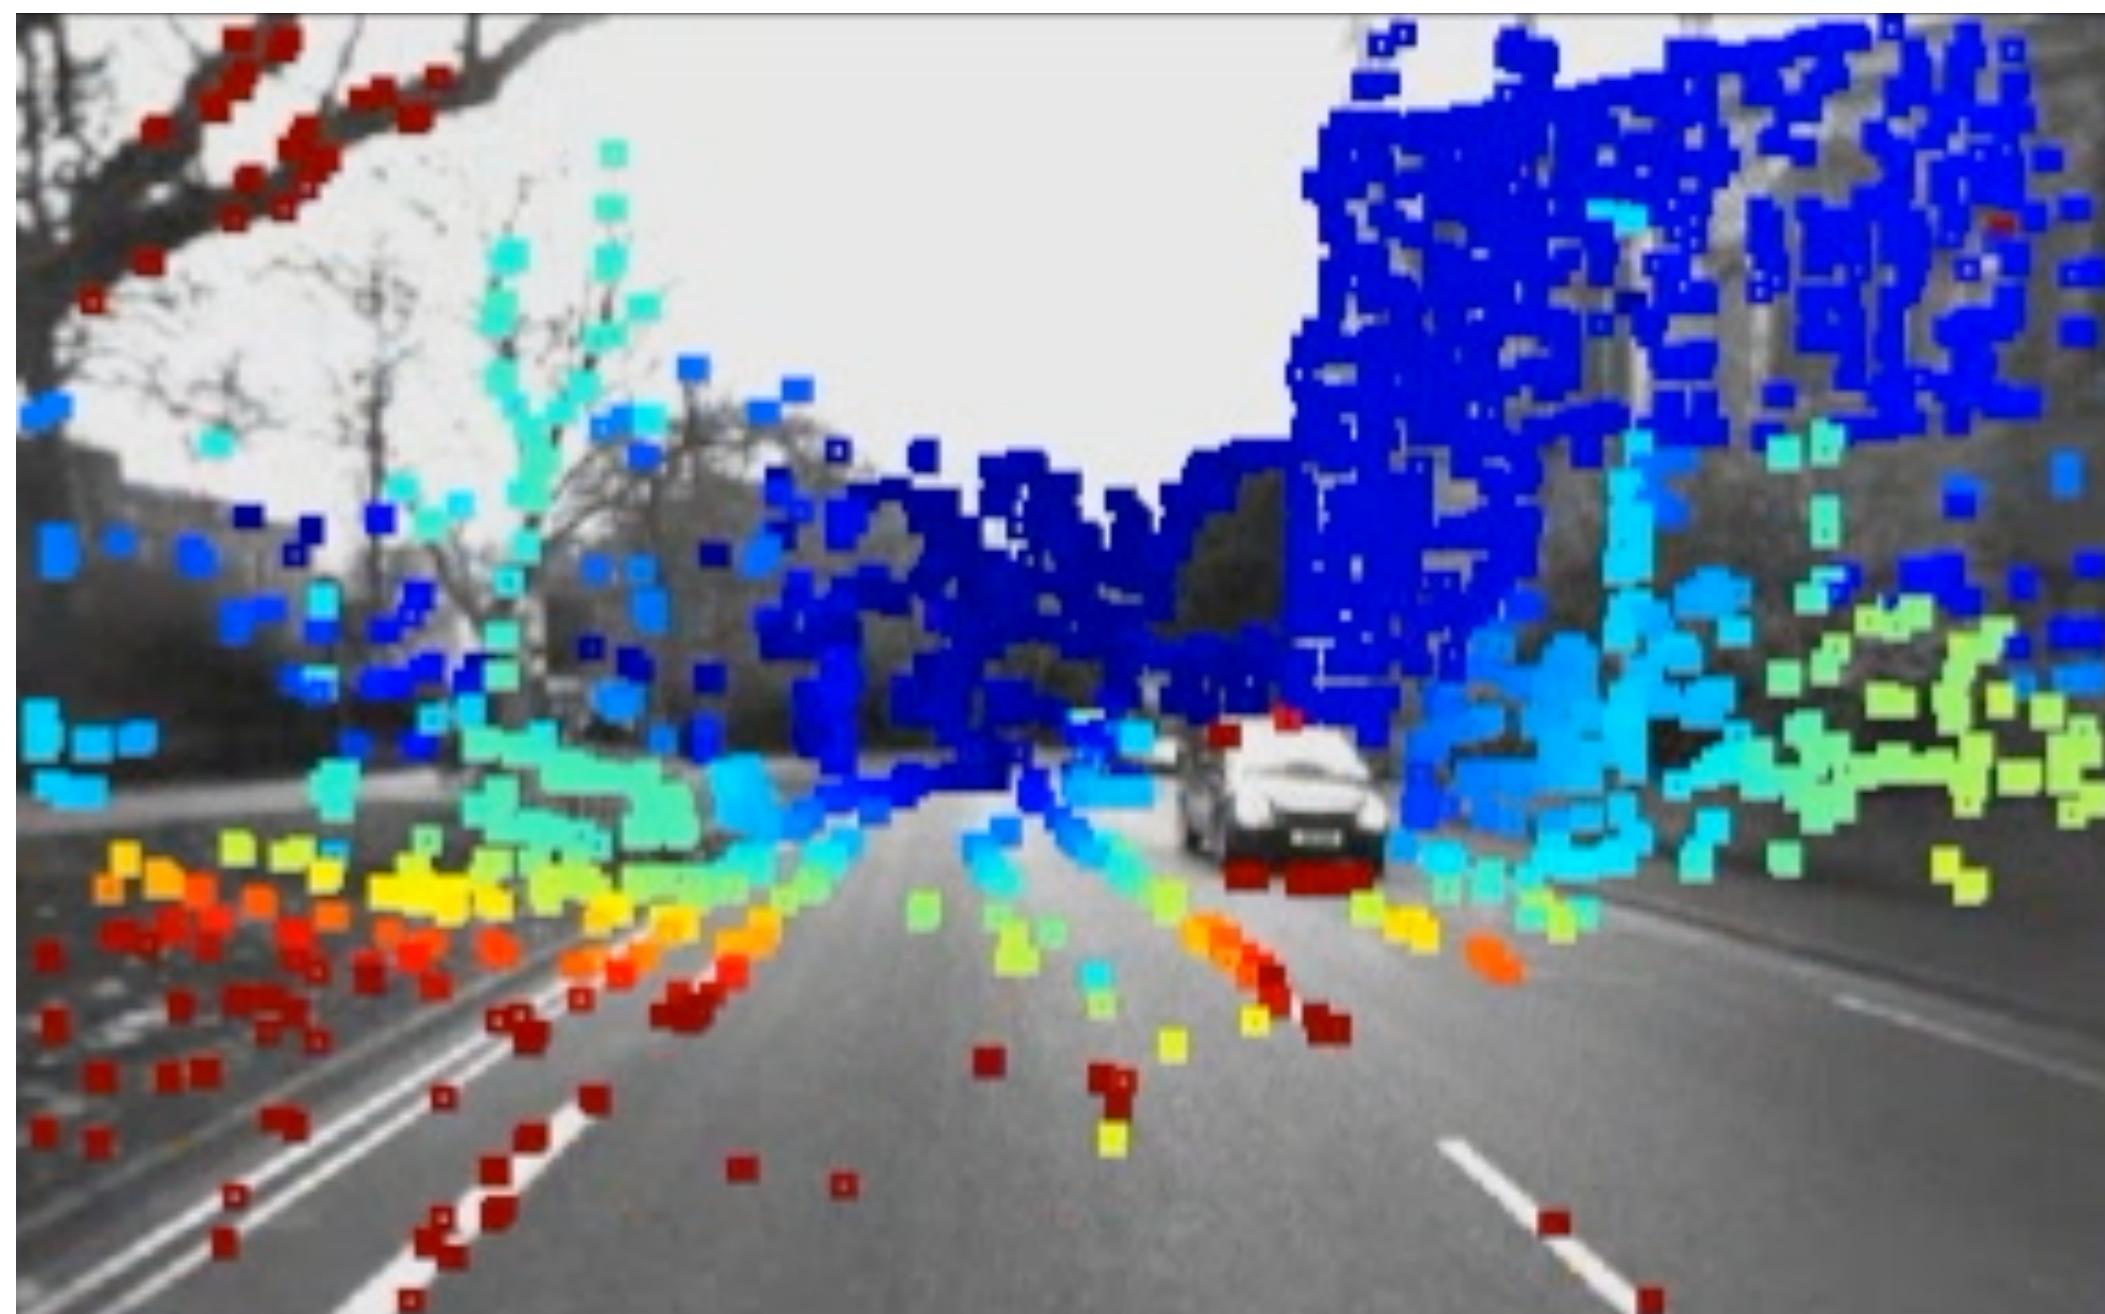
Direct Methods

- Drawbacks:
 - Need a good initialization
 - Cannot handle strong lighting / weather changes

→ Relocalization predominantly done with indirect methods for a long time



Direct Image Alignment

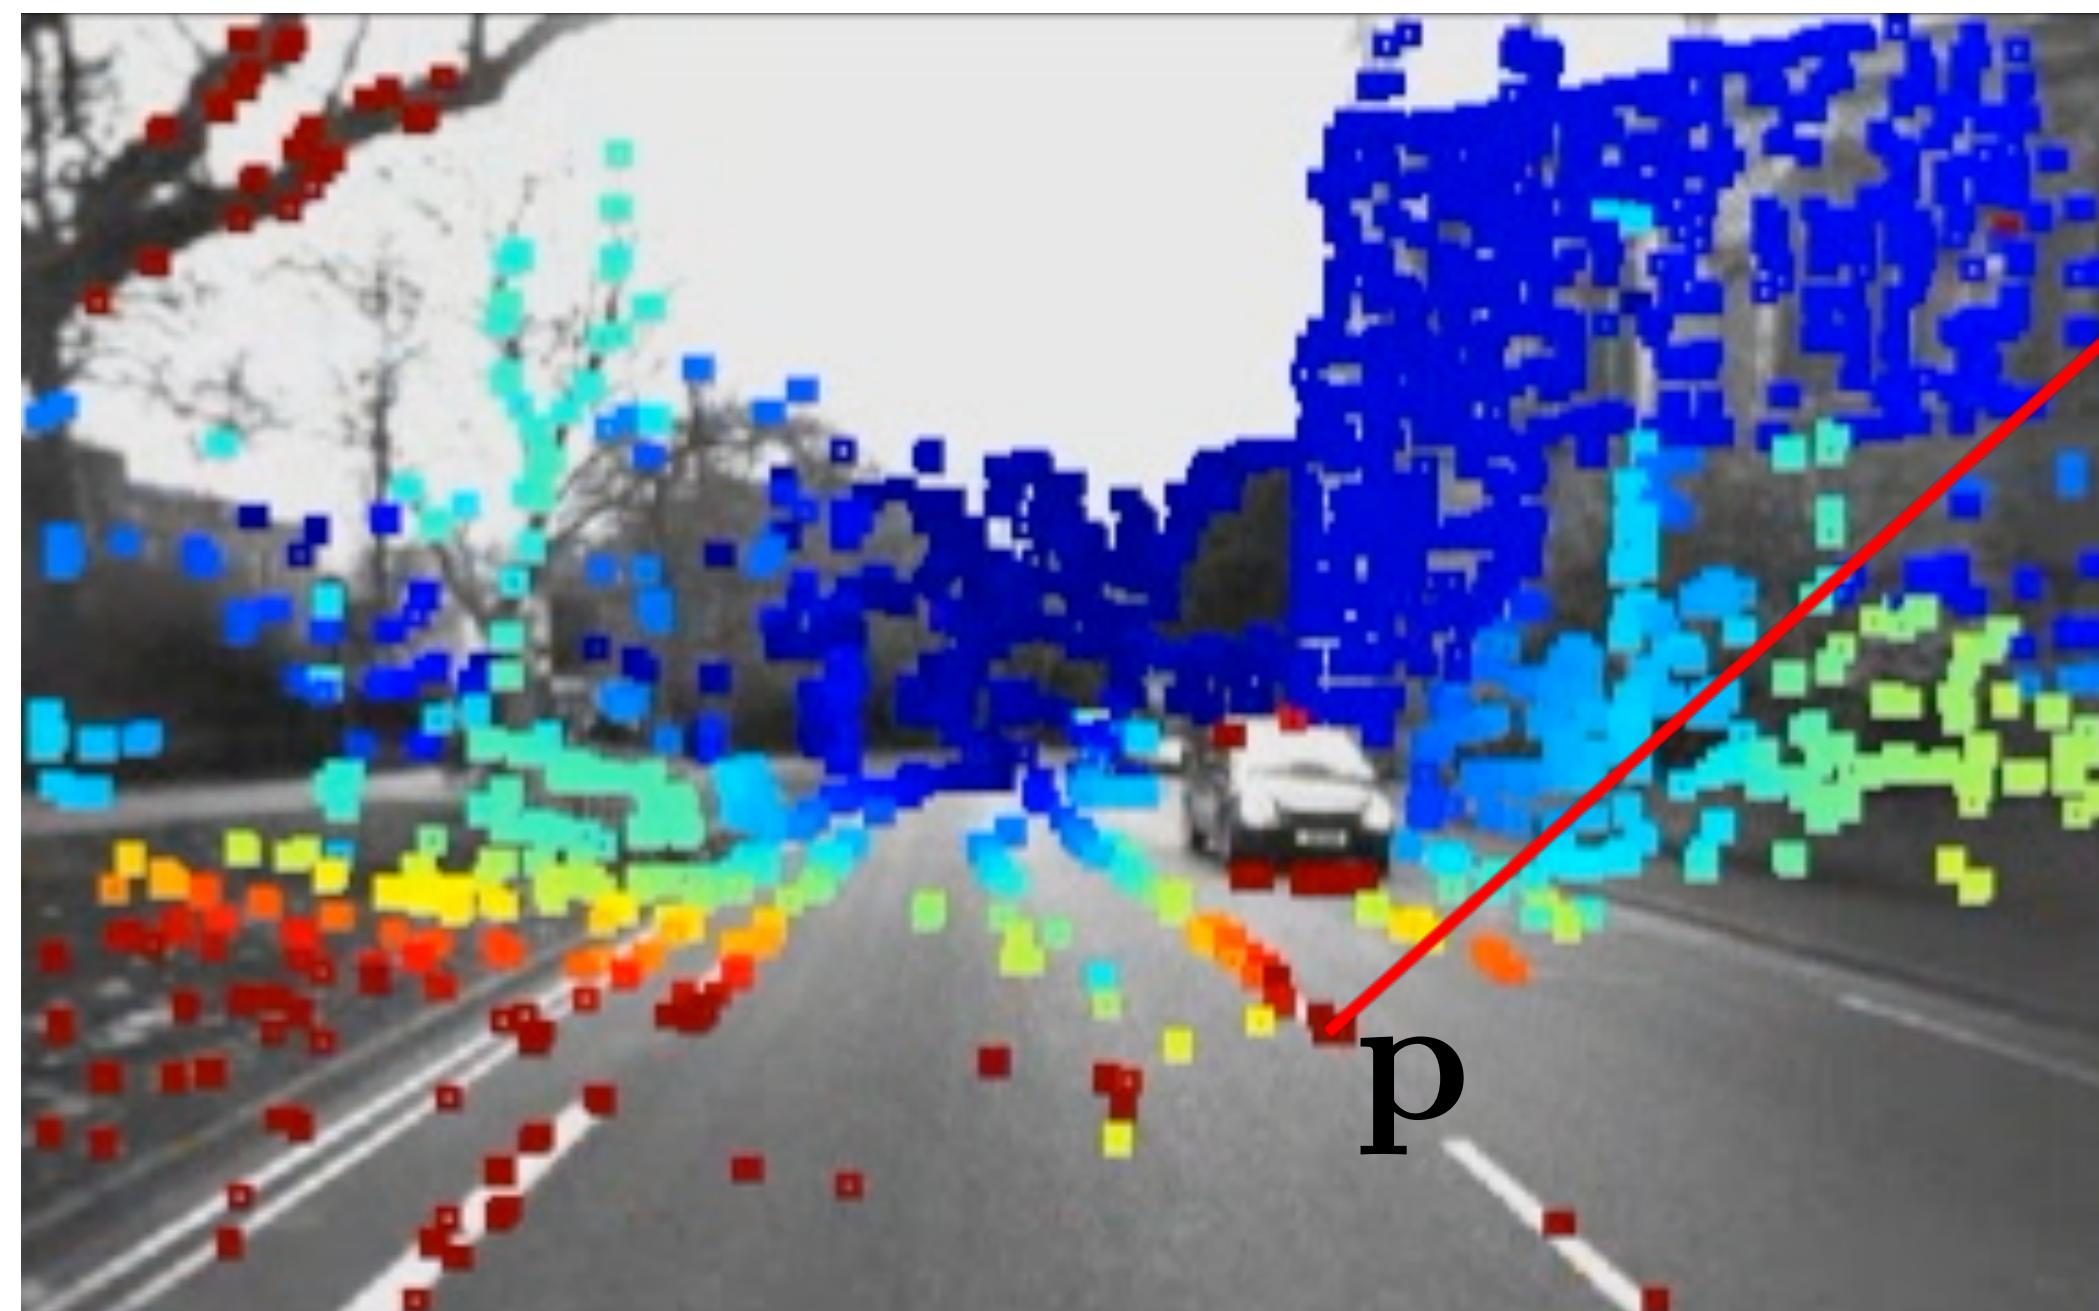


I

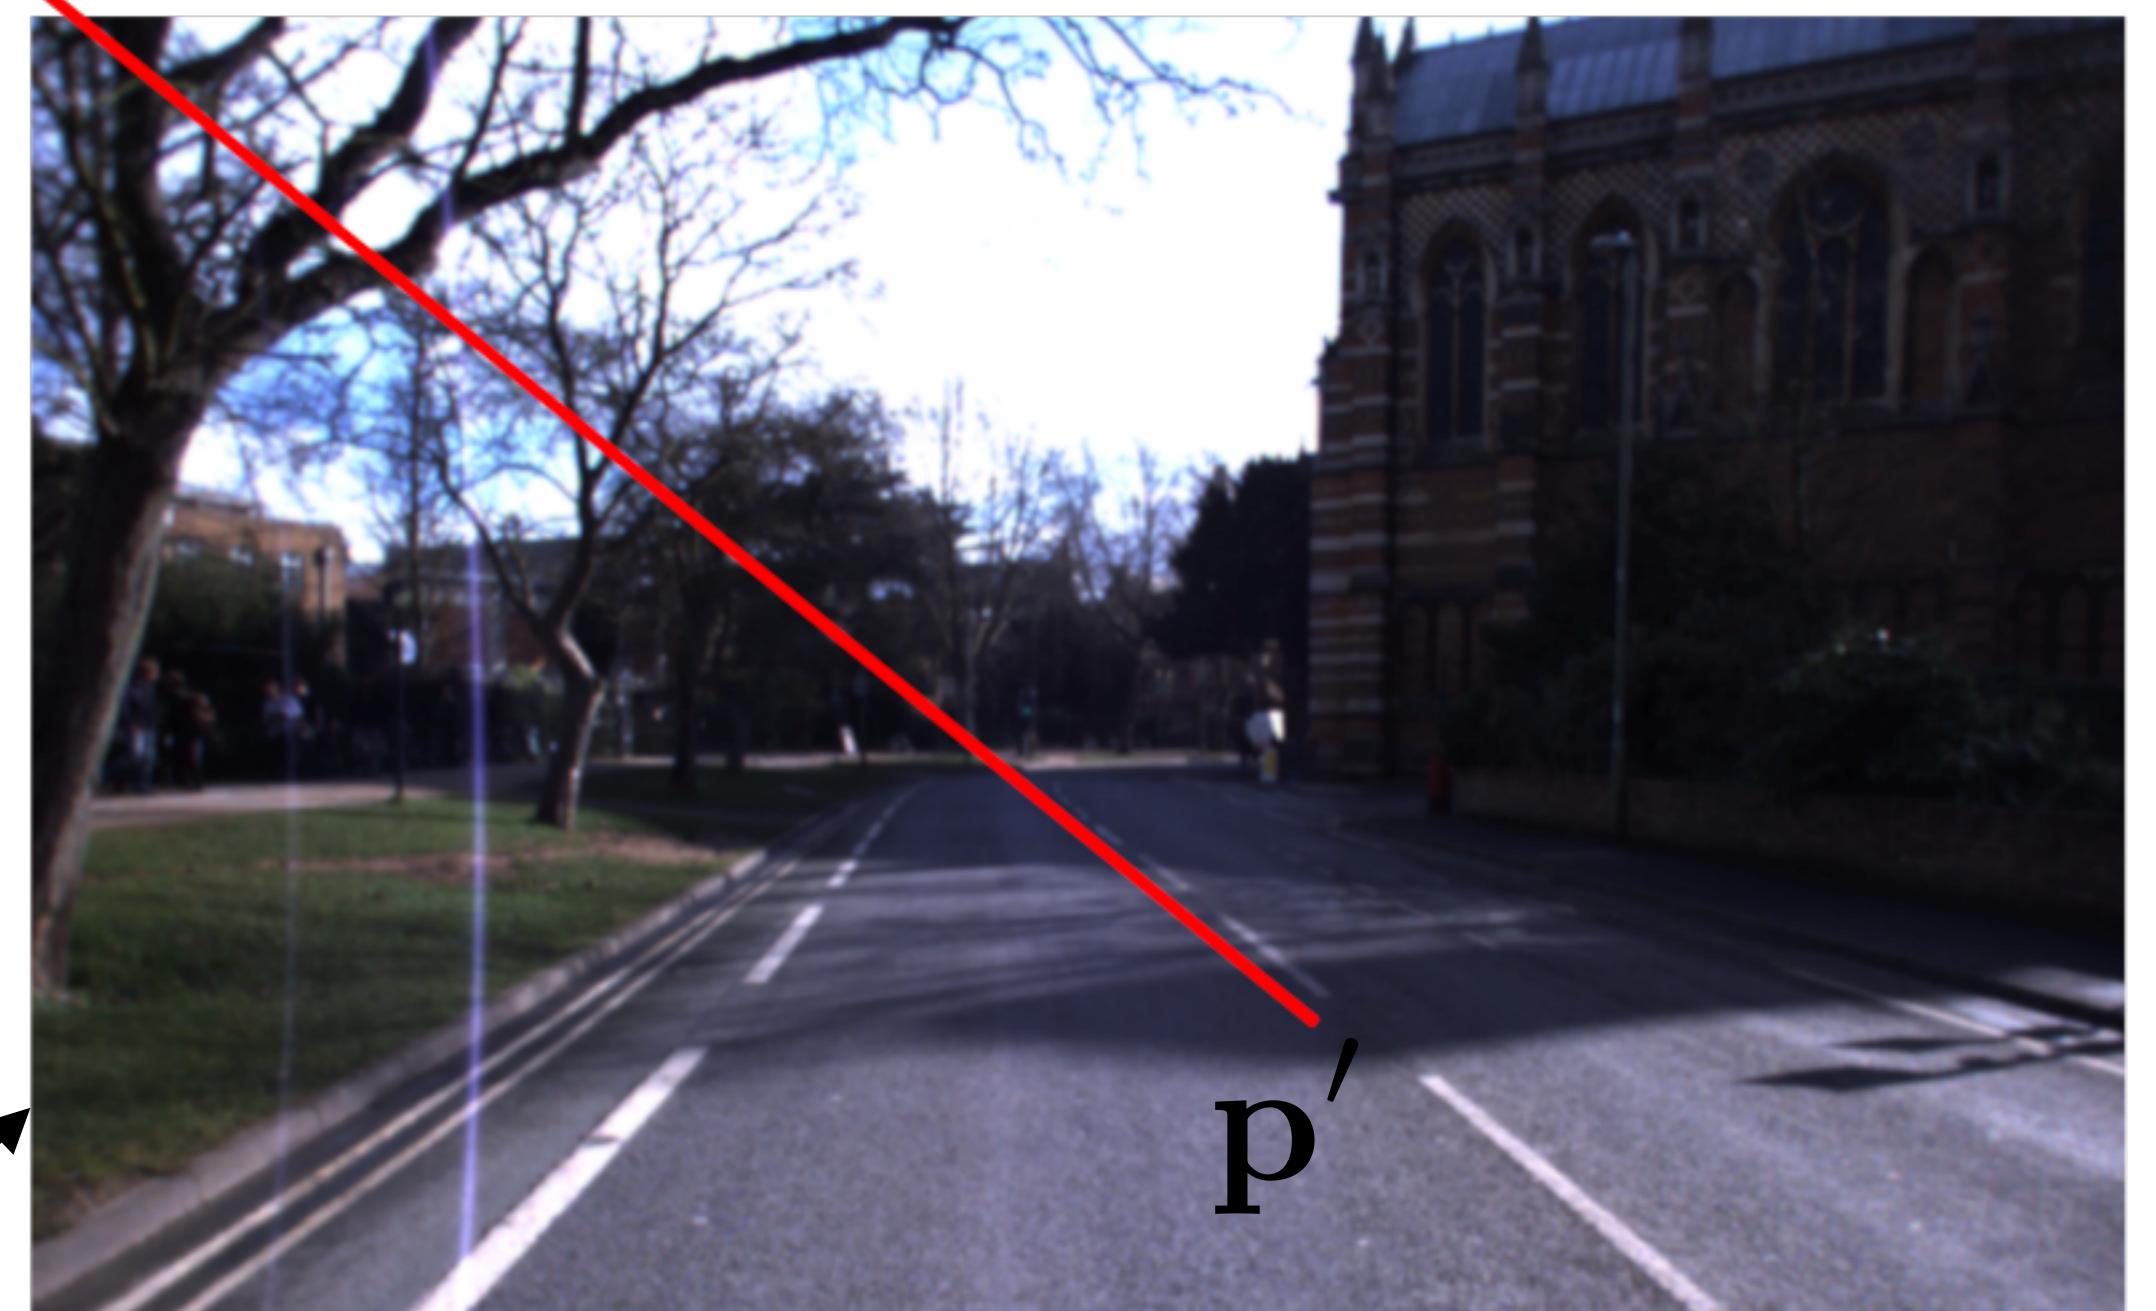


I'

Direct Image Alignment



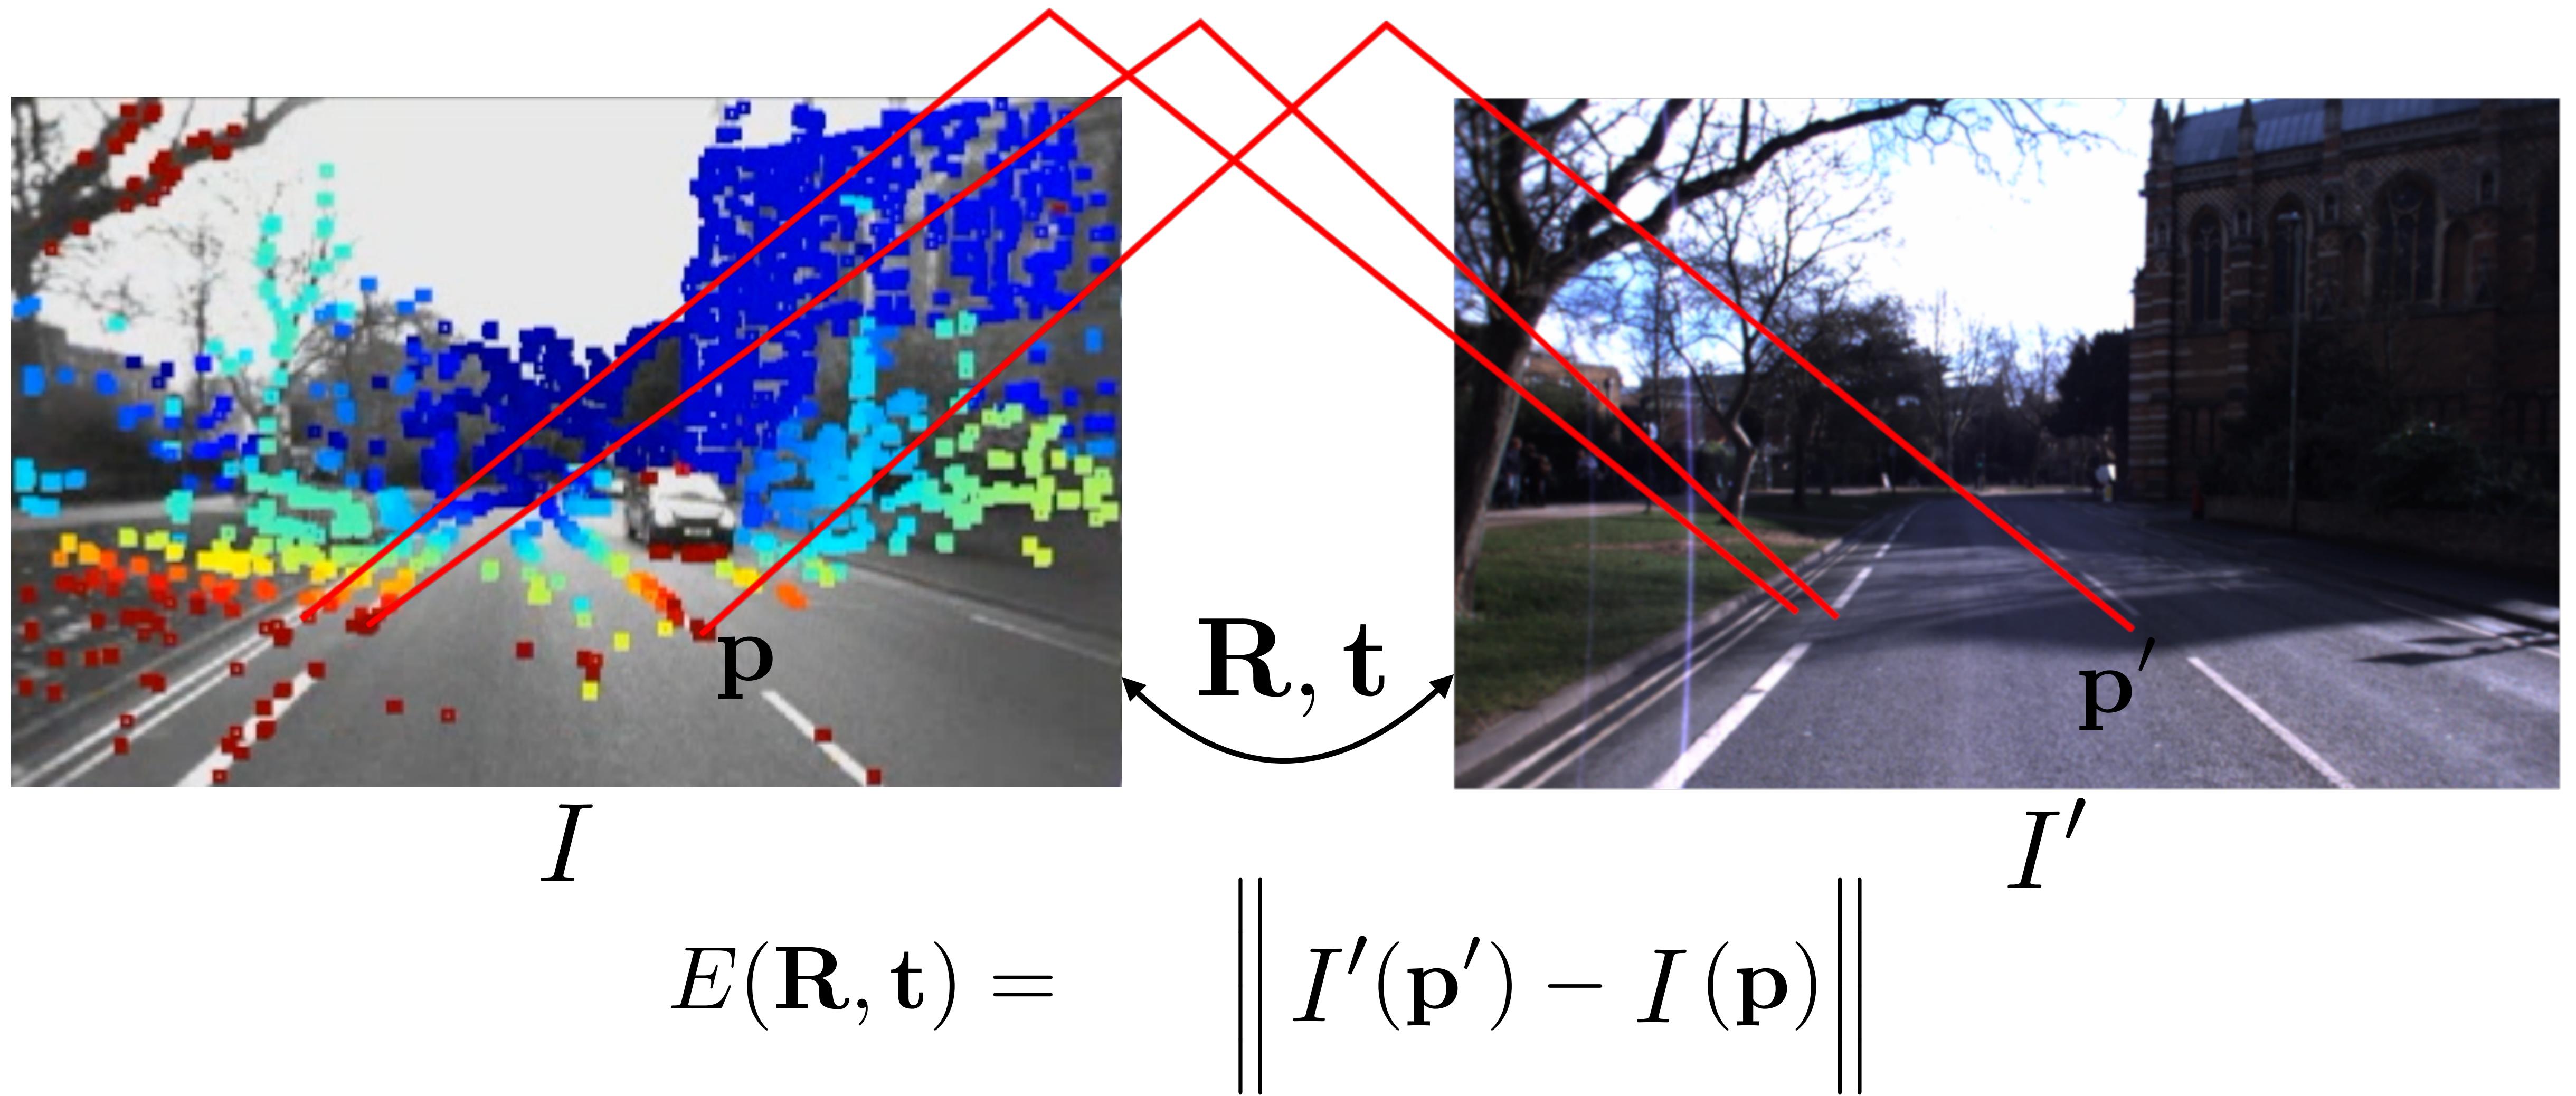
I



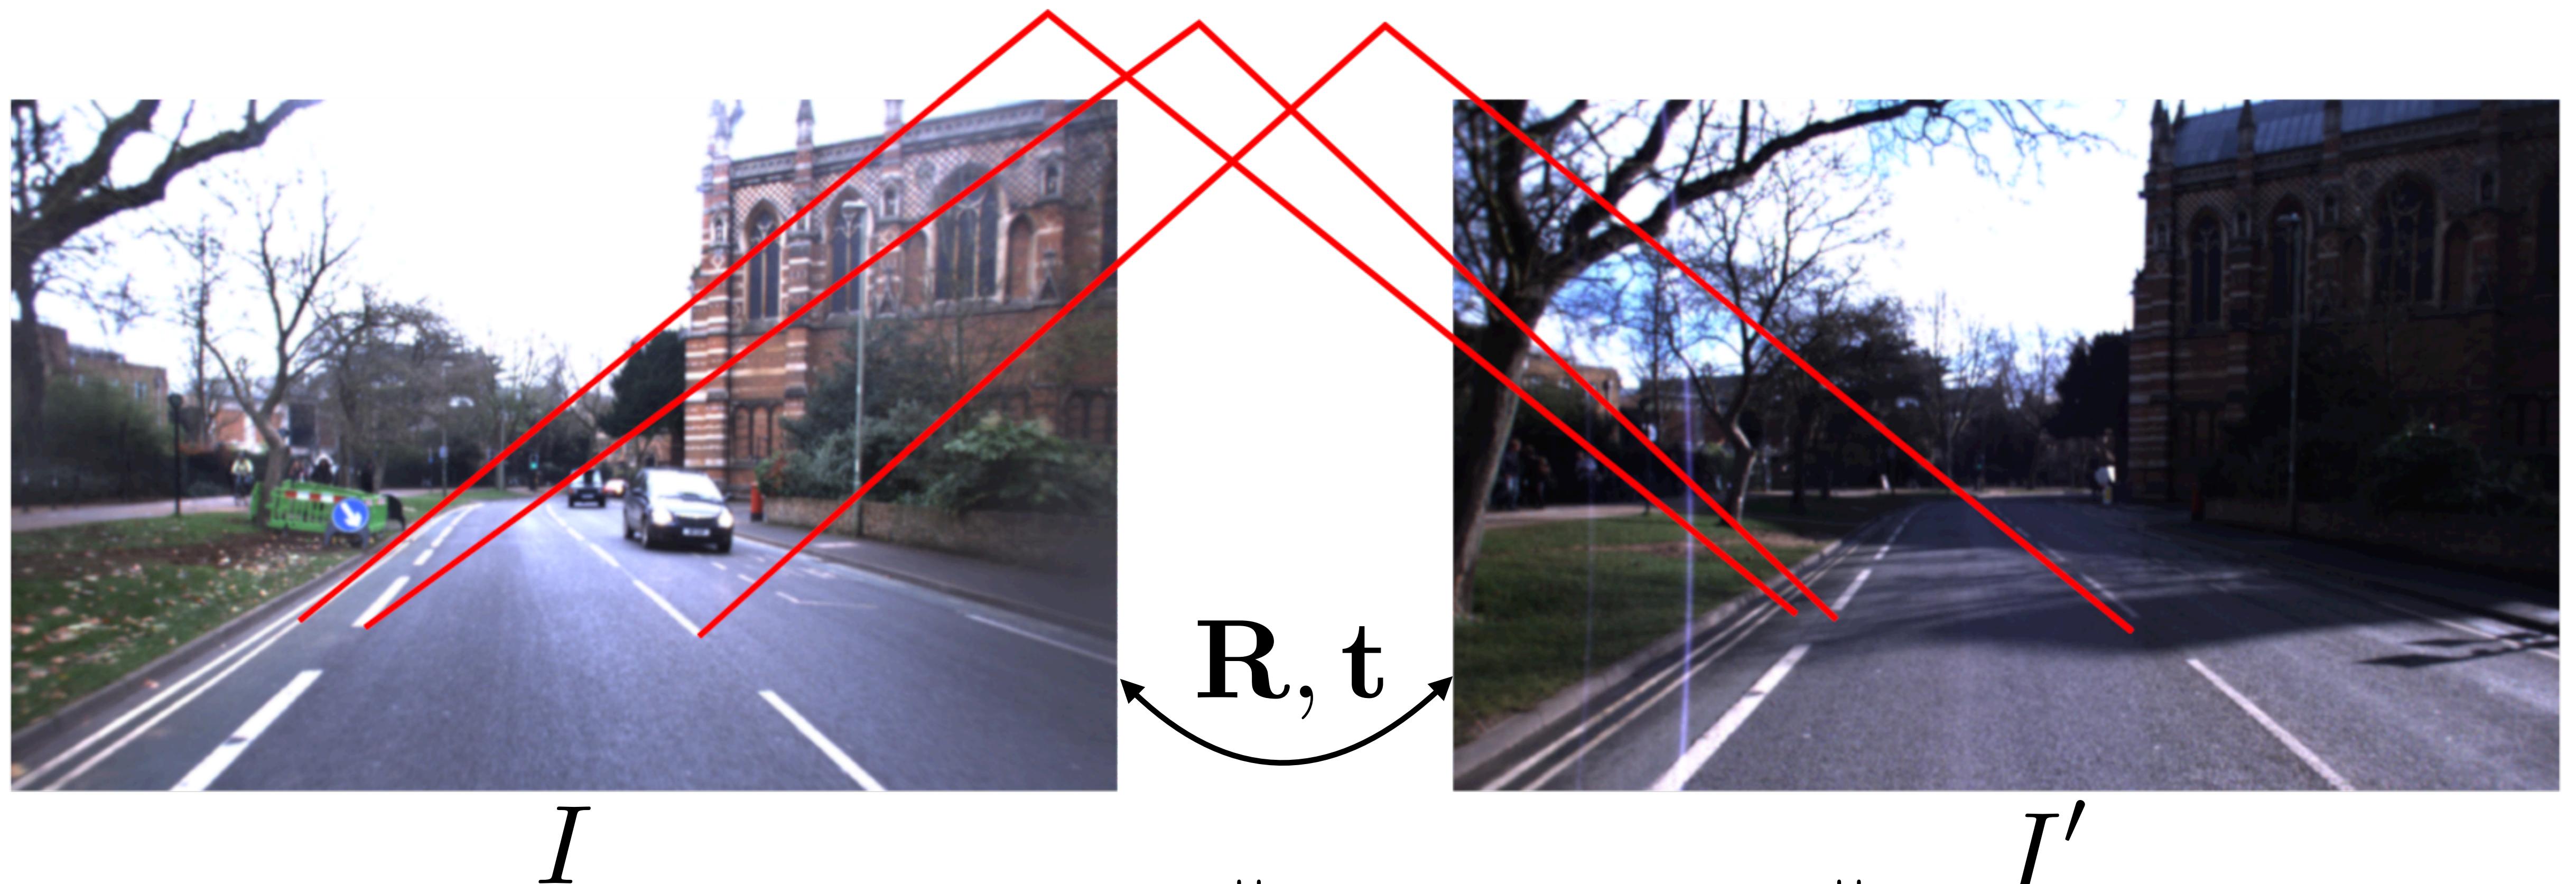
I'

$$E(R, t) = \| I'(p') - I(p) \|$$

Direct Image Alignment



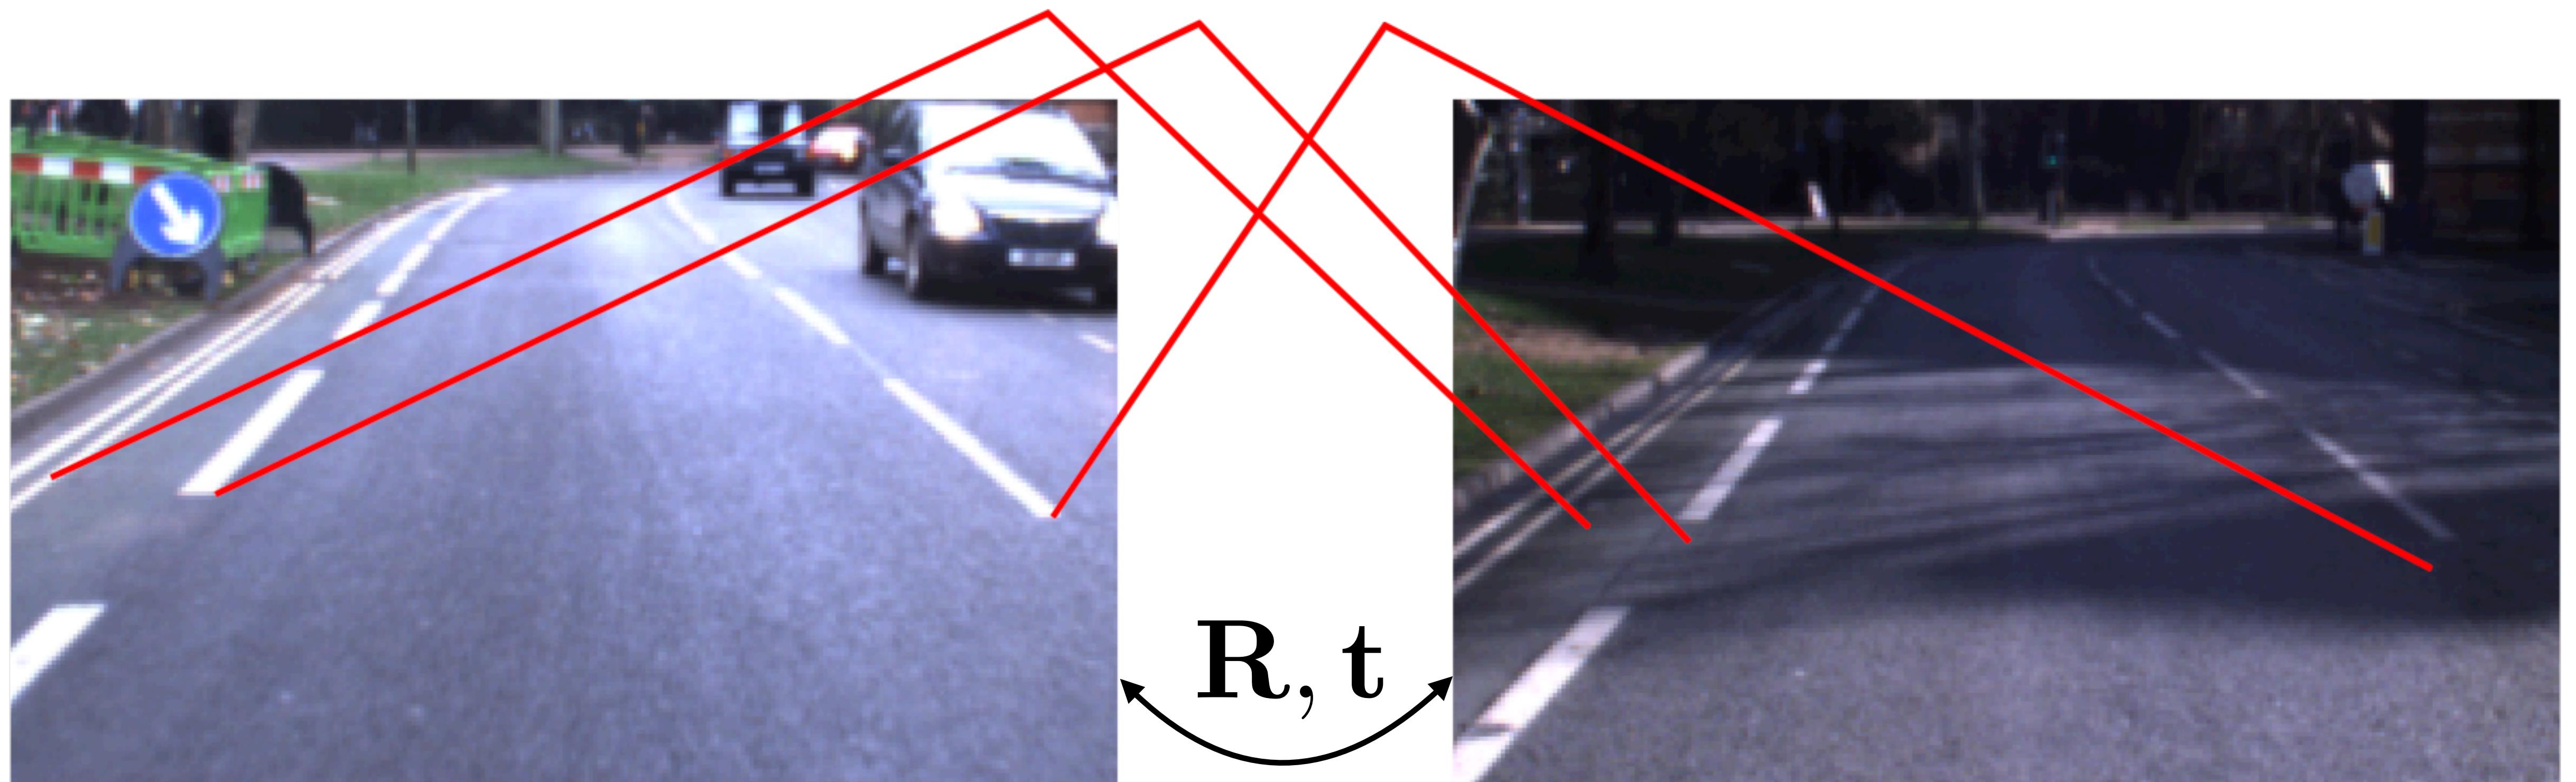
Direct Image Alignment



LM / Gauss-Newton Optimization

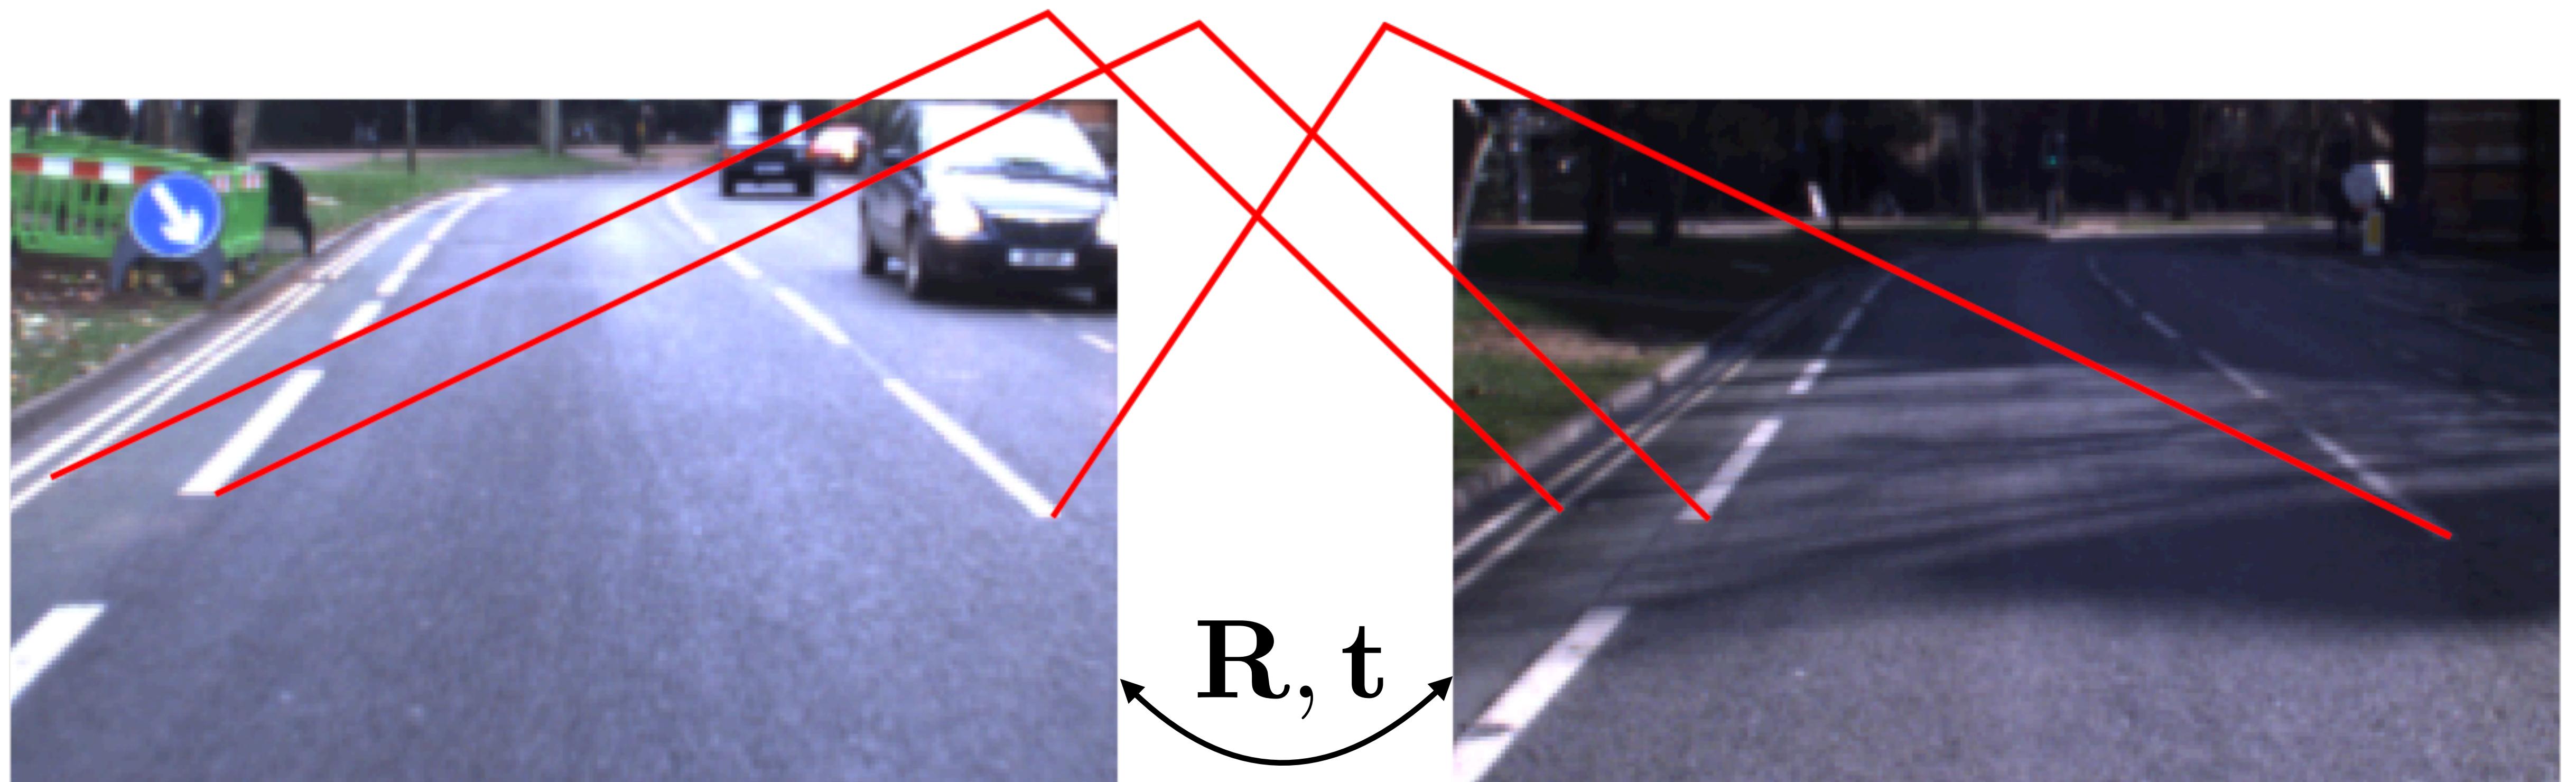
$$E(R, t) = \sum_p \| I'(p') - I(p) \|$$

Direct Image Alignment

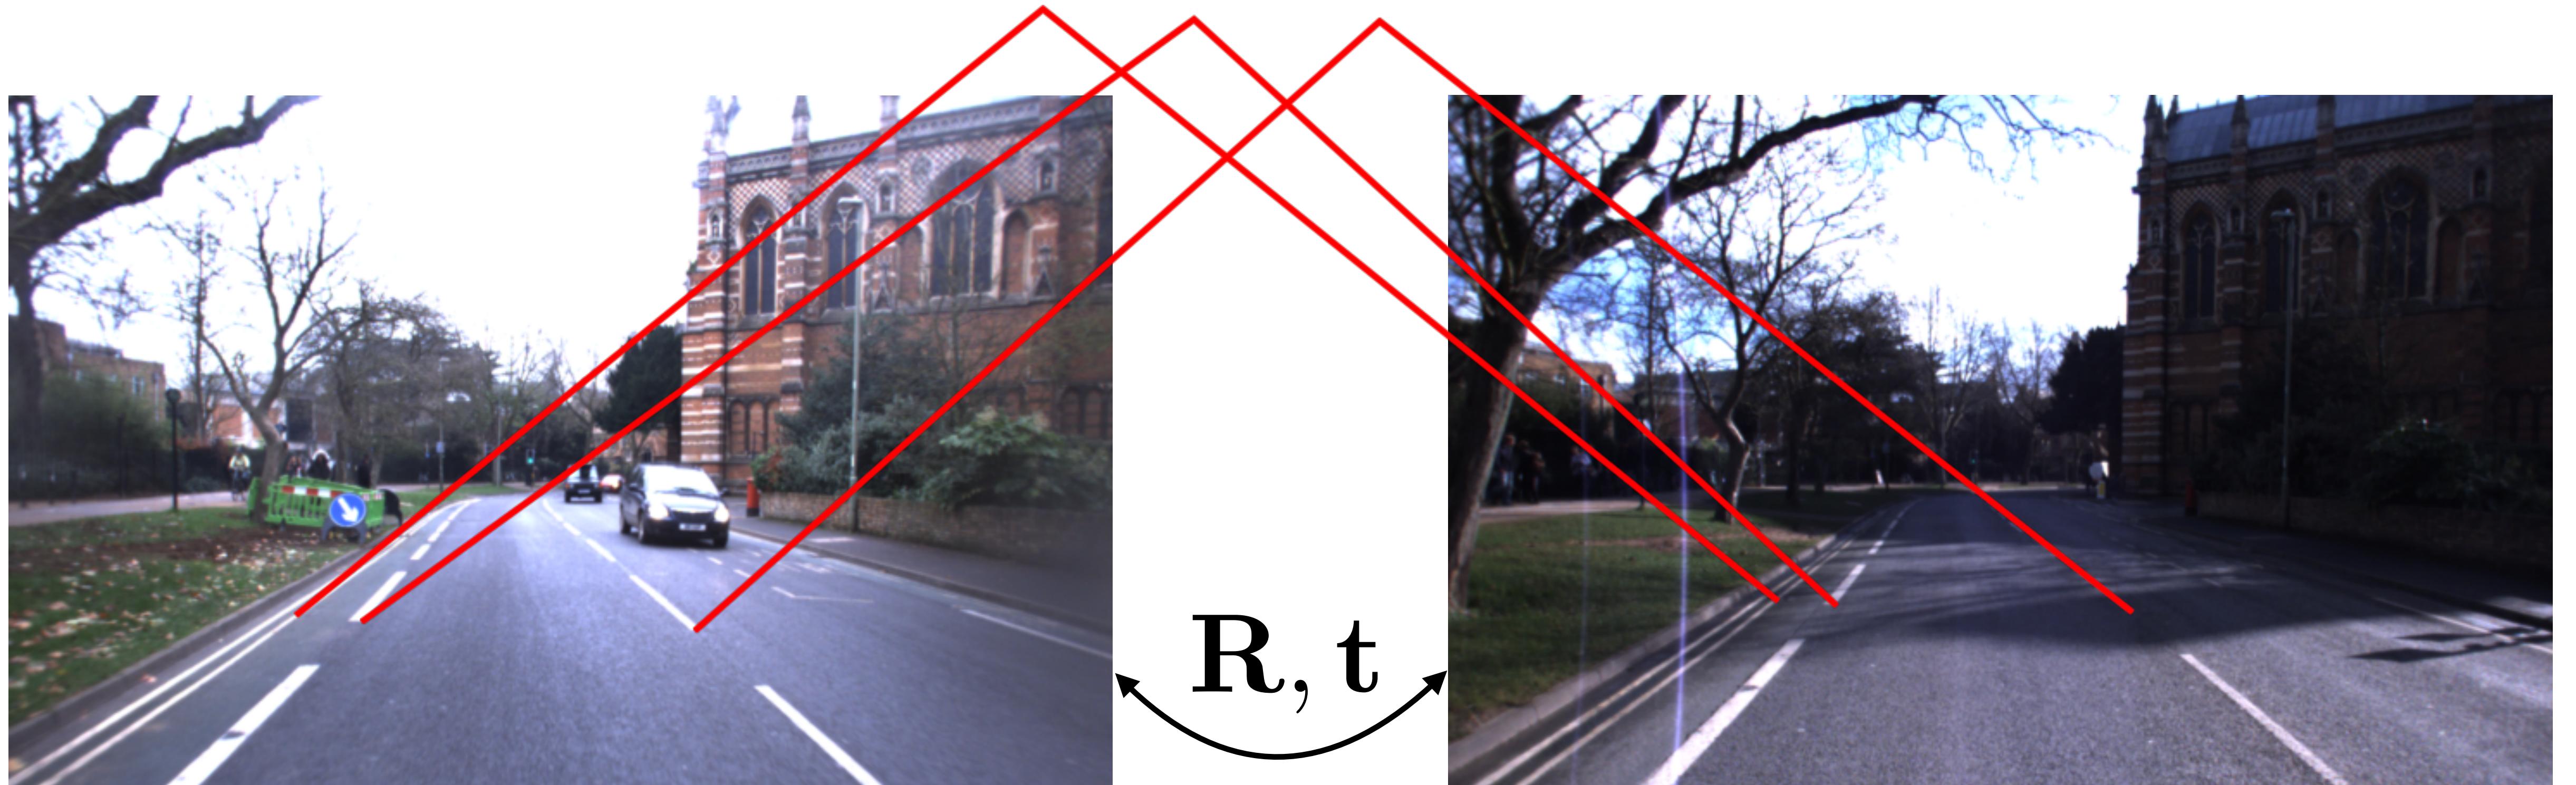


$$\text{LM / Gauss-Newton Optimization} \quad E(R, t) = \sum_p \| I'(p') - I(p) \|$$

Direct Image Alignment

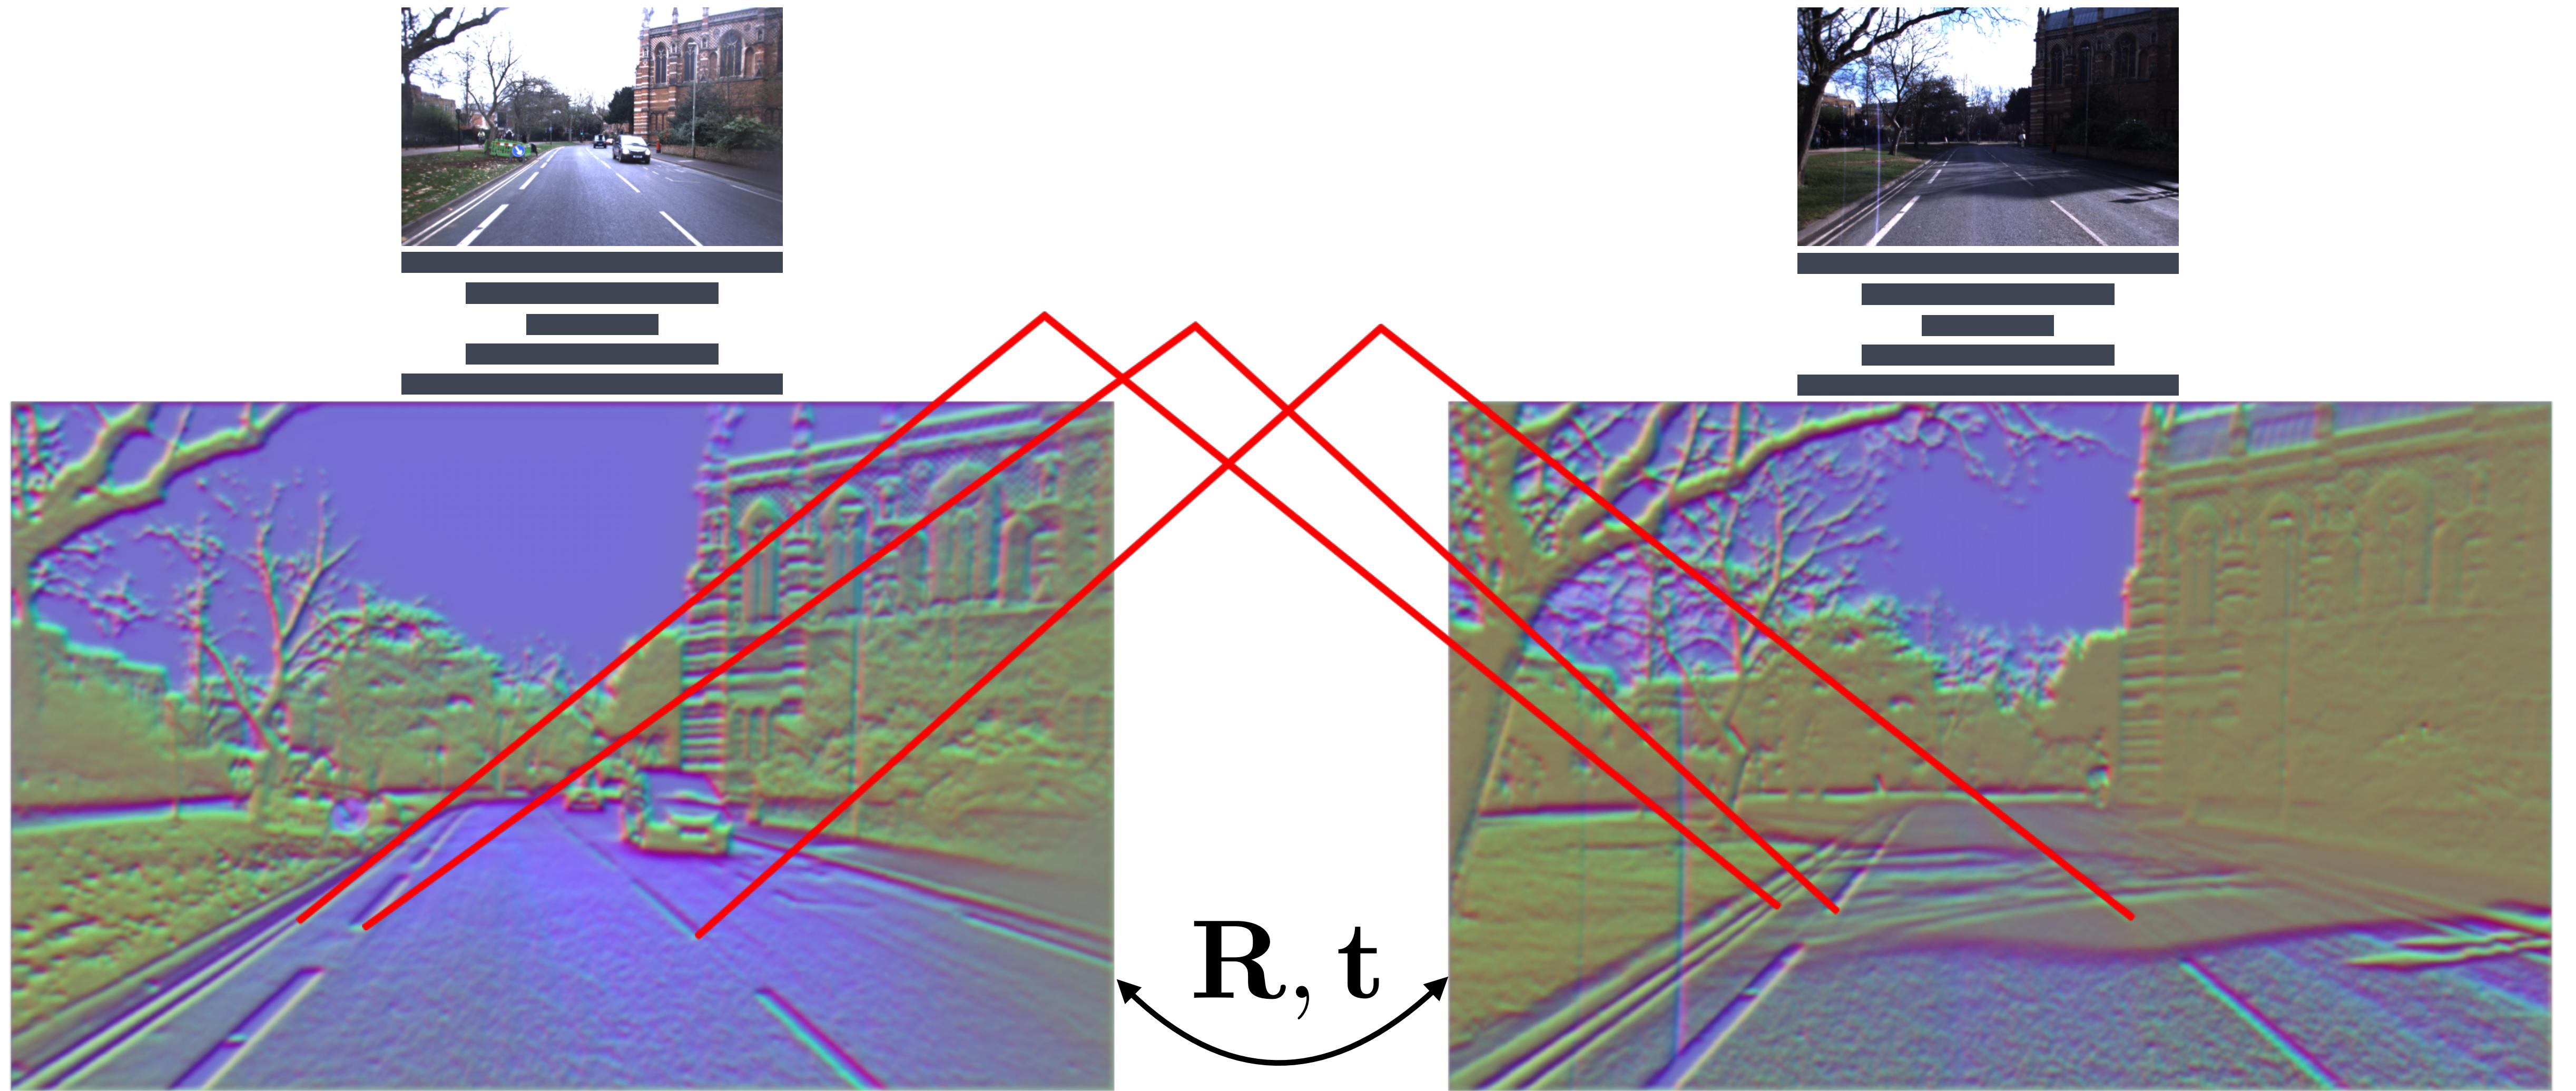


$$\text{LM / Gauss-Newton Optimization} \quad E(\mathbf{R}, \mathbf{t}) = \sum_{\mathbf{p}} \left\| I'(\mathbf{p}') - I(\mathbf{p}) \right\|$$



LM / Gauss-Newton Optimization

$$E(R, t) = \sum_p \| I'(p') - I(p) \|$$

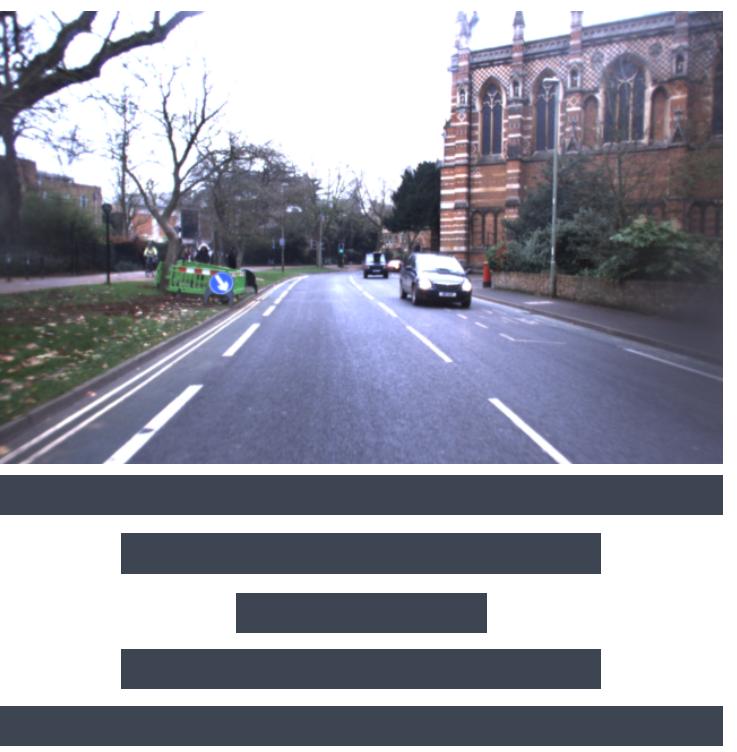


F

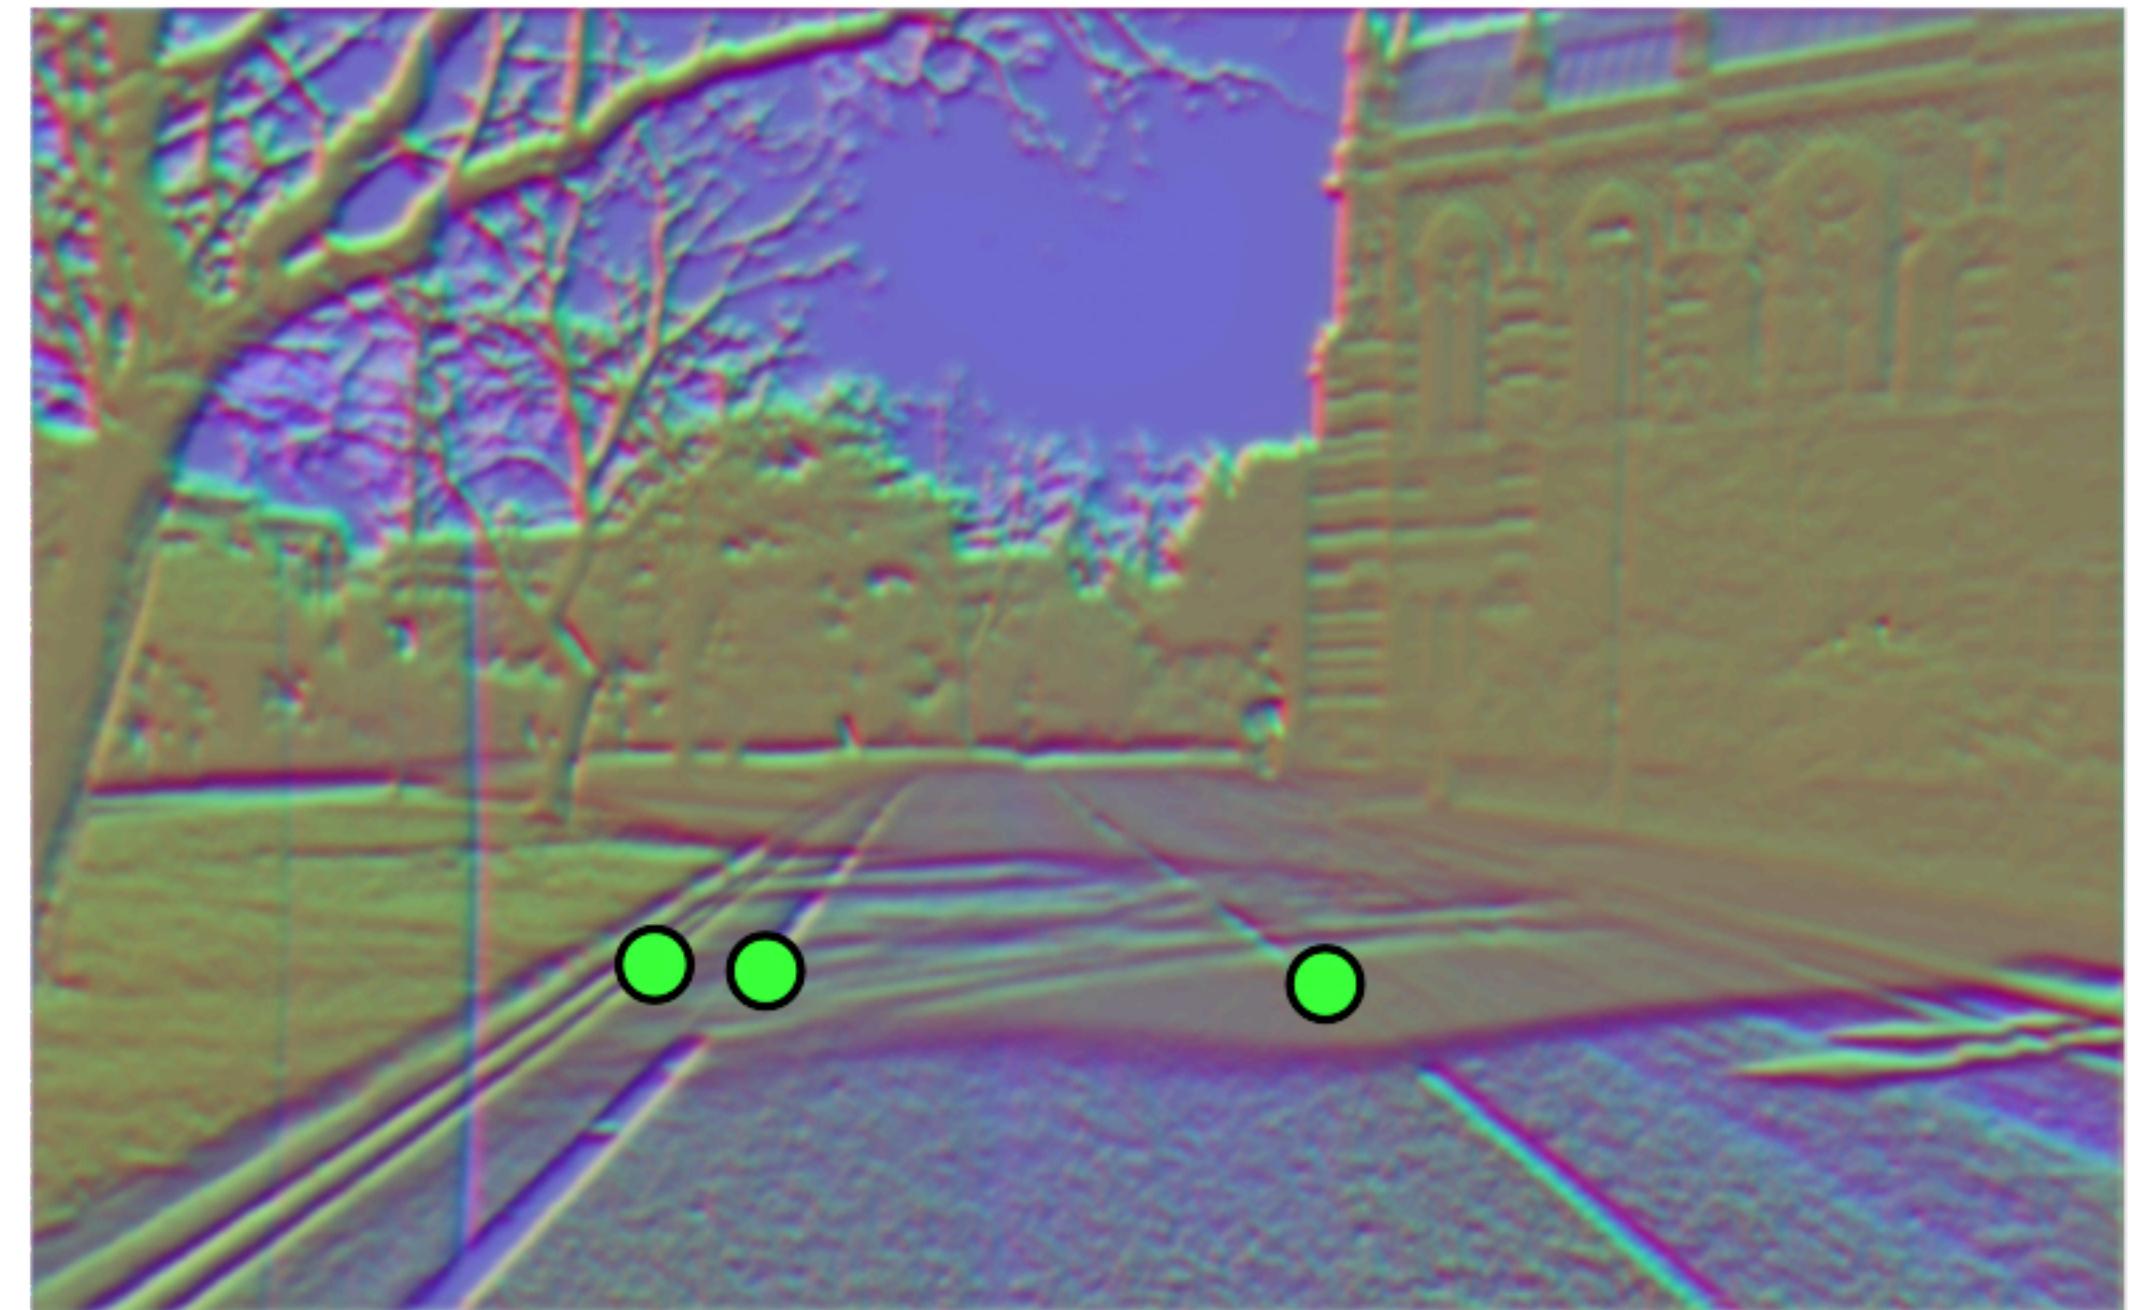
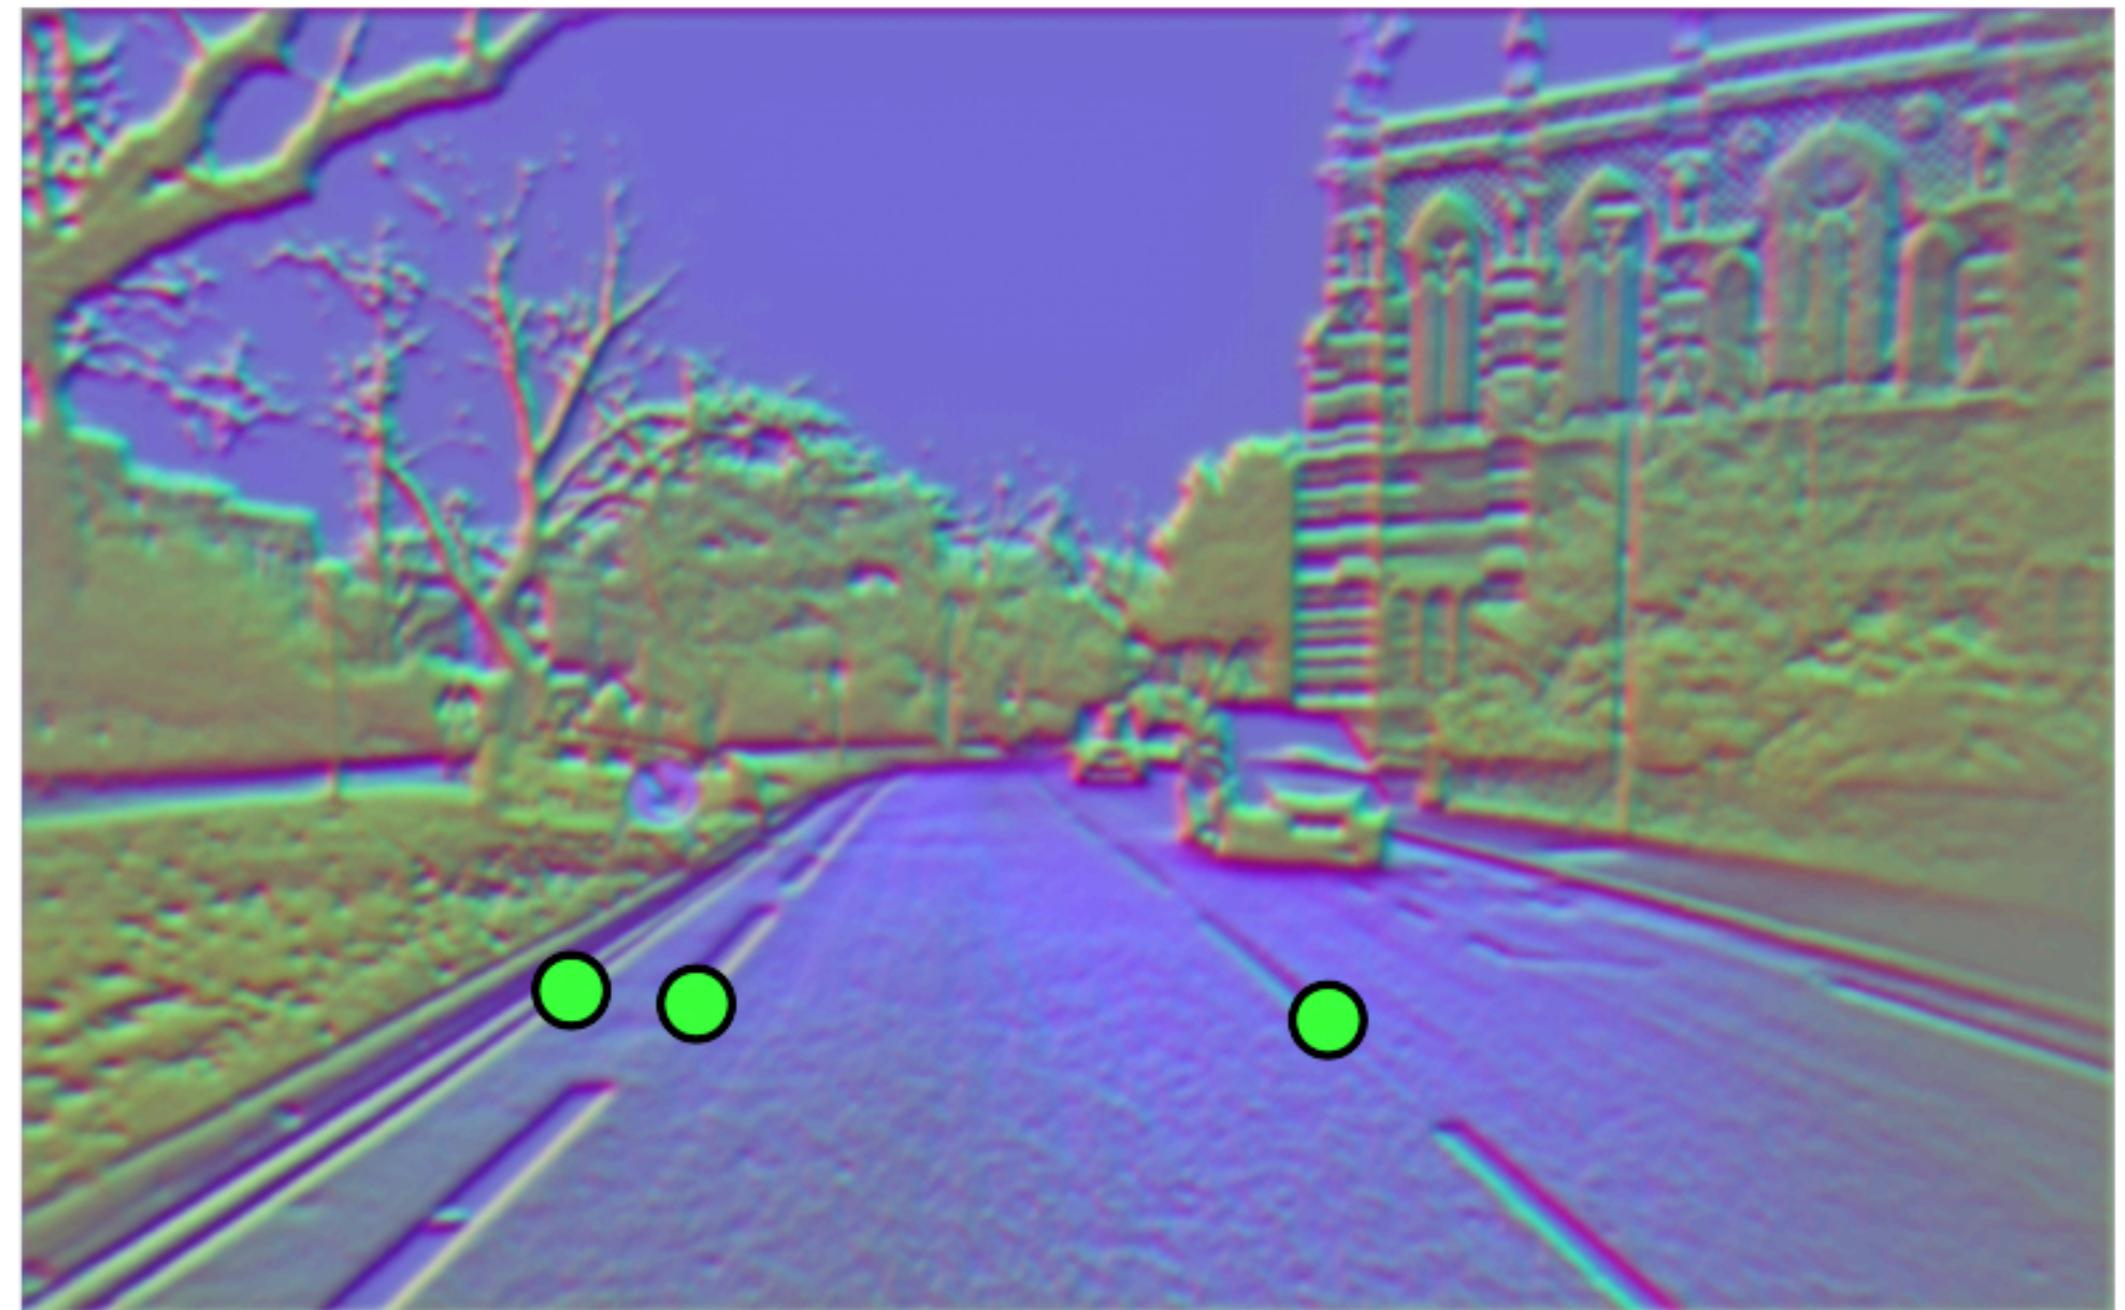
LM / Gauss-Newton
Optimization

$$E(\mathbf{R}, \mathbf{t}) = \sum_{\mathbf{p}} \left\| I'(\mathbf{p}') - I(\mathbf{p}) \right\|$$

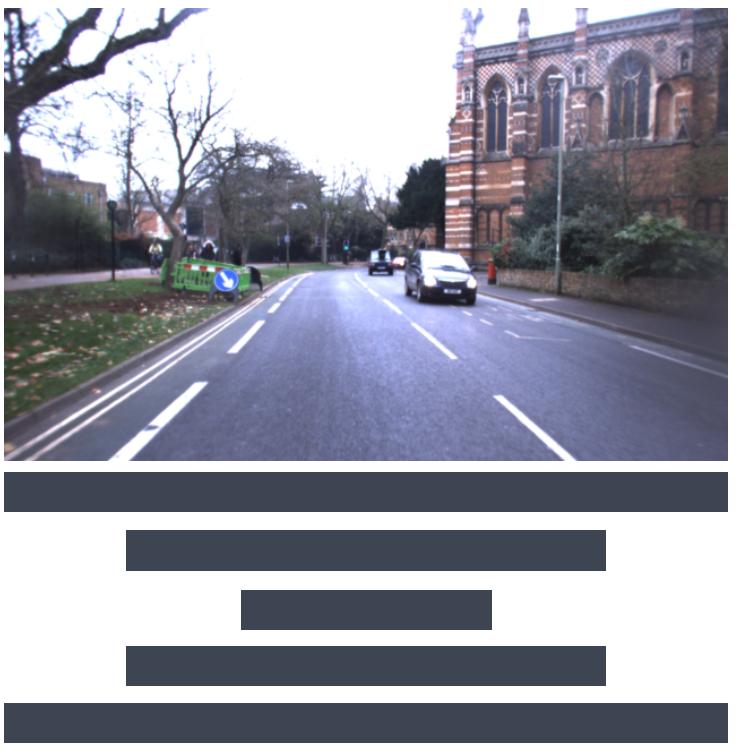
F'



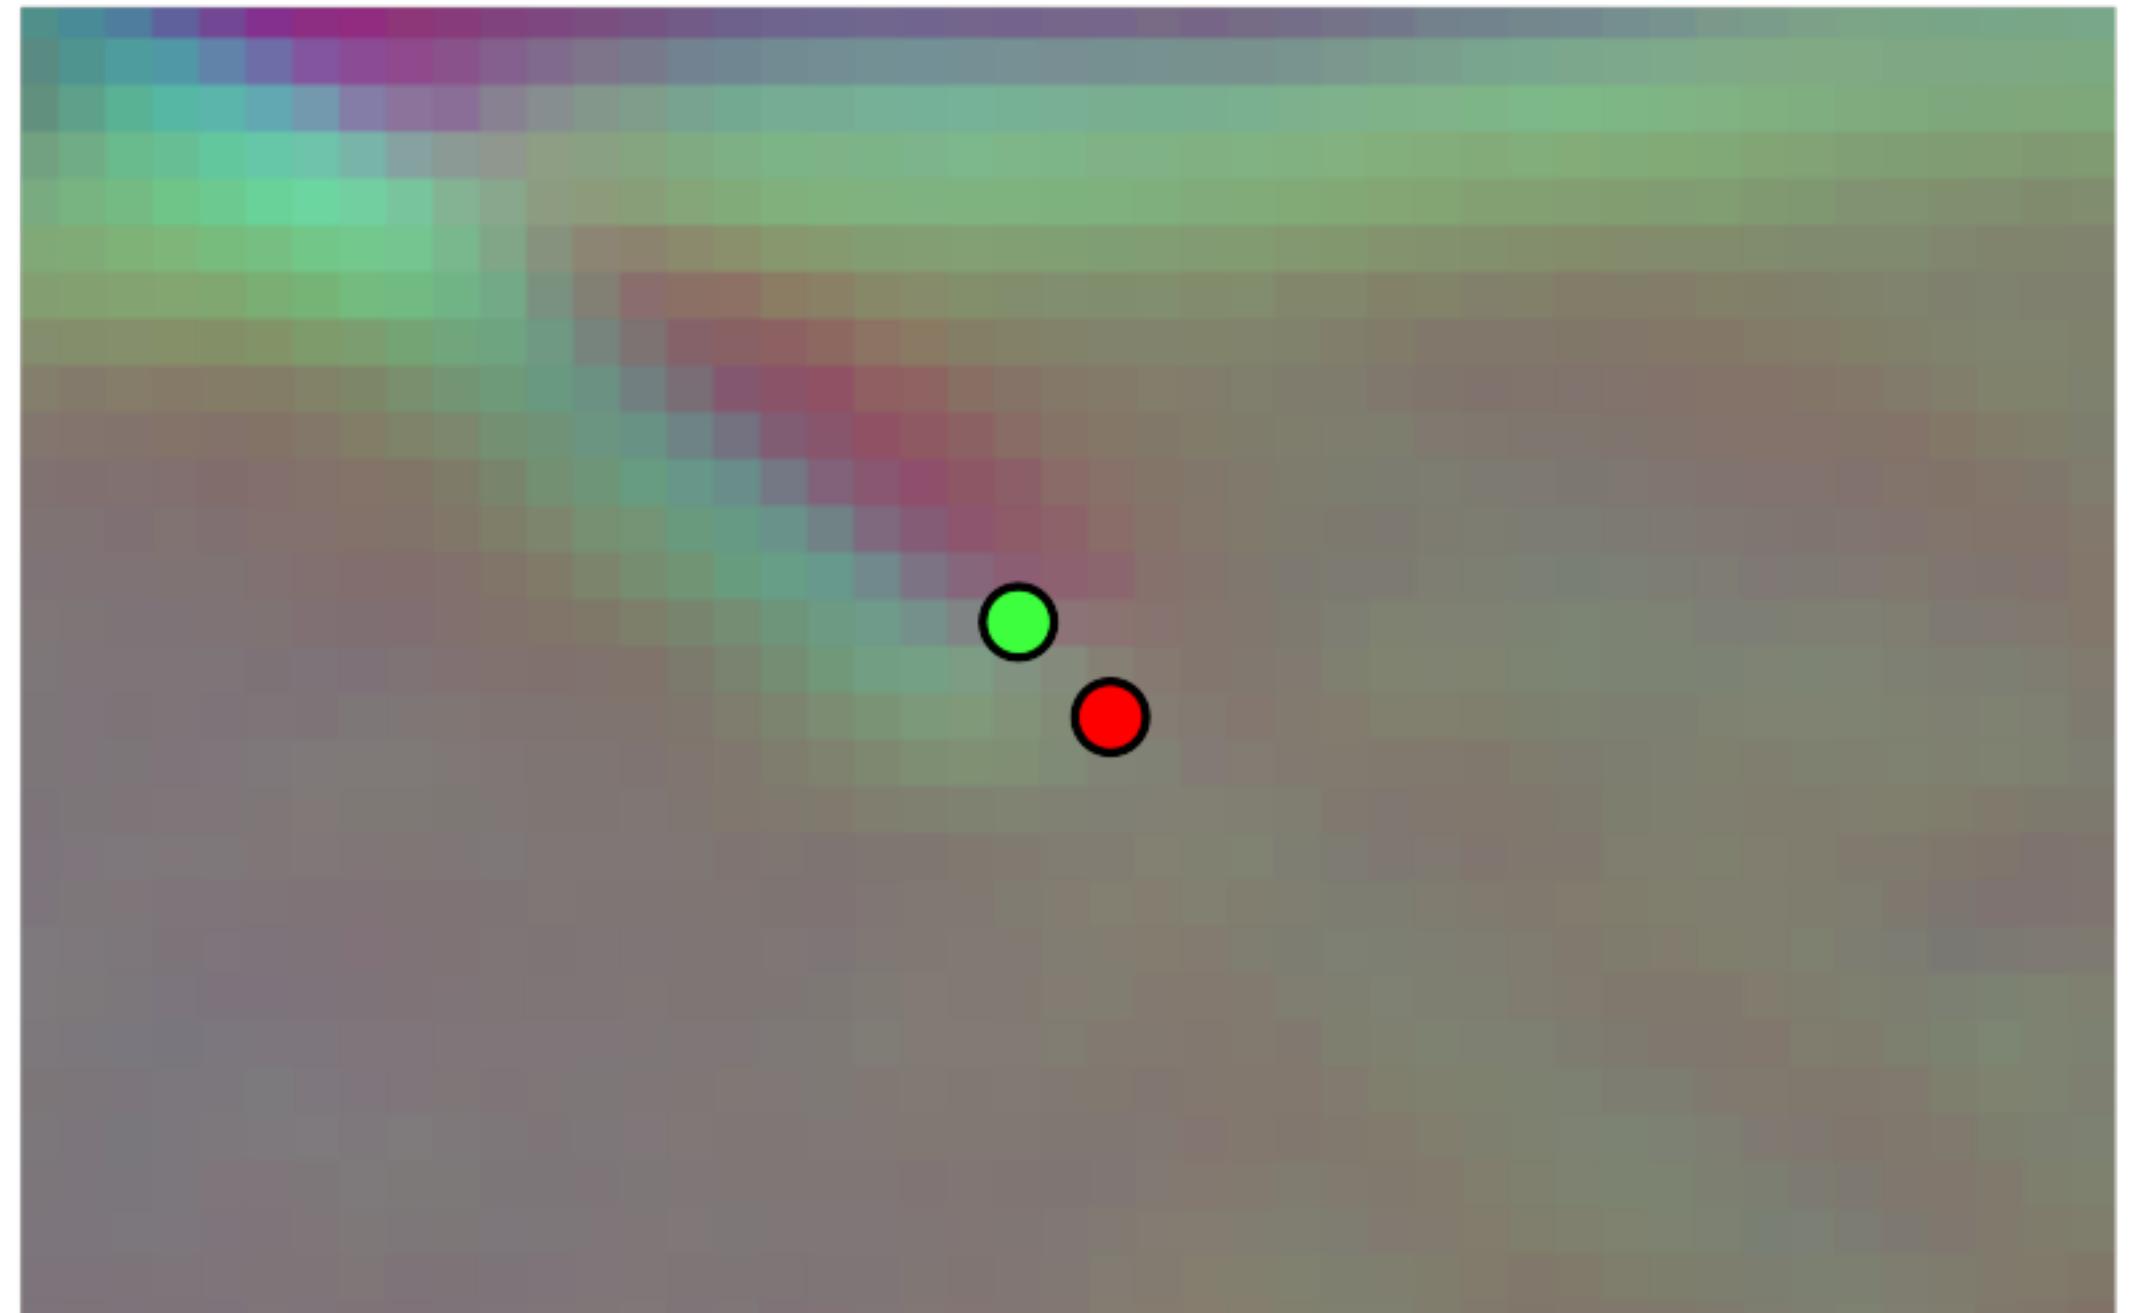
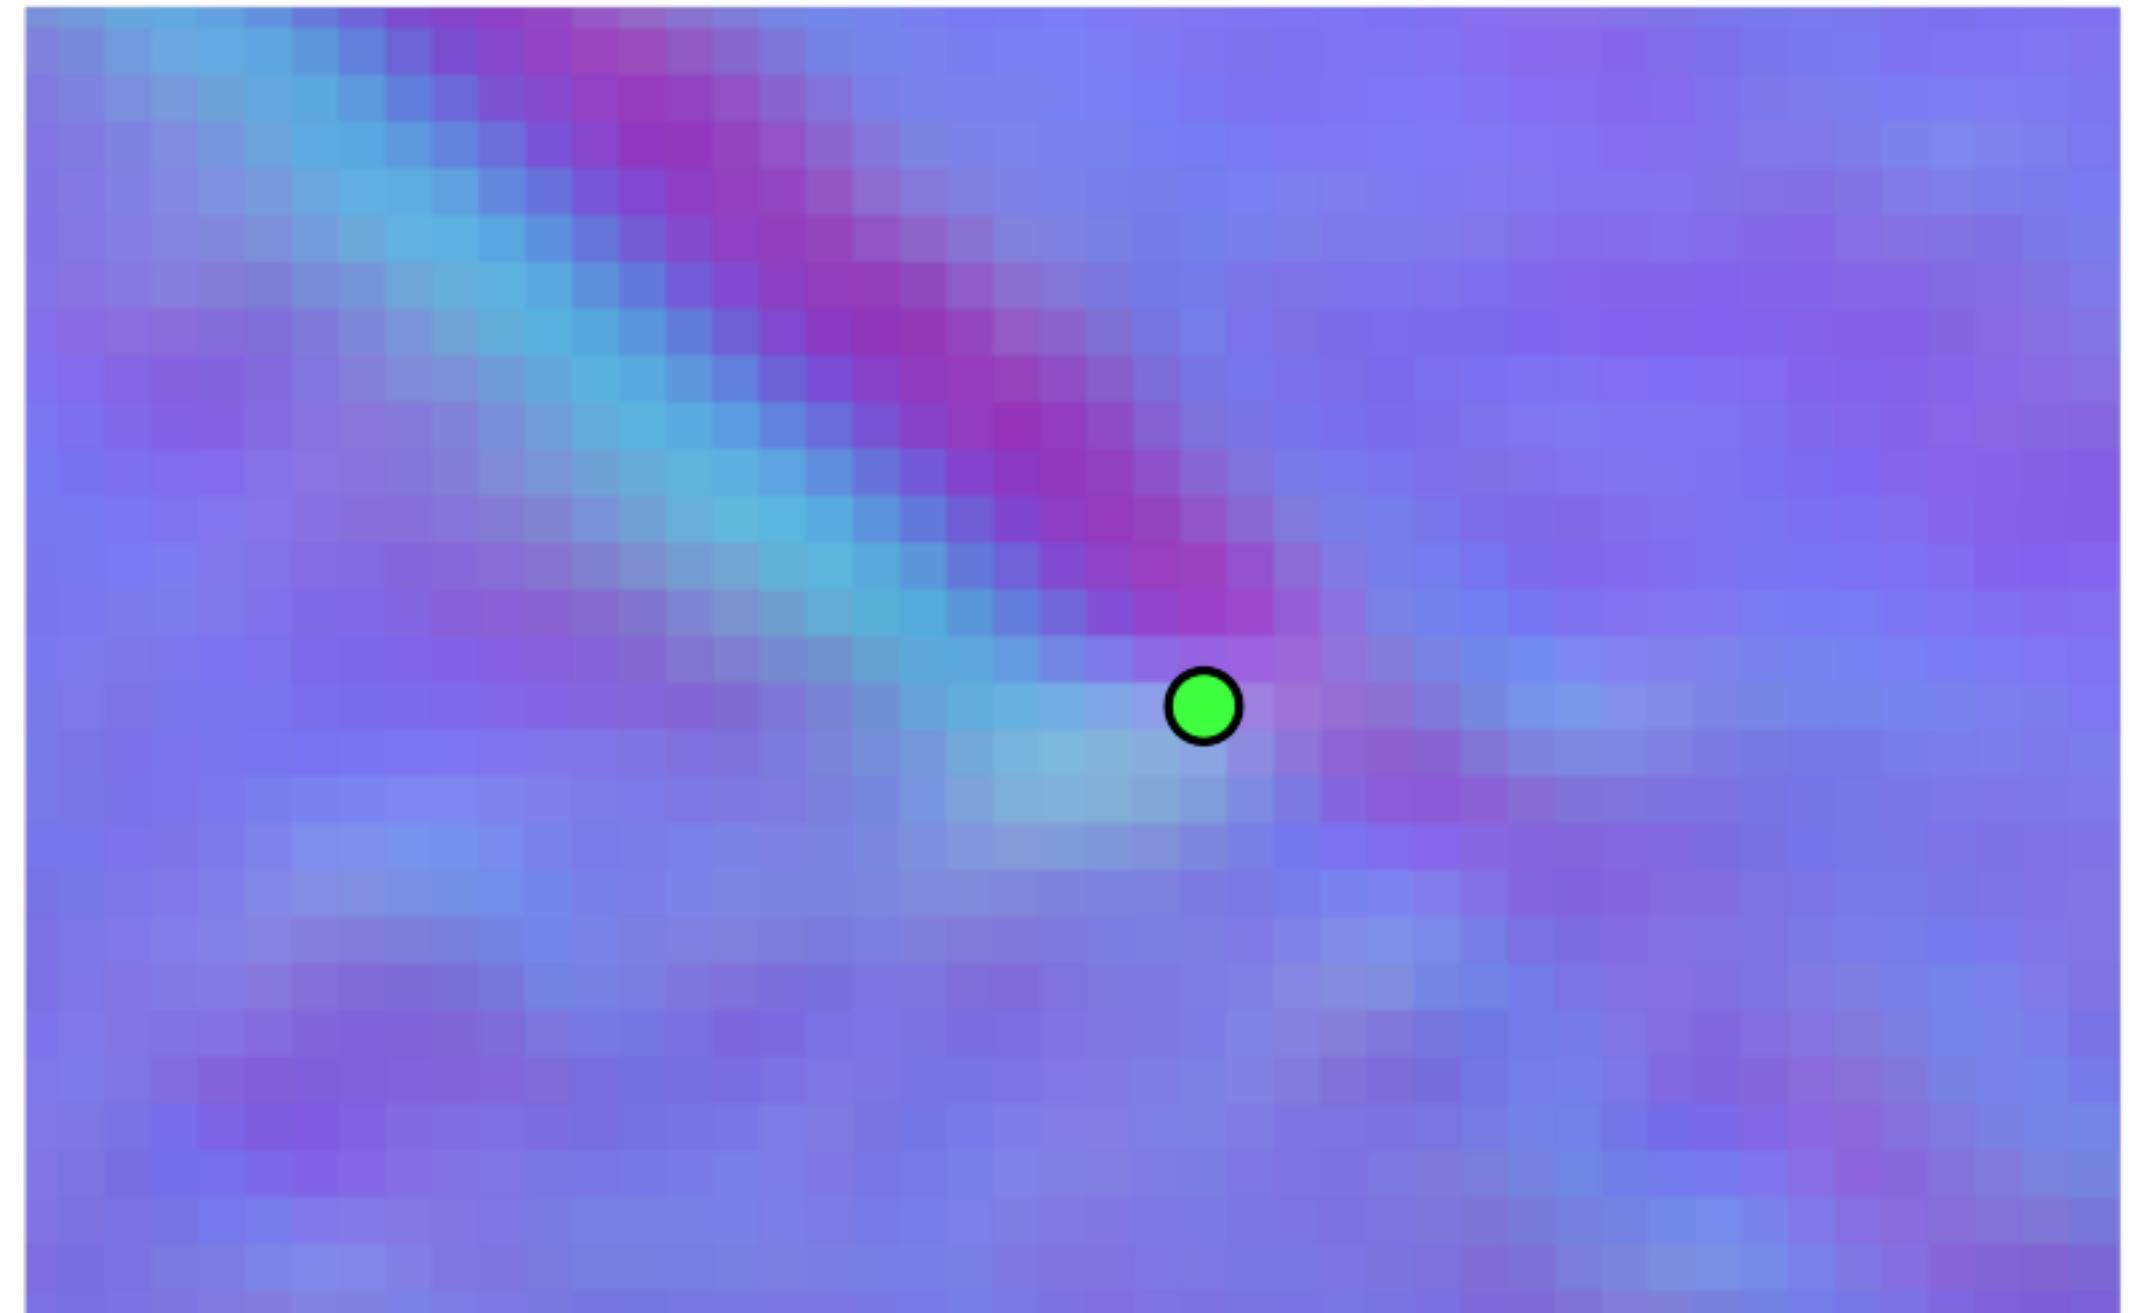
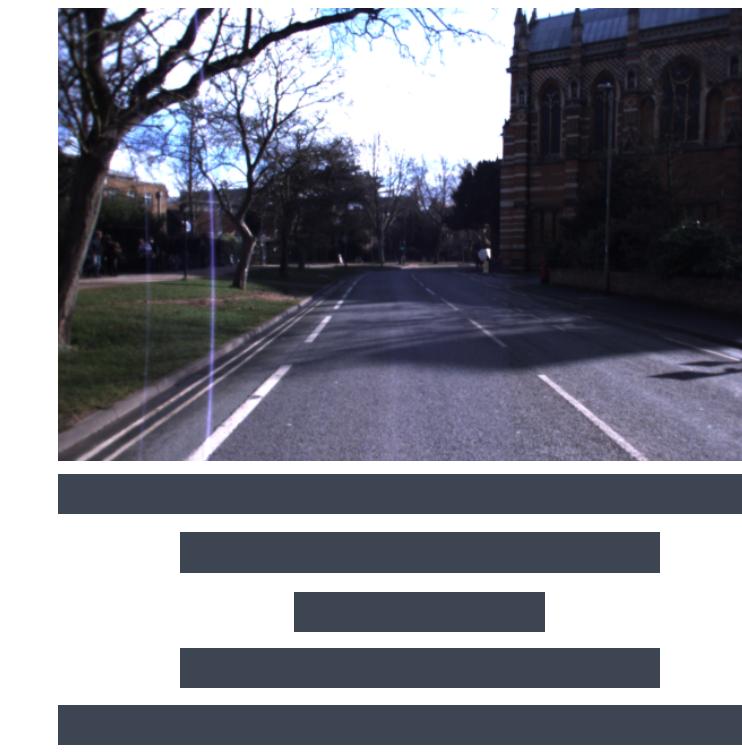
Gauss-Newton Loss



Green: Ground truth correspondence



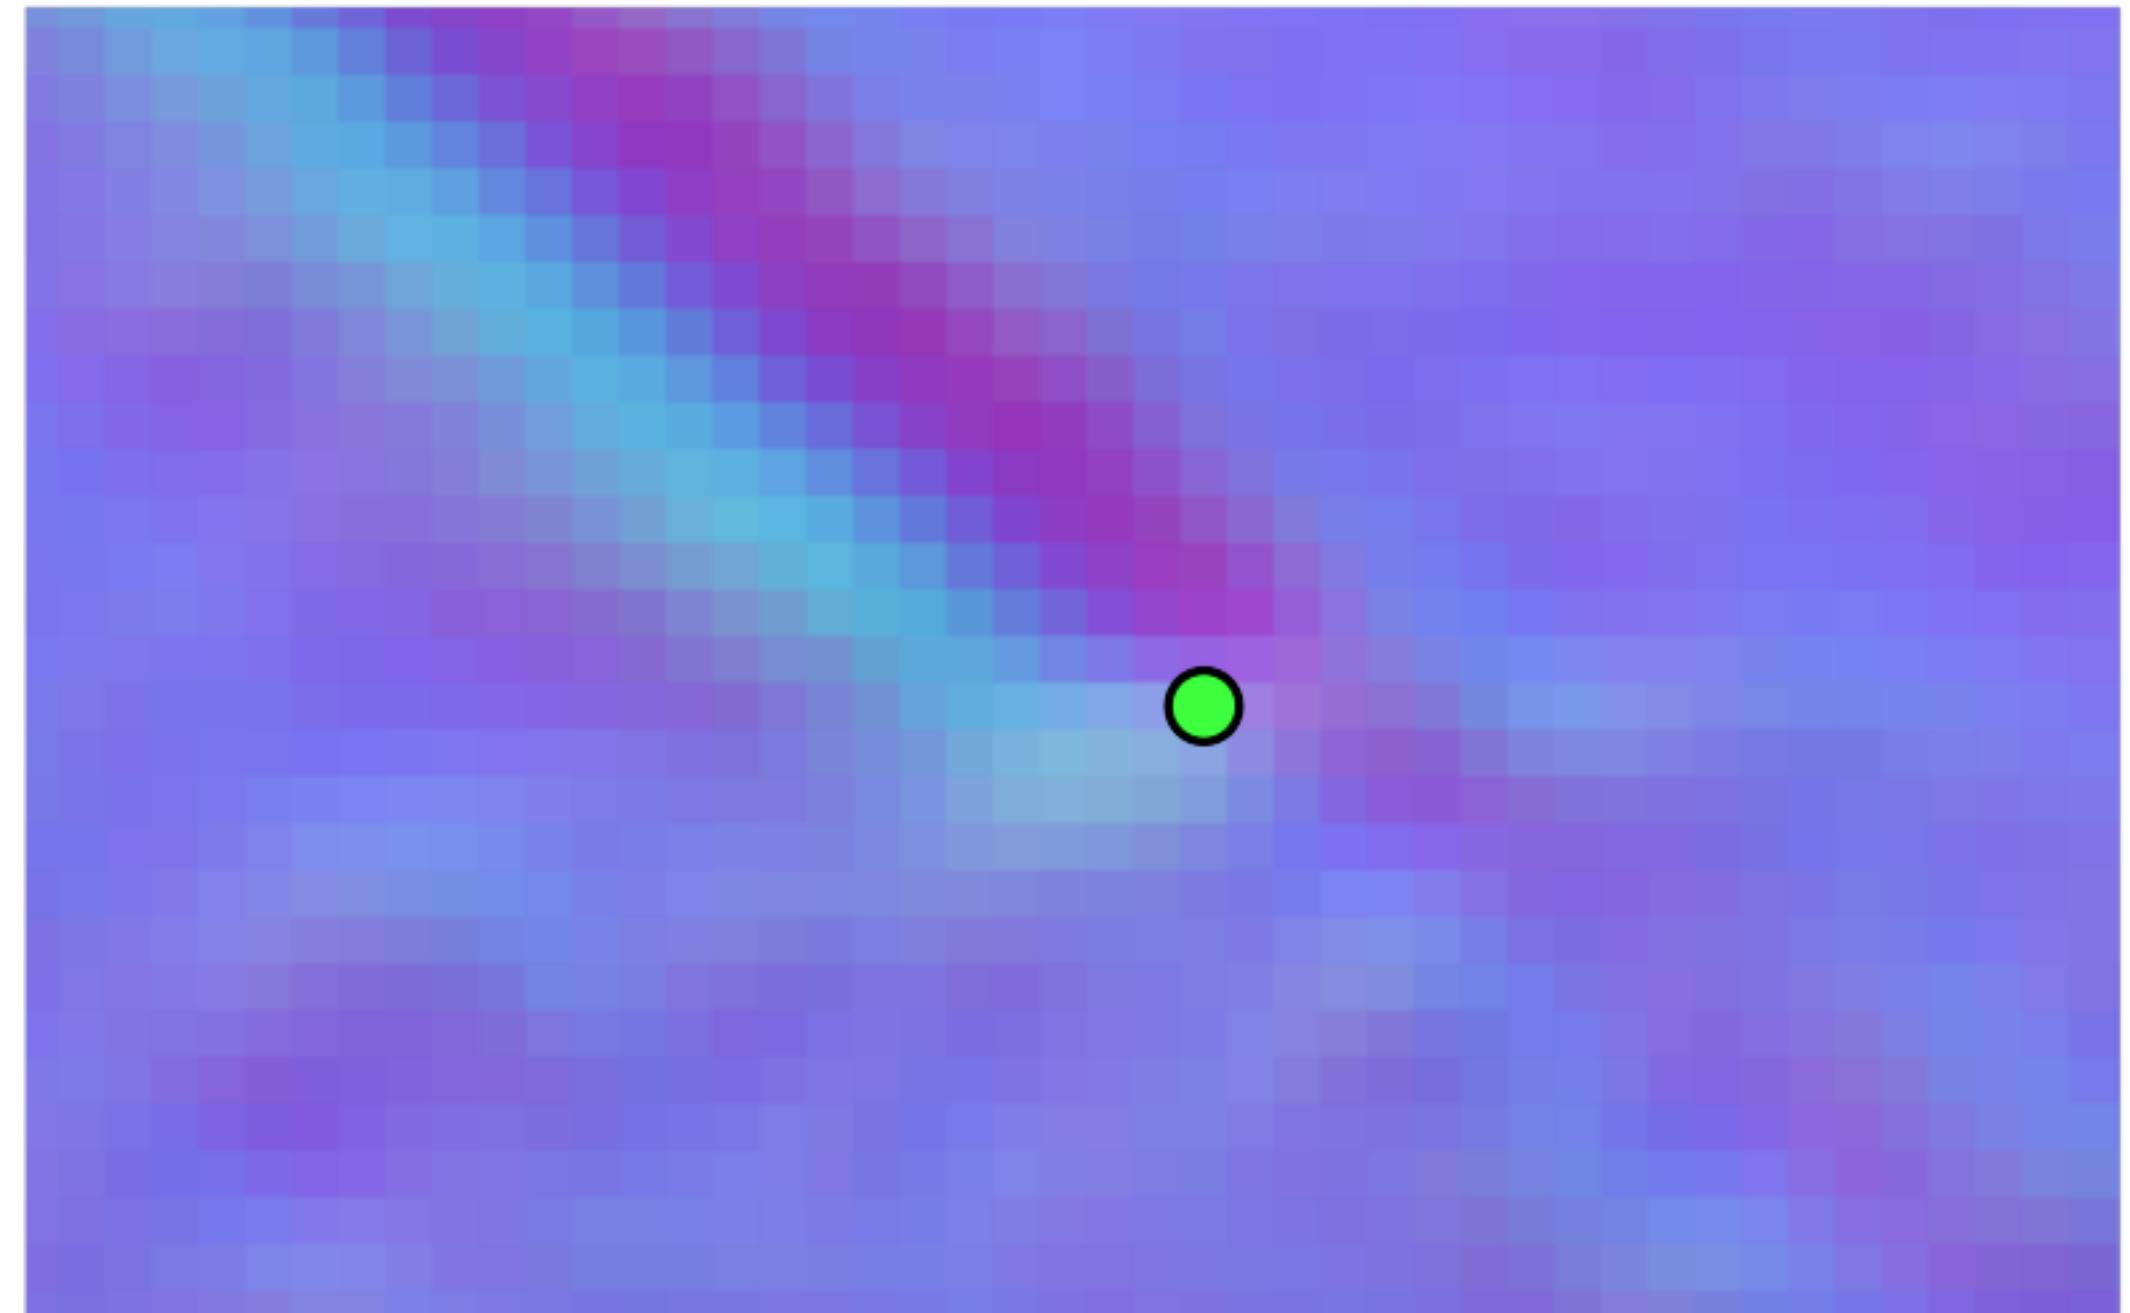
Gauss-Newton Loss



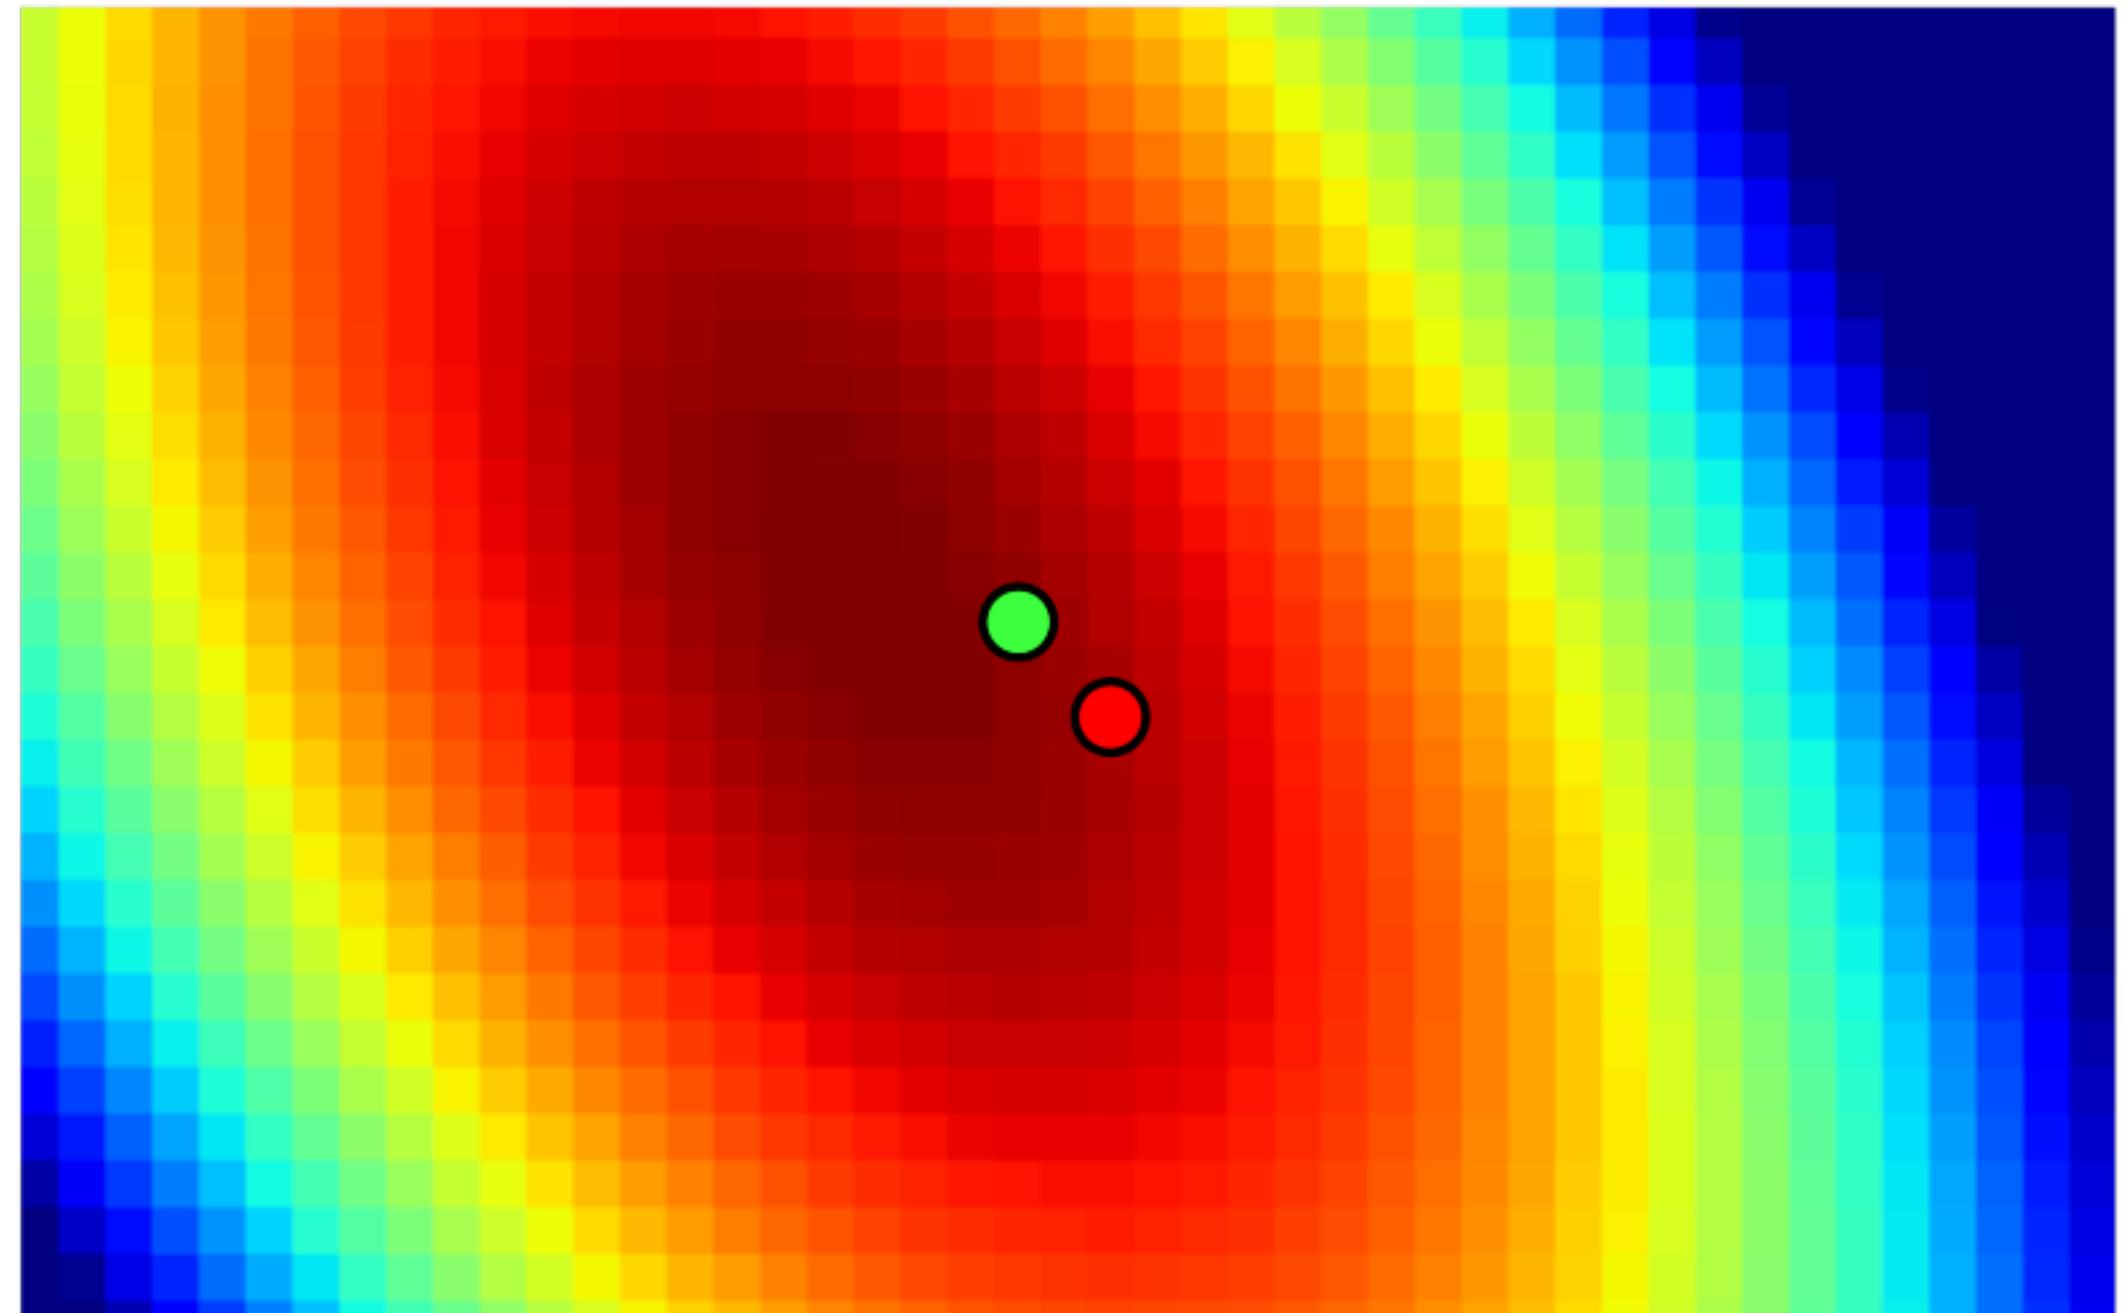
Green: Ground truth correspondence, Red: Initial solution



Gauss-Newton Loss



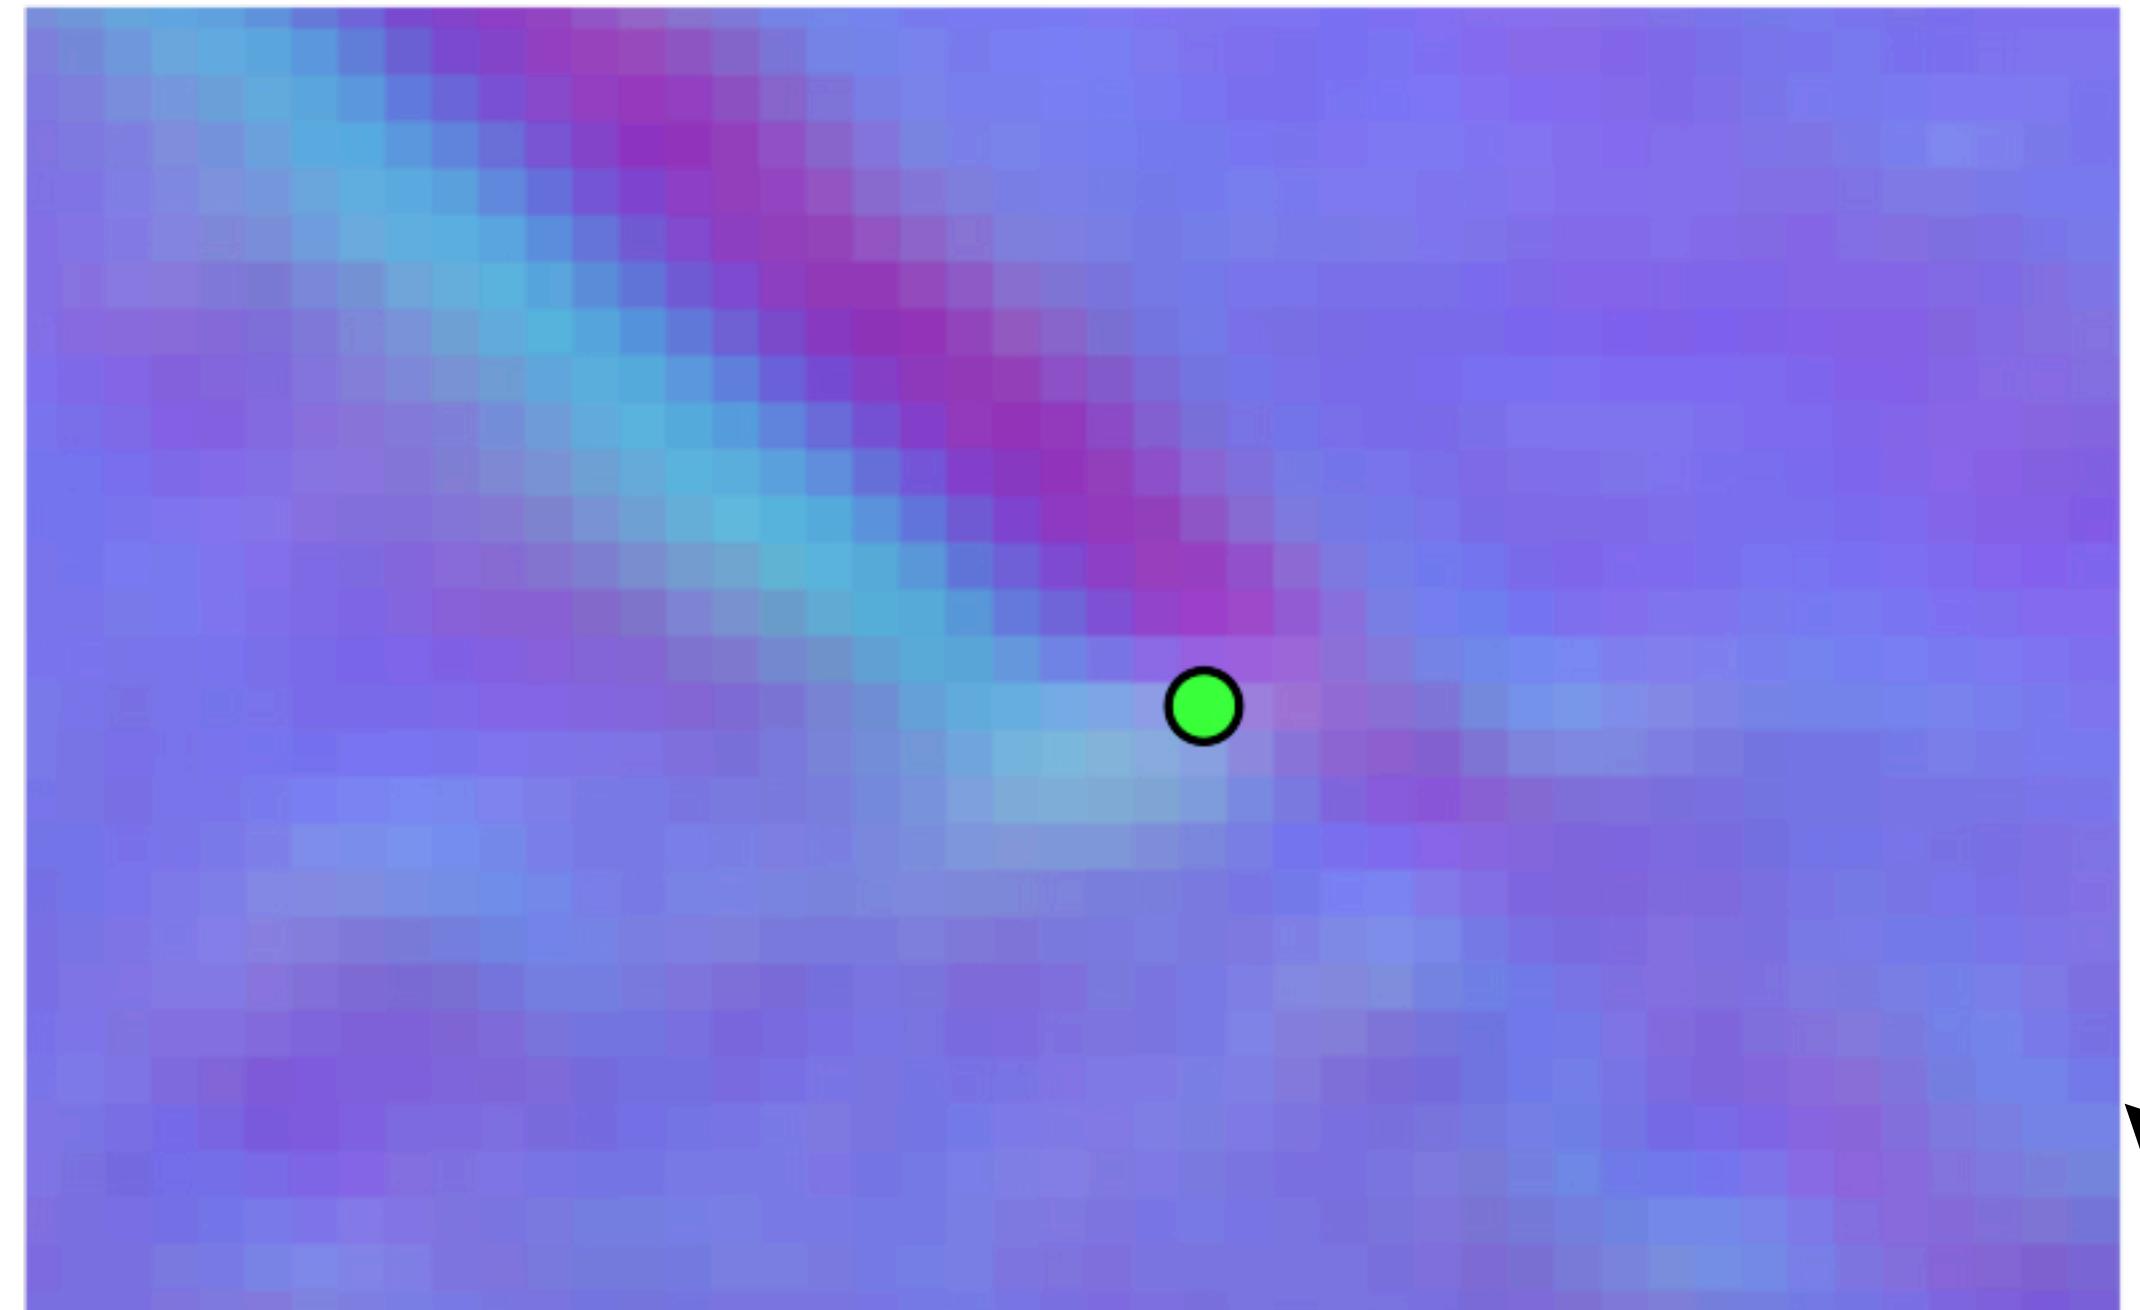
Green: Ground truth correspondence, Red: Initial solution



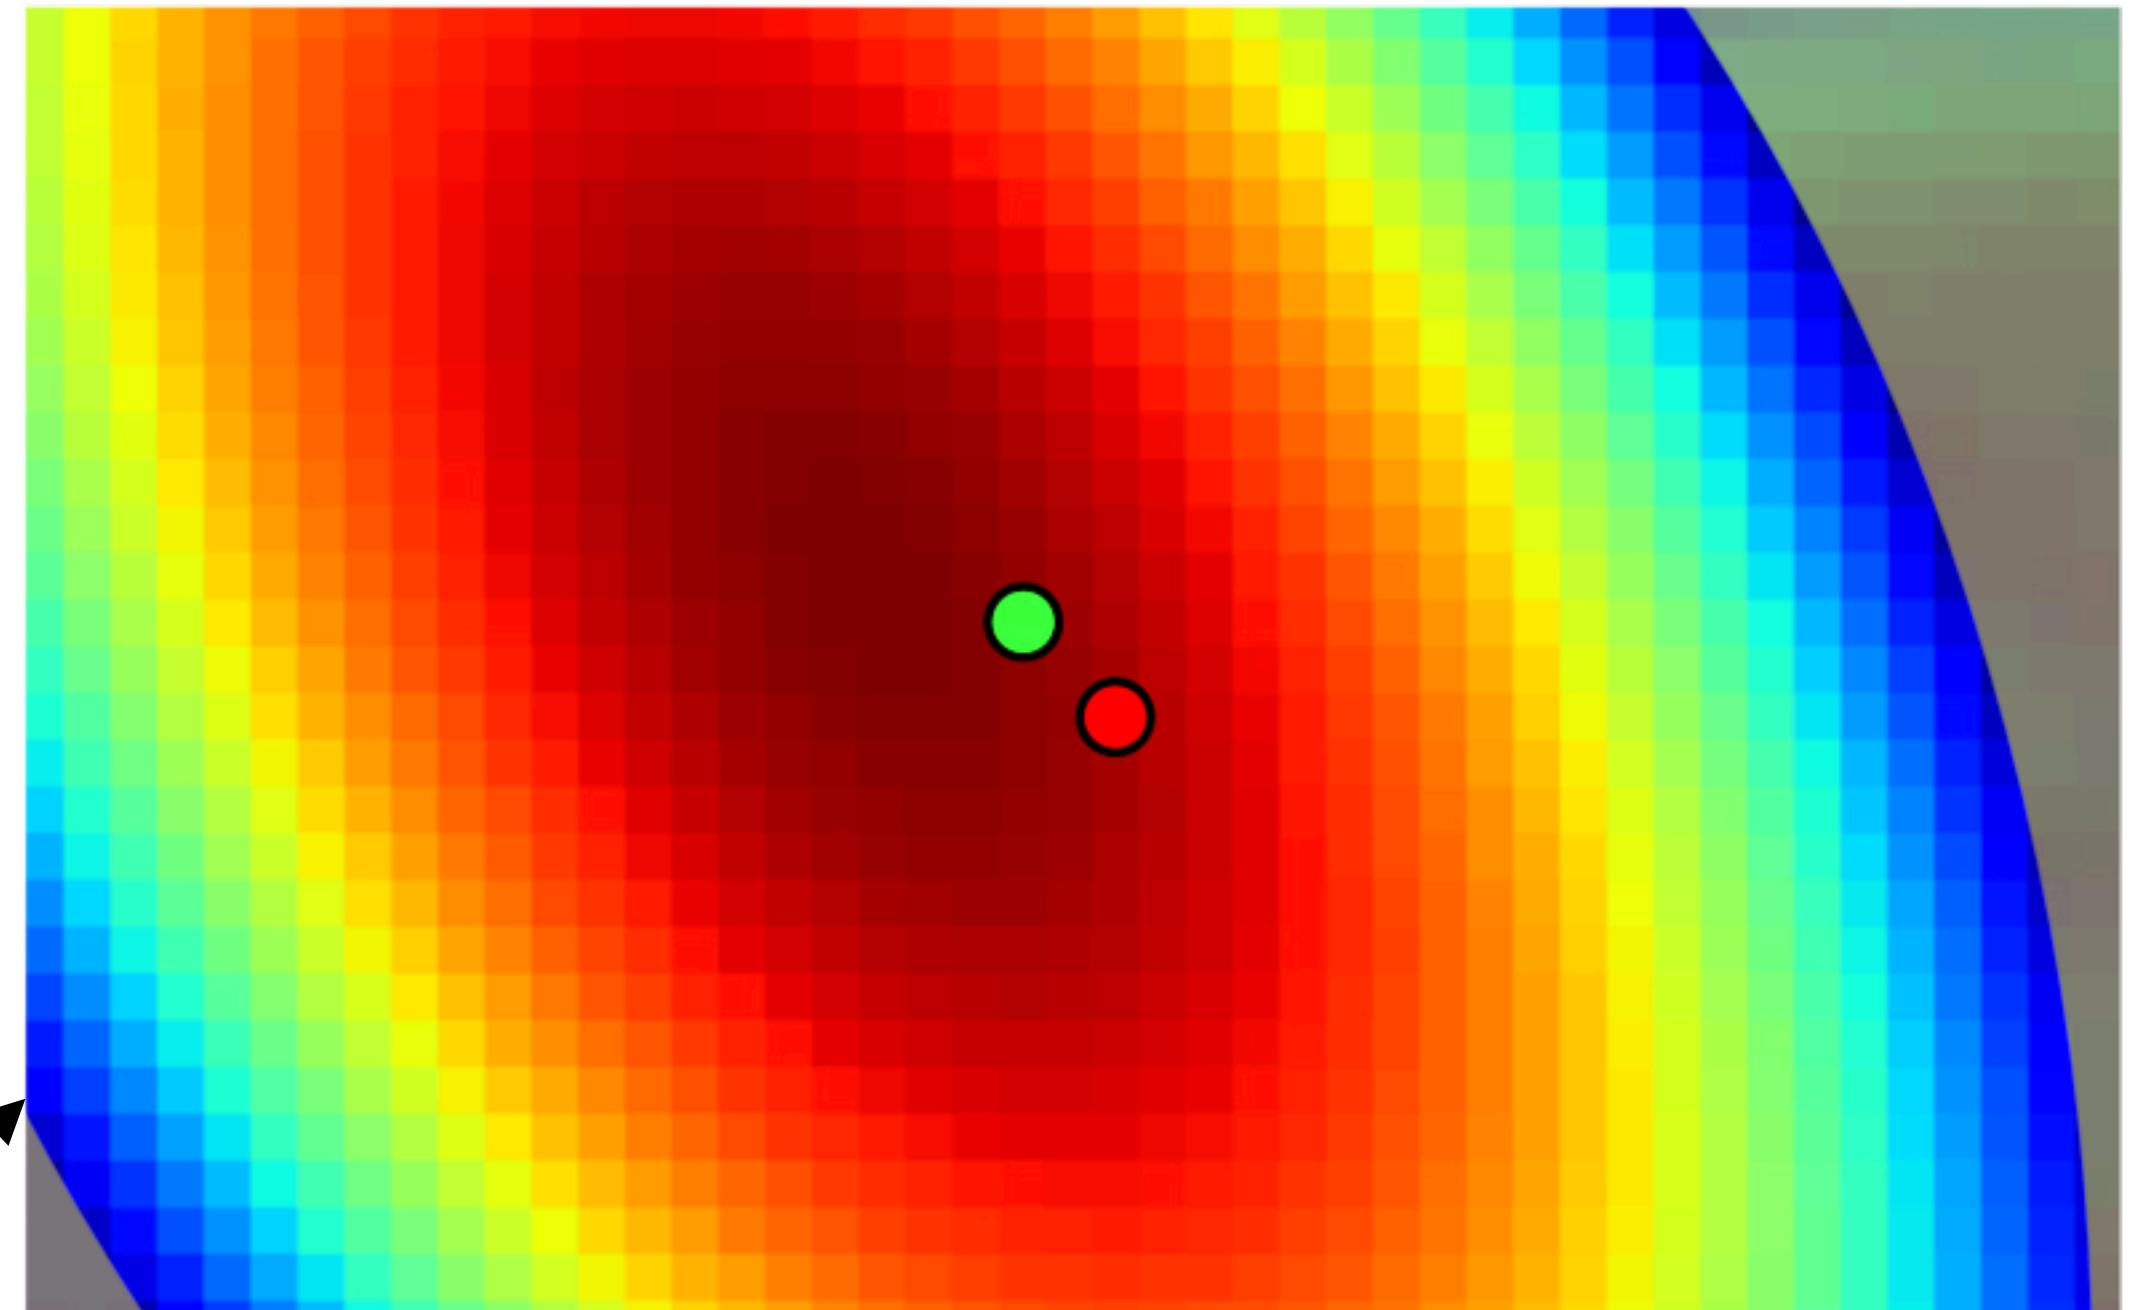
Ellipse: Gaussian distribution



Gauss-Newton Loss

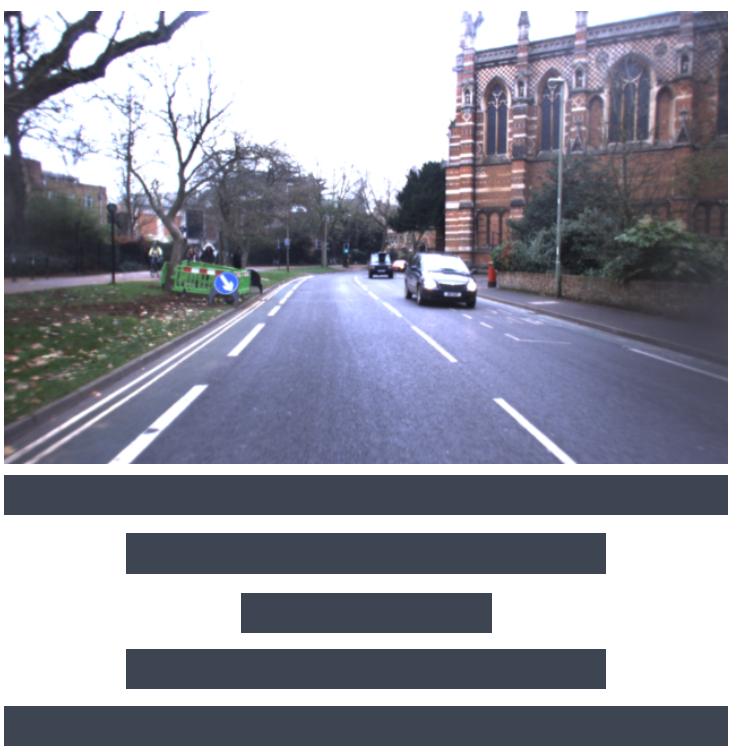


R, t



Green: Ground truth correspondence, Red: Initial solution

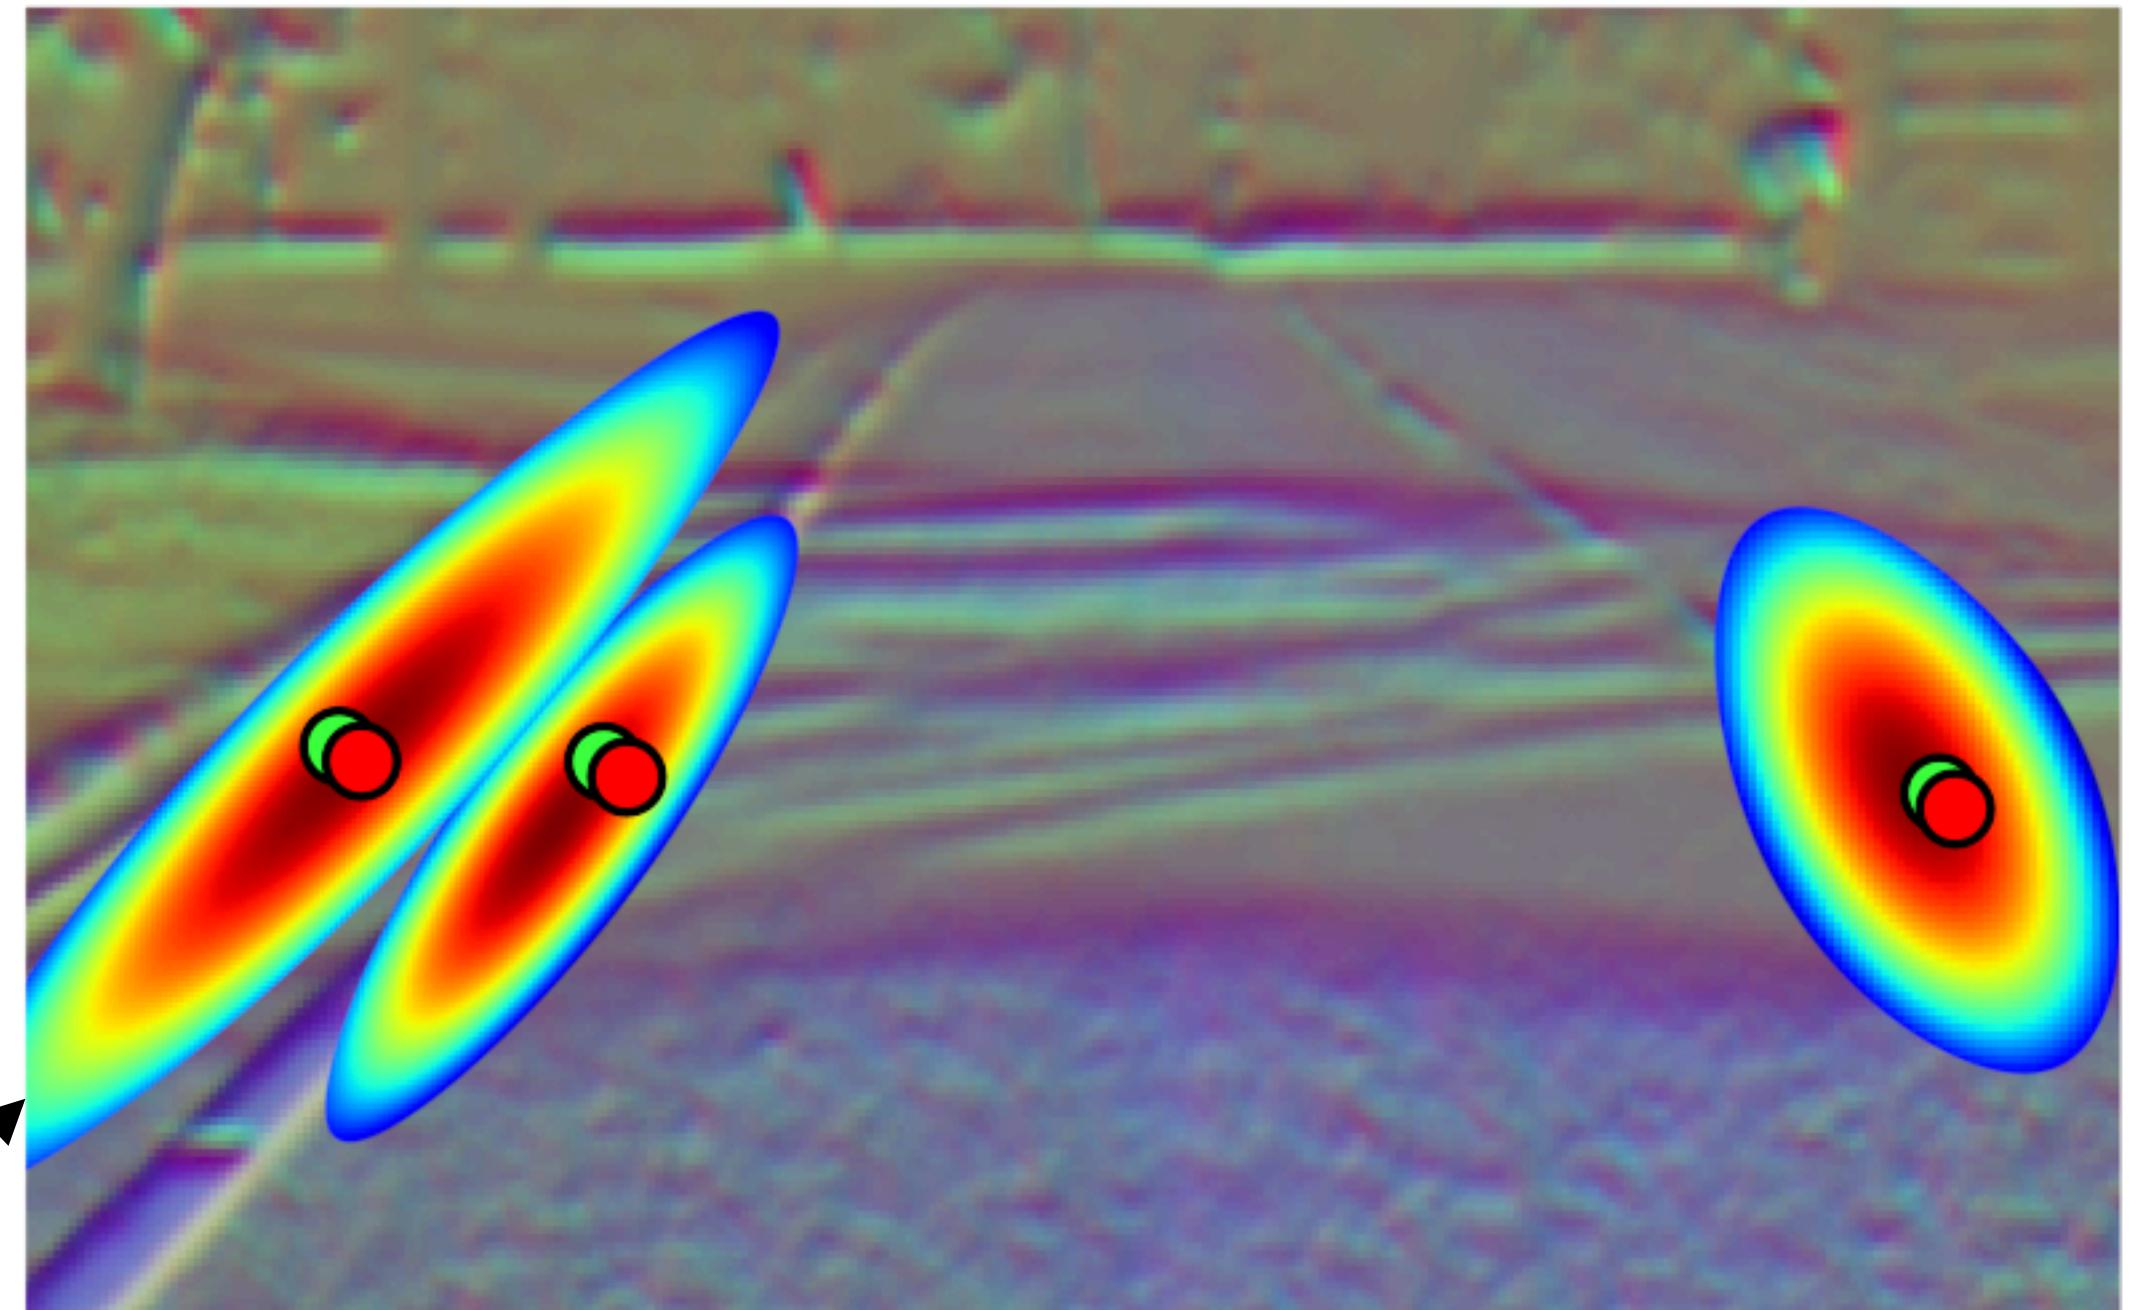
Ellipse: Gaussian distribution



Gauss-Newton Loss



R, t



Green: Ground truth correspondence, Red: Initial solution

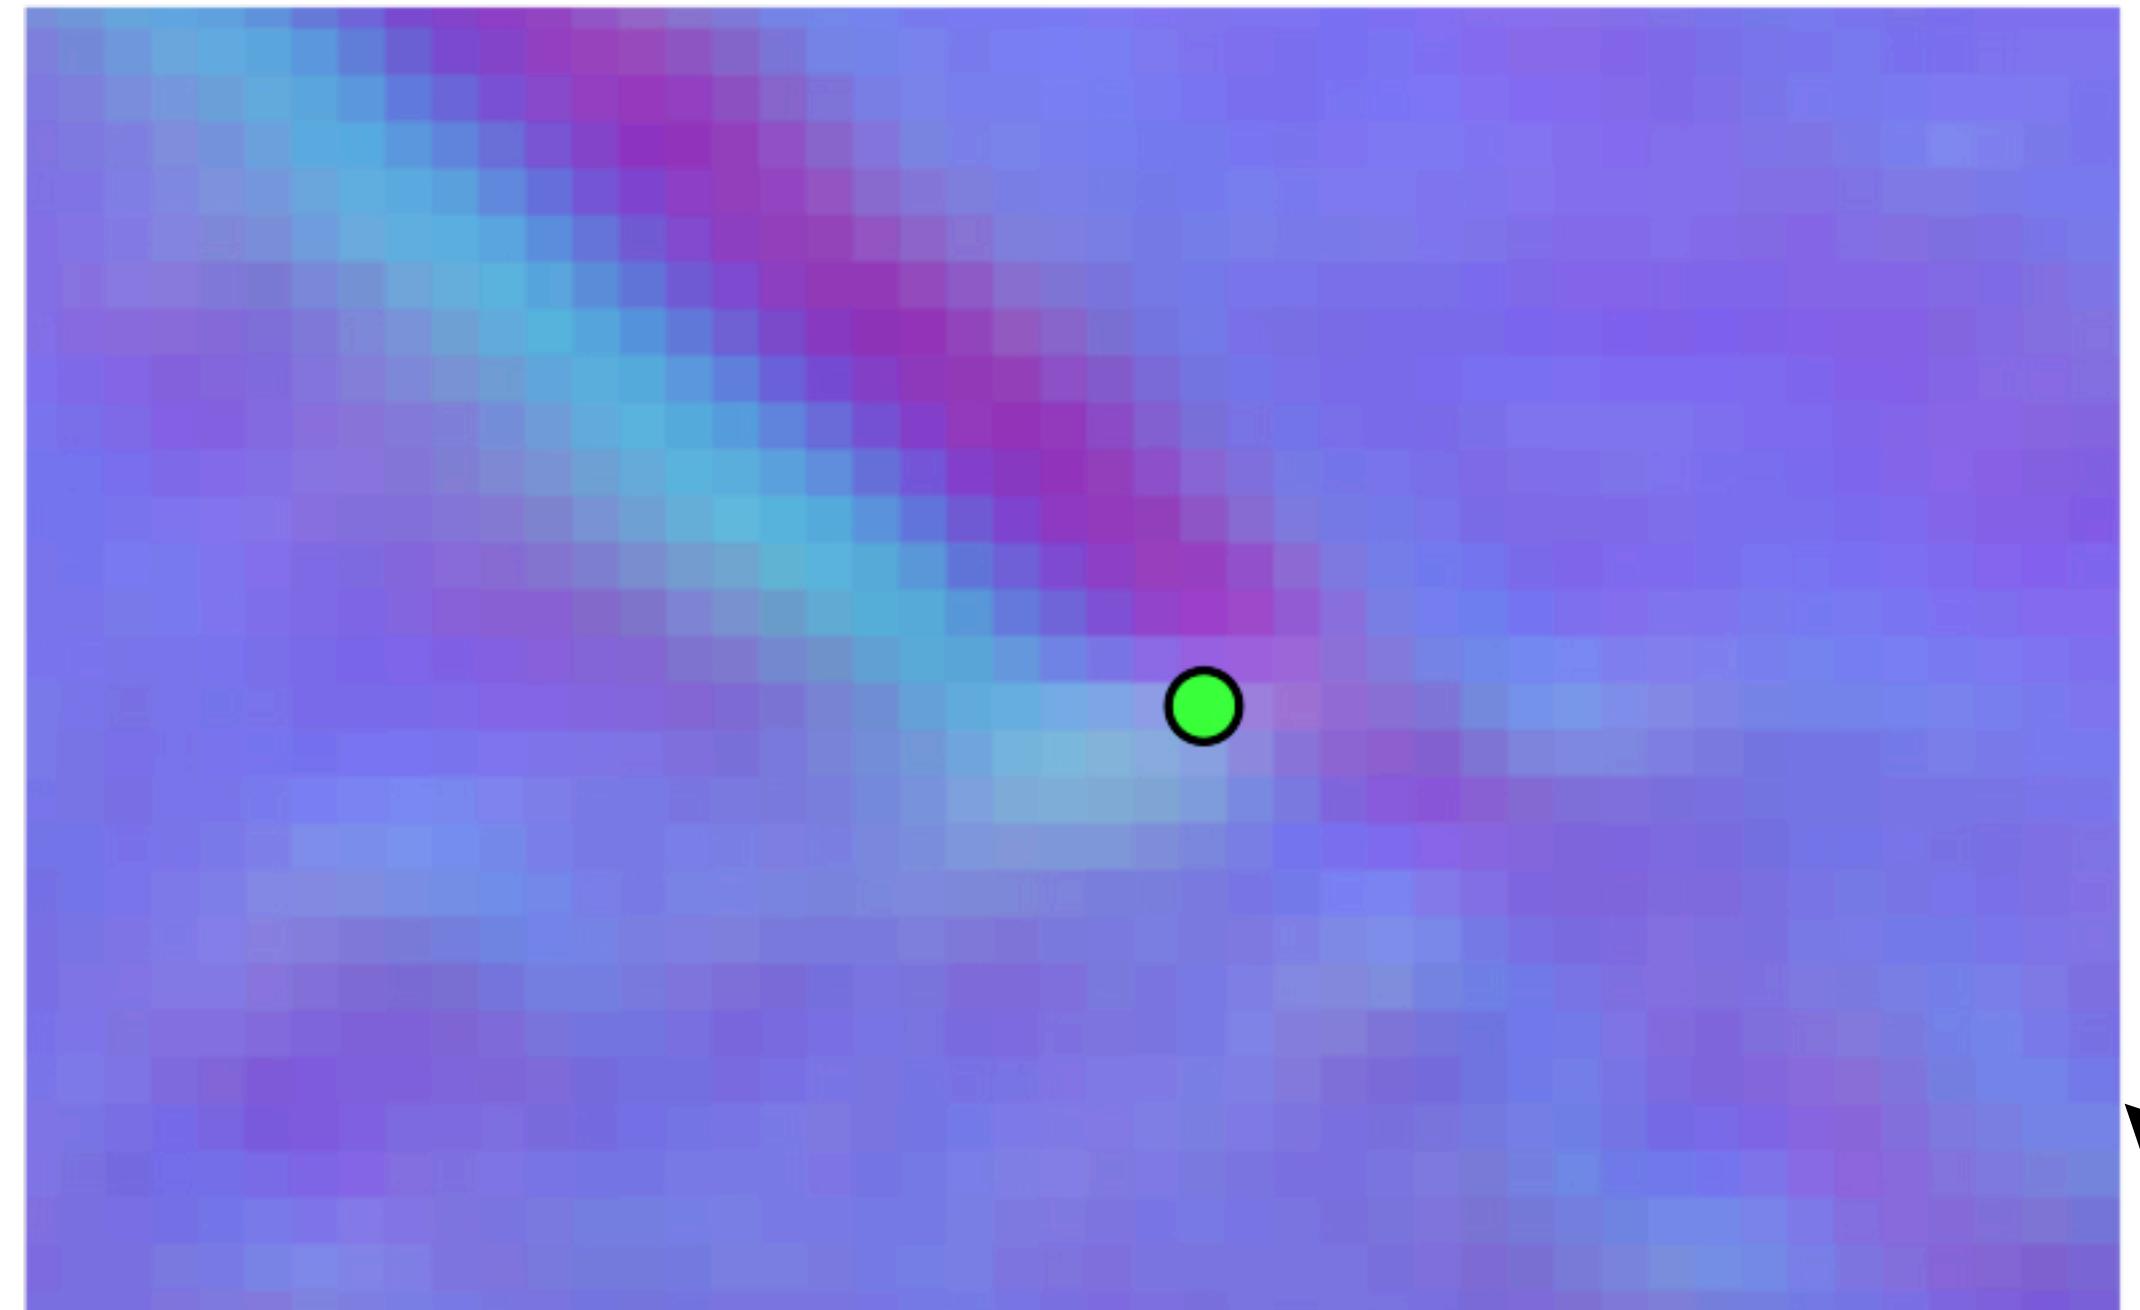
Ellipse: Gaussian distribution

Gauss-Newton Loss:

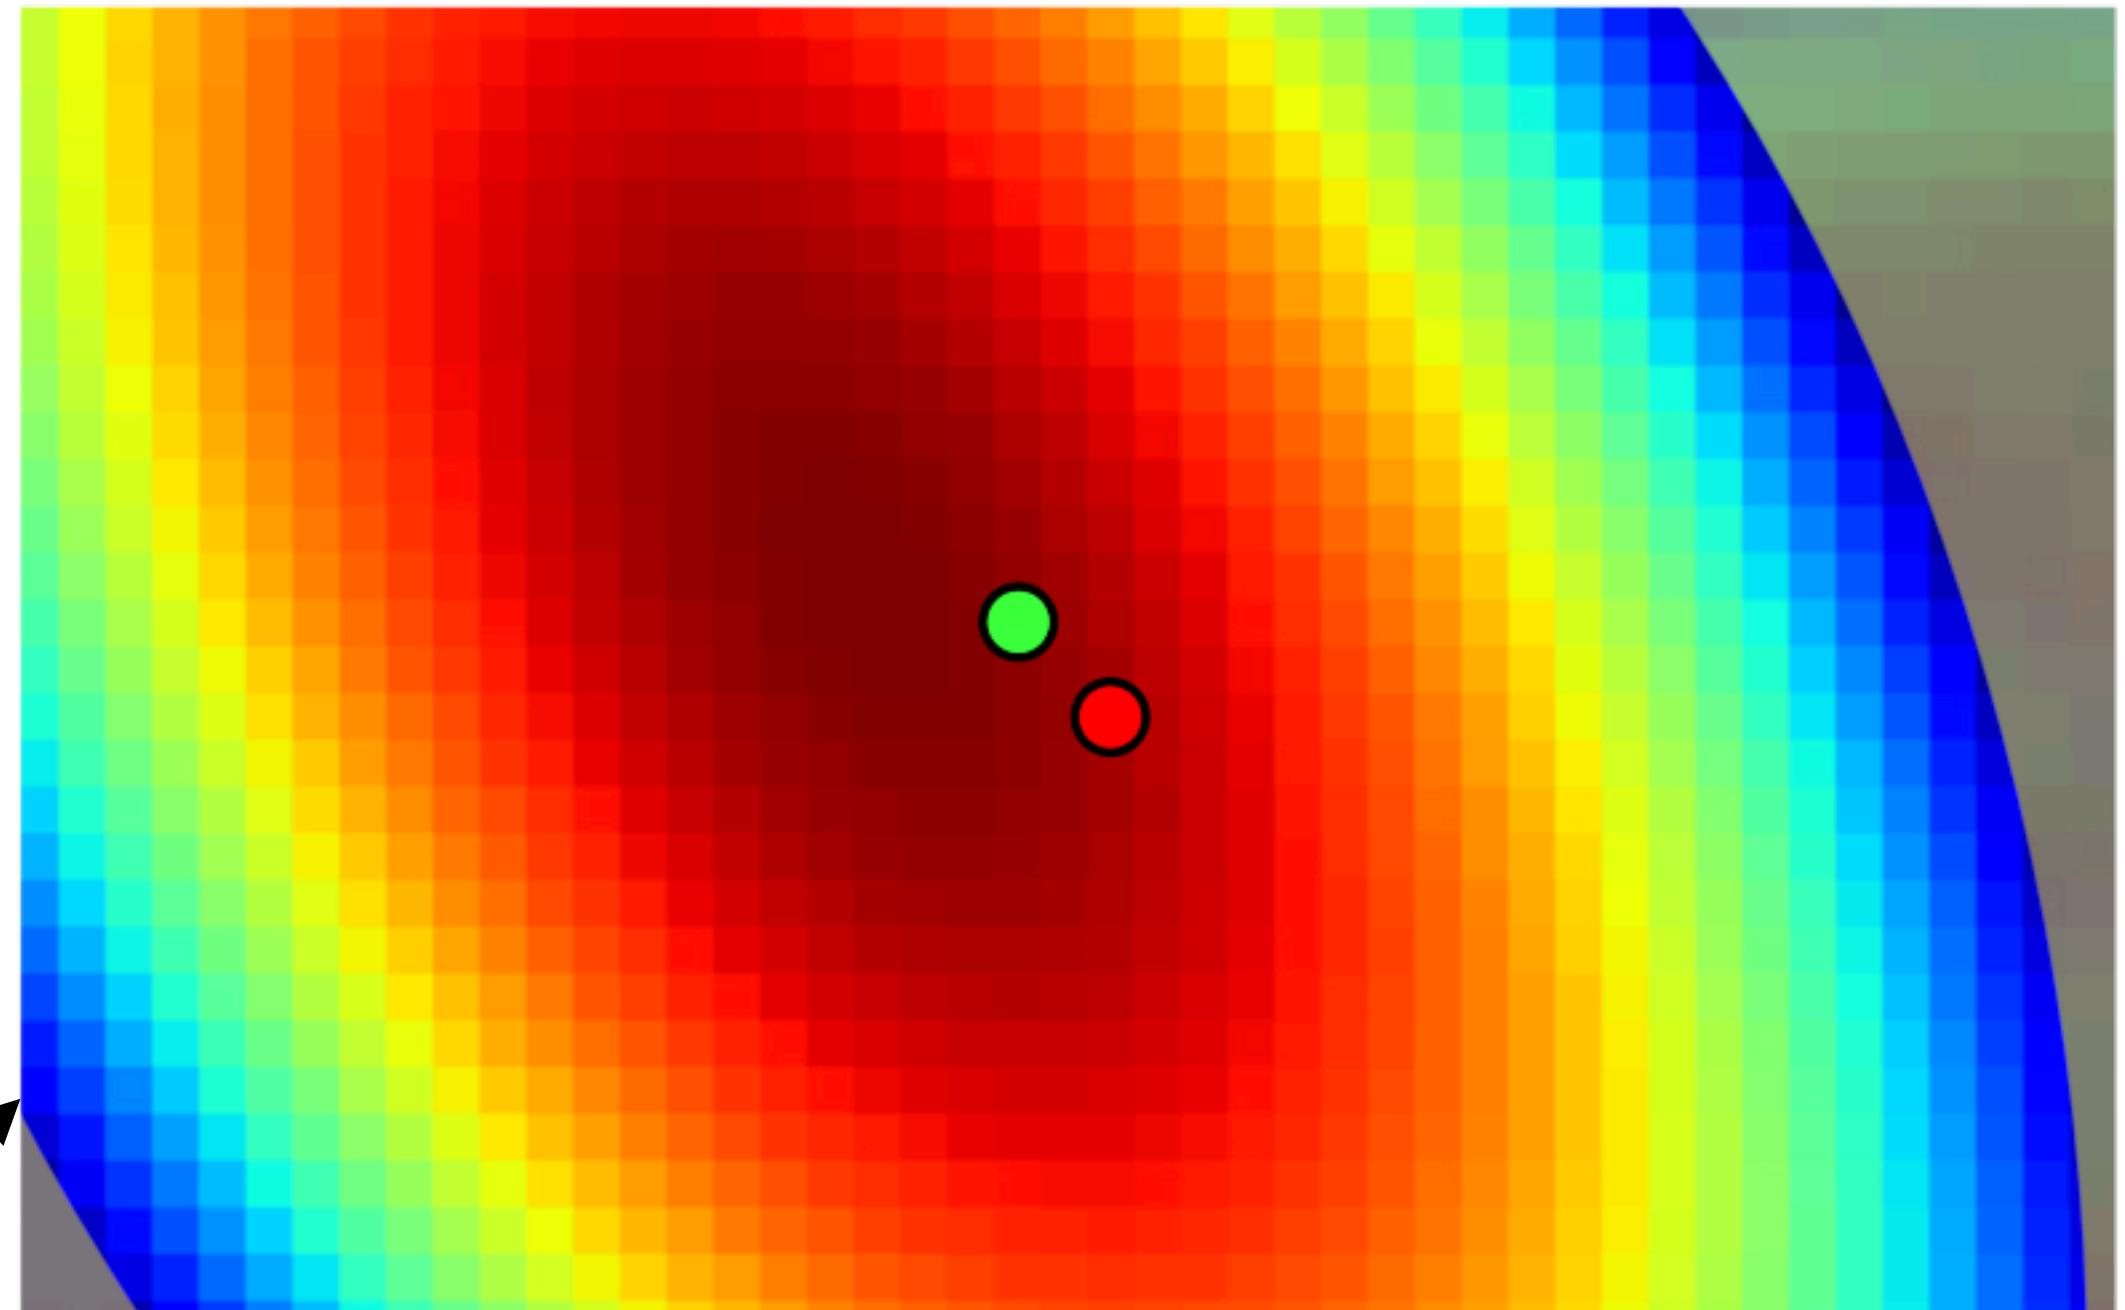
Maximize estimated probability
of ground truth correspondence



Gauss-Newton Loss



Green: Ground truth correspondence, Red: Initial solution



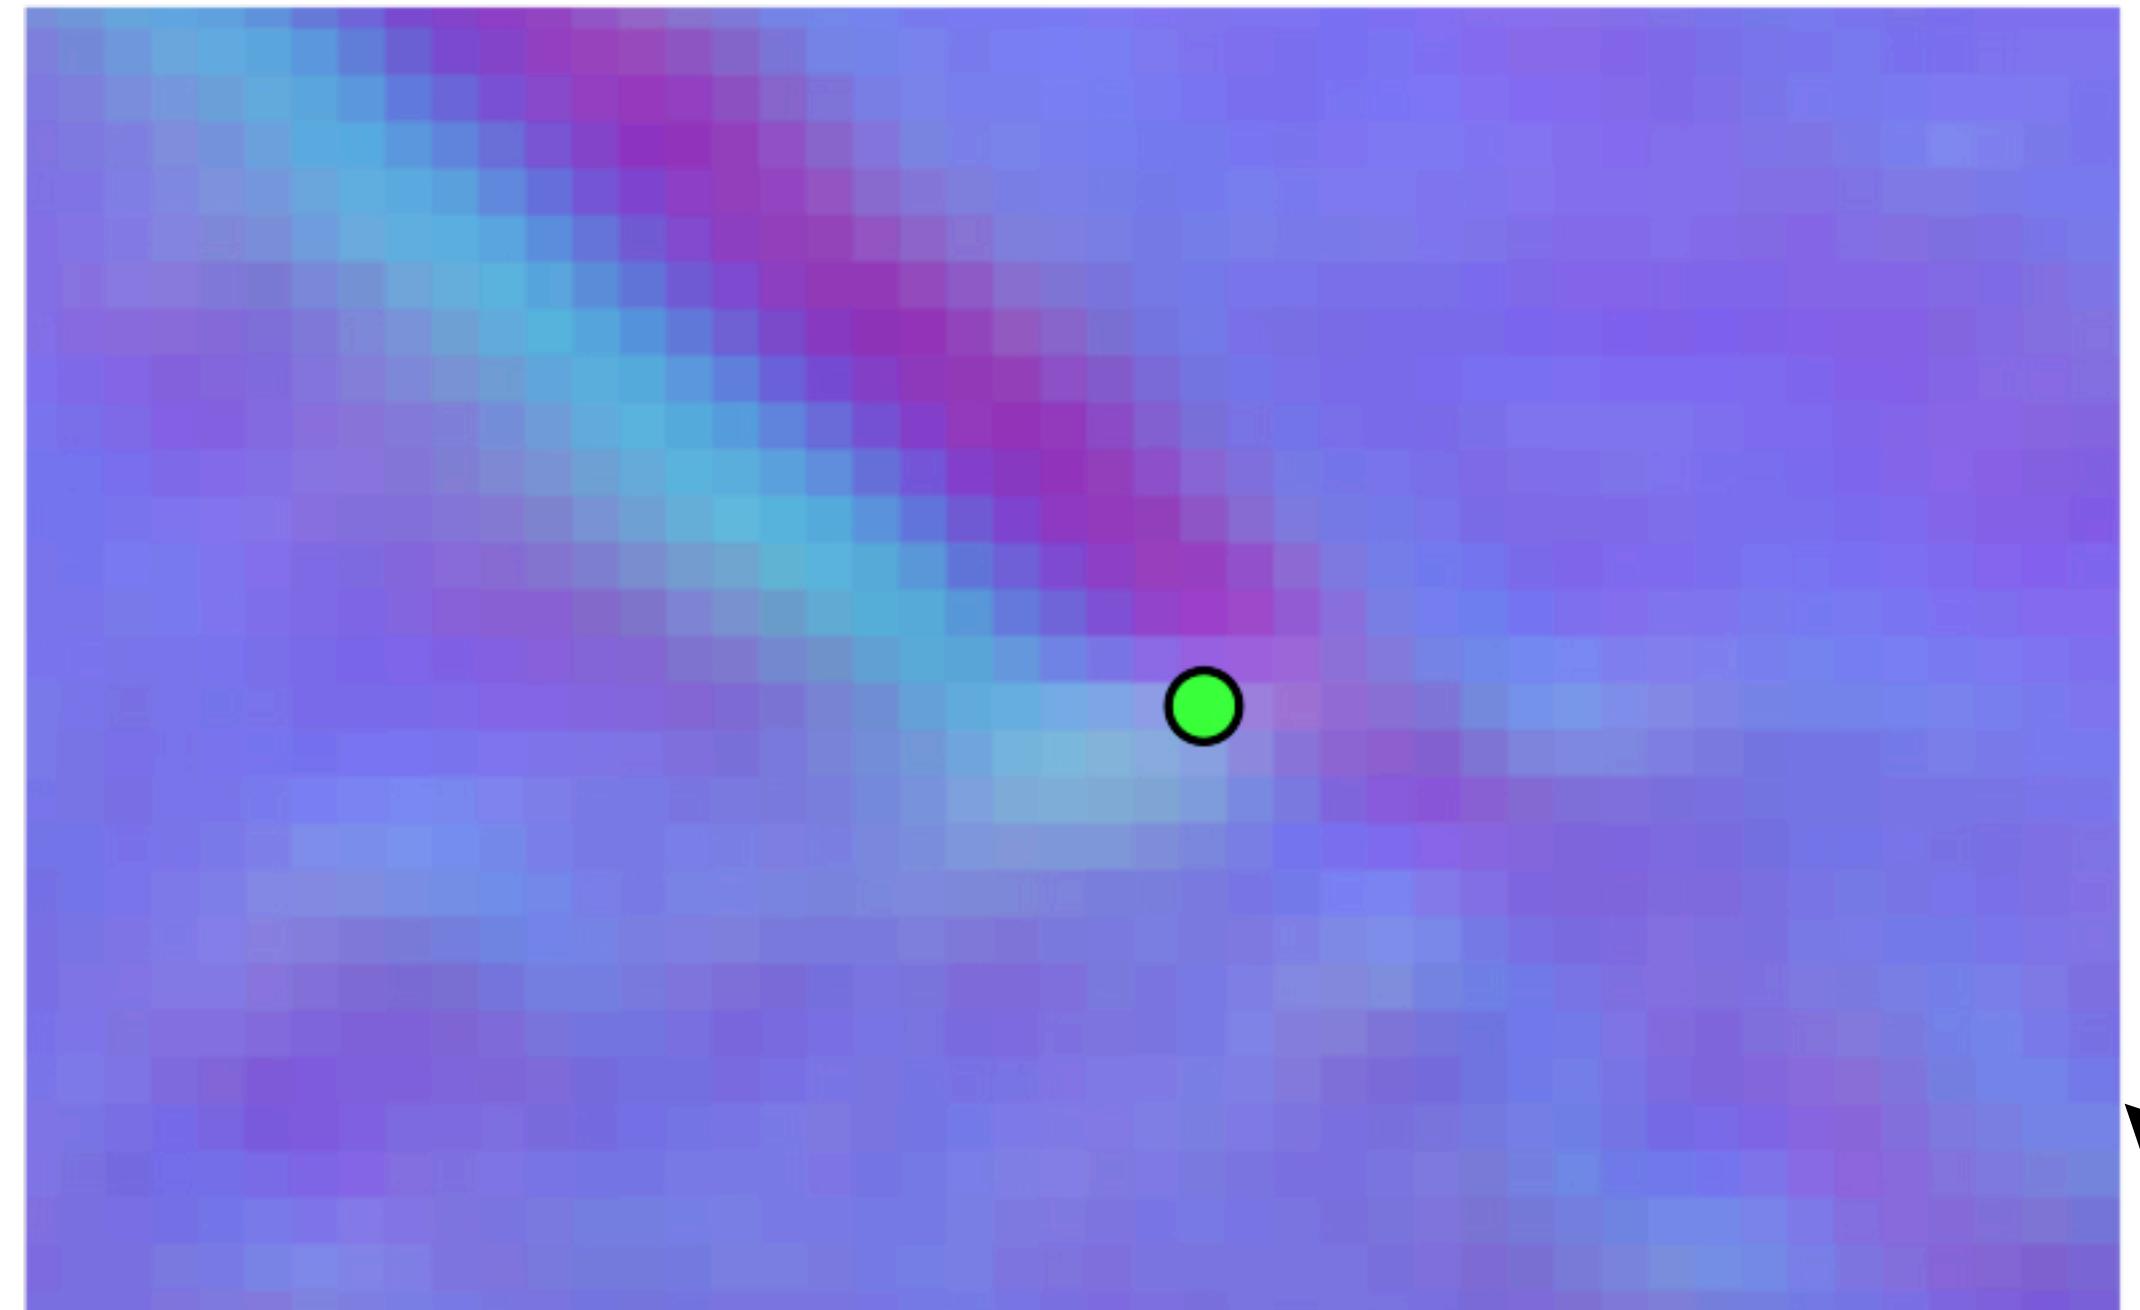
Ellipse: Gaussian distribution

Gauss-Newton Loss:

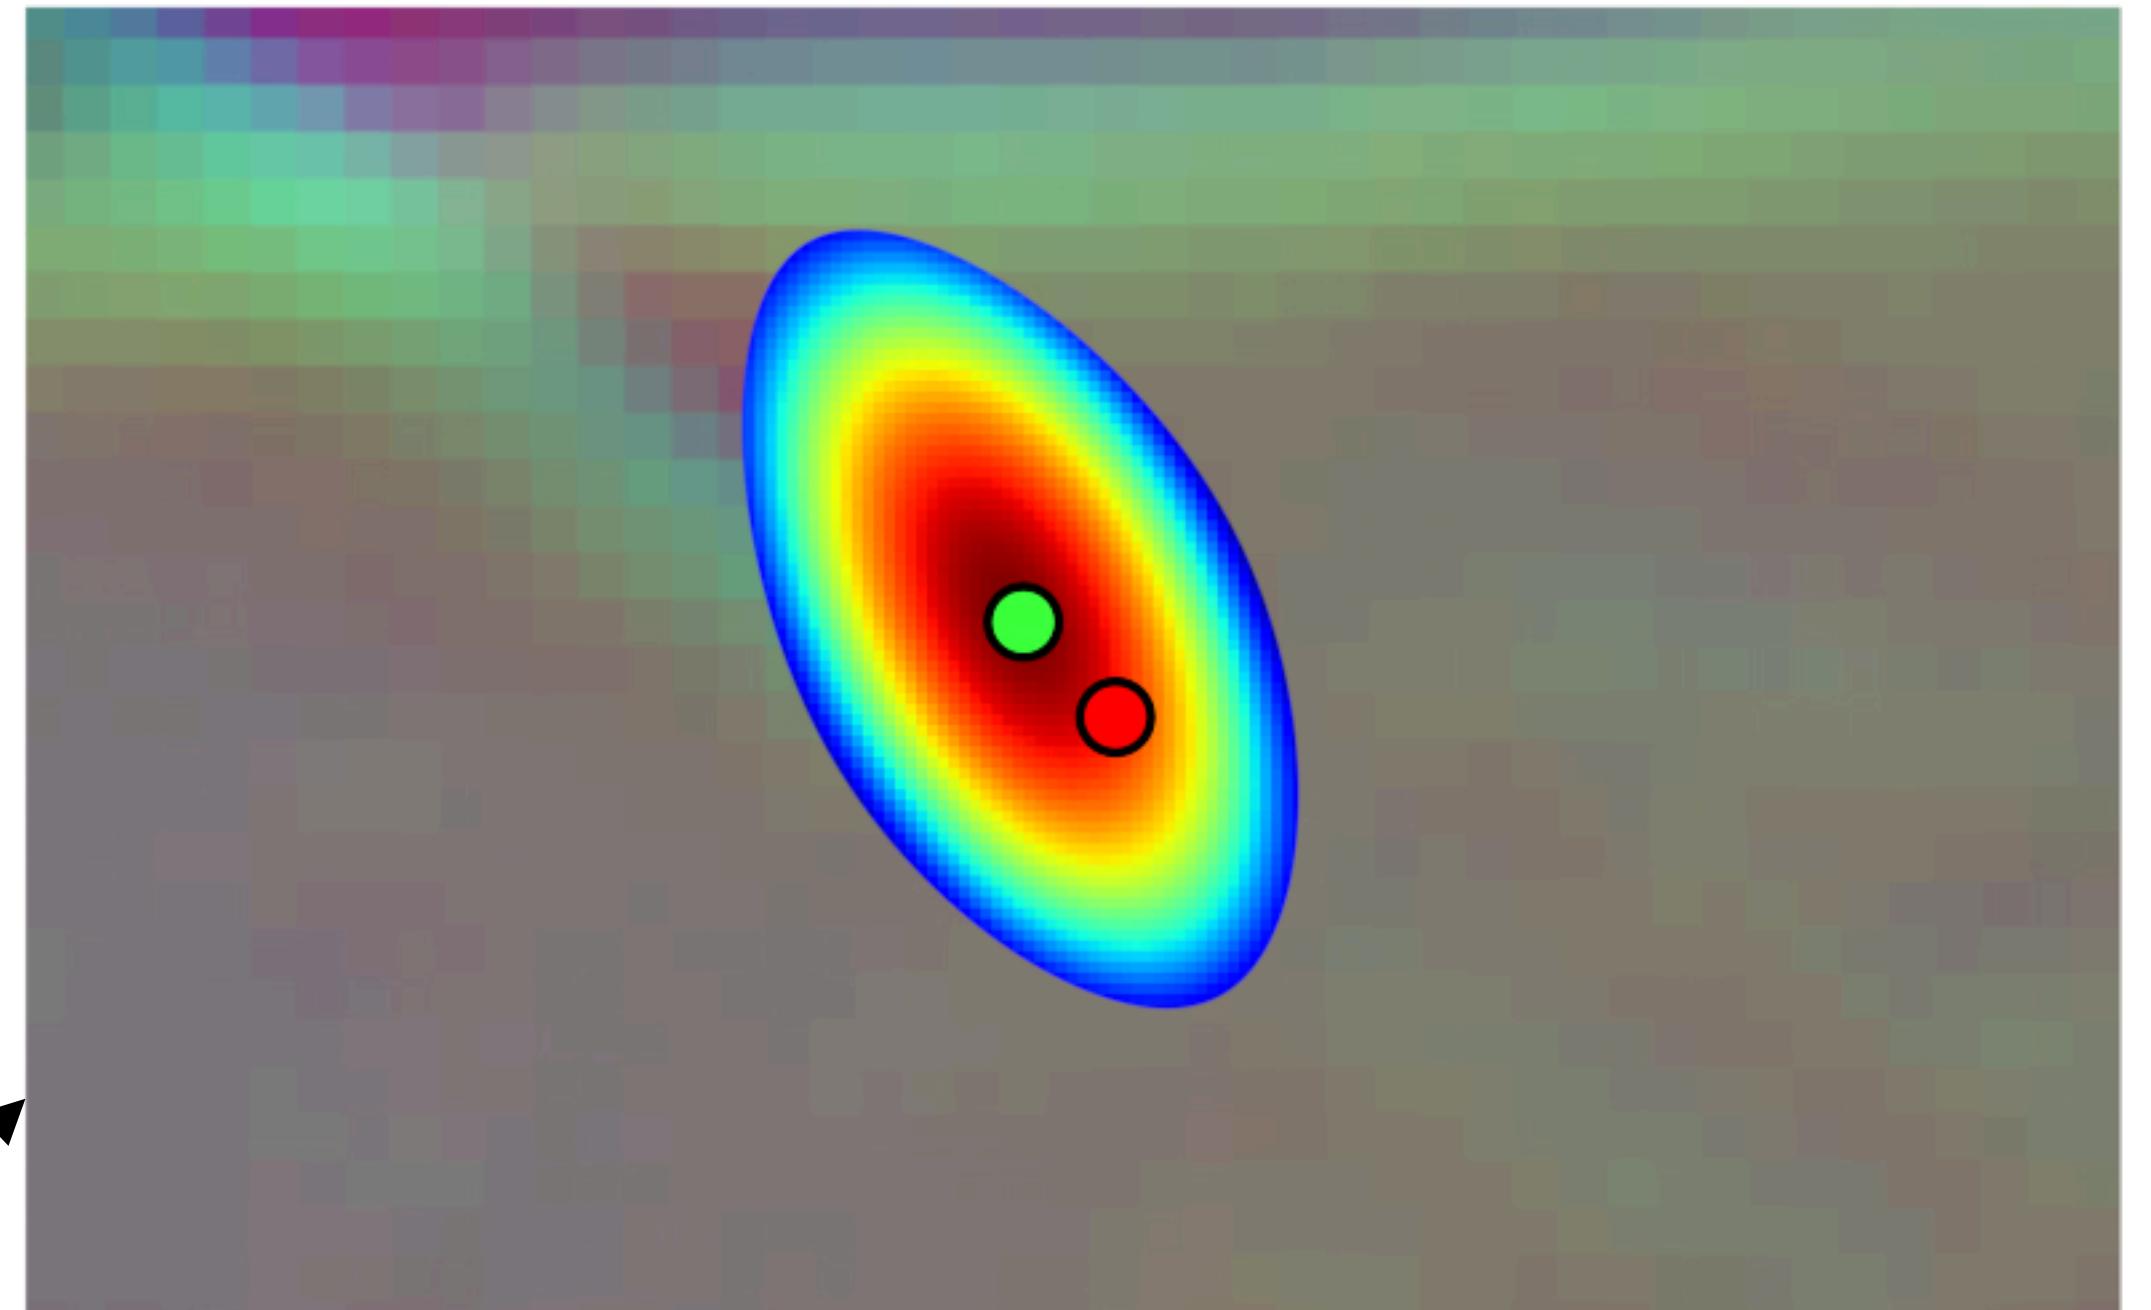
Maximize estimated probability
of ground truth correspondence



Gauss-Newton Loss



R, t



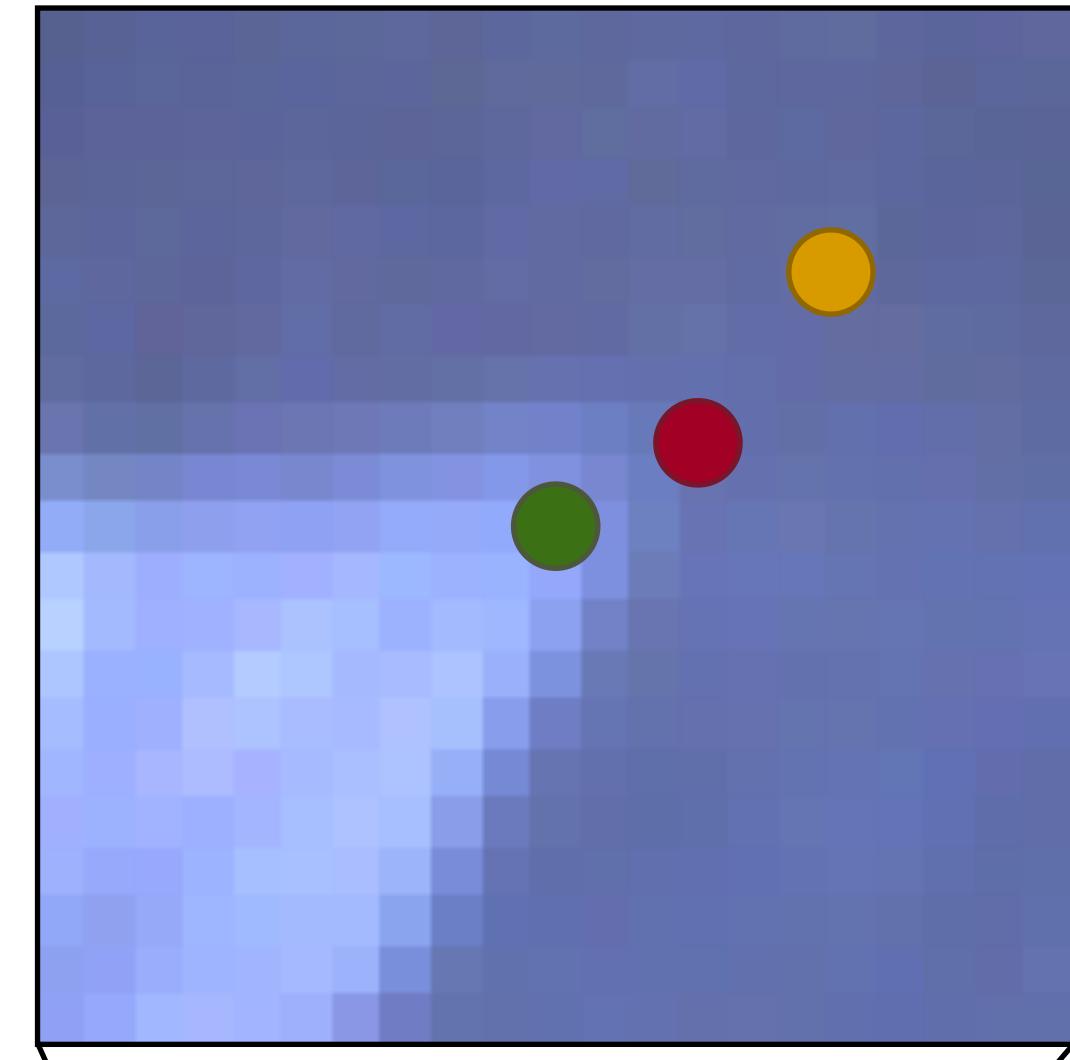
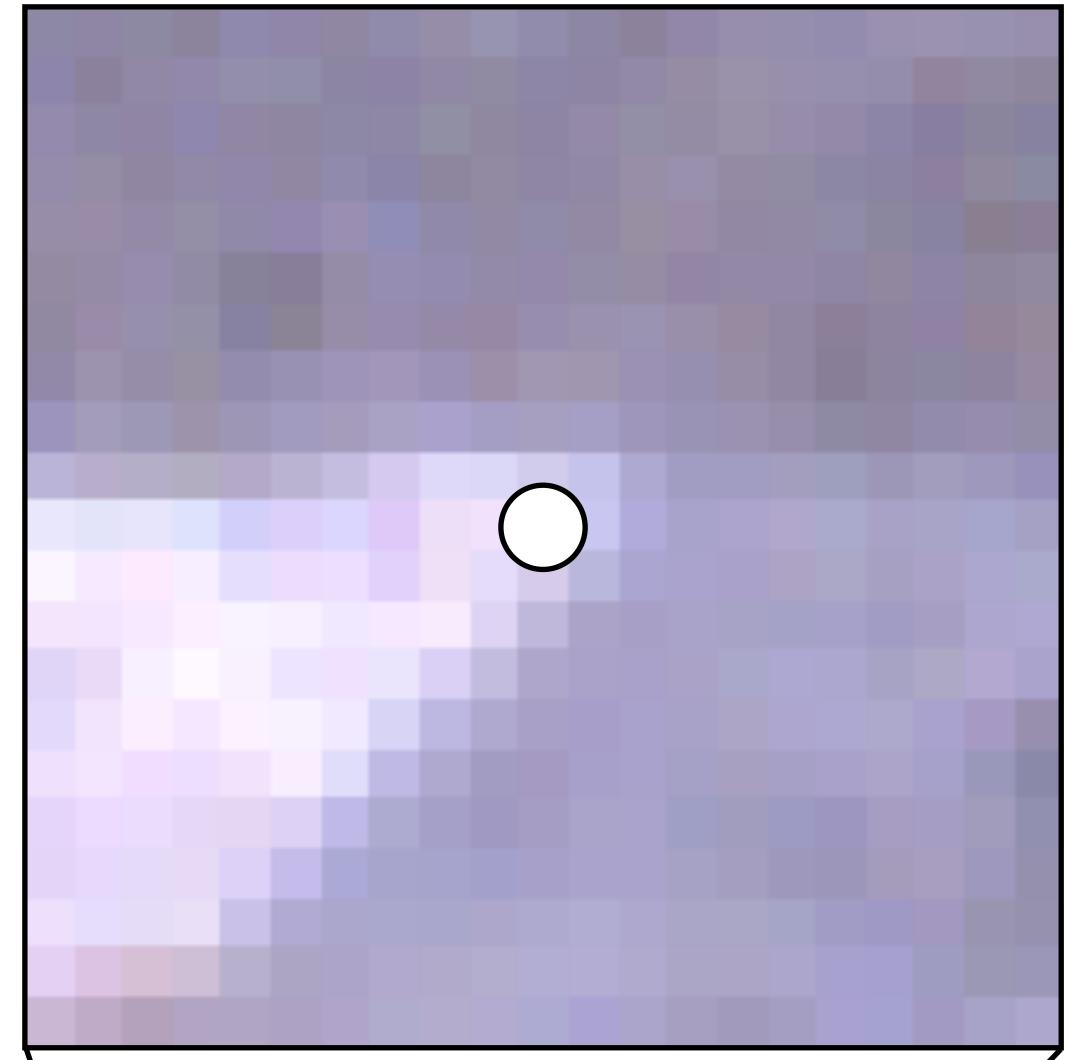
Green: Ground truth correspondence, Red: Initial solution

Ellipse: Gaussian distribution

Gauss-Newton Loss:

Maximize estimated probability
of ground truth correspondence

Loss Formulation:



Loss Formulation:

1. The point is at the correct location

→ The residual should be small!

$$E_{\text{pos}} = \|F(\bigcirc) - F'(\bullet)\|^2$$

2. The point is an outlier

→ The residual should be large!

$$E_{\text{neg}} = \|F(\bigcirc) - F'(\bullet)\|^2 > M$$

3. The point is relatively far

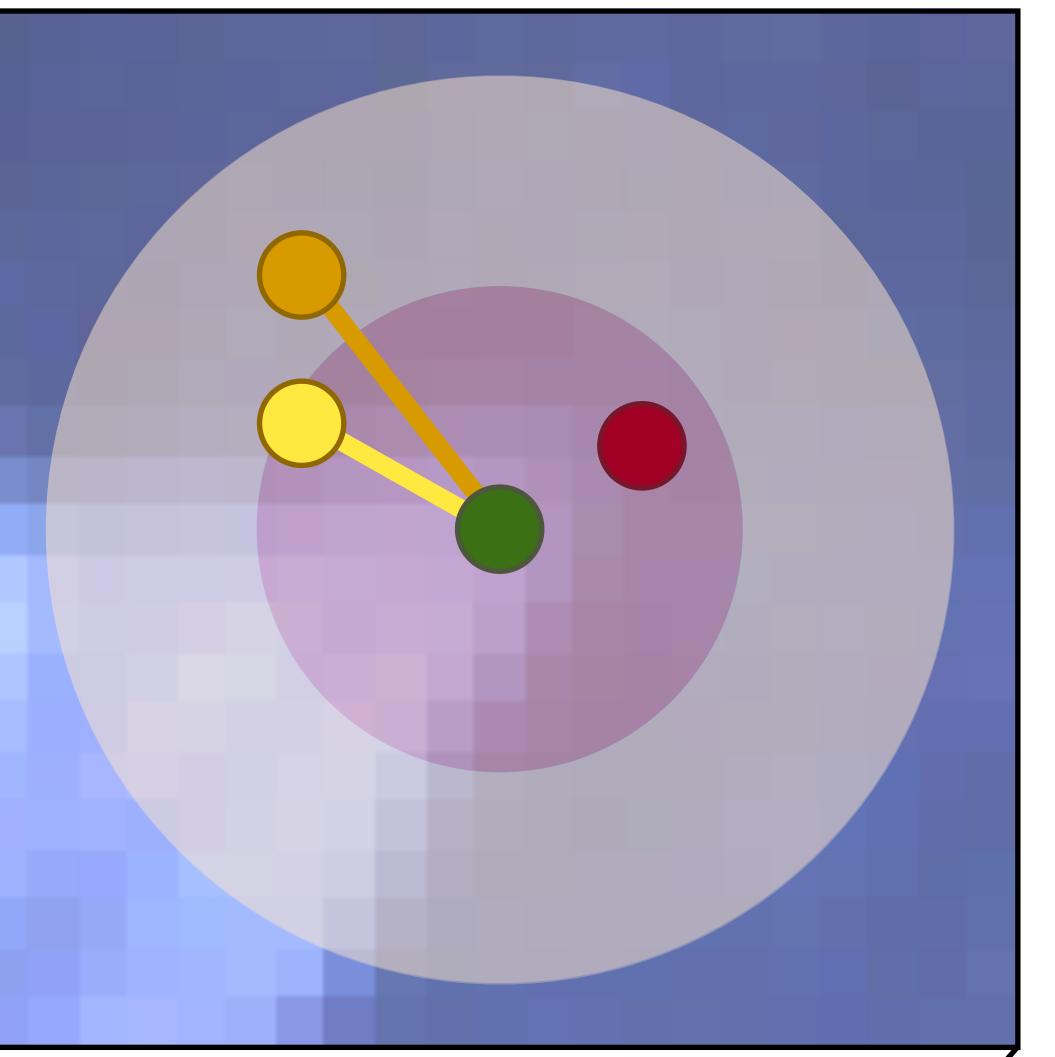
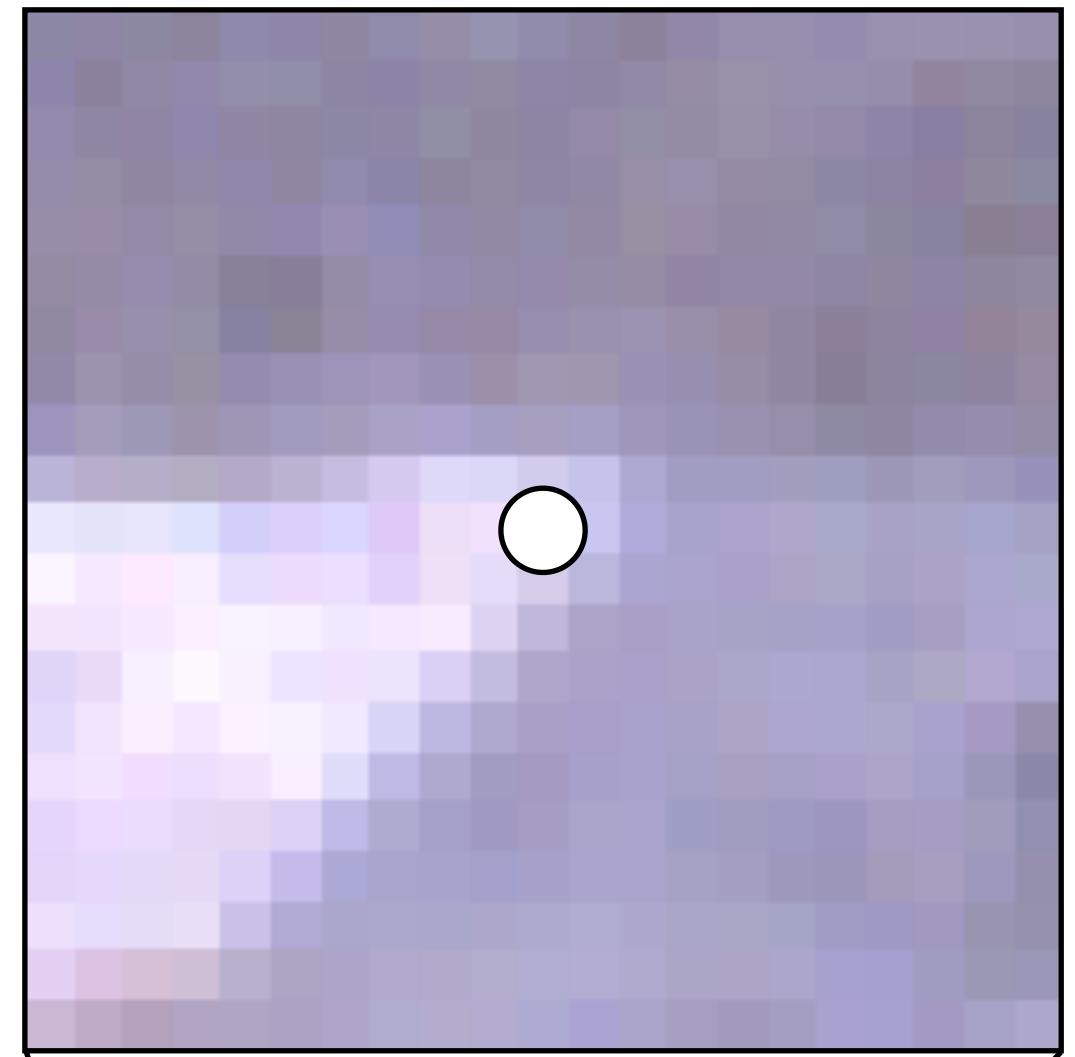
→ Gradient shall point in right direction

$$E_{\text{GD}} = \underline{\text{dist}_{\text{after}}} < \underline{\text{dist}_{\text{before}}} - \delta$$

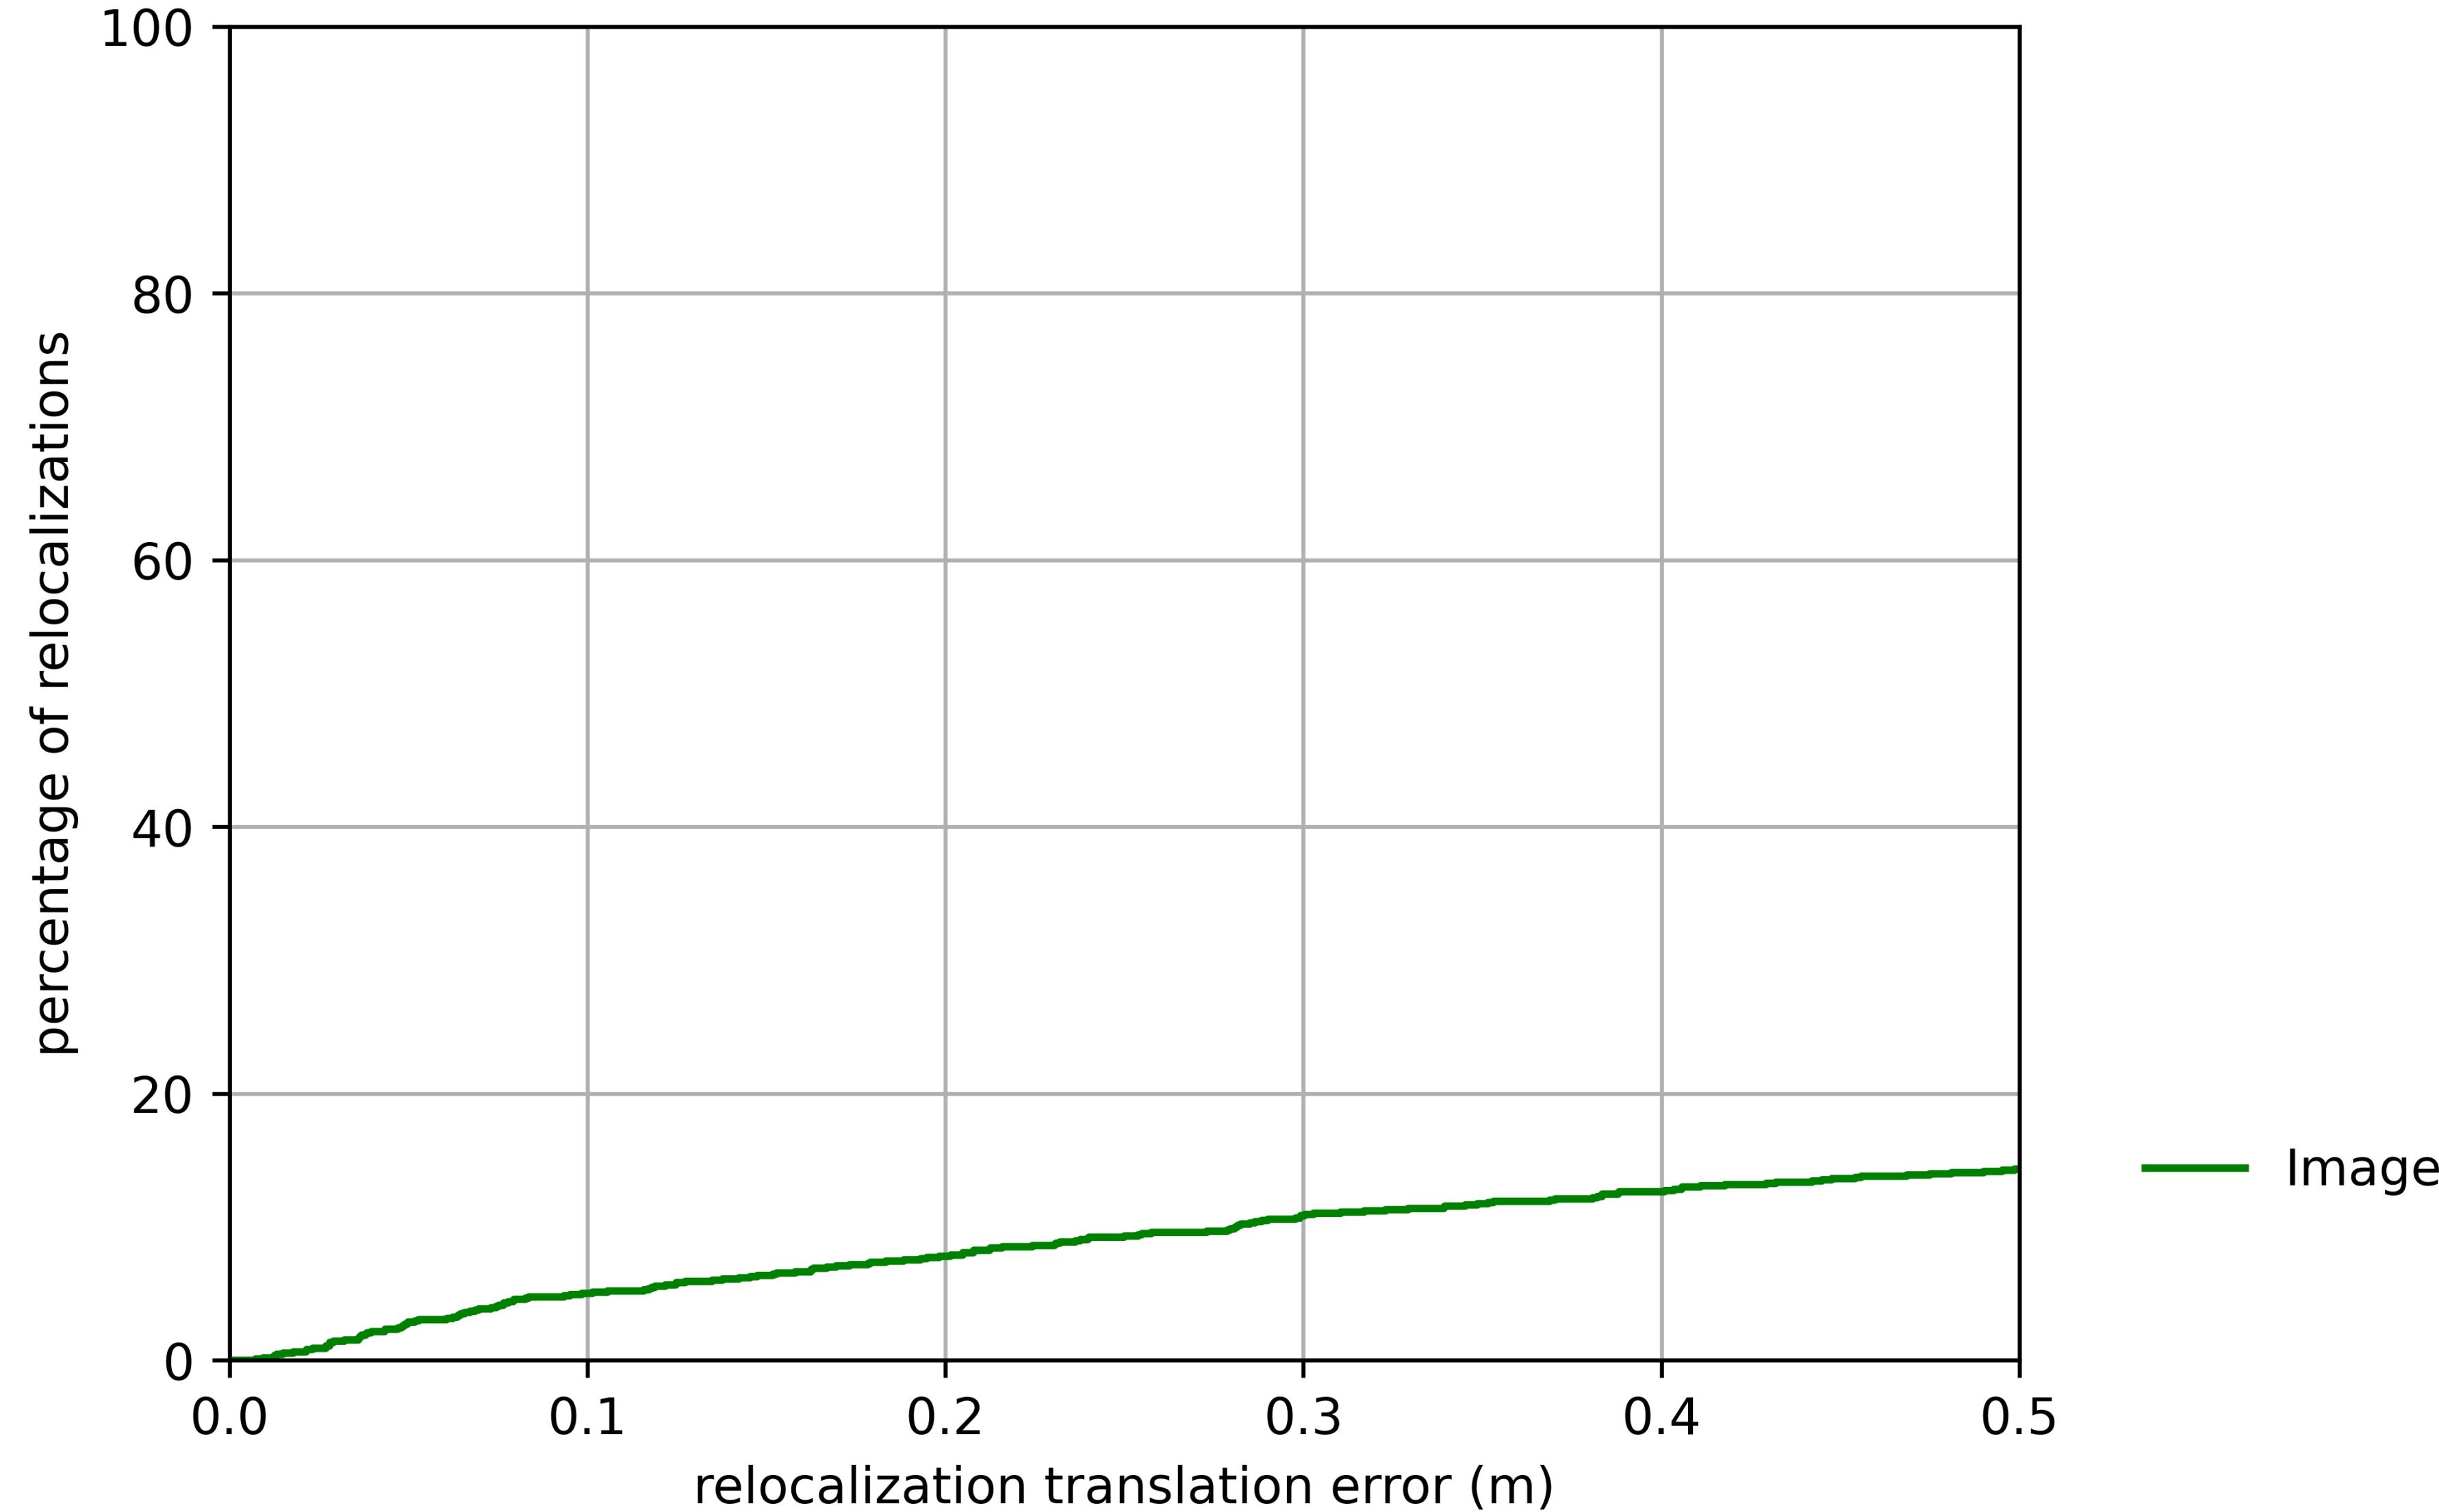
4. The point is very close

→ Now we should converge quickly.

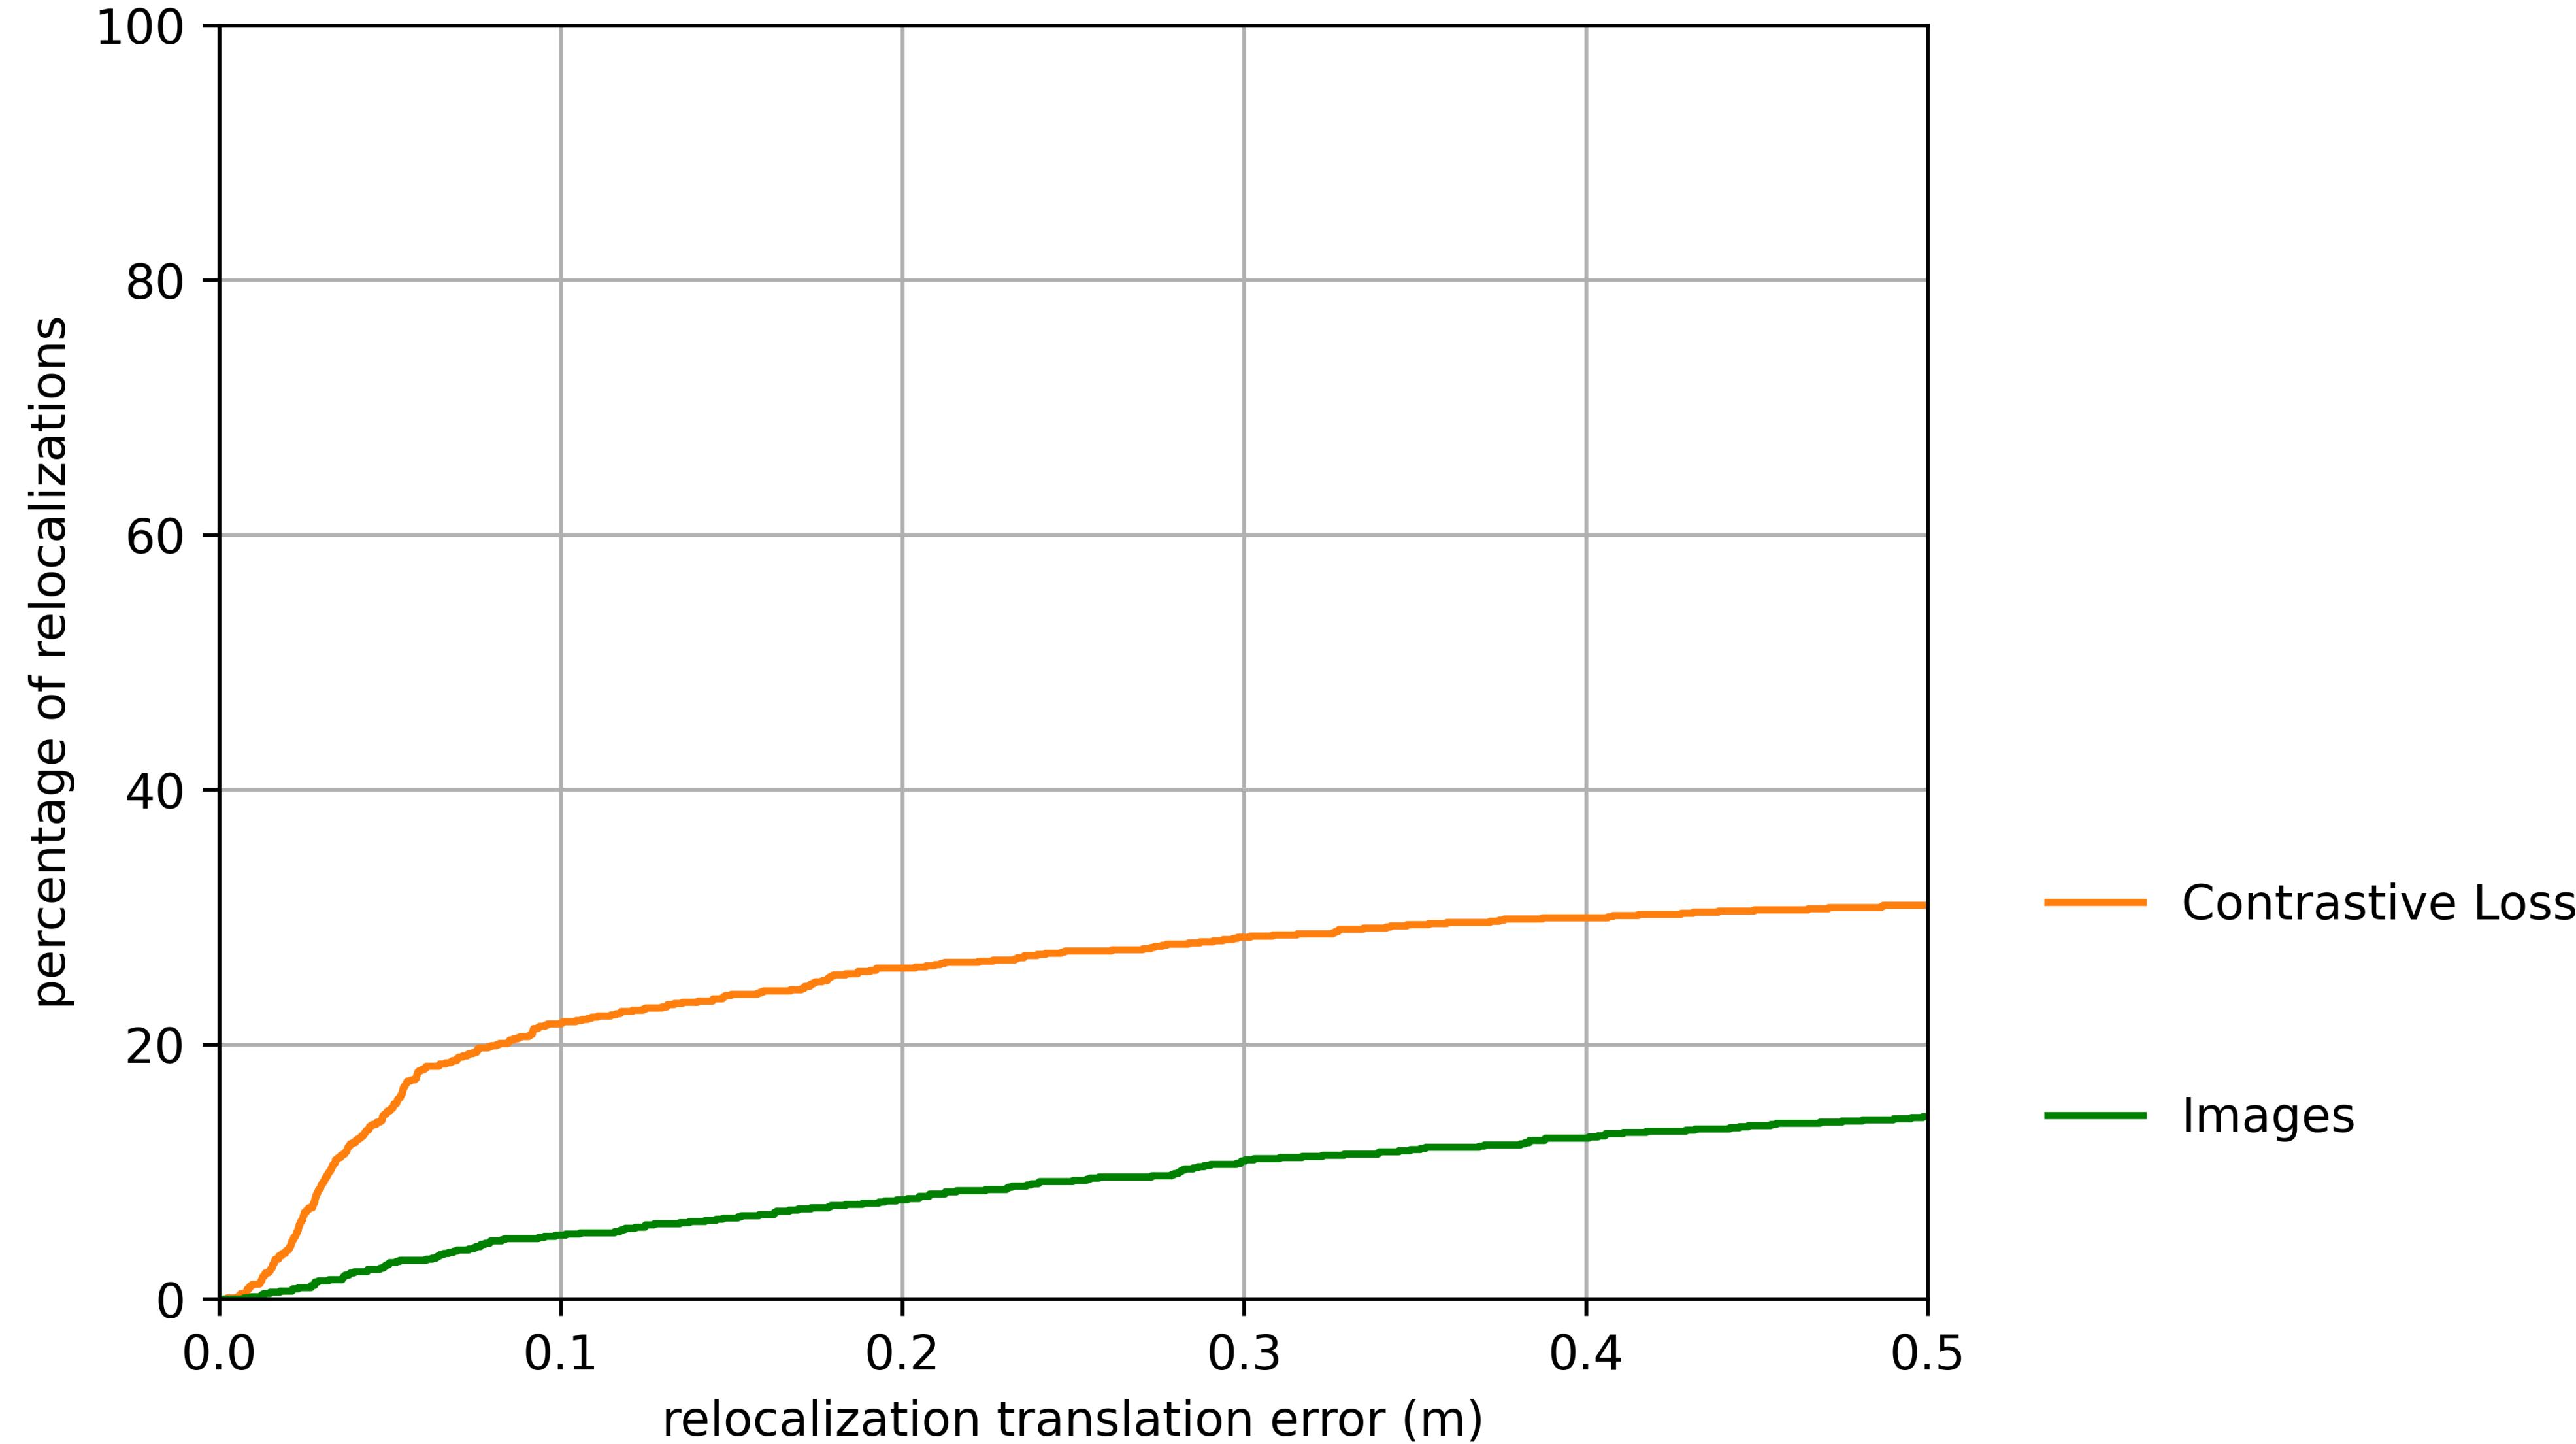
$$E_{\text{GN}} = \text{Gauss-Newton Loss}$$



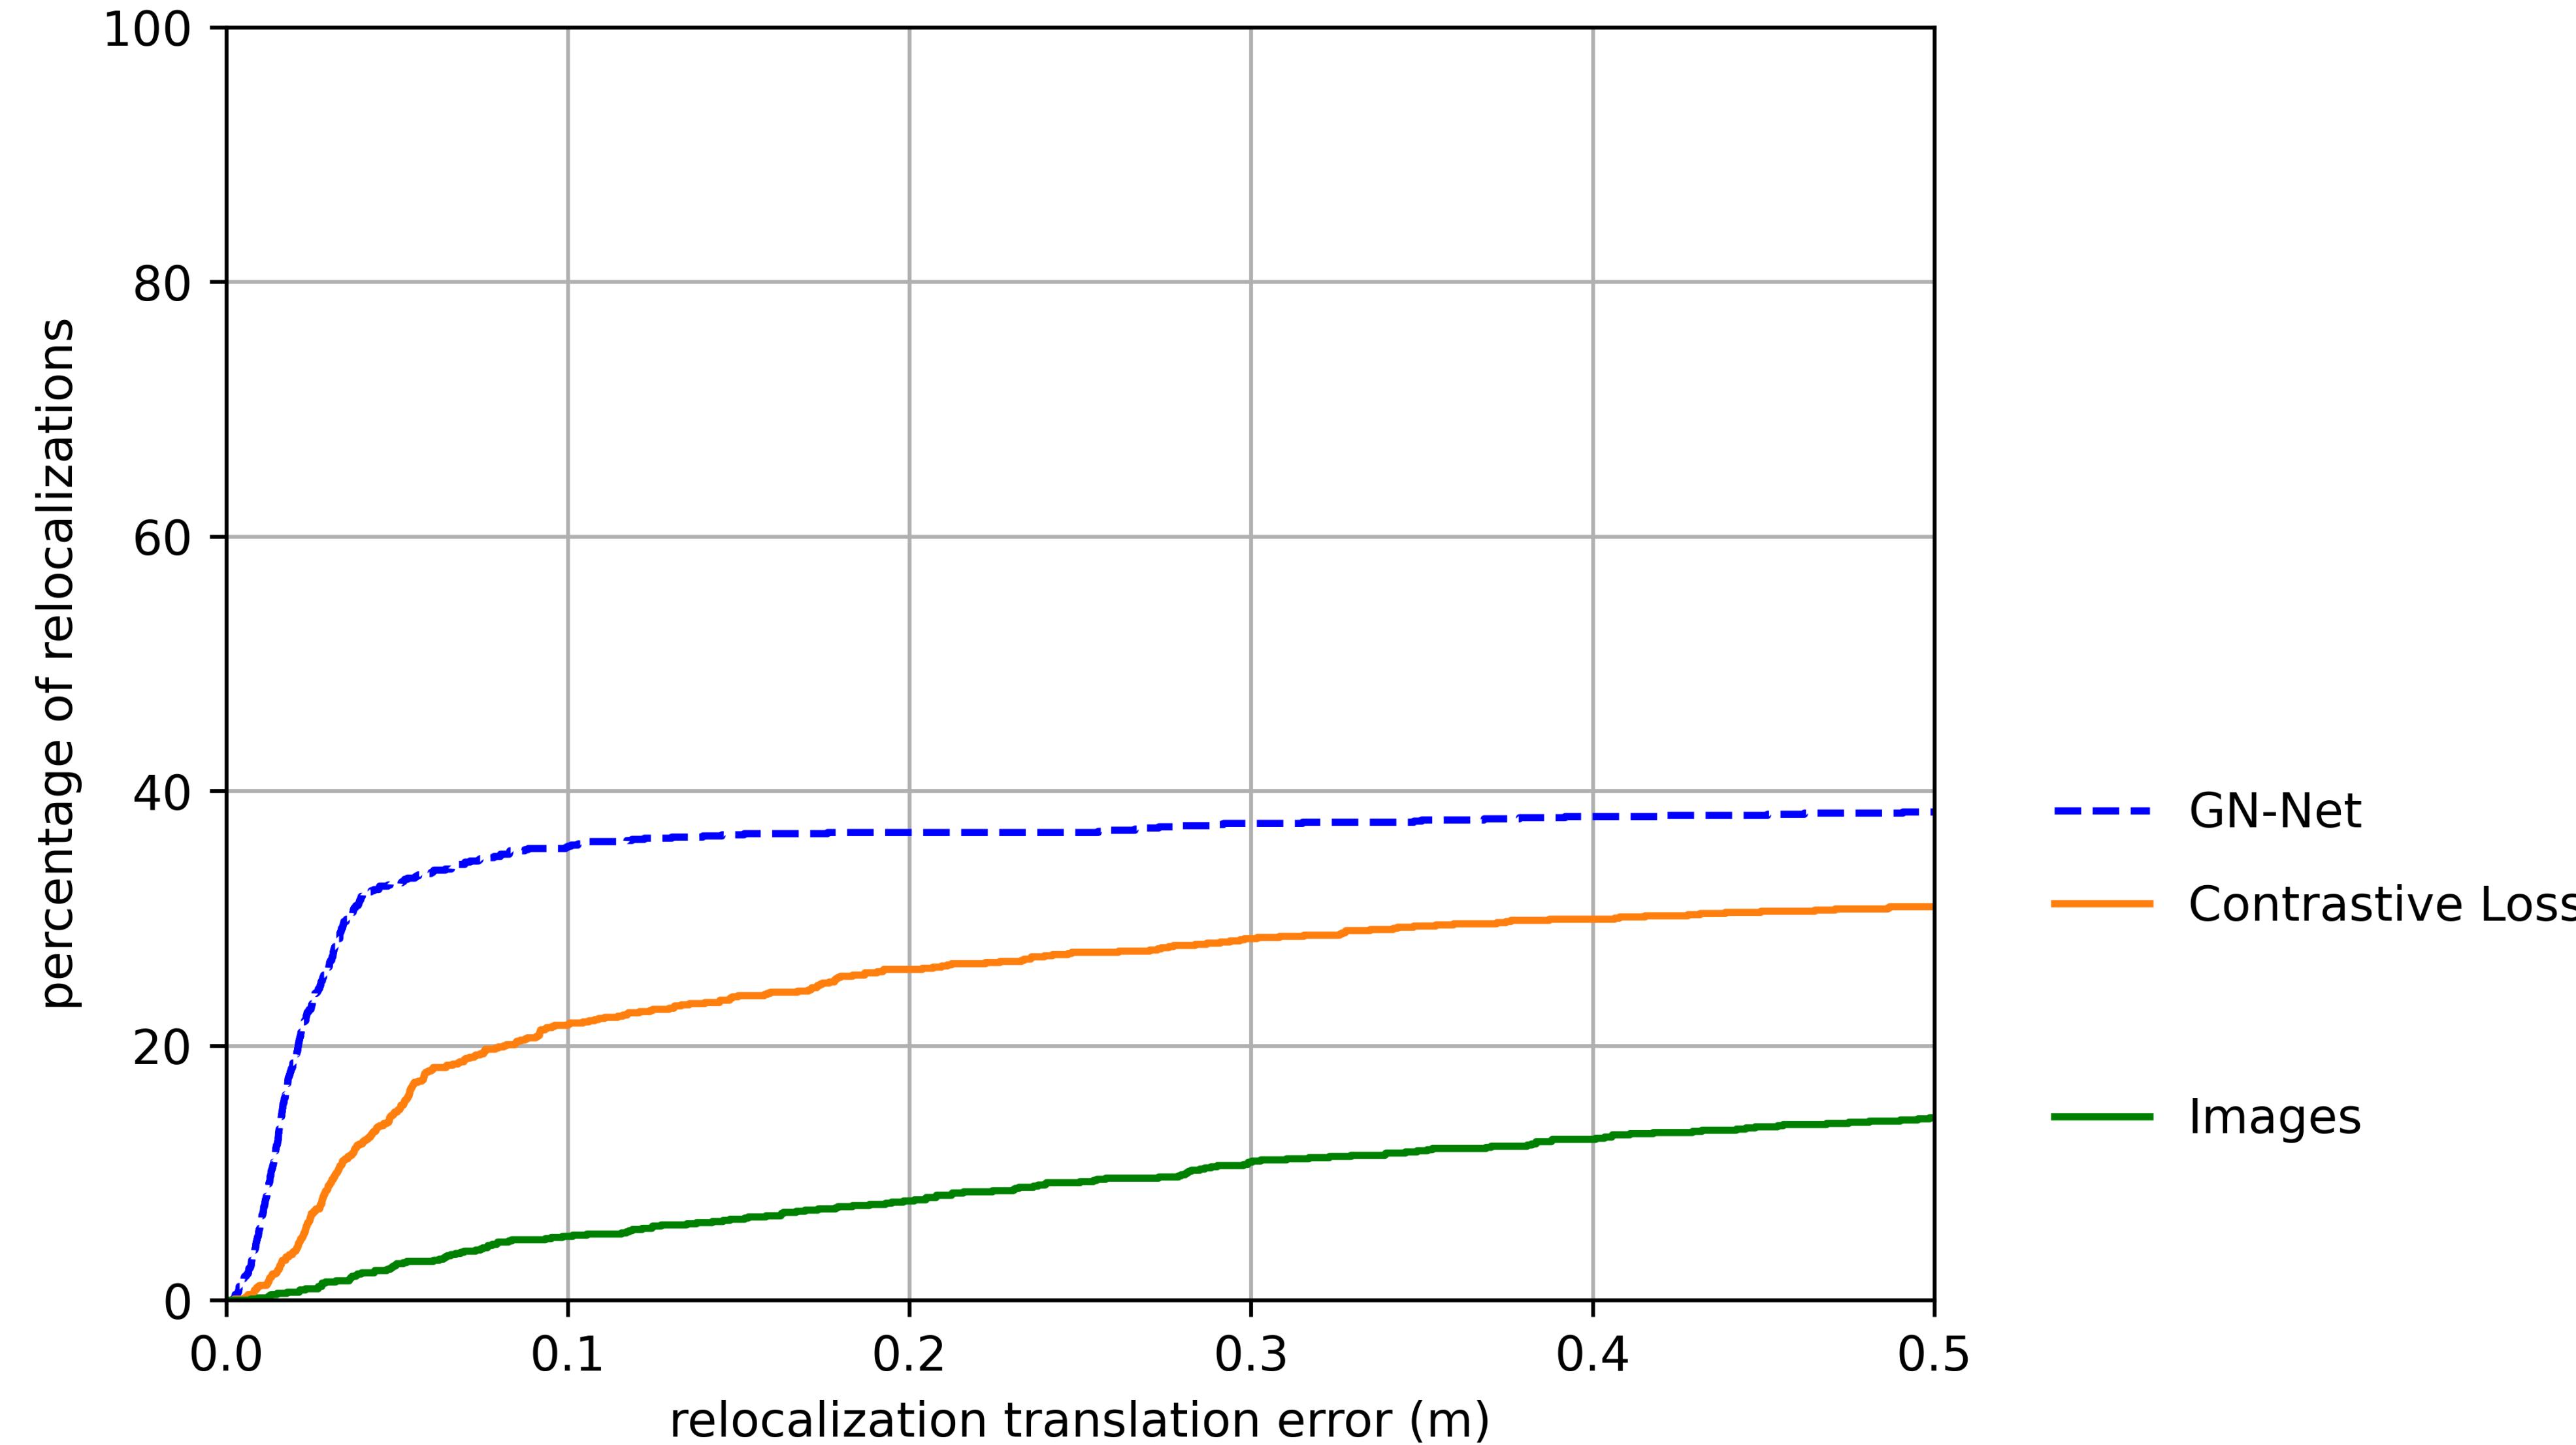
Relocalization Tracking Benchmark



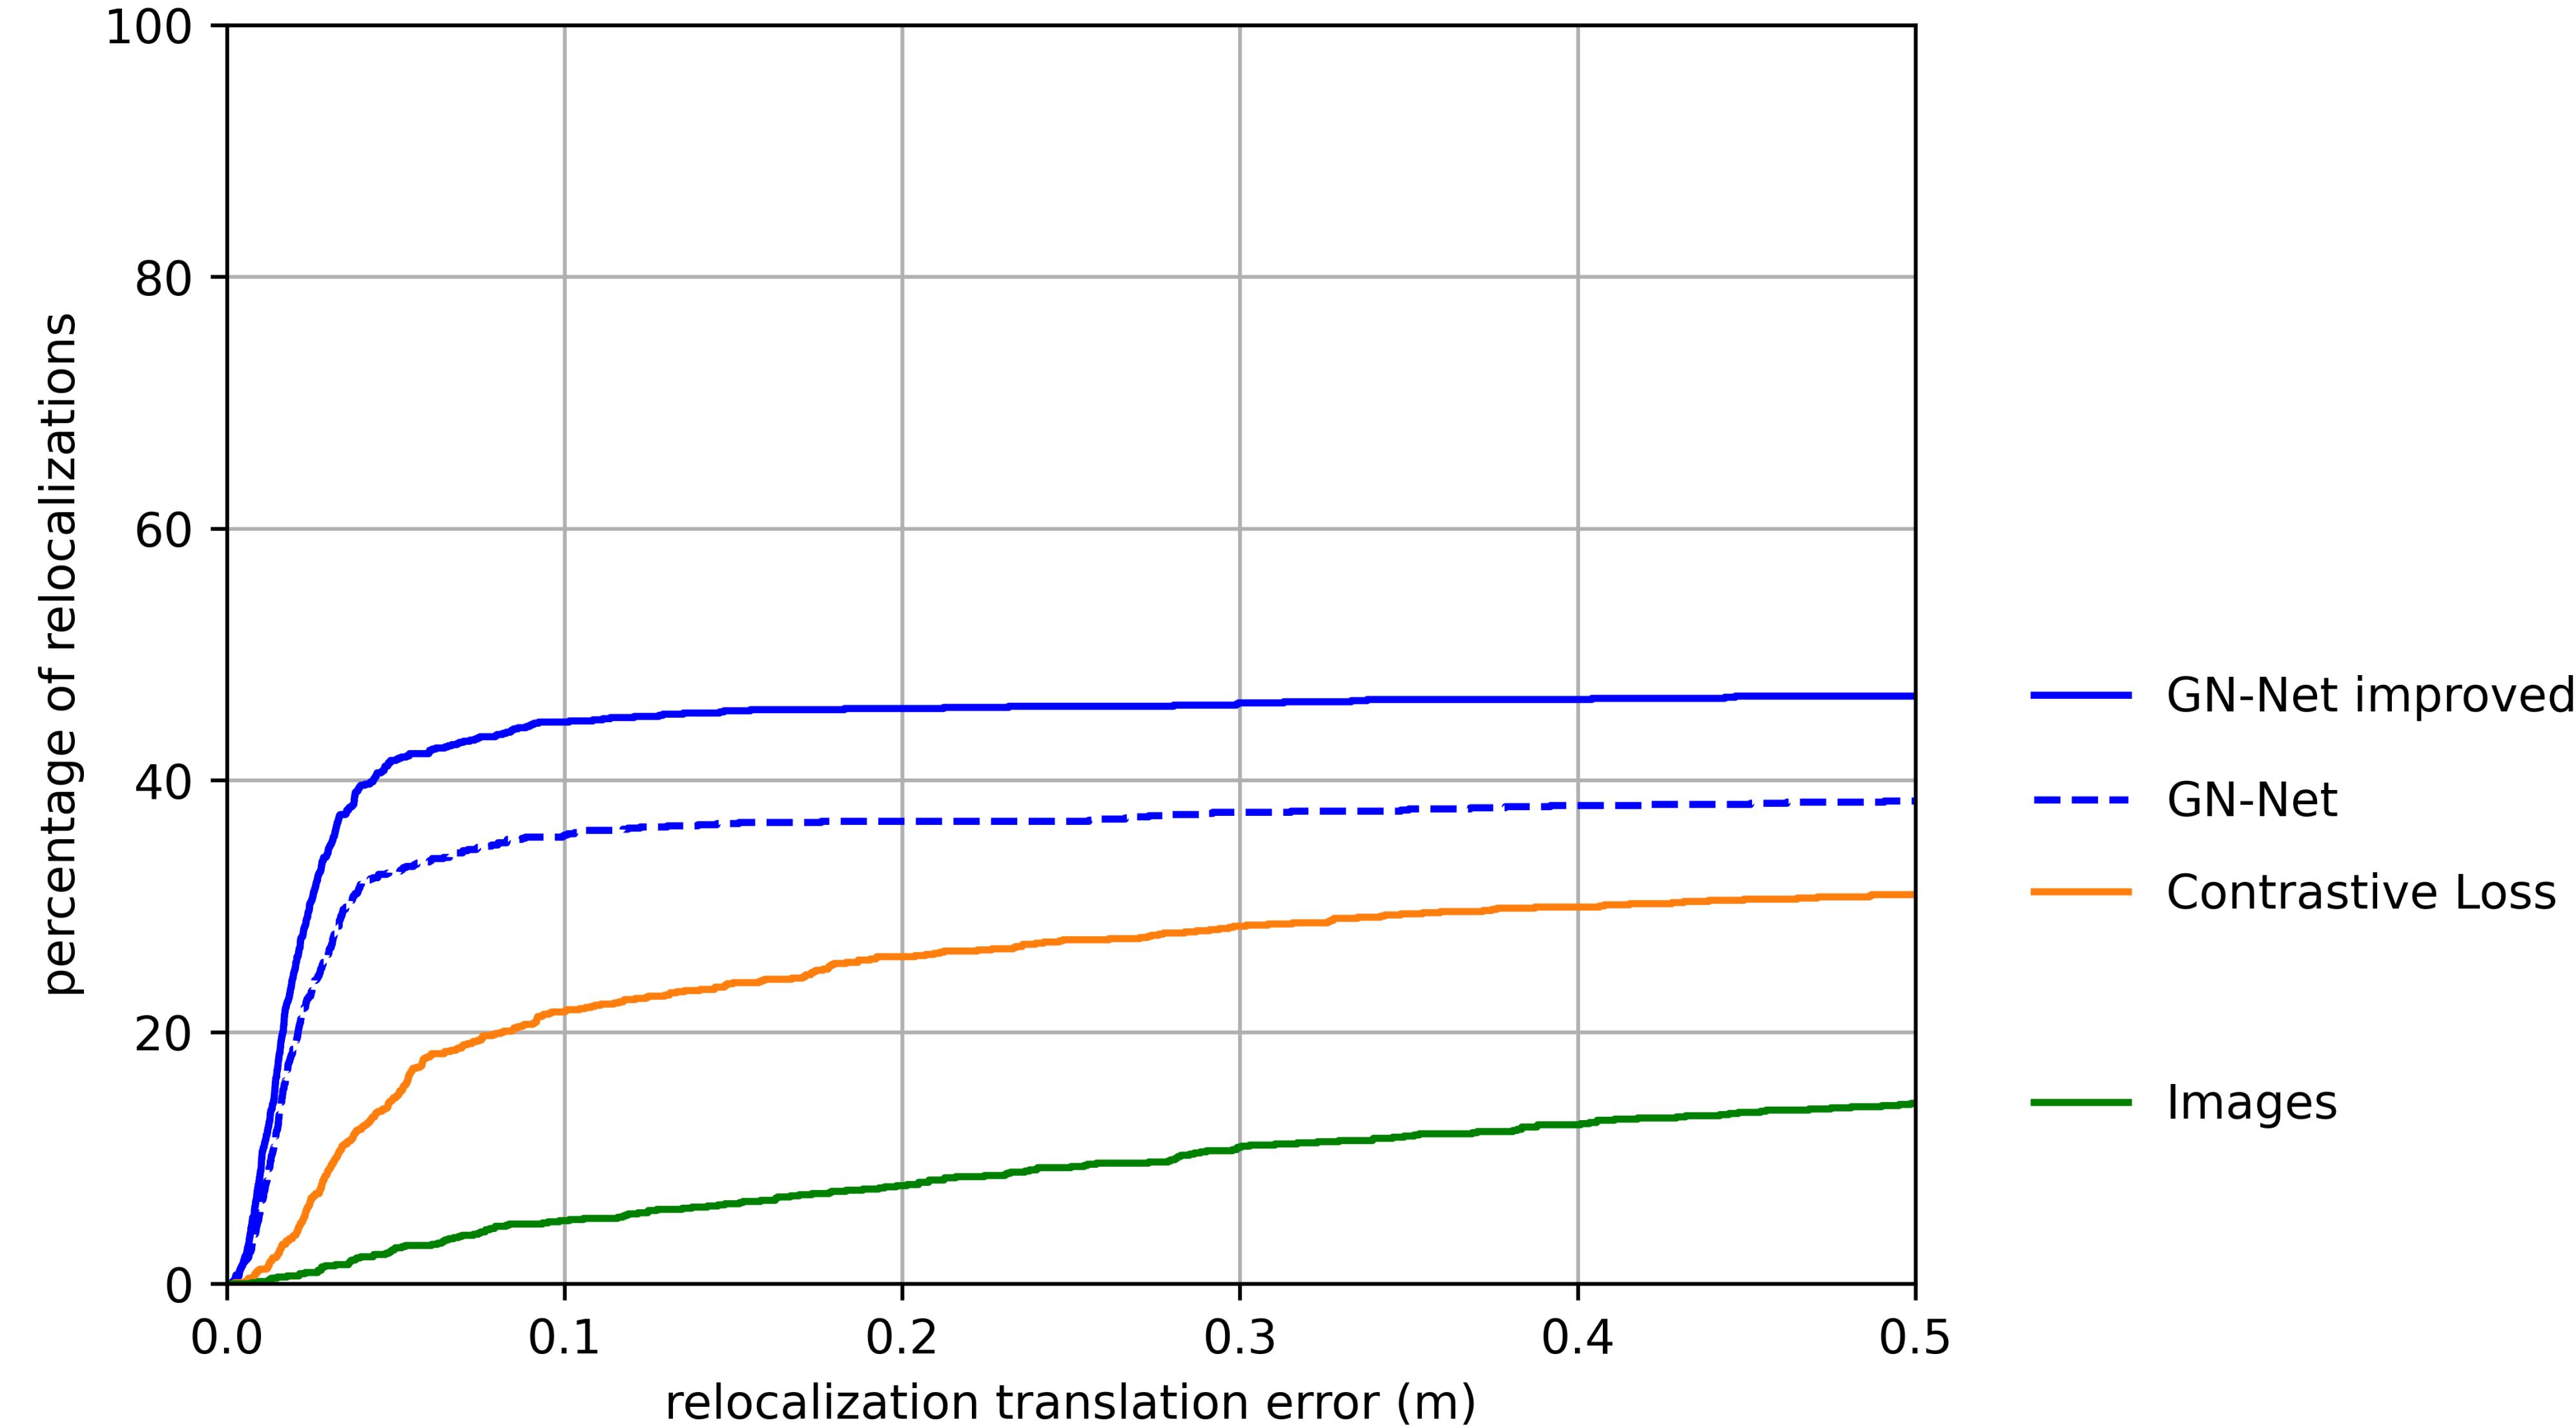
Relocalization Tracking Benchmark



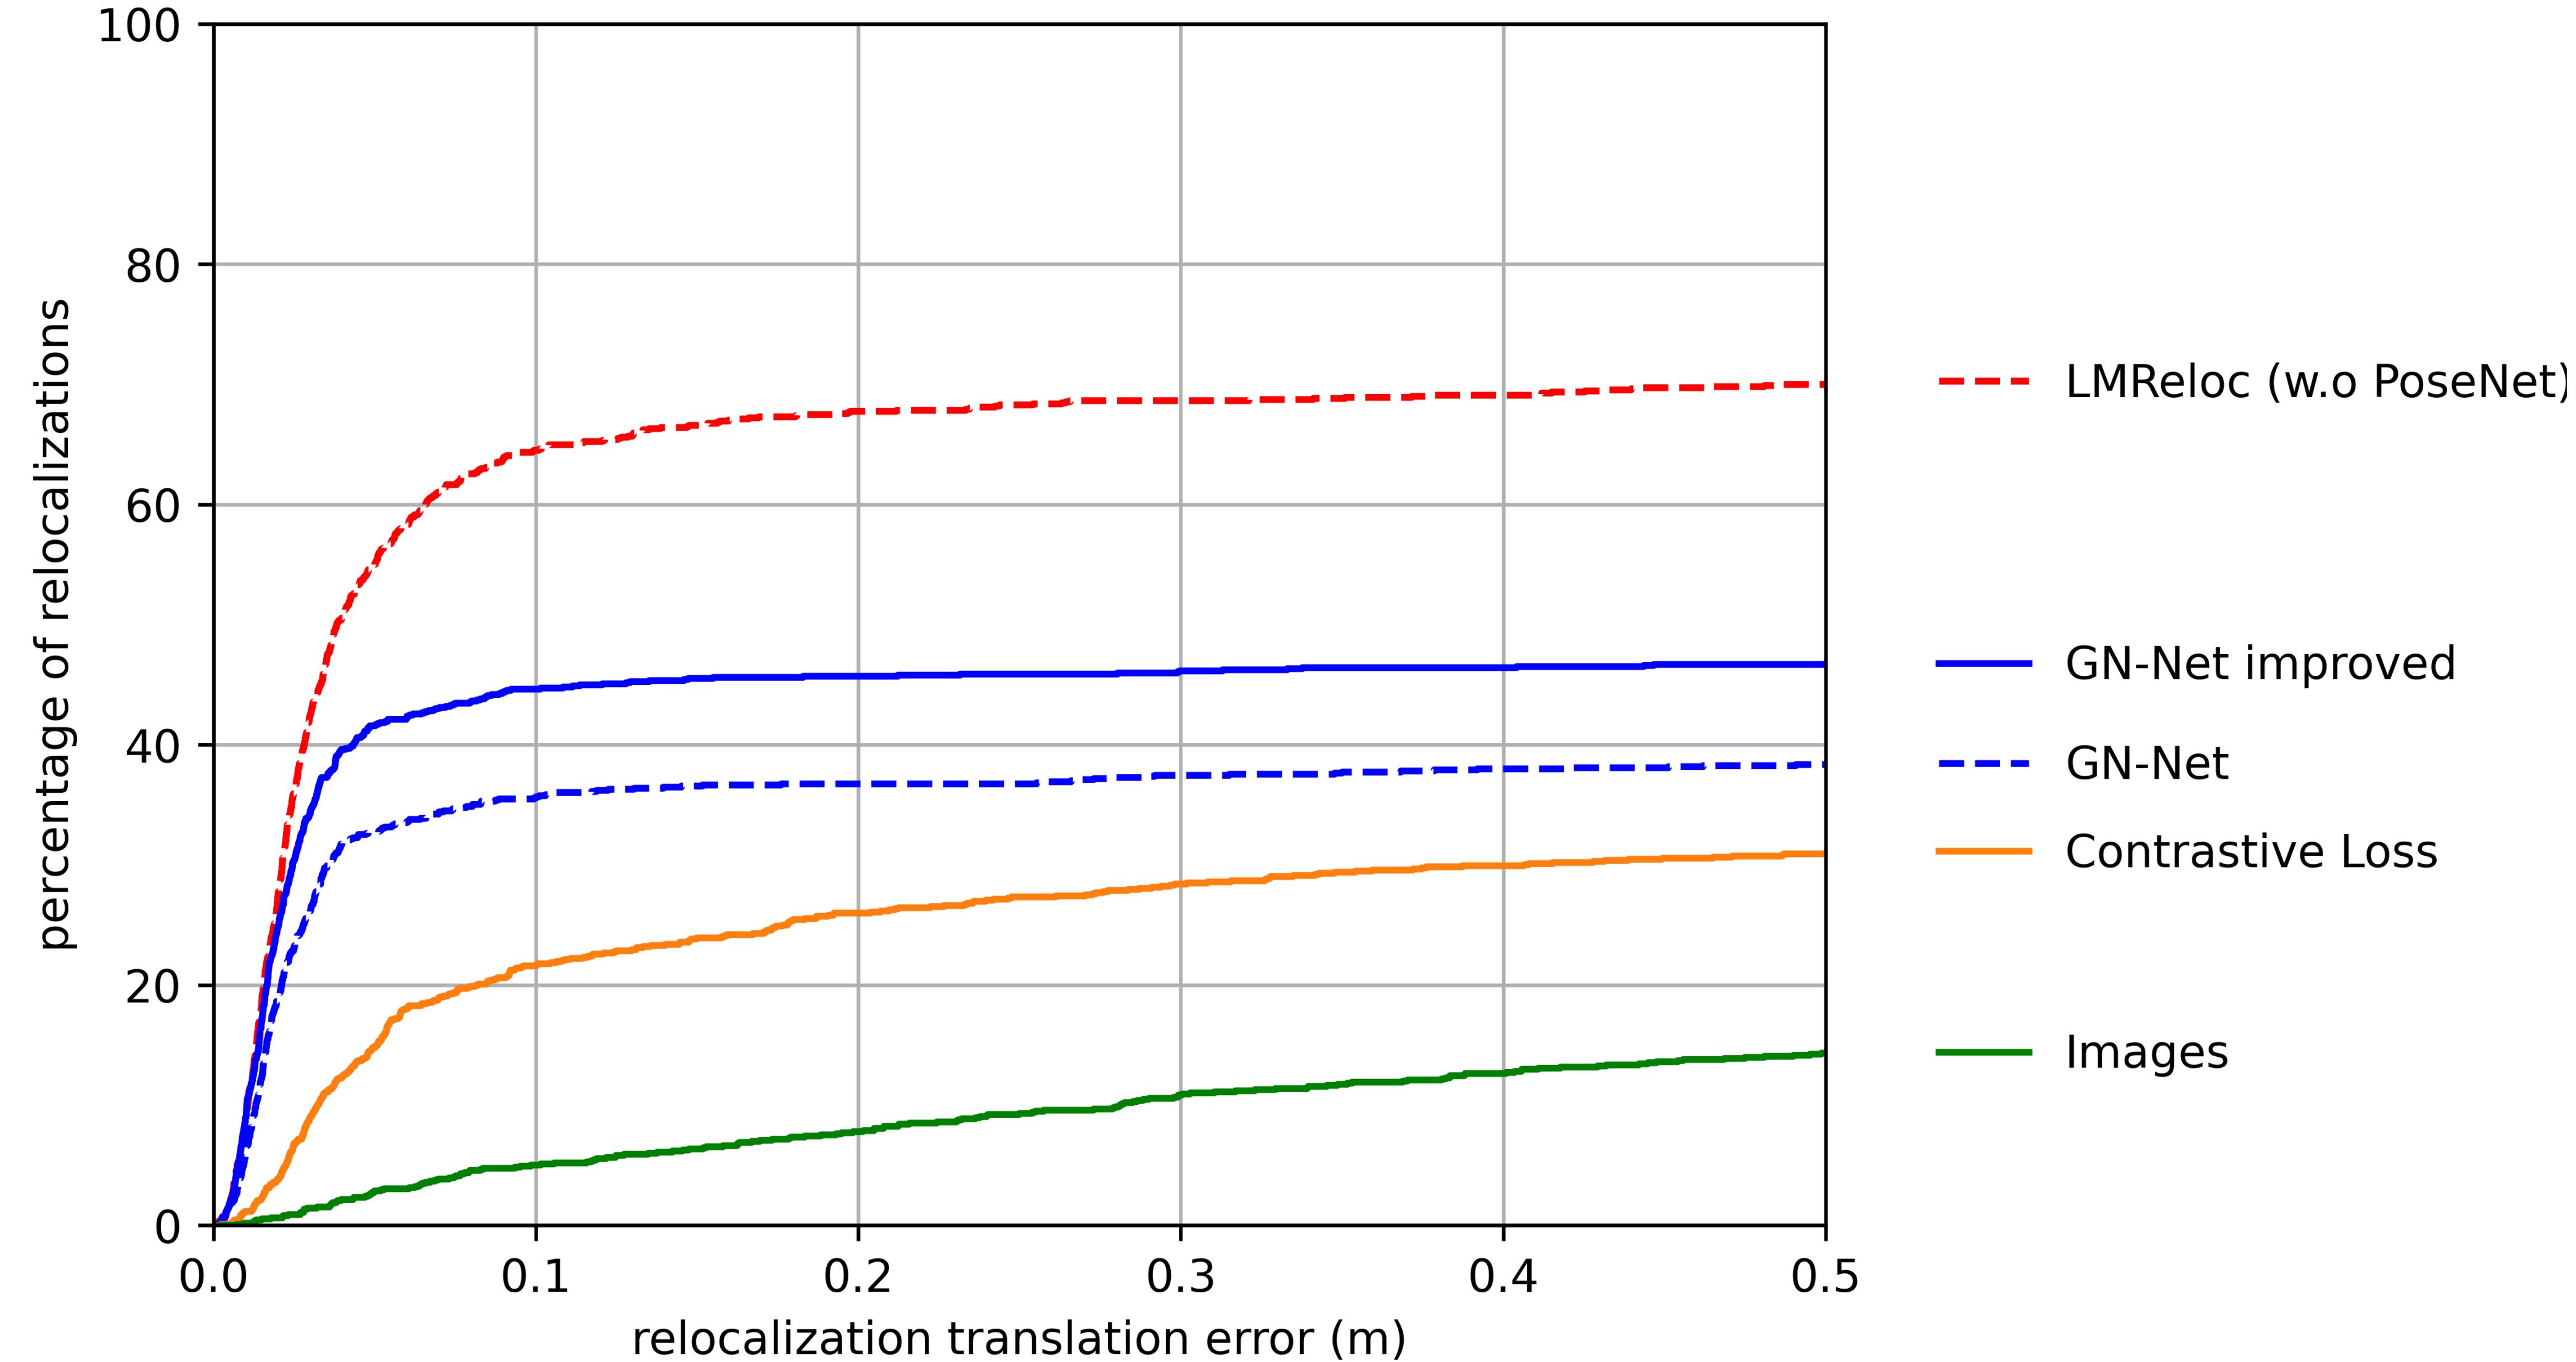
Relocalization Tracking Benchmark



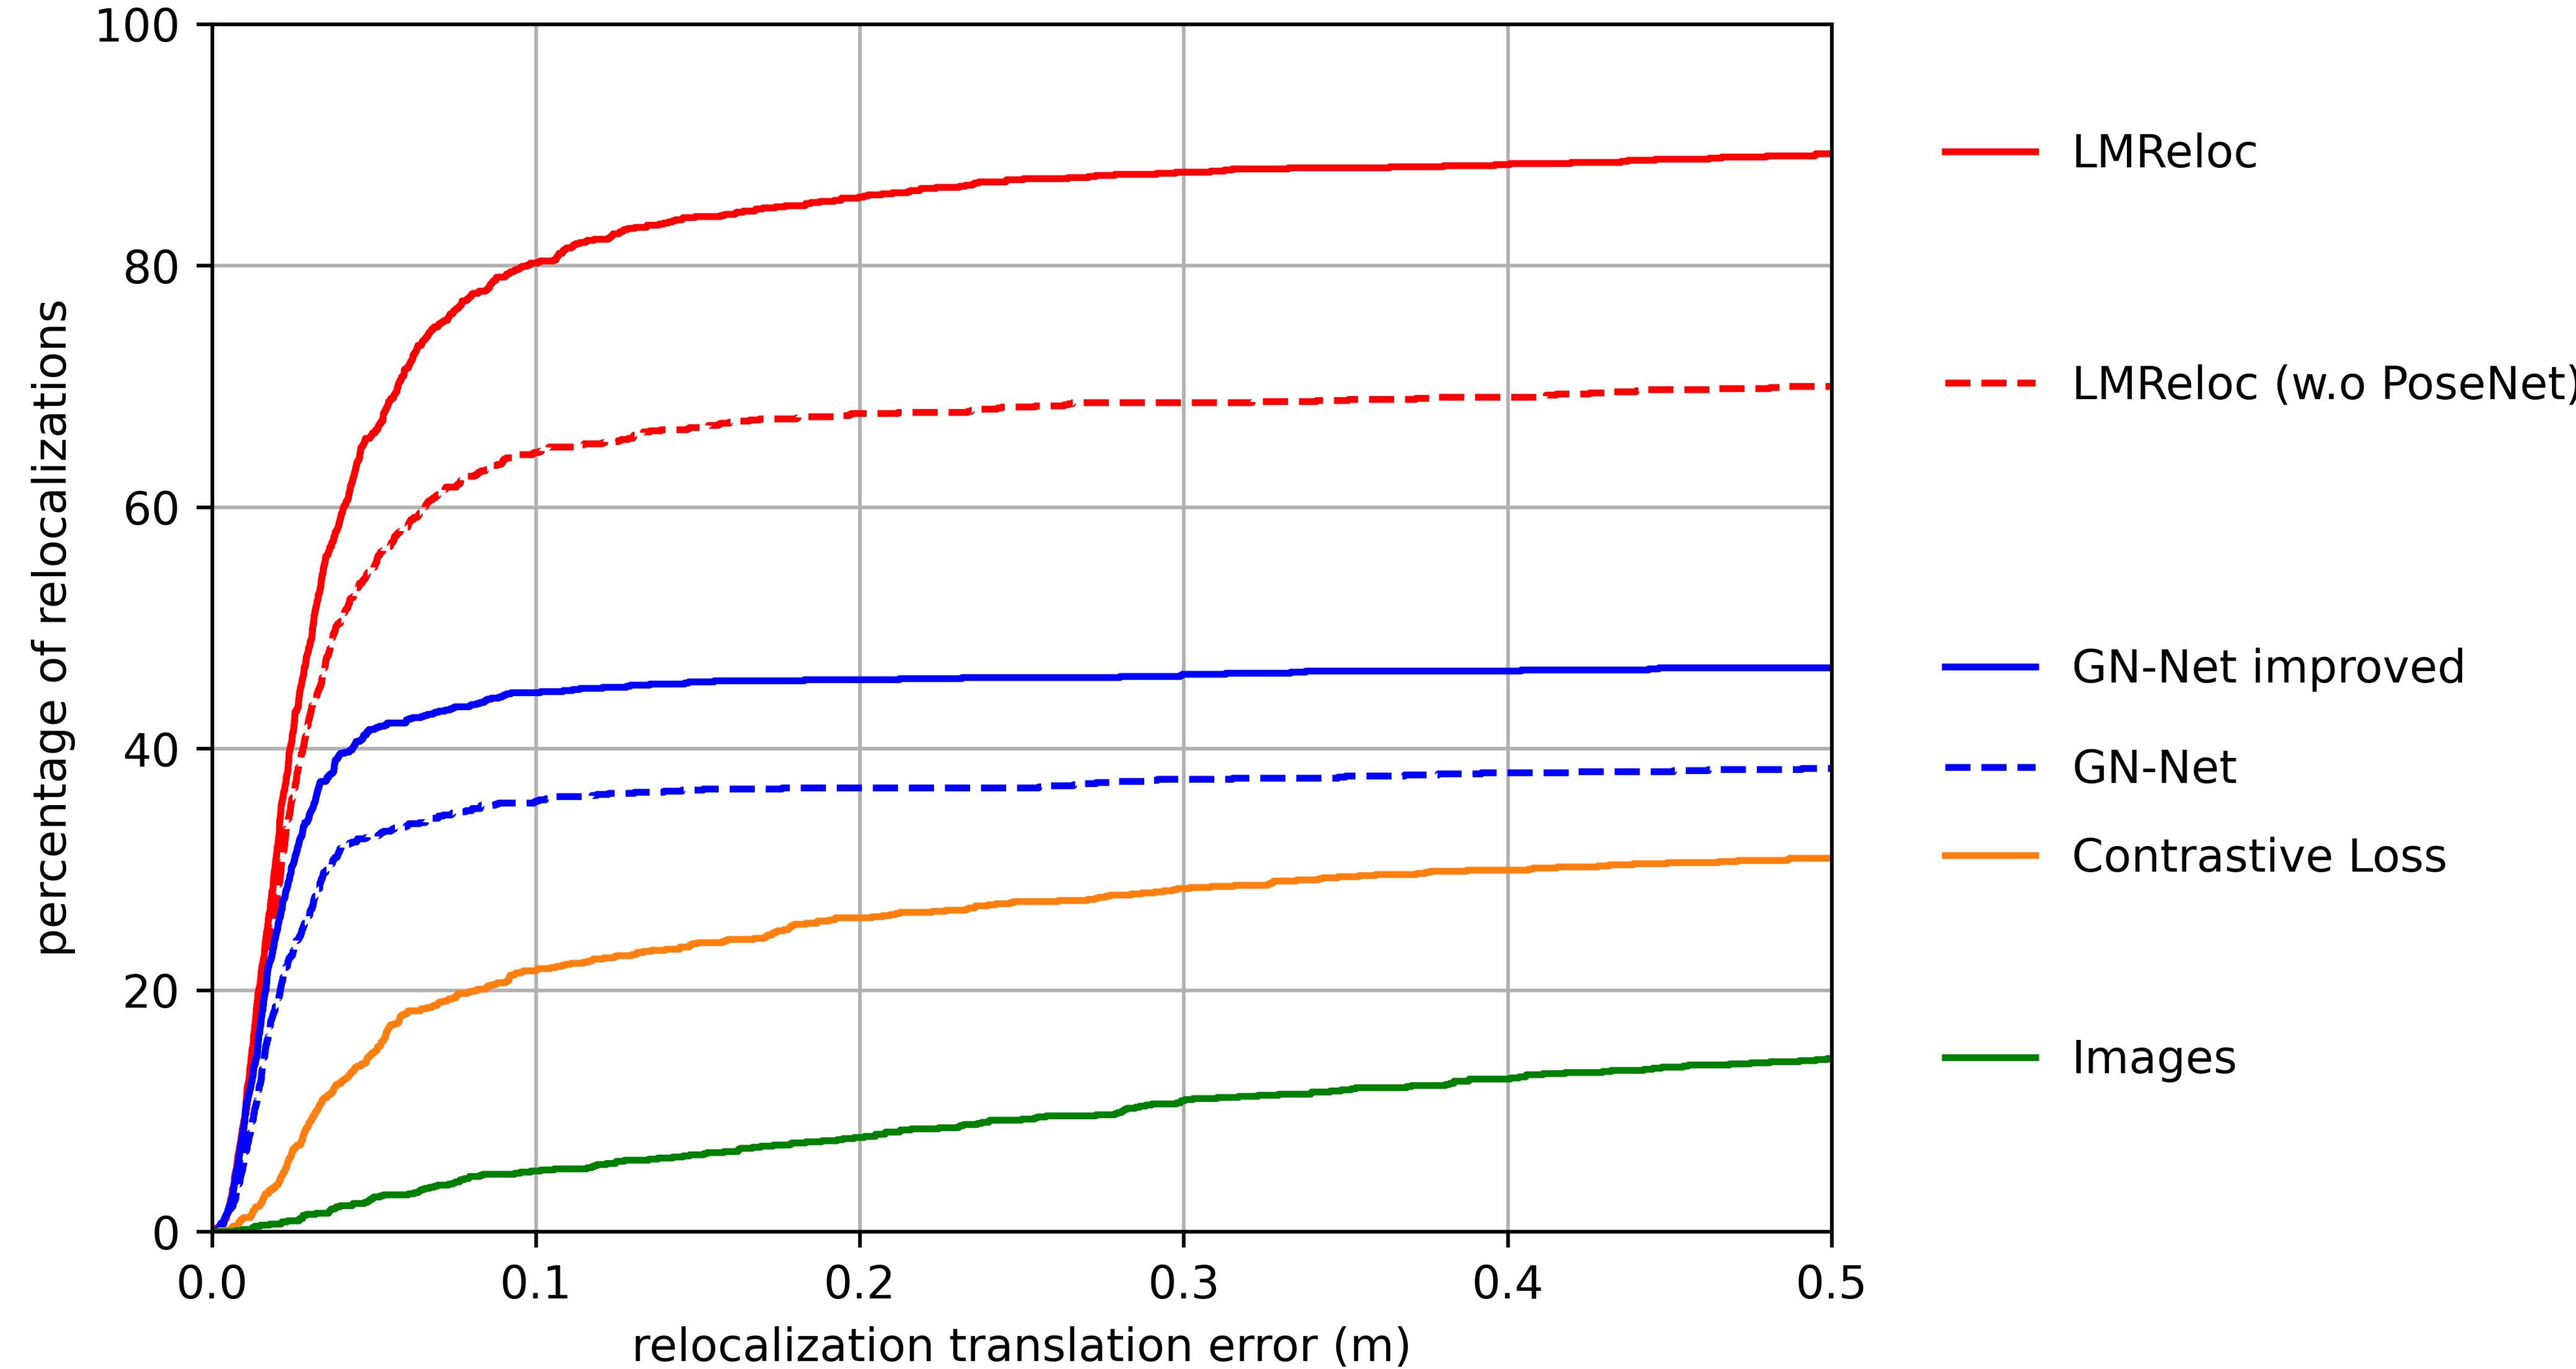
Relocalization Tracking Benchmark

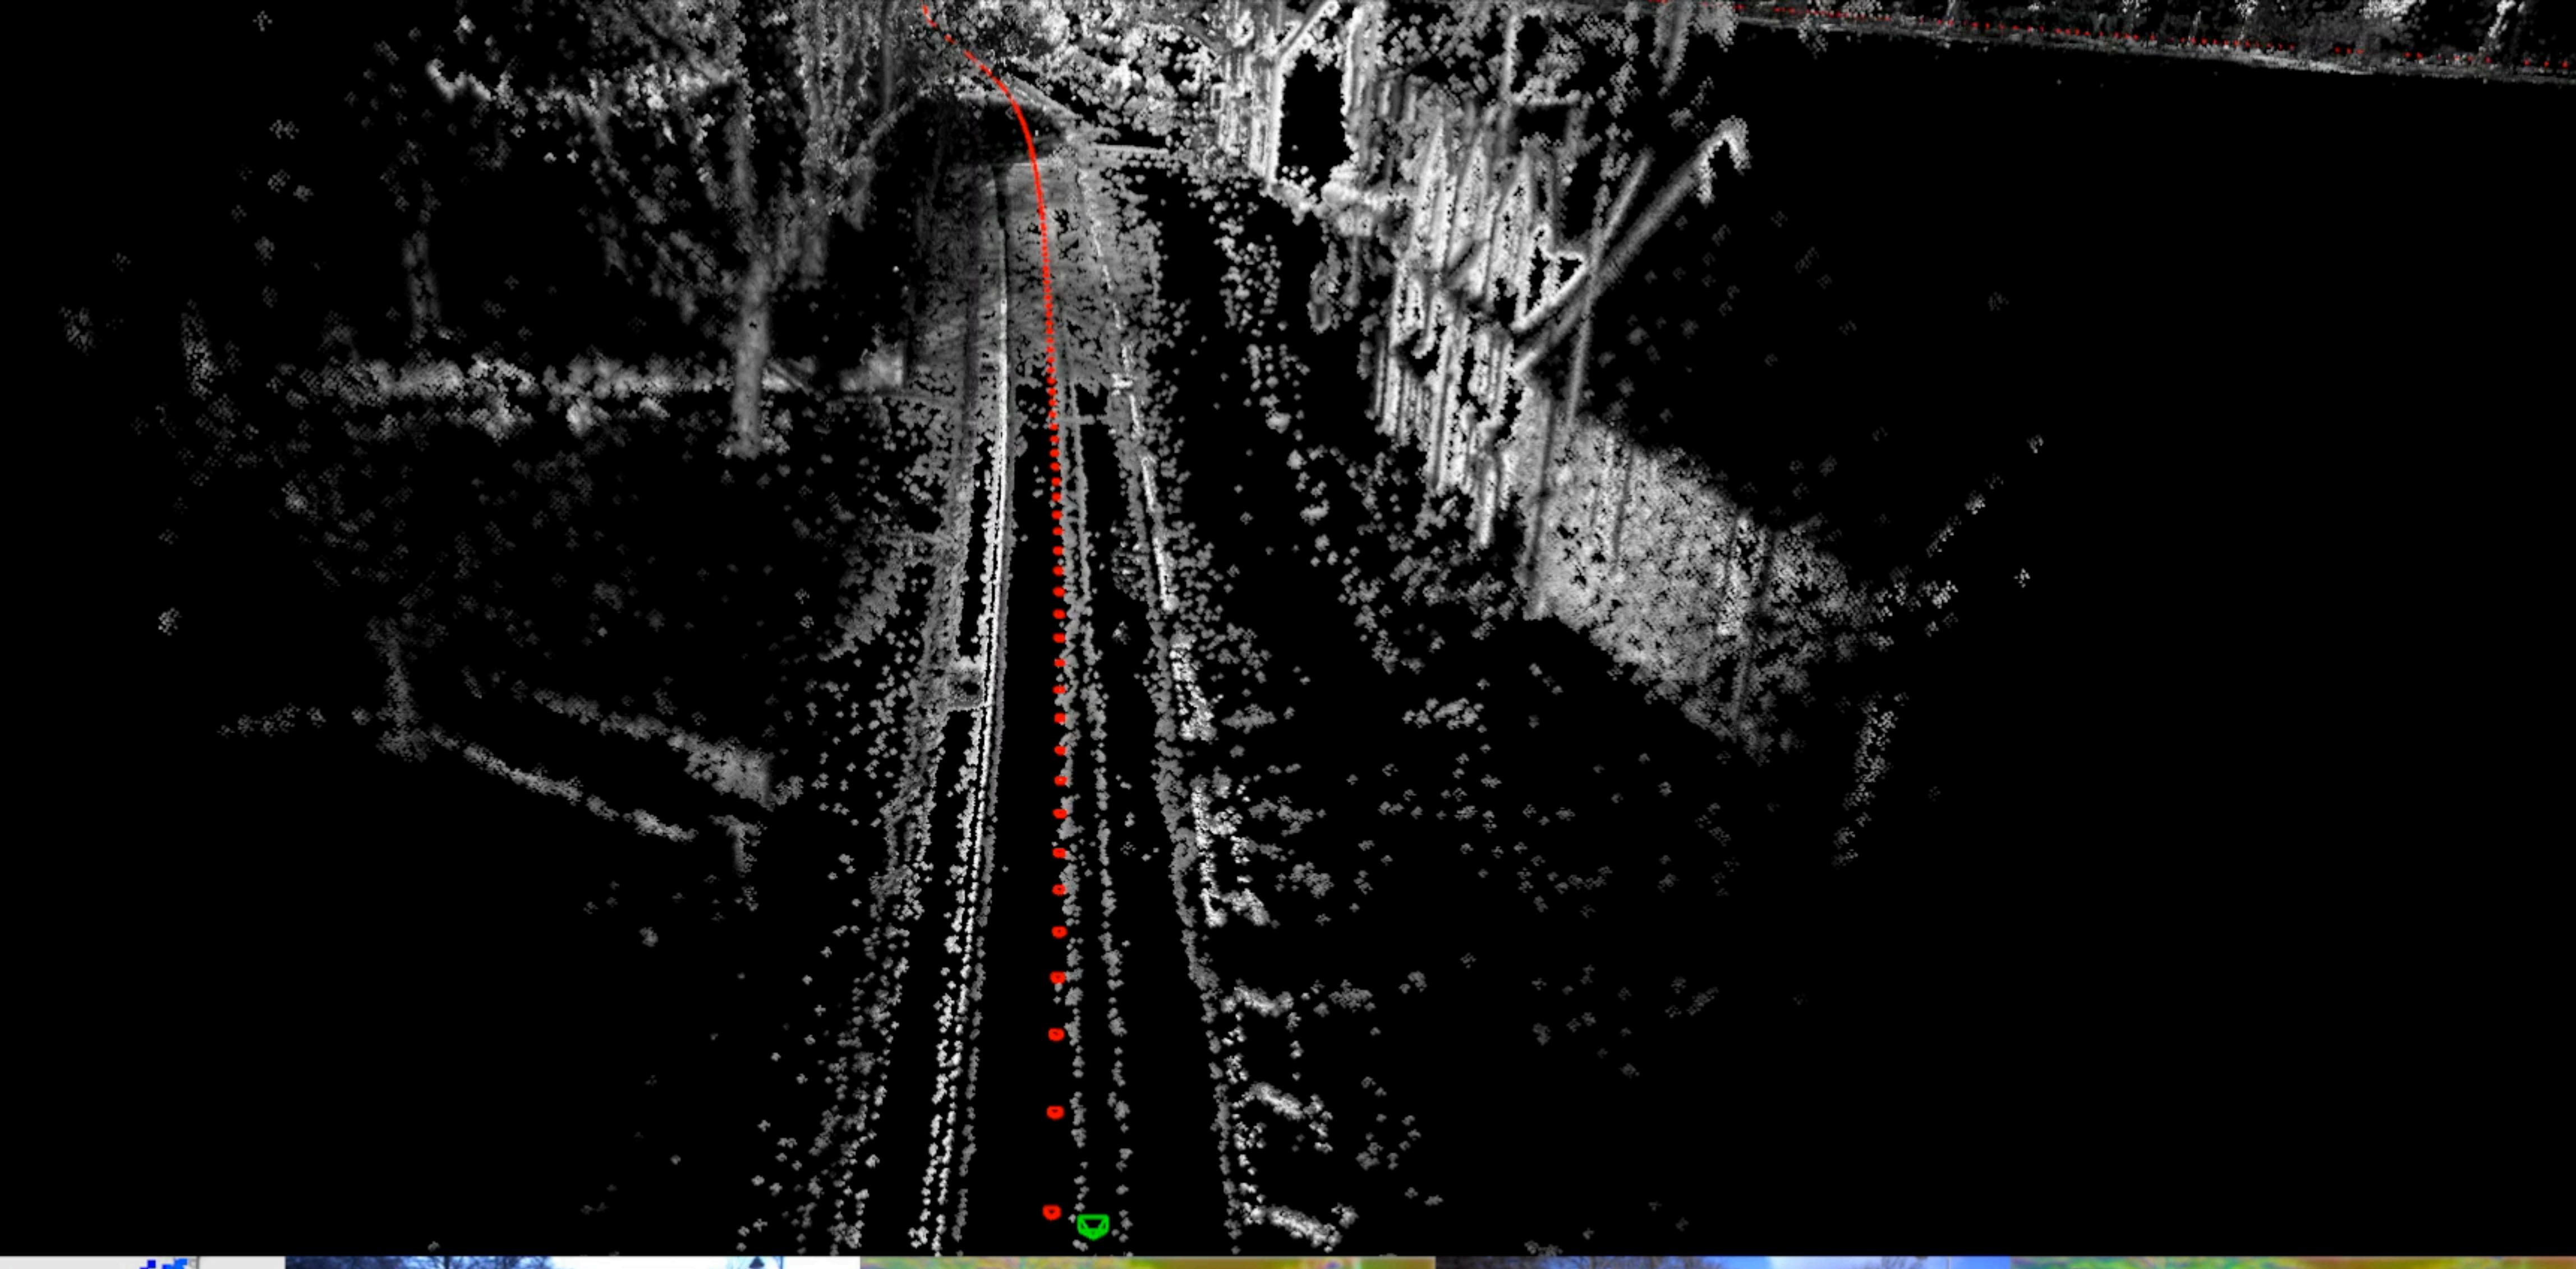


Relocalization Tracking Benchmark



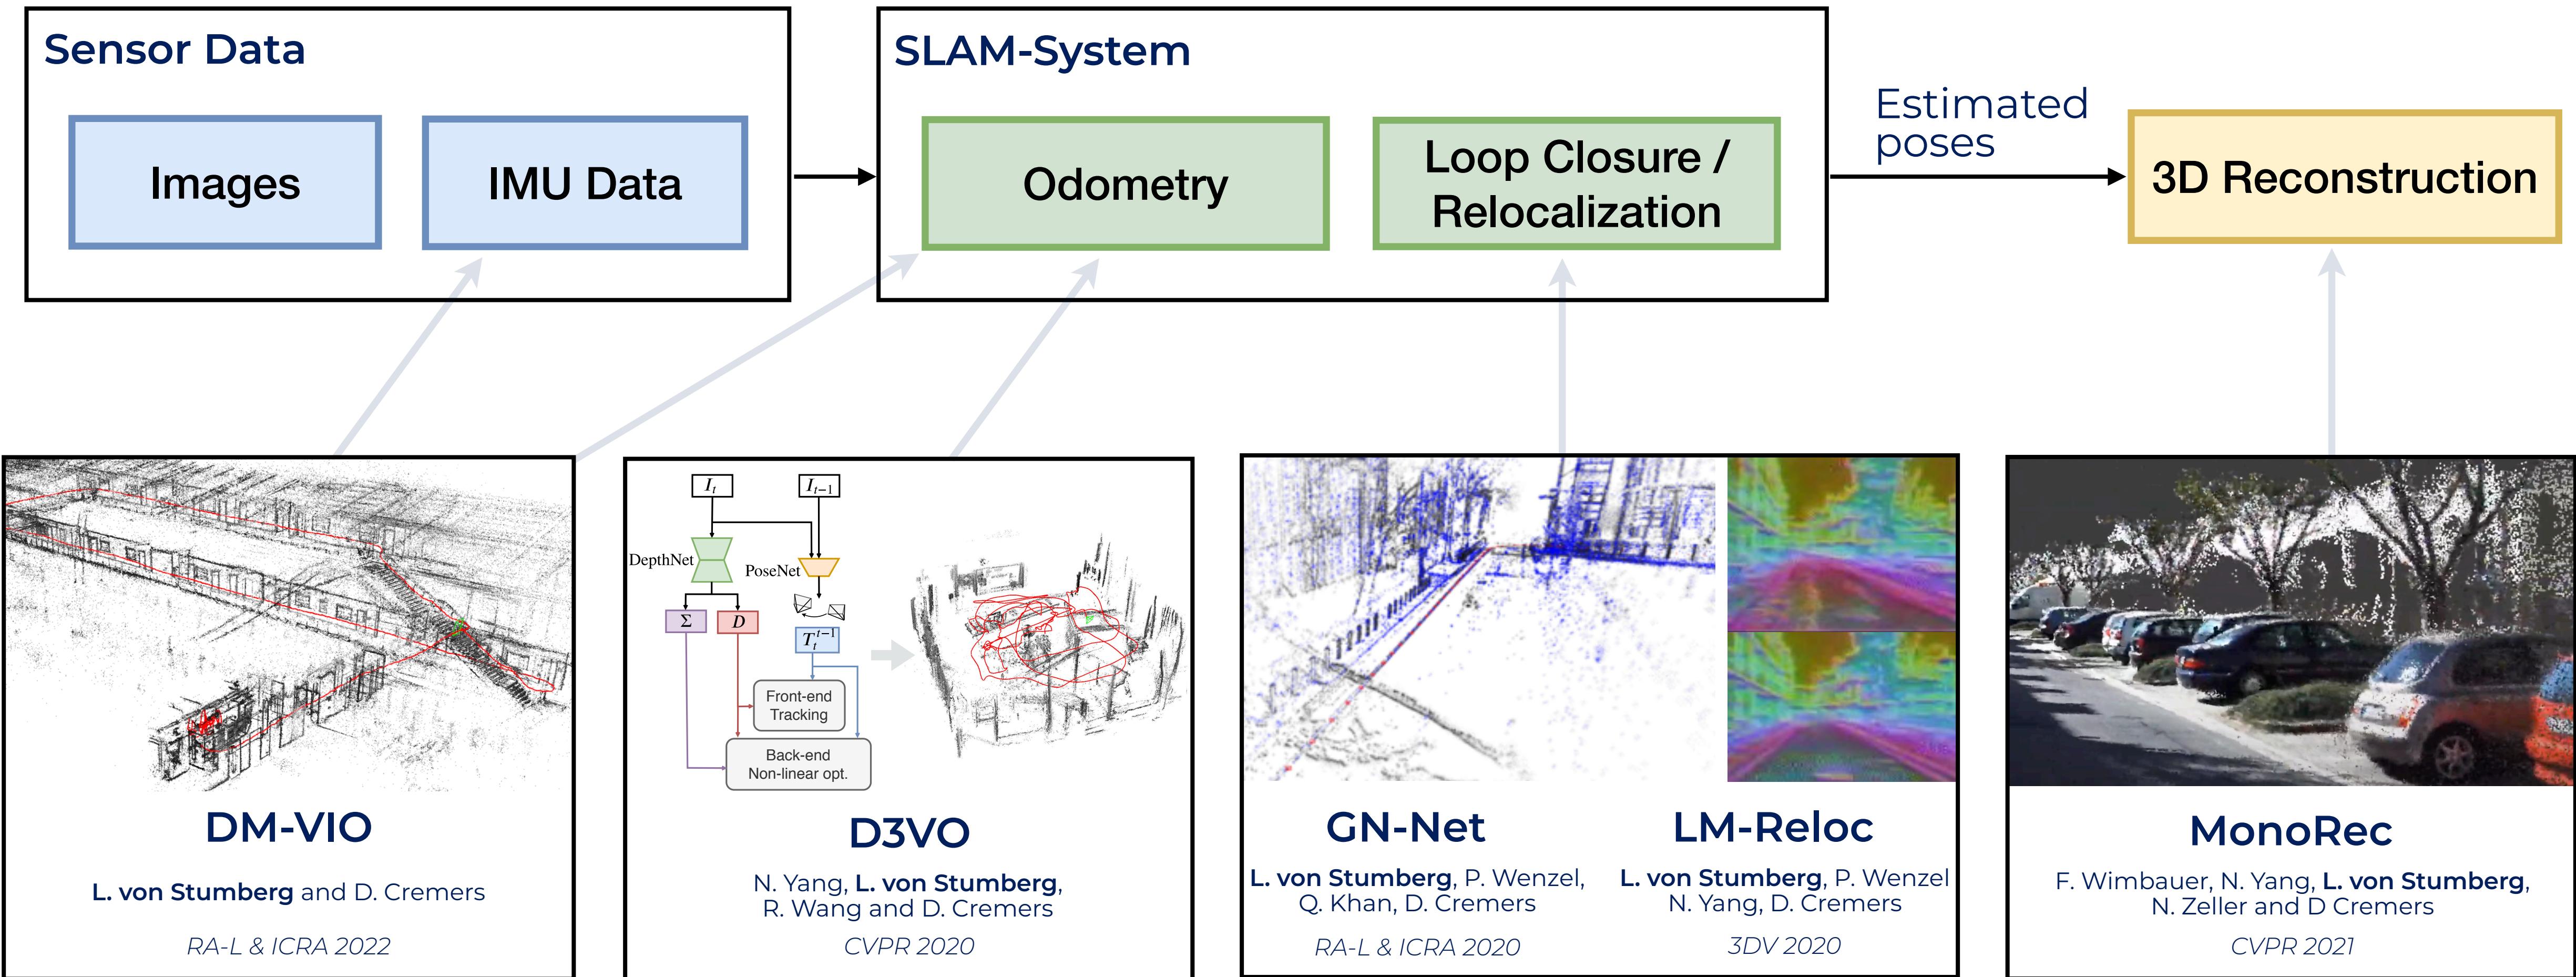
Relocalization Tracking Benchmark





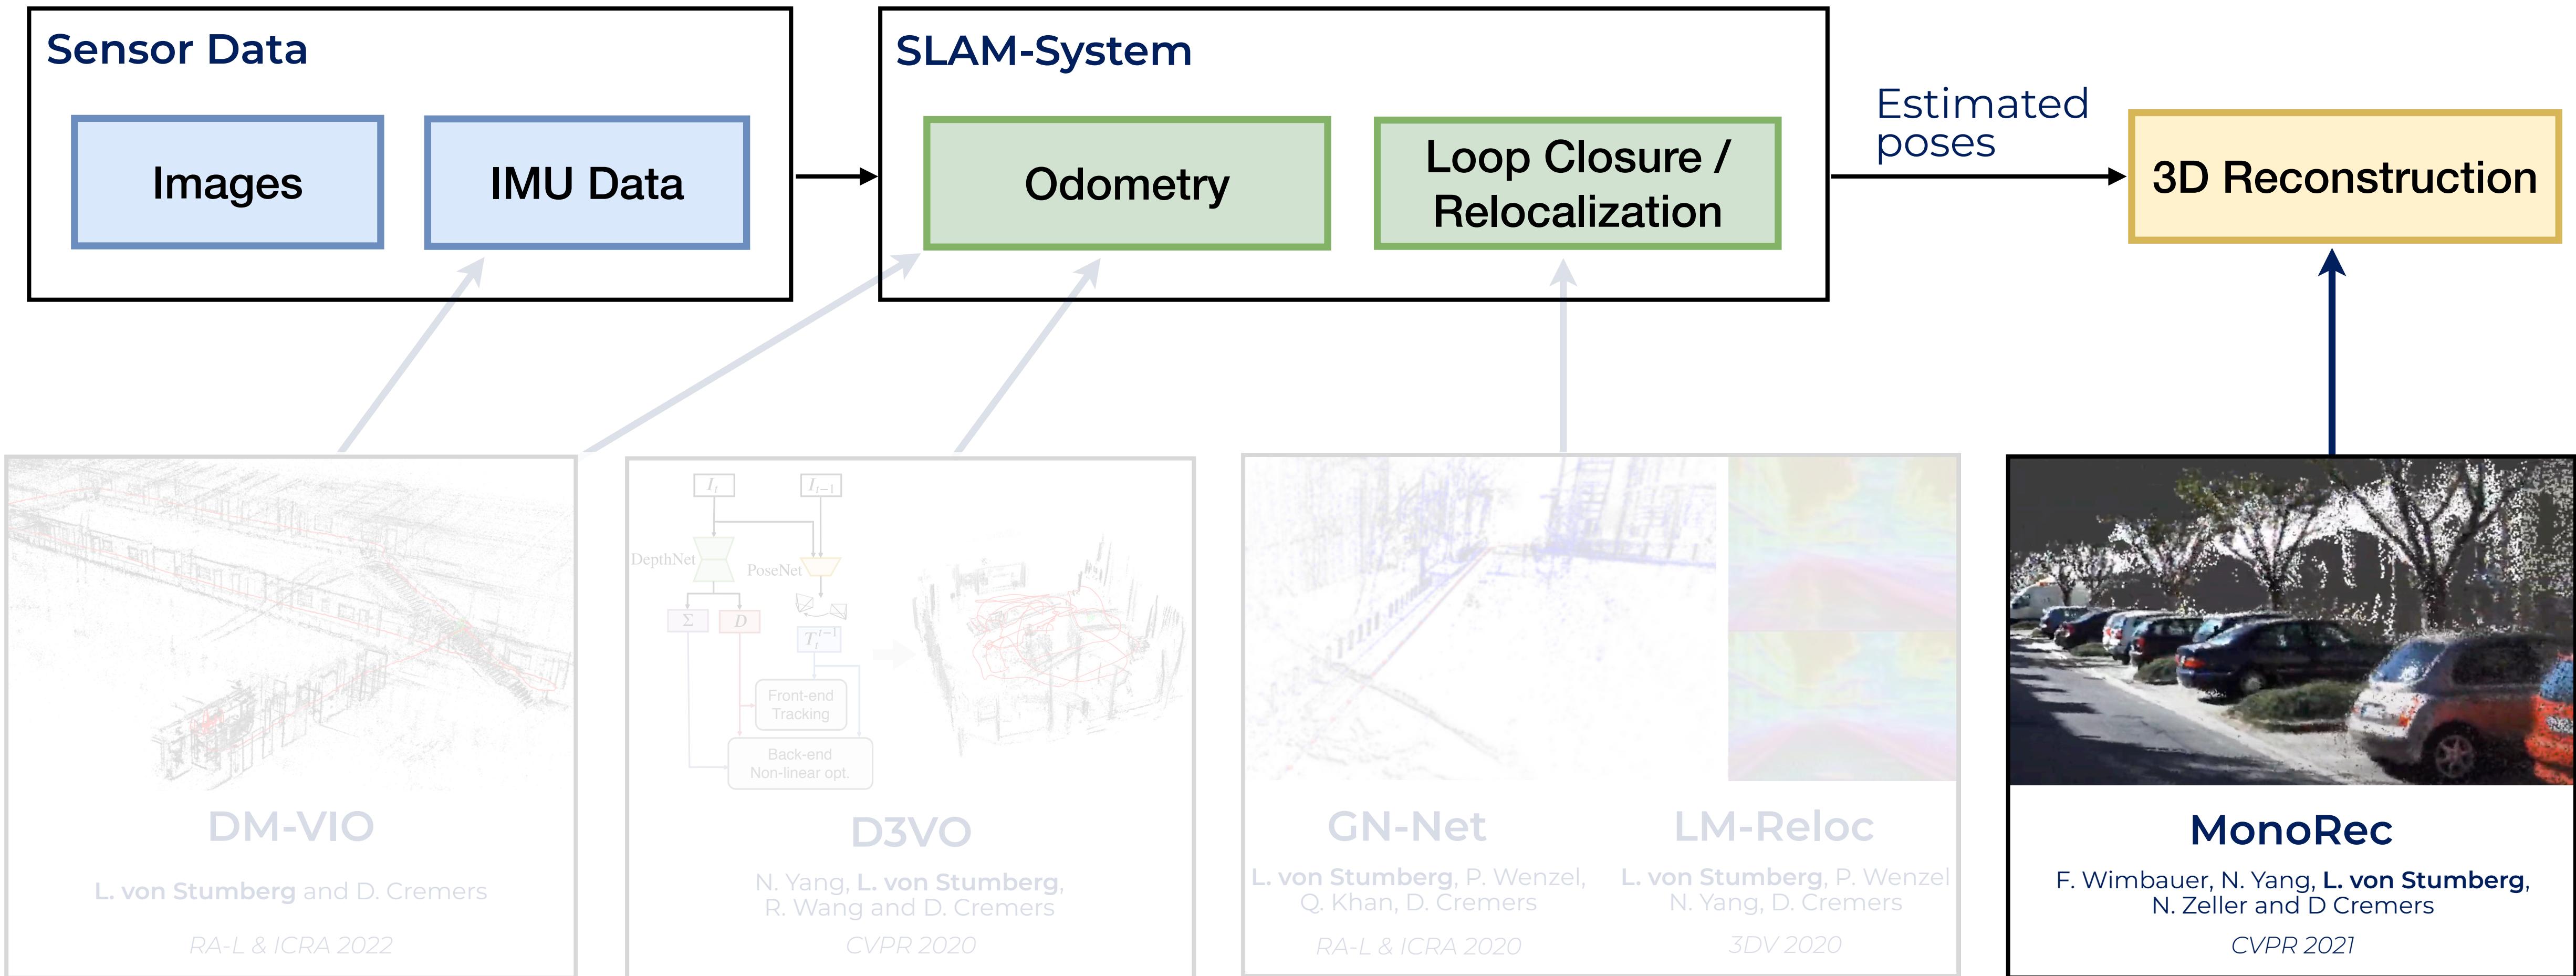
Visual SLAM

From Optimization to Learning



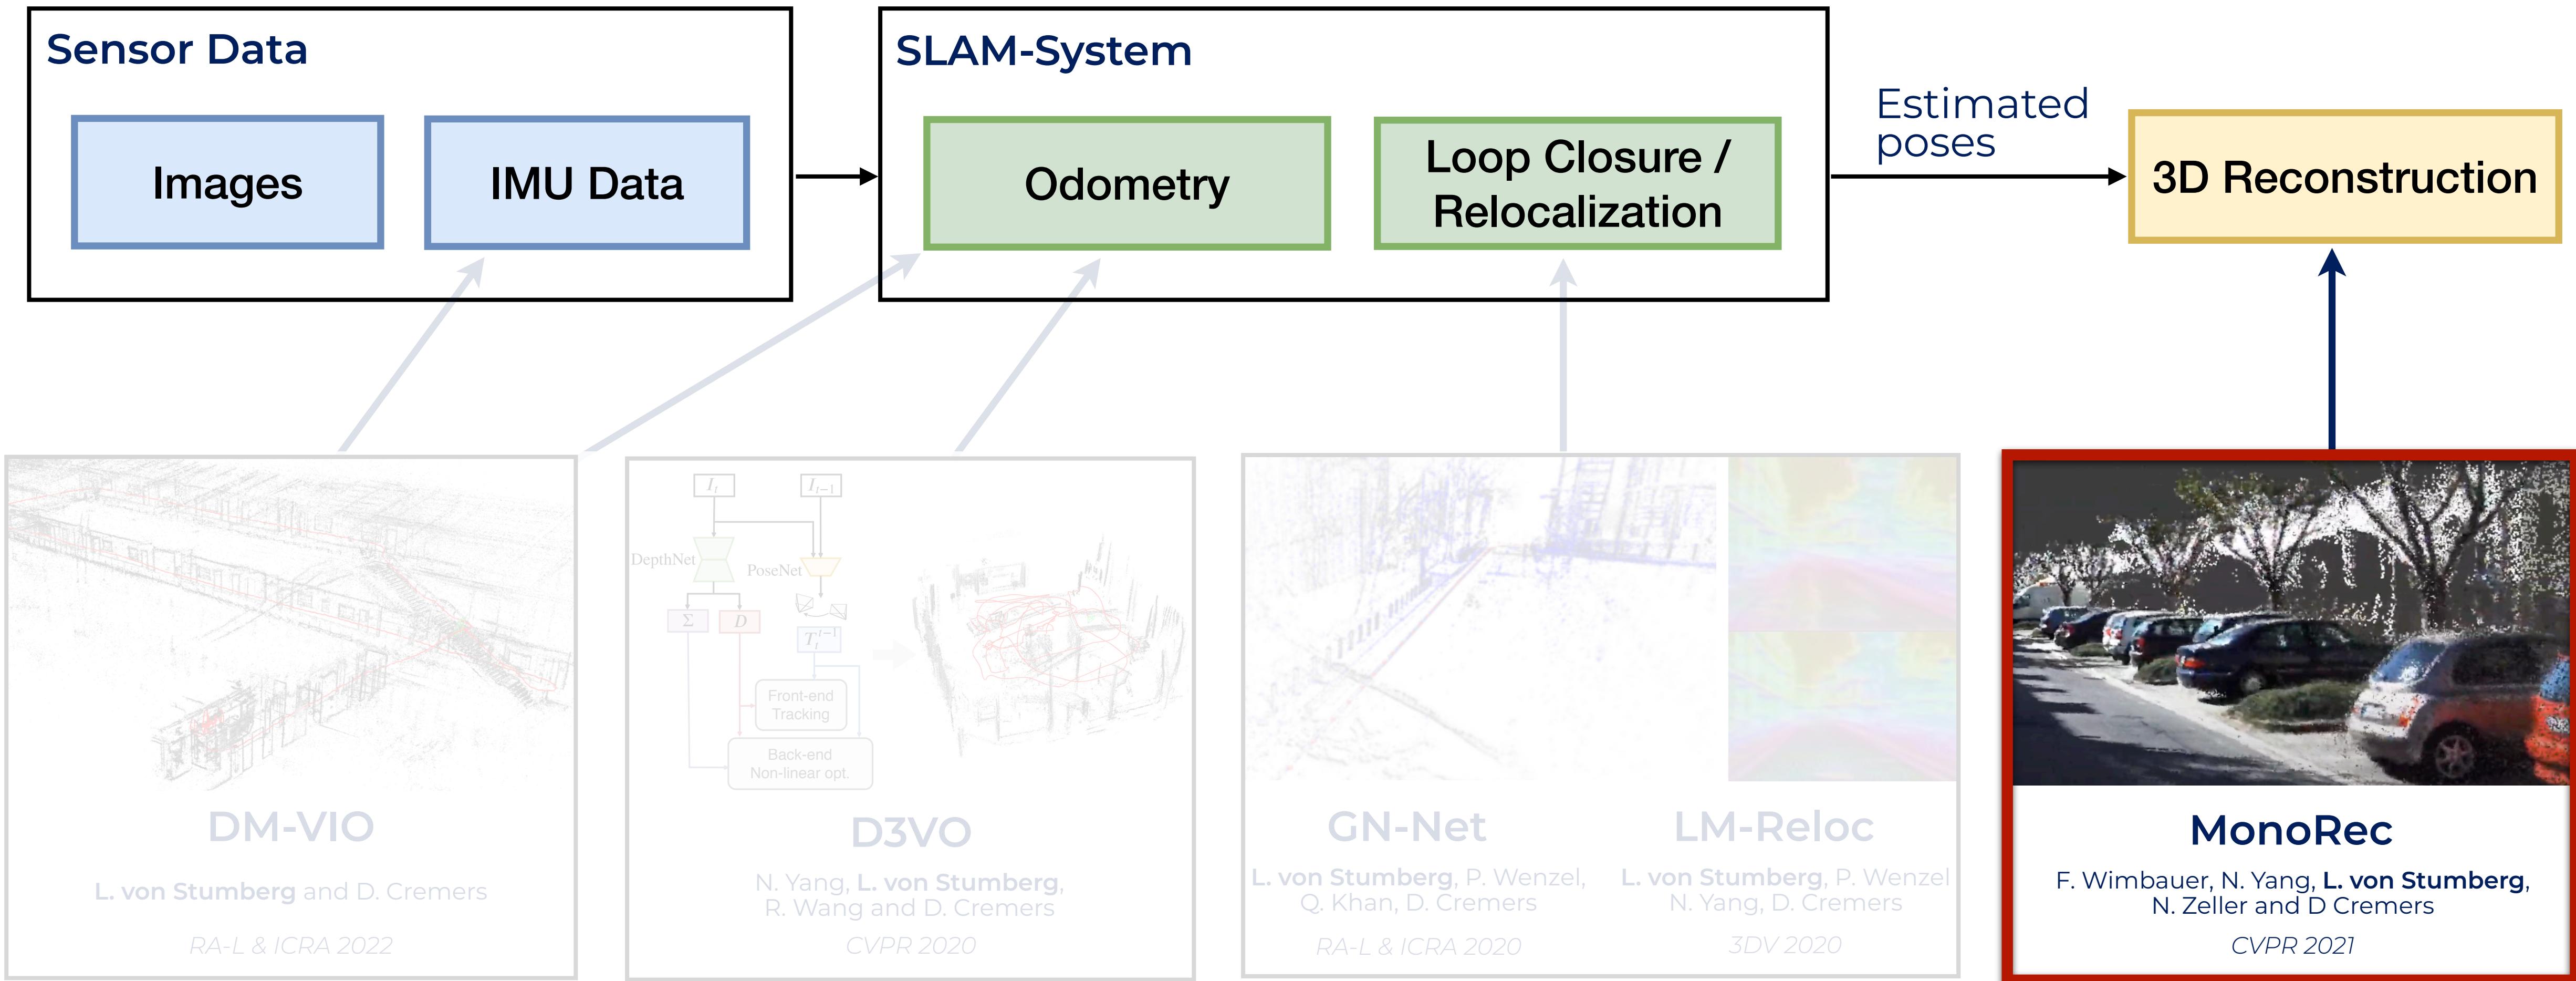
Visual SLAM

From Optimization to Learning



Visual SLAM

From Optimization to Learning





MonoRec: Semi-Supervised Dense Reconstruction... CVPR 2021 F. Wimbauer, N. Yang, L. von Stumberg, N. Zeller and D Cremers



MonoRec: Semi-Supervised Dense Reconstruction... CVPR 2021 F. Wimbauer, N. Yang, L. von Stumberg, N. Zeller and D Cremers

Visual SLAM

From Optimization to Learning

